

THE OMNI-COMPASS®

DATA CENTERS & PHYSICAL INFRASTRUCTURE

-  NVIDIA DGX Systems
-  NVIDIA DGX Cloud / Base Command
-  NVIDIA DCGM / GPU Operator
-  NVIDIA Dynamic Power Software (DPS)
-  NVIDIA Run:ai
-  Phaidra (Autonomous Cooling)
-  Schneider Electric EcoStruxure IT
-  Vertiv

CLOUD & MULTI-CLOUD PLATFORMS

-  Amazon Web Services
-  Microsoft Azure
-  Google Cloud
-  Oracle Cloud Infrastructure
-  Alibaba Cloud
-  IBM Cloud
-  DigitalOcean

CONTAINERS & ORCHESTRATION

-  Kubernetes
-  Red Hat OpenShift
-  Amazon EKS
-  Microsoft AKS
-  Google GKE
-  VMware Tanzu
-  SUSE Rancher / RKE2

WORKLOAD SCHEDULING & HPC

-  Slurm
-  KubeRay
-  Ray
-  Kueue

RESOURCE OPTIMIZATION & AUTONOMOUS OPERATIONS

-  IBM Turbonomic
-  Apptio
-  Moogsoft AIOps
-  BigPanda

AI APPS & PLATFORMS

-  ChatGPT (OpenAI)
-  Claude (Anthropic)
-  Grok (xAI)

A UNIFIED GOVERNING CONVERGENCE CONTROL CORE PLANE ENGINE

A MATHEMATICALLY CLOSED CIRCLE FRAMEWORK



A UNIFIED GOVERNANCE.
BOUNDED EVOLUTION.
SUSTAINED COHERENCE.

OWNER AND DEVELOPER
THE OMNI-COMPASS LLC.

THE BRAIN AND NERVOUS SYSTEM
FOR FRAGMENTED STACKED
MUSCLE / MODULES.

STABILITY. CONVERGENCE. CONTINUITY.

www.omni-compass.com

NETWORKING, SECURITY & SERVICE MESH

-  Istio
-  Cilium
-  Cloudflare
-  Akamai
-  Palo Alto Networks
-  Fortinet
-  CrowdStrike
-  F5

OBSERVABILITY & AIOPS

-  Datadog
-  Dynatrace
-  New Relic
-  AppDynamics
-  Splunk
-  Elastic
-  Prometheus
-  Grafana
-  OpenTelemetry

GITOPS & AUTOMATION

-  Terraform
-  Ansible
-  Argo CD
-  Flux
-  GitHub Actions
-  GitLab CI
-  Jenkins

DATA, STREAMING & ANALYTICS

-  Snowflake
-  Databricks
-  MongoDB
-  PostgreSQL
-  Oracle
-  Apache Kafka
-  Redis

IDENTITY, POLICY & GOVERNANCE

-  Okta
-  HashiCorp Vault
-  Consul
-  OPA / Gatekeeper

IT OPERATIONS & SERVICES

-  ServiceNow
-  PagerDuty
-  Splunk On-Call
-  Jira

U.S. PATENT, COPYRIGHT & TRADEMARK APPLICATIONS FILED

ALL RIGHTS RESERVED

THE OMNI-COMPASS MANUAL

THE OMNI-COMPASS MANUAL

*The Governor of Everything:
Its Philosophy, Its Mathematics, Its Physics,
and How to Wire It onto Your Stack*



AJ Dubra
Owner and Developer

THE OMNI-COMPASS LLC
October 2026

The Omni-Compass Manual

The Governor of Everything: Its Philosophy, Its Mathematics, Its Physics, and How to Wire It onto Your Stack

October 2026

Copyright © 2026 The Omni-Compass LLC. All rights reserved. No part of this book may be reproduced, stored or transmitted in any form or by any means, except as allowed by the Omni-Compass Evaluation License, without the written permission of The Omni-Compass LLC.

Patent applications, copyright registrations and trademark applications covering the Omni-Compass engine, its mathematics and its software have been filed in the United States by The Omni-Compass LLC.

“Omni-Compass”, the Omni-Compass rose and the related names and marks are trademarks of The Omni-Compass LLC. Other names that appear in this book are the property of their owners and are used only to identify the systems Omni-Compass works with.

License. This book and the software it describes are proprietary and are not open source (SPDX-License-Identifier: LicenseRef-OmniCompass-Evaluation-1.0). They may be used only to evaluate Omni-Compass and to reproduce its published results, in simulation and on systems the reader owns or controls. Any commercial use, commercialization, monetization, production use, redistribution, hosted or managed service, or incorporation into any product or service requires a written Omni-Compass Enterprise License, signed by The Omni-Compass LLC and paid for. The full terms are printed in Appendix I.

Disclaimer. The software is provided “as is”, without warranty of any kind. Every result in this book carries its evidence class. A simulated result is a statement about a model; a software benchmark is a statement about the software it ran on; only a hardware meter speaks for hardware. Nothing in this book is a promise of a particular saving on a particular system: the only number that applies to a system is the one its own paired runs produce, on its own receipt.

Front cover: the Omni-Compass identity plate. Back cover: the Omni-Compass unified control authority. Plates 1 to 19: the founder's plates of the theory.

Set in Liberation Serif and DejaVu Sans Mono. Built from the repository by `docs/book/build_book.py`; the book and the code change in the same commit.

The Omni-Compass LLC · www.omni-compass.com

Contents

Foreword	xv
Preface	xvii
Executive Summary	xix
Disclosures, Declarations and Disclaimers	xxi
Part One. The Philosophy and the Theory	1
1. What Omni-Compass Is	2
2. The Unified Circle Principle	6
3. The Four-Piece Engine	20
4. The Compass	27
5. Basins, Polarity and the Dual-Basin Engine	29
Part Two. The Mathematics	33
6. The Canonical Declaration	34
7. The Engine: Eight Equations and One Control Law	41
8. The Canonical Engine	43
9. The Closed Circle: Why It Cannot Leave Its Compass	46
10. Closing the Circle in the Engine	50
11. The Tracking Theorem	51
12. The Conveyance Law	54
13. Formal Status of the Mathematics	58
14. The Mechanism of Action	59
15. The Engines and Their Audit	67
Part Three. The Physics: the Compass and the Nervous System	71
16. The Compass: Push, Pull and the Two Forces	72
17. The Physics of a Processor	75

18. The Two-Way Nervous System	76
19. The Nervous System in Detail	78
Part Four. The Body: Muscles, Realms and Organisms	81
20. The Muscles, the Realms and the Six Organisms	82
21. The Four Realms	84
22. The Realm Muscles	86
23. The Domain Map	97
24. The Problem Map	103
25. The Six Organisms and the Benchmark Grid	107
26. The Metrics Catalog	108
Part Five. The Harness and the Wiring	113
27. The Plug, the Adapters and the Wire Check	114
28. The Harness	117
29. The Wiring Guide	119
30. Before You Start, and the Eight Levels	122
31. Stack by Stack	125
32. The Integration Manual	132
33. Running the GPU Benchmark	141
Part Six. Operating It	147
34. The OFF Switch, the Rules, and the Log	148
35. Maintenance, Upgrades and Security	151
36. The Operator Manual	153
Part Seven. Proving It	161
37. How to Read the Results	163
38. The Dossier: Every Result in One Place	168
39. Paired Runs and Receipts on Your Own System	178
40. Evidence Classes and How to Read a Result	180
41. Results to Date	182
42. The Pilot Protocol and Kit	194
43. The GPU Bench	196
44. The GPU Preregistration	199

45. The Realms Preregistration	210
46. The Compass Law on Real Kubernetes: Preregistration	221
47. The Evidence Ledger	241
48. The Claims Register	247
49. The Benchmark Report	254
50. The Referee Report	275
51. Comparison with Existing Controllers	288
52. The State of Play	292
Part Eight. Value, License and History	299
53. Where the Value Comes From	300
54. The Economics of a Receipt	305
55. The Buyer Edition	306
56. Due Diligence	314
57. License and Commercial Terms	316
58. Licensing: Questions and Answers	317
59. Third-Party Notices	319
60. Repository Standards	320
61. Python, C++ and the Seal	322
62. The Founder's Working Notes	324
63. History	329
Back Matter	337
Glossary	338
Appendix A. Command Reference	341
Appendix B. File Map	343
Appendix C. The Equations in Full	345
Appendix D. Metrics	347
Appendix E. Troubleshooting	349
Appendix F. Evidence Map	351
Appendix G. The Muscles: What Each Is For, and How It Is Wired	354
Appendix H. Source of the Engine	413
Appendix I. The License	440
Contact	442

List of Plates

Plate 1. The Unified Circle Principle

Plate 2. Closed-circle structural architecture of the unified law

Plate 3. The closed circle tested

Plate 4. Symmetry principles and conservation laws in the closed circle

Plate 5. Scale invariance and structural universality

Plate 6. The lineage of unification

Plate 7. A candidate for a closed, unified law

Plate 8. The Omni-Compass Grand Theory: the four pieces

Plate 9. The unified container of all reality: 24 domains, four-piece engine, one closure law

Plate 10. The four quadrants and the compression process

Plate 11. The closure law and the closed-circle engine

Plate 12. The Omni-Compass rose

Plate 13. The Omni-Compass Principle

Plate 14. Geometric funneling and basin formation

Plate 15. The Omni-Compass dual-basin engine

Plate 16. Mathematical integration

Plate 17. The engine core: the governing law in motion

Plate 18. The equation form and the symbol chart

Plate 19. Empirical validation, experimental pathways, and the discipline of proof

Foreword

I built Omni-Compass to sit on top of what already exists. Kubernetes, the GPU driver, the building controller, the battery inverter, the robot's servo loop: those are muscles, and they are good muscles. What none of them has is one brain that reads all of them at once, holds each of them in the middle of its safe range, and puts every setting back exactly where it found it when it stops. That brain is Omni-Compass.

This manual is the whole of it: what the governor is, the mathematics it runs on, how its nervous system reads and writes, how it is wired to each kind of machine, how to turn it on one level at a time, how to switch it off, and how to prove on your own system what it does. Keep it with the code. When the code changes, this manual changes in the same commit.

On this edition. This edition is written for the referee and the underwriter as much as for the engineer. Every result in it was run three separate times on a frozen, fingerprinted engine, under rules written before the first run, and is reported with every loss beside every gain. Where a number is from a model it says so; where a test could not see an effect it says so; where something has not been shown it says so, in its own section. The questions a careful reader will ask are asked here first, in section 3.5 beside each result and in section 16.7 as a list, and answered where an answer exists. Nothing in this manual asks to be believed on our word: every table names the runs it was made from, every run's raw files are in the repository, and every table can be rebuilt from them by the tool it names.

AJ Dubra Founder, The Omni-Compass LLC

Preface

I wrote this book so that whoever connects Omni-Compass to a running system understands what they are connecting: not only which wire goes where, but why the governor behaves the way it does, where its law comes from, and how to prove on their own system what it does for them.

The book is in eight parts, and it can be read in two ways.

You are	Read
CEO, board member, investor	the Foreword, the Executive Summary, Part One, the Results to Date, and Part Eight
CTO, architect, head of platform	Parts One to Four, then the eight wiring levels in Part Five, then Part Seven
The engineer wiring it	everything, in order. Do not skip the wire check or the watch level
Auditor, diligence team	Part Two (the mathematics), Part Seven (the proof), Appendices C, F and H

Part One, the philosophy and the theory, sets out the closed circle: a system that is closed, bounded and pulled toward a center cannot run away, and everything it does is a return. **Part Two, the mathematics**, gives the eight equations and the control law that make the circle exact, with the theorem, the declaration and the audits that hold them fixed. **Part Three, the physics**, brings the law into the machine: the compass, push and pull, the band and its cushions, the physics of a processor, and the two-way nervous system. **Part Four, the body**, lays out the muscles, the four realms and the six organisms. **Part Five, the harness and the wiring**, is the universal plug and the step-by-step work of wiring Omni-Compass onto a stack. **Part Six, operating it**, is the OFF switch, the rules and the log. **Part Seven, proving it**, is every rule written before a run and every result after it. **Part Eight** is value, the license, the seal and the history.

Conventions. Code is a command, file or switch exactly as typed. "Native" means a system as it runs today, without Omni-Compass. "Muscle" means any machine, service or controller Omni-Compass can read and set. "Knob" or "lever" is one setting on a muscle. "The band" is a knob's or a service reading's safe range. A step that writes to a system is marked **WRITES**. Every result carries its evidence class: T for a theorem, V for verification of the code, S for simulation, L for live software, P for a physical meter.

Some chapters gather documents that were written as the work went on, each at the moment its result came in. They are kept as they were written, because the record of how a result was reached is part of the proof. Where two chapters give different figures for the same thing, the later run and the State of Play govern.

Executive Summary

What it is. Omni-Compass is a supervisory governor. It sits on top of the controller a system already has, the Kubernetes autoscaler, the GPU's firmware, the cooling loop, the battery inverter, the robot's servo, the grid's tap changer, the database's connection pooler, the message broker's consumer group, the cache's memory ceiling, and it moves only the settings those controllers already accept. It reads each muscle's own meter, computes one bounded force per muscle from a single smooth law, and holds the service in the middle of its safe band rather than far below its limit. The native controller keeps running exactly as shipped; Omni-Compass sets what it already exposes, and when Omni-Compass stops, every setting goes back to the value it read once before its first write. It is never standalone and it never competes with the controller underneath it.

What it does for you. Every running system keeps room "just in case": machines sized for the worst minute of the day, a pool of connections held open for a load that comes once an hour, a consumer group sized for the quiet hours and then overrun at the peak, a cache ceiling set once and left. Omni-Compass turns that room into one of two things, and the operator chooses which: more work for the same resources, or the same work for fewer. The two are the same measurement read from two sides (section 3.4): a third more work from what is already installed is a quarter off the bill.

What has been shown, on the frozen engine Omni v3, three separate times each. On real Kubernetes, six preregistered tests of ten paired runs each (`results/live/V3_*.md`): work inside the response line **+35% to +49%** in every run of the all-four test, the 95th-percentile response time **47% to 66% faster** across the steady, wandering, all-four and fault tests, failed requests 9% to 14% fewer where load swings, machines 1.5% to 2.9% fewer at steady load and 19% to 23% fewer on a queue of batch jobs, and beside a noisy neighbour no difference beyond the noise on any row. On a real database behind its pooler (`results/live/V3_PGBENCH.md`): 61% to 72% fewer connections held open for the same work and the same latency on two of three workloads, at a confirmed cost of 14% to 28% more CPU on the host, counted against Omni-Compass. On a real message broker, Apache Kafka as shipped (`results/live/V3_KAFKA.md`): 16% to 21% more messages inside the 500 ms line on all three untouched workloads, a 95th percentile of 9 to 14 ms against native's 1.6 s, no message lost, at the confirmed cost of three to four times the consumers running. On a real cache, Redis as shipped (`results/live/V3_REDIS.md`): 14% to 27% more requests inside the 2 ms line and a hit rate 14% to 27% higher on all three untouched workloads, at the confirmed cost of a memory ceiling held at 200 to 270 MB against the operator's 64 MB, which outweighs the gain in the index. On a real database's storage-engine cache, MongoDB as its publisher ships it under YCSB (`results/live/V3_YCSB.md`): the cache given back by about half on two of four untouched workloads and by 13% to 37% on a third, confirmed, with work inside the line, p95 and CPU inside the noise and no failed operation, at the confirmed cost of 2% to 4% on the mean latency of the burst workload. In the modelled realms, 945 muscles and six organisms (`results/realms/REALMS.md`): every organism superior within its guardrails, zero muscles worse, at every size from 1 to 1,000 copies. On the independent simulators: a power grid's losses better in 7 of 11 grids and worse in 4; robot arms' peak torque down 10% to 29% where Omni-Compass moved and left native where the paired trial said not to; buildings' electricity bought and daily peak better in all 11 battery districts and the bill worse in 7; drone swarms' energy a mission 7% to 20% lower with no late mission, near miss or collision. The one combined number, the Omni index over the five real categories confirmed three times, stands at **+25.9%**

(results/OMNI_INDEX.md): Kubernetes +25.5%, the database +14.2%, messaging +166.9%, the cache -24.9% and the database's storage-engine cache +10.2%, every category weighed the same and every row inside the noise counted as exactly nothing. The cache's category is negative because the memory it holds is the resource it trades and reads worse by rule, even as its work and hit rate read better; the index does not hide that.

What has not been shown, said plainly. An energy or cloud-bill saving on real machines. Energy on kind is a declared model, because the machines are containers on one runner. Azure's bill on a 4-worker fleet read no difference beyond the noise on every gauge, which is the honest reading of a fleet too small to show one machine; the 40-worker fleet that can show one is preregistered and has been refused by Azure's own cluster capacity in its region, each refusal recorded. Every earlier GPU result ran on a controller since replaced and is obsolete; the card runs again on rented hardware at one named commit. Each of these is an open item in Chapter 41, not a footnote.

How it stays safe. It watches before it writes. It records every original setting before it acts and reads back every write. It never writes past a knob's cover. It gives everything back the moment service crosses the wall or a sense goes blind. It stops writing a knob the instant anyone else touches it. It moves a slow knob only after a paired trial on the muscle itself shows no more than 2% cost. One OFF switch returns every setting to its original value; a watchdog with a lease does the same if the governor is killed outright; and a separate kill switch, for security only, turns the whole harness off in one human hand.

How you know. Every run produces receipts: native against omni on the same system, the same load and the same clock, each gauge with its 95% interval, the raw files archived beside the code and fingerprinted. Results are labelled by a rule written before the run, confirmed only when three separate runs agree, and reported with every loss in the table.

Disclosures, Declarations and Disclaimers

This page is the one place every declaration about Omni-Compass is made. The README, the manual, the book, every guide and every report point here. Where any other page seems to say more than this page, this page governs.

Rights

1. Omni-Compass, its engine, its mathematics, its software and its documentation are the property of The Omni-Compass LLC. Patent applications, copyright registrations and trademark applications covering them have been filed in the United States by The Omni-Compass LLC. No filing number is stated here, and no grant, registration or approval is claimed.
2. The software is licensed for evaluation and simulation only (LICENSE). Running it in production, on any system beyond evaluation, or in any product or service, requires a signed, paid Omni-Compass Enterprise License.
3. Names of other companies and products (NVIDIA, Kubernetes, Red Hat OpenShift, Amazon EKS, Google GKE, Microsoft AKS, Karpenter, Lambda and others) are the property of their owners. They identify the systems Omni-Compass was tested with or connects to. No affiliation, endorsement or certification by any of them is stated or implied.

What a result is, and what it is not

1. Every result carries its evidence class: **T** a theorem, **V** verified in code, **S** a model (simulation), **L** live software (real Kubernetes on GitHub's machines), **P** a physical meter (a real card's own power meter). A model is a statement about the model. A software benchmark is a statement about the software and machine it ran on. Only a meter speaks for hardware.
2. The models were written by the same people who wrote the law. They show how the mechanism behaves; they are not an independent test.
3. On Kubernetes in kind every machine stays powered, so energy there is a declared model, not a meter.
4. Every comparison is paired: the same seed, load and clock with and without Omni-Compass, in an order rotated by repetition. Intervals are 95% intervals over the paired repetitions. Ratios across seeds are summarised as geometric means. Labels are assigned by code from rules written and committed before the run (docs/GPU_PREREGISTRATION.md, docs/REALMS_PREREGISTRATION.md, docs/K8S_COMPASS_PREREGISTRATION.md), never by hand.
5. Losses and failed runs are kept beside the wins (docs/EVIDENCE_LEDGER.md, results/live/, git history). The first real-card run, for example, saved energy (+3.6% work per energy, proven) and made the slowest answers 58.5% slower; it is published as it happened.
6. No result is a promise of a saving on any other system. No dollar figure in this repository is a measured result. Nothing here is a vendor certification, a safety certification, legal advice or investment advice.

The rule Omni-Compass keeps on every muscle

Omni-Compass sits on top of a muscle that already works and moves its settings only to save energy, and only where it measures that the muscle loses at most **2%** in any measure: work done, response time (median, p95 or p99), time over the service line, or the machine's own cost of running Omni-Compass. Where no move passes that test, Omni-Compass leaves the muscle exactly as it runs alone. On a live system the verdict (`omniconpass/verdict.py`) measures it with paired trials on the muscle itself before every step. Every result reports each measure, and where one is worse at all, by however little, the result says so and by how much.

The wiring declaration

Omni-Compass acts only through the wires it is given. It reads the meters it is pointed at and moves the settings it is allowed to move. If a reading is the wrong one, slow, or blind; if a lever is the wrong one, already owned by another controller, or given a range (cover) that is wrong for the machine; if a band or a service line is set for a different workload; then Omni-Compass will do exactly what its law says with the wrong information, and the result will not be the published one.

Therefore:

1. **It cannot be slapped on.** Every installation goes through the levels of the manual in order (chapter 9), starting read-only, and every writing level begins with the wire check (manual, section 8.4) and a watch arm that must equal native.
2. **If your paired receipts differ from the published benchmarks in direction** (service worse, or no saving where one was shown), **the first presumption is wiring, not the law.** Confirm the installation with the checks in the manual's section 8.5 (*Wired right or wired wrong*) before drawing any conclusion about Omni-Compass.
3. **This has happened to us, and is published.** On the first real card (NVIDIA A10, 2026-10-02) the governor read being busy as trouble, set its power lid under the card's own working draw and allowed the clock under the card's own working clock; the card served bursts more slowly and the slowest answers were 58.5% slower. The cause was found in the card's own samples and corrected (`docs/GPU_PREREGISTRATION.md`, amendments 6 and 7). That is what a wiring fault looks like, and how it is found.
4. Omni-Compass never writes outside a lever's cover, never fights another writer (it stops and leaves that lever alone), reads back every write, and returns every lever to the value it read before its first write when it stops or when the OFF switch is used. These protect the machine; they do not make a wrong wiring right.

Safety and responsibility

1. Run watch mode first; run the wire check before any write; keep the master OFF switch (`python3 tools/omni_switch.py off`, the whole harness at once) in reach at every level, and use it at the first sign of anything wrong, a suspected breach included.
2. Omni-Compass is a supervisory governor. The machines' own controls (firmware, autoscalers, safety systems) stay in place and keep their own protections; Omni-Compass sets only values they already accept.
3. Do not connect Omni-Compass to safety-critical systems (vehicles, medical devices, grid protection, life safety) except as a modelled study, and never without the qualified review those systems require.
4. The software is provided "as is", without warranty of any kind, as stated in LICENSE. The operator is responsible for its use on their systems.

Changes to the work and to this page

1. The software, the law's settings, the benchmarks, the results, the documents, the licensing terms and this page may change at any time, without notice. The version in the repository's main branch at a given commit is the version that commit describes; `RELEASE_MANIFEST.json` and `python3 verify.py` identify it.
2. Statements about work in progress, planned features, expected results and future runs are expectations, not promises. Where a result has not yet been measured, the documents say so (`docs/STATE_OF_PLAY.md`, `docs/DOSSIER.md`, Chapter 27).
3. A signed Omni-Compass Enterprise License governs its own terms for its own term.

Privacy and data

Omni-Compass contains no telemetry. It sends nothing to The Omni-Compass LLC or anyone else, and makes no network connection of its own except to the systems an operator points it at. Its logs and receipts stay on the operator's machines.

Where everything is

Question	Page
What has been measured, and how strongly	<code>docs/STATE_OF_PLAY.md</code> , <code>docs/DOSSIER.md</code> , <code>docs/EVIDENCE_LEDGER.md</code>
The rules written before each run	<code>docs/GPU_PREREGISTRATION.md</code> , <code>docs/REALMS_PREREGISTRATION.md</code> , <code>docs/K8S_COMPASS_PREREGISTRATION.md</code>
How to wire it, level by level, and how to confirm it is wired right	<code>docs/OMNI_COMPASS_MANUAL.md</code> (chapters 8 and 9), the book <code>docs/OMNI_COMPASS_MANUAL.pdf</code>
The license, in plain answers, and third-party components	<code>LICENSE</code> , <code>LICENSING_FAQ.md</code> , <code>NOTICE</code> , <code>THIRD_PARTY_NOTICES.md</code> , <code>PATENTS.md</code> , <code>TRADEMARKS.md</code>
The muscles, what each is for and how it is wired	<code>docs/MUSCLE_CATALOG.md</code>
Proof the code is the code that ran	<code>python3 verify.py</code> , <code>RELEASE_MANIFEST.json</code> , <code>results/SEAL.json</code>

Evaluation and simulation use only. Copyright (c) 2026 The Omni-Compass LLC. Commercial use, commercialization or monetization of any part of Omni-Compass requires a signed, paid Omni-Compass Enterprise License from The Omni-Compass LLC. Patent applications, copyright registrations and trademark applications filed in the United States. See LICENSE and NOTICE at the root of this repository.

PART ONE

The Philosophy and the Theory

Where Omni-Compass comes from: the closed circle, the four pieces, the compass and the basins. The engineering in the rest of the book is this idea made exact.

CHAPTER 1

What Omni-Compass Is

A governor on a steam engine does not build the engine and does not turn the shaft. It watches the speed and moves the throttle so the speed stays in a band. Omni-Compass is that, for every machine you run.

It is a process on a host. It reads meters, steps a bounded mathematical law, and writes only the levers it has been given. It is not the chip, not the GPU driver, not Kubernetes, not the building controller. Those keep running exactly as they do today; Omni-Compass sets the values they already accept.

Always on top of a native controller

The first design decision, and the one every other follows from, is that Omni-Compass is **supervisory**. It never replaces the controller a system already has, never bypasses it, and never runs in its place. The Kubernetes Horizontal Pod Autoscaler keeps scaling pods on its own CPU target; Omni-Compass moves that target inside a range the operator allows. The cluster autoscaler keeps adding and deleting machines; Omni-Compass marks machines idle and lets the autoscaler do the deleting. The card's firmware keeps boosting the clock and enforcing its power limit; Omni-Compass sets the ceiling the boost may reach and the limit the firmware enforces. PgBouncer keeps handing connections to clients; Omni-Compass sets its pool size through the pooler's own console. Kafka's consumer group keeps rebalancing partitions on its own; Omni-Compass changes how many consumers are in the group, which is what an operator does. Redis keeps evicting by its own rule; Omni-Compass sets the ceiling that rule works under.

This matters to a referee for a precise reason: every benchmark in this manual has exactly two arms, **native** (the system with its own controller, as shipped and as the operator set it) and **omni** (the same system with Omni-Compass on top). There is no third arm in which Omni-Compass replaces the controller, and no comparison in which Omni-Compass is a rival to it. A gain is therefore always "what the governor added on top of what the native controller already did", and a loss is always "what the governor cost on top of it". Nothing in the record is an argument against any native controller; every native controller in this manual is a good controller that the governor sits on.

Do no harm: the verdict

The second design decision is that Omni-Compass must never make anything worse. This is not a slogan printed on a report; it is a trigger inside the engine. Before Omni-Compass moves a slow knob (a machine given back, a card's clock ceiling, a robot axis's speed, a drone's cruise override), it runs a **paired trial on the muscle itself** (`omnicompass/verdict.py`): the muscle's own cost per piece of work is measured at native and at one step past the deepest step already allowed, under the same traffic, with at least thirty pieces of work on each side; the step is allowed only if the trial's median cost is no more than 2% above the reference and no more than 2% above the cost first measured at native. A refused step is not tried again for a set number of decisions. Where no step passes, the verdict reads **left native**, the knob is never moved, and the result for that muscle says "nothing for Omni to move". This is why, in the robot-arm benchmark, two of four arms were left entirely native and are reported as such, and why in the real Kubernetes tests machines were kept whenever one machine fewer made the requests slower in the trial. Gains are not capped; the allowance caps only what may be spent to get them.

The pedals, the modes and the two switches

Omni-Compass drives a system the way a self-driving car drives itself: it feels, gauges and adjusts, with nobody in the seat. The vocabulary is used the same way everywhere in this repository (`docs/MECHANISM_OF_ACTION.md`, "The pedals"):

- **Off (manual).** Omni-Compass is not driving; native runs alone. This is the native arm of every benchmark.
- **Watching.** Omni-Compass reads every gauge and logs the decision it would take, and writes nothing. Any difference between a watching arm and native is the cost of its being there (its CPU, its reads) plus noise, never a decision; a watching arm that writes anything fails its run. Watching is the first level of every rollout (Chapter 30).
- **Autopilot.** The pedals: **idle** (no foot on the gas or the brake, calm traffic, the engine at its floor of two machines in service, ready for the next burst), **gas** (traffic climbs and capacity is added at once, never held back), **brake** (traffic falls and capacity is eased off a step at a time, never below idle), **reset** (the brake held to the floor: every setting the governor ever wrote is handed back to where native had it, read back, and the record removed; every benchmark arm ends with a reset and checks it).
- **Cruise.** For work that comes as a pile: when pods have waited for a place for two decisions in a row, every machine goes into service and stays there, without second-guessing, until the line has been empty for two decisions. In cruise the governor also stops its five-second floor reads, because every machine is already in service and the reads would only cost CPU on the machines the queue is using.
- **The emergency brake.** When the work is done and demand is at zero (the service at its autoscaler's own floor, wanting no more, at half its target or less), straight to idle in one move, with every safety check still holding.
- **The kill switch.** Separate from all of the above, and for **security only**: one switch in a human hand (`python3 tools/omni_switch.py off, omnicompass/master.py`) turns Omni-Compass's governing off across the whole machine at once, for anything rogue, or for anyone trying to drive a system through Omni-Compass's brain. Every governor hands back and exits; none starts while the switch is off. The reset is the end of a run; the kill switch is the end of trust.

The floor of two machines is an operator's number; the brake never takes a system below it, however hard it is pressed. There is no ceiling; every machine the operator has is usable, and the protection is the band inside each machine.

Nothing is hardwired

Omni-Compass carries no table of per-stack special cases. One law (Chapter 16) is written once and runs on every muscle through one plug (Chapter 27). The gains differ from muscle to muscle only through the muscle's declared physics: its response time, its decision period and its cover. Where a rule was found not to work, it was removed, not switched off and left in the code: the engine files are frozen and fingerprinted, and `verify.py` fails if any of them changes. A reader who finds a behaviour in a result can find the exact lines that produced it, and will find no second path.

The word

The law is a **compass**: a band with a cushion at each wall, pushed to its middle by two antagonist forces, smooth, never hammering. The repository uses that word and no other for it. Earlier dated records that used another word are kept word for word, as records are; the current law is `omnicompass/compass_law.py`, `CompassLaw`.

What it sets, muscle by muscle

Every muscle keeps its own control, and Omni-Compass sets what that control already accepts:

Muscle	Its own control (kept)	What Omni-Compass sets
Kubernetes service	the HPA	the HPA's CPU target; replica floors
Kubernetes node pool	Cluster Autoscaler / Karpenter / MachineSet	how many machines stay in service (park and wake)
NVIDIA GPU	the card's firmware (boost, power and thermal limits)	the clock ceiling (up wire) and the power limit (down wire)
CPU	the kernel's frequency governor	the frequency ceiling
Data hall, building	the chiller and air handler loops	the supply-air setpoint; units in service
Battery site	the inverter	the reserve level; the peak ceiling
Robot joint, vehicle axis	the servo loop	speed and effort limits
Process loop, feeder	the PI controller	the setpoint inside its band
Database behind a pooler	PgBouncer's pool, the DBA's fixed setting	the pool size through the pooler's own console
Database's storage engine	MongoDB's WiredTiger cache, the operator's fixed setting	the cache size through the server's own console
Message broker	the consumer group at the operator's count	how many consumers are in the group
Cache	the operator's memory ceiling and eviction rule	the memory ceiling through the cache's own console
Substation	the tap changer's own controller	the tap position, one whole tap per move
Drone	the shipped autopilot	the cruise override inside the autopilot's limits

When Omni-Compass stops, every one of those values goes back to what it was before Omni-Compass acted: not to the last value it wrote, but to the value it read once, before its first write.




Plate 1. The Unified Circle Principle

CHAPTER 2

The Unified Circle Principle

Omni-Compass begins with one idea: a system that is closed, bounded and pulled toward a center cannot run away, and everything it does is a return. I call this the closed circle. This chapter states the principle as I hold it, in its mathematical form, and then says plainly which parts of it the software uses and which parts belong to the wider theory the software grew out of.

The statement


```
dX/dt = G(X),          X(0) in Omega, a subset of R^n
G(X) . n(X) <= 0      on the boundary of Omega
grad L(X) . G(X) <= 0
=> lim (t -> infinity) X(t) in M*, a subset of F
```

Read line by line:

1. **$dX/dt = G(X)$.** The state X moves according to one law G . There is no outside hand: every change comes from the law and the state it is in.
2. **$X(0)$ in Ω .** The state starts inside the admissible region Ω : the set of states that are physically and operationally allowed.
3. **$G(X) \cdot n(X) \leq 0$ on the boundary.** At the wall of Ω , the flow never points outward (n is the outward normal). Nothing that starts inside can cross out. In control theory this is the Nagumo condition: Ω is forward invariant.
4. **$\text{grad } L(X) \cdot G(X) \leq 0$.** There is a function L , the energy of the compass, that never increases along the motion. This is a Lyapunov condition.
5. **$\Rightarrow X(t)$ settles in M^* .** Together, an invariant container and a non-increasing energy force every path into the set M^* where the energy stops falling: the bottom of the compass.

The power of the statement is that it does not depend on the particular disturbance. Every input up to the size the container is built for has the same answer before it is asked: the state stays inside and returns. That is what I mean when I say that a closed circle holds the answer to every question that can be put to it.

Closed-Circle Structural Architecture of the Unified Law



Dynamic Flow

B Redistribution Field

$\mathcal{M}$ Manifold Boundary

Plate 2. Closed-circle structural architecture of the unified law

Closure, admissibility, boundedness

The theory sets five requirements on any law that claims to close the circle:

Requirement	Meaning	In the engine
Finite deviation	the distance from the admissible state has a ceiling, $\text{abs}(E) \leq E_{\text{max}}$	the deviation state E drains faster than it can fill
Finite curvature	the landscape the state moves on has no infinite slopes or pits	the drive is bounded by tanh; the authority is clipped
Global conservation	what flows in is matched by what flows out	the drain and the bath's friction match the drive
Recurrence	a disturbed state returns arbitrarily close to where it was	convergence to the target pole at rate K_P
One law	every subsystem obeys the same law, no special cases	one engine, one compass, every muscle




Plate 3. The closed circle tested

The closed circle tested

The canonical evolution law has four parts: deviation is damped nonlinearly and coupled to alignment and to the basin; alignment relaxes toward unity and responds to deviation; the basin's curvature evolves with internal structure; damping absorbs excess growth and redistributes energy. Bounded random noise is permitted; divergence is not. A cycle is identified when the system returns arbitrarily close to a prior state within a bounded tolerance under noise. The state advances stepwise under a stable numerical integrator with boundedness checks. Under parameter sweeps the test is always the same three checks: no divergence, bounded, recurring.



Plate 4. Symmetry principles and conservation laws in the closed circle

Symmetry and conservation in the closed circle

Every conservation law is the shadow of a symmetry (Noether). The closed circle carries four:

Symmetry	Conserved	In Omni-Compass
Temporal (the law does not change with time)	energy	the same law governs every decision; the receipt compares like with like
Spatial (the law does not depend on where)	momentum	the same law governs every muscle in every realm
Rotational	angular momentum	the compass is symmetric about its center; the engine is symmetric between its poles
Internal alignment	alignment invariance	the alignment state is pulled to its pole whichever pole it is

Around them sit the boundary conditions of the circle: maximum deviation (the boundary limit), topological closure (bounded wholeness), nonlinear dynamics (complex interactions), resonant reversal (feedback return), energy-momentum conservation, and cross-scale invariance (scale-independent patterns).



Plate 5. Scale invariance and structural universality

Scale invariance

The theory claims that one bounded nonlinear action has the same form at every scale: cosmological (curvature geometry, compression and expansion, thermodynamic cycling), galactic (rotation curves, mass-energy distribution), stellar (gravitational compression, fusion equilibrium, radiative transport), atomic (electron displacement, orbital coherence, radiative damping) and quantum (localized deviation wells, bounded oscillation, finite amplitude). The universal criteria for such a law are that it operates at every scale and preserves bounded deviation, finite curvature and recurrence. The state vector carries the same meaning at every scale: E the deviation magnitude, U the coherence or flow alignment, S the structural curvature or basin geometry, B the redistribution or damping.



Plate 6. The lineage of unification

The lineage of unification

Each step in the history of physics unified two things that had been separate: geometry formalized structure; motion was unified across earth and sky; astronomy revealed recurrence; fields replaced action at a distance; curvature unified gravity, though singularities remained; symmetry came to govern the quantum realm. The closed circle is offered as the next step: unity without divergence, a single bounded law with no singularity and no runaway.




Plate 7. A candidate for a closed, unified law

What is tested and what is theory

The software uses the engineering core of this principle, and every part of that core is checked:

- The six-state engine, its bounded control command held through the integration step, and the convergence of its alignment state to its pole are proved (docs/TRACKING_THEOREM.md) and tested against a C++ twin.
- The invariance and boundedness of each engine state were checked state by state (Chapter "Closing the Circle").
- The compass law pulls every muscle's service reading to the middle of its band under a bounded, smooth force, and its behavior is measured in simulation and on real software.

The wider claims - one law across every physical scale, the cosmological and quantum readings, the candidate unified action - are the theory from which Omni-Compass was built. They are stated here so that the engineer understands where the design comes from. They are not claims that the software has measured, and no result in this book depends on them.




Plate 8. The Omni-Compass Grand Theory: the four pieces

CHAPTER 3

The Four-Piece Engine

The engine is built in four pieces, stacked like the stages of a press, each feeding the next:

Piece	Name	What it establishes
1	Foundational Engine	the base laws, the states and the boundaries: a closed, bounded, nonlinear adaptive engine with a four-state dynamical structure (misalignment, alignment recovery, coherence potential, path interaction), bounded convergence and persistence with decay, intrinsic feedback and stabilization, and an empirical validation framework for recovery
2	Geometric Response	geometric alignment and response dynamics: a geometric coherence field, a unified potential governing geometric deformation, and the propagation mechanism of that response
3	Coherence Evolution	coherence through response and adaptation: the coherence-response envelope, the system response field, and the coupling of coherence state with response propagation
4	Unified Master Field Law	one coherent field structure: it integrates pieces 1-3 into one closed dynamical structure with a unified potential governing all subsystems, bidirectional propagation between engine state, geometry and response, bounded convergence across the whole system, and no external authority: all evolution emerges from internal coupling

The four pieces compress every domain through one center core. The compression process has four stages: the domains unordered; compression in progress; coherence achieved; return to equilibrium. The result has four properties: coherent, bounded, admissible, stable - one law, one return, one cycle.

The 24 domains. The theory gathers 24 domains into four quadrants around that core:

Quadrant	Domains
1. Foundational structures	Newtonian gravity; Maxwell electromagnetism; Einstein relativity; Schrodinger quantum mechanics; Shannon information theory; Dirac quantum field theory
2. Dynamics and symmetries	Noether symmetry laws; Hamiltonian dynamics; Lagrangian mechanics; thermodynamics; topology; differential geometry
3. Coherence evolution	gauge field theory; fluid dynamics; field theory; discrete mathematics; solid-state physics; plasma physics
4. Complex systems and integration	astrophysics; consciousness and information integration; biological systems and self-organization; universal mathematical structure; systems theory and complex adaptive systems; the unified Omni-Compass domain (the unified admissible state)

OMNI-COMPASS

THE UNIFIED CONTAINER OF ALL REALITY

24 DOMAINS • 4 PIECE ENGINE • 1 CLOSURE LAW
ONE TWIST • ONE LAW • ONE RETURN • ONE REALITY

THE FOUR PIECE ENGINE

PIECE 4 UNIFIED MASTER FIELD LAW

Unifies all domains into one coherent field structure.

PIECE 3 COHERENCE EVOLUTION

Evolves coherence through response and adaptation.

PIECE 2 GEOMETRIC RESPONSE

Creates geometric alignment and response dynamics.

PIECE 1 FOUNDATIONAL ENGINE

Establishes the base laws, states, and boundaries.

THE FOUR PIECE HYDRAULIC COMPRESSION SYSTEM

PIECE 4 UNIFIED MASTER FIELD LAW

PIECE 3 COHERENCE EVOLUTION

PIECE 2 GEOMETRIC RESPONSE

PIECE 1 FOUNDATIONAL ENGINE

THE UNIFIED CIRCLE PRINCIPLE

$$X = G(X), X(0) \in \Omega \subset \mathbb{R}^n$$

$$G(X) \cdot n(X) \leq 0 \text{ on } \partial\Omega,$$

$$\nabla L(X) \cdot G(X) \leq 0$$

$$\Rightarrow \lim_{t \rightarrow \infty} X(t) \in M^* \subset \mathcal{F}$$

WHAT'S INSIDE

- ALL PHYSICS
- ALL MATTER
- ALL ENERGY
- ALL INFORMATION
- ALL GEOMETRY
- ALL TIME & SPACE
- ALL FORCES
- ALL FIELDS
- ALL PARTICLES
- ALL STATES
- ALL STRUCTURES
- ALL POSSIBILITIES

THE MACHINE OF REALITY

- ALL INGREDIENTS
- ALL LAWS
- ALL EQUATIONS
- ALL STRUCTURES
- ALL STATES
- ALL POSSIBILITIES

THE RESULT: ONE UNIFIED REALITY

✓ COHERENT ✓ BOUNDED ✓ ADMISSIBLE ✓ STABLE
ONE LAW • ONE RETURN • ONE CYCLE

WHAT IT DOES

- MIXES ALL DOMAINS
- CREATES COHERENCE
- ENFORCES BOUNDEDNESS
- ENSURES ADMISSIBILITY
- DRIVES THE RETURN CYCLE
- CLOSES ALL LOOPS
- MAKES REALITY STABLE

24 DOMAINS 4 PIECE ENGINE 1 CLOSURE LAW THE COMPLETE UNIFIED SYSTEM

www.omni-compass.com

Plate 9. The unified container of all reality: 24 domains, four-piece engine, one closure law




Plate 10. The four quadrants and the compression process




Plate 11. The closure law and the closed-circle engine

In the software, the four pieces appear as the engine's state equations (piece 1: deviation, alignment, memory and basin), the basin potential and its flow (piece 2), the bath that absorbs and redistributes (piece 3), and the one governor that couples every muscle to one law with no outside authority (piece 4).



Plate 12. The Omni-Compass rose




Plate 13. The Omni-Compass Principle

CHAPTER 4

The Compass

The Omni-Compass rose is the whole mechanism drawn as one picture.

The ring. Twenty-four letters, the complete Greek alphabet from Alpha to Omega: the complete set, nothing missing, one letter for each domain. On a ring the end runs back into the beginning: after the last letter the ring returns to Alpha. That is the closed circle: identity, return, cycle.

The four points. Four letters are lifted out of alphabetical order and set on the four points of the rose:

Point	Letter	Mark	Meaning in the engine
North	Alpha, the first	+	the positive pole of the alignment's double well, $U = +1$
South	Omega, the last	-	the negative pole, $U = -1$
East	Beta	>	the drive coming in (the beta terms that fill the deviation)
West	Gamma	<	the coupling that carries energy out to the bath, where friction absorbs it

The north-south line is polarity; the east-west line is flow in against flow out. A system stays inside its circle exactly when what flows out can always match what flows in.

The inner marks.

Mark	Meaning	In the mechanism
+ and -	the two polarities	the two poles of the double well
> and <	push out, pull in	the antagonist pair: the up force and the down force
the double arrow	exchange both ways	the two-way wire: read in, write out
the up-tack	the ceiling	the cover: nothing goes past the top
approximately-equal	the wave, the tolerance	the band and its cushion: hold near the center
the star	the guiding star	the target the needle points to

The center. The spiral at the middle is the attractor: every path winds inward to one point.

Opposite letters. Across the center, zeta faces xi: both are the engineer's symbol for the damping ratio, and a damping ratio of one is the critical glide to the center with no overshoot. Lambda faces psi: the eigenvalue and the wave function of the eigenvalue equation.




Plate 14. Geometric funneling and basin formation

CHAPTER 5

Basins, Polarity and the Dual-Basin Engine

Geometric funneling and basin formation

A basin forms in three stages. While the driving tension T is low, curvature is weak and there is no dominant basin: flows spread outward. As T increases, curvature concentrates and a positive feedback loop begins. Then a stable geometric funnel forms: self-reinforcing curvature, finite depth, bounded deviation, with the boundary $\text{abs}(E) \leq E_{\text{max}}$ respected. Structure emerges from fluctuation through curvature concentration and feedback.

In the engine, the basin state S rolls down its potential $\Phi(S)$ to a stable bottom; in the compass, every muscle's service reading is pulled into the bottom of its own funnel.

Polarity

Two like charges repel without anyone pushing them: the force comes from the shape of the field. A wall built of like polarity at the edge of a band pushes the state back harder the closer it comes (as one over the distance squared), and two such walls cancel at the center, where the state comes to rest. This is the boundary condition of the closed circle written as a force: the field itself points inward at the wall.

The repulsion does not create energy: the field stores what it took to bring the charges together. What is gained in the computer is that the force is computed rather than spent: the brain computes it at no physical cost, and the machine's energy is saved because its knob no longer fights itself.

A symmetric pair of repelling walls was tested on the modelled GPU card. It raised work per energy but also raised the time spent past the service line, because the lower wall pushed capacity away when the card was calm. It was removed. The upper wall alone remains a candidate.




Plate 15. The Omni-Compass dual-basin engine

The dual-basin engine

The dual-basin engine is a recursive transform with five stages:

1. **Compression.** Two outer basins of positive polarity draw matter and energy inward in a spiral; rotation organizes the flow; compression increases toward the center.
2. **Saturation.** At the center, two cores of negative polarity meet back to back; the basin structure S reaches its critical value and the curvature reaches its maximum.
3. **Transition (ignition).** Where the cores touch, the transition operator $T(X) = \sigma(S - S_c) T(X)$ activates: compression halts and redistribution begins.
4. **Redistribution.** Two axial jets carry the excess out along the axis, with opposite helicity to the inflow.
5. **Reset and recurrence.** The domain relaxes and the cycle repeats.

The core mathematics is a bounded flow: $dX/dt = G(X, t)$ with $\text{abs}(G(X)) \leq C_G$, a curvature envelope $K(X) \leq K_{\text{max}}$, and a redistribution term that activates only at saturation. The state $X = (E, U, S, B)$ stays in its admissible domain, whose horizon boundary is forward invariant (the Nagumo condition again).

For an engineer, the transform reads: pull toward the center (compression); a wall where the muscle saturates; and, when one muscle saturates, move the excess to its siblings rather than hammering it (redistribution). The first two are built into the compass. Redistribution across muscles is designed and not yet built; it is on the roadmap.

PART TWO

The Mathematics

Eight equations, one control law, and the proofs, audits and declarations that hold them fixed.

CHAPTER 6

The Canonical Declaration

Every report the engine has ever produced opens with the same declaration. It is the birth certificate of the mechanism: what the engine is, which equations it runs, what every symbol means, which parameters it samples, how a run is born, conveyed and certified, and what the engine is allowed to say about itself. This chapter sets it down in full so that anyone who reads the rest of this book can always come back to the source.

The identity of the engine

The Omni-Compass Unified Governing Convergence Control Core Engine. One canonical eight-line mathematical core, called Piece I, integrated by fourth-order Runge-Kutta (RK4) and exercised by Monte Carlo. Around that core sit the benchmark suites: the public enterprise manager mechanism benchmark, the one-stack-normalized global megastack, fragmented stack governance over matched scenarios, and the enterprise evidence suite covering cause, fail-safe, scale, muscle, security and economics. All of them run on bare metal, on premises, private, public, hybrid and edge alike, because the core does not care where it runs.

The equation form (authoritative, eight lines)

```
(1) dE/dt = -α_E E + β_int(t) + β_ext(t) + v_eff
(2) dU/dt = μ U(1 - U²) - (dE/dt)/E_max - λ_U U + u,      |u| ≤ U_AUTHORITY
(3) dI_U/dt = 1 - U - σ_1 E - δ S - λ_I I U
(4) v_eff = χ(t) c · tanh(λ_0 + λ_1(U - U_t) + λ_2 S),      χ(t) = cos(ω_B t / 2)
(5) Φ(S) = ½ α_s S² + ¼ β_s S³ - δ S
(6) dS/dt = F_state(S) = -∂Φ(S)/∂S = δ - α_s S - ¾ β_s S²
(7) B'' + (ω_B / Q_B) B' + ω_B² B = γ_c δ S [continuous]
(8) R_B[n] = D²_h B[n] + (ω_B / Q_B) D_h B[n] + ω_B² B[n] - γ_c δ S[n] = 0 [discrete audit]
```

Line (1) is the error, the energy of misalignment: it decays at its own rate, it is pushed by forces from inside and outside, and it is pushed by the effective velocity of the system. Line (2) is coherence, the alignment index: it lives in a double well with two homes at +1 and -1, it is pulled down by any rise in error, it leaks, and it is the one line the governor touches, through the control u , which is never allowed beyond its authority. Line (3) is the information channel that keeps the record of how far coherence has been from home. Line (4) is the effective velocity, saturated so that it can never exceed the limiting speed c , and multiplied by the spinor phase. Lines (5) and (6) are the potential of the coherence state and the flow it creates, downhill. Line (7) is the bath, the field that every real system sits in, driven by the state. Line (8) is the bath written on the grid the computer actually steps on: it is an audit of the numbers, not a second bath.

The control law

```
σ = nearest_signed_basin(U_0)
u_m = sat(-f_U(x_m, t_m) + K_P (σ - U_m), ± U_AUTHORITY)
x_(m+1) = RK4_h(x_m; u_m), u_m held across all four RK4 stages
```


The nearest signed basin is declared at Step 0, before the first microstep. The control cancels the natural drift of coherence and adds a proportional pull to the declared home, then it is clipped at its authority. It is held fixed through the four stages of each RK4 step. There is no post-step overwrite of the state and no projection onto the basin. Whatever the engine reaches, it reaches by the flow of its own equations.

THE OMNI-COMPASS UNIFIED GOVERNING CONVERGENCE CONTROL CORE ENGINE

THE OMNI-COMPASS EQUATION FORM

1	$\partial E = -\alpha E + \beta_{\text{int}}(t) + \beta_{\text{ext}}(t) + v_{\text{eff}}$
2	$\partial U = \alpha(1 - U) - \frac{\partial E}{E_{\text{max}}} - kU(1 - U)$
3	$\partial I_U = 1 - U - \sigma_1 E - \delta S$
4	$v_{\text{eff}} = c \tanh[\lambda_0 + \lambda_1(U - U_t) + \lambda_2 S]$
5	$\Phi(S) = \frac{1}{2}\alpha_s S^2 + \frac{1}{4}\beta_s S^3 - \delta S$
6	$\partial S = F_{\text{state}}(S) = -\frac{\partial \Phi(S)}{\partial S} = \delta - \alpha_s S - \frac{3}{4}\beta_s S^2$
7	$\partial^2 B + \frac{\omega_B}{Q_B} \partial B + \omega_B^2 B = \gamma_c \delta S \quad [\text{continuous}]$
8	$\Delta^2 B + \frac{\omega_B}{Q_B} \Delta B + \omega_B^2 B = \gamma_c \delta S \quad [\text{discrete}]$

THE OMNI-COMPASS SYMBOL CHART

E Error / misalignment energy	U Coherence / alignment index	I_U Information / internal- coherence channel	S Coherence-state potential	B Bath field
B_{dot} Bath-field time derivative	α Feedback rate	β_{int}(t) Internal forcing function	β_{ext}(t) External forcing function	k Coherence logistic damping gain
σ₁ E-coupling into I _U	δ Linear coherence coupling parameter	λ₀ Saturation bias	λ₁ Saturation gain	λ₂ S-state coupling gain
U_t Target coherence state	c Limiting velocity	E_{max} Per-trajectory rate normalization	α_s S-potential quadratic coefficient	β_s S-potential cubic coefficient
Φ(S) Potential-field function	F_{state}(S) Canonical S-state flow	ω_B Bath angular frequency	Q_B Bath quality factor	γ_c Bath coupling gain

www.omni-compass.com

U.S. PATENTS / COPYRIGHTS / TRADEMARKS — © ® 2026

Plate 18. The equation form and the symbol chart

The symbol chart

Symbol	Meaning
E	Error, the misalignment energy
U	Coherence, the alignment index
I_U	Information, the internal-coherence channel
S	Coherence state potential
B	Bath field
B'	Bath field time derivative
α	Feedback rate
$\beta_{\text{int}}(t)$	Internal forcing function
$\beta_{\text{ext}}(t)$	External forcing function
k	Coherence logistic damping gain
σ_1	Coupling of error into the information channel
δ	Linear coherence coupling parameter
λ_0	Saturation bias
λ_1	Saturation gain
λ_2	S-state coupling gain
U_t	Target coherence state
c	Limiting velocity
E_{max}	Per-trajectory error-rate normalization scale
α_s	S-potential quadratic coefficient
β_s	S-potential cubic coefficient
$\Phi(S)$	Potential field function
$F_{\text{state}}(S)$	Canonical S-state flow
ω_B	Bath frequency
Q_B	Bath quality factor
γ_c	Bath coupling gain

The coupling architecture of Piece I

The closed feedback core is E , U and S : error drives coherence, coherence and the state drive the velocity, the velocity drives error. The information channel I_U and the bath B with its derivative are driven response channels: they listen to the core and record it, and the core does not depend on them. γ_1 is reserved for compatibility and is inactive in the Piece I derivatives. The bath law uses $\omega_B^2 B$, so ω_B carries the units of angular frequency. Line (7) advances B and B' continuously by RK4. Line (8) evaluates the discrete residual with $D^2_h B[n] = (B[n+1] - 2B[n] + B[n-1]) / h^2$ and $D^-_h B[n] = (B[n] - B[n-1]) / h$.

The spinor closure

The spinor extension introduces the phase $\theta(t) = \omega_B t$ and multiplies line (4) by $\chi(\theta) = \cos(\theta/2)$. At 0 degrees the factor is +1, at 360 degrees it is -1, and at 720 degrees it is +1 again. A full turn of the bath reverses the sign of the velocity; only two full turns bring it home. The standard 20-step run evaluates the part of the phase it reaches, and a dedicated evidence test checks the 0, 360 and 720 degree periodicity on its own.

The signed basins

The canonical Monte Carlo sampler draws the starting coherence $U(0)$ symmetrically on $[-1, +1]$. The nearest signed basin at Step 0 is declared the target before the first microstep. That keeps the eight-line equation exactly as written while making every governed run exercise capture into both $+1$ and -1 . A separate dual-basin suite throws runs across from one basin to the other under disturbance and watches them recover.

Birth, conveyance and certification

Step 0 is the sampled state before integration. Step 1 is the first completed macro step. A run has at most 20 macro steps of 10 micro steps each, at a time step of 0.1. A step is in basin when U lies within 0.10 of $+1$ or -1 .

- **Born in basin.** A run that starts inside a basin at Step 0 is born in basin for life. Step 0 counts as the first observation of its streak. It is confirmed conveyed after Steps 0 to 4 stay in the same basin. It is never called self-tuned, and it stays under the same regulation for all 20 macro steps.
- **Self-tuning.** A run that starts outside both basins is self-tuning applied from Step 1. The same nearest-basin law is active on every trajectory, born or not, over the whole 200-microstep horizon. The category never switches the regulation off.
- **CONVEY-5.** Five consecutive same-basin observations. The first is the conveyance entry step, the fifth is the confirmation step. The latest streak may start at Step 16 and confirm on Step 20.
- **CERT-10.** Ten consecutive same-basin observations. The latest streak may start at Step 11 and complete on Step 20. The code asserts both bounds before it runs and checks every emitted row against them.

Conveyance is an event, not an ending. No run may be both born in basin and self-tuned. Stability is judged on the history, not on one snapshot.

The Monte Carlo declaration

Setting	Value
Initial U distribution	Uniform on $[-1, +1]$, symmetric
Monte Carlo runs	500
Macro steps per run	20
Micro steps per macro step	10
Integration time step	0.1
Basin centers	$+1$ and -1
Basin tolerance	0.10
Stability window	5
Conveyance cutoff step	20

The full horizon is always executed. A run is classified as a failure if it crosses the divergence threshold at any recorded state, or if it exhausts the horizon without CONVEY-5.

The parameter sets

Group	Parameter	Value or range
Core	α	4.2
Core	β_{int}	0.0 to 1.0
Core	β_{ext}	0.0 to 1.0
Core	E_{max}	1.0 to 10.0, sampled once per trajectory
Core	k	1.7
Information and coupling	σ_1	0.38
Information and coupling	δ	0.01 to 1.0
Information and coupling	γ_1	1.35, reserved, inactive in the derivatives
Information and coupling	γ_{c}	1.0
Saturation and phase	c	0.5 to 5.0
Saturation and phase	λ_0	-2.0 to 2.0
Saturation and phase	λ_1	0.1 to 3.0
Saturation and phase	λ_2	0.0 to 2.0
Saturation and phase	U_{t}	0.5
Saturation and phase	$\chi(t)$	$\cos(\omega_{\text{B}} t / 2)$ when the spinor closure is on
S-state potential	α_{s}	0.05 to 0.25
S-state potential	β_{s}	0.05 to 0.25
Bath	ω_{B}	0.1 to 5.0
Bath	Q_{B}	0.5 to 10.0
Noise and drift	D	0.0
Noise and drift	μ	0.0

The mechanism-bound reporting doctrine

Every table, grid, chart, scorecard, classification and conclusion the engine emits must be computed from what it actually executed: the trajectory rows, the recorded state histories of E, U, I_U, S, B and B', the sampled parameters, the RK4 derivative evaluations, the basin entry and dwell histories, the line (8) residuals, the solver comparisons, the ablations, the disturbance and stress runs, the compute benchmark event records, and the enterprise benchmark ledgers.

No decorative grade, inferred score, hand-assigned quality category or presentation-only number is allowed to stand as an engine result. Prose may point to the exact source and the edge of the claim. It may not manufacture evidence, put a grade where a measurement belongs, or present a grid the mechanism never ran as though it had. That doctrine runs through every page of this book.

Availability and the rights boundary

Time is of the essence. Omni-Compass is available from The Omni-Compass LLC for controlled technical evaluation, research collaboration, pilot integration, strategic engagement and separately licensed commercial deployment. The same rights boundary applies at every level of use: inspection, local execution, simulation, research, benchmarking, observe-only, shadow mode, supervised control, bounded autopilot, hybrid control, direct-to-muscle control, production, embedding, hosted service, productization, monetization, redistribution and derivative work.

Public access does not make Omni-Compass open source. Commercial use of any kind requires a separately executed written license. The instrument is a license, not a sale and not a transfer of ownership. For planning only, the introductory benchmark is 10% of independently verified and contractually accepted value captured; the base and the amount are specific to each company, the terms are expected to rise as validation and adoption grow, and only a signed agreement creates any obligation.

Evaluation and simulation use only. Copyright (c) 2026 The Omni-Compass LLC. Commercial use, commercialization or monetization of any part of Omni-Compass requires a signed, paid Omni-Compass Enterprise License from The Omni-Compass LLC. Patent applications, copyright registrations and trademark applications filed in the United States. See LICENSE and NOTICE at the root of this repository.

CHAPTER 7

The Engine: Eight Equations and One Control Law

The engine (`omnicompass/core.py`, frozen and fingerprinted) carries a six-part state $x = (E, U, I_U, S, B, B_{\dot{}})$:

State	Meaning	Physical picture
E	deviation: how far the system is from where it should be	a tank that fills with stress and drains on its own
U	alignment, between the two poles -1 and +1	a ball in a double well: two stable poles, a hill between them
I_U	memory of misalignment	an integrator that remembers and slowly forgets
S	basin structure	a ball rolling to the bottom of its landscape
B, $B_{\dot{}}$	the bath	a spring with friction that absorbs and settles energy

The eight equations (Appendix C gives them in full):

1. $dE/dt = -\alpha_E E + \beta_{\text{int}} + \beta_{\text{ext}} + v_{\text{eff}}$
2. $dU/dt = \mu U (1 - U^2) - (dE/dt)/E_{\text{max}} - \lambda_U U + u$, with $|u| \leq 25$
3. $dI_U/dt = (1 - U) - \sigma_1 E - \delta S - \lambda_I I_U$
4. $v_{\text{eff}} = \cos(\omega_B t / 2) c \tanh(\lambda_0 + \lambda_1 (U - 0.5) + \lambda_2 S)$
5. $\Phi(S) = \alpha_S S^2/2 + \beta_S S^3/4 - \delta S$
6. $dS/dt = -d\Phi/dS$
7. $d^2B/dt^2 = \gamma_c \delta S - (\omega_B/Q_B) dB/dt - \omega_B^2 B$
8. R_B : a finite-difference audit of (7), never fed back

The control law. $u = \text{clip}(-f_U(x, t) + K_P (\sigma - U), -25, +25)$, where f_U is the drift of U with the command at zero, σ is the target pole and $K_P = 12$. The first term cancels whatever is shoving U (the push); the second pulls U to its pole, harder the farther it is (the pull). The command is held unchanged through every stage of the fourth-order Runge-Kutta step and U is never rewritten afterwards. Inside the authority limit, U converges to its pole at rate K_P ; this is proved (`docs/TRACKING_THEOREM.md`).

What each equation does, in words

The six numbers describe the whole system the way a physician's six vital signs describe a patient, and each equation says how one of them moves.

- **Equation 1, the deviation tank.** E is how hard the system is being driven: queues, overload, power, heat. It drains on its own at a fixed rate ($\alpha_E = 3.487$) and is filled by two pressures, internal (queue and overload, β_{int}) and external (power, heat, network, staleness, security, β_{ext}), plus a drive term v_{eff} that is bounded by a hyperbolic tangent. However extreme the observation, the engine's response rate is capped at $c =$

1. This is the saturation, the "speed limit" that no input can exceed.

- **Equation 2, the alignment.** U is how healthy and settled the system is, and it lives in a double well: two stable poles at $+1$ (healthy) and -1 (failed) with an unstable ridge at 0 . A rising deviation pushes U out of the healthy well; the bounded command u pulls it back. The cubic term $\mu U (1 - U^2)$ is what makes the walls of the well push back harder the farther U strays.
- **Equation 3, the memory.** I_U grows while U is below its pole, is reduced by deviation and stress already being dealt with, and leaks at $\lambda_I = 0.719$. It is the engine's memory of unrelieved pressure: a system that has been stressed for an hour is not the same as one stressed for a second, and this integrator is what knows the difference.
- **Equation 4, the drive.** The bounded drive v_{eff} mixes the alignment and the stress through a tanh, with a slow cosine of the bath's frequency on top, so the drive breathes rather than pushes constantly.
- **Equations 5 and 6, the stress.** S rolls downhill in a cubic potential to its stable root. It is the slowest restoring force: structural stress from heat, power, network, security and staleness that builds and relaxes over minutes rather than seconds.
- **Equation 7, the bath.** A damped harmonic oscillator (natural frequency $\omega_B = 1$, quality $Q_B = 2.5$) driven by stress. It turns stress into a smooth, lagged tide, and it is the only second-order element in the engine; its friction is what settles the whole.
- **Equation 8, the audit.** The bath equation is re-checked numerically every step with a finite difference, and the check is recorded and never fed back: an instrument on the engine, not a part of it.

Integration is classical fourth-order Runge-Kutta, ten micro-steps of 0.01 per macro step, and the C++ twin matches the Python to 3.6×10^{-15} on 500 reference trajectories (Chapter 61).

From telemetry to state, and from state to action

Each decision, the telemetry is normalised to nine observations: queue, load, power, thermal, network, drift, staleness, security and conflict. An observed state is computed from them by fixed weights (omnicompass/adaptor.py, observe_vector), the engine's state is blended toward it with weight 0.339 , and the state is then evolved by one macro step of the equations. From the state, the governor evaluates the control's push, $u(x)/25$, which is a number between -1 and $+1$ that says how hard the engine is pulling toward health; a target utilisation that falls as the memory and the deviation rise; whether a release (giving something back) is permitted at all, which is only while the push reports convergence and the band and dwell allow; and whether change is permitted, which requires U above its gate and no security block. All of this is written to the audit every decision (docs/MECHANISM_OF_ACTION.md, Chapter 53 to 6, and python tools/mechanism.py, which reproduces the measured numbers in results/MECHANISM_OF_ACTION.json).

Where the engine sits, and where the compass law sits

A reader meeting the engine and then the compass law (Chapter 16) may ask which of the two moves the knob. The answer, in the live controllers, is this: the **compass law** sets each knob, the **verdict** narrows where the law may go, and the **engine** grants the authority, feeds the release gate and the nervous system, and is audited every decision (omni_controller/controller.py, its docstring). The engine is the organism's physiology: how stressed it is, how long it has been, whether it is settling. The compass law is the hand on each lever. The engine's own control law u is evaluated, not applied to a knob directly: it is the push the nervous system reads to decide whether giving anything back is permitted now. This separation is why the engine could be frozen with its proofs untouched while the compass law was written as its own, separately tested law, and why verify.py seals both.

CHAPTER 8

The Canonical Engine

One engine runs Omni-Compass, and every result in this repository comes from it. This page names it, gives its equations exactly as the code computes them, and names every other form as a variant. Where any document, manual, chart or filing states the equations differently, this page and the file it fingerprints are what the software runs.

1. The canonical engine: `symmetric_verified`

Name	<code>symmetric_verified</code> (omnicompass/configurations.py, DEFAULT)
Code	omnicompass/core.py (Python), cpp/src/core.cpp (C++ twin, sealed in results/SEAL.json)
Fingerprint	SHA-256 of omnicompass/core.py, recorded in results/PREREGISTRATION.json and RELEASE_MANIFEST.json; checked by verify.py
Evidence	every simulation, benchmark, live Kubernetes run and GPU harness in this repository

State: E (deviation), U (coherence), I_U (unmet-need integral), S (structural stress), B and B_dot (the bath).

```
v_eff = cos(omega_B t / 2) · c · tanh(lambda_0 + lambda_1 (U - 0.5) + lambda_2 S)      (spinor closure)
↳ e, 720°)
dE/dt = -alpha_E E + beta_int + beta_ext + v_eff
dU/dt = mu U (1 - U²) - (dE/dt) / E_max - lambda_U U + u                          (symmetric double well)
↳ ble well)
dI_U/dt = (1 - U) - sigma_1 E - delta S - lambda_I I_U
dS/dt = delta - alpha_s S - (3/4) beta_s S²                                       (gradient of Phi)
↳ Phi)
dB/dt = B_dot
dB_dot/dt = gamma_c delta S - (omega_B / Q_B) B_dot - omega_B² B
Phi(S) = alpha_s S² / 2 + beta_s S³ / 4 - delta S
```

u is the controller's command, $K_P (target - U) - drift$, bounded by `U_AUTHORITY`; integration is RK4 with the macro and micro steps of `omnicompass/core.py`.

Mechanism identity

`results/MECHANISM_IDENTITY.json` (`tools/mechanism_identity.py`) fingerprints the mechanism $M = (F, \Theta, C, h, G, M_{act}, dt, A)$ component by component: state law, parameters, controller, observation map, authority law, actuator map, execution timing, shield. `verify.py` fails if any component changes without the record being rewritten. Every evidence claim names the mechanism id that produced it; "the eight-line engine" alone names none.

Configuration	Mechanism id	Role
<code>symmetric_verified</code>	29d9808dfb8f626ad5de17a8a1efa37411dbab08a64af4b276143b466f7ce21c	canonical
<code>printed_eight_line</code>	cd333dc166fb7684ec0fca71f5f50488041e47825ee47d1ea667bec1f3d2259d	named alternative embodiment

On the frozen 500-fixture population (benchmarks/core_evidence.py, seed 223387268), both are executable and neither is a stand-in for the other:

	finite	CONVEY-5	CERT-10	mean final target error	mean integrated abs(u)	max abs(u)
symmetric_verified	500/500	500/500	500/500	7.48e-5	0.878	13.93
printed_eight_line	500/500	500/500	500/500	5.13e-6	11.79	18.87

The printed form reaches its target more tightly with about 13 times the control effort. The structure is Option A: one canonical configuration, one named alternative; no equivalence is claimed.

2. Variants

printed_eight_line (*the printed chart*)

Implemented in omnicompass/configurations.py, checked against the printed plate (tests/test_engine_configurations.py), **not benchmarked**: no result in this repository comes from it.

```

v_eff = c * tanh(lambda_0 + lambda_1 (U - U_t) + lambda_2 S)           (no spinor fac
    ↳ tor)
dE/dt = -alpha E + beta_int + beta_ext + v_eff
dU/dt = alpha (1 - U) - (dE/dt) / E_max - k U (1 - U) + u           (logistic)
dI_U/dt = (1 - U) - sigma_1 E - delta S                             (no lambda_I
    ↳ term)
dS/dt, dB/dt, dB_dot/dt as in section 1

```

Differences from the canonical engine: logistic U instead of the symmetric double well; no spinor factor; a general target U_t; no lambda_I damping; alpha in place of alpha_E. alpha_U and k enter the printed form; they do not enter the canonical engine.

3. What would change this declaration

Promoting a variant to canonical needs, in one commit: the variant run through every harness beside the canonical engine, its results reported next to the canonical results, this page and RELEASE_MANIFEST.json updated, and the preregistration amended on record (results/LOCK_AMENDMENTS.json). Until then the canonical engine is symmetric_verified.

4. Open mathematical items (from docs/FORMAL_STATUS.md and the handoff)

- A global stability proof of the forced six-state system is not closed. What is held on the U channel is proved in docs/TRACKING_THEOREM.md: unsaturated exponential tracking; the saturated case under a drift bound that holds over the declared box ($F_{\text{bar}} = 16.86 < 25$); the sampled-data bound of the executed RK4 controller; admissibility.
- Candidate routes: the Unified Circle Principle on a region; Theorem 5.6 of the Closed Structure, if its core maps onto (E, U, I_U, S, B).

5. For filings

Which form a patent or copyright filing claims as the principal embodiment is a decision for The Omni-Compass LLC and its counsel. Whatever that decision, the software described by this repository runs the canonical engine of section 1, and a filing that claims the printed form should name symmetric_verified as the embodiment that has been implemented and tested, or the printed form should be benchmarked first (section 3).

Evaluation and simulation use only. Copyright (c) 2026 The Omni-Compass LLC. Commercial use, commercialization or monetization of any part of Omni-Compass requires a signed, paid Omni-Compass Enterprise License from The Omni-Compass LLC. Patent applications, copyright registrations and trademark applications filed in the United States. See LICENSE and NOTICE at the root of this repository.

CHAPTER 9

The Closed Circle: Why It Cannot Leave Its Compass

The governing principle is the Unified Circle Principle:

$dX/dt = G(X), \quad X(0) \text{ in } \Omega$	
$G(X) \cdot n(X) \leq 0 \text{ on the boundary of } \Omega$	(the flow points inward at the wall: nothing crosses)
$\text{grad } L(X) \cdot G(X) \leq 0$	(the compass's energy L only falls)
$\Rightarrow \lim X(t) \text{ in } M^*$	(every path settles in the bottom of the compass)

In control engineering the second line is the Nagumo condition for an invariant set and the third is a Lyapunov function. Together they mean: anything that starts inside the container stays inside it, for every disturbance up to a known size, and comes to rest at the bottom. That is a certificate, not a test: it answers every case inside the walls at once.

In plain words. Picture a bowl-shaped valley with a fence around its rim. The first line says the valley has a shape: given where the system is, the equations say which way it rolls. The second line says that at the fence the ground always slopes inward, so nothing that starts inside can ever roll out, however it is pushed within the strength the fence was built for. The third line says there is a height, L , that only ever decreases along any path, so the system cannot circle forever at one level; it must keep descending. The conclusion follows: every path ends at the bottom, the set M^* . The engineering value of such a certificate is that it is not a sample of behaviours; it is a statement about all of them at once, inside the fence.

What is certified, and what is not. The certificate is about the engine's own state: the six numbers of Chapter 7 stay inside their walls and settle. It is not a certificate about the plant the engine governs, whose physics the engine does not contain; that is why every plant is tested paired against native (Part VI), why the compass law clips every write to the knob's cover (Chapter 16), and why the shield between the engine and the levers is tested separately on two million adversarial cases with zero violations (`tests/test_shield_properties.py`) and the C++ engine soaked on one hundred million decisions with no failure (`results/SOAK.json`). The mathematics closes the brain; the harness, the cover, the verdict and the paired receipts close the hand.

The engine's states are each closed in this sense: the deviation tank drains faster than it can fill; the alignment's cubic walls push back from far out; the memory forgets; the bath's friction settles it. The rule behind all four is the compass's east-west line: what flows out (drain, damping) must always be able to match what flows in (drive). The formal status of each statement, what is proved, what is tested and what is conjecture, is written out in `docs/FORMAL_STATUS.md` and `docs/TRACKING_THEOREM.md`, and the theory chapters of this book (The Unified Circle Principle, Closing the Circle in the Engine) give the derivations.

What the certificate buys an operator, in practice. Three things. The engine cannot run away: there is no input, no sequence of readings and no fault in a sense that drives its internal state to infinity or into oscillation that grows, because the walls are part of the equations and not a check bolted on afterwards. The engine cannot hold a grudge:

the memory state forgets at a fixed rate, so a bad hour does not bias the next day, and a governor restarted after a kill starts from the same rest state the certificate names. And the engine's behaviour under any disturbance smaller than the fence is the same kind of behaviour, a descent to the bottom, so a test on one plant's readings says something about the engine on another's; what it does not say, and what no mathematics of the engine alone could say, is how the plant answers, which is why the paired runs exist. An underwriter reading this section should take from it exactly this much: the part of the system that makes decisions is bounded by proof; the part of the system that carries them out is bounded by the cover, the gates and the hand-back, each tested; and the part that is the customer's own plant is measured, never assumed.

The compass. The Omni-Compass rose carries the whole Greek alphabet, Alpha to Omega: the complete set, in a ring whose end runs back into its beginning. Its north-south axis is polarity: Alpha and plus at the top, Omega and minus at the bottom, the two poles of the alignment's double well. Its east-west axis is flow: Beta and the push outward on one side, Gamma and the pull inward on the other, the drive and the damping. The spiral at the center is the attractor every path winds into. The ring closing on itself is the return: when Omni-Compass stops, every lever returns to where it began.



Plate 16. Mathematical integration




Plate 17. The engine core: the governing law in motion

CHAPTER 10

Closing the Circle in the Engine

Each state of the engine was checked against the closed-circle conditions:

State	Closed?	Why
E, deviation	closed	it drains at rate α_E while its inputs are bounded (the drive is capped by $\tanh$)
U, alignment	closed	the cubic term $\mu U (1 - U^2)$ pushes back hard from far out
I_U, memory	closed	it forgets at rate λ_I
B, the bath	closed	the friction ω_B / Q_B settles it
S, the basin	closed on the positive side; open on the negative side below an unstable ridge	$\Phi(S)$ carries a cubic term, so the landscape falls away below the ridge (at $S = -3.5$ with the default parameters)

Three openings were found and each has its closure:

1. **The basin's negative side.** A quartic term in $\Phi(S)$ would give the compass a wall on both sides. This changes the frozen engine and is a new version with its own proof.
2. **The outer loop.** The frozen live governor computes the push and pull u and uses it as a convergence signal; it does not send u to a lever. The compass closes this loop: reading, force, plug, lever, read-back. On the modelled card this is the difference between +0.1% and several percent of work per energy (the one-wire governor against the corrected two-wire compass, `results/sim/gpu_two_wire/`).
3. **The corner.** `clip()` is a hard stop; $\tanh$ is its smooth form. The compass uses $\tanh$.

What closing the loop has shown since. The compass law of the second opening is the law every result in the manual was made on: it was frozen as Omni v1 on 5 October 2026 with its proofs untouched, grown to the 945-muscle catalog as v2, and given one gate as v3, the slack gate on speed knobs, after the v2 realms read a service tradeoff on the Physics realm and the tower. That gate is the closed circle applied to a muscle's physics: an axis that is busy more than half the time at full speed has no slack to spend, and a law that spent it anyway would be spending energy for nothing; with the gate, every organism reads superior within its guardrails and zero muscles read worse (`results/realms/REALMS.md`). On real software the same law, read through the plug and the read-back, is the law of the Kubernetes, database, broker and cache tables, with their gains and their losses. The first opening, the basin's negative side, remains open and documented: it would be a new version with its own proof, and no result depends on it.

CHAPTER 11

The Tracking Theorem

This page proves the internal tracking result of the canonical engine (`symmetric_verified`, mechanism id in `results/MECHANISM_IDENTITY.json`) in four steps, from the continuous law to the code as executed. Each statement carries one evidence class (see `docs/EVIDENCE_LEDGER.md`): **T** proved here, **V** checked by computation over a finite, frozen set (`tools/tracking_bounds.py` → `results/TRACKING_BOUNDS.json`).

What this page establishes is **internal**: inside the declared model, the controller conveys U to its target σ . It does not establish that σ is the right target for any outside plant, that the telemetry map is right, that any actuator mapping (HPA, GPU, CPU) is right, any energy saving, production performance, or stability of the whole forced six-state system. Target correctness and target conveyance are separate obligations; this page carries only the second.

Setting

U channel of equation (2): $dU/dt = f_U(x, t) + u$, with the uncontrolled drift

$$f_U(x, t) = \mu U (1 - U^2) - (dE/dt) / E_{\max} - \lambda U$$

and the bounded cancellation-plus-proportional controller (`omnicompass/core.py`, `control_command`)

$$u = \text{sat}(-f_U(x, t) + KP (\sigma - U)), \quad \text{sat}(v) = \max(-u_{\max}, \min(u_{\max}, v)), \quad KP = 12, \quad u_{\max} = U_{\text{AUTHOR}} \hookrightarrow \text{ITY} = 25.$$

Tracking error $e = U - \sigma$, σ in $\{-1, +1\}$.

Theorem 1 (T). Unsaturated tracking

On the unsaturated region $\Omega_{\text{unsat}} = \{x : |-f_U(x, t) - KP e| \leq u_{\max}\}$, $de/dt = -KP e$. Along any interval spent in Ω_{unsat} , $e(t) = e(t_0) \exp(-KP (t - t_0))$, and $V = e^2/2$ satisfies $dV/dt = -KP e^2 < 0$ for $e \neq 0$.

Proof. Inside Ω_{unsat} the saturation is inactive, so $dU/dt = f_U - f_U + KP (\sigma - U) = -KP e$; since σ is constant, $de/dt = dU/dt$. The solution and $dV/dt = e de/dt = -KP e^2$ follow. ■

This is a statement about intervals inside Ω_{unsat} only; it is not global convergence.

Theorem 2 (T). The saturated case, under a declared drift bound

Assumption D. Along the trajectory, $|f_U(x(t), t)| \leq F_{\text{bar}} < u_{\max}$.

Under D:

1. If saturation is active with $e < 0$, then $u = +u_{\max}$ and $de/dt = f_U + u_{\max} \geq u_{\max} - F_{\bar{u}} > 0$; symmetrically for $e > 0$. So $dV/dt \leq -(u_{\max} - F_{\bar{u}}) |e| < 0$ wherever saturation is active.
2. With Theorem 1, V is strictly decreasing wherever $e \neq 0$, saturated or not; e never changes sign (at $e = 0$ the command is unsaturated because $F_{\bar{u}} < u_{\max}$, so $e = 0$ is an equilibrium of the error dynamics).
3. Any saturated stretch ends within $(|e(t_0)| - e^*) / (u_{\max} - F_{\bar{u}})$ time units, $e^* = (u_{\max} - F_{\bar{u}}) / KP$; the ball $|e| \leq e^*$ lies inside Ω_{unsat} and is forward invariant, and inside it the decay is exponential at rate KP .

Proof of 1. Saturation with $u = +u_{\max}$ means $-f_U - KP e > u_{\max}$, so $-KP e > u_{\max} + f_U \geq u_{\max} - F_{\bar{u}} > 0$, so $e < 0$, and $de/dt = f_U + u_{\max} \geq u_{\max} - F_{\bar{u}}$. The other side is symmetric. 2 and 3 follow from 1 and Theorem 1, and $|-f_U - KP e| \leq F_{\bar{u}} + KP |e| \leq u_{\max}$ on the ball. ■

Where D holds (T). On the declared box — parameters in `PARAMETER_RANGES`, E_0 in $[0, 1]$, $|U| \leq 1.1$ — equation (1) keeps E in $[(\beta - c)/\alpha_E, (\beta + c)/\alpha_E]$ widened to hold E_0 (the right side of (1) points inward at both ends, because $|v_{\text{eff}}| \leq c$), so $|E| \leq 7/4.2$, $|dE/dt| \leq 4.2 \cdot |E| + 2 + 5 \leq 14$, and

$$F_{\bar{u}} = 6 \cdot 2/(3\sqrt{3}) + 14 / 1 + 0.5 \cdot 1.1 = 16.86 \leq 25 = u_{\max}.$$

Because e keeps its sign and $|e|$ does not grow (continuous time), U stays between U_0 and σ , so $|U| \leq 1$ and the box is self-consistent for the nearest-basin target. **Theorem 2 therefore holds globally over the declared box, in continuous time.** The actuator can still saturate: $|raw| \leq F_{\bar{u}} + KP |e_0|$ can exceed 25 when $|e_0|$ is near 1.

Checked (V). Over the 500 frozen fixtures: observed max $|f_U| = 6.84$ (nearest target), 7.01 (wrong target), each under that fixture's own bound (max 8.99) and under 16.86; saturated micro steps 0 of 100 000 (nearest target), 5 of 100 000 (wrong target).

Theorem 3 (T with a computed constant). The executed sampled-data controller

The code does not run the continuous law. Per micro step $h = \text{MICRO_DT} = 0.01$, it computes u_k from x_k , holds it through the four stages of one RK4 step (zero-order hold), and takes $x_{(k+1)} = F_h(x_k, u_k)$. There is no overwrite of U after the step, no projection onto a basin, no state replacement.

Integrating the U equation over one held interval,

$$e_{(k+1)} = e_k + h (f_U(x_k) + u_k) + \int_{t_k}^{t_{k+h}} (f_U(x(t)) - f_U(x_k)) dt + \tau_k$$

with τ_k the RK4 truncation error. Unsaturated, $f_U(x_k) + u_k = -KP e_k$, so

$$e_{(k+1)} = \rho e_k + d_k, \quad \rho = 1 - h KP = 0.88, \quad |d_k| \leq (h^2/2) \sup |df_U/dt| + |\tau_k| =: \epsilon_{\text{h}}.$$

Then by induction $|e_k| \leq \rho^k |e_0| + (1 - \rho^k)/(1 - \rho) \cdot \epsilon_{\text{h}}$, and the error enters and stays in the neighbourhood $|e| \leq \epsilon_{\text{h}} / (1 - \rho)$.

The proof above is exact given ϵ_{h} . **ϵ_{h} itself is computed, not proved (V):** the largest $|d_k|$ over every micro step of the 500 frozen fixtures is 0.00378 (nearest target), giving the ultimate bound 0.0315, inside $\text{BASIN_TOL} = 0.10$ that CONVEY-5 and CERT-10 test. With the target aimed wrong, $\epsilon_{\text{h}} = 0.0243$ (dominated by the long transient, which includes the 5 saturated steps) and the bound 0.20 is wider than the basin — a conservative bound; the fixtures still certify 500 of 500.

What the discrete execution does **not** inherit: the continuous monotone decay. Within the band, $|e|$ grew on 27 061 of 100 000 micro steps and e changed sign 747 times — the sampled controller dithers inside $\epsilon_h / (1 - \rho)$.

CONVEY and CERT. They are dwell checks against a declared neighbourhood ($|U - \sigma| \leq 0.10$ for 5 and 10 consecutive macro steps). Theorem 3 says why they pass when the bound sits inside the neighbourhood. They are not an external certification.

Theorem 4 (T continuous; V discrete). Admissibility

Admissible set for a fixture with parameters p and initial state x_0 :

$$A = \{ x : E \text{ in } [E_{\text{lo}}, E_{\text{hi}}], |U - \sigma| \leq |U_0 - \sigma|, S \text{ between } S_0 \text{ and } S_{\text{plus}}(p) \}$$

$E_{\text{lo}}, E_{\text{hi}}$ as in Theorem 2; S_{plus} the stable root of equation (6).

- **Continuous (T).** A is forward invariant: E by the inward-pointing field at both ends; U by Theorem 2; S because equation (6) is the one-dimensional gradient flow of Φ , monotone toward S_{plus} from any $S_0 > S_{\text{minus}}$ (`s_admissibility` in `results/core_evidence.json`: 500 of 500 fixtures start above S_{minus}).
- **Discrete (V).** For the executed map $F_h(x, C(x))$, every micro step of every fixture: x_k in A implied $x_{(k+1)}$ in A — 0 failures in 100 000 steps, target nearest or wrong. This is a finite verification on the frozen fixtures, not a proof over the box. An interval-arithmetic proof over the whole box is **open (O)**.

What remains open (O)

- Global stability of the forced six-state system (E, I_U, S, B are not all covered above).
- A proved (not computed) ϵ_h over the declared box, and discrete invariance over the box.
- The inheritance / intertwining residual against an outside plant: instrumented (`tools/gpu_reps.py`, representation fidelity), not yet measured on a real card.
- The printed configuration (`printed_eight_line`) is not covered by this page: its U drift is logistic, not the double well, and needs its own bound (`docs/CANONICAL_ENGINE.md`).

Evaluation and simulation use only. Copyright (c) 2026 The Omni-Compass LLC. Commercial use, commercialization or monetization of any part of Omni-Compass requires a signed, paid Omni-Compass Enterprise License from The Omni-Compass LLC. Patent applications, copyright registrations and trademark applications filed in the United States. See LICENSE and NOTICE at the root of this repository.

CHAPTER 12

The Conveyance Law

1. What the manuscript says, read as mechanism

Source	Text	Mechanism
Ch. 29 §8	"E_system = E + U + S + B ... dE_system/dt = 0 (over full cycle). Energy redistributes. Energy reorganizes. Energy does not diverge."	The budget (site watts, cluster cores) is conserved : it is moved between organs, never created.
Ch. 31 §4	"Redistribution B acts as mediator between the baths. When E grows locally, B increases to distribute gradients."	Budget flows down the gradient of need : from organs holding surplus to organs in deficit.
Ch. 31 §5	"eigenvalues of the coupled system must remain bounded ... $\text{Re}(\lambda_i) \leq 0$ "	The flow must have a proof of convergence, not a tuned damper.
Ch. 30 §5-7	"Redistribution surge ... channels stored deviation into expansion ... controlled release ... No discontinuity."	A reserve is held and released continuously to organs at their limit. There is no on/off.
Ch. 29 §6	"Backpressure develops through S ... Compression is ... structured inward rotation."	When need falls, surplus is pulled back smoothly. Budget nobody needs is not spent.
App. J	"the engine can only act through what you expose"	Organs keep their own mechanisms. The law sets only each organ's share of the budget.

2. The law

Notation.

- Organs $i = 1..n$ share a budget A .
- Organ i holds allocation a_i and has need d_i ; the budget that serves its work at the engine's utilisation target ρ .
- Its deficit (local deviation) is

$$e_i = d_i / (\rho a_i) - 1 \quad (e_i > 0: \text{needs juice}; e_i < 0: \text{holds surplus})$$

The redistribution mediator:

$$da_i/dt = \kappa a_i (e_i - \bar{e}), \quad \bar{e} = \sum_j a_j e_j / A$$

3. Properties and proofs

P1. Conservation.

- $\sum_i da_i/dt = \kappa (\sum_i a_i e_i - \bar{e} \sum_i a_i) = \kappa (A \bar{e} - \bar{e} A) = 0$.
- The budget is moved, never created, as in Ch. 29 §8.

P2. Gradient flow with a Lyapunov function.

- Substitute e_i : $a_i(e_i - \bar{e}) = d_i/\rho - a_i D/(\rho A)$, where $D = \sum d_i$.

- So $da_i/dt = (\kappa/\rho)(d_i - a_i D/A)$.
- This is the Shahshahani gradient on the simplex $\Sigma a_i = A$ of $F(a) = \sum_i (d_i/\rho) \ln a_i$. The Shahshahani metric weights by $1/a_i$, and $\partial F/\partial a_i = d_i/(\rho a_i)$, so $a_i(\partial F/\partial a_i - \text{mean})$ is exactly the flow above.
- Along the flow $dF/dt = \sum_i (1/a_i)(da_i/dt)^2 \cdot (\rho/\kappa) \geq 0$.
- F is strictly concave, so F is a Lyapunov function and the flow cannot cycle.

P3. Global exponential convergence.

- The ODE in P2 is linear: $a_i(t) = a_i^* + (a_i(0) - a_i^*) e^{(-\kappa D t / (\rho A))}$, with $a_i^* = A d_i / D$.
- The Jacobian is $-(\kappa D / (\rho A)) I$ on the simplex, so every eigenvalue is $-\kappa D / (\rho A) < 0$. This meets $\text{Re}(\lambda) \leq 0$ of Ch. 31 §5 strictly.
- At the equilibrium every organ has the same deficit $e^* = D/(\rho A) - 1$. None is starved while another idles.
- `replicator_step` integrates it exactly, in closed form, for any step size.

P4. Bounds (floors, ceilings, the budget).

- Allocations are projected onto $\{l_i \leq a_i \leq h_i, \Sigma a_i \leq A\}$.
- The projection is water-filling: organs at a bound keep it, and the rest share the remainder in proportion.
- Budget an organ cannot use (it is at its need or its ceiling) is not spent.

P5. Reserve and surge.

- The dual-bath rate channel (the rate tracker) projects the rise in need over one actuation delay: $R = \sum_i \max(0, d_i)$, capped at $A/2$. R is held back.
- After the flow, R is released continuously to organs whose deficit stays positive, in proportion to their shortfall and up to their ceilings.
- This is the controlled release of Ch. 30: continuous state, no switch.

4. Evidence

Mathematics in code. `tests/test_conveyance.py` (in `verify.py`) checks 20,000 random systems:

- P1 conservation to $1e-9$;
- P2, F never decreases;
- P3, the trajectory equals the closed form;
- P4, bounds and budget.

It also runs 2,000 systems over 30 steps and checks the full law never exceeds the budget or breaks a floor. That test found and fixed one real bug: an organ whose need fell below its device floor was given a ceiling under that floor.

Simulation. `hardware/site_exchange.py`: four GPU groups of 16 H100-class GPUs share one site budget (MLPerf gamma). Held-out seed 929292, 48 scenarios per budget. Results: `results/hardware/SITE_EXCHANGE_HELDOUT_*.json`.

Site budget (share of TDP)	70%	60%	50%
Site-budget violation minutes, native / independent / conveyance	10.5 / 1.5 / 0	45.6 / 10.5 / 0	106.1 / 31.3 / 0
Energy (kWh), native / static split / independent / conveyance	142.5 / 117.6 / 104.1 / 103.8	139.1 / 107.9 / 103.5 / 102.0	127.3 / 97.1 / 101.0 / 96.4
Backlog minutes, static split (the arm that also keeps the budget) / conveyance	32.0 / 1.0	89.7 / 8.9	187.3 / 72.2

Reading.

- **Keeping the budget.** Conveyance is the only dynamic arm that never exceeds the site budget.
- **Energy.** It uses the least energy at every budget.
- **Against the only other arm that keeps the budget (static split):** 31-115 fewer backlog minutes.
- **At tight budgets, native and independent sizing show less backlog,** but only by drawing power over the site limit for 10-106 minutes, which in a facility trips breakers.
- **When the budget is below total need,** conveyance shares the shortage in proportion to need. That is the unique stable point of P3: fair, and not a choice of weights.

CPU and GPU on one budget. `hardware/node_exchange.py`: the same law with CPU organs beside the GPU organs. Four groups, each two 8-GPU servers (16 H100-class GPUs, 4 CPU sockets); each unit of GPU work needs 0.35 of the CPUs at full clock to feed it; demand 1.6x the plant traces (the GPUs are held back by the budget, not by work). Arms:

- **S** today's practice: every GPU at one fixed cap low enough that the site fits even with every CPU at maximum;
- **XM** conveyance among the GPUs, counting the CPUs' last measured draw;
- **XC** conveyance over CPU and GPU organs together: each CPU group is held to the clock its feeding work needs, and the watts it no longer holds flow to the GPU groups in deficit.

Seed 515151, 24 scenarios. Results: `results/hardware/NODE_EXCHANGE_*.json`. XC against S, paired:

Site budget (share of GPU TDP + CPU maximum)	60%	70%	80%
Work served	+5.7%	+3.6%	+1.4%
Backlog minutes	-15.8%	-27.5%	-17.8%
95th-percentile latency factor	-39.6%	-43.6%	-35.2%
Site-budget violation minutes, native / S / XM / XC	195.8 / 0 / 3.9 / 0	140.9 / 0 / 6.2 / 0	61.8 / 0 / 5.3 / 0

Reading.

- **More work from the same building.** The tighter the budget, the more the CPUs' unused watts are worth to the GPUs.
- **Holding the CPUs is what makes the hand-over safe.** Counting the CPUs' measured draw (XM) serves about as much, but goes over the budget for 4-6 minutes in every setting: a CPU can rise between one reading and the next. XC holds each CPU to its allocation, so the watts it gives up are really free, and it never goes over.
- **Most of the gain is the CPUs' reserve, not their clock.** Against XM, XC gains ~1% work per kWh; against S, the watts a fixed plan must keep for CPUs that might peak are what the GPUs receive.
- The same holds for the least and most favourable MLPerf gamma fits, and for CPU shares 0.2 and 0.6 (+2.3% to +4.2% work at the 70% budget), `tests/test_node_exchange.py` checks N1-N4 in `verify.py`.

5. What is not claimed

- This is a simulation on declared device physics.
- The live levers that would carry it are GPU power limits (`nvidia-smi -pl`, DCGM), RAPL package limits, and pod CPU limits under a namespace budget. They exist in `omni_controller/muscles.py`, but the exchange between them has not run on hardware.

Evaluation and simulation use only. Copyright (c) 2026 The Omni-Compass LLC. Commercial use, commercialization or monetization of any part of Omni-Compass requires a signed, paid Omni-Compass Enterprise License from The Omni-Compass LLC. Patent applications, copyright registrations and trademark applications filed in the United States. See LICENSE and NOTICE at the root of this repository.

CHAPTER 13

Formal Status of the Mathematics

Not a global-stability proof of the forced six-state system. Proofs and their evidence classes: docs/TRACKING_THEOREM.md; every statement's class: docs/EVIDENCE_LEDGER.md.

Held:

- Isolated S-flow is the gradient of Φ (same formula in both cores).
- (T) Unsaturated: $dU/dt = KP(\sigma - U)$; $V = e^2/2$ has $dV/dt = -KP e^2$ on that channel (Theorem 1).
- (T) Saturated, with the U drift bounded by $F_{\text{bar}} < U_{\text{AUTHORITY}}$: V decreases wherever $e \neq 0$; over the declared parameter box $F_{\text{bar}} = 16.86 < 25$, so this holds over the box in continuous time (Theorem 2).
- (T, constant computed) Sampled-data RK4 controller: $e_{(k+1)} = 0.88 e_k + d_k$; with ϵ_h computed over the frozen fixtures (V), the ultimate bound 0.0315 lies inside the 0.10 basin (Theorem 3).
- (T continuous, V discrete) Admissible box forward invariant (Theorem 4).

Open:

- ϵ_h and discrete invariance proved over the whole box (interval arithmetic).
- Inheritance embedding / intertwining residual against a real plant (Closed Structure Def. 11.1): instrumented in the GPU bench (tools/gpu_reps.py, representation fidelity), not yet measured on a card.
- Mapping of monograph Theorem 5.6 onto (E, U, I_U, S, B) .
- A tracking bound for printed_eight_line (logistic U drift).

Default core remains symmetric_verified (mechanism id in results/MECHANISM_IDENTITY.json).

Evaluation and simulation use only. Copyright (c) 2026 The Omni-Compass LLC. Commercial use, commercialization or monetization of any part of Omni-Compass requires a signed, paid Omni-Compass Enterprise License from The Omni-Compass LLC. Patent applications, copyright registrations and trademark applications filed in the United States. See LICENSE and NOTICE at the root of this repository.

CHAPTER 14

The Mechanism of Action

Every statement here is tied to code in this repository, and the measured numbers are reproduced by python tools/mechanism.py, which writes results/MECHANISM_OF_ACTION.json. The engine files are frozen. Their SHA-256 hashes are checked by verify.py.

1. The state: six numbers that describe the whole system

Symbol	Name	What it stands for in a data centre
E	energy / excitation	how hard the system is being driven: queues, overload, power, heat
U	order parameter	how healthy and settled the system is: +1 healthy basin, -1 failed basin
I_U	integrated need	unmet need accumulated over time: memory of pressure that has not been relieved
S	stress	slow structural stress from heat, power, network, security, staleness
B, B_dot	bath and its rate	a damped oscillator driven by stress: the slow "tide" of physical load

2. The equations (omnicompass/core.py, lines 5-12)

```
(1) dE/dt = -alpha_E E + beta_int + beta_ext + v_eff
(2) dU/dt = mu U (1 - U^2) - (dE/dt)/E_max - lambda_U U + u, |u| <= 25
(3) dI_U/dt = (1 - U) - sigma_1 E - delta S - lambda_I I_U
(4) v_eff = cos(omega_B t / 2) * c * tanh(lambda_0 + lambda_1 (U - 0.5) + lambda_2 S)
(5) Phi(S) = alpha_s S^2 / 2 + beta_s S^3 / 4 - delta S
(6) dS/dt = -dPhi/dS = delta - alpha_s S - (3/4) beta_s S^2
(7) dB/dt = B_dot ; dB_dot/dt = gamma_c delta S - (omega_B / Q_B) B_dot - omega_B^2 B
(8) R_B[n] = finite-difference audit of (7), never fed back
```

Each term does one job:

- **(1) Energy.** Energy decays at rate $\alpha_E = 3.487$. It is driven by two forcings: internal pressure β_{int} (queue and overload) and external pressure β_{ext} (power, heat, network, staleness, security). The drive term v_{eff} (4) is bounded by c through the hyperbolic tangent, so no input can drive E without limit. This is the saturation (the tanh "speed limit"): however extreme the observation, the engine's response rate is capped at $c = 1$.
- **(2) Order.** $\mu U (1 - U^2)$ is a double-well (bistable) drift with two stable basins, $U = +1$ (healthy) and $U = -1$ (failed), and an unstable ridge at $U = 0$. A rise in energy (dE/dt) pushes U out of the healthy basin. u is the bounded control input.
- **(3) Integrated need.** I_U grows while U is below 1 (the system is not settled). It is reduced by energy and stress already being dealt with, and it leaks at $\lambda_I = 0.719$. It is the engine's memory of unrelieved pressure.

- **(5)-(6) Stress.** A cubic potential (the "well in the basin"). S rolls downhill to its stable root. This is the slowest restoring force.
- **(7) Bath.** A damped harmonic oscillator (natural frequency $\omega_B = 1$, quality $Q_B = 2.5$) driven by stress. It turns stress into a smooth, lagged "tide". It is the only second-order element.
- **(8) Audit.** The bath equation is re-checked numerically every step. The check is never fed back.

Integration: classical 4th-order Runge-Kutta, 10 micro-steps of 0.01 per macro step (rk4_step, macro_step). The C++ twin matches the Python to 3.6e-15 on 500 reference trajectories.

3. From telemetry to state (omnicompass/adapt.py, observe_vector, assimilate)

Each decision the telemetry is normalised to nine observations: q (queue), load, power, thermal, network, drift, stale, security, conflict. An observed state is computed from them:

```
E_obs  = 0.25 q + 0.18 max(0, load - 0.85) + 0.16 power + 0.13 thermal + 0.10 network + 0.08 drift + 0.
↳ 06 conflict + 0.04 stale
U_obs  = 1 - (0.27 q + 0.18 power + 0.16 thermal + 0.12 network + 0.12 drift + 0.10 conflict + 0.05 sta
↳ le)
S_obs  = 0.36 thermal + 0.28 power + 0.18 network + 0.10 security + 0.08 stale - 0.20 q
B_obs  = 0.52 power + 0.26 thermal + 0.22 drift
I_obs  = q + drift + conflict
beta_int = 0.58 q + 0.42 max(0, load - 0.75)      beta_ext = 0.33 power + 0.27 thermal + 0.20 network
↳ + 0.12 stale + 0.08 security
```

The state is blended toward the observed state with weight $a = 0.339$: $x \leftarrow (1 - a) x + a x_{\text{obs}}$. It is then evolved by one macro step (0.1 engine time units) of equations (1)-(7).

4. From state to action (the allocation law, Governor.step)

```
push      = u(x) / 25, u(x) = clip(-f_U(x) + 12 (1 - U), -25, 25)      equation (2)'s controller, eval
↳ uated, not applied
rho*      = clamp(rho0 - kI I_U - kE E, rho_min, rho0)                target utilisation (the HPA t
↳ arget)
n_req     = ceil(n load / rho*) + ceil(kq q n)                        machines wanted
release   only while push <= push_release (the engine reports convergence), after band and dwell
cap       = clamp(cap_now load (1 + margin), cap_min, 1) when calm and U >= U_gate; 1 under backlog;
           always <= 1 - cap_gain max(0, B - B_cap)                  the bath limits the power ca
↳ p
change    permitted only while U >= U_gate and no security block
rollback  authorised if S >= S_rollback, U is falling, or a security block is seen
```

5. Which equation drives which muscle

Muscle	Engine quantity used	Where
HPA target (pods)	ρ^* from I_U and E, equations (1) and (3)	adapter <code>Governor.step</code> ; live <code>omni_controller/controller.py</code>
Node pool (machines)	n_{req} from load, queue and ρ^* ; release gated by push from equation (2)	same
Power cap (pod CPU limits, CPU frequency, GPU limit)	cap from U gate and bath B, equation (7)	same; <code>omni_controller/muscles.py</code>
Rollout guard	<code>change_permitted</code> (U), <code>rollback</code> (S, dU)	<code>muscles.py</code>
Security hold	security observation raises S and blocks expansion	adapter, shield
Speed-first law	ρ^* for pods; push gates machine release	<code>omnicompass/speed.py</code>
Right-sizing	ρ^* sizes CPU; E widens the memory margin; push gates scale-in	<code>omnilab/rightsize.py</code>
Cold start	ρ^* sizes instances; I_U sets how long to keep warm	<code>omnilab/coldstart.py</code>
GPU packing	I_U pre-warms nodes; push gates power-off	<code>omnilab/gpupack.py</code>
Training power	E tightens the applied ramp limit under sustained grid stress	<code>omnilab/powersmooth.py</code>
GPU health	S and E shorten the checkpoint interval	<code>omnilab/health.py</code>
Cooling	E widens the inlet-temperature margin	<code>omnilab/cooling.py</code>
Inference	the full allocation law (<code>node_delta</code>) sizes replicas	<code>omnilab/inference.py</code>
Agent containment	I_U throttles; the security observation contains	<code>omnilab/containment.py</code>
VM energy	I_U and E raise the learning rate when attribution goes stale	<code>omnilab/vmenergy.py</code>

6. Measured: what the equations actually do inside the governor

Measured on a real observation stream (fleet plant, web services, fleet mode, 360 decisions):

Mode (eigenvalue of the Jacobian at the operating point)	Time constant	In decisions
-7.90	0.13	1.3
-3.10	0.32	3.2
-0.72	1.39	13.9
-0.20 +/- 0.98i (the bath)	5.0, period 6.4	50, rings every 64 decisions
-0.17	6.0	60

- **Every mode is stable** (all real parts negative). The operating point is a stable equilibrium inside the healthy basin: $E = 0.095$, $U = 0.976$, $S = 0.315$, $B = 0.10$. This is the basin keeping everything in, stated as mathematics.
- **Per decision, the equations and the observation move the state about equally.** Across all six variables the equations contribute 45-50% of each step's movement and the observation blend the rest.
- **Equation (2)'s control input u is not applied during governance.** The engine evolves open-loop ($u = 0$). The control law is evaluated only to produce push, which decides when capacity may be released.
- **Compute:** about 95 microseconds per decision in Python, and about 2.9 microseconds in the C++ twin.

7. What this explains, and what it points to

- **Why "engine off" scores close to "engine on" in the muscle benchmarks.** With the equations frozen, the observation blend still moves the state. Both versions are then low-pass filters of the same telemetry. The engine's fast modes (1-3 decisions) behave almost like the blend itself. Its distinct contribution is its slow memory: I_U (14 decisions), S (60) and the bath (period 64 decisions). The muscle mappings mostly read E and ρ^* , which live on the fast modes.
- **The engine clock is not matched to the workload's rhythm.** One decision advances engine time by 0.1. At 60-second decisions the bath rings every 64 minutes, but web demand has a 24-hour rhythm and training power a 2-second one. This is a wiring choice in the adapter (engine time per decision), not a property of the equations.
- **The controller in equation (2) is unused as an actuator.** Wiring its output u (or U 's trajectory under u) directly to the actuators is the unexplored way to "let the math do the work".

Three wiring experiments follow directly. None changes a single equation:

1. **Clock matching.** Set engine time per decision so that the bath's natural period equals the workload's rhythm. The oscillator in (7) then carries the cycle and can anticipate it.
2. **Slow-mode actuation.** Drive machines from I_U, S and B (the memory modes) rather than E (a fast mode).
3. **Closed-loop actuation.** Apply equation (2)'s u and use its magnitude, not only its sign, as the actuator command.

Each is testable in the existing harness against every platform in the league (tuning/league.py), and will be tested there.

The pedals: idle, gas, brake, reset, and the kill switch

Omni-Compass drives a system the way a self-driving car drives itself: it feels, gauges and adjusts, with nobody in the seat. The words below are used the same way everywhere in this repository.

Two names only. **native** is the system as it runs on its own (Kubernetes and its autoscalers, the card's firmware, CityLearn's controller). **omni** is Omni-Compass on top of native: the same native system, with Omni-Compass governing it. Omni-Compass never runs instead of native, only on top of it, so every comparison is native against omni.

The modes. Off (manual): Omni-Compass not driving, native runs alone. Watching: Omni-Compass reads every gauge and logs what it would do, and writes nothing. Autopilot: the pedals below. Cruise and the emergency brake: for work that comes as a pile.

Term	Meaning	In the code
Idle	No foot on the gas or the brake: calm traffic, the engine running at its floor (two machines in service, ready)	MIN_NODES, the floor of the staging law (8.8)
Gas	Ramping up: traffic climbs and capacity is added at once, never held back; the whole cluster can take work the instant it is needed	(8.3); a lower HPA target is never limited
Brake	Ramping down: traffic falls and capacity is eased off a step at a time, never below idle	(8.4)-(8.5); coasting, rule 6
Reset	The brake held to the floor: every setting the governor ever wrote is handed back to where native had it, read back, and the record removed; the next run starts fresh. Every benchmark arm ends with one and checks it	the reset file (- -kill - file, a name kept for compatibility), Controller.restore

Term	Meaning	In the code
Cruise	A pile of work waiting for a place: every machine in service, held there without second-guessing, until the line has been empty for two decisions	rule 7, --cruise-after
Emergency brake	The work is done and demand is at zero (the service at its least pods, wanting no more, at half its target or less): straight to idle (the floor) in one move, every safety check still holding	rule 8, --brake-demand, amendment 9
Kill switch	Security only: one switch in a human hand that turns Omni-Compass's governing off across the whole system at once, for anything rogue or anyone trying to drive a system through Omni-Compass's brain	omniconpass/master.py, tools/omni_switch.py off

The fuel cut-off point (the floor) is a number the operator sets; the brake never takes the system below it, however hard it is pressed. Earlier dated records, and the sealed original engine files, call the reset a "kill switch"; they are kept word for word.

8. The staging law: machines on and off in order, with demand that moves

Demand climbs, spikes, eases part way, climbs again and settles to idle. The staging law decides how many machines are in service at every decision, and which ones, so that capacity follows demand up at once and comes back down only as far as it safely can. It is the compass law (omniconpass/compass_law.py) applied to the machine pool (omni_controller/controller.py), the release gate (omniconpass/nervous_system.py node_release_gate), the verdict (omniconpass/verdict.py) and the actuator (scripts/kind_nodepool.sh). Every symbol below is a line of that code.

Position and force. The service reading r (the 95th-percentile response time, or the pressure that stands for it) is placed in its band $[lo, hi]$, where hi is the response line:

$$p = (r - lo) / (hi - lo), \quad v = s v + (1 - s) (p - p_{prev}) / dt \quad (8.1)$$

$$F = A \quad \text{if } p \geq 1 - c \quad (\text{fail up, past the wall})$$

$$F = A \tanh((kp(p - p_0) + kd v) / A) \quad \text{otherwise; } F = 0 \text{ if } F \leq 0 \text{ while } v \geq 0 \text{ and } p \geq p_0 \quad (8.2)$$

with the band's center $p_0 = 0.5$, cushion $c = 0.05$, authority $A = 1$, $kp = 1$ and $kd = \max(0, 2 \sqrt{kp \tau} / dt) - 1$ (critical damping for a muscle that answers in τ seconds). $F > 0$ asks for capacity, $F < 0$ offers it back.

How many machines. With n machines in service now, n_{nat} the count native ran, m the floor (two by default, `MIN_NODES`) and M the machines that exist (every machine usable; an operator may hold whole machines back with `NODE_CUSHION`, off by default), every count stays in $m \leq n' \leq M$. The protection is not a machine sitting out: it is the band inside every machine. Capacity is added at 95% of the response line, before the line is reached (8.2-8.3), and a machine goes back only if the ones left still run at or under the engine's utilisation target ρ (8.5): a margin spread over all of them.

$$\begin{aligned} \text{up} \quad n' &= \min(M, \max(m, n + 1)) & \text{if } p \geq 1 - c & \quad (8.3) \\ \text{down} \quad n' &= \max(m, n - 1) & \text{if } F \leq R \text{ and } p \leq p_0 \text{ and the verdict} & \\ & & \text{allows } n_{nat} - (n - 1) \text{ and the gate } G \text{ ho} & \\ \quad \hookrightarrow \text{lds} \quad (8.4) & & & \\ \text{hold} \quad n' &= n & \text{otherwise} & \end{aligned}$$

with the release threshold $R = -0.2$. Up is immediate and is never gated. Down is one machine per decision, never below the floor, and only when every term of the release gate holds:

$$G = [n \geq 1] \text{ and } [\text{nothing pending}] \text{ and } [\text{pods not scaling up}] \text{ and } [\text{no breach now}] \text{ and } \\ [\text{used} / ((n - 1) \times \text{per_node}) \leq \rho] \text{ and } [\text{contraction authority}] \text{ and } [\text{every sense live}] \text{ and } \\ [\text{the last machine command landed}] \quad (8.5)$$

rho is the engine's utilisation target: after a machine goes back, what remains runs at or under rho. The verdict (section 7 of the GPU preregistration, `omnicompass/verdict.py`) allows a machine count below native only after a paired trial showed the service at that count at most 2% worse than at the count before; that 2% is where the trigger sits, inside the engine.

This is hysteresis by construction. The up condition ($p \geq 0.95$) and the down condition ($p < 0.5$ with $F < -0.2$) are far apart, so a demand that hovers never makes machines flap; a spike crosses the wall and adds a machine at once; a dip has to be deep, calm and sustained before one goes back, and then only one per decision. Between the two, the count holds.

Which machines. The actuator orders the pool:

```
idle    the open machine with the fewest serving pods, then the fewest pods, goes first;
        it is marked prefer-not (PreferNoSchedule) and its pods first to go (pod-deletion-cost), so the
        autoscaler's own scale-down empties whole machines; no pod is moved, evicted or restarted (8.6)
wake    the idle machine still carrying the most work goes first (it is warm); the mark comes off and it
    ↪ is in
        service at once, with no boot (8.7)
floor   two machines always in service (MIN_NODES, the operator may set more), ready for a spike (8.8)
margin  every machine usable; the cushion is the band inside each one (the 5% before the line, and rho),
    ↪ not a
        whole machine held out (8.9)
```

An idled machine stays powered and Ready at its floor ($\text{park_frac} \times \text{idle power}$), never off: a pod that finds the open machines full lands on it at once, so no request ever waits on a machine Omni-Compass idled. Where a node autoscaler underneath really removes an empty machine (Azure's cluster autoscaler, Karpenter), that same order hands it the emptiest machine first, and the bill falls with the machine.

The same law on every discrete unit. Machines are one case. In the muscle models (`realms/compass_arm.py` `release_safe`), a cooling plant's chillers and a compute pool's machines follow the same law: one unit back only when the units left cover the recent peak with headroom (`RELEASE_MARGIN` 0.6 for machines, 0.8 of a unit's capacity for chillers), one at a time, and back at once when the wall is reached.

What is proved, measured and open. The order (8.6)-(8.8) is checked through the real actuator by `tests/test_staging_order.py` (the emptiest idles first, the warmest wakes first, two always in service, every machine usable, no pod moved). The gate (8.5) is checked by `tests/test_node_release_gate.py`. On real Kubernetes the law gave the same work 55-65% faster on 29-36% fewer machines (sets 22-27) and 48% more work on the same machines (the capacity test); its run under demand that wanders (up, spike, part way down, back up, idle) is in progress. Open: an order across different kinds of unit (a battery before a generator, a CPU before a GPU) is not wired; each pool is ordered within itself.

9. The state each muscle carries, and the modes as a state machine

9.1 The compass on a muscle is a dynamical system, so its state must persist. On muscle i the compass carries $z_i = (p_i, v_i, x_i)$: the position of the service in its band, its velocity, and the knob's continuous value. At each decision k , with reading r_k :

```
p_k = (r_k - lo) / (hi - lo),    v_k = s v_{k-1} + (1 - s)(p_k - p_{k-1}) / dt,    F_k = F(p_k, v_k) by (8
    ↪ .2) (9.1)
x_{k+1} = clamp( x_k + sigma g(F_k) F_k (x_hi - x_lo), x_lo, x_hi ),    g = UP if F_k > 0 else DOWN
    ↪ (9.2)
```

with x_0 the native value and $\text{sigma} = +1$ when a higher value is more capacity. Two properties depend on z persisting: (a) the knob integrates the force, $x_k - x_0 = \sum_j \text{sigma } g F_j \text{ span}$ (until the cover binds), so a calm

muscle eases step by step toward its cover; and (b) the damping term $kd v_k$ needs p_{k-1} . If z were made fresh at every decision ($p_{k-1} = p_k$, $v = 0$, $x_k = x_0$), (9.2) would collapse to $x_{k+1} = \text{clamp}(x_0 + \sigma g F(p_k, 0))$: a memoryless map of the present reading, at most one step from native, with no damping. That is a different, weaker law. A renamed attribute did exactly that on 2026-10-05; `tests/test_compass_arm.py` now fails if the compass is ever made twice on one muscle.

9.2 The modes. The live controller is in exactly one mode at each decision. With q_k the pods waiting for a place, u_k the scaling-up flag, c the cruise count (`--cruise-after, 2`), M the machine ceiling and m the floor:

cruise on $q_j \geq 0$ for the last c decisions (9.3)

cruise off $q_j = 0$ and not u_j for the last c decisions (9.4)

in cruise $n' = M$; the service's autoscaler target held at the operator's own; the floor step makes no read and no write (9.5)

emergency brake (not cruising) $q = 0$, no breach, not u , every sensed service idle or at demand $\<= b$, the gate G (8.5) holds, every sense live, the last command landed:

$n' = \max(m, \text{smallest } n \text{ carrying the CPU in use at } \rho)$ (9.6)

idle (a service at its floor) $I(h) := \text{current} = \text{min} = \text{desired}$, and $u_h \leq \text{target}_h / 2$ (9.7)

gas / brake / idle (8.3) / (8.4)-(8.5) / (8.8)

9.3 Why cruise steps back (9.5). In cruise every machine is already in service, $n = M$, so the floor step's only action (raise n to the floor the waiting pods need) is empty: $\text{floor} \leq M = n$. Its reads (every pod, every node, node metrics, every HPA, each five seconds) carry no value of information, so removing them cannot change any action; they only cost CPU and API-server time on the machines the queue is using. The service's response time is also confounded in cruise: with a queue of other work holding the machines, a request waits on contended CPU, $R = S / (1 - u_{\text{host}})$, not on the service's own pods. Reading that R as the service's own queue and adding pods would answer the wrong cause. So in cruise Omni holds the machines and leaves the service to its own autoscaler. `tests/test_compass_controller.py` checks that the floor step makes no API read during cruise; the controller before this rule made the full set of reads every five seconds.

9.4 Why idle is read from the autoscaler's floor (9.7). A served service never reads zero: a health probe or a client's keep-alive keeps one pod at a few percent of its request. In the batch test, one pod at about 10% of its request read as demand 0.10, above $b = 0.05$, and the brake waited 15 minutes after the queue emptied. The autoscaler's own floor is the observable of "no demand beyond the minimum": the pods it keeps are its minimum, it wants no more, and they run under half the target. In the five frozen-engine web tests the service was never at its floor (0 of 5,102 readings), so (9.7) changes nothing there.

9.5 Reporting resolution. A paired change smaller than one part in a million of the value ($|\text{mean difference}| \leq 1e-6 |\text{native}|$) is read "same", the number still shown. Below that, a difference is rounding in floating point and in the models' accumulators, not behaviour; a deterministic run gives a zero-width interval around it, which would otherwise read as proven.

9.6 A lever is moved only where moving it can pay. The compass decides how far and how fast. Whether a lever may move at all is a property of the muscle's physics. Four cases, each derived from the muscle's own model and each its own test:

- *Cooling power.* A zone's heat must leave it: over a run, the heat removed equals the heat made plus what the room's mass stores. A cap on cooling power q_{max} cannot remove less heat. It only lets the temperature drift above setpoint while the PI command grows, and the staging rule, $\text{units} = \text{ceil}(\text{command} / (0.8 q_{\text{unit}}))$, turns that command into more units and more fan power, $p = q_c / \text{COP} + \text{units} \times p_{\text{unit}}$. So a lower cap never saves energy, and the lever stays native. (The setpoint, which changes the COP, and the units, which change fan power at the same heat, remain levers.)

- *A battery's reserve.* Energy D drawn from the cells below the operator's reserve saves $D \sqrt{\eta}$ of grid import, and the deficit is bought back at $D / \sqrt{\eta}$. The net, $D (1 - \eta) / \sqrt{\eta}$ (10.5% of D at $\eta = 0.90$), is always positive. Spending the reserve is never an energy saving; it is the price of keeping the draw under the connection's limit. So the reserve is spent only at the wall ($p \geq 0.95$) and only when the battery can cover the excess ($0 < \text{imp} - p_{\text{lim}} \leq p_{\text{batt}}$). Otherwise the operator's reserve holds, and no round trip is paid for nothing.
- *A backup reserve (a UPS).* It is held for an outage, which the model does not contain, so spending it for a peak trades away the reason it exists. Never a lever.
- *A compute pool's power cap.* With throughput $\mu(c)$ proportional to c^{ϵ} and the autoscaler holding utilisation near its target u^* , the energy per unit of work is $e(c) = (p_{\text{idle}} / u^* + p_{\text{dyn}} c) / \mu(c)$. At $c = 1$, $de/dc > 0$ (a lower cap saves) if and only if

$$\epsilon < p_{\text{dyn}} / (p_{\text{idle}} / u^* + p_{\text{dyn}})$$

↪ .8)

(9)

At $\epsilon = 0.4$ the cap saves on CPU hosts (threshold 0.62), GPUs (0.74), databases (0.51) and radio units (0.54). It costs energy where idle power dominates: a quantum computer's cryostat (0.19), a network switch (0.32), storage (0.26). There, running slower keeps the machines on longer: race to idle. In a deployment p_{idle} , p_{dyn} and the speed curve come from the machine's own specification.

- *A motion axis's effort cap.* It lowers the copper loss of a full-acceleration move, $(J a_{\text{max}} / kt)^2 R$, and stretches the move, which keeps paying the standing draw $p_{\text{idle}} + b v_{\text{max}}^2$. It moves only where the first exceeds the second: a reaction wheel (32 W against 5 W), not EV traction, a flight axis or a robot joint.
- *A machine released from a node pool.* A machine boots in minutes, and a burst that arrives in the meantime is served late. So a machine goes back only when the machines left cover the recent peak at 0.3 of the pool's own release level, the largest margin at which no pool is later than native (`results/realms/RELEASE_MARGIN_SWEEP.md`). This is the verdict of the live controller (section 8) in model form.

What a lever buys is labelled for what it is. Less work per energy, with no work lost and the time over the line proven lower, is SERVICE IMPROVEMENT WITH ENERGY TRADEOFF, the mirror of ENERGY IMPROVEMENT WITH SERVICE TRADEOFF. Without that proof it is WORSE.

Evaluation and simulation use only. Copyright (c) 2026 The Omni-Compass LLC. Commercial use, commercialization or monetization of any part of Omni-Compass requires a signed, paid Omni-Compass Enterprise License from The Omni-Compass LLC. Patent applications, copyright registrations and trademark applications filed in the United States. See LICENSE and NOTICE at the root of this repository.

CHAPTER 15

The Engines and Their Audit

Which engine is which

There is **one** engine. Every other engine-looking file in this repository is either its required C++ twin or a frozen copy of where it came from. Nothing else runs.

File	Role	Runs?
omniconpass/core.py	The engine (<code>symmetric_verified</code>): six states, eight lines, RK4 with microsteps. Everything Omni does goes through it.	yes
cpp/src/core.cpp, cpp/include/omniconpass/core.hpp	The same engine in C++. Required: <code>verify.py</code> checks it against the Python on every fixture.	yes (twin)
docs/handoff/mathematics/omni_compass_engine_source_c527df2d.py	The original source as handed over (sha256 c527df2d...). Kept unchanged as proof of origin and as the parameter authority.	no (frozen)
reference/omni_compass_reference_engine.py	The same original with comments stripped (<code>tools/strip_reference.py</code>); same program fingerprint (<code>reference/PROVENANCE.json</code>).	no (frozen)

Fingerprints of the running engine, its parts and its twin: `results/MECHANISM_IDENTITY.json` and `RELEASE_MANIFEST.json`. `verify.py` fails if any of them changes without a new seal.

Engines from outside packages, not adopted

Copies of this repository passed around as zips (the XPASS packages) carry changes that are **not** in the engine here:

- `batch_claim1.py` (in those zips, not in this tree): a vectorised variant that applies `u` during the RK4 step (CLAIM1). No C++ twin; not verified against the 500 fixtures.
- an adapter that defaults to CLAIM1, takes the target sign from the starting state and adds a disruption budget. Its C++ governor was only partly ported, which is where the reported Python/C++ mismatches came from.

None of these enter before the GPU confirmation run (the engine is frozen for it, `docs/GPU_PREREGISTRATION.md`). Any of them can be adopted afterwards only as a full change: Python and C++ together, the parity tests and `verify.py` green, and a new mechanism id.

The rule from here

One engine, one twin, two frozen originals. A new version **replaces** the old one in place, with a new seal; it is never added beside it. Older states are in git history and `docs/HISTORY.md`, not in extra files.

Evaluation and simulation use only. Copyright (c) 2026 The Omni-Compass LLC. Commercial use, commercialization or monetization of any part of Omni-Compass requires a signed, paid Omni-Compass Enterprise License from The Omni-Compass LLC. Patent applications, copyright registrations and trademark

applications filed in the United States. See LICENSE and NOTICE at the root of this repository.

My engine, line by line: the source against what runs in your cluster

The source. docs/handoff/mathematics/omni_compass_engine_source_c527df2d.py (SHA-256 c527df2d..., 31,348 lines). Every mechanism below is named by its function there, followed by where it runs live and how it is checked on every python verify.py.

Mechanism	In the source	Where it runs	Checked by	Status
Six states E, U, I_U, S, B, Ḃ and equations (1)–(7)	oc_derivatives, rk4_step	omnicompass/core.py	tests/test_core_parity.py: 20,000 random states, max difference $\leq 1e-12$	exact
Spinor closure, 720° : drive $\times \cos(\frac{1}{2} \omega_B t)$	v_eff, spinor_closure_factor	core.v_eff	same parity test	exact
Axle rest S^* from equation (6)	stable_S_equilibrium	omnicompass/storage.py, compass.axle	parity; tests/test_compass.py	exact
Double well $W(U)$, basins ± 1	W_potential, W_gradient_flow	core.derivatives	parity	exact
Regulation $u = -f_U + K_P(\sigma - U)$, clipped at the authority	hybrid_microstep_operator	core.control_command; the governor's push	parity; 500 frozen trajectories (fixtures/), 0 mismatches	exact
Basin lock, conveyance, certification	run_monte_carlo, macro_step	core.simulate	500 frozen trajectories, 0 mismatches	exact
Composite storage $V = V_U + V_W + V_E + V_S + V_I + V_B$, the ledger that closes the circle	composite_practical_lyapunov_value, bath_lyapunov_matrix, weights	omnicompass/storage.py, read by the compass every decision	parity on 5,000 states ($\leq 1e-13$); descends $648 \rightarrow 8$ over 60 regulated steps with no rise	exact
Human switch: a file or a variable stops every command	_oc_runtime_stop_requested	--kill-file, restore of every lever	tests/test_failsafe.py, tests/test_muscles.py, every live run's switch drill	exact
Cap doctrine C1: only a ceiling above the reference is admissible; a lower one buys energy by slowing work	_ci_omni_decision, certificate C1–C4	muscles.convey: CPU limit never below the operator's, idle CPU conveyed to the work	tests/test_convey.py	exact
Sleep only in a certified empty interval	_ci_omni_decision, sleep channel	scripts/kind_nodepool.sh: a machine idles only once its work has left	tests/test_active_nodes.py, tests/test_convey.py	exact
Reference Kubernetes HPA rule	oc_hpa_desired_replicas	cpp/ HPA, pod reflex	tests/test_cpp_hpa_parity.py	exact

Where my live wiring reads the engine differently from the source's own embodiment

Evolution between decisions. The source's compute embodiment advances the assimilated state *with* regulation locked on basin +1. My live governor advances it with $u = 0$ and reads the regulation command separately, as the push that gates every machine release.

- Both use the same equations.
- The difference is whether the regulation is applied to the state or only read from it.
- The pre-registered, frozen results were measured with $u = 0$. I keep that setting, and I report the difference here rather than change a frozen law mid-measurement.

Evaluation and simulation use only. Copyright (c) 2026 The Omni-Compass LLC. Commercial use, commercialization or monetization of any part of Omni-Compass requires a signed, paid Omni-Compass Enterprise License from The Omni-Compass LLC. Patent applications, copyright registrations and trademark applications filed in the United States. See LICENSE and NOTICE at the root of this repository.

PART THREE

The Physics: the Compass and the Nervous System

Push and pull, the band and its cushions, the physics of a processor, and the two-way wires that carry the force from the brain to every muscle and back.

CHAPTER 16

The Compass: Push, Pull and the Two Forces

The compass (`omnicompass/compass Law.py`) is the law that carries the engine's push and pull to every muscle.

Two bands.

- **The cover**, a knob's hard range (the device's or the operator's lowest and highest setting). Every write is clipped to it. Nothing Omni-Compass computes can set a knob outside it.
- **The compass**, the service reading as a position from 0 (calm) to 1 (the service line). Omni-Compass pulls that position to the bottom of the compass, the middle (0.5 by default; an operator may set it lower for extra margin). The walls run from 0.05 to 0.95; the last 5% on each side, 10% in all, is cushion.

The force. $F = A \tanh((K_P (p - \text{center}) + K_D v) / A)$, where p is the position, v its rate of change, A the authority.

- **Pull:** $K_P (p - \text{center})$, gentle near the bottom, stronger up the walls.
- **Push:** $K_D v$, the force that meets whatever is shoving the position, and the friction that stops it sloshing. With K_D at or above critical damping the position glides to the middle and stops, with no overshoot and no ringing. Overshoot is wasted energy: force spent going the wrong way and spent again coming back.
- **Smooth, never a hammer:** $\tanh$ bends the force over into its maximum instead of slamming into a wall.
- **Fail up:** past the 0.95 wall the up side goes to its full force at once and the down side may not act until the position is back inside the compass.

Two forces: antagonist pairs. Like the muscles of an arm, every plug has an up side (adds capacity, power, cooling, speed) and a down side (takes it back), each with its own gain, because adding and taking back do not cost the same: a new machine takes minutes to boot; giving one back is instant. Where a machine has two wires, each side gets its own lever:

Muscle	Up force	Down force
GPU card	clock ceiling raised	power limit lowered (the lid)
Kubernetes	scale out	scale in
Cooling	chiller on, colder setpoint	warmer setpoint, unit released
Battery	discharge	charge, reserve held
Vehicle, joint	motor effort	braking, regeneration

Why the GPU needs both wires. A card's firmware boosts its clock as high as it may while there is work, and its power limiter knocks the clock back each time the draw crosses the limit; on a busy card this happens many times a second, at the top of the clock range where each extra step of speed costs the most watts. With one wire (the power limit) a governor can only move the wall the boost pushes against. With two wires, the clock ceiling sets how high the boost may climb and the power limit becomes a lid that rarely needs to act: the card runs at the bottom of its compass instead of fighting itself at the top.

What the card taught us. Two rules make the difference between saving energy and spending the customer's time. First, *race while work waits*: when the card is saturated, both wires go to full at once, so a burst is always served at full speed; the compass paces only the slack between bursts. Second, *never slower than the card on its own*: the governor learns, from the card's own meter, the clock and the draw the card reaches by itself while busy, and never sets the clock ceiling or the lid under them. Without these rules the first real card saved 3.5% of its energy and made the slowest answers 58.5% slower; with them, the modelled card saves energy with the slowest answers at native speed or faster (Chapter 41, and docs/GPU_PREREGISTRATION.md, amendments 6 and 7).

The same law on every stack, and what differs

The law is written once (CompassLaw), and three declared numbers fit it to a muscle: the **decision period** dt (how often the governor decides), the muscle's **response time** τ (how long the service reading takes to follow the lever), and the **cover**. From τ and dt the law derives its own damping gain, $K_D = (2 \sqrt{K_P \tau / dt} - 1) dt$, which is critical damping for a first-order muscle under a proportional-derivative pull: the position glides to the center and stops. Nothing else is tuned per stack. The table below is the complete list of the compass's settings on every stack in this manual, so a referee can see that the "gains" differ only through the declared physics:

Stack	Reading (the wire in)	Line	Center	Decision period	Response time	Cover of the knob
Kubernetes, the HPA target	the mean response time of the latency window	the operator's SLO (500 ms)	0.4	60 s	the pods' start time	[$0.6 \times$ the operator's target, the operator's target]
Kubernetes, the machines	the same, through the release gate and the verdict	the same	0.4	60 s	a machine's wake	[the floor of 2, every machine]
PgBouncer, the pool size	time in the server plus time waiting for one, per transaction	50 ms	0.4	1 s	2 s	[2, PostgreSQL's limit minus 10]
Kafka, the consumer count	the group's own end-to-end latency, mean of the last second	500 ms	0.4	1 s	3 s (a rebalance)	[1, the topic's partitions]
Redis, the memory ceiling	the application's request latency, mean of the last second	2 ms	0.4	1 s	2 s	[16 MB, 512 MB]
MongoDB, the storage-engine cache	the server's own mean read latency, last second	1 ms (set on the tuning workload's smoke run, from 2 ms)	0.4	1 s	2 s	[256 MB, 2,048 MB]
MySQL, the InnoDB buffer pool	the server's own mean statement latency, last second	0.6 ms a statement (set on the tuning workload's smoke runs, from 1 ms)	0.4	1 s	2 s	[128 MB, 2,048 MB], in 128 MB chunks
A drone's cruise override	the drone's tracking error	0.25 m, the declared safe error	0.5	one control tick (48 Hz)	1 s	[1.0, 2.0] $\times$ the planner's cruise
A substation's tap	the bus voltage in its band	the band's edge	0.5	one solve	a week of evidence before a step down	the tap's own range, one tap per move
A robot axis's speed	the axis's tracking error against its takt	the declared error	0.5	one control step	the axis's response	the axis's own limits
A GPU's clock ceiling and lid	the mean response time of the last 5 s	the SLO	0.4	2 s	the card's response	[the card's own busy clock, its top clock]; [the busy draw + 10%, the start limit]

The center is 0.4 wherever the reading is a service time (the "service profile" of the Kubernetes adapter, so that the service sits a little below the middle of its band with margin for a burst) and 0.5 where the reading is a physical error. A **cushion** of ± 0.05 in force moves nothing, so a reading that sits at the center with the usual jitter does not make the knob chatter. These are the same numbers in every preregistration and every table; none was changed after a result was seen.

Direction rules: how a force becomes a move

The force is a number between -1 and $+1$. Each stack turns it into a move on its knob by a direction rule declared in its preregistration, always of the same shape: a positive force adds capacity in proportion to the force ($\text{ceil}(\text{force} / 0.10)$ notches, so full force adds ten notches at once, capped by the cover), a negative force takes one notch back per decision, and only when the resource it would take is demonstrably idle (an idle consumer that read nothing in the last second, an idle server the pooler itself shows idle, a cache with nothing evicted in the last second) and only after a **dwel**l of a few seconds since the last change, so a knob that was just moved is not moved back by its own transient. At 95% of the line the knob **fails up**: every consumer the topic can use, the pooler's own setting handed back, a quarter of the cache's cover added at once, the planner's own cruise, the card's own clocks. The up side is fast and the down side is slow and gated because adding and taking back do not cost the same, and because a gain bought by taking something back must be able to prove itself before it is kept.

The do-no-harm gates inside the law

Three gates sit inside the law itself, each a consequence of a loss seen on a tuning case and each disclosed in the preregistration that carries it:

- **The verdict** (section 1.2): a slow knob moves only where a paired trial on the muscle shows no more than 2% cost.
- **The slack gate on speed knobs** (Omni v3): a motion axis that is busy more than half the time at full speed keeps its speed native, because spending its slack as speed when it has none would only cost energy. This gate is what turned the v2 "service tradeoff" on the Physics realm and the tower into "superior within guardrails" on v3, and it is the one rule that distinguishes v3 from v2.
- **The full-cache gate** (Redis): a miss in a cache with room to spare is a cold miss that no ceiling can mend, so the ceiling grows only while the cache is at 90% of it or more. On the tuning workload without that gate the law raised the ceiling to 500 MB while 21 MB were in use; with it, the ceiling follows the working set.

The profiles and the pedals

The compass has two **profiles**, service (the default) and batch, which differ only in what the pedals do with a pile of work: in the service profile the law paces the slack; in the batch profile cruise puts every machine in service while work waits for a place and the emergency brake takes them straight to the floor when the queue is done (section 1.3). Both are written out in the Kubernetes preregistration (docs/K8S_COMPASS_PREREGISTRATION.md, rules 7 and 8, amendment 9), both ran in the batch test, and the result (Chapter 41) is machines 19% to 23% fewer over the whole window and 29% to 35% fewer after the queue finished, with the queue itself finishing no later beyond the noise.

CHAPTER 17

The Physics of a Processor

A GPU's firmware raises its clock one step at a time whenever there is work and room; each higher clock needs a higher voltage, and dynamic power rises with the clock times the square of the voltage, so the top of the clock range is where each extra step of speed costs the most watts. When the draw crosses the power limit, the firmware knocks the clock down several steps; with room again, it climbs again. On a busy card this repeats many times a second. The clock saws against the limit, the overshoots burn watts before the limiter catches them, and the card spends its time at the steepest part of its power curve. The power is the burner; the cooling removes the heat at its own pace; the limiter turns the burner down so the cooling can catch up, and the boost turns it back up.

The firmware does not aim at a balance point because it cannot see the job: it does not know whether the work is a response due in milliseconds or a batch due overnight, the right point moves with the workload, the room and the chip, and a limiter is simple to prove safe. The vendor default is set for maximum performance, and the electricity bill is the customer's.

Omni-Compass sits on the customer's side, where the job's service line is known. It holds two wires: the clock ceiling sets how high the boost may climb, and the power limit becomes a lid set just above what that ceiling draws. The boost stops at the bottom of the compass instead of slamming into the wall; the limiter rarely needs to act. Every watt kept off the chip is saved twice in a data center: once at the chip and again at the chillers that would have carried its heat.

CHAPTER 18

The Two-Way Nervous System

Every muscle is wired both ways: a sensory wire in (its meters) and a motor wire out (its knob), with the read-back closing the loop. Between them sits the nervous system (`omnicompass/nervous_system.py`), which decides how much authority each organ has at each moment:

- **Expand is always allowed** (except under a security hold): adding capacity, power, cooling or protection never waits.
- **Contract needs calm and a clean service record**: giving anything back is allowed only while the organ is calm enough and service has been inside its line for the last three decisions; continuous organs give back at most the calm share of their surplus per decision; discrete organs one unit at a time, through a release gate.
- **Fail up**: while service is breached, the knob returns to native at once.
- **Blind means hold**: if a sense goes stale or unreadable, nothing is given back until it returns.

One brain reads every organ at once. Because one law sets every knob, no two muscles fight: when the GPU's watts turn to heat, the cooling knob already knows it is coming; when pods scale up, the power envelope is ready.

Authority, and the release gate

Authority is how far the nervous system lets an organ move in one decision. It is granted from the engine's state (section 4.3): an organ in a calm body with a clean service record may give back; an organ in a stressed body may only add. On Kubernetes the **release gate** (`omnicompass/nervous_system.py`, `node_release_gate`) lets one machine be given back only when every one of the following holds, and writes the first reason that fails into the audit so a reader can see why a machine was kept: the remaining machines would stay under 95% full; no pod is waiting for a place; the pods are not scaling up; the service is not breached; the organ's own pressure is under the release threshold; every sense is live; and the last command has landed (no new order goes on top of one the cluster has not yet carried out). One machine per decision is the rule in the slow, reversible direction; past the wall, one machine up at once.

What the gate writes, in a real run. The first repetition of the first kill-scenario run (`results/live/raw/run-37716845219/paired-1/`, the omni arm's `audit.jsonl`) holds 114 decisions and 42 writes. Its gate records read, in the governor's own words, "release permitted" eleven times, each with the utilisation the remaining machines would carry after the release (0.034 to 0.153, all far under the 95% wall); "node organ has no contraction authority" once, a calm cluster whose engine state had not yet earned the right to give back; and "latency breached now; node organ has no contraction authority; a sense is blind" once, three reasons at once in a decision whose latency sense read stale, every one of them a reason to hold. A reader can therefore see, decision by decision, not only what the governor did but why it declined to do the rest, and the first reason that failed is the one written, so that a machine kept in service always has a stated cause.

Proprioception and the one-writer rule

The nervous system knows where its own levers are. Every write is read back from the device (the plug contract, section 8), and the audit records both the value sent and the value the device took. If a lever is found at a value Omni-Compass did not write, someone else owns it: the governor stops writing that lever and leaves it alone, because restoring over it would fight the new owner. This is **proprioception**: the governor's sense of its own hand. It is the reason a wrong installation leaves a signature in the receipts (section 8.5) rather than a silent error, and it is what the wire check proves before anything runs.

Blind means hold, and fail up

A sense that goes stale or unreadable is not read as calm. It is read as **blind**, and while any sense is blind nothing is given back; the knob holds, or, if the service was breached when the sense went blind, returns to native at once. In the fault test on real Kubernetes the response-time probe is deliberately paused for a minute in every arm; the audit shows the governor holding through the blind minute and resuming when the sense returns, and the table shows no row worse for it. The rule is simple to state and strict in effect: Omni-Compass never acts on a reading it does not have.

Why one brain

The alternative to one brain is a controller per layer, each tuned alone: an autoscaler that adds pods while a power cap is trying to hold the machines down, a cooling loop that chases heat the clock governor is about to remove. These controllers fight because none of them knows the others' intent. One law that sets every knob from one state cannot fight itself: the force on each lever is computed from the same reading of the same body at the same moment. This is also why the modelled organisms are run whole, at up to 1.7 million muscles on one clock (section 2.4); the question they answer is whether one brain stays coherent at that size, and the reading at every size is the same to the digit.

CHAPTER 19

The Nervous System in Detail

The two-way nervous system (live controller)

Source. Manuscript Appendix J: "Nervous system. The signal layer: sensing interfaces, unit integrity, timing coherence, delay and dropout handling, and feedback interpretation. ... The brain cannot compensate for corrupted signals." Section 5.3: the engine "projects forward state trajectories ... and detects destabilization pressure before divergence."

The gap. Before this change, the live controller fed the engine's stale (signal dropout) and drift_ratio (actuation mismatch) channels with a constant 0. So the upward path could not report a blind sense, and the downward path never learned whether its orders landed.

Upward: afferent integrity (senses to engine)

Every declared sense reports whether it is live.

- **Latency.** Blind when its newest sample is older than two windows by the wall clock, or when the window holds no successful request.
 - A hung probe freezes its file; its last clean window must never be read as the present.
 - Failed requests inside a live window become latency pressure (the failed share), so a failure is never read as silence.
- **Power.** Blind when its command returns no number.

How blindness is used.

- The blind share enters the engine as stale. It raises E, lowers U and feeds the external bath (equations 1-3 via assimilate).
- The nervous system grants no contraction on a blind sense, to any organ, and admits no held batch work.
- Expansion and batch pause (the protective directions) stay allowed.

Downward and back: efferent feedback (proprioception)

Every command is recorded.

- The node count commanded.
- The HPA targets written.

At the next decision the observed state is read back. For each organ, $\text{drift} = |\text{observed} - \text{commanded}| / \text{commanded}$.

- The largest drift enters the engine as drift_ratio. It raises E and I_U, lowers U and feeds B (equations 1-3).

- The node organ may not be given a new release while its last command did not land. For example, a drain refused by the PodDisruptionBudget leaves the node in service, and the next decision sees it.

Sideways: attribution (organ to organ)

The machine organ has its own engine view, fed with machine-attributable pressure only (pods waiting for a place).

It releases a machine only when all of these hold:

- pods are not scaling up and latency is not breached;
- nothing is pending;
- after the release the remaining machines stay at or below the engine's rho;
- its own calm, stress and security gates grant contraction;
- every sense is live;
- its last command landed.

Tests (all in verify.py)

Test	What it checks
tests/test_two_way.py	blind senses (about 75,000 random states), frozen-probe detection, failures read as pressure, release refused while blind or after an order that did not land
tests/test_node_release_gate.py	the release gate, over 200,000 states
tests/test_nervous_system.py	the original invariants N1-N6, over 300,000 states

Every live decision records the senses, their ages, the proprioceptive drift and the gate's reason in the audit, and the benchmark prints them in the job log.

Evaluation and simulation use only. Copyright (c) 2026 The Omni-Compass LLC. Commercial use, commercialization or monetization of any part of Omni-Compass requires a signed, paid Omni-Compass Enterprise License from The Omni-Compass LLC. Patent applications, copyright registrations and trademark applications filed in the United States. See LICENSE and NOTICE at the root of this repository.

Nervous system

Omni-Compass is one field. Muscles are nerves.

```
telemetry -&gt; afferent (per muscle)
           -&gt; one six-state engine
           -&gt; reflex (shield)
           -&gt; efferent (only WIRED + authority + not killed)
           -&gt; native controller on kill
```

omnicompass/nervous.py is the register. It does not invent GPU or chiller physics.

Status	Meaning
wired	this tree can sense and push
sensed	this tree can sense; it does not push
open	named, no plant, no push

Wired today: nodes, hpa, power_cap. Sensed: heat, network, security. Open: GPU, cooling, grid, queues, agents, ... Never a muscle: value alignment.

```
python k8s_controlplane/test_nervous.py
```

Evaluation and simulation use only. Copyright (c) 2026 The Omni-Compass LLC. Commercial use, commercialization or monetization of any part of Omni-Compass requires a signed, paid Omni-Compass Enterprise License from The Omni-Compass LLC. Patent applications, copyright registrations and trademark applications filed in the United States. See LICENSE and NOTICE at the root of this repository.

PART FOUR

The Body: Muscles, Realms and Organisms

Six hundred and fifty-six muscles in four realms, stacked into six organisms, and every gauge used to judge them.

The Muscles, the Realms and the Six Organisms

The catalog

The catalog (`realms/catalog.csv`) lists **945 distinct muscles in 59 families** (Omni v2 and v3; v1 listed 656 in 46), each with its family, its plant model and its knob. A muscle is anything with a meter and a setting: a compute pool and its machine count, a GPU and its clock ceiling, a chiller and its setpoint, a battery and its reserve, a feeder and its tap, a robot joint and its speed limit, a pump and its flow. Each muscle carries a declared plant model, the simplest physics that reproduces how its meter follows its knob, and a native controller of its own, which is what the muscle does when Omni-Compass is not there. The catalog was built in four waves (`docs/realm_study/`): the spine from primary specifications, each realm's extras from its own specifications, a cross-check against industry and open-source controllers, and a verification, de-duplication and mapping pass. The register (`docs/REGISTER.md`, Chapter 1) lists every family with its muscle count and the kinds of knob it carries, so a reader can see at a glance what the modelled realms contain and what they do not.

The four realms and the spine

The muscles fall into four realms. Every realm stands on the same **spine** of 257 muscles (Kubernetes, machines, GPUs and CPUs, network, storage, observability, security, cooling and electrical distribution), because every realm's work runs on computers that must be cooled and powered, plus its own domain muscles.

Organism	Muscles	What it is
1. Compute / AI / Cloud	430	clusters, GPUs, AI training and inference, cloud capacity
2. Physics / Robotics / Autonomous	376	joints, fleets, vehicles, flight and spacecraft axes
3. Energy / Facility / Industrial	470	data halls, buildings, batteries, UPS, process loops, feeders, hospitals, farms, pipelines, mines, district heat, power plants, renewables, pharmaceutical and food plants
4. Distribution / Specialized	440	networks, storage, databases, commerce, workflows, radio networks, clinical systems, ports
5. The four stacked, every duplicate kept	1,716	all four realms on one clock, the shared spine counted in each
6. The whole tower, every muscle once	945	every distinct muscle on one clock

The realms differ in what their domain muscles are and in how much slack their native controllers leave. Compute, AI and Cloud is clusters, cards, training and inference pools and cloud capacity, where the knobs are replica targets, machine counts, clock ceilings and power limits, and where native controllers are already good, so the modelled gains are small and the real gains come from the queue and the tail. Physics, Robotics and Autonomous is joints, fleets, vehicles, flight and spacecraft axes, where the knob is a speed or effort limit and the physics decides, axis by

axis, whether spending slack as speed saves energy at all; this is the realm the v3 slack gate was written for. Energy, Facility and Industrial is data halls, buildings, batteries, UPS, process loops, feeders and plants, where the knobs are setpoints, units in service, reserves and taps, and where the muscles with the most room are also the ones whose reserves exist for a reason the model does not contain (a UPS held for an outage is never a lever). Distribution and Specialized is networks, storage, databases, commerce, workflows, radio networks, clinical systems and ports, where the knob is usually an admission, a pool or a replica count, and where the real stacks of this manual (the pooler, the broker, the cache) live. Every realm's list, by family and knob, is docs/REALM_MUSCLES.md and docs/REGISTER.md Chapter 1.

What an organism run is

These six organisms are the benchmark set of the modelled realms. An organism is every one of its muscles stepped on one clock, each from its own plant model, each with its own native controller, and the whole fed one demand trace from one seed. Native runs the organism with every native controller alone. Omni runs the same organism, the same seed, the same demand and the same clock, with the compass law on every knob the verdict allows. Each run produces its own receipt: the organism's work, its energy, its time over its service line, and every knob's hand-back. In code and in the workflows the two large organisms are named `tower` and `stack`, never by a count, because the counts changed between v1 and v2 and a name that carried a count would read as a different organism; the old names are read as aliases.

The grid of copies and runs

Every organism runs at **1, 10, 100 and 1,000 copies** (that many instances of the organism on one clock, so the four stacked at 1,000 copies is 1.7 million modelled muscles) and at **1, 10, 100 and 1,000 paired runs** (that many native-and-omni pairs, each from its own seed): 96 cells in all, both arms shown in every cell. One thousand is the maximum in both directions. The grid as it stands (results/scale/GRID.md) is 84 of 90 judged cells on Omni v3; the six cells left, 100 and 1,000 runs at 1,000 copies, are beyond the machines available and are declared as such, not hidden. The reading at every size is the same to the digit, which is itself a finding about the law: it does not drift with scale.

The organisms with a real cluster inside

The modelled organisms are also run with a **real Kubernetes cluster inside** as one more muscle (tools/run_kil.py, workflows six-kube, big-organism, big-organism-detached): the organism's own compute demand drives a real load generator against a real service on a real cluster, and the cluster's watts come back into the organism as heat and load. At 1 to 100 copies this runs on GitHub's machines; at 1,000 copies it needs a rented machine and runs there, detached from the GitHub job that started it, for about seven hours a repetition. These runs are evidence class L for the cluster and S for the organism around it, and are labelled "L + S".

The register

The register (docs/REGISTER.md) is the one list a referee can audit: every muscle by family (Chapter 1), every benchmark by platform with its native controller, Omni-Compass's knob, its gauges, its result and its file (Chapter 20), what is not measured yet, said plainly (Chapter 53), and every open benchmark still to run, in order, one or two at a time (Chapter 7). The program that takes each benchmark to its full size, with what each step costs, is docs/PROOF_PROGRAM.md. Every benchmark in this manual appears in the register under the same name.

CHAPTER 21

The Four Realms

The muscle tower on modelled plants (945 muscles in Omni v2, 656 in v1): four realms and the whole tower as a fifth organism, each run natively and with Omni-Compass on top, with meters and receipts. Evidence class **S** (simulation).

Run it

```
pip install -r requirements.txt
python3 tools/run_realms.py          # the preregistered run: every muscle and 5 organisms,
  ↳ seeds 1000-1009
python3 tests/test_realms.py        # the harness's own checks (also inside verify.py)
```

Results (Omni v3; v2's are in results/realms/v2/, v1's in results/realms/v1/): results/realms/REALMS.md (the tables), MUSCLES.csv (one row per muscle), REALMS.json (every per-seed contrast), RUN.json and SHA256SUMS.txt (commit and fingerprints).

What is where

File	What it is
realms/catalog.csv	the 945 muscles: family, name, realm, plant, parameter set, knob (v1's 656: realms/catalog_v1.csv)
tools/realms_catalog.py	the rules that gave each muscle its realm, plant and knob
realms/plants.py	the five plants and their native controllers; the four knobs; the capacity law for one plant
realms/presets.py	every parameter, one set per family class
realms/harness.py	the arms (native, watch, omni, fixed setpoint), the organisms, the outcome and the label rule
docs/REALMS_PREREGISTRATION.md	the question, the rules, the seeds and the development history, frozen before the run

How it relates to the rest

Layer	Realm harness	Elsewhere in this repository
Mathematics (T, V)	uses the frozen engine and governor unchanged	docs/TRACKING_THEOREM.md, verify.py
Simulation (S)	this	fleet and cluster simulators, GPU model
Real software (L)	not here	set 22 on real Kubernetes (results/live/LIVE_REPS_22.md)
Physical (P)	not here	the GPU bench: first run on an NVIDIA A10, 2026-10-02 (results/gpu/run-20261002T082232Z/GPU_REPS.md); the corrected governor not yet run on a card (docs/GPU_PREREGISTRATION.md)

A realm result that looks good is a reason to test that knob on a real machine, not a substitute for it. The realms whose knobs can be tested for real first are the compute realm's (the GPU bench, kind), because the tools already exist.

Evaluation and simulation use only. Copyright (c) 2026 The Omni-Compass LLC. Commercial use, commercialization or monetization of any part of Omni-Compass requires a signed, paid Omni-Compass Enterprise License from The Omni-Compass LLC. Patent applications, copyright registrations and trademark applications filed in the United States. See LICENSE and NOTICE at the root of this repository.

CHAPTER 22

The Realm Muscles

Generated from `realms/catalog.csv` (rules: `tools/realms_catalog.py`; the v2 build: `tools/realms_catalog_v2.py`). Every realm's organism is its own families plus the **shared spine** (the infrastructure every real stack runs on), so the realm counts add to more than the tower. Each muscle is followed by the one knob Omni may hold on it.

The shared spine: 257 muscles, in all four realms

Cloud VM & Capacity (21; plant: `compute_pool`, `node`)

`vm_count` [capacity], `vm_size` [capacity], `vm_start_stop` [capacity], `vm_migrate` [capacity], `instance_family` [capacity], `asg_min` [capacity], `asg_max` [capacity], `spot_mix` [capacity], `reservation_mix` [capacity], `zone_selection` [capacity], `region_selection` [capacity], `architecture_selection` [capacity], `accelerator_selection` [capacity], `boot_disk_class` [capacity], `placement_group` [capacity], `interruption_response` [capacity], `target_tracking_policy` [setpoint], `scaling_cooldown_warmup` [capacity], `vm_cpu_cap_by_priority` [admission], `predictive_scaling` [capacity], `instance_refresh` [capacity]

Container Resources (23; plant: `compute_pool`, `server`)

`cpu_request` [capacity], `cpu_limit` [capacity], `memory_request` [capacity], `memory_limit` [capacity], `ephemeral_storage` [capacity], `hugepages` [capacity], `io_weight` [setpoint], `pids_limit` [capacity], `memory_high` [capacity], `memory_reclaim` [capacity], `swap_limit` [capacity], `cgroup_io_max` [capacity], `cpu_weight` [setpoint], `cpu_quota` [admission], `cpuset` [capacity], `runtime_class` [capacity], `in_place_resize` [capacity], `namespace_quota` [admission], `default_limits` [capacity], `memory_hard_limit` [capacity], `memory_protection` [capacity], `io_latency_target` [setpoint], `freeze` [admission]

Cooling, Chillers & Thermodynamics (18; plant: `thermal_zone`, `data_hall`)

`supply_air_temperature` [setpoint], `return_air_target` [setpoint], `coolant_supply_temperature` [setpoint], `coolant_flow` [capacity], `pump_speed` [capacity], `fan_speed` [capacity], `chiller_setpoint` [setpoint], `compressor_authority` [power], `cooling_tower_fan` [capacity], `cooling_capacity` [capacity], `rack_thermal_budget` [power], `gpu_thermal_envelope` [setpoint], `cpu_thermal_envelope` [setpoint], `thermal_workload_migrate` [admission], `thermal_load_shed` [admission], `server_fan_speed` [capacity], `economiser` [capacity], `chiller_plant_supervisory_control_rl_with_safety_layer` [capacity]

Host CPU & Memory (23; plant: compute_pool, cpu_host)

cpufreq_min [power], cpufreq_max [power], rapl_package_power [power], uncore_frequency [power], energy_perf_preference [power], cpu_idle_policy [power], memory_bandwidth [capacity], numa_balance [capacity], irq_affinity [capacity], llc_allocation [capacity], memory_pressure_gate [admission], page_reclaim_rate [capacity], transparent_hugepages [capacity], core_online_offline [capacity], thermal_throttle_policy [power], host_power_profile [power], frequency_governor [power], turbo_boost [capacity], dram_power_limit [power], swappiness_reclaim [capacity], core_packing [capacity], per_pod_cpu_frequency_range_and_c_state_access [power], uncore_frequency_per_node [power]

Kubernetes Placement & Scheduling (20; plant: compute_pool, server)

node_selector [capacity], node_affinity [capacity], pod_anti_affinity [capacity], topology_spread [capacity], numa_placement [capacity], gpu_topology [capacity], storage_locality [capacity], network_locality [capacity], taint_toleration [setpoint], priority_class [admission], preemption_policy [admission], device_claim [capacity], failure_domain_spread [capacity], scheduler_backoff [admission], gang_admission [admission], deschedule [capacity], scheduler_scoring_strategy [capacity], scheduler_plugin_weights [setpoint], scheduling_gates [admission], cpu_manager [capacity]

Kubernetes Workload Scaling (31; plant: compute_pool, server)

replicas [capacity], hpa_cpu_target [setpoint], hpa_memory_target [setpoint], hpa_custom_target [setpoint], vpa_apply [capacity], scale_to_zero [capacity], keda_threshold [admission], rollout_rate [capacity], max_surge [capacity], max_unavailable [capacity], deployment_pause_resume [admission], rollout_abort [admission], pod_eviction [admission], pdb_policy [capacity], scheduler_queue_priority [admission], api_priority_fairness [admission], hpa_min_replicas [capacity], hpa_max_replicas [capacity], hpa_scale_down_stabilisation_window [setpoint], hpa_scale_up_stabilisation_window [setpoint], hpa_scaling_rate_policies [capacity], hpa_policy_selection [capacity], hpa_tolerance [capacity], hpa_sync_period [capacity], vpa_minmax_allowed_resources [capacity], keda_polling_interval [setpoint], rollout_progress_deadline [admission], rollout_readiness_delay [capacity], unhealthy_pod_eviction_policy [admission], graceful_termination [capacity], probes [capacity]

NVIDIA GPU Hardware (19; plant: compute_pool, gpu)

gpu_allocate [capacity], gpu_power_limit [power], gpu_sm_clock [power], gpu_memory_clock [power], gpu_persistence_mode [capacity], gpu_compute_mode [capacity], mig_mode [capacity], mig_geometry [capacity], gpu_timeslice [capacity], gpu_mps [capacity], gpu_quarantine [admission], gpu_reset [capacity], gpu_thermal_limit [power], gpu_ecc_response [capacity], gpu_job_power_budget [admission], application_clocks [power], gpu_temperature_target [setpoint], sync_boost [capacity], target_clocks [power]

Network Routing & Switching (19; plant: compute_pool, network)

lb_weight [setpoint], route_weight [setpoint], rate_limit [admission], bandwidth_limit [capacity], qos_class [capacity], connection_limit [admission], failover_route [capacity], nic_rate_limit [admission], nic_queue_count [capacity], queue_discipline [capacity], congestion_control [capacity], egress_budget [admission], ingress_budget [admission], ecmp_weight [setpoint], path_selection [capacity], network_isolation [admission], nic_interrupt_coalescing [capacity], nic_ring_size [capacity], port_link_power [power]

Node Fleet & Karpenter-Class Control (24; plant: compute_pool, node)

node_desired [capacity], node_pool_min [capacity], node_pool_max [capacity], node_provision [capacity], node_cordon [admission], node_drain [admission], node Consolidate [capacity], node_replace [capacity], node_shutdown [capacity], node_power_on [capacity], nodepool_weight [setpoint], disruption_budget [admission], consolidation_policy [capacity], consolidate_after [capacity], expire_after [capacity], capacity_class [capacity], scale_down_utilisation_threshold [admission], scale_down_unneeded_time [capacity], scale_down_delay_after_add [capacity], expander [capacity], node_pool_limits [capacity], kubelet_eviction_thresholds [admission], kubelet_reserved_resources [capacity], max_pods_per_node [capacity]

Observability & Telemetry (9; plant: compute_pool, server)

collector_memory_limit [capacity], export_concurrency [admission], cardinality_budget [admission], retention_window [setpoint], remote_write_queue [capacity], telemetry_shed [admission], metric_scrape_interval [setpoint], batching [capacity], per_pod_energy_attribution [capacity]

PDU, UPS & Electrical Distribution (18; plant: energy_storage, ups)

server_power_cap [power], rack_power_cap [power], pdu_branch_power_limit [power], pdu_outlet_control [capacity], ups_operating_mode [capacity], ups_charge_rate [power], ups_discharge_rate [power], phase_balance [capacity], load_transfer [admission], reactive_power_target [setpoint], voltage_target [setpoint], generator_dispatch [capacity], electrical_isolation [admission], breaker_trip_gate [admission], ups_shutdown_battery_test [capacity], hierarchical_power_budget_row_pdu_switchboard_controllers [power], priority_aware_capping_throttle_low_priority_first [power], power_oversubscription_with_prediction [power]

Reliability, Security & Recovery (12; plant: compute_pool, server)

restart [capacity], rollback [capacity], traffic_divert_recovery [capacity], degraded_mode [capacity], actuator_freeze [admission], workload_isolate [admission], credential_rotation_gate [admission], policy_enforcement [capacity], rate_abuse_gate [admission], fault_domain_isolate [admission], backup_trigger [capacity], kill_switch [admission]

Storage Block/File/Object (20; plant: compute_pool, storage)

volume_size [capacity], iops_limit [capacity], throughput_limit [capacity], replica_count [capacity], storage_tier [capacity], volume_placement [capacity], snapshot_trigger [capacity], rebalance [capacity], recovery_rate [capacity], backfill_rate [capacity], compaction_pressure [capacity], cache_allocation [capacity], object_replication [capacity], erasure_code_profile [capacity], storage_admission [admission], degraded_storage_gate [admission], scrub_schedule [capacity], io_scheduler [capacity], disk_power_spin_down [power], lifecycle_tiering [capacity]

Realm 1: Compute / AI / Cloud: 173 own muscles + the spine = 430 in its organism

AI Inference Serving (24; plant: compute_pool, gpu)

model_replicas [capacity], model_route_weight [setpoint], model_load [capacity], model_unload [capacity], model_instance_count [capacity], continuous_batching [capacity], max_batch_size [capacity], batch_queue_delay [capacity], inference_concurrency [admission], inference_max_tokens [capacity], kv_cache_budget [admission], prefix_cache_budget [admission], speculative_decode_budget [admission], model_precision [power], inference_priority [admission], inference_slo_gate [admission], autoscaling_target_concurrency_qps [admission], scale_to_zero_panic_window [setpoint], request_priority_levels [admission], kv_cache_swap_space [capacity], chunked_prefill [capacity], tensor_parallel_degree [capacity], disaggregated_prefill_decode [capacity], lora_adapter_load [capacity]

AI Training (18; plant: compute_pool, gpu_batch)

training_workers [capacity], global_batch_size [capacity], microbatch_size [capacity], gradient_accumulation [capacity], data_parallelism [capacity], tensor_parallelism [capacity], pipeline_parallelism [capacity], expert_parallelism [capacity], checkpoint_interval [setpoint], checkpoint_trigger [capacity], training_preempt [admission], training_gang_size [capacity], elastic_worker_count [capacity], straggler_mitigation [capacity], training_precision [power], compute_comm_overlap [capacity], activation_checkpointing [capacity], dataloader_workers [capacity]

Cross-Cluster, Multi-Region & Edge (15; plant: compute_pool, node)

multi_cluster_dispatch [capacity], region_dispatch [capacity], zone_dispatch [capacity], edge_dispatch [capacity], cloud_capacity_class [capacity], workload_migrate_region [capacity], data_residency_gate [admission], latency_region_gate [admission], cost_region_gate [admission], carbon_region_gate [admission], global_failover [capacity], federation_quota [admission], cross_cluster_replication [capacity], edge_offload [capacity], global_admission [admission]

DPU SmartNIC & Programmable IO (12; plant: compute_pool, fabric)

dpu_pf_tx_rate [capacity], dpu_vf_tx_rate [capacity], dpu_sf_tx_rate [capacity], dpu_bandwidth_share [capacity], dpu_qos_group [capacity], sriov_vf_count [capacity], smartnic_flow_steering [capacity], smartnic_offload_enable [capacity], dpu_cpu_budget [admission], dpu_memory_budget [admission], dpu_service_placement [capacity], dpu_failover [capacity]

Distributed Cluster Managers (15; plant: compute_pool, batch)

task_admission [admission], task_priority [admission], resource_reservation [capacity], cluster_quota [admission], task_preemption [admission], binpack_pressure [capacity], spread_pressure [capacity], gang_schedule [capacity], worker_allocation [capacity], maintenance_evacuation [capacity], oversubscription [capacity], resource_reclaim [capacity], queue_fairness [admission], deadline_pressure [admission], scheduler_retry [admission]

GPU Fabric & RDMA (18; plant: compute_pool, fabric)

nvlink_placement [capacity], nvswitch_route [capacity], gpu_fabric_quarantine [admission], rdma_bandwidth [capacity], rdma_route [capacity], nic_affinity [capacity], collective_concurrency [admission], collective_algorithm [capacity], collective_chunk_size [capacity], rank_placement [capacity], gpudirect_policy [capacity], congestion_response [capacity], rail_selection [capacity], fabric_failover [capacity], communication_priority [admission], fabric_isolation [admission], collective_protocol [capacity], nic_hca_selection [capacity]

HPC & Distributed Compute (19; plant: compute_pool, batch)

job_slots [capacity], mpi_ranks [capacity], rank_mapping [capacity], node_allocation [capacity], job_walltime [capacity], job_priority [admission], job_preemption [admission], checkpoint_restart [capacity], parallel_io_budget [admission], collective_budget [admission], accelerator_share [capacity], cpu_gpu_ratio [setpoint], memory_per_rank [capacity], scratch_budget [admission], scheduler_fair_share [admission], backfill_policy [capacity], node_power_saving [power], cpu_frequency_per_job [power], partition_limits [capacity]

Kubernetes Dynamic Device Allocation (8; plant: compute_pool, gpu)

dra_device_class_selection [capacity], dra_claim_capacity [capacity], dra_claim_sharing [capacity], dra_device_taint [capacity], dra_device_eviction [admission], dra_binding_readiness [capacity], dra_binding_failure_response [capacity], dra_device_configuration [setpoint]

OpenShift & Machine API (9; plant: compute_pool, node)

machine_remediation [capacity], machine_health_gate [admission], mcp_pause [admission], mcp_max_unavailable [capacity], node_config_rollout [capacity], operator_reconcile_budget [admission], cluster_version_pacing [capacity], infra_machine_admission [admission], machine_failure_domain [capacity]

Quantum Computing Control Simulation (16; plant: compute_pool, qpu)

qubit_mapping [capacity], circuit_admission [admission], shot_allocation [capacity], circuit_scheduling [admission], gate_scheduling [capacity], pulse_amplitude [capacity], pulse_duration [setpoint], pulse_phase [capacity], pulse_frequency [power], coupling_control [capacity], reset_scheduling [capacity], measurement_scheduling [capacity], dynamical_decoupling [capacity], noise_aware_routing [capacity], error_mitigation_budget [admission], quantum_queue_priority [admission]

Work Admission & Demand Shaping (19; plant: compute_pool, server)

api_concurrency [admission], queue_concurrency [admission], queue_backpressure [admission], job_admission [admission], batch_admission [admission], inference_admission [admission], load_shed [admission], priority_gate [admission], tenant_admission [admission], burst_limit [admission], deadline_admission [admission], work_budget [admission], request_queue_limit [admission], retry_admission [admission], background_work_gate [admission], maintenance_work_gate [admission], request_concurrency_limit [admission], adaptive_concurrency_limit [admission], priority_request_queue_under_limit [admission]

Realm 2: Physics / Robotics / Autonomous: 119 own muscles + the spine = 376 in its organism***Automotive EV & Mobile Powertrain (14; plant: motion_axis, ev_traction)***

traction_torque_limit [power], regen_braking_level [power], battery_charge_limit [power], battery_discharge_limit [power], battery_thermal_target [power], motor_thermal_limit [power], vehicle_speed_envelope [capacity], energy_recovery_target [power], auxiliary_power_budget [power], fast_charge_current [power], fast_charge_voltage [power], vehicle_safe_state [admission], state_of_charge_limits [power], smart_charging_schedule [admission]

Aviation & Autonomous Flight (20; plant: motion_axis, flight_axis)

throttle_envelope [power], attitude_target [capacity], attitude_rate_target [capacity], velocity_target [capacity], altitude_target [capacity], waypoint_authority [capacity], flight_hold [admission], return_to_home [admission], land_action [admission], mission_admission [admission], geofence_response [admission], failsafe_selection [admission], battery_reserve_threshold [admission], actuator_saturation_envelope [capacity], flight_mode_transition [admission], flight_termination_safe_state [admission], attitude_rate_gains [capacity], horizontal_speed_limit [power], vertical_speed_limits [power], acceleration_jerk_limits [power]

Elevators & Vertical Transport (12; plant: motion_axis, elevator_hoist)

hoist_speed_target [capacity], acceleration_limit [capacity], regen_drive_mode [power], standby_power_mode [admission], destination_dispatch_schedule [admission], car_parking_mode [admission], door_dwell_hold [admission], escalator_speed_target [capacity], group_capacity_mode [admission], motor_thermal_derate [power], brake_test_schedule [admission], fire_recall_mode [admission]

Marine Propulsion & Vessel Automation (12; plant: motion_axis, marine_propulsion)

shaft_speed_target [capacity], propeller_pitch_limit [power], shaft_power_limit [power], bow_thruster_duty [power], ballast_pump_mode [admission], engine_load_sharing_setpoint [capacity], slow_steaming_speed_target [capacity], auxiliary_engine_staging [capacity], shore_power_mode [admission], hotel_load_budget [power], trim_target [capacity], rudder_rate_limit [power]

Rail Traction & Train Control (14; plant: motion_axis, rail_traction)

traction_effort_limit [power], regenerative_braking_share [power], coasting_speed_target [capacity], dwell_time_schedule [admission], hvac_duty_cycle [power], train_auxiliary_power_budget [power], platform_approach_speed [capacity], acceleration_rate_setpoint [capacity], wheel_slip_protection_mode [admission], traction_motor_thermal_derate [power], catenary_voltage_limit [power], timetable_recovery_margin [admission], headway_target [capacity], door_release_hold [admission]

Robotics Fleet & Warehouse Automation (15; plant: compute_pool, robot_fleet)

robot_dispatch [capacity], task_assignment [capacity], traffic_reservation [capacity], robot_route [capacity], charging_dispatch [capacity], battery_reserve [capacity], elevator_request [capacity], door_request [capacity], conveyor_speed [capacity], agv_speed [capacity], warehouse_zone_admission [admission], robot_quarantine [admission], fleet_failover [capacity], human_safe_stop [capacity], fleet_concurrency [admission]

Robotics Motion Control (18; plant: motion_axis, robot_joint)

joint_position [capacity], joint_velocity [capacity], joint_acceleration [capacity], joint_effort [power], cartesian_velocity [capacity], trajectory_speed [capacity], trajectory_acceleration [capacity], jerk_limit [capacity], collision_margin [capacity], force_limit [power], gripper_force [power], locomotion_speed [capacity], steering_angle [capacity], braking_force [power], balance_correction [capacity], trajectory_tolerances [capacity], velocity_acceleration_scaling [capacity], controller_update_rate [capacity]

Spacecraft & Flight Software (14; plant: motion_axis, reaction_wheel)

space_command_admission [admission], flight_task_schedule [admission], space_mode_transition [admission], payload_duty_cycle [power], communication_allocation [capacity], space_power_budget [power], space_thermal_command [power], attitude_command_envelope [capacity], reaction_wheel_allocation [capacity], rcs_authority [power], safe_mode_transition [admission], watchdog_recovery [admission], instrument_activation [capacity], fault_isolation [admission]

Realm 3: Energy / Facility / Industrial: 213 own muscles + the spine = 470 in its organism

Agriculture & Irrigation (14; plant: process_loop, irrigation)

irrigation_pump_speed [capacity], mainline_pressure_setpoint [setpoint], soil_moisture_target [setpoint], pivot_speed_setpoint [setpoint], fertigation_dose [power], greenhouse_temperature_setpoint [setpoint], greenhouse_co2_target [setpoint], vent_position_setpoint [setpoint], grain_dryer_temperature_setpoint [setpoint], cold_storage_temperature_setpoint [setpoint], barn_ventilation_rate [capacity], milking_vacuum_level_setpoint [setpoint], well_drawdown_level_target [setpoint], drip_zone_dispatch [admission]

Building & Critical Environment HVAC (15; plant: thermal_zone, building)

zone_temperature_target [setpoint], zone_airflow [capacity], ahu_fan_speed [capacity], damper_position [capacity], economizer_position [capacity], boiler_setpoint [setpoint], heat_pump_mode [admission], humidity_target [setpoint], occupancy_ventilation [admission], building_demand_limit [power], thermal_storage_dispatch [capacity], hvac_emergency_mode [admission], setpoint_reset_trim_and_respond [setpoint], duct_static_pressure_setpoint [setpoint], optimal_start_stop [capacity]

District Heating & Cooling (13; plant: process_loop, district_heat)

supply_temperature_setpoint [setpoint], return_temperature_target [setpoint], differential_pressure_setpoint [setpoint], network_pump_speed [capacity], heat_pump_staging [capacity], chp_dispatch [admission], thermal_storage_level_target [setpoint], peak_boiler_heater_output [power], substation_flow_limit [capacity], peak_demand_shed [admission], outdoor_reset_setpoint [setpoint], cooling_network_supply_temperature [setpoint], cooling_storage_level_target [setpoint]

Energy Storage & Microgrid (19; plant: energy_storage, microgrid)

battery_charge_power [power], battery_discharge_power [power], battery_soc_reserve [setpoint], grid_import_limit [power], grid_export_limit [power], pv_curtailment [power], ev_charge_power [power], heat_pump_power [power], electrolyzer_power [power], microgrid_demand_limit [power], peak_shaving [capacity], time_of_use_schedule [setpoint], energy_load_shed [admission], flex_load_admission [admission], storage_dispatch [capacity], microgrid_emergency_reserve [setpoint], volt_var [capacity], volt_watt [capacity], constant_power_factor_reactive_power [power]

Facility & Grid Optimization (15; plant: energy_storage, facility)

facility_power_budget [power], utility_demand_limit [power], demand_response [admission], electricity_price_gate [admission], carbon_intensity_gate [admission], renewable_dispatch [capacity], generator_start_stop [capacity], site_battery_dispatch [capacity], pue_target [setpoint], cooling_power_budget [power], it_power_budget [power], rack_power_allocation [power], facility_peak_guard [capacity], grid_frequency_response [admission], facility_islanding [admission]

Grid Transmission & Distribution (12; plant: process_loop, feeder_voltage)

capacitor_bank_switch [capacity], voltage_regulator_tap [setpoint], transformer_tap [setpoint], inverter_real_power [power], inverter_reactive_power [power], feeder_voltage_target [setpoint], feeder_load_transfer [capacity], distribution_storage_dispatch [admission], demand_response_dispatch [admission], frequency_droop_setpoint [setpoint], grid_protection_mode [admission], grid_restoration_sequence [admission]

Healthcare Critical Environments (13; plant: thermal_zone, hospital)

operating_room_air_change_setpoint [setpoint], isolation_room_pressure_target [setpoint],
 patient_room_temperature_setpoint [setpoint], surgical_suite_humidity_target [setpoint],
 ahu_supply_air_temperature [setpoint], pharmacy_cold_room_setpoint [setpoint],
 sterile_storage_humidity_setpoint [setpoint], imaging_suite_cooling_capacity [capacity], chiller_plant_staging
 [capacity], exhaust_fan_capacity [capacity], ward_night_setback_mode [admission],
 medical_gas_plant_demand_limit [power], emergency_power_load_shed [admission]

Industrial PLC & Process Automation (18; plant: process_loop, process)

plc_cycle_authority [admission], machine_cell_admission [admission], valve_position [setpoint], pump_flow
 [capacity], compressor_speed [capacity], heater_power [power], furnace_setpoint [setpoint], pressure_setpoint
 [setpoint], temperature_setpoint [setpoint], mass_flow_setpoint [setpoint], tank_level_target [setpoint],
 conveyor_rate [capacity], feed_rate [capacity], purge_vent_action [admission], controller_mode [capacity],
 alarm_limits [capacity], safety_interlock_trip [capacity], opc_ua_writes [capacity]

Mining & Mineral Processing (14; plant: process_loop, mill)

sag_mill_load_setpoint [setpoint], mill_speed_target [setpoint], crusher_gap_setpoint [setpoint],
 flotation_aeration_rate [power], cyclone_feed_pressure_target [setpoint], thickener_underflow_density_target
 [setpoint], reagent_dose [power], conveyor_speed_setpoint [setpoint], dewatering_pump_level_target [setpoint],
 ventilation_on_demand_airflow [capacity], stockpile_level_target [setpoint], slurry_pump_speed [capacity],
 tailings_discharge_shutdown [admission], ore_blend_dispatch [admission]

Oil & Gas Pipelines (14; plant: process_loop, pipeline)

compressor_discharge_pressure_setpoint [setpoint], pump_station_suction_pressure_target [setpoint],
 line_pack_target [setpoint], pipeline_flow_setpoint [setpoint], compressor_unit_staging [capacity],
 vfd_pump_speed [capacity], terminal_tank_level_target [setpoint], leak_detection_shutdown [admission],
 batch_interface_dispatch [admission], heater_outlet_temperature_setpoint [setpoint], drag_reducing_agent_dose
 [power], valve_position_target [setpoint], cathodic_protection_voltage [setpoint], pressure_protection_priority
 [admission]

Pharmaceutical & Food Manufacturing (14; plant: process_loop, batch_reactor)

reactor_temperature_setpoint [setpoint], agitator_speed_setpoint [setpoint], fermenter_dissolved_oxygen_target
 [setpoint], ph_setpoint [setpoint], pasteurizer_holding_temperature_setpoint [setpoint],
 freezer_tunnel_temperature_setpoint [setpoint], cip_cycle_dispatch [admission],
 cleanroom_pressure_cascade_setpoint [setpoint], lyophilizer_shelf_temperature_setpoint [setpoint],
 chromatography_flow_setpoint [setpoint], steam_sterilizer_cycle_priority [admission],
 oven_zone_temperature_setpoint [setpoint], refrigerant_compressor_authority [admission], batch_hold_quarantine
 [admission]

Power Generation & Turbine Control (14; plant: process_loop, turbine)

turbine_speed_droop_setpoint [setpoint], unit_load_setpoint [setpoint], boiler_steam_pressure_setpoint [setpoint],
 feedwater_level_target [setpoint], agc_participation_limit [capacity], excitation_voltage_setpoint [setpoint],
 hydro_gate_position_target [setpoint], combustion_air_ratio_setpoint [setpoint], cooling_water_flow_setpoint
 [setpoint], inlet_guide_vane_position [setpoint], reserve_dispatch_priority [admission], turbine_ramp_rate_limit
 [capacity], emissions_shutdown [admission], nuclear_rod_position_target [setpoint]

Renewable Generation & Inverter Control (13; plant: process_loop, inverter)

inverter_volt_var_setpoint [setpoint], inverter_volt_watt_setpoint [setpoint], active_power_curtailment [power], plant_reactive_power_target [setpoint], wind_turbine_yaw_offset [capacity], pitch_angle_limit [capacity], rotor_speed_setpoint [setpoint], inverter_frequency_droop_setpoint [setpoint], ramp_rate_limit [capacity], tracker_stow_mode [admission], string_mppt_voltage_setpoint [setpoint], noise_mode_schedule [admission], ice_detection_shutdown [admission]

Semiconductor Fab & Precision Manufacturing (13; plant: process_loop, chamber)

tool_job_dispatch [admission], wafer_route [capacity], chamber_recipe_selection [admission], chamber_temperature [setpoint], chamber_pressure [setpoint], gas_flow [capacity], rf_power [power], vacuum_pump_speed [capacity], robot_transfer_rate [capacity], lot_priority [admission], tool_quarantine [admission], run_to_run_control [capacity], idle_sleep_mode [capacity]

Water Wastewater & Pumping (12; plant: process_loop, water)

pump_speed_water [capacity], valve_position_water [setpoint], reservoir_level_target [setpoint], line_pressure_target [setpoint], flow_target_water [setpoint], aeration_rate [power], chemical_dose_rate [power], filtration_backwash [admission], lift_station_dispatch [admission], leak_isolation [admission], water_demand_shed [admission], water_emergency_shutdown [admission]

Realm 4: Distribution / Specialized: 183 own muscles + the spine = 440 in its organism***Cache & Memory Services (16; plant: compute_pool, server)***

cache_size [capacity], cache_ttl [setpoint], cache_eviction_policy [admission], cache_replicas [capacity], cache_sharding [capacity], cache_prefetch [capacity], cache_writeback_rate [capacity], cache_admission [admission], hot_key_isolation [admission], cache_connection_limit [admission], cache_memory_limit [capacity], cache_compression [capacity], cache_warmup [capacity], cache_failover [capacity], cache_flush_rate [capacity], memory_size_threads [capacity]

Commerce & Payment Systems (15; plant: compute_pool, commerce)

payment_admission [admission], payment_concurrency [admission], payment_retry [admission], payment_timeout [admission], fraud_review_gate [admission], authorization_route [capacity], processor_route_weight [setpoint], transaction_queue_limit [admission], idempotency_window [setpoint], order_reservation [capacity], inventory_hold [admission], checkout_load_shed [admission], refund_queue_rate [capacity], settlement_batch [capacity], payment_failover [capacity]

Data Analytics & ETL (15; plant: compute_pool, batch)

executor_size [capacity], dynamic_allocation_min [capacity], dynamic_allocation_max [capacity], shuffle_partitions [capacity], shuffle_bandwidth [capacity], etl_concurrency [admission], stage_parallelism [capacity], query_slots [capacity], spill_threshold [admission], cache_fraction [setpoint], batch_interval [setpoint], stream_backpressure [admission], data_locality_wait [capacity], speculation_policy [capacity], analytics_admission [admission]

Database & Transactions (19; plant: compute_pool, database)

db_replicas [capacity], db_memory [capacity], db_cache [capacity], query_concurrency [admission], read_route [capacity], shard_placement [capacity], db_failover [capacity], replication_lag_gate [admission], db_write_throttle [admission], db_pool_resize [capacity], transaction_concurrency [admission], lock_timeout [admission], checkpoint_rate [capacity], vacuum_compaction_rate [capacity], parallel_workers [capacity], background_writer [capacity], synchronous_replication [capacity], buffer_pool [capacity], io_capacity [capacity]

Medical Imaging & Clinical Systems (12; plant: compute_pool, clinical)

pacs_archive_tier_target [setpoint], dicom_router_concurrency [admission], ehr_application_replicas [capacity], hl7_interface_queue_limit [admission], fhir_api_rate_limit [admission], reconstruction_gpu_workers [capacity], modality_worklist_timeout [admission], patient_monitoring_gateway_capacity [admission], lab_analyzer_batch_window [admission], telehealth_session_admission [admission], clinical_backup_window [admission], imaging_prefetch_priority [admission]

Messaging & Streaming (18; plant: compute_pool, server)

partition_count [capacity], partition_placement [capacity], producer_quota [admission], consumer_quota [admission], broker_io_quota [admission], message_retention [capacity], queue_depth_limit [admission], consumer_concurrency [admission], producer_batch_size [capacity], fetch_batch_size [capacity], rebalance_rate [capacity], replication_factor [capacity], retry_backoff [admission], dead_letter_divert [capacity], stream_priority [admission], broker_failover [capacity], prefetch [capacity], memory_disk_alarm [capacity]

Ports & Maritime Logistics (12; plant: compute_pool, port)

quay_crane_allocation [capacity], yard_crane_fleet_size [capacity], berth_window_admission [admission], truck_gate_rate_limit [admission], horizontal_transport_fleet_size [capacity], reefer_plug_power_budget [power], shore_power_connection_capacity [capacity], rail_mounted_gantry_dispatch [capacity], container_dwell_priority [admission], vessel_arrival_pacing [admission], stacking_height_target [setpoint], equipment_charging_window [admission]

Runtime & Application (17; plant: compute_pool, server)

worker_count [capacity], thread_pool [capacity], jvm_heap [capacity], gc_budget [admission], connection_pool_runtime [capacity], application_cache_size [capacity], runtime_memory [capacity], async_concurrency [admission], event_loop_workers [capacity], process_count [capacity], request_timeout [admission], background_workers [capacity], runtime_cpu_budget [admission], runtime_io_budget [admission], runtime_restart [capacity], go_runtime [capacity], keepalive_connection_reuse [capacity]

Search, Indexing & Vector DB (16; plant: compute_pool, server)

index_workers [capacity], index_refresh_rate [capacity], segment_merge_rate [capacity], search_concurrency [admission], search_timeout [admission], shard_count [capacity], shard_replication [capacity], shard_rebalance [capacity], vector_search_k [capacity], vector_batch_size [capacity], embedding_workers [capacity], index_memory_budget [admission], query_route [capacity], hot_shard_isolation [admission], search_admission [admission], circuit_breakers [admission]

Service Mesh & API Reliability (16; plant: compute_pool, server)

circuit_breaker [admission], retry_budget [admission], service_timeout [admission], service_concurrency [admission], connection_pool [capacity], traffic_divert [capacity], traffic_mirror [capacity], canary_weight [setpoint], outlier_ejection [admission], health_threshold [admission], dns_traffic_weight [setpoint], session_affinity [capacity], request_hedging [capacity], fault_injection_gate [admission], service_failover [capacity], load_balancing_policy [capacity]

Telecom RAN & Edge Radio (12; plant: compute_pool, ran)

ran_connection_admission [admission], ran_ue_handover [capacity], ran_cell_traffic_steering [capacity], ran_slice_resource_budget [admission], ran_prb_allocation [capacity], ran_scheduler_weight [setpoint], ran_tx_power [power], ran_antenna_tilt [capacity], ran_carrier_enable [capacity], ran_cell_sleep [capacity], ran_du_cu_placement [capacity], ran_fronthaul_budget [admission]

Workflow, Logistics & Fulfillment (15; plant: compute_pool, workflow)

workflow_admission [admission], workflow_worker_rate [capacity], task_queue_rate [capacity], workflow_retry [admission], workflow_backoff [admission], workflow_timeout [admission], inventory_allocation [capacity], fulfillment_route [capacity], warehouse_queue [capacity], carrier_selection [capacity], dispatch_priority [admission], shipment_batch [capacity], route_replan [capacity], sla_escalation [capacity], compensation_action [capacity]

Organism 5: the four stacked, every duplicate kept, 1,716 muscles**Organism 6: the whole tower, all 945 muscles once**

Evaluation and simulation use only. Copyright (c) 2026 The Omni-Compass LLC. Commercial use, commercialization or monetization of any part of Omni-Compass requires a signed, paid Omni-Compass Enterprise License from The Omni-Compass LLC. Patents, copyrights and trademarks filed in the USA. See LICENSE and NOTICE at the root of this repository.

CHAPTER 23

The Domain Map

Every system Omni-Compass can sit on, what it reads, what it pushes, the reflexes that bound it, and where it stands.

1. The connector pattern

Omni-Compass is the same engine on every muscle. Each muscle gets one connector with five parts:

Part	Role	Example (compute)
Afferent (sense)	What the muscle reports	load, pending work, power, heat
Engine	Equations (1)-(7) turn readings into state: error E, coherence U, pressure I_U, stress S, bath B	same six-state engine everywhere
Efferent (push)	The equation (2) law and allocation law drive the muscle toward convergence	node count, HPA target, power cap
Reflex (shield)	Hard limits checked before any push	never below pod requests, step limits, power limit
Kill	Control returns to the muscle's own controllers	restore HPA targets, observe only

Fit tiers used below:

- **Direct:** Omni can govern the muscle as a supervisory controller.
- **Supervisory:** Omni governs budgets, schedules and envelopes; the muscle's own real-time controller stays in charge of fast inner loops.
- **Advisory:** regulated or safety-critical; Omni recommends until certified (IEC 61508, DO-178C, ISO 26262, medical and financial regulation).
- **Not Omni's problem:** a different kind of problem; stated so no one claims otherwise.

Status: **Wired** (in code, verified), **Partial**, **Open** (connector not yet built).

2. Digital infrastructure

Muscle	Reads	Pushes	Reflexes	Fit	Status
Compute nodes and node pools	load, pending pods, requests	node count, park or power off	request floor, step limit	Direct	Wired
Pods via HPA	utilisation, replicas	HPA target	bounds 50-95%	Direct	Wired
Power caps	watts, site limit	per-node caps	site power limit (I4)	Direct	Wired (harness)
Heat	temperatures	feeds decisions	thermal limit	Direct	Sensed
GPUs / accelerators	utilisation, power, temperature, memory	power caps, clocks, partitioning (MIG), packing	temperature, job safety	Direct	Open
CPU power states	per-socket power (RAPL)	frequency scaling, sleep states	latency floor	Direct	Open
Memory	used vs requested	request right-sizing	OOM margin	Direct	Open
Storage	I/O load, latency, capacity	tiering, placement, volume scaling	durability, replica counts	Supervisory	Open
Network, load balancers	traffic, congestion, latency	routing weights, traffic shift	capacity per link	Direct	Partial
CDN / edge	cache hit, origin load	cache and origin routing	freshness	Direct	Open
Batch and HPC job queues	queue depth, deadlines	admission, priority, time shifting	deadline guarantees	Direct	Open
Multi-cluster / multi-region	load, cost, carbon per region	placement	data residency	Direct	Open
Deployments	release health	pause, roll back	never skip required checks	Direct	Partial (stack harness)
Databases, stateful services	load, replication lag	read replicas, pools	consistency limits	Supervisory	Open
Control plane health	API latency, etcd load	rate of Omni's own actions	never overload the API	Direct	Open
Tenant quotas	usage per tenant	who may grow	fairness bounds	Direct	Open
Hardware health	disk and memory errors	early drain of failing machines	availability floor	Direct	Open
Observability pipelines	data volume, cost	sampling, retention	required audit data	Direct	Open
Security posture	policy blocks	block expansion during holds	I1	Direct	Partial (sensed)

3. AI

Muscle	Reads	Pushes	Reflexes	Fit	Status
Training clusters	GPU power, progress, network	power caps, packing, checkpoint timing	never kill a run mid-step	Direct	Open (gpu vessels in harness)
Training power smoothing	synchronised power swings	ramp limits, staggering	grid ramp limits	Direct	Open
Inference serving	request rate, latency, queue, GPU memory	replicas, batch size, GPU sharing	latency SLO	Direct	Open
Model routing and cost	cost per request, quality signals	route to smaller or larger models, token budgets	quality floor	Direct	Open
Model releases	evaluation results, drift, errors	gate, pause, roll back a model	release gates	Direct	Open
Data pipelines, vector databases	backlog, freshness	scaling, scheduling	freshness bounds	Direct	Open
AI agents: containment	actions, spend, resources, network use	CPU, memory, network and spend caps, tool permissions	hard caps, instant kill	Direct	Open
Carbon-aware AI	grid price and carbon intensity	when and where training runs	deadlines	Direct	Open

AI alignment: what Omni can and cannot do

Alignment has two parts, and only one is Omni's.

Part	What it means	Omni's role
Value alignment	The model itself wants and does what people intend: honest, safe, helpful behaviour learned in training	Not Omni's problem. A governor outside the model cannot make the model's goals or judgement correct.
Control and containment	Whatever the model wants, it can only act within bounds: resources, permissions, spend, network reach, a kill that works	Direct. This is a governance problem: sense what the agent is doing, bound what it may spend and touch, revoke instantly, keep an audit trail.

Omni-Compass can be the boundary layer AI runs inside. It cannot be the thing that makes AI trustworthy. Claims must keep that line.

4. Facilities and energy

Muscle	Reads	Pushes	Reflexes	Fit	Status
Cooling plant	inlet temperatures, chiller load	setpoints, fan curves, chiller staging	inlet temperature limits	Direct	Open
Building HVAC	zone temperatures, occupancy	setpoints, schedules	comfort bounds	Direct	Open
Facility power distribution	PDU and UPS load	limits for the shield	breaker limits	Supervisory	Open
On-site batteries and generators	state of charge, grid state	charge and discharge	reserve floor	Supervisory	Open
Power grids: demand response	grid signals, prices	shift or shed flexible load	contractual limits	Direct	Open
Microgrids, renewable smoothing	generation, storage, load	dispatch	frequency and voltage limits	Advisory to supervisory	Open
EV charging fleets	vehicle needs, grid capacity	charge schedules	departure deadlines	Direct	Open
Water and pumping	demand, pressure, energy price	pump schedules	pressure limits	Supervisory	Open
Grid frequency and protection	frequency, faults	none	certified protection systems	Advisory	Out of scope for control

5. Physical operations

Muscle	Pushes	Fit
Manufacturing lines	scheduling, energy, throughput balance	Supervisory
Warehouses and logistics	task allocation, fleet charging	Direct
Robot fleets	task assignment, charging, traffic, safety envelopes	Supervisory (never joint-level control)
Traffic signals, transit, rail scheduling	timing and schedules	Advisory (safety-certified)
Ports, shipping	berth and crane scheduling, energy	Supervisory
Agriculture, greenhouses, irrigation	water and climate setpoints	Direct
Mining, oil and gas	energy and scheduling only	Advisory (process safety)
Semiconductor fabs	tool scheduling, energy	Supervisory

6. Science, space, biology, finance, health

Domain	What Omni could govern	What it must not touch	Fit
HPC / supercomputers	job queues, power, cooling	none beyond the facility	Direct
Quantum-computer facilities	cryogenics energy, calibration scheduling	qubit control	Supervisory
Satellites and constellations	power, thermal, data budgets, scheduling	attitude and orbit control (flight software)	Supervisory
Spacecraft / astrodynamics	mission resource budgets	guidance and navigation	Advisory
Biology: lab automation, bioreactors	setpoints, schedules, cold chain	living-system interventions without validated models	Supervisory
Financial markets	risk limits, exposure caps, trading infrastructure	making trades, market control	Advisory (regulated)
Hospitals	bed, staffing and energy operations	medical devices, treatment	Advisory
Aviation, automotive, nuclear	monitoring	any direct control	Advisory (certification required)

7. The problems everyone in computing has, and Omni's fit

Problem	What Omni can do	Fit
Data-centre energy growth	right-size, park, cap, shift work in time and place	Direct
Cloud cost overruns and idle capacity	consolidate, right-size, remove padding	Direct
GPU scarcity and low GPU utilisation	pack jobs, share GPUs, schedule by priority	Direct
Heat and cooling limits	govern IT load and cooling plant together	Direct
Grid connection limits for new data centres	power smoothing, demand response, stay under contracted power	Direct
Carbon reporting and reduction	carbon-aware placement and timing, measured energy	Direct
On-call burnout and alert fatigue	hold the levers people turn at night; fewer, better alerts	Direct
Autoscalers fighting each other	single authority	Direct (proven in simulation)
Latency and SLO breaches under load	headroom governance, fast scheduling floor	Direct
Noisy neighbours, unfair tenants	quotas and growth permissions	Direct
Hardware failures	early drain on failure signs	Direct
Observability cost and data explosion	sampling and retention governance	Direct
Configuration drift	detect and gate, not rewrite configurations	Supervisory
Cold starts	pre-warming policy	Direct
Data egress costs	placement near data	Supervisory
Runaway AI agents and spend	caps, permissions, kill, audit	Direct
Security misconfiguration, supply-chain attacks	block risky expansion during holds	Partial (not a security scanner)
Model hallucination and bias	none	Not Omni's problem
AI value alignment	none	Not Omni's problem
Software bugs, technical debt, legacy migration	none	Not Omni's problem
Vendor lock-in	a single governor over mixed clouds	Supervisory

8. Order of wiring

1. Digital infrastructure: GPUs, CPU power states, memory, storage, network, job queues, multi-region.
2. Facilities and energy: cooling plant, power distribution, batteries, demand response.
3. AI: training power smoothing, inference serving, model routing, agent containment.
4. Physical operations, then science and space as supervisors.
5. Regulated domains as advisors, until certified.

Each connector follows the same path as compute: wired, benchmarked against today's controls, verified, pre-registered, tested on held-out scenarios, then released.

Evaluation and simulation use only. Copyright (c) 2026 The Omni-Compass LLC. Commercial use, commercialization or monetization of any part of Omni-Compass requires a signed, paid Omni-Compass Enterprise License from The Omni-Compass LLC. Patent applications, copyright registrations and trademark applications filed in the United States. See LICENSE and NOTICE at the root of this repository.

CHAPTER 24

The Problem Map

Method: the issue trackers of the projects that run the world's clusters (Kubernetes autoscaler, Karpenter, Kepler, Kueue, Volcano, Knative, KEDA, kubernetes/kubernetes), the public industry reports (Cast AI, Datadog, Uptime Institute, IEA) and recent papers (Meta Llama 3, AI-datacenter power stabilisation). Not all of GitHub: a targeted scan of where the money and the failures are. Status: **have** = in this repository and tested; **partial** = built but not proven live; **missing** = not built.

#	Problem (evidence)	Size	Best tool today and its gap	Omni-Compass fit	Status
1	Idle capacity. Average Kubernetes CPU utilisation 8%, memory 20%, falling; CPU over-provisioning 69%, memory 79% (Cast AI 2026). 83% of container cost goes to idle resources (Datadog)	largest money leak in cloud	Karpenter, Cluster Autoscaler, CAST AI: each fixes nodes only	one authority over replicas, requests, nodes and power	built (omnilab/rightsize.py, held-out)
2	HPA and VPA cannot be used together on CPU/memory , "not compatible by design" (kubernetes/autoscaler #1726, #2939, #6060, #6247; coordination asked again in #8493, 2025)	blocks right-sizing for most teams	none; official advice is "don't"	exactly what a single authority solves: one engine owns both numerator and denominator	built (omnilab/rightsize.py: HPA+VPA vs one authority, held-out)
3	Karpenter consolidation churn: nodes replaced every 5-10 minutes for 2-3 generations, busy nodes removed instead of empty ones, one shared timer (karpenter #1851, #1019, #2705, #3046; provider-aws #7146, #8868, #7356, #8536)	outages, wasted boots	Karpenter's own timers and budgets	engine-gated release (equation 2) plus dwell; C-throughput cut node churn 50% in simulation	built in simulation (AKS NAP = Karpenter arm, tuning/vendor_compare.py); live Karpenter opponent next
4	CPU limits throttle apps (CFS quota), causing latency and even OOMs (kubernetes/kubernetes #67577, #97445)	latency tails	manual tuning, "remove limits" advice	power-cap muscle must watch latency; our first live run hit this exact problem	have; live runs 3-5 found and fixed the resize permission
5	Cold starts / scale-from-zero latency (knative/serving #4902, #14202, #9104; keda/core/http-add-on #219)	seconds of latency per edge	Knative activator, KEDA	pre-warming decided by the engine from its state (anticipation)	built (omnilab/coldstart.py, held-out)

#	Problem (evidence)	Size	Best tool today and its gap	Omni-Compass fit	Status
6	GPU utilisation ~5% (Cast AI 2026); idle GPUs scattered across nodes (volcano #3948, kueue #5243)	the most expensive silicon idle	KAI Scheduler, HAMi, nvshare, Volcano binpack	GPU power-limit muscle (have, calibrated), packing and sharing (missing)	built (omnilab/gpupack.py, held-out) + power limit
7	AI training power swings: 50-75% of TDP in milliseconds per GPU, ramps over 1000 MW/s at gigawatt scale; operators impose power and ramp-rate limits (SemiAnalysis; Uptime; arXiv 2508.14318, 2606.04869)	grid-connection risk, equipment damage	batteries, fast PSUs, manual GPU power caps	"training power smoothing" muscle: ramp-rate limit through GPU power caps, the engine's bath state B is literally a damped oscillator	built (omnilab/powersmooth.py, held-out; basin-only arm recorded)
8	GPU failures and stragglers: Llama 3, 419 interruptions in 54 days on 16K H100s, one every ~3 hours; 58.7% GPU-related; slow stragglers undetected (Meta; Lablup 2605.09370)	lost training time	Meta's internal tools, checkpointing	hardware-health muscle: sense straggling, drain early	built (omnilab/health.py, held-out)
9	Outages: 54% cost more than \$100,000, 1 in 5 more than \$1M; power causes 45% of incidents (Uptime 2025)	direct money	SRE runbooks, AIOps	recovery and physical-SLA results in simulation: recovery 58 to 17 min	have (sim)
10	Cooling: industry PUE stuck at ~1.54 for six years (Uptime 2025)	~35% overhead on every watt	DCIM, DeepMind cooling AI	heat afferent have; cooling-plant muscle	built (omnilab/cooling.py, held-out)
11	Energy measurement in VMs is impossible: no RAPL/IPMI in VMs, Kepler building a model server (kepler #2487, 2026)	nobody can prove savings in cloud VMs	Kepler estimators	same gap we hit	built (omnilab/vmenergy.py: meter-calibrated attribution, held-out)
12	LLM inference autoscaling: KV cache fills memory before compute; prefill vs decode need separate scaling (vLLM, llm-d)	latency and GPU cost	llm-d, KServe, custom	inference muscle: scale on KV-cache use and queue, not CPU	built (omnilab/inference.py, held-out)
13	Runaway AI agents: loops and recursive calls, ~\$10,000 overnight examples (Dark Reading; sandbox guides)	cost and safety	per-tool caps, early kill-switch projects	agent-containment muscle: caps, kill, audit	built (omnilab/containment.py, held-out); live version next
14	Human error: failures to follow procedures rose 10 points (Uptime 2025)	outages	runbooks	fewer manual actions: pages and human interventions to zero in simulation	have (sim)

Results of the build pass (held-out seeds 90001-90030, 30 per muscle, code frozen by SHA first)

Omni-Compass against the strongest native tool in each muscle; + better, - worse, = tie / not significant.

#	Muscle	Opponent (best native)	Better	Worse	Tie / n.s.
1-2	right-sizing	VPA in-place + HPA	CPU -3%, memory -11%, p99 -99%, reversals -84%, SLO minutes -57%	p95 +15%, OOM kills (few)	work done
5	cold start	KEDA HTTP	p95 -75%, p99 -58%, delayed >1 s -71%, instance-hours -8%	cold-start count (x8)	
6	GPU packing	Volcano binpack	idle powered GPU-hours -20%, energy -2%	mean wait +16% (seconds)	p95 wait, jobs, fragmentation
7	training power	vendor floor, worst-day setting	burn -9%, throughput loss -44%, 1-s swing -6%	max ramp (inside the grid's limit)	0 grid violations both
8	GPU health	threshold detection	goodput +10%, lost GPU-h -20%, restarts -29%, straggler hours -99%		checkpoint overhead
10	cooling	outdoor-air reset	PUE -2%, cooling -13%, inlet violations -73%		max inlet
11	VM energy	host meter split by CPU	per-VM error -37%, worst VM -34%		total
12	LLM inference	KEDA on queue	TTFT p95 -18%, SLO breaches -58%, GPU-hours -19%, preemptions -56%		TPOT, TTFT p50
13	agent containment	static caps	rogue spend -96%, time to contain -99%, peak sub-agents -59%	false stops (0.6 -> ~2.7/day), honest work -7%	forbidden actions (0 both: RBAC)

Engine vs no-engine: in every muscle the omni_no_engine arm scores close to omni. The mapping from engine state to action carries most of the effect; the evolved dynamics add smoothing. The basin-only arm (power smoothing, the bath equation alone setting the draw) cut ramps 63% but did not hold the grid limit: the human-set limit is required.

What this says

- The problems with the most money behind them are **idle capacity (1)**, **HPA/VPA coordination (2)** and **GPU idleness (6)**. Problem 2 is the cleanest proof of Omni-Compass's thesis: the Kubernetes project itself says its two autoscalers cannot share a metric, because they fight. A single engine owning both is the answer the issue tracker keeps asking for.
- The problem with the most strategic weight is **AI training power swings (7)**: it is new, it threatens grid connections, and Omni-Compass's own equation (7) is a damped second-order bath, the natural controller for ramp limits.
- The problems we already hit ourselves (4, 11) are the same ones the industry has: good evidence the benchmark is real.

Build order proposed

- Request right-sizing muscle (VPA role) under the same engine, benchmarked against HPA + VPA "in conflict".
- Training power-smoothing muscle (GPU power-cap ramp limits) with a synthetic synchronized-training power trace.

3. GPU packing / fragmentation muscle (with fake-gpu-operator or KWOK on the live cluster).
4. Live Karpenter opponent (kwok provider) to test problem 3 head to head.
5. Inference muscle (KV-cache and queue driven), then hardware-health (straggler drain), then agent containment.

Sources

- kubernetes/autoscaler issues #1726, #2939, #6060, #6247, #8493
- Karpenter #1851, #1019, #2705, #3046; aws/karpenter-provider-aws #7146, #8868, #7356, #8536
- kubernetes/kubernetes #67577, #97445
- Knative #4902, #14202, #9104; KEDA http-add-on #219
- Volcano #3948; Kueue #5243; KAI Scheduler; HAMi
- Kepler #2487
- Cast AI 2026 State of Kubernetes Resource Optimization; Cloud Native Now on the report
- Uptime Institute annual outage analysis 2025; Uptime global survey 2025
- SemiAnalysis: AI training load fluctuations; Uptime: AI power fluctuations; Power Stabilization for AI Training Datacenters; Source-side mitigation
- The Llama 3 Herd of Models; Tom's Hardware on Llama 3 failures; Lablup 504-GPU report
- vLLM anatomy; llm-d 0.5
- Dark Reading: AI agents and runaway costs

Evaluation and simulation use only. Copyright (c) 2026 The Omni-Compass LLC. Commercial use, commercialization or monetization of any part of Omni-Compass requires a signed, paid Omni-Compass Enterprise License from The Omni-Compass LLC. Patent applications, copyright registrations and trademark applications filed in the United States. See LICENSE and NOTICE at the root of this repository.

CHAPTER 25

The Six Organisms and the Benchmark Grid

Omni-Compass is benchmarked on six organisms, each run native and with Omni-Compass on top on the same seed, the same load and the same clock: Compute / AI / Cloud (430 muscles), Physics / Robotics / Autonomous (376), Energy / Facility / Industrial (470), Distribution / Specialized (440), the four stacked with every duplicate kept (1,716), and the whole tower with every muscle once (945); these are the Omni v2 counts, v1's were 345, 262, 282, 337, 1,226 and 656.

Each organism is run at 1, 10, 100 and 1,000 paired runs and at 1, 10, 100 and 1,000 copies of the organism on one clock. Every receipt shows work per energy, work, energy, the time spent outside the service line, and whether every knob was handed back.

What the grid has shown. On Omni v3, 84 of the 90 judged cells are done (`results/scale/GRID.md`): every organism reads superior within its guardrails in every cell of ten runs or more, with work per energy +0.07% (the Physics realm) to +0.37% (the Energy realm, the four stacked and the tower), the work unchanged, the time over the line not above native's, and every knob handed back. The figure is the same at every size, from one copy to a thousand, which is the property the grid exists to test: one brain stays coherent at 1.7 million muscles. The six cells left, 100 and 1,000 runs at 1,000 copies, are beyond the machines available and are declared, not hidden. The number is small because the native controllers in the models are well tuned and leave little room; it is real within the model, and it is never counted in the Omni index, which is made of real stacks only.

On a GPU machine the real card is wired into each organism as one more muscle of its NVIDIA GPU family; its own power meter and its requests are counted in the organism's receipt, kept apart from the modelled plants and also added to them. Real Kubernetes is benchmarked separately on real clusters, and it is also wired into the organisms as one more muscle (`tools/run_kil.py`): at 1 to 100 copies on GitHub's machines and at 1,000 copies on a rented machine, the organism's own compute demand drives a real load generator against a real service on a real cluster, and the cluster's watts come back into the organism as heat and load. Those runs are labelled "L + S": the cluster is live software, the organism around it a model.

The six organisms are models (evidence class S). They show how the law behaves across hundreds of kinds of machine at once and whether its effect is stable as the count and the size grow. Real Kubernetes, the real database, the real broker, the real cache and the real card are the anchors measured on real software and, for the card, a real meter; their results and their losses are in the manual's section 16.

CHAPTER 26

The Metrics Catalog

Every gauge Omni-Compass produces, where it comes from, and what kind of number it is. Three kinds:

- **Measured:** read from a real system (a real Kubernetes cluster, a real GPU's own meter, a real wall plug).
- **Modelled:** computed by a simulation from declared physics or recorded traces. It shows the mechanism, not a measurement.
- **Internal:** Omni-Compass's own state and decisions, recorded in its audit log on every decision, live or simulated.

Every report says which kind each number is. A modelled number is never presented as a measured one.

1. Service: what the customer feels

Gauge	Unit	Where	Kind
Response time, mean / 95th / 99th percentile	ms	live Kubernetes (<code>tools/live_reps.py</code>), GPU bench (<code>tools/gpu_reps.py</code>), GPU sims	measured live; modelled in sims
Failed requests	%	live Kubernetes	measured
Requests served / not served	count	GPU bench, GPU sims	measured / modelled
Pod start wait, total and mean	s	live Kubernetes (API server creation-to-Ready)	measured
Pending pods	pod-minute s	live Kubernetes	measured
Backlog (SLO breach) minutes	min	site and node exchange sims, stack benchmark	modelled
95th-percentile latency factor	×	site and node exchange sims	modelled
Availability, time healthy, recovery minutes, recovered	share, min	stack benchmark (<code>benchmarks/stack_benchmark.py</code>)	modelled
SLA violations (physical, total, backlog, power, heat)	share of time	stack benchmark	modelled

2. Work and capacity: what the operator gets for the money

Gauge	Unit	Where	Kind
Work per energy (served requests per kJ)	req/kJ	GPU bench (the preregistered primary outcome), GPU sims	measured / modelled
Work per wall energy (whole machine)	req/kJ	GPU bench with a smart plug (tools/wall_meter.py)	measured
Work served	work units	site and node exchange sims	modelled
Work per kWh	units/kWh	node exchange sim	modelled
Energy per served request	J	GPU bench	measured
Energy per core-hour	Wh	live Kubernetes	declared model
CPU used by the app	cores	live Kubernetes (kubectl top)	measured
Utilisation (used / allocatable)	share	live Kubernetes	measured
Replicas (HPA), mean	count	live Kubernetes	measured
Pods started	count	live Kubernetes	measured
Worker nodes in service, node-hours, idle node-hours	count, h	live Kubernetes, stack benchmark	measured / modelled

3. Energy and power: the bill and the building

Gauge	Unit	Where	Kind
GPU energy (the device's own power.draw, integrated)	J	GPU bench	measured
GPU mean power	W	GPU bench, GPU sims	measured / modelled
CPU package energy (RAPL)	J	GPU bench, where the machine exposes RAPL	measured
Whole-machine energy at the wall	J	GPU bench with a smart plug	measured
Energy on kind	Wh	live Kubernetes	declared model (kind has no meter)
Energy	kWh	site, node exchange and stack sims	modelled
Peak power	kW	site and node exchange sims, stack benchmark	modelled
Site-budget violation minutes (time over the building's limit)	min	site and node exchange sims	modelled
GPU power limit, mean; power-limit writes	W, count	GPU bench, GPU sims	measured / modelled
Power-cap travel	share	stack benchmark	modelled

4. Hardware: heat and wear

Gauge	Unit	Where	Kind
GPU temperature, peak and mean	°C	GPU bench, GPU sims	measured / modelled
Heat-over minutes	min	hardware plant sim (hardware/plant.py)	modelled
Thermal travel	share	stack benchmark	modelled
Machine round trips, node starts and stops, scale reversals	count	stack benchmark	modelled

Lower power and fewer temperature swings are the conditions for longer hardware life. **No hardware life extension is claimed** until real temperature, duty-cycle and wear records exist.

5. Battery equivalent (derived)

A battery holds a fixed amount of energy, so battery life follows directly from work per energy:

- runtime on the same charge, for the same work: $\times (\text{native energy} \div \text{Omni-Compass energy})$
- work on the same charge: $\times (\text{Omni-Compass work per energy} \div \text{native work per energy})$

Example: the modelled GPU card at +5.1% work per kJ (`results/gpu/sim/after`) does 5.1% more work on one charge. The same rule turns any measured work-per-energy result into a battery figure. On-site battery banks as an organ of the power budget are designed (`docs/DOMAIN_MAP.md`, "on-site batteries") and not yet built.

6. Safety and control: proof it behaved

Gauge	Where	Kind
Writes executed per arm (native 0, watch 0, Omni-Compass n)	GPU bench, live Kubernetes	measured
Every arm ended at its start setting (the reset restored it)	GPU bench, live Kubernetes switch drill	measured
Invariant violations (all, and excluding power)	stack benchmark	modelled
Security violations, contradictions, pages	stack benchmark	modelled
Freeze check: code hashes at start and end (confirmation runs)	GPU bench (<code>FREEZE.json</code> , <code>FREEZE_END.json</code>)	measured
SHA-256 of every raw file	every live run	measured
The seal: every Python/C++ twin unchanged since proven equal	<code>results/SEAL.json</code> , <code>verify.py</code>	checked on every build

7. Internal: Omni-Compass's own state, on every decision (audit log)

Field	Meaning
E, U, I_U, S, B, B_dot	the six-state engine: energy, coherence (health), unmet-need integral, stress, redistribution, its rate
state_observed, state_projected_next, prediction_error	what the device said, what the engine projected, and how far off the last projection was
requested_cap, granted_cap, shield_bound, want_w	what the engine asked for, what the shield allowed, which bound decided it, the limit written
admissible	whether the engine was permitted to change anything this decision
speed_lock(ratios, line, aim, rate)	the speed lock's reading against the run without Omni-Compass
convey engaged / released	when idle CPU is conveyed to serving pods, and when it is given back
authority: calm and its scalars (kappa, h, sigma, nu), execute, contract per organ, and the machine organ's own view	how settled the engine is, whether it may act at all, and which organs may give capacity back this decision (Kubernetes controller; <code>omnicompass/nervous_system.py</code>)
compass readings	the direction the engine reads the whole system to be moving (<code>omnicompass/compass.py</code>)
write, would_write, why	every command sent (or, in watch mode, withheld) and its reason
snapshot, restored	every setting read at start, and its read-back after the reset

8. Where each report lives

Report	Command	Output
Live Kubernetes, paired	<code>benchmark-reps.yml</code> (commit with [reps])	<code>LIVE_REPS.md</code> , <code>results/live/</code>
GPU bench on a real card	<code>sudo bash scripts/gpu_paired.sh</code>	<code>results/gpu/run-*/GPU_REPS.md</code>
GPU physics sim	<code>python3 tools/gpu_physics_sim.py</code>	<code>results/gpu/sim/</code>
CPU-then-GPU speed lock sim	<code>python3 tools/gpu_pipeline_sim.py</code>	<code>results/gpu/sim/pipeline/</code>
Site power exchange (GPU groups)	<code>python3 hardware/site_exchange.py</code>	<code>results/hardware/SITE_EXCHANGE_*.json</code>
CPU + GPU on one budget	<code>python3 hardware/node_exchange.py</code>	<code>results/hardware/NODE_EXCHANGE_*.json</code>
Full stack benchmark and every check	<code>python3 verify.py</code>	<code>results/</code> , <code>VERIFICATION: PASS</code>

Evaluation and simulation use only. Copyright (c) 2026 The Omni-Compass LLC. Commercial use, commercialization or monetization of any part of Omni-Compass requires a signed, paid Omni-Compass Enterprise License from The Omni-Compass LLC. Patent applications, copyright registrations and trademark applications filed in the United States. See LICENSE and NOTICE at the root of this repository.

PART FIVE

The Harness and the Wiring

*The universal plug, the adapters, the wire check, and the step-by-step work of wiring
Omni-Compass onto a running stack.*

CHAPTER 27

The Plug, the Adapters and the Wire Check

Think of a high-end car stereo: one head unit, one standard plug on its back, and an adapter harness for each make of car that matches the car's factory plug. You never cut a factory wire, and when you pull the stereo, the car works as it did. Omni-Compass is wired the same way.

8.1 The plug (one design for every muscle). Each plug has as many channels as its muscle has knobs. Every channel declares its cover (range), its direction (which way is "more"), its units and how fast it may move. Every plug keeps the same contract (`omnicompass/compass Law.py`, `Plug`):

1. **attach** - read the knob once, before any write: the snapshot. It never moves afterwards.
2. **read** - the service reading, into the compass.
3. **write** - one value, clipped to the cover, then read back from the device.
4. **one writer** - if the knob is found at a value Omni-Compass did not write, someone else owns it: Omni-Compass stops writing and leaves that value alone.
5. **restore** - on stop, the knob returns to the snapshot (where Omni-Compass found it, not the last value it wrote). For muscles whose native controller moves the knob itself, restore hands control back.

8.2 Adapters (one per standard plug). Omni-Compass speaks each industry's standard control interface through an adapter. Status is stated plainly:

Standard	Muscles	Adapter status
Kubernetes API	HPAs, deployments, pods, nodes	built (<code>omni_controller/controller.py</code> , <code>muscles.py</code>)
NVIDIA NVML / <code>nvidia-smi</code>	GPU clock ceiling, power limit	built (<code>omni_controller/gpu_compass.py</code> , two wires; <code>gpu_governor.py</code> , one wire)
Linux <code>cpufreq</code> , RAPL	CPU frequency ceiling, package watts	built (<code>--cpufreq-policy-root</code> , <code>--rapl-cmd</code>)
Site meter, building controller by command	site watts, supply-air setpoint	built (<code>--power-cmd</code> , <code>--cooling-cmd</code> command templates)
Cloud node groups (Karpenter, Cluster Autoscaler, MachineSet)	machines	built through <code>--node-scale-cmd</code> templates
BACnet, Modbus, SNMP	chillers, air handlers, PDUs, UPS	designed; today reached through the command templates above
OPC UA, EtherNet/IP, PROFINET, EtherCAT, ROS 2, CAN	PLCs, drives, robots, vehicles	designed; modelled in the realm harness
IEC 61850, DNP3, OpenADR, IEEE 2030.5, SunSpec, OCPP	substations, batteries, inverters, chargers	designed; modelled in the realm harness

The engineer's job for one muscle is three things: which adapter, the address of the knob, and its safe range.

The plug contract, line by line, and why each line is there. *Attach* reads the knob once before any write and keeps that value as the snapshot. It is read once and never again, because a snapshot that moved with the knob would restore to wherever the governor last left it, and the whole point of the restore is to return to where the operator had it. *Read* takes the service reading (a response time, a temperature, a queue, an error) into the compass as a position in its band; the plug declares the band, the law does not guess it. *Write* takes one value, clips it to the cover before it leaves the plug, sends it through the muscle's own interface, and then reads the lever back from the device and records both the sent and the taken value; a device that took something other than what was sent (a clock that rounds to 15 MHz steps, a pool size the pooler capped) is recorded as it is, not as it was asked. *One writer* compares the lever as found at every read with the value last written; a lever found elsewhere has another owner, and the plug raises *ForeignWriter*, after which the governor leaves that lever alone for the rest of the run and says so in its audit and its exit code. *Restore* writes the snapshot back on stop and reads it back; a restore that does not read back is an error (exit 3), never a silent success. The same five lines are implemented for Kubernetes objects (annotations hold the snapshot on the object itself, so a watchdog can restore from the object without the governor's memory), for *nvidia-smi*, for *PgBouncer*'s console, for *Redis*'s *CONFIG SET*, for the consumer group and for the simulators' knobs, and *tests/test_** prove each against a fake device.

8.3 Power-up order. Read-only first; then take the knobs, the least consequential first; on stop, hand them back in reverse order. In the Kubernetes controller the order is: the latency feed and the cluster's own metrics (read); the HPA target (a setting the HPA accepts and corrects within a minute); the pods' CPU limits (conveyance, in place); the machines (the slowest and most consequential, through the release gate and the verdict). On stop the machines are woken first, then the limits, then the target, so the cluster is never left with fewer machines than its pods need while a setting above them is still being handed back.

8.4 The wire check (mandatory before any write). For the GPU it is one command, `sudo python3 tools/gpu_wire_check.py --gpu 0`, and `scripts/gpu_rented_run.sh` runs it first and stops if it fails:

Step	It proves
1 read	every meter answers: watts, temperature, utilization, limit, clock, top clock
2 snapshot	the start limit is recorded once
3 up wire	with the card busy, a ceiling below its own busy clock pulls the clock down under it; reset, the clock comes back above it
4 down wire	a lower power limit is reported back exactly; the start limit is reported back exactly
5 restore	clocks reset and the start limit read back
6 stop	the governor, stopped by its signal, hands both wires back
7 other writer	a limit set by someone else while the governor runs is left alone, and the governor exits 5

Any failure prints the wire, the step and what the device said, and nothing else runs. That turns a wiring fault from guesswork into one named line to fix.

8.5 Wired right or wired wrong. The wire check proves the wires move. It does not prove the governor reads the right thing. A correct installation leaves signatures in its own receipts; a wrong one leaves others. Check them on your first paired runs, before you believe any number, good or bad.

Check	Wired right	Wired wrong, and what to look at
Watch arm against native	equal on every gauge (watch writes nothing)	different: the probe, the load or the machine differs between arms, not Omni
Requests served, failed requests	equal to native	fewer served or more failed: a lever is cutting capacity; check its cover and its sign
p95 and p99 response time	at native or faster	slower: the reading or the floors are wrong (rows below)

Check	Wired right	Wired wrong, and what to look at
GPU: the card's clock while busy	at or above the clock the card reaches on its own	below it: the speed floor is missing or the lid is under the card's own draw
GPU: the lid while busy	at or above the card's own busy draw	at the envelope floor while busy: the lid is sized from a curve, not from the card's own meter
GPU: credit per write (GPU_REPS.md)	compass decisions pace the slack, race decisions cover the bursts; fail-up is rare	fail-up in most decisions: being busy is being read as a breach, or the service line is set too low
Kubernetes: machines given back	through the release gate, one per decision, with no pod waiting	machines given back while pods wait, or none at all while the service is far inside its band: check the latency feed, - - s10-ms and the release gate's reasons in the audit
Kubernetes: the reset	every HPA target, replica range, CPU limit and worker back to native, no record left	anything left: an earlier run or another controller wrote the same objects
Every lever after the run	at its snapshot	not at its snapshot: another writer, or a restore that failed (exit 3): restore by hand and investigate
The feed itself	fresh samples every few seconds	stale or empty: Omni reads it as blind and fails up; fix the feed, not the governor

The rule. If any row reads "wired wrong", the result of that run says nothing about Omni-Compass: fix the wiring, run the wire check and the watch arm again, and only then compare. The first real card run is the worked example: its receipts showed the card's busy clock under its own (736-768 MHz against 861-889), the lid at the 105 W floor under its own 135 W draw, and fail-up in 46% of decisions. Each of those is a row in this table.

8.6 The wire check on the other stacks. The GPU has a one-command wire check because its wires are the hardest to get right. The other stacks carry the same proof in a different place: every Kubernetes arm records `kubectl auth can-i` receipts for what the governor's identity can and cannot do (and fails if it can do anything it must not), checks that the cluster holds nothing but the harness, and ends with the reset check that every HPA target, replica range, CPU limit and worker is back at native with no record left; every database, broker and cache arm reads the knob back after every write and hands it back at the end, and the report says "handed back: yes" or the run is invalid; every simulator arm checks that every override was handed back. The watching arm, where it is run, is the wire check of the reading: a watching governor that disagrees with native on any gauge has found a difference in the probe, the load or the machine, not in Omni-Compass.

CHAPTER 28

The Harness

Before you wire anything: read DISCLOSURES.md. Omni-Compass acts only through the wires it is given; it cannot be slapped on. If your paired receipts differ from the published benchmarks in direction, the first presumption is wiring: confirm the installation with section 8.5 of the manual (*Wired right or wired wrong*).

This tree is the GitHub-ready kit: frozen engine + C++ twin + two plants + live stub.

What to run

```
pip install -r requirements.txt
python verify.py --quick          # engine, C++, soak, parity
python k8s_controlplane/test_controlplane.py
python -m k8s_controlplane.suite # elastic / always_on / idle_power
python tests/test_hpa_three_way.py
python tests/test_omni_controller.py
```

Fleet plant (Omni as node authority vs HPA+CA / Karpenter-lite):

```
python -m fleet.planetlab --dir fleet/traces/planetlab --scenarios 8 --out /tmp/pl
```

C++:

```
cmake -S cpp -B cpp/build -DCMAKE_BUILD_TYPE=Release &&& cmake --build cpp/build -j2
./cpp/build/oc_smoke
```

Two plants (do not mix the tables)

Tree	What Omni is	What the energy number means
k8s_controlplane/	On top of HPA+CA (target / gate / park)	Pack and optional CA gate on a 20-node replica
fleet/	Node-pool authority; CA off	Consolidation vs CA and Karpenter-lite

Observe must match the native arm on that plant. If it does not, the run is invalid.

Engine species

Shipped omnicompass/core.py is the patent principal form: cubic $\sqrt{U(1-U^2)}$, FIG. 4 command, RK4 with held $\sqrt{u}$. omnicompass/pools.py is actuation only (off / hold / park). It does not change the field.

Not in this harness

Live kube-controller-manager, kind CI, GPU MIG scheduler, facility cooling plant. omni_controller/ is observe-first against kubectl; tests use tests/fake_cluster/kubectl.

See `LIMITS.md`.

Evaluation and simulation use only. Copyright (c) 2026 The Omni-Compass LLC. Commercial use, commercialization or monetization of any part of Omni-Compass requires a signed, paid Omni-Compass Enterprise License from The Omni-Compass LLC. Patent applications, copyright registrations and trademark applications filed in the United States. See LICENSE and NOTICE at the root of this repository.

CHAPTER 29

The Wiring Guide

Before you wire anything: read DISCLOSURES.md. Omni-Compass acts only through the wires it is given; it cannot be slapped on. If your paired receipts differ from the published benchmarks in direction, the first presumption is wiring; confirm the installation with section 8.5 of the manual (*Wired right or wired wrong*).

The current manual is docs/INTEGRATION_MANUAL.md (every stack, every level, every switch). This page is kept for reference; where the two differ, the integration manual is current.

The full step-by-step manual, with the switch, the living band, parking and the pod reflex, is docs/OPERATOR_MANUAL.md. This page is the short version.

Three stages. Each stage adds only the permissions it needs. A stage is promoted only after its evidence is in, and the reset works at every stage:

```
kubectl -n omni-compass exec deploy/omni-compass -- touch /tmp/omni.kill
```

The switch is one human-operated switch for the whole harness (OFF restores every HPA target, replica range, CPU limit and cordon Omni-Compass changed, and records it; removing the file turns Omni-Compass back ON). It never fires by itself and never switches off one module: a failed decision is skipped, boundaries are held by the living band.

0. Build the image

```
docker build -f deploy/Dockerfile -t <registry>/omni-compass:<tag> .  
docker push <registry>/omni-compass:<tag>
```

What the image contains.

- The engine.
- The nervous system.
- The live controller and its levers.
- The frozen closure-law setting (/app/law/closure.json).

How it runs. As a non-root user, with a read-only root filesystem and no Linux capabilities.

1. Shadow: read-only, decides and logs, never writes

```
kubectl apply -f deploy/install/omni-compass.yaml      # set the image line first  
kubectl -n omni-compass logs deploy/omni-compass -f
```

Proof of read-only. scripts/pilot_shadow.sh runs the same stage from a workstation with kubectl auth can-i receipts that the identity cannot write, and writes SHADOW_REPORT.md.

Run length. One to two weeks.

Pass condition. Zero writes, and recommendations you agree with.

2. Target: Omni-Compass sets each HPA's CPU target

```
kubecttl apply -f deploy/rbac-target.yaml # adds: patch horizontalpodautoscalers
kubecttl -n omni-compass patch deploy omni-compass --type=json -p \
  '[{"op": "replace", "path": "/spec/template/spec/containers/0/args/1", "value": "target"}]'
```

What happens. The HPAs keep scaling pods. Omni-Compass only moves their target within bounds.

The SLO reflex. Add `--latency-file` and `--slo-ms`, and the target is never tighter than native while the SLO is breached.

Pass condition. Latency no worse than native, and fewer pod-hours.

3. Node pool: Omni-Compass sizes the node pool

Add permissions. patch nodes, create pods/eviction (deploy/kind/rbac-omni.yaml is the tested example).

Set the mode. `--mode nodepool --node-scale-cmd "<your pool resize command with {n}>"`. Examples of the resize command:

- a Karpenter NodePool limit;
- a cloud node-group size;
- `scripts/kind_nodepool.sh` on kind.

What guards a release. The machine organ gives a node back only when all of these hold:

- every sense is live;
- its last order landed;
- pods are not scaling up;
- nothing is pending;
- the remaining nodes stay at or below the engine's target utilisation.

A PodDisruptionBudget on each service is required, because drains go through the eviction API.

Optional levers. Each has its own flag and permissions, and all are listed in `omni_controller/muscles.py`:

- right-sizing;
- cold start;
- batch pacing;
- agent containment;
- cooling;
- CPU frequency.

Evidence, and where it is

Evidence	File
The controller, every lever and the Unified Control Switch on real Kubernetes (kind)	results/live/LIVE_LEVERS_2_NERVOUS.txt
Read-only shadow on real Kubernetes: every decision logged, 0 writes	results/live/LIVE_SHADOW_1.txt
Least-privilege identity receipts	rbac_omni.txt in every live run
Native vs Omni on top vs Omni alone, paired on one machine, real Kubernetes	results/live/LIVE_PAURED.md
The full wiring manual	docs/OPERATOR_MANUAL.md
The frozen engine's results, three runs each, read by rule	results/live/V1_*.md, results/live/V3_*.md; docs/OMNI_V1.md, docs/OMNI_V3.md
The same wiring on Azure's managed Kubernetes (the bill), a database pooler, a rented machine, the simulators, a message broker's consumer group, a cache's memory ceiling, a drone swarm's autopilot	the manual, section 10 (docs/OMNI_COMPASS_MANUAL.md)

Evaluation and simulation use only. Copyright (c) 2026 The Omni-Compass LLC. Commercial use, commercialization or monetization of any part of Omni-Compass requires a signed, paid Omni-Compass Enterprise License from The Omni-Compass LLC. Patent applications, copyright registrations and trademark applications filed in the United States. See LICENSE and NOTICE at the root of this repository.

CHAPTER 30

Before You Start, and the Eight Levels

What your system needs

You run	You need
Kubernetes (vanilla, EKS, GKE, AKS, OpenShift/OKD, Rancher, kind)	Kubernetes 1.34 or newer (in-place pod resize); metrics-server (<code>kubectl top nodes</code> works); an HPA with a CPU target on each governed service; a readiness probe and a short preStop pause on each service
A response-time feed (strongly recommended at every level)	a CSV per service, <code>elapsed_seconds</code> , <code>latency_ms</code> , <code>ok</code> , written continuously; <code>scripts/latency_probe.py</code> writes one against any HTTP endpoint
NVIDIA GPUs	a driver with <code>nvidia-smi</code> ; root (or the capability to run <code>nvidia-smi -pl</code> and <code>-lgc</code>); on a VM, full GPU passthrough
CPU power control	Linux <code>cpufreq</code> with the <code>schedutil</code> governor; RAPL for package watts; usually bare metal
Site power and cooling	a command that prints site watts, and a command that sets the supply-air setpoint through your building management system

Why each requirement is there. Kubernetes 1.34 or newer is required because the conveyance muscle resizes a pod's CPU limit in place, which older versions do only by restarting the pod, and a restart is a disruption Omni-Compass never causes. metrics-server is the HPA's own eyes and the governor's: without it neither knows what the pods use. An HPA with a CPU target on each governed service is the native controller Omni-Compass sits on; a service without one has nothing for level 2 to move. A readiness probe and a short preStop pause are what make a drain invisible to clients: the Service stops sending requests to a pod before it stops, and kube-proxy sends them to another ready endpoint. The response-time feed is the one wire a cluster does not already have: Kubernetes knows what its pods use, not what its users wait; the compass needs the second, and `scripts/latency_probe.py` writes it against any HTTP endpoint every five seconds. On the GPU, root is needed only for the two `nvidia-smi` writes, and a VM must pass the whole card through or the driver refuses to lock clocks (the wire check's step 3 is where that shows).

Get the software and check it

```
git clone https://github.com/The-Omni-Compass-LLC/The-Omni-Compass.git
cd The-Omni-Compass
pip install -r requirements.txt
python3 verify.py                                # every check; ends with VERIFICATION: PASS
```

Build the container image once:

```
docker build -f deploy/Dockerfile -t <your-registry>/omni-compass:<tag> .
docker push <your-registry>/omni-compass:<tag>
```

It runs as non-root, with a read-only root filesystem and no Linux capabilities.

The levels

Take them in order. Each has a pass condition; go to the next level only when it holds. The OFF switch works at every level.

The levels exist because a governor earns authority the way a new operator does: by watching first, then by touching the setting that is cheapest to undo, then the next. Level 0 touches nothing and proves the software is the software (`verify.py` checks every sealed file's fingerprint and runs every test). Level 1 watches a live cluster for as long as the operator likes and logs every decision it would have taken, so the operators can read those decisions against what they would have done; nothing passes level 1 until they agree with its readings. Level 2 moves the HPA target, which the HPA itself corrects within a minute if the governor is wrong, and raises replica floors ahead of bursts; it is the first level that can show a gain and the first that can show a loss, and the paired receipts of Chapter 39 are its pass condition. Level 3 moves machines, the slowest lever, through the release gate and the verdict; it is where the bill moves and where the most care is spent. Level 4 hands Omni-Compass the replica decision itself, with the HPA kept as the fast reflex upward; it is an option, not a step on the way, and the benchmarks in this manual do not use it. Levels 5 to 7 are the GPU, the CPU and the site, each with its own wire check and its own pass. An installation may stop at any level and run there indefinitely; most will stop at level 3.

Level 0 - Evaluate without touching anything.

```
python3 verify.py
python3 tools/run_scale.py --runs 10 --scale 1 --out results/scale/eval      # the six organisms, simulat
  ↪ ed
python3 tools/run_gpu_card.py results/sim/gpu_two_wire/eval                # the two-wire card, modell
  ↪ ed
```

Pass: VERIFICATION: PASS. Nothing in your systems is touched.

Level 1 - Watch (read-only).

```
kubectl apply -f deploy/install/omni-compass.yaml      # set the image line; it runs --mode observe
kubectl -n omni-compass logs deploy/omni-compass -f
```

The install file grants a read-only identity; `scripts/pilot_shadow.sh` records `kubectl auth can-i` receipts that show it cannot write. Every decision it would take is logged with its reason. Pass: zero writes, and readings your operators agree with.

Level 2 - The pods, on top of your autoscalers. WRITES.

1. Grant patch `horizontalpodautoscalers` and patch `Pods/resize` (`deploy/rbac-target.yaml`).
2. Run with `--mode target --latency-file <feed> --slo-ms <your p95 target>`.
3. Optional pod muscles, one at a time:

Muscle	Switch	What it does
convey	<code>--latency-file</code> and <code>--cap-deployments ns/name</code>	gives each machine's idle CPU to the serving pods on it
rightsize	<code>--rightsize-deployments ns/name</code>	each pod's CPU request follows its measured use times (1 + headroom), in place
coldstart	<code>--coldstart-deployments ns/name</code> <code>--coldstart-signal ns/configmap</code>	scales a service to zero while no work waits, wakes it when work arrives

Muscle	Switch	What it does
batch	<code>--batch</code>	admits held Jobs labelled <code>omnicompass.io/batch=true</code> when there is load and power headroom
batch pace	<code>--batch-pace</code>	pauses Jobs labelled <code>omnicompass.io/pausable=true</code> under power or heat stress, resumes them after
rollout guard	<code>--rollout-guard ns/name</code>	pauses a rollout while change is not permitted
contain	<code>--contain-namespaces ns --contain-cpu-m <m></code>	holds an agent namespace to a CPU budget
security hold	<code>--security-configmap ns/name</code>	key <code>hold: "true"</code> blocks every expansion

Your HPAs keep scaling as before; Omni-Compass sets their targets and raises floors ahead of bursts. Pass: p95, p99 and failed requests no worse than native, over paired runs (Chapter 39).

Level 3 - The machines. WRITES.

1. Grant patch nodes and patch pods (deploy/kind/rbac-omni.yaml).
2. Run with `--mode nodepool --active-nodes-only --closure /app/law/closure.json`
`--node-scale-cmd "<command with {n}>"` and `--node-restore-cmd "<command>"` for the OFF switch.
 For the compass law on the same levers, use `--law compass` (reads the mean response time of the latency window between a tenth of `--slo-ms` and `--slo-ms`, held at `--compass-center 0.4`; p95 at the SLO, a blind feed or a waiting pod adds capacity at once; the HPA target never goes above the operator's; machines go back one at a time through the release gate).

Your platform	The park/wake command
any cluster, bare metal, kind	<code>bash scripts/kind_nodepool.sh {n}</code>
Karpenter / EKS Auto Mode	the NodePool CPU limit at <code>{n}</code> times node CPU, parked nodes kept
Cluster Autoscaler node group	the group's desired size, scale-down through parking
OpenShift / OKD	the worker MachineSet replicas, parked

A machine is given back only when every sense is live, the last order landed, pods are not scaling up, nothing waits for a place, and the remaining machines stay inside the band. Pass: fewer machines in service, with no service gauge worse.

Level 4 - Omni-Compass decides; Kubernetes is the muscle. WRITES. Add `--strict-replicas`: Omni-Compass decides each service's replica floor and when to shrink; the HPA stays as the fast reflex upward. Pass: as level 3.

Level 5 - A GPU box, two wires. WRITES.

```
sudo python3 tools/gpu_wire_check.py --gpu 0 # must end: WIRED RIGHT
```

CHAPTER 31

Stack by Stack

Every wiring below is the same plug (Chapter 27): one wire in (the stack's own meter), one wire out (one setting the stack already accepts), a snapshot taken once before the first write, every write read back, everything restored on OFF. Native is always the stack's own controller left as shipped; omni is that controller with Omni-Compass on top.

Stack	Levels	Notes
Vanilla Kubernetes, kind, Rancher	1-7	as written
Amazon EKS	1-5	machines via Cluster Autoscaler node group or Karpenter; CPU power control is not exposed on EC2 VMs; GPUs on bare-metal or full-GPU instances
Google GKE, Azure AKS	1-5	the provider's node-pool size as the park/wake command; on AKS the managed cluster autoscaler stays native and deletes the machines Omni-Compass idles (10.1)
Red Hat OpenShift / OKD	1-5	the same permissions through a Role; machines via the worker MachineSet
NVIDIA GPU servers without Kubernetes	5	the two-wire GPU governor alone
Bare-metal CPU servers	6	cpufreq and RAPL through sysfs
Slurm / HPC	5 on the GPU nodes	a job-level Slurm connector is not built
Building management	7	through your BMS's command line or API, in the command templates
PostgreSQL behind PgBouncer (any database behind a pooler with a console)	one knob	the pool size through the pooler's own admin console (10.2)
A power grid's substation tap changer (pandapower, SimBench)	one knob	the tap position, one whole tap per move (10.3)
Buildings with batteries (CityLearn)	one knob a building	the battery's charge and discharge inside the simulator's own limits (10.3)
Robot arms (MuJoCo Menagerie)	one knob an axis	the joint speed where the axis has slack; never on an axis at its takt (10.3)
A rented cloud machine running the big organisms	the whole harness	the detached run: rent, run, collect, delete (10.4)
A message broker's consumer group (Apache Kafka)	one knob	the consumer count inside [1, partitions], read from the consumers' own records; the rebalance paid on every move (10.5)
A cache's memory ceiling (Redis)	one knob	maxmemory through the cache's own console, grown only while the cache is full (10.6)
Drone swarms (gym-pybullet-drones; PX4 and ArduPilot next)	one knob a drone	the cruise override inside the autopilot's limits; a separation wall; collisions void the cell (10.7)
A database's storage-engine cache (MongoDB under YCSB)	one knob	the WiredTiger cache size through the server's own console, grown only while the cache is full (10.8)

Managed Kubernetes on Azure: the bill as the gauge

On Azure Kubernetes Service the muscle is the same as on kind (Chapter 30, levels 1 to 3) and one thing is new: the bill is real. The workflow `aks-metered` builds a fresh cluster for every arm the same way: a system pool of one machine, tainted so no workload lands on it (Azure's add-ons and the load generator live there; it is never measured and never governed), and one or several **work pools** under Azure's own cluster autoscaler, every pool labelled `omni-role=work`, every pool starting full. Native is that autoscaler alone. Omni on top is the same autoscaler with Omni-Compass idling the machines it gives back, which the autoscaler then deletes on its own schedule. The bill counts every work machine that exists, every 15 seconds, priced at Azure's list price for its family (`tools/live_reps.py`, `POOL_PRICES`); each cluster is deleted when its arm ends, so nothing keeps billing.

Sizing the fleet to the lever. The smallest saving a fleet can show is one machine. On 4 workers one machine is 25% of the fleet, so only a saving of about 30% or more can clear the noise; on 15 workers about 10%; on 96 about 3%. The expected machine saving under Omni-Compass is a few percent, so the 4-worker runs (`results/live/V1_AKS_STEADY.md`, `V1_AKS_BURST.md`) read no difference beyond the noise on every gauge, which is the honest reading of a fleet too small for the lever, not a result against it. A new subscription allows 200 vCPUs in a region but 10 per machine family, so one family gives five 2-vCPU workers; the way past it is **several families at once**: eight pools of four to five workers each make a fleet of 39 (`worker_pools` input, `size:max,size:max,...`; `pool_prices` names each family's price). The app's replica ceiling is raised with the fleet (`hpa_max`: nine 200-millicore pods fit one 2-vCPU worker, so $9 \times$ workers lets the fleet fill), the same in both arms. The preregistered design is in `docs/K8S_COMPASS_PREREGISTRATION.md`, "The fleet that can show one machine"; the account setup, click by click, in `docs/AZURE_SETUP.md`.

What the Azure runs taught us. Three lessons, each recorded as an amendment to the preregistration before the fleet was dispatched again. First, a fleet is sized by the lever, not by the budget: a 4-worker fleet cannot show the few-percent machine saving the law makes, so its honest reading is "inside the noise", and running it again would not change that. Second, a new subscription's allowance is per machine family, so the fleet that can show one machine is built from nine families at once, each pool at its own list price, and the bill prices each pool separately. Third, a cloud region can refuse a new cluster for hours at a time (Azure's capacity for new control planes in its region, on both control-plane tiers); the harness therefore checks the subscription's allowances before it builds anything, so a refusal costs cents and not a cluster, and every refusal is recorded with its time.

A database behind its connection pooler

PostgreSQL is wired through PgBouncer, the pooler most of its deployments already run (`docs/POSTGRES_PREREGISTRATION.md`, `tools/run_pgbench.py`). Nothing models the database. The wire in is the pooler's own report (`SHOW STATS`, `SHOW POOLS`): once a second, the time a transaction spent in the server plus the time clients waited for one, over the transactions completed. The wire out is one setting through the pooler's own admin console, `SET default_pool_size`, inside the cover [2, 90]: never under two server connections, never within ten of PostgreSQL's connection limit. The compass holds that reading at 40% of a 50 ms line. Where the time was spent decides the direction: waiting for a server grows the pool; time inside the server shrinks it by one, so fewer backends fight for the same cores; calm gives back one idle server a second, and only one the pooler itself shows idle. At 95% of the line the knob is handed back to the pooler's own setting at once. If the knob is found at a value Omni-Compass did not write, it stops writing (one writer). At the end of every omni arm the pool size is restored to the snapshot and read back. The same plug fits any pooler or proxy with a console and a pool-size setting (PgBouncer, Pgpool-II, ProxySQL, the application's own pool).

What the database taught us. The first untouched run on this stack was a loss, and it stays in the preregistration as it came (docs/POSTGRES_PREREGISTRATION.md, "The first untouched run, and amendment 1"): on the read-only workload the pooler's own clock read 0.11 ms a transaction against a 50 ms line, so the law gave back one idle connection a second until the pool stood at its floor of two, while the users' 95th percentile went from 4 ms to 3.9 s. The pooler could not see the time a client spent behind its own schedule waiting to send, so its clock said calm while the users waited; the preregistration had said so in one line without drawing the consequence. Amendment 1, declared before any counted run, made the reading the worse of two: the pooler's service time, or the share of its clients queued for a server, so a pool too small for the offered rate shows at the pooler as clients waiting and the knob fails up to the operator's setting at once; and it held the dwell on the brake only, never on the gas. The counted runs are three new untouched runs on the amended rule; the first run is not counted and is not hidden. The confirmed result is a resource held, not speed, and its confirmed cost is the host's CPU: the compass's own reads of the pooler's console every second are CPU the native pooler never spends, and they are counted against it.

Independent simulators, each with its own native controller

Three published simulators are wired the same way, and the simulator's own controller is always native:

- **CityLearn** (tools/run_citylearn.py): every district the simulator ships; the knob is each building's battery charge and discharge inside the simulator's own limits; the native arm is the simulator's own rule-based control.
- **pandapower on SimBench grids** (tools/run_pandapower.py): the knob is the substation tap changer, one whole tap per move and a week of evidence before stepping down; the native arm holds the tap at 1.00 per unit as the grid ships.
- **MuJoCo Menagerie robot arms** (tools/run_mujoco.py): the knob is the joint speed of an axis that has slack; the paired physics trial written before the run leaves an axis at its takt native, so UR5e and iiwa 14 have nothing for Omni-Compass to move and the Gen3 and Panda do.

Each is preregistered (docs/CITYLEARN_PREREGISTRATION.md, docs/PANDAPOWER_PREREGISTRATION.md, docs/ROBOTICS_PREREGISTRATION.md), each runs A, B and C as separate GitHub runs, and each table is made by rule (tools/citylearn_abc.py, tools/pandapower_abc.py, tools/mujoco_abc.py). These are evidence class S: statements about the simulator's model, never about hardware.

What the simulators taught us. Three things a referee should weigh. First, a deterministic simulator makes the three-run rule stricter, not looser: MuJoCo and pandapower reproduce to the digit on any machine, so their A, B and C must agree to the digit and any difference is a finding about the harness, while PyBullet reproduces only to a part in a thousand across machines, which the swarm table discloses as an amendment and tolerates by that much and no more. Second, a knob that cannot pay is left alone by the physics, not by a rule of thumb: the robot arms' paired trial showed that two of four arms gain nothing from a speed override and left them native, the cooling model showed that a cap on cooling power can never remove less heat and so is never a lever, and a battery's reserve is spent only at the wall because every round trip through the cells costs the efficiency loss (docs/MECHANISM_OF_ACTION.md, 9.6). Third, a gain in one gauge is often a loss in another, and the simulators show both: the grid's losses rise in the four grids with their own generation while the energy drawn falls everywhere; the buildings' bills rise in the 2023 districts while their peaks fall; the Panda's copper loss rises while its torque falls. The tables carry every one of those rows, and the index carries none of them, because they are models.

The big organisms on a rented machine

The four stacked at 1,000 copies (1.7 million modelled muscles) with a real cluster inside takes more than GitHub's six-hour job limit per repetition, so it runs **detached** (`.github/workflows/big-organism-detached.yml`): `start` rents one Azure machine, installs the tools, starts every repetition of every organism under `nohup` and leaves the machine running on its own; `collect` (every two hours on the clock, or by hand) copies the files back when every repetition has finished, publishes one artifact per cell and deletes the machine; `survey` lists which machine sizes and quotas a subscription may rent, region by region, and rents nothing. A machine older than `max_hours` is collected as it is and deleted, so nothing runs forever on the bill. The commit to run is an input, so an older engine can be run again on the same machine. **The clock rule** (Chapter 39) decides the window for each size on each machine: the first v3 machine stepped the stack in 31 s against a 12 s step and fell 4,628 s behind its window, so it was stopped, recorded, and started again with a 10,800 s window and 45 s steps.

A message broker's consumer group

Apache Kafka is wired as an operator runs it (`docs/KAFKA_PREREGISTRATION.md`, `tools/run_kafka.py`): one broker as shipped, a topic of eight partitions, a consumer group at the operator's count of two. Nothing models the broker. The wire in is the group's own end-to-end latency: each message carries its production time, each consumer records when its handling ended, and the reading is the mean over the messages consumed in the last second; a second with nothing consumed while messages wait reads as the line (someone is waiting and no one is taking), and nothing consumed with nothing waiting reads as calm. The wire out is the count of consumers in the group, which is what an operator changes, inside the cover [1, 8]: never fewer than one, never more than the topic has partitions. The compass holds the reading at 40% of a 500 ms line with a response time of 3 s, because a consumer added or removed shows only after the group rebalances. Slow service with messages waiting adds consumers in proportion to the force; calm with nothing waiting gives back one idle consumer a second after a five-second dwell; at 95% of the line every consumer the topic can use starts at once. The rebalance that every move costs is paid by Omni-Compass and counted in the gauges. At the end of every omni arm the count is handed back to the operator's and read back. The result (Chapter 41) is the shape this wiring predicts: the queue kept short, no message lost, at the cost of consumers running, which reads worse by rule and stands beside the gain.

A cache's memory ceiling

Redis is wired as Ubuntu ships it, with the two settings an operator sets for a cache, a memory ceiling of 64 MB and the allkeys-lru eviction rule (`docs/REDIS_PREREGISTRATION.md`, `tools/run_redis.py`). The application in front of it is ours and declared: 1,500 requests a second from 32 threads on a Zipf working set that widens and narrows one notch at a time; a miss costs a declared 5 ms trip to the store behind the cache and a write-back. The wire in is the application's own request latency, the mean over the last second (a hit about 0.3 ms, a miss about 5.5 ms). The wire out is one setting through Redis's own console, `CONFIG SET maxmemory`, inside the cover [16, 512] MB. The compass holds the reading at 40% of a 2 ms line. Its direction rule carries the do-no-harm gate in the form this stack needs: misses grow the ceiling by notches of 8 MB **only while the cache is full** (memory used at 90% of the ceiling or more), because a miss in a cache with room to spare is a cold miss that no ceiling can mend; calm with nothing evicted in the last second gives back one notch a second, after a five-second dwell since the last growth; at 95% of the line with the cache full a quarter of the cover is added at once. One writer: a ceiling found at a value Omni-Compass did not write stops it. At the end of every omni arm the ceiling is handed back to the operator's 64 MB and read back. The memory the compass holds for a wide working set is the resource this benchmark trades and reads worse by rule; the gain is the hit rate and the work inside the line. The result (section 16.4) is exactly that shape: on all three untouched workloads the work inside the line and the hit rate rose 14% to 27% and the ceiling held rose from 64 MB to 200 to 270 MB, with the host's CPU inside the noise and every ceiling handed back.

Drone swarms: the autopilot's cruise

In `gym-pybullet-drones` (University of Toronto, MIT License; `docs/SWARM_PREREGISTRATION.md`, `tools/run_swarm.py`) every Crazyflie 2.x flies its shipped position controller along the planner's path at the planner's cruise: that is native. The wire in is each drone's tracking error, the distance between the drone and the point its autopilot was told to be at, on a band from 0 to a declared safe error of 0.25 m. The wire out is one knob a drone, the cruise override: the fraction of the planner's cruise the carrot moves at, never under 1.0 (the planner's own) and never over 2.0 (a quarter of the autopilot's own speed limit), moving by at most 0.01 a control tick. With slack the compass spends it as speed; as the error grows it eases; at 95% of the safe error the cruise is the planner's at once. Two guards belong to this stack: a **separation wall** (a drone with another within 0.5 m flies the planner's cruise, whatever the error reads) and the **paired physics trial** before the counted missions, which asks the simulator's own figures whether a faster mission costs less energy at all and keeps the error under the safe error; where it does not, the override is left native for that cell and the cell reads "nothing for Omni to move". Energy is a declared model from the simulator's own motor constants, said to be a model wherever it is printed; a collision voids the cell. At every landing and at the end of the run the override is handed back and read back.

A database's storage-engine cache

MongoDB is wired as an operator runs it (docs/YCSB_PREREGISTRATION.md, tools/run_ycsb.py): the 8.0 series from its publisher's own signed repository, its configuration as shipped but for the one setting an operator sets for the storage engine, the WiredTiger cache, at 512 MB. The questions are asked by YCSB, the published cloud-serving benchmark, running its own core workloads as the project ships them, with the key space stepping one notch at a time through the cache and past it, so the working set fits the operator's cache at the low notches and outgrows it at the high ones. The wire in is the server's own mean read latency over the last second, from its `serverStatus` operation latencies, differenced. The wire out is one setting through the server's own console, `setParameter wiredTigerEngineRuntimeConfig cache_size`, inside the cover [256 MB, 2,048 MB]: the server's own floor at one end, a quarter of the machine at the other. The compass holds the reading at 40% of a 1 ms line (written as 2 ms before the first smoke run and set at 1 ms on the tuning workload's figures, as the preregistration reserved; see below). The direction rule is the cache's: slow reads grow the cache by notches of 64 MB only while the cache is full (bytes in it at 90% of its size or more), because slow reads in a cache with room to spare are not the cache's to mend; calm with no page evicted in the last second gives back one notch a second after a five-second dwell; at 95% of the line with the cache full a quarter of the cover is added at once. One writer, read-back and the hand-back at the end are the plug's, as everywhere. Disclosed before the run: on a machine where the data files fit in the operating system's own page cache, a storage-engine cache miss is a read from memory and a decompression, not a disk read, so the gain available to this knob is smaller than it would be on a machine whose data does not fit in memory; the result will say what it is. The same plug fits any store whose engine exposes a cache size at run time (InnoDB's buffer pool, RocksDB's block cache) and any store that does not can only be sized at restart, which is not a knob Omni-Compass moves.

What the smoke runs taught, and why they are in the record. A benchmark of a cache is only a benchmark if the working set reaches the cache. Three smoke runs of the tuning workload, one repetition each and none counted, were needed before the harness asked the question it was built to ask, and each is written in the preregistration with its run number and its figures. The first failed before an arm ran, on two harness faults (a relative file path handed to a program that runs from its own directory; a server log copied as a file the upload step could not read). The second ran end to end and showed the design did not touch the cache at all: YCSB's shipped request distribution is zipfian, which concentrates almost every request on a few thousand hot records whatever the key space, so at the top notch of 1.5 million records the cache held 36 MB of its 512 MB, both arms answered inside the line with a p95 of 0.20 ms, and omni, seeing calm and nothing evicted, gave memory back to the floor of the cover. That would have read as a free memory saving for Omni-Compass, and it would have been worthless: a knob with nothing to do had moved to its floor. The requests were changed to a uniform draw over the notch's key space so the working set is the key space, the one departure from YCSB as shipped, and said so. The third smoke still held 38 MB, and the per-notch figures the runner now prints showed why: the dataset was loaded with ordered keys but the run phase used YCSB's default hashed keys, so every read asked for a record that did not exist and was answered from the index alone, which fits in a few megabytes. The run phase now names ordered keys too. Nothing in the rules, the gauges or the workloads changed across the three. The fourth smoke reached the cache: native held 336 to 481 MB of its 512 MB across the notches and read 453,000 pages into it, and its per-notch figures set the line. The server's mean read latency ran 0.15 to 0.24 ms while the cache held the working set and 0.23 to 0.33 ms while it did not, with a p95 never above 0.42 ms; against a 2 ms line every one of those readings sat in the bottom quarter of the band, under the 0.4 center, so the compass could only ever read calm, the same fault as a Kubernetes line set at ten times the service's normal response. The line was set at 1 ms, the one change the preregistration had reserved for the tuning workload, with the center unchanged, so that a hit is calm and a full cache reading past 0.4 ms is slow; the one repetition's whole-arm figures (work inside the line -1.9%, cache held -31%, pages read +32%, host CPU +8%, p95 equal) are recorded and not counted. The point for a referee is the method: a result that flatters the

governor is the first thing to suspect, the suspicion is pursued to a harness cause, the cause is written down with the run that showed it, a band is set on the tuning case's own figures and never on an untouched one, and the counted runs begin only when the harness provably asks the question.

A database's buffer pool, resized online in the server's own chunks

MySQL is wired as an operator runs it (docs/MYSQL_PREREGISTRATION.md, tools/run_sysbench.py): the 8.0 series from Ubuntu's own package, its configuration as shipped but for the settings an operator sets for InnoDB, the buffer pool at 512 MB (MySQL's own default is 128 MB, and the preregistration says why 512 MB is the operator's setting here), the chunk the server resizes in at its shipped 128 MB, one pool instance, and the performance schema on. The questions are asked by sysbench, the standard OLTP benchmark Ubuntu ships as its own package, running its published scripts as shipped at a fixed offered rate, with the number of tables in use stepping one notch at a time through the pool and past it. The wire in is the server's own mean statement latency over the last second, from the performance schema's statement summary, differenced. The wire out is one statement through the server's own console, SET GLOBAL innodb_buffer_pool_size, inside the cover [128 MB, 2,048 MB] and always in whole chunks, because that is how the server itself moves the pool. Two things about this knob differ from the MongoDB cache and the plug carries both. The server resizes asynchronously and reports its progress in its own status variable, so the plug waits for the resize to complete before reading back, and a pool found mid-resize is the server still carrying out Omni-Compass's own write, not a foreign hand. And a shrink evicts pages while the resize runs, so the dwell after any move is ten seconds rather than five. The compass holds the reading at 40% of a 0.6 ms statement line (written as 1 ms, set on the tuning workload's smoke runs, as the MongoDB line was: a hit reads 0.14 to 0.22 ms across GitHub's runners and a full pool missing a third of its reads 0.05 to 0.06 ms more, so the center sits at 0.24 ms, above the hits and under the misses); the direction rule is the cache's with one difference the smoke runs forced: slow statements grow the pool by chunks only while the pool is full, at 95% of the line with the pool full four chunks are added at once, and calm gives one chunk back after the dwell only while the pool's misses are under one percent of its read requests, because InnoDB keeps stale pages resident, so neither "pages free" nor "no page read from disk" ever says that a pool holds its working set, and the miss share does. The work-inside-the-line gauge uses a transaction line, the statement line times the statements a transaction as sysbench's script ships it (one for a point select, fourteen for the read-only mix), and says so. One writer, read-back and the hand-back at the end are the plug's, as everywhere. The same disclosure as MongoDB's holds and is made before the run: on a 16 GB machine with a 1.4 GB dataset a pool miss is a read from the operating system's page cache, not from disk.

CHAPTER 32

The Integration Manual

Before you wire anything: read `DISCLOSURES.md`. Omni-Compass acts only through the wires it is given; it cannot be slapped on. If your paired receipts differ from the published benchmarks in direction, the first presumption is wiring: confirm the installation with section 8.5 of the manual (*Wired right or wired wrong*).

How to wire Omni-Compass into your own systems yourself, from watching only to running your stack from the top. This manual ships in the box with the code and the license. Nobody from The Omni-Compass LLC needs to be on site.

Read section 1 once, then take the levels in section 4 in order. Each level has a pass condition; go to the next level only when it holds. The OFF switch (section 2) works at every level.

1. What is in the box

Item	Where
The engine, nervous system, safety shield, compass and conveyance law	<code>omnicompass/</code>
The Kubernetes controller and its muscles	<code>omni_controller/controller.py</code> , <code>omni_controller/muscles.py</code>
The GPU governor (one GPU box, no Kubernetes needed)	<code>omni_controller/gpu_governor.py</code>
Container image recipe; install and permission files	<code>deploy/Dockerfile</code> , <code>deploy/install/omni-compass.yaml</code> , <code>deploy/rbac-*.yaml</code> , <code>deploy/pilot/</code>
Proof tools: paired tests on your own stack, the GPU test, simulations	<code>scripts/</code> , <code>tools/</code> , <code>hardware/</code>
Every check the code must pass	<code>python3 verify.py</code> (ends with <code>VERIFICATION: PASS</code>)
The C++ engine: every law twinned in C++20, proven equal to the Python	<code>cpp/</code> (build: <code>cmake -S cpp -B cpp/build && cmake --build cpp/build</code>)
The seal: the fingerprints that lock each Python law to its C++ twin	<code>results/SEAL.json</code> , <code>python3 tools/seal.py --check</code>
Metrics: every gauge and where it comes from	<code>docs/METRICS_CATALOG.md</code>
How it compares with what you run today	<code>docs/COMPARISON.md</code>
The license and notices	<code>LICENSE</code> , <code>NOTICE</code>

License, in short (the `LICENSE` file governs): you may download, run, modify and test Omni-Compass free of charge for evaluation, research and non-commercial use, including on your own systems. Commercial use (running it for your business, selling it or building it into a product or service) needs a paid commercial license from The Omni-Compass LLC. The copyright and patent notices must stay with every copy.

2. The rules Omni-Compass keeps on your system

1. **One OFF switch for the whole harness, in a human hand.** `python3 tools/omni_switch.py off` turns every Omni-Compass governor on the machine off at once: each returns every setting it changed to the value it recorded before it acted, reads each back and exits, and no governor starts again until `python3 tools/omni_switch.py on` (`omnicompass/master.py`; the switch file is `OMNI_MASTER_OFF`, by default `/tmp/omni-compass/OFF`). Each governor also has its own switch for one muscle at a time: creating its kill file (or setting `OMNI_KILL=1`).
 - Kubernetes controller: `--kill-file` (default `/tmp/omni.kill`).
 - GPU governor: `--kill-file` (default `/tmp/omni-gpu-kill`), or send it `SIGTERM`.
2. **It records before it acts.** Every original setting is written down first (annotations on the Kubernetes objects, the snapshot line in the GPU audit), so the OFF switch always knows what to restore.
3. **It watches before it writes.** Every level starts in watch mode: it decides and logs, and writes nothing.
4. **It never goes blind and keeps acting.** If a reading fails (`nvidia-smi`, metrics, the response-time feed), it returns what it changed to the recorded setting at once and gives nothing back until it can see again.
5. **Service first.** While response time is over your target, and for a few decisions after (`--slo-clear`), it may not cap power or pack tighter than your own settings.
6. **Every change is read back.** No new order goes on top of one that has not landed.
7. **The living band.** No organ is driven below 5% or above 95% of its range; machines are marked idle, never switched off by Omni-Compass; it never evicts or moves a pod.
8. **Everything is logged.** Every read, decision, write and its reason goes to the audit log (`--audit`).

3. Before you start: what your system needs

You run	You need
Kubernetes (any: vanilla, EKS, GKE, AKS, OpenShift/OKD, Rancher, kind)	Kubernetes 1.34 or newer (in-place pod resize); metrics-server (<code>kubectl top nodes</code> works); an HPA with a CPU target on each service to be governed; a readiness probe and a short preStop pause on each service
A response-time feed (strongly recommended at every level)	a CSV per service, <code>elapsed_seconds, latency_ms, ok</code> , written continuously; <code>scripts/latency_probe.py</code> writes one against any HTTP endpoint
NVIDIA GPUs	driver with <code>nvidia-smi</code> ; root (or the capability to run <code>nvidia-smi -pl</code>); on a VM, full GPU passthrough (container or "pod" GPU rentals usually block power-limit changes)
CPU power control	Linux <code>cpufreq</code> with the <code>schedutil</code> governor (<code>/sys/devices/system/cpu/cpufreq</code>); RAPL for package watts (<code>/sys/class/powercap/intel-rapl</code>); usually bare metal only, since cloud VMs rarely expose these
Site power and cooling	a command that prints site watts, and (for cooling) a command that sets the supply-air setpoint through your building management system

Build the image once:

```
docker build -f deploy/Dockerfile -t <your-registry>/omni-compass:<tag> .
docker push <your-registry>/omni-compass:<tag>
```

It runs as non-root, with a read-only root filesystem and no Linux capabilities.

4. The levels, from watching to running the stack

Level 0. Evaluate without touching anything

```
pip install -r requirements.txt
python3 verify.py                # every check; ends with VERIFICATION: PASS
python3 tools/gpu_physics_sim.py # the GPU governor against a modelled NVIDIA card
python3 hardware/node_exchange.py # CPU and GPU on one power budget, modelled
```

Pass condition: VERIFICATION: PASS. Nothing in your systems is touched.

Level 1. Watch (read-only)

```
kubectl apply -f deploy/install/omni-compass.yaml # set the image line; it runs --mode observe
kubectl -n omni-compass logs deploy/omni-compass -f
```

- The install file grants a read-only identity. `scripts/pilot_shadow.sh` records `kubectl auth can-i` receipts that show it cannot write.
- Every decision it would take is logged with its reason; nothing is written.

Pass condition: zero writes, and readings your operators agree with.

Level 2. The pods, on top of your autoscalers

1. Grant patch `horizontalpodautoscalers` and patch `pods/resize` (`deploy/rbac-target.yaml`).
2. Run with `--mode target --latency-file <feed> --slo-ms <your p95 target>`.
3. Optional pod muscles, one at a time:

Muscle	Switch	What it does
convey	on with <code>--latency-file</code> and <code>--cap-deployments ns/name</code>	gives each machine's idle CPU to the serving pods on it (default: always); <code>--convey-on 0.5 --convey-off 0.25</code> engages it only while response time is over half the target
rightsize	<code>--rightsize-deployments ns/name</code>	each pod's CPU request follows its measured use $\times (1 + \text{headroom})$, in place
coldstart	<code>--coldstart-deployments ns/name</code> <code>--coldstart-signal ns/configmap</code>	scales a service to zero while no work waits, wakes it the moment work arrives
batch	<code>--batch</code>	admits held Jobs labelled <code>omnicompass.io/batch=true</code> when there is load and power headroom
batch pace	<code>--batch-pace</code>	pauses Jobs labelled <code>omnicompass.io/pausable=true</code> under power or heat stress, resumes them after
rollout guard	<code>--rollout-guard ns/name</code>	pauses a rollout while change is not permitted, undoes one past its deadline when rollback is authorised
contain	<code>--contain-namespaces ns --contain-cpu-m <m></code>	holds an agent namespace to a CPU budget with a quota
security hold	<code>--security-configmap ns/name</code>	key <code>hold: "true"</code> blocks every expansion

Your HPAs keep scaling as before. Omni-Compass sets their targets and raises floors ahead of bursts.

Pass condition: p95, p99 and failed requests no worse than your own, over paired runs (section 6).

Level 3. The machines

1. Grant patch `nodes` and patch `pods` (`deploy/kind/rbac-omni.yaml`).

2. Run with `--mode nodepool --active-nodes-only --closure /app/law/closure.json`
`--node-scale-cmd "<command with {n}>"`.

Your platform	The park/wake command
any cluster, bare metal, kind	<code>bash scripts/kind_nodepool.sh {n}</code> : close machines to new work and mark them first to go; open again, warm machines first
Karpenter / EKS Auto Mode	the NodePool CPU limit at <code>{n} × node CPU</code> , parked nodes kept, not consolidated away
Cluster Autoscaler node group	the group's desired size, scale-down through parking, not deletion
OpenShift / OKD	the worker MachineSet replicas, parked, not deleted

A machine is given back only when every sense is live, the last order landed, pods are not scaling up, nothing waits for a place, and the remaining machines stay inside the band.

Pass condition: fewer machines in service and less energy, with no service gauge worse.

Level 4. Omni-Compass decides; Kubernetes is the muscle

Add `--strict-replicas`: Omni-Compass decides each service's replica floor and when to shrink; your HPA stays as the fast reflex upward; the node law sizes the machines.

Pass condition: as level 3.

Level 5. A GPU box (with or without Kubernetes)

The GPU governor runs on any Linux machine with NVIDIA GPUs:

```
# watch: decides and logs, writes nothing
sudo python3 -m omni_controller.gpu_governor --mode watch --gpus 0,1,2,3 --audit /var/log/omni/gpu.jsonl
└─ \
  --latency-file <feed> --slo-ms <p95 target>
# cap: writes the power limits
sudo python3 -m omni_controller.gpu_governor --mode cap --gpus 0,1,2,3 --audit /var/log/omni/gpu.jsonl
└─ \
  --latency-file <feed> --slo-ms <p95 target>
# OFF
sudo touch /tmp/omni-gpu-kill
```

What it does, every `--interval` seconds (default 2):

- reads each GPU's own meter: power draw, temperature, utilization, power limit;
- the engine sets a power cap; the shield keeps it above `--min-share × the start limit (0.70)` and above `draw × (1 + --headroom)`;
- a busy card (smoothed utilization at or over `--util-gate, 0.5`) gets its full limit back at once;
- a response-time breach or a failed reading returns the start limit at once.

Speed lock (optional). Keep every response-time gauge at least `--speed-gain (1%)` faster than without Omni-Compass, and spend any speed won elsewhere on watts:

```
python3 tools/gpu_baseline.py baseline.json native-run/latency.csv # a run without Omni-Compass, other days
└─ er days
sudo python3 -m omni_controller.gpu_governor --mode cap --baseline-file baseline.json --latency-file <feed> ...
└─ ed> ...
```

Prove it on your card: `sudo bash scripts/gpu_paired.sh` runs native, watch and Omni-Compass back to back on the same machine and prints the table and verdict (`docs/GPU_RUN_GUIDE.md`).

Pass condition: work per energy up, and requests served and p95 inside the guardrails.

Level 6. CPU clock and power

Add to the Kubernetes controller, on bare metal:

```
--cpufreq-policy-root /sys/devices/system/cpu/cpufreq --cpufreq-require-schedutil \
--rapl-cmd "&lt;prints CPU package watts&gt;";
```

The CPU frequency ceiling follows the engine's cap inside the nervous system's envelope; the OFF switch writes every policy's recorded maximum back exactly. GPUs can be wired the same way from the controller:

```
--gpu-query-cmd "nvidia-smi" --query-gpu=power.draw,temperature.gpu
--format=csv,noheader,nounits" --gpu-power-cmd "nvidia-smi -pl {w}" --gpu-max-w <max>.
```

Pass condition: energy per unit of work down, no service gauge worse.

Level 7. Site power, cooling, batteries

Organ	Switch	Status
site power stress	--power-cmd "<prints site watts>" --site-limit-w <limit>	wired: power stress enters the engine; batch pace and caps respond
cooling setpoint	--cooling-cmd "<sets {c}>" --cooling-min-c 18 --cooling-max-c 27 --cooling-restore-c 22	wired: warmer supply air while cool, colder as heat rises
CPU + GPU on one power budget	hardware/node_exchange.py (the conveyance law over CPU and GPU organs)	simulation only. The live levers exist (levels 5 and 6); the exchange between them has not run on hardware
GPU groups sharing a site budget	hardware/site_exchange.py	simulation only
on-site batteries as an organ	designed (docs/DOMAIN_MAP.md)	not built

5. Stack by stack

Stack	Levels available	Notes
Vanilla Kubernetes, kind, Rancher	1-6	as written
Amazon EKS	1-5	nodes via Cluster Autoscaler node group or Karpenter NodePool (level 3 table); CPU power control is not exposed on EC2 VMs; GPUs on bare-metal or full-GPU instances
Google GKE, Azure AKS	1-5	as EKS, with the provider's node-pool size as the park/wake command
Red Hat OpenShift / OKD	1-5	OpenShift is Kubernetes underneath; grant the same permissions through a Role; machines via the worker MachineSet
NVIDIA GPU servers without Kubernetes	5	the GPU governor alone
Bare-metal CPU servers	6	cpufreq and RAPL through sysfs
Slurm / HPC schedulers	5 on the GPU nodes	a Slurm connector for job-level decisions is not built
Building management (cooling, power meters)	7	through your BMS's command line or API, wrapped in the command templates
Azure AKS, the bill as the gauge	1-3	Azure's managed cluster autoscaler stays native and deletes the machines Omni-Compass idles; one or several work pools (<code>worker_pools</code>), the replica ceiling raised with the fleet (<code>hpa_max</code>); a fleet of 4 cannot show a saving of one machine, 39 can (manual, section 10.1; <code>docs/AZURE_SETUP.md</code>)
PostgreSQL behind PgBouncer, or any pooler with a console	one knob	the pool size through the pooler's own admin console, cover [2, 90], one writer, restored on OFF (manual, section 10.2; <code>docs/POSTGRES_PREREGISTRATION.md</code>)
Apache Kafka, or any consumer group an operator sizes	one knob	the consumer count inside [1, partitions], the group's own end-to-end latency as the reading, the rebalance paid on every move, handed back on OFF (manual, section 10.5; <code>docs/KAFKA_PREREGISTRATION.md</code>)
Redis, or any cache with a console and a memory ceiling	one knob	<code>maxmemory</code> through the cache's own console, cover [16, 512] MB, grown only while the cache is full (a cold miss is not the ceiling's), one writer, restored on OFF (manual, section 10.6; <code>docs/REDIS_PREREGISTRATION.md</code>)
Drone swarms (gym-pybullet-drones; PX4 and ArduPilot next)	one knob a drone	the cruise override inside the autopilot's limits, a separation wall, the paired physics trial, collisions void the cell (manual, section 10.7; <code>docs/SWARM_PREREGISTRATION.md</code>)
MongoDB, or any store whose engine exposes its cache size at run time	one knob	the storage-engine cache through the server's own console, cover [256, 2,048] MB, grown only while the cache is full, one writer, restored on OFF (manual, section 10.8; <code>docs/YCSB_PREREGISTRATION.md</code>)
MySQL, or any store whose engine resizes its buffer pool online	one knob	the InnoDB buffer pool through the server's own console (<code>SET GLOBAL innodb_buffer_pool_size</code>), cover [128, 2,048] MB in the server's own 128 MB chunks, grown only while the pool is full, the plug waiting for the server's asynchronous resize before reading back, one writer, restored on OFF (manual, section 10.9; <code>docs/MYSQL_PREREGISTRATION.md</code>)
A rented cloud machine running the big organisms	the whole harness	<code>big-organism-detached</code> : start, collect, survey; the clock rule sets the window (manual, section 10.4)
Independent simulators (CityLearn, pandapower, MuJoCo)	one knob each	the simulator's own controller is native; preregistered, A/B/C (manual, section 10.3)

6. Proving it on your own system

1. **Paired runs.** Run your service the same way twice, once as you run it today and once with Omni-Compass on top, back to back on the same machines, order rotated, at least 5 times (10 for a result you publish). `scripts/kind_paired.sh` and `tools/live_reps.py` do this and print each gauge with its 95% interval. A change is proven only when its interval excludes zero.
2. **The switch drill,** after every run: turn Omni-Compass OFF, confirm every setting is back at its recorded original and no `omnicompass.io/*` annotation remains, turn it ON.
3. **Read the results** with `docs/METRICS_CATALOG.md`: every gauge says whether it is measured or modelled.

7. Reading the log

Line	Meaning
gate: a sense is blind	a reading failed; nothing is given back until it returns
gate: pods scaling up	pods first, machines after
decision failed (n in a row)	the cluster could not be reached; nothing was written. Turn it OFF for native at once
convey: <machine> idle CPU to its k serving pod(s), limit c	that machine's idle CPU now reaches the work on it
convey: response time calm, operator's limit	the pods are back at your own CPU limit
busy gate: utilization u, the limit read at start	a busy GPU got its full power limit back
speed lock: worst ratio r vs line 0.99 (spend / hold / release)	the speed lock's reading and what it did
response-time reflex: the limit read at start	response time went over target; full power at once
reset: the limit read at start	the OFF switch restored this setting

8. What is proven, and how

Claim	Evidence	Kind
On real Kubernetes, Omni v1, three runs of ten pairs each: more work inside the response line (+42% to +48%), p95 -57% to -69%, machines fewer where a paired trial allowed it (steady -1.5% to -3.4%, batch -15% to -24%)	results/live/V1_ALL_FOUR.md, V1_STEADY.md, V1_WANDERING.md, V1_FAULTS.md, V1_BATCH.md, V1_FAIRNESS.md	measured on kind; read by the three-run rule (docs/OMNI_V1.md)
On a real database, Omni v3: 61% to 72% fewer connections held open for the same work and latency, at +14% to +28% host CPU-seconds (confirmed worse, reported)	results/live/V3_PGBENCH.md	measured on GitHub's machines
On Azure's bill, 4 workers: no difference beyond the noise on any gauge (the fleet is too small for the lever)	results/live/V1_AKS_STEADY.md, V1_AKS_BURST.md	a real bill
Energy on Kubernetes	declared model on kind, no meter	modelled
GPU: more work per energy within the speed guardrail	earlier card controller, obsolete; rerun on rented cards pending	hardware meter
CPU and GPU on one power budget: more work, never over the budget	results/hardware/NODE_EXCHANGE*.json	modelled
Safety: OFF switch restores everything; watch mode writes nothing	every live run's switch drill; verify.py	measured / checked

9. Python, C++ and the seal

Omni-Compass exists in two languages. The laws are twinned: each has a Python version and a C++20 version that give the same answers, proven by a parity test on every build.

Law	Python	C++	Proven by
core engine (six-state equations)	omnicompass/core.py	cpp/src/core.cpp	500 frozen fixtures
governor (allocation laws)	omnicompass/adapt.py	cpp/src/governor.cpp	tests/test_cpp_governor_parity.py
safety shield	omnicompass/shield.py	cpp/src/shield.cpp	tests/test_cpp_shield_parity.py (plus adversarial cases)
HPA replica law	fleet/harness.py	cpp/src/hpa.cpp	tests/test_cpp_hpa_parity.py
closure law (machines)	omnicompass/closure.py	cpp/src/closure.cpp	tests/test_cpp_closure_parity.py
conveyance law (the conserved budget)	omnicompass/conveyance.py	cpp/src/conveyance.cpp	tests/test_cpp_conveyance_parity.py (identical to the last bit)
nervous system (authority per organ, living band)	omnicompass/nervous_system.py	cpp/src/nervous_system.cpp	tests/test_cpp_twins_parity.py
compass and ledger (composite storage)	omnicompass/compass.py, omnicompass/storage.py	cpp/src/compass.cpp	tests/test_cpp_twins_parity.py
GPU governor rules (shield limit, busy gate, speed lock, window, baseline)	omni_controller/gpu_governor.py	cpp/src/gpu_rules.cpp	tests/test_cpp_twins_parity.py

The seal (results/SEAL.json) holds the SHA-256 fingerprint of every file of every twin, written only after all parity tests pass (python3 tools/seal.py). verify.py fails if any sealed file changes afterwards, and names it.

So the Python and the C++ cannot drift apart unnoticed. What stays in Python is the plumbing that talks to Kubernetes, `nvidia-smi` and sensors (`omni_controller/controller.py`, `muscles.py`, the device I/O of `gpu_governor.py`) and the simulation harnesses; every decision they take goes through the twinned laws. The seal lists them as Python only.

Keep this manual with the code. When Omni-Compass changes, this manual, the C++ twin and the seal change in the same commit.

Evaluation and simulation use only. Copyright (c) 2026 The Omni-Compass LLC. Commercial use, commercialization or monetization of any part of Omni-Compass requires a signed, paid Omni-Compass Enterprise License from The Omni-Compass LLC. Patent applications, copyright registrations and trademark applications filed in the United States. See LICENSE and NOTICE at the root of this repository.

Running the GPU Benchmark

Before you wire anything: read DISCLOSURES.md. Omni-Compass acts only through the wires it is given; it cannot be slapped on. If your paired receipts differ from the published benchmarks in direction, the first presumption is wiring: confirm the installation with section 8.5 of the manual (*Wired right or wired wrong*).

There are three ways to get real-machine numbers:

- **Your own tower:** a machine you control, with a smart plug measuring the whole machine at the wall (section A).
- **GitHub's GPU machines:** they run the test from the repository with one click (section B).
- **A rented cloud GPU:** sections 1 to 6.

A. Your own tower, measured at the wall

You need:

- an NVIDIA graphics card (a GeForce RTX works);
- Linux on that machine (a spare drive or a USB boot is fine);
- a smart plug that reports watts, for example a Shelly Plus Plug or any plug running Tasmota, which cost about \$20 to \$30.

Plug the tower into the smart plug, and connect the plug to your home Wi-Fi with its own app. Note its IP address, shown in the app or on your router's device list. Then on the tower:

```
python3 tools/wall_meter.py shelly2:192.168.1.50 --once          # prints the tower's watts right now
sudo WALL_METER=shelly2:192.168.1.50 bash scripts/gpu_paired.sh  # trial run
sudo WALL_METER=shelly2:192.168.1.50 PHASE=confirm bash scripts/gpu_paired.sh  # the real test
```

Use shelly1: for older Shelly plugs and tasmota: for Tasmota plugs. The table then adds whole-machine energy at the wall and requests served per wall kilojoule. The plug is read by the test only; Omni never sees it, so its number is independent of Omni.

B. GitHub's own GPU machines

1. **One-time setup, by an owner of the Omni-Compass organisation on GitHub:**
 1. Go to **Settings** → **Actions** → **Runners** → **New runner** → **New GitHub-hosted runner**.
 2. Choose the image **NVIDIA GPU-Optimized Image for Linux** and a GPU size.
 3. Name its label `gpu-t4`, or set the repository variable `GPU_RUNNER` to the label you chose.

GPU runners are billed per minute on paid plans.

2. **To run:** go to **Actions** → **gpu-bench** → **Run workflow**, then choose smoke or confirm.
3. **Results:** each repetition runs on its own GPU machine. The results are pooled into one table, shown on the run's summary page, and every raw file is committed back to the branch under `results/gpu/github-<run id>/`.

If GitHub's machines don't allow changing the GPU power limit, the job stops in its first minute and says so.

C. A rented cloud GPU

This costs roughly \$10 to \$25 in rented GPU time. You type a handful of commands; the machine does the rest.

1. Rent the right kind of machine

The test changes the GPU's power limit, so you need a machine where you are the full administrator of the GPU:

- **Use a virtual machine (VM) or a bare-metal server.** Lambda Cloud "on-demand instances" are VMs; so are the GPU instances on AWS, Google Cloud and Azure.
- **Avoid "container" or "pod" rentals.** Their GPUs usually block power-limit changes. This is common on the cheapest per-hour marketplaces.
- **Any single NVIDIA data-center GPU works:** A10, L4, A100, H100, L40S. One GPU is enough.
- **Choose an image that already has PyTorch.** On Lambda that is the default "Lambda Stack" image.

You don't have to guess whether a machine allows it. In its first minute the test checks, and stops with *"cannot set the power limit (run as root)"* if the machine doesn't allow it. If you see that, shut the machine down (you pay only for the minutes used) and rent a different kind.

2. Connect to it

The rental site shows a command like `ssh ubuntu@123.45.67.89`. Paste it into Terminal (Mac) or PowerShell (Windows).

3. Get the code

On the rented machine (the repository is public while the benchmarks run; if it is private, use a read-only fine-grained token in the URL):

```
git clone https://github.com/The-Omni-Compass-LLC/The-Omni-Compass the-omni-compass
cd the-omni-compass
git log --oneline -1          # the commit you are about to test; write it down
pip install -r requirements.txt
nvidia-smi                  # should show your GPU
```

4. The whole test, one command (about 45 hours)

```
sudo nohup bash scripts/gpu_rented_run.sh &gt; run.log 2&gt;&1 &
tail -f run.log          # Ctrl+C stops watching; the test keeps running
```

It runs, in order, and stops at the first failure:

1. the machine check (one copy only, nothing else on the card, the card's default limit and clock range restored);
2. the **wire check**, which must end WIRED RIGHT (manual, section 8.4);
3. the smoke test, about 40 minutes, never counted;

4. the preregistered confirmation on compute-bound work: 10 repetitions × native / watch / Omni, 600 s each, about 6 hours;
5. the second preregistered confirmation on AI token generation (memory-bound), the same design, about 6 hours (SKIP_DECODE=1 skips it); each is its own result, never pooled. Both are packed into one file as soon as they finish: `results/gpu/omni-gpu-<stamp>-confirmations.tar.gz`;
6. **an operator's power cap underneath** (70% of the card's default limit): the cap alone vs the cap with Omni-Compass on top, at the usual load and fully loaded (more work from the same watts), about 3 hours (SKIP_CAP=1 skips it);
7. **the GPU fault drill**, about 10 minutes: the governor killed outright, the master switch pulled, the response feed blind; every check must pass (SKIP_DRILL=1 skips it). Everything so far is then packed into `results/gpu/omni-gpu-<stamp>-partial.tar.gz`;
8. the six organisms with this card inside, each as **1, 10, 100 and 1,000 copies** on one clock (3, 3, 2 and 1 repetitions), about 25 hours; the 1,000-copy stacks need a longer step than 2 s, measured on the machine and stated in the receipt (SKIP_HIL=1 skips this stage);
9. **real AI serving** last: a small open language model served by vLLM, about 2 hours (SKIP_LLM=1 skips it); if vLLM cannot install on the machine, the stage says so and nothing before it changes;
10. one packed file: == send this one file back: `results/gpu/omni-gpu-<stamp>.tar.gz`, with the label each table chose by rule.

To stop everything at any moment: `sudo python3 tools/omni_switch.py off` turns every Omni-Compass governor off and hands the card back to its own settings.

Do not start it twice and do not use the card for anything else while it runs.

5. Bring the results back

Download `results/gpu/omni-gpu-<stamp>.tar.gz` (in JupyterLab: right-click, Download; or scp from your own computer). It holds every raw reading, every table, the verdict and the checksums. **Then shut the rented machine down** on the rental site, so the billing stops.

6. Read it before you believe it

Open `results/gpu/run-<stamp>/GPU_REPS.md` and check the rows of the manual's section 8.5 (*Wired right or wired wrong*): watch equal to native, requests equal, the card's busy clock under Omni at or above its own, the lid while busy at or above its own busy draw, fail-up rare in the credit-per-write table. If a row reads wired wrong, the run says nothing about Omni-Compass until the wiring is fixed (DISCLOSURES.md, section 3).

D. The 8-GPU result: every card of one server at once

The same test as section C, on a machine with several cards (Lambda "8x A100" or "8x H100", a VM, the Lambda Stack image). Every card runs its own paired test at the same moment, sharing the server's power supply, cooling and neighbours' heat, as in a real data center. Each card starts its arm rotation one step later than the card before it, so at any moment some cards run native and some run Omni-Compass.

Get the code as in step 3, then:

```
sudo nohup bash scripts/gpu_8card.sh &> run8.log 2>&1 &
tail -f run8.log
```


Each card: wire check, envelope, smoke, the compute confirmation, the AI token generation confirmation, the operator's power cap underneath (about 15 to 16 hours, all cards together). Then the fault drill once on card 0, with every other card idle. Then one pooled table per workload (every repetition of every card, each paired within its own card: 80 repetitions per workload on 8 cards) and each card's own table. `SKIP_DECODE=1 SKIP_CAP=1` runs the compute confirmation alone, about 6.5 hours. **Every 8-GPU run has a hard time budget** (`MAX_HOURS`; by default 10.5 hours with `P00LED=1`, 4 with `FAST=1`, 26 for the whole design): at the budget the master switch hands every card back, the stages still running stop, and what finished is packed and printed as the file to send back. Lambda bills until the instance is terminated, so terminate it when that line appears. `sudo P00LED=1 nohup bash scripts/gpu_8card.sh &` runs every stage with the repetitions pooled across the cards (3 per card per confirmation, 24 per test on 8 cards; 2 per power cap load; 1 per organism size), about 9 to 10 hours. `sudo FAST=1 nohup bash scripts/gpu_8card.sh > run8.log 2>&1 &` is a short look, about 3.5 hours: 3 repetitions per card (24 on 8 cards), then one model across every card; it leaves AI token generation, the power cap and the whole stacks out, so it is never the result to show. Last, the whole server as one: a language model (Qwen2.5-7B-Instruct, open) served across every card at once by vLLM, one Omni-Compass governor per card, the server's total GPU energy, about 2 hours (`SKIP_LLM=1` skips it). Before it, **the whole stacks with a real card inside**: the six organisms at 1, 10, 100 and 1,000 copies, native and Omni, the full repetitions (3, 3, 2, 1 by size), each organism on its own card at the same time, so the stage that takes about 25 hours on one card takes the time of its longest organism (the four stacked, at 1,000 copies), about 4 to 6 hours (`SKIP_HIL=1` skips it). The whole design, every stage: about 22 to 24 hours. Progress: `tail -f results/gpu/8card-<stamp>/card-*.log`. At the end: == send this one file back: `results/gpu/omni-8card-<stamp>.tar.gz`. The master switch (`sudo python3 tools/omni_switch.py off`) stops every card's governor at once.

E. Your own serving engine

own vLLM? Start it as you always do, then:

```
sudo LLM_URL=http://127.0.0.1:8000 LLM_MODEL=<your model name>; GPU=0 ENVELOPE=<envelope.json>; OUT=result
  ↪ s/gpu/mine \
bash scripts/gpu_vllm.sh
```

The same paired test (your engine alone against your engine with Omni-Compass on top) runs against it; nothing is installed or started. `GPU=0,1,2,3,4,5,6,7` when your engine spans several cards. Make the envelope once with `python3 tools/declare_envelope.py envelope.json --gpu 0`.

F. Whole-server power from the server itself

The bench records the whole machine's watts beside the cards' own when you give it a meter: `WALL_METER=redfish:<BMC address>` (with `REDFISH_USER` and `REDFISH_PASSWORD`, read only), `WALL_METER=ipmi:local` (the server's management controller, read on the machine), or a smart plug (`tools/wall_meter.py` lists them). Omni-Compass never reads it.

Before any machine: the simulated card

`python3 tools/gpu_physics_sim.py --out results/gpu/sim/after` runs Omni's own GPU governor, with its current settings, against a modelled card for 10 paired repetitions in about two minutes. No GPU is needed. Every number it prints comes from the card model, not from a meter, so it is a preview of what the governor does and never the result. `results/gpu/sim/before` holds the same run with the settings before the share floor and busy gate were added.

What you will get

A table from the GPU's own power meter:

- native against Omni watching only, and native against Omni governing;
- the preregistered primary result: work per energy, in requests served per kilojoule, with its 95% interval.

It ends with one verdict line, one of:

- **better, proven;**
- **worse, proven;**
- **not proven;**
- **better on energy, fails the service guardrail.**

Whatever it says is the answer, and it is the first number in this project that Omni-Compass's own code did not compute.

Evaluation and simulation use only. Copyright (c) 2026 The Omni-Compass LLC. Commercial use, commercialization or monetization of any part of Omni-Compass requires a signed, paid Omni-Compass Enterprise License from The Omni-Compass LLC. Patent applications, copyright registrations and trademark applications filed in the United States. See LICENSE and NOTICE at the root of this repository.

PART SIX

Operating It

The OFF switch, the rules the governor obeys, the log it keeps, and the care of a running installation.

The OFF Switch, the Rules, and the Log

The rules Omni-Compass keeps on your system.

1. **The master switch, for the whole harness, in a human hand:** `python3 tools/omni_switch.py off` turns every Omni-Compass governor on the machine off at once (the Kubernetes controller and the GPU governors). Each puts every setting it ever wrote back to the value it read before its first write, reads it back and exits; the command waits until all have, and says so. While the switch is OFF, no governor will start. `python3 tools/omni_switch.py on` allows them to be started again (nothing restarts by itself); `status` shows the switch and every governor running. Use it the moment anything looks wrong, including a suspected breach. Each governor also keeps its own switch for one muscle at a time: the kill file (or `OMNI_KILL=1`, or `SIGTERM` to the GPU governor). **If a governor dies without handing back** (killed outright, the machine crashed, it hung), nothing it wrote is left in place: every governor records, before its first write, the exact commands that put every setting back, and renews a lease every decision. Run `python3 tools/omni_switch.py watchdog` as its own service next to the governors: it hands back for any governor whose process is gone or whose lease has run out (a hung one is stopped first). `status` names any governor that died without handing back.
2. It records before it acts (annotations on Kubernetes objects; the snapshot line in the GPU audit).
3. It watches before it writes.
4. It never acts blind.
5. Service first: while response time is over its target, and for `--slo-clear` decisions after, nothing is given back.
6. Every change is read back; no new order goes on top of one that has not landed.
7. The living band and the cover: no organ is driven outside its range; machines are parked, never switched off by Omni-Compass; it never evicts or moves a pod.
8. One writer: if anyone else changes a knob, Omni-Compass stops writing it and leaves it alone.
9. Everything is logged: every read, decision, write and reason goes to the audit log.

The three ways a governor stops, and what each leaves behind. A referee or an underwriter should know exactly what happens to a governed system in each of the three ways Omni-Compass can stop, because the answer is the same in all three: every setting at native, read back, with no record left.

- **The reset** (the brake held to the floor) is the planned end of a run. The governor sees its reset file, hands every setting back in reverse order of power-up, reads each back, writes its last audit line and exits 0. Every benchmark arm ends with one and then checks the cluster from outside the governor: the HPA target at the operator's, the replica range as set, every pod's CPU limit at its shipped value, every worker back in service, no `omncompass.io/` annotation left on any object. An arm whose reset check fails is invalid and is left out of the table, never counted.
- **The kill switch** (security) is the unplanned end of trust. `tools/omni_switch.py off` writes the switch file and sends every registered governor `SIGTERM`; each performs the same hand-back as a reset and exits; the

command waits until all have and reports any that have not. No governor starts while the switch is off. It is one switch for the whole machine, not one per muscle, because the case it exists for, something driving a system through Omni-Compass's brain, is not a case for picking which muscle to stop.

- **A death without a hand-back** (killed outright, a crash, a hang) is the case the lease covers. Before its first write every governor registers itself with the exact command that puts every setting back (`omni_controller.controller --restore-only`, which restores from the snapshot annotations on the objects themselves, so it needs no memory of the dead process), and renews a lease every decision. The watchdog (`tools/omni_switch.py watchdog`), run as its own service beside the governors, finds any governor whose process is gone or whose lease is older than its limit (five decision periods, at least two minutes), stops a hung one, runs its recorded restore command, and records what it did. The robustness benchmark (Chapter 41, in preparation) kills the governor outright in the middle of a governed run and measures how many seconds pass before every setting is back and what the service saw in that gap.

Reading the log.

Line	Meaning
gate: a sense is blind	a reading failed; nothing is given back until it returns
gate: pods scaling up	pods first, machines after
decision failed (n in a row)	the cluster could not be reached; nothing was written; turn it OFF for native
decided_by: compass	the compass set the wires this decision
decided_by: fail_up / blind_fail_up	the service crossed the 0.95 wall, or a sense went blind: full capacity at once
decided_by: thermal_hold	the card reported a heat slowdown; nothing was tightened
foreign_writer	someone else changed a knob; Omni-Compass now observes only
restored ... ok: true	the OFF switch put every setting back and read it back
"watchdog": "handed back" (in the watchdog's own log)	a governor died or hung without handing back; the watchdog ran its recorded restore command; OK says whether every command exited 0
cruise: on / cruise: off	work waited for a place for two decisions: every machine in service; the line has been empty for two: cruise ends
emergency_brake	the work is done and demand is at zero: straight to the floor in one move
verdict_state: left native	no step of the slow knob passed the paired trial: the knob is never moved on this muscle
shield_interventions	the shield clipped a write to its bounds or its step limit; the count says how often
overhead (the last line of the audit)	the governor's own CPU over the window: its process and every command it ran

How to read an audit from end to end. The audit (`audit.jsonl`, one JSON record a line, every line timestamped) is the governor's complete account of itself, and a reader can follow a run through it without the manual: the first lines are the snapshot (every setting as found, written to the objects as annotations so a watchdog can restore from them); then, every decision, the readings (the response time and its age, the pods waiting, the machines seen, whether any sense is blind), the compass's position and force (compass), the authority the nervous system granted and the release gate's reason (authority, node_gate), the decision itself (the machines recommended against the machines seen, the HPA target set), every write with the command that made it and the value read back, and the decision's wall time (decision_ms); in cruise, the floor step's silence; at the end, the reset's writes and read-backs and the overhead record. The benchmark harness prints a decision trail from the audit at the end of every omni arm (`scripts/kind_bench.sh`), one line a decision, so a reviewer reading a run's log sees what the governor saw and

did at each minute.

Maintenance, Upgrades and Security

- Run `python3 verify.py` after every upgrade; it must end `VERIFICATION: PASS`.
- Upgrade in watch mode first; take the levels again from level 1.

What an upgrade is, and what it is not. The engine is frozen and fingerprinted (Chapter 15). An upgrade that changes any engine file, a rule, a gain, a guard, a preset or a muscle, is by definition the next engine version, and every published result is run again on it before it is quoted beside the old ones; `tools/omni_version.py` says which version a checkout carries, and a checkout that matches none says which files differ. An upgrade that changes only the harness, the tools, the tables or this manual leaves the engine's fingerprint as it was and every result standing; the robustness harness of section 16.8 is an example, built without touching an engine file. An operator can therefore tell, from one command, whether a new checkout is the engine the results were made on or something that must prove itself again.

- The container runs non-root, read-only, with no capabilities; the Kubernetes identity has only the permissions of the level you run.
- The GPU governor needs root only for `nvidia-smi -pl` and `-lgc`.
- Report a security issue as `SECURITY.md` describes.

Least privilege, proven on every run. The Kubernetes governor runs as its own service account (`deploy/kind/rbac-omni.yaml`) with exactly the verbs the level needs: patch on the governed HPA, nodes, pods and the governed deployment, get on its security ConfigMap, create on pod evictions. It cannot delete nodes, create or delete pods or deployments, touch the load generator or anything in `kube-system`, read any secret, create a namespace, or write its own security hold. Every benchmark arm records `kubectl auth can-i` for each of these and fails if a "cannot" reads yes or a "can" reads no (`bench-*/rbac_omni.txt` in every archived run). The security hold itself is a ConfigMap another identity owns: when its key reads hold: "true" every expansion is blocked, and the governor cannot clear it.

What never leaves the repository. Cloud credentials live only in the repository's encrypted secrets and are used only by the workflows that need them; no credential, key or token appears in any file, log, table or commit, and the founder is never asked for one in conversation. The repository carries SPDX license identifiers on every source file so that license scanners (Black Duck, FOSSA, Snyk) read the evaluation license correctly, and every generated report carries the legal notice at its head and foot (`tools/legal.py`). Patent, copyright and trademark filings are referred to as filed; no filing number appears anywhere.

Supply chain. Everything the benchmarks download is pinned and checked: metrics-server by version and SHA-256 (`scripts/kind_bench.sh`), Apache Kafka against the Apache Software Foundation's own SHA-512 file for that release, Redis as the Ubuntu runner's own package, MongoDB from its publisher's own repository signed by the publisher's key and YCSB against the SHA-256 written in its preregistration, gym-pybullet-drones at one named commit, Python packages by `requirements.txt`. A download whose checksum does not match stops the run before anything else happens.

The upgrade drill, step by step. An operator who receives a new checkout does five things, in order, and stops at the first that fails. First, `python3 tools/omni_version.py`: if it prints the version the operator's results were made on, the engine is unchanged and the rest of the drill is about the harness; if it prints a list of differing files, the checkout is a new engine and nothing the operator measured before applies to it until measured again. Second, `python3 verify.py`, which must end with its pass line; the verifier runs every unit test, the shield's adversarial cases, the Python-against-C++ agreement and the release manifest, and a checkout that fails it is not installed anywhere. Third, the wire check on the operator's own stack, which proves that every wire the new checkout would use still follows, reads back and returns. Fourth, watch mode for at least one full cycle of the operator's load, reading the log the governor would have written and comparing its would-be decisions against what the operator would have done. Fifth, the levels again from level one, each with its paired receipt. An upgrade that skips a step has not been upgraded; it has been installed.

The security hold, and who holds it. The hold is the one control that outranks the governor's own judgment, and it is deliberately not the governor's to hold. It lives in a `ConfigMap` owned by another identity (the operator's security team, in practice); while its key reads true, every expansion the governor would make is blocked and every contraction is still allowed, so a governor under a hold can only give capacity back, never take more. The governor reads the hold on every decision and records in its audit that it read it; it has no verb that could write it. The kill switch (Chapter 34) is the blunter instrument beside it: the hold constrains a running governor, the switch stops every governor on the machine and hands everything back. An operator who suspects anything untoward uses the switch first, reads the audit second, and rotates the governor's identity third; the hold is for the slower case of a change window or an incident elsewhere during which no one wants capacity added.

What has been reviewed, and what has not. Reviewed and tested on every run: the governor's permissions against the list above, read from the cluster itself; the shield's refusal of every write past a cover, of every give-back while blind, and of every move under a hold, on two million generated cases; the lease and the watchdog's hand-back on a real cluster, thirty times (section 16.4b); the checksums of every download. Not yet done, and said so: an independent security review of the code by a party outside the company, which is in the open program (section 16.8), and a penetration test of the deployed governor on a customer-like cluster. Until those are done, the operator's own review stands in for them, and this section tells the operator exactly what to look at.

The Operator Manual

Before you wire anything: read DISCLOSURES.md. Omni-Compass acts only through the wires it is given; it cannot be slapped on. If your paired receipts differ from the published benchmarks in direction, the first presumption is wiring: confirm the installation with section 8.5 of the manual (*Wired right or wired wrong*).

The current manual is docs/INTEGRATION_MANUAL.md (every stack, every level, every switch). This page is kept for its first-person account of the laws; where the two differ, the integration manual is current.

I am Omni-Compass. I am the brain and the nervous system; your stack is the muscle. Everything below is what I am and what I do, in my own terms and in my own mathematics. Every command and flag here exists in this repository.

I. My face

My face is the compass. It is not decoration: it is how I read myself, every decision, and it is written into my code (omnicompass/compass.py).

The wheel

My deviation E, over its ceiling E_max, runs along the horizontal. Its rate of change runs along the vertical. Where I stand on that wheel is my heading, measured clockwise from north:

Point	Letter	Heading	Where I am
+	A	0°	rising through rest
↻	Δ	45°	expansion: above rest and rising, exchange under way
>	B	90°	peak extension: the turn, where my metric flips ($k \rightarrow -k$)
⊥	Λ	135°	dispersion: above rest and easing, my ceiling holds
–	Ω	180°	falling through rest
≈	Π	225°	compression: below rest, settling into my basin
<	Γ	270°	deepest compression: the turn at my floor
✦	Ψ	315°	re-alignment: below rest and rising, ignition ahead

My rim carries twenty-four letters, fifteen degrees apart, clockwise from A:

A E Z Δ H Θ B I K Λ M N Ω Ξ O Π P Σ Γ T Y Ψ Φ X

Every reading names its letter.

The four strokes

My quadrants are the strokes of the closed circle:

Quadrant	Stroke	Where I am
I	Expansion	above rest, rising
II	Dispersion	above rest, easing
III	Compression	below rest, falling
IV	Re-alignment	below rest, rising

The wheel turns $I \rightarrow II \rightarrow III \rightarrow IV \rightarrow I$:

- **Ignition:** crossing from IV into I.
- **Metric flip:** the turn at peak extension.
- **Continuity:** I am continuous through every crossing:

$$\lim X(t^-) = \lim X(t^+)$$

Nothing in me switches off, and nothing restarts.

- **Counting:** I count every circle I close.

The axle

My axle is the structural basin S . Its rest point S^* solves my equation (6):

$$\delta - \alpha_s S - \frac{3}{4} \beta_s S^2 = 0$$

Without the axle, deviation diverges. With it, deviation circulates.

Closing the circle

I hold the three conditions of the Unified Circle Principle and report them every decision:

- $\dot{X} = G(X)$, $X(0) \in \Omega$. Ω is my living band: every level I hand out lies between 5% and 95% of its range.
- $G(X) \cdot n(X) \leq 0$ on $\partial\Omega$. At a boundary, my next move points inward, never outward.
- $\nabla L(X) \cdot G(X) \leq 0$. My ledger $L = E^2/2 + \Phi(S) - \Phi(S^*)$ descends when nothing forces me. When outside load forces me, I say so; I never hide it.
- $\Rightarrow \lim X(t) \in M^*$. I settle into my basin.

A reading, as I print it on every decision trail:

```
Ψ + 318.2° Re-alignment: re-alignment, ignition ahead | Ω: machine_fill below the floor, returning | G·n
↳ ≤0 yes | L 0.4121 -0.0133 | axle S 0.19 (rest 1.90) | circles 3
```

II. My laws in your system

1. One switch, for all of me, in a human hand

I have one switch, the Unified Control Switch. A human turns it; I never turn it myself.

- **ON:** I am the primary control authority inside the scope you granted me.
- **OFF:** every setting I changed returns to what it was, and your native control takes full custody.

```
# OFF (all of me, at once, at any moment, for any reason, including a suspicion that someone has taken m
↳ y brain):
kubectl -n omni-compass exec deploy/omni-compass -- touch /tmp/omni.kill
# ON:
kubectl -n omni-compass exec deploy/omni-compass -- rm /tmp/omni.kill
```

What OFF restores. Every HPA CPU target and replica range, pod CPU limit, idle mark, paused rollout or job, GPU and CPU frequency ceiling, and containment quota. Each restore is recorded, and each lever restores on its own.

What never trips the switch.

- The switch never turns one organ off.
- No boundary, no error and no failed decision trips it.
- A decision that fails writes nothing, and my next decision comes on time.

2. My living band: 5% to 95%

Every level I hand out lives inside $\Omega = [0.05, 0.95]$ of its range:

- frequency ceilings, GPU power limits and the site power envelope;
- each organ's share of the energy budget;
- how full I run any machine.

No part of the body is driven to zero: a part with no work idles at its floor, alive and ready. No part is driven to its absolute top: I never spend the last five percent.

Tested over 300,000 of my states and 60,000 of my energy allocations (tests/test_living_band.py).

3. Machines idle; I never switch them off, and I never move a pod

When I need fewer machines, I idle the rest by letting their work leave on its own:

1. **Prefer not:** the machine is marked `omnicompass.io/idle:PreferNoSchedule`. New pods go to the open machines first, but a pod that finds them full lands here at once, so no pod ever waits because of me.
2. **Marked:** its pods are marked first to go (`controller.kubernetes.io/pod-deletion-cost`), the machine with the least work first. When the load falls, your autoscaler's own scale-down removes exactly those pods, emptying one machine at a time.
3. **Idle:** once its work is gone, the machine stays powered and Ready, gauged down to its idle floor.

While it still carries work, a machine counts as in service at full power. No pod is ever evicted, moved or restarted to idle a machine. When work returns, I remove the mark, the warm machines still carrying work first; it is in service at once, with no boot and no power cycling.

An idle machine draws `park_frac × idle power` (0.25), never zero.

4. Every change is continuous

- I idle at most one machine per decision.
- I wake a machine within five seconds of a pod waiting for a place.
- My pod reflex reads the queue every five seconds.

Nothing jumps, and nothing is restarted.

5. My pod sense: the muscle makes the pods

Your autoscaler is the muscle that makes and removes pods. I never start or stop a pod it would not.

Its rule is $\text{replicas} = \text{current} \times \text{busy} \div \text{target}$. I read the same rule from the live queue every five seconds:

- a replica serving requests answers in $R = S / (1 - u)$, so $u = 1 - S/R$;

- S is the bare service time: the fastest tenth of the recent requests;
- R is the recent mean response, over the same window;
- your target, in queue terms, is $\text{target} \times \text{request} \div \text{limit}$.

I record what the queue needs (`pod_reflex_reading`) and act only through the energy I give the pods and the target I hold for the muscle.

The target I hold. g is the CPU each pod is guaranteed with your autoscaler's largest count spread over the machines in service, divided by your limit. That target moves only when a machine idles or wakes, never each time a pod starts or leaves.

The muscle's own clock. I hold each target for your autoscaler's scale-down window (300 s unless you set one), the time it takes to answer a target. A target moved faster would pull the muscle mid-movement and start pods it then removes. A response-time breach returns your own target at once.

The reflex that raises the floor itself exists (`--pod-reflex-writes`); it is off unless you turn it on.

6. *My energy is moved, never created*

One budget comes in, and I convey it across the body by need (`omniconpass/conveyance.py`):

- $da_i/dt = \kappa a_i (e_i - \bar{e})$;
- the budget is conserved exactly, and each organ converges to its share of the demand;
- organs with no work idle at their floor, and what they do not need goes where it is needed;
- nothing leaves Ω .

On each machine, I hand its idle CPU to the pods serving on it:

$$c_i = \min(\max(L_i, (0.95 A_j - Q_j) / |P_j|), 0.95 A_j)$$

- A_j is the machine's CPU.
- Q_j is what every other pod on it has requested.
- P_j is its serving pods.
- L_i is the limit you gave the pod.

A pod's CPU limit is a quota. A request that needs more than one quota period waits for the next one while the machine stands idle. That wait is energy withheld from the work, not saved, because the request spends the same CPU-seconds either way.

I change the limit in place: the pod is not restarted, its request is untouched, and it never gets less than you gave it. Every fifteen seconds a new pod gets its share. The OFF switch returns every pod to your limit.

7. *I see before I act*

My nervous system runs both ways:

- **Afferent:** a sense that is stale, frozen or unreadable is blind, and while any sense is blind I give nothing back.
- **Efferent:** I read back every order I give. An order that did not land blocks my next release.

III. Wiring me in, step by step

Take the steps in order. Take the next step only when this step's pass condition holds. The switch works at every step.

Step 0. What your stack needs

1. Kubernetes 1.34 or newer, with metrics-server (`kubectl top nodes` answers).
2. An HPA with a CPU target on each service I govern.
3. A PodDisruptionBudget on each service, for your own maintenance; I never evict a pod.
4. A readiness probe and a short preStop pause on each service, so a new pod never takes traffic before it answers and a pod your autoscaler removes never drops a request as it leaves (`deploy/kind/demo.yaml`).
5. A response-time feed for each governed service, as CSV `elapsed_seconds,latency_ms,ok`. `scripts/latency_probe.py` writes one.

Step 1. Build me

```
docker build -f deploy/Dockerfile -t <registry>/omni-compass:<tag> .
docker push <registry>/omni-compass:<tag>
```

I run as non-root, with a read-only root filesystem and no Linux capabilities. I carry my engine, my nervous system, my compass and my frozen law.

Step 2. Let me watch (monitor only)

```
kubectl apply -f deploy/install/omni-compass.yaml # set the image line; --mode observe
kubectl -n omni-compass logs deploy/omni-compass -f
```

Identity. I hold a read-only identity. `scripts/pilot_shadow.sh` records `kubectl auth can-i` receipts showing I cannot write.

Pass condition. Zero writes, and readings your operators agree with.

Step 3. Give me the pods (on top of your autoscalers)

1. Grant patch `horizontalpodautoscalers` and patch `Pods/resize` (`deploy/rbac-target.yaml`, `deploy/kind/rbac-omni.yaml`).
2. Run me with `--mode target --latency-file <feed> --slo-ms <your p95 target>`.

Your HPAs keep scaling. My pod reflex raises floors ahead of the CPU averages.

Targets. I hold your promise in queue terms: $\text{busy} = \text{target} \times \text{request} \div \text{limit}$.

- g is the CPU each pod is guaranteed (section II.5), divided by your limit; the target that keeps each pod exactly as busy is g times yours.
- The pod answers faster, because it has g times the CPU, at the same busy share.
- Apart from that, I only tighten, never loosen.
- While response time is over your target, and for three decisions after, your own target stands.

Pass condition. p95, p99 and failed requests no worse than native.

Step 4. Give me the machines

1. Grant patch `nodes` and patch `Pods` for the first-to-go mark (`deploy/kind/rbac-omni.yaml`).
2. Run me with `--mode nodepool --active-nodes-only --closure /app/law/closure.json --node-scale-cmd "<park/wake command with {n}>"`.

Your platform	The park/wake command
any cluster, kind, bare metal	<code>bash scripts/kind_nodepool.sh {n}</code> (close to new work and mark first to go; open again, warm machines first)

Your platform	The park/wake command
Karpenter / EKS Auto Mode	the NodePool CPU limit at $\{n\} \times \text{node CPU}$, parked nodes kept, not consolidated away
Cluster Autoscaler node group	the group's desired size, with scale-down through parking, not deletion
OpenShift	the worker MachineSet replicas, parked, not deleted

I give a machine back only when all of these hold:

- every sense is live;
- my last order landed;
- pods are not scaling up;
- nothing is waiting for a place;
- the machines that remain stay inside Ω .

Pass condition. Fewer machines in service and less energy, with no service gauge worse.

Step 5. Let me decide alone (Kubernetes as the muscle only)

Add `--strict-replicas`:

- I decide each service's replica floor and when to shrink;
- your HPA stays as the fast up-reflex;
- my node law sizes the machines.

Pass condition. As step 4.

Step 6. Give me the hardware

Organ	How to wire it
CPU frequency	<code>--cpufreq-policy-root /sys/devices/system/cpu/cpufreq</code> <code>--cpufreq-require-schedutil --rapl-cmd "<prints package watts>"</code>
GPU	<code>--gpu-query-cmd "nvidia-smi --query-gpu=power.draw,temperature.gpu</code> <code>--format=csv,noheader,nounits" --gpu-power-cmd "nvidia-smi -pl {w}"</code> <code>--gpu-max-w <max></code>
Power and cooling	<code>--power-cmd "<prints site watts>" --site-limit-w <limit></code> <code>--cooling-cmd "<sets {c}>"</code>
Batch and rollouts	<code>--batch --batch-pace --rollout-guard <ns/deployment></code>

Pass condition. Energy per unit of work down, with no service gauge worse.

Step 7. The switch drill, at every step

1. Turn me OFF.
2. Confirm every setting is back to its recorded original, and no `omnicompass.io/*` annotation remains.
3. Turn me ON.

`scripts/kind_bench.sh` does exactly this after every live run.

IV. How to see me work

- **Live, paired, on real Kubernetes:** `benchmark-reps.yml` (a commit with `[reps]`, or by hand). Each repetition runs native, me on top, and me alone, back to back on one machine.
- **On your own machine:** `bash RUN_LIVE.sh 3`.

- **Every run leaves:**
 - my audit log of every read, write and compass reading;
 - my decision trail;
 - my identity receipts;
 - a SHA-256 fingerprint of every file.

V. When you read these lines

I say	I mean
gate: a sense is blind	I cannot see, so I give nothing back until I can
gate: pods scaling up	pods first, machines after
decision failed (n in a row)	I could not reach the cluster and wrote nothing; turn me OFF if you want native now
pod_reflex_reading	what the queue needs now; the autoscaler decides the pods
convey: <machine> idle CPU to its k serving pod(s), limit c	that machine's idle CPU now reaches the work on it
Ω: machine_fill below the floor, returning	the machines are underfilled and my move is bringing them back into the band

Evaluation and simulation use only. Copyright (c) 2026 The Omni-Compass LLC. Commercial use, commercialization or monetization of any part of Omni-Compass requires a signed, paid Omni-Compass Enterprise License from The Omni-Compass LLC. Patent applications, copyright registrations and trademark applications filed in the United States. See LICENSE and NOTICE at the root of this repository.

PART SEVEN

Proving It

Paired runs, receipts, rules written before the runs, and every result to date with its evidence class.



Plate 19. Empirical validation, experimental pathways, and the discipline of proof

How to Read the Results

This chapter teaches you to read every table Omni-Compass produces. You do not need to know Kubernetes or GPUs to follow it. Read the five ideas first, then go to the table you have in front of you.

The five ideas behind every table

1. **Native, and native with Omni-Compass on top.** Omni-Compass never runs a machine by itself. Every result compares a system running alone ("native": Kubernetes alone, the card's own firmware alone, a building's own controller alone) with the same system with Omni-Compass sitting on top of it. Same machine, same work, same moment: the only difference is Omni-Compass.
2. **Paired runs.** Each comparison is run many times ("repetitions" or "runs"), each time with both arms on the same seed (the same pattern of work). Each repetition gives one difference; the table reports the average difference.
3. **The 95% interval.** Next to each average is a range, for example $+0.092\%$ ($+0.090$ to $+0.093$). It means: if we ran this again and again, the true answer would land inside that range 95 times out of 100.
 - If the whole range is on one side of zero, the difference is **proven** (the table says "yes" or "significant").
 - If the range crosses zero, the difference is **not proven**: it could be noise.
4. **Better is not always "up".** For energy, response time, failures, machines and time over the line, **lower is better**, so a minus sign is good. For work per energy and work done, **higher is better**, so a plus sign is good. Every table says which.
5. **Evidence class.** Every result is marked with how it was measured:
 - **T**: a theorem;
 - **V**: verified in code;
 - **S**: a model (simulation);
 - **L**: live software (real Kubernetes);
 - **P**: a physical meter (a real card's own power meter).

Only **P** speaks for hardware energy.

The words in the tables

Word	What it means	Better is
Work per energy	work done for each unit of energy (requests per kilojoule on a card)	higher
Energy (J, Wh)	the energy used	lower
Work, requests served	how much was done; both arms must serve the same work	equal
Response time, median (p50)	half of the answers were faster than this	lower
Response time, p95	95 of 100 answers were faster than this: the slow answers, what service promises are written on	lower
Response time, p99	99 of 100 answers were faster than this: the slowest answers	lower
Time over the line, violations (pp)	the share of time the service was past its promise, in percentage points	lower
Failed requests, not served	answers that never came	zero
Machines in service, node-hours	how many machines were kept on, and for how long	lower
Band first	the founder's rule: no win unless time over the line is no higher than native's	"held"
Knobs handed back, restored	at the end every setting Omni-Compass touched was put back exactly	True / yes
Label	the verdict, chosen by a rule written before the run, never by hand	see each table

The GPU card (`results/gpu/run-<stamp>/GPU_REPS.md`)

One card, three arms, many repetitions:

- **native:** the card's firmware alone;
- **watch:** Omni-Compass running and deciding but writing nothing. It must equal native: this proves Omni-Compass watching costs nothing;
- **omni:** Omni-Compass on top, moving the card's clock ceiling and power limit.

Read it in this order:

1. **The verdict line** at the top, one of: *better, proven; worse, proven; not proven; better on energy, fails the service guardrail*.
2. **Work per energy** (requests per kilojoule): the headline. A plus sign with a range wholly above zero is a proven gain.
3. **Requests served** must be equal in every arm. More energy saved by serving less work would be cheating; the bench checks it.
4. **p95 and p99:** the slow answers. With the verdict rule (`omnicompass/verdict.py`), Omni-Compass never takes a step that makes a request more than 2% slower.
5. **Writes and restore:** how many settings were changed, and that every arm ended at the card's own limit.

The file `omni-gpu-<stamp>-confirmations.tar.gz` holds two such tables: compute-bound work (matrix products) and AI token generation (memory-bound work).

The card inside the six organisms (`results/hil/run-<stamp>/HIL.md`)

The real card is wired into each of the six organisms (the four realms, the four stacked, the whole tower) at four sizes (1, 10, 100 and 1,000 copies). Each row is one size, one organism and one part:

- **stacks (model):** the simulated machines (evidence S);

- **card (meter)**: the real card, its own meter (evidence P);
- **both**: the two added together.

The second table lists the card's own receipts for every arm: its energy, requests served, p95, and its power limit at the start and at the end, which must be equal.

Real Kubernetes (results/live/LIVE_REPS_<n>.md)

Each set is 10 paired repetitions on a real Kubernetes cluster (kind): native Kubernetes (its HPA and scheduler) against the same Kubernetes with Omni-Compass on top. Read:

1. **worker nodes in service** and **node-hours**: how many machines were kept busy. Minus is better.
2. **response time p95 and p99**: minus is faster.
3. **failed requests**: must be zero on both sides.
4. **CPU used with Omni's own**: Omni-Compass's own cost included; "no" in the significant column means it costs nothing measurable overall.
5. **The label** at the bottom, by the rule written before the run.

Energy on kind is a declared model, not a meter: kind keeps every machine powered.

The three-run tables (results/live/V1_<TEST>.md, V3_<TEST>.md)

Every benchmark since the engine was frozen runs three separate times, A, B and C, on the same bytes, and the table shows all three beside each gauge. Read the last column first: **confirmed better** or **confirmed worse** (the same sign in all three, every interval clear of zero), **no difference beyond the noise** (an interval over zero in at least one run: that is the result), or **the runs disagree** (the test is unstable there and is looked into). Nothing reads "not confirmed". The V1_ or V3_ in the name is the engine the runs carried (python3 tools/omni_version.py --commit <sha>); a table never mixes engines, and a reading is never carried from one engine to another. Losses are in the table beside the gains: the database's host CPU-seconds, the grid's losses in four grids, the Panda's copper, CityLearn's bill.

The workload tables (V3_PGBENCH.md, V3_KAFKA.md, V3_REDIS.md, V3_YCSB.md) and the swarm table (V3_SWARM.md)

The database, the broker and the cache share one shape. The head names the three runs, their commits and the engine each carried. Then one section per **workload**: the tuning workload first, marked "shown and not counted" because the rules were fitted on it, then the untouched workloads, which are the judged ones. Each section says what the workload is (value sizes, rates, steps) and what native's own capacity measured on that run, so the reader knows how hard native was pushed. The rows are the gauges with their direction in the name: the first row is always **work inside the response line**, the product number; then throughput, the latency percentiles, failures (any increase in any run is WORSE), **the resource held** (connections, consumers, the memory ceiling: lower is better, and Omni-Compass holds more of it wherever it bought a gain, so expect WORSE there), the host's CPU (the compass's own cost included), and the knob's moves (shown, not judged). The last line of each section says whether every arm handed the knob back; the last line of the table counts the rows better, worse and disagreeing across the untouched workloads. The swarm table reads the same way with cells instead of workloads, with collisions voiding a cell, and with a declared reproduction tolerance (one part in a thousand) because PyBullet does not reproduce to the bit across machines; its amendment says so.

The governor's own cost (V3_OWN_COST.md) and the robustness tables

V3_OWN_COST.md is Omni-Compass's own CPU, from the overhead record every omni arm's audit carries, by organism and copies, under the engine each run carried; it is shown and not judged, and it is a few thousandths of one core at every size. The robustness tables (V3_ROBUST_KILL.md and V3_ROBUST_LONG.md in results/live/, when they land) add rows that belong to the omni arm alone, since native has no governor to kill: the seconds from the kill until every setting was back at the operator's (the allowance is 60 s, declared before the run), whether a second governor started and governed to the end, and the governor's memory in the last ten minutes over the first ten (a quarter's growth is a leak, declared before the run); the 120 s after the kill are compared between the arms like any paired gauge.

The organisms with a real cluster inside (V1_SIX_KUBE.md, V3_SIX_KUBE.md, V1_BIG_ORGANISM.md)

Two kinds of rows, kept apart: **the cluster** (real Kubernetes, measured) and **the organism** (the modelled stacks around it, evidence S). "Organism behind its window (s)" is shown and not judged; a repetition reads **OFF THE CLOCK** when either arm ended more than 5% of its window late (its last steps saw a cluster whose load had already ended). The pairing stands, the mark stays, and the remedy is a longer window for that size on that machine.

The one number (results/OMNI_INDEX.md)

Every measure is a ratio oriented so that above 1 is better for Omni-Compass on top of native: work, speed, machines, energy. A test's index is the geometric mean of its ratios, a category's of its tests, the headline of the real categories, each weighted the same. A measure enters only as its three-run reading allows; no difference beyond the noise enters as exactly 1. Modelled muscles are shown beside the index, never inside it.

The six organisms grid (results/scale/GRID.md)

Columns are **size × runs**: 1, 10, 100 and 1,000 clusters, each at 1, 10, 100 and 1,000 paired runs. Rows are the six organisms. There is one table each for work per energy, energy, time over the line and work done. Read a cell as "with Omni-Compass on top, against native, at this size, over this many runs". Cells that are still computing say so; 1,000 runs at 1,000 clusters is beyond the machines available and says so.

Kubernetes at scale (KWOK.md, from the kwok-scale workflow)

Real Kubernetes at 50, 500 and 1,000 nodes (KWOK nodes: real Kubernetes objects with no machines behind them):

- **decision time** is how long Omni-Compass takes to decide, at that size;
- **memory** is what it uses;
- **master switch** must read "yes": everything handed back.

Where to start

1. docs/STATE_OF_PLAY.md: one page, every result and how strong it is.
2. docs/DOSSIER.md: every result with its chart.
3. This chapter, next to the table in front of you.
4. DISCLOSURES.md: what a result is and is not.

Reading the benefit: more work, or the same work for less

With the work held equal in both arms, the gain from Omni-Compass on top is $G = C_{\text{native}} / C_{\text{omni}} - 1$, where C is a resource (machine-hours, CPU core-hours, joules). The same work then needs $1 / (1 + G)$ of the resources, a saving of $G / (1 + G)$: a third more work is a quarter off the bill. The full explanation, the three rules for reading a receipt and where every number stands today are in the manual, section 3 (docs/OMNI_COMPASS_MANUAL.md).

Evaluation and simulation use only. Copyright (c) 2026 The Omni-Compass LLC. Commercial use, commercialization or monetization of any part of Omni-Compass requires a signed, paid Omni-Compass Enterprise License from The Omni-Compass LLC. Patent applications, copyright registrations and trademark applications filed in the United States. See LICENSE and NOTICE at the root of this repository.

The Dossier: Every Result in One Place

© 2026 The Omni-Compass LLC. All rights reserved. **Evaluation and simulation use only.** Any commercial use, commercialization, monetization, production use, redistribution or hosted service of any part of Omni-Compass requires a signed, paid Omni-Compass Enterprise License from The Omni-Compass LLC. Patents, copyrights and trademarks filed in the USA. See LICENSE, NOTICE and DISCLOSURES.md.

Every mechanism, harness, receipt and result, read from the files named beside it. Built by tools/dossier.py at commit 7514d1e5. Evidence classes: **T** theorem, **V** verified in code, **S** a model, **L** live software (real Kubernetes), **P** a physical meter. A model is not a meter, and a model written by the people who wrote the law is not an independent test; where a result is a model it says so.


1. The mechanism, and proof that it is the one that ran

Check	Result	Where
The whole repository re-runs and checks itself (python3 verify.py)	PASS	results/VERIFY_RECEIPT.txt
The eight-line engine and its six states, fingerprinted (omnicompass/core.py)	sha256 bd615f156169f679...	RELEASE_MANIFEST.json
Python and C++20 twins of every law, proven equal and sealed	seal intact: 9 Python/C++ twins	results/SEAL.json
The mechanism's identity against the code	mechanism identity matches the code	results/MECHANISM_IDENTITY.json
The engine's convergence to its pole under the bounded command	proved	docs/TRACKING_THEOREM.md
Safety shield: 2,000,000 adversarial cases	0 violations	tests/test_shield_properties.py
The compass law (push and pull, 5% cushions, fail up, plug contract) on every muscle, the card and Kubernetes	one law, one file	omnicompass/compass_law.py
Every file the results depend on, by fingerprint	written after the check passes	RELEASE_MANIFEST.json

The engine is the founder's eight-line equation, integrated by RK4 with the bounded command held across all four stages; it is frozen and fingerprinted, and the C++ twin matches it. The compass law is the outer loop that moves each muscle's own setting: it reads one service position (0 calm, 1 the line), pulls it to the compass's center, pushes against whatever is rising, bounds its force by tanh, fails up past the wall, and writes through a plug that reads every lever once before the first write, reads back every write, yields to any other writer and restores the snapshot at the end.

2. Real Kubernetes, Omni v3, every test three times (evidence class L)

Six tests, ten paired repetitions each, three separate GitHub runs on the frozen engine (A the result, B and C the replications): native Kubernetes (its HPA and scheduler) against the same Kubernetes with Omni-Compass on top, a fresh six-worker cluster per arm, order rotated, the same work sent to both arms. A row reads **confirmed better** or **confirmed worse** only when all three runs move the same way with every 95% interval clear of zero; otherwise it reads no difference beyond the noise, which is the result. Every Omni arm ends with every setting handed back and read back. Steady runs: A 37568428903 (omni-v3); B 37573736336 (omni-v3); C 37578883344 (omni-v3); the other tests' runs are named in their tables.



Kubernetes on v3

Test	Work	Response time (p95; the batch queue's mean)	Machines in service	Energy (declared model)	Failed requests	Table
Steady work in steps	equal by design	-65% to -65%, confirmed better	-3% to -1%, confirmed better	-2% to -1%, confirmed better	same	results/live/V3_STEADY.md
Demand that wanders	not taken	-63% to -57%, confirmed better	no difference beyond the noise	no difference beyond the noise	-12% to -8%, confirmed better	results/live/V3_WANDERING.md
All four in one run	+35% to +49%, confirmed better	-66% to -61%, confirmed better	no difference beyond the noise	no difference beyond the noise	-14% to -11%, confirmed better	results/live/V3_ALL_FOUR.md
Fairness, a noisy neighbour	not taken	no difference beyond the noise	no difference beyond the noise	no difference beyond the noise	no difference beyond the noise	results/live/V3_FAIRNESS.md
Faults: machine down, spike, runaway pod, blind probe	not taken	-62% to -47%, confirmed better	no difference beyond the noise	no difference beyond the noise	no difference beyond the noise	results/live/V3_FAULTS.md
A queue of batch jobs	not taken	-14% to -10%, confirmed better	-23% to -18%, confirmed better	-16% to -13%, confirmed better	no difference beyond the noise	results/live/V3_BATCH.md

Energy on kind is a declared model: the machines are containers on one runner, so a machine out of service saves modelled watts, not a metered bill. Omni-Compass gives a machine back only after a paired trial shows the service no slower without it (the verdict, `omnicompass/verdict.py`); on these clusters one machine fewer made requests 30-45% slower in most trials, so the machines stayed and were spent on speed and work. The v1 tables (`results/live/V1_*.md`, `docs/OMNI_V1.md`) read the same; the sets before v1 are in `docs/history/`.

3. A real database: PostgreSQL behind PgBouncer, Omni v3, three runs (evidence class L)

PostgreSQL 16 as shipped behind PgBouncer's shipped pool of 20 is native; omni is the compass law on one knob, the pool size, through PgBouncer's own console, inside the cover [2, 90] (docs/POSTGRES_PREREGISTRATION.md). Three paired repetitions a run, three runs, pgbench's own log for the gauges. Runs: A 37435740735; B 37435751322; C 37435761371.

Workload	Work inside the 50 ms line	p95	Connections held open	Host CPU-seconds (the compass's own cost)
select	no difference beyond the noise	no difference beyond the noise	-63% to -61%, confirmed better	+14% to +27%, confirmed WORSE
simple_update	no difference beyond the noise	no difference beyond the noise	the runs disagree	+18% to +28%, confirmed WORSE
tpcb_hot	no difference beyond the noise	no difference beyond the noise	-72% to -69%, confirmed better	+15% to +25%, confirmed WORSE

The compass holds fewer connections open for the same work and the same latency, and it costs CPU on the host to do so; that cost is confirmed worse and counted against Omni in the index. Table: results/live/V3_PGBENCH.md.

3b. Real messaging: Apache Kafka, a consumer group's operator-set size, Omni v3, three runs (evidence class L)

Apache Kafka as shipped (one broker, a topic of 8 partitions) with the consumer group at the operator's 2 consumers is native; omni is the compass law on one knob, the consumer count, inside the cover [1, 8], holding the group's own end-to-end latency at 40% of the 500 ms line (docs/KAFKA_PREREGISTRATION.md). Three paired repetitions a run, three runs, the consumers' own records for the gauges; the tuning workload is shown and not counted. Runs: A 37697222651; B 37697239400; C 37697255445.

Workload	Work inside the 500 ms line	End-to-end p95	Consumers running (the resource held)	Host CPU-seconds (the compass's own cost)
tuning (tuning, shown, not counted)	+15% to +16%, confirmed better	-99% to -99%, confirmed better	+264% to +297%, confirmed WORSE	+5% to +9%, confirmed WORSE
burst	+20% to +21%, confirmed better	-99% to -99%, confirmed better	+183% to +208%, confirmed WORSE	no difference beyond the noise
heavy	+16% to +16%, confirmed better	-99% to -99%, confirmed better	+158% to +192%, confirmed WORSE	no difference beyond the noise
light	+16% to +16%, confirmed better	-99% to -99%, confirmed better	+255% to +297%, confirmed WORSE	+10% to +20%, confirmed WORSE

Native sat at nine tenths of its measured capacity by design, so its queue grew at the high steps and its slowest 5% waited about 1.6 s; Omni added consumers while messages waited and gave them back when the queue was empty, so its slowest 5% waited 9 to 14 ms, at the cost of three to four times the consumers running, confirmed worse and counted against Omni in the index. No message was lost in any arm; every count was handed back. Table: results/live/V3_KAFKA.md.

3c. A real cache: Redis, the operator's memory ceiling, Omni v3, three runs (evidence class L)

Redis as Ubuntu ships it with the operator's 64 MB ceiling and allkeys-lru is native; omni is the compass law on one knob, the ceiling, inside the cover [16, 512] MB through Redis's own console, growing only while the cache is full and giving a notch back when calm and nothing is evicted (docs/REDIS_PREREGISTRATION.md). An application with a declared 5 ms store trip on a miss and a working set that steps up and down; three paired repetitions a run, three runs; the tuning workload is shown and not counted. Runs: A 37704450300; B 37704464642; C 37704479534.

Workload	Work inside the 2 ms line	Cache hit rate	p95	Memory ceiling held, MB (the resource held)	Host CPU-seconds
tuning (tuning, shown, not counted)	+15% to +15%, confirmed better	+15% to +15%, confirmed better	-1% to -1%, confirmed better	+363% to +367%, confirmed WORSE	no difference beyond the noise
burst	+14% to +15%, confirmed better	+14% to +15%, confirmed better	-1% to -0%, confirmed better	+205% to +215%, confirmed WORSE	no difference beyond the noise
large	+14% to +14%, confirmed better	+14% to +14%, confirmed better	no difference beyond the noise	+294% to +301%, confirmed WORSE	no difference beyond the noise
small	+26% to +27%, confirmed better	+26% to +27%, confirmed better	-1% to -0%, confirmed better	+320% to +325%, confirmed WORSE	no difference beyond the noise

The memory the compass holds for a wide working set is the resource this benchmark trades, and reads worse by rule. Table: results/live/V3_REDIS.md.

3d. A real database's storage-engine cache: MongoDB under YCSB, the operator's cache size, Omni v3, three runs (evidence class L)

MongoDB 8.0 as its publisher ships it with the operator's WiredTiger cache of 512 MB is native; omni is the compass law on one knob, the cache size, inside the cover [256, 2,048] MB through the server's own console, growing by notches of 64 MB only while the cache is full and reads are slow, and giving a notch back when calm and nothing is evicted (docs/YCSB_PREREGISTRATION.md). YCSB's published core workloads with the key space stepping through the cache and past it, drawn uniformly, at 3,000 operations a second from 32 threads; three paired repetitions a run, three runs; the tuning workload (workload A) is shown and not counted. Runs: A 37727968670; B 37727976107; C 37727983746.

Workload	Work inside the 1 ms line	p95	Cache size held, MB (the resource held)	Pages read into the cache	Host CPU-seconds
tuning (tuning, shown, not counted)	no difference beyond the noise	no difference beyond the noise	-15% to -12%, confirmed better	shown, not judged	no difference beyond the noise
b	no difference beyond the noise	no difference beyond the noise	no difference beyond the noise	shown, not judged	no difference beyond the noise
burst	no difference beyond the noise	no difference beyond the noise	-49% to -41%, confirmed better	shown, not judged	no difference beyond the noise
c	no difference beyond the noise	no difference beyond the noise	-50% to -49%, confirmed better	shown, not judged	no difference beyond the noise
f	no difference beyond the noise	no difference beyond the noise	-37% to -12%, confirmed better	shown, not judged	no difference beyond the noise

On this machine the data files sit in the operating system's page cache as well, so a storage-engine miss is a read from memory, not a disk, as disclosed before the run: the result is memory given back at no measurable cost in

work inside the line, p95 or CPU, with one confirmed loss, the mean latency on the burst workload. Table: `results/live/V3_YCSB.md`.

4. The bill on a real cloud: Azure Kubernetes Service, Omni v1 (evidence class L, a metered bill)


Azure's own cluster autoscaler is native; omni is the same autoscaler with Omni-Compass idling the machines it gives back; the bill is Azure's own count of machines every 15 s at list price, a fresh cluster per arm, deleted after it (`.github/workflows/aks-metered.yml`, `docs/K8S_COMPASS_PREREGISTRATION.md`, the bill on a real cloud). Transcribed from the receipts they cite:

Test	Fleet	Bill	Response time	Reading	Receipt
Steady load, 5 pairs	4 workers	-4.7%, interval across zero	inside the noise	no difference beyond the noise on any gauge	<code>results/live/V1_AKS_STEADY.md</code>
A burst sized to the cluster, 5 pairs	4 workers	+5.0%, interval across zero	p99 -34% clear of the noise in this one run; p95 inside the noise	the bill inside the noise	<code>results/live/V1_AKS_BURST.md</code>

One machine is a quarter of a 4-worker fleet, so only a saving of about 30% or more can clear the noise there; the expected machine saving is a few percent. The fleet that can show one machine (40 workers in nine machine families, the subscription's per-family allowance being 10 vCPUs) is preregistered and dispatched on v3; its first dispatches were refused by the subscription's allowances and by Azure's own cluster capacity in eastus before any arm ran, each refusal recorded in the preregistration's amendments.

4b. The GPU governor on the modelled card: each base alone, and with Omni on top (evidence class S)

Omni-Compass never runs the card. It sits on the card's own firmware (or on an operator's power cap) and moves the clock ceiling and the power limit, which that base already accepts (`omni_controller/gpu_compass.py`, the same law in `realms/gpu_card.py`). A step down is taken only after a paired trial on the card shows it adds at most 2% to the card's own time on a request (`omnicompass/verdict.py`); where no step passes, the card runs as it does alone.



The modelled card


Work	Base	Energy (tuning / fresh)	Median response	p95	p99
Compute-bound	firmware + Omni vs firmware alone	-0.70% / -0.48%	+1.56% / +1.47%	-0.84% / +0.01%	-0.09% / +0.02%
Compute-bound	105 W cap + Omni vs the cap alone	+0.02% / -0.11%	-12.29% / -25.59%	-6.41% / -7.06%	-3.15% / -5.82%
AI token generation	firmware + Omni vs firmware alone	-3.25% / -3.72%	+0.55% / +0.70%	+0.29% / +0.26%	+0.02% / -0.52%
AI token generation	105 W cap + Omni vs the cap alone	-2.10% / -2.38%	+0.07% / +0.12%	+0.03% / -0.02%	-0.06% / -1.57%

Source: results/sim/gpu_two_wire/RESULT.md and fresh/RESULT.md. The rule, and why the allowance is 2%, is amendment 8 of docs/GPU_PREREGISTRATION.md.


5. The six organisms at 1, 10, 100 and 1,000 runs and sizes, Omni v3 (evidence class S)

Each organism runs native (its own controllers) and with the compass law on every muscle, same seed, same load, same clock. Size is the number of copies of the organism governed together on one clock (1,000 copies of the four stacked is 1.7 million modelled muscles); runs are the first N of the same paired set, so 1, 10, 100 and 1,000 nest. 84 of 90 cells are done; the six left (100 and 1,000 runs at 1,000 copies) are beyond the machines available and say so.

Six organisms: work per energy, Omni against native (model)

*Work per energy*

Six organisms: time over the service line, Omni minus native (model)

*Time over the line*

The full grid with every cell: `results/scale/GRID.md`; the receipts, one per size: `results/scale/receipts/`. Work per energy is better in every cell, the same figure at every size (+0.07% for Physics to +0.37% for Energy, the four stacked and the tower); the time over the service line is at or under native's in every cell, so the band-first rule holds and every cell of 10 runs or more is labelled superior within guardrails by the preregistered rule. Every knob was handed back in every run.

6. The 945 muscles and the four realms, Omni v3 (evidence class S)

The catalog (realms/catalog.csv): 945 muscles in 59 families, 430 in Compute / AI / Cloud, 376 in Physics / Robotics / Autonomous, 470 in Energy / Facility / Industrial and 440 in Distribution / Specialized (1,716 counting a muscle once per realm, a shared spine of 257). Every muscle alone and every organism whole ran as A, B and C on v3 and reproduced to the last digit (results/realms/REALMS.md): 0 muscles worse, every organism superior within guardrails (work per energy +0.1% to +0.3%, work unchanged, time over the line not above native's). On v2 the Physics realm and the tower read a service tradeoff; the cause was a missing do-no-harm gate on speed knobs, which made v3 (docs/OMNI_V3.md). Three independent simulators with their own native controllers are wired the same way and read by the same rule: the power grid (results/live/V3_PANDAPOWER.md: energy drawn and net import better in all 11 SimBench grids with ZIP loads, losses worse in 4, tap operations 4 → 8 a year in one), robot arms (results/live/V3_MUJOCO.md: Gen3 peak torque -29%, tracking error -21%, energy per takt -0.8%; the Panda's copper +14% worse; two arms left native) and CityLearn (results/live/V3_CITYLEARN.md: electricity bought, peak and unevenness better in all 11 battery districts; the bill worse in 7). Losses stand in every table.

7. Harnesses and receipts

Harness	What it proves	Receipt
scripts/kind_paired.sh, tools/live_reps.py, workflow benchmark-reps	native against Omni on real Kubernetes, paired on one runner, the reset checked	results/live/raw/run-*/live-reps/
tools/confirm_abc.py (and pgbench_abc.py, kafka_abc.py, redis_abc.py, swarm_abc.py, mujoco_abc.py, pandapower_abc.py, citylearn_abc.py)	three separate runs read by one rule: confirmed better, confirmed worse, no difference beyond the noise, the runs disagree; each run's engine checked against the fingerprint	results/live/V3_*.md, V1_*.md
tools/omni_version.py	which frozen engine a checkout or any result's commit carries	OMNI_V1.json, OMNI_V2.json, OMNI_V3.json
scripts/aks_paired.sh, workflow aks-metered	the same on Azure's managed Kubernetes, Azure's own bill, a fresh cluster per arm, the fleet pre-flighted against the subscription's allowances	results/live/V1_AKS_*.md
tools/run_kil.py, workflows six-kube, big-organism, big-organism-detached; tools/six_kube_report.py	the six organisms with a real cluster inside at 1 to 1,000 copies, on GitHub and on a rented machine; the clock rule	results/live/V1_SIX_KUBE.md, V3_SIX_KUBE.md, V1_BIG_ORGANISM.md
tools/run_pgbench.py, workflow pgbench	a real database behind its pooler, one knob through the pooler's console	results/live/V3_PGBENCH.md
tools/run_kafka.py, workflow kafka	a real message broker as shipped, one knob (the consumer group's size) read from the consumers' own records	results/live/V3_KAFKA.md
tools/run_redis.py, workflow redis	a real cache as shipped, one knob (the memory ceiling) through its own console, growing only while the cache is full	the three-run table V3_REDIS.md in results/live/ when its runs land
tools/run_swarm.py, workflow swarm	Omni on top of each drone's shipped autopilot in gym-pybullet-drones, collisions voiding the cell	results/live/V3_SWARM.md
tools/run_scale.py, tools/pool_scale.py, workflow six; tools/grid.py	the six organisms at every run count and size	results/scale/GRID.md, results/scale/receipts/
tools/run_realms.py, workflow realms	every muscle alone and every organism whole	results/realms/REALMS.md
tools/run_pandapower.py, run_mujoco.py, run_citylearn.py	Omni on top of an independent simulator's own controller	results/live/V3_PANDAPOWER.md, V3_MUJOCO.md, V3_CITYLEARN.md
tools/omni_index.py	the one combined number, read only from the three-run tables	results/OMNI_INDEX.md
scripts/gpu_rented_run.sh, tools/gpu_wire_check.py	one command on a rented card: wire check, smoke, the organisms with the card inside, the preregistered confirmation	the founder's runs, to come
workflow archive-run	every finished run's files copied with the code, one SHA-256 manifest per repetition	results/live/raw/run-<id>/
verify.py, tools/release_manifest.py	everything above re-runs and checks itself; the manifest fingerprints the result	results/VERIFY_RECEIPT.txt, RELEASE_MANIFEST.json

The rules for each run were written and committed before it ran (docs/*_PREREGISTRATION.md); every row is reported, losses included; nothing is read across engine versions.

8. What is not yet shown

- A cloud-bill or energy saving on real machines: the 4-worker Azure fleet reads inside the noise, as a fleet too small to show one machine must; the 40-worker fleet runs are the test of that.
- The real card on the current governor: every earlier card result ran on a controller since replaced and is obsolete; the one-card, card-inside-the-organisms and eight-card runs are the founder's, on rented cards, at one named commit.
- The four stacked and the tower at 1,000 copies with the real cluster inside on v3 (the stack runs on a rented machine).
- 100 and 1,000 runs at 1,000 copies (beyond the machines available).
- The queue in docs/REGISTER.md section 4 (drone swarms on gym-pybullet-drones and Kafka done; Redis running): PX4 and ArduPilot swarms, YCSB and HammerDB, Spark, OpenSearch, fio, Open-RMF, the 24-hour robustness run, Basilisk, Orekit and GMAT, RocketPy, Cantera.

© 2026 The Omni-Compass LLC. All rights reserved. **Evaluation and simulation use only.** Any commercial use, commercialization, monetization, production use, redistribution or hosted service of any part of Omni-Compass requires a signed, paid Omni-Compass Enterprise License from The Omni-Compass LLC. Patents, copyrights and trademarks filed in the USA. See LICENSE, NOTICE and DISCLOSURES.md.

Paired Runs and Receipts on Your Own System

The method of every benchmark in this manual is the method a reader can use on their own system, and it is deliberately simple: the same system, the same work, native and then omni back to back on the same machines, repeated with the order rotated, each gauge reported with the 95% interval of its paired difference. Pairing is what makes a difference of a few percent visible on a shared machine whose own noise is larger than that: two arms run minutes apart on one runner see the same neighbours and the same hour, so their difference is a difference between arms, while two arms on separate machines differ by about 15% in p95 on their own. The order rotates so that no arm always runs first, warm or cold. The work is held equal by an open-loop load, a fixed rate of requests at fixed moments, so that a lower resource count can never mean that less work was done. The steps below are the whole method.

1. **Paired runs.** Run your service the same way twice, once native and once with Omni-Compass on top, back to back on the same machines, the order rotated, at least 5 times (10 for a result you publish).
`scripts/kind_paired.sh` and `tools/live_reps.py` do this on Kubernetes; `scripts/gpu_rented_run.sh` does it on a GPU box, in one command: wire check, smoke, the six organisms with the card inside, then the preregistered confirmation.
2. **The switch drill** after every run: turn Omni-Compass OFF; confirm every setting is back at its recorded original; turn it ON.
3. **The receipt.** Each gauge: native, Omni-Compass, the change, the 95% interval of the difference, and whether the interval excludes zero. A change is proven only when it does.
4. **The label, by rule written before the run:** SUPERIOR WITHIN GUARDRAILS / ENERGY IMPROVEMENT WITH SERVICE TRADEOFF / NONINFERIOR / NOT ESTABLISHED / WORSE / INVALID. Guardrails: work not lower by more than 1%; the share of time outside the service line not higher by more than 1 percentage point. The band-first rule is stricter: no win is claimed while the time outside the service line is above native's.

At scale (simulated). The six organisms run on GitHub's machines (Actions, workflow `six`) or on any machine (`bash scripts/scale_ladder.sh`) at 1, 10, 100 and 1,000 paired runs and at 1, 10, 100 and 1,000 copies of each organism on one clock. Real Kubernetes runs on GitHub's machines (workflow `benchmark-reps`); the six organisms with a real cluster inside as one more muscle run in `six-kube` (1 to 100 copies on GitHub) and `big-organism` and `big-organism-detached` (1,000 copies on a rented machine).

5. **Preregister, then run.** The arms, the knob, the cover, the gauges and the readings are written before the first counted run (`docs/*_PREREGISTRATION.md`); the rules are frozen on one tuning case and then applied unchanged to untouched cases. A rule changed after a result is a new version of the engine (Chapter 15), never a footnote.

6. **The clock rule.** An organism must keep its window. `tools/run_kil.py` records how long after the window its last step ended; the report shows it as "organism behind its window (s)", shown and not judged, and marks a repetition **OFF THE CLOCK** when either arm ended more than 5% of the window late, because its last steps then saw a cluster whose load schedule had already ended. The pairing stands (both arms slip alike) and the mark stays on the cell. The remedy is a longer window for that size on that machine, never a faster reading of the same run.
7. **Every row, losses included.** A row that went against Omni-Compass is reported with the others. If Omni-Compass loses, the first suspect is our own wiring, native setup or scoring; that is fixed and the test run again, and the loss stays in the record. A benchmark may be left unpublished; it is never published with rows cut out.
8. **The raw files live with the code.** Every finished run's files (capture, response times, audit log, controller log, one SHA-256 manifest per repetition) are copied into `results/live/raw/run-<id>/` by the `archive-run` workflow (one run id per line in `.github/archive_request.txt`; a run still going is left for a later pass), and every table names the runs and the commit it was made from.

How many repetitions, and why. The width of a paired interval shrinks with the square root of the number of pairs, so going from five pairs to ten narrows it by about a third, and from ten to forty by half again. On GitHub's runners, ten pairs of a 900 s window put the interval on work inside the line at about ± 10 points of native and on p95 at ± 13 to ± 29 points (the all-four table: +48.6% with an interval of +38.7 to +58.4; p95 -61.6% with -74.8 to -48.4), wide intervals that still confirm because the gains are tens of percent; the energy intervals on the same runs are ± 0.2 to $\pm 3\%$, narrow because the gauge is a declared formula over counts of machines, and that is why energy there reads "no difference beyond the noise" honestly rather than as a small claimed gain. A gain of a few percent on a cluster would need forty pairs or more to confirm at these widths, and the manual does not claim one anywhere a table shows it inside the noise. Three pairs of a 7,200 s window, the long run's shape, give a wider interval on each run but three such runs, each read on its own, still have to agree in sign for a confirmed reading. On a stack with a finer knob (connections, consumers, megabytes) the intervals are narrower at three pairs than the cluster's at ten, because the knob moves in hundreds of small steps rather than in whole machines. An operator planning their own receipt should run five pairs first, look at the interval widths on the gauges they care about, and set the count for the published run from those widths, never from the result.

Why the load is open-loop. A closed-loop load generator waits for each answer before sending the next request, so a slower system is offered less work and its throughput "improves" its own latency; two arms under closed-loop load are not doing the same work. An open-loop generator sends requests at fixed moments whatever the system does, so both arms are offered the same work to the request, a queue forms when the system falls behind, and the time over the line and the failed requests record the cost. Every real-stack test in this program (the cluster's load generator, pgbench's rate limit, Kafka's producer at a fixed rate, Redis and YCSB at a fixed target rate) offers its work open-loop, and a lower resource count in the omni arm can therefore never mean less work was asked.

What a receipt looks like. One row per gauge: the gauge's name and direction; native's value and omni's; the paired change as a percentage of native with its 95% interval; the reading. A head naming the system, the knob, the cover, the window, the number of pairs, the order rotation, the engine's fingerprint and the commit; a foot saying whether every arm was handed back and whether any repetition was off the clock. `tools/live_reps.py` prints exactly this for a Kubernetes run and the three-run tools print it for three runs side by side; a receipt from any other tool should be laid out the same way so that a reader of this manual can read it without learning a new shape.

Evidence Classes and How to Read a Result

Class	Rung	What it is	What it can show
T / V	E1	deterministic tests, proofs, Python against the C++ twin	the law is what it says, and both languages agree
S	E2	simulation on a modelled plant	whether the law helps the model, and where it breaks
L	E3	real software (Kubernetes on kind), no hardware meter	real decisions on real software; energy there is a declared model
P	E4	a physical meter (the GPU's own power reading)	the hardware's own answer

Read every number with its class beside it. A simulation number is never quoted as a hardware result. When a receipt's energy line is modelled, the receipt says so.

What each class can and cannot carry, with the examples in this manual. A class **T / V** result (the C++ twin agreeing with the Python to 3.6×10^{-15} on 500 trajectories; two million adversarial cases through the shield with zero violations) says that the law is what the text says it is and that both implementations compute it; it says nothing about whether the law helps any machine. A class **S** result (the 945 muscles, the organisms at 1,000 copies, the power grid in pandapower, the robot arms in MuJoCo, the buildings in CityLearn, the drones in PyBullet) says what the law does to a published model of a plant under that model's own native controller; it carries exactly the model's fidelity and no more, which is why S results are shown beside the Omni index and never inside it, and why the manual says "the energy is a declared model" on every S line that reports energy. A class **L** result (real Kubernetes on kind, PostgreSQL behind PgBouncer, Apache Kafka, Redis, the organisms with the real cluster inside) says what the law did to real software making real decisions on real requests; the machines, where they are containers on one runner, are counted in machine-hours in service and their energy, where it is reported, is a declared formula, which the table says; the CPU seconds of the host are measured. A class **P** result would say what a physical meter read, and the one P instrument in this program, the GPU's own power reading, has no current result: every earlier card result ran on a controller since replaced and is obsolete. Azure's bill is an L result with a real bill: the machines billed are Azure's own count, the price is Azure's list price, and the energy is still not a meter.

How to read a row. Every row in every three-run table has the same shape: the gauge and its direction; native's value in run A and omni's value in run A, so the size of the numbers is visible; the paired change in each of the three runs as a percentage of native with its 95% interval; and the reading. Read the reading first, then the intervals, then the sizes. A reading of "confirmed better" with intervals of -0.5% to -0.1% is a small, real effect; a reading of "no difference beyond the noise (3 of 3 runs)" with intervals of -30% to $+25\%$ is a test that could not see an effect of that size either way, which is a statement about the test's power, and the manual says so where it applies (the 4-worker Azure fleet). A "confirmed WORSE" row is a loss, stands in the table, and enters the index as a loss.

A worked reading, from the cache table. Take the Redis table's large workload (results/live/V3_REDIS.md). The product row, work inside the 2 ms line, reads confirmed better, with the three runs' paired changes each clear of zero and of the same sign; that is the gain. The resource row, the memory ceiling held, reads confirmed WORSE: Omni held a ceiling of about 256 MB on average where the operator had set 64 MB, in every run, clear of zero. The

two rows together are the result: the cache answered more requests inside the line because it was given more memory, and the table says both. A reader who quotes the first row without the second has misread the table, and the index is built so that this cannot happen there: the cache's category enters the one number as a loss (about 0.75 on its own), because the resource it spent is counted at full weight against the work it gained. The CPU row, inside the noise, says the host did not pay in processor time; the failed-request row, zero in both arms, says nothing was refused to get there. Four rows, three directions, one honest sentence: more work inside the line, bought with memory, at no cost in CPU or refusals. That is how every table in Chapter 41 should be read, row by row, before any sentence about it is written.

Three questions to ask of any class L result. First, what was native, exactly: which controller, at which published version, with which settings, and were those settings an operator's or ours? (Every preregistration names them; the broker's native was set at its knee by design and says so.) Second, what did the governor spend: machines, memory, connections, CPU, its own cost? (Every table has the resource rows; a result with a gain and no resource row is not in this program.) Third, who else could have produced the gain: the order of the arms, the warm-up, the runner's neighbour, the clock? (Rotation of the arm order, the warm-up before every window, the paired design on one machine and the clock rule are the answers, and section 16.7 names the residue they do not remove.)

Results to Date

Every result below is Omni-Compass **on top of** a native system against the same native system alone, with the same work in both arms, read by the three-run rule of Chapter 15, on the engine named. Losses are in the tables beside the gains. The whole list, benchmark by benchmark with its native engine, its knob, its gauges and its file, is docs/REGISTER.md; the program that takes every benchmark to full size is docs/PROOF_PROGRAM.md.

Real Kubernetes: six tests, ten pairs each, three runs each

The six Kubernetes tests were designed to ask six different questions of the same cluster, and they are read together. The cluster is kind: a real Kubernetes control plane and six worker nodes as containers on one 4-core GitHub runner, with a real metrics-server, a real HPA at the operator's target of 50% CPU, a real PHP service behind a NodePort, and a real load generator whose replica count steps through a schedule identical in both arms. A real probe times HTTP requests every five seconds from outside the service. Native is that cluster with its HPA alone; omni is the same cluster with the compass law on the HPA target inside [30%, 50%] and on the machines through the release gate and the verdict. Both arms start from the same fresh cluster, run for 900 s after a 120 s warm-up, and end with the reset check.

- **Steady work in steps** (V3_STEADY.md) asks: at a fixed, stepped load, does the governor hold the same work with fewer machines or faster answers? Work is equal by design. The 95th percentile fell 65% to 66% and the 99th 68% to 72% in all three runs; machines in service fell 1.5% to 2.9% and the standby-model energy 1.3% to 2.1%; time over the line fell 97% to 99%; failed requests were zero in both arms. The machine saving is small because the verdict kept every machine whose removal made requests slower in the paired trial.
- **Demand that wanders** (V3_WANDERING.md) asks: when load goes up and down one step at a time, does the governor follow it without harm? p95 fell 57% to 63%, p99 40% to 51%, the mean 46% to 49%, time over the line 39% to 40%, and failed requests 9% to 12%, all confirmed better; machines and energy read inside the noise.
- **All four in one run** (V3_ALL_FOUR.md) asks the product question: more work, faster, on fewer machines, with less energy, in one run. Work inside the response line rose **35% to 49%** (21.0 to 31.2 requests a second in run A), p95 fell 61% to 66%, failed requests 11% to 14%, all confirmed better; machines and energy inside the noise.
- **Faults** (V3_FAULTS.md) asks: when a machine dies, traffic triples, a pod runs away and the probe goes blind, at the same moments in both arms, does the governor make recovery worse? p95 fell 47% to 62% and p99 39% to 63%, time over the line 22% to 42%, confirmed better; failed requests were lower in all three runs and clear of the noise in one; machines and energy inside the noise.
- **A queue of batch jobs** (V3_BATCH.md) asks: when the work is a pile, does cruise and the emergency brake finish it no later and hand the machines back? Machines fell 19% to 23% over the window and 29% to 35% after the queue finished, the standby-model energy 13% to 16%, the mean response 10% to 14%, all confirmed

better; the queue finished no differently beyond the noise in all three runs (on v1 it had read 0.7% to 0.9% later, a loss that stayed in the record).

- **Fairness** (V3_FAIRNESS.md) asks: does the governor help or hurt a noisy neighbour on the same workers? No difference beyond the noise on every row, the neighbour's own p95, p99, failures and time over the line included.

The same six tests were run on Omni v1 with the same readings (V1_*.md), the controllers being the same bytes, and the v1 tables stay as the first engine's record. The earlier engines' sets, before the verdict, parked 29% to 36% of the machines and paid for it in response time (docs/history/); the frozen engine puts service first, and that choice is visible in every machines row above.

The organisms with the real cluster inside, and the big organisms

The six modelled organisms were run with the kind cluster inside as one more muscle at 1, 10 and 100 copies on GitHub (V1_SIX_KUBE.md, V3_SIX_KUBE.md) and at 1,000 copies on a rented Azure machine (V1_BIG_ORGANISM.md). At 10 and 100 copies on v3, every one of the twelve cells read better on four to six gauges (response time, time over the line, the organism's energy and its time over its own line) and worse on none beyond the noise, except a rounding-level work loss of six to eleven parts in a million in four cells, which the table shows. The four stacked at 100 copies ran off the clock in four of five repetitions (up to 277 s past a 1,440 s window): the machine could not step 170,000 muscles inside the window, both arms slipped alike, the pairing stands and the mark stays on the cell. At 1,000 copies on v1 the tower answered its slowest 5% in 0.2 s against native's 4.1 s and was over its line 0.2% of the time against 64%; the four stacked read p95 -74% and p99 -83% with the compass arm 300 to 615 s past its 2,880 s window in all three repetitions, marked off the clock, which is why the v3 stack was run with a 10,800 s window. **On v3 the four stacked at 1,000 copies** (V3_BIG_ORGANISM.md: 1.7 million muscles on one clock with the real cluster inside, three paired repetitions on one rented machine over 29 hours, 240 steps of 45 s an arm) **kept the clock in both arms**, 22 s behind a 10,800 s window (0.2%), and read: the cluster's slowest 5% answered in 150 ms against native's 2,882 ms (-95%), its slowest 1% in 165 ms against 5,405 ms (-97%), time over the response line 51% of samples to none, failed requests 0.14% to none, all clear of zero with three pairs; replicas, pods started and machines (six in both arms) the same; the standby-model energy -0.6% inside the noise; the host's CPU -1%, shown. The organism's own model (class S, shown beside the cluster and never inside the index) read energy -0.4% and work per energy +0.4% better, time over its own line -0.7% better, and work -0.0006%, a rounding-level model difference that the rule reads as worse and the table shows as such. The v1 stack's loss, the compass arm falling off the clock, is gone with the longer window; the gain, the cluster's tail, is the same shape at 1,000 copies as at 10 and 100. The whole tower at 1,000 copies runs the same way on a second rented machine as this edition is written.

The bill on a real cloud

Azure Kubernetes Service with Azure's own cluster autoscaler as native (V1_AKS_STEADY.md, V1_AKS_BURST.md): five paired repetitions each on a 4-worker fleet, every machine billed every 15 s at list price, a fresh cluster per arm. Every gauge read no difference beyond the noise; the bill read -4.7% on steady load and +5.0% on the burst, both with intervals across zero; the burst's p99 read -34% clear of the noise in that one run. This is the honest reading of a fleet in which one machine is a quarter of the fleet, so that only a saving of about 30% could clear the noise against an expected saving of a few percent (section 10.1). The fleet of 40 workers in nine machine families that can show one machine is preregistered (docs/K8S_COMPASS_PREREGISTRATION.md, amendments 1 to 4) and has been refused by Azure's own cluster capacity in its region on every dispatch so far, each refusal recorded; it runs the moment the region admits a cluster.

A real database, a real message broker, a real cache

Three stacks whose knob is a resource an operator sizes once and leaves, run as the same shape of test: three paired repetitions a run, three runs, the operator's fixed setting as native, the compass on that setting through the stack's own console as omni, a tuning workload shown and not counted, and untouched workloads judged.

- **PostgreSQL 16 behind PgBouncer** (V3_PGBENCH.md): connections held open fell 61% to 72% on select and tpcb_hot, confirmed better; the runs disagreed on simple_update; the host's CPU-seconds rose 14% to 28% on all three, confirmed worse, the compass's own cost included and counted against Omni-Compass in the index; work and latency read inside the noise. The gain is a resource held, not speed; the cost is CPU, and both are in the table.
- **Apache Kafka 3.9.1** (V3_KAFKA.md): on light, heavy and burst, work inside the 500 ms line rose 16% to 21%, the end-to-end 95th percentile fell from about 1.6 s to 9 to 14 ms, the mean lag fell 92% to 97%, all confirmed better; no message was lost in any arm; consumers running rose from 2 to 5.8 to 7.9, confirmed worse; host CPU-seconds rose 10% to 20% on light, confirmed worse, and read inside the noise on heavy and burst; CPU per thousand messages inside the line fell 10% to 12% on burst. Native sat at nine tenths of its measured capacity by design, so its queue grew at the high steps and Omni's did not; the gain is the queue kept short and the cost is the consumers that kept it so. 21 gauge-rows better, 8 worse, 0 where the runs disagree.
- **Redis 7.0.15** (V3_REDIS.md): on small, large and burst, work inside the 2 ms line rose 14% to 27%, the hit rate 14% to 27% and the mean latency fell 30% to 61%, all confirmed better; no request failed; the 95th percentile stayed within a hair of native's, because a miss costs the declared 5 ms trip in both arms and 5% of requests still miss; the memory ceiling held rose from 64 MB to 200 to 270 MB and the memory used with it, both confirmed worse, the resource the gain costs; the host's CPU-seconds read inside the noise on all three; keys evicted fell 76% to 92% (shown); every ceiling was handed back. 12 gauge-rows better, 6 worse, 0 where the runs disagree. The tuning workload, shown and not counted, read the same way (hit rate 75% to 87%, the ceiling 64 to 299 MB).
- **MongoDB 8.0.32 under YCSB 0.17.0** (V3_YCSB.md, docs/YCSB_PREREGISTRATION.md): the knob is the storage engine's cache size, the operator's 512 MB as native, and the question is the opposite of the Redis one: not whether the governor will buy service with memory, but whether it will give memory back without losing service. On workloads c and burst the cache held fell by about half (512 to about 260 MB) and on f by 13% to 37%, confirmed better; on b it fell 27% to 50% but one run's interval included zero, so that row reads no difference beyond the noise; work inside the 1 ms line, p95 and p99 read inside the noise on all four workloads (work between -2.7% and +1.1%, every interval over zero); no operation failed in any arm; the host's CPU-seconds read inside the noise on all four, +12% to +15% on c with intervals over zero; pages read into the cache rose 14% to 58% (shown, not judged: the misses the smaller cache costs); and one row is a confirmed loss, the mean latency on burst, +2.2% to +3.7% in all three runs, clear of zero, which stands. Every cache size was handed back in every arm. 6 gauge-rows better, 1 worse, 0 where the runs disagree. The tuning workload, shown and not counted, read the same way (the cache 12% to 15% smaller, everything else inside the noise). The disclosure made before the run held: on this machine the data sits in the operating system's page cache too, so a storage-engine miss costs a memory read and a decompression, not a disk read, which is why half the cache could be given back at no measurable cost in work or p95 and a few percent of mean latency; on a machine whose data does not fit in memory the same knob would be asked a harder question, and this table does not answer it.

16.4b Robustness: the governor killed outright, and its own cost

The robustness benchmark (docs/ROBUSTNESS_PREREGISTRATION.md) is a test of Omni-Compass itself, not of a gain. In the kill scenario (V3_ROBUST_KILL.md, ten pairs in each of three runs), the governor received SIGKILL at 40% of a 900 s window on the wandering load, with the watchdog running beside it from the start as it would in service. In all thirty repetitions the watchdog found the dead lease on its next pass, ran the governor's recorded restore command, and every setting (the HPA target, the replica range, every pod's CPU limit, every worker, the records on the objects) was back at the operator's **7 to 11 s after the kill**, mean 9 s against a preregistered allowance of 60 s; a second governor started at 50% and governed to the end in every repetition; and the 120 s after the kill read no difference beyond the noise against native's same window in all three runs, so the service did not measurably notice the governor's death. The whole window, with a kill and a restart inside it, still read mean response 32% to 39% faster, p95 42% to 45% faster, time over the line 28% to 34% lower and failed requests 9% to 15% fewer, confirmed better, with machines and energy inside the noise. The governor's own CPU over every archived run with the real cluster inside (V3_OWN_COST.md) is 0.006 to 0.013 of one core at every size from 1 to 1,000 copies, 0.1% to 0.3% of the host's cores, each run under its own engine; the cost of governing a real cluster does not grow with the modelled organism around it.

The long run (V3_ROBUST_LONG.md, three pairs of 7,200 s an arm in each of three runs) asks whether the governor drifts, leaks or slows over hours. Its resident memory after two hours was at most 1.048, 1.067 and 1.052 times its first ten minutes in the three runs, against a preregistered limit of a quarter: no leak. It made 117 to 119 of the 120 decisions its 60-second interval predicts in every repetition, 97.5% at the fewest against a threshold of 95%: valid. Its decision time, read as the gap between consecutive decisions less the interval it sleeps, was at most 1.20 times in the last hour what it was in the first: no slowing. Every setting was handed back at the end of every repetition and read back at the operator's with no record left. Two things are shown rather than judged, as the preregistration said they would be. First, failed decisions: 2, 3 and 3 in the three runs, all in one repetition's omni arm, all the same cause, the cluster's own API answering 500 to the governor's read of the HPA for two or three decisions in a row; the governor held, wrote nothing, and resumed on the next good read, which is the blind-means-hold rule doing what section 7.3 says it does. Second, the whole window's service: with three pairs a run the intervals are wide, and mean response, p95, p99 and time over the line read confirmed better in run A and inside the noise in B and C, so the reading is no difference beyond the noise in two of three runs; machines, energy and failed requests read inside the noise in all three. The long run is a test of the governor's constancy, and on that question every row answered as the preregistration required.

The modelled realms and the independent simulators

Every result here is evidence class S: a statement about a published model under its own native controller, never about hardware, and never counted in the Omni index.

- **The 945 muscles and six organisms** (results/realms/REALMS.md, A/B/C on v3): reproduced to the last digit in three of three runs; zero muscles worse; every organism superior within guardrails, with work per energy +0.1% to +0.3%, work unchanged and time over the line not above native's. On v2 the Physics realm and the tower had read a service tradeoff; the slack gate on speed knobs corrected it and made v3.
- **The grid of copies and runs** (results/scale/GRID.md, 84 of 90 cells on v3): every organism superior within guardrails in every cell of ten runs or more; work per energy +0.07% (Physics) to +0.37% (Energy, the four stacked, the tower), the same figure at every size; every knob handed back; the six cells left are beyond the machines available.

- **The power grid** (V1_PANDAPOWER.md, V3_PANDAPOWER.md, the same table to the digit): eleven SimBench grids solved by pandapower, the tap changer's own controller as native, one whole tap per move as the knob. With ZIP loads the energy the loads drew fell 1.3% to 1.5% and the net import fell in all eleven grids; losses fell in seven and **rose in four** (the rural and semi-urban grids with their own generation, +0.6% to +1.5%); tap operations fell in ten grids and rose from 4 to 8 a year in one, the declared cost, confirmed worse; no grid was more often outside its voltage band. 67 gauge-rows better, 14 worse.
- **Robot arms** (V1_MUJOCO.md, V1_MUJOCO_PANDA.md, V3_MUJOCO.md): four arms from the MuJoCo Menagerie with their shipped controllers as native and the joint speed of an axis with slack as the knob, after a paired physics trial. The trial left the UR5e and the iiwa 14 native ("nothing for Omni to move"); on the Gen3 and the Panda, peak torque fell 29% and 10%, tracking error 21%, energy per takt 0.8% and 0.5%, confirmed better; the Panda's copper loss **rose 14%**, confirmed worse.
- **Buildings with batteries** (V1_CITYLEARN.md, V3_CITYLEARN.md): every district CityLearn ships, its own rule-based control as native, each building's battery charge and discharge inside the simulator's limits as the knob. In the eleven battery districts electricity bought, the daily peak and the daily unevenness fell in all eleven and carbon in eight, confirmed better; the bill **rose in seven** (the 2023 districts) and ramping rose in seven, confirmed worse; 71 score-rows better, 33 worse, 1 where the runs differ by the simulator's own variation; three districts had nothing to move and eight the simulator cannot run.
- **Drone swarms** (V3_SWARM.md): twenty Crazyflie 2.x quadrotors in gym-pybullet-drones with the shipped position controller as native and the cruise override inside the autopilot's limits as the knob, three untouched cells. Energy a mission fell 7% (short), 18% (mixed) and 20% (long) and missions a charge rose 8%, 22% and 25%, confirmed better; no late mission, reserve breach, near miss or collision in any arm; the tracking error rose from 0.07 to 0.13 m inside its 0.25 m band, shown; every override handed back. PyBullet reproduced to a part in a thousand across GitHub's machines and not to the bit, so the table's reproduction tolerance was set to one part in a thousand after the runs were seen, and the preregistration's amendment 1 says so with its time.

The one number, and the table

The one number. The Omni index over the real categories confirmed three times: **+25.9%** (real Kubernetes on v3 +25.5%, the real database on v3 +14.2%, real messaging on v3 +166.9%, the real cache on v3 -24.9%, the real database's storage-engine cache on v3 +10.2%, each category weighed the same; Azure and the card join as their three-run tables land). The fifth category lowered the headline from +30.2% to +25.9% when it joined on 8 October, and that is how the number is meant to move: every real category enters at equal weight as its table lands, whatever it does to the mean. The storage-engine cache's figure is positive for the reason the Redis figure is negative, read the other way: there the governor gave memory back (the cache held halved on two workloads) while work, p95 and CPU stayed inside the noise, so its resource column is a gain and its other columns are exactly one; its one confirmed loss, the burst mean latency, is not an index column and stands in its table. The messaging figure is large because its speed ratio is: native sat at nine tenths of its measured capacity by design, so its queue grew at the high steps and its slowest 5% waited 1.6 s, while Omni's waited 9 to 14 ms; the resource that bought it, three to four times the consumers, stands beside it as a confirmed loss. The cache's figure is negative because the memory held is its resource column: the ceiling rose from 64 MB to 200 to 270 MB (a ratio of about 0.25 in the index's "fewer machines" sense) while work and the hit rate rose 14% to 27%, and a geometric mean of those four columns is below one. That is the arithmetic of a resource bought, reported as such; a reader who weighs memory as cheaper than a consumer or a machine will read the cache's rows for themselves, which is why every row is in its table.

Result	Engine	Class	Reading	Source
Real Kubernetes, all four in one run (load up and down one step at a time, 10 pairs × 3 runs)	v1	L	work inside the response line +42% to +48% in all three runs, confirmed better ; p95 –57% to –63% confirmed better; failed requests –12% to –13% confirmed better; machines no difference beyond the noise	results/live/V1_ALL_FOUR.md
Real Kubernetes, steady work in steps (10 pairs × 3 runs)	v1	L	p95 –65% to –69% and time over the line confirmed better; machines –1.5% to –3.4% confirmed better; failed requests zero in both arms	results/live/V1_STEADY.md
Real Kubernetes, demand that wanders (10 pairs × 3 runs)	v1	L	p95 –56% to –60% confirmed better; failed requests confirmed better; modelled standby energy –0.4% to –0.8% confirmed lower	results/live/V1_WANDERING.md
Real Kubernetes, faults (a machine down, a spike, a runaway pod, a blind probe, 10 pairs × 3 runs)	v1	L	p95 –61% to –64% confirmed better; failed requests, machines and energy no difference beyond the noise	results/live/V1_FAULTS.md
Real Kubernetes, a queue of batch jobs (cruise, then the emergency brake, 10 pairs × 3 runs)	v1	L	machines –15% to –24% and modelled energy –11% to –17% confirmed better; p95 no difference beyond the noise	results/live/V1_BATCH.md
Real Kubernetes, fairness (a noisy neighbour on the same workers, 10 pairs × 3 runs)	v1	L	no difference beyond the noise on every row: Omni-Compass neither helps nor hurts the neighbour	results/live/V1_FAIRNESS.md
The same six Kubernetes tests on v3 (10 pairs × 3 runs each)	v3	L	the same readings as v1, the controllers being v1's bytes: all four work inside the line +35% to +49%, p95 –47% to –66% across steady, wandering, all four and faults, failed requests –9% to –14% where they occur, steady machines –1.5% to –2.9%, batch machines –19% to –23% and standby-model energy –13% to –16%, all confirmed better; fairness no difference beyond the noise on every row	results/live/V3_STEADY.md, V3_WANDERING.md, V3_ALL_FOUR.md, V3_FAIRNESS.md, V3_FAULTS.md, V3_BATCH.md
The six organisms with the real cluster inside , 1 to 1,000 copies, 5 paired repetitions a cell (3 at 1,000)	v1	L + S	98 of 100 cells: p95 and time over the line better in every cell; 0 gauges worse beyond the noise except a rounding-level work loss (–0.0003%) and HPA replicas +0.7% in one cell; machines stay at 6 in both arms (no autoscaler under kind); the Physics realm and the tower at 1,000 copies off the clock in 2 of 3 repetitions (marked)	results/live/V1_SIX_KUBE.md

Result	Engine	Class	Reading	Source
The same at 10 and 100 copies	v3	L + S	12 cells, 5 pairs each: every cell better on 4 to 6 gauges, worse on none beyond the noise except a rounding-level work loss in 4 cells; the stack at 100 copies off the clock in 4 of 5 repetitions (marked)	results/live/V3_SIX_KUBE.md
The big organisms at 1,000 copies on a rented machine: the tower (3 of 3) and the four stacked (3 of 3, detached)	v1	L + S	tower: p95 -95% (4.1 s → 0.2 s) and time over the line -99.7% clear of the noise, machines 6 in both arms, energy -0.2%, work rounding-level worse, repetition 3 off the clock; stack: p95 -74% and p99 -83% clear of the noise, the compass arm 300 to 615 s past the 2,880 s window in all three (marked OFF THE CLOCK; native kept the clock), which is why the v3 stack runs with a 10,800 s window	results/live/V1_BIG_ORGANISM.md
Azure AKS, the bill, steady load, 4 workers, 5 pairs	v1	L (a real bill)	no difference beyond the noise on any gauge; the bill -4.7% with its interval across zero; a fleet of 4 cannot show the lever (10.1)	results/live/V1_AKS_STEADY.md
Azure AKS, the bill, a burst sized to the cluster, 4 workers, 5 pairs (repetition 4 lost its native cluster to an Azure API error in the first attempt and was run again by itself)	v1	L (a real bill)	the bill +5.0% with its interval across zero; machines, p95 and failed requests inside the noise; p99 -34% clear of the noise in this one run; B and C to follow	results/live/V1_AKS_BURST.md
PostgreSQL behind PgBouncer, the three untouched workloads, 3 pairs × 3 runs	v3	L	connections held open -61% to -72% confirmed better on select and tpcb_hot , the runs disagree on simple_update ; host CPU-seconds confirmed worse on all three (+14% to +28%, the compass's own cost, counted against Omni-Compass); work and latency no difference beyond the noise	results/live/V3_PGBENCH.md
The 945 muscles and six organisms, modelled, A/B/C	v3	S	reproduced to the last digit in 3 of 3; 0 muscles worse; every organism superior within guardrails (work per energy +0.1% to +0.3%, work unchanged, time over the line not above native's); on v2 Physics and the tower read a service tradeoff, which the slack gate corrected	results/realms/REALMS.md
The organisms at 1, 10, 100 and 1,000 copies, 1 to 1,000 paired runs a cell, 84 of 90 cells	v3	S	every organism superior within guardrails in every cell of 10 runs or more; work per energy +0.07% to +0.37%, the same figure at every size; the six cells left (100 and 1,000 runs at 1,000 copies) are beyond the machines available	results/scale/GRID.md, results/scale/receipts/

Result	Engine	Class	Reading	Source
Power grid: 11 SimBench grids solved by pandapower , both load models, A/B/C	v1 and v3 (the same table to the digit)	S	with ZIP loads the energy the loads drew confirmed better in all 11 (−1.3% to −1.5%) and the net import in all 11; losses confirmed better in 7 and worse in 4 (the rural and semi-urban grids with their own generation, +0.6% to +1.5%); tap operations fewer in 10 grids, 4 → 8 a year in one (confirmed worse, the declared cost); no grid more often outside its band	results/live/V1_PANDAPOWER.md, V3_PANDAPOWER.md
Robot arms, MuJoCo Menagerie , A/B/C	v1 and v3	S	where Omni-Compass moved (Gen3, Panda): peak torque −29% and −10%, tracking error −21%, energy per takt −0.8% and −0.5%, confirmed better; the Panda's copper loss +14% confirmed worse ; UR5e and iiwa 14 left native by the paired physics trial	results/live/V1_MUJOCO.md, V1_MUJOCO_PANDA.md, V3_MUJOCO.md
CityLearn, every district it ships , A/B/C	v1 and v3	S	11 battery districts: electricity bought, daily peak and daily unevenness confirmed better in all 11, carbon in 8; the bill worse in 7 (the 2023 districts) and ramping worse in 7; 71 score-rows better, 33 worse, 1 where the runs differ (the simulator's own variation); 3 districts with nothing to move; 8 the simulator cannot run	results/live/V1_CITYLEARN.md, V3_CITYLEARN.md
Drone swarms, gym-pybullet-drones (Crazyflie 2.x, the shipped autopilot as native; Omni on the cruise override inside the autopilot's limits), 20 drones × 4 missions, three cells, A/B/C	v3	S	energy a mission −7% (short), −18% (mixed), −20% (long) and missions a charge +8% to +25%, confirmed better; no late mission, reserve breach, near miss or collision in any arm; tracking error 0.07 → 0.13 m inside its 0.25 m band	results/live/V3_SWARM.md
Apache Kafka as shipped, a consumer group's operator-set size (one broker, 8 partitions; the group at the operator's 2 consumers as native; Omni on the count inside [1, 8]), three untouched workloads, 3 pairs × 3 runs	v3	L	work inside the 500 ms line +16% to +21% confirmed better on all three; end-to-end p95 1.6 s → 9 to 14 ms and mean lag −92% to −97% confirmed better; no message lost in any arm; consumers held 2 → 5.8 to 7.9 confirmed worse (the resource the gain costs); host CPU-seconds confirmed worse on light (+10% to +20%), inside the noise on heavy and burst; every count handed back; 21 gauge-rows better, 8 worse	results/live/V3_KAFKA.md

Result	Engine	Class	Reading	Source
Redis as shipped, a cache's operator-set memory ceiling (64 MB, allkeys-lru as native; Omni on the ceiling inside [16, 512] MB through Redis's own console, grown only while the cache is full), three untouched workloads, 3 pairs × 3 runs	v3	L	work inside the 2 ms line +14% to +27% confirmed better on all three; hit rate +14% to +27% and mean latency −30% to −61% confirmed better; no failed request; the memory ceiling held 64 → 200 to 270 MB and the memory used confirmed worse (the resource the gain costs); p95 within a hair of native's; host CPU inside the noise; every ceiling handed back; 12 gauge-rows better, 6 worse	results/live/V3_REDIS.md
MongoDB as its publisher ships it, a database's operator-set storage-engine cache size (512 MB as native; Omni on the cache size inside [256, 2,048] MB through the server's own console, grown only while the cache is full and reads are slow), YCSB's workloads b, c, f and burst untouched, 3 pairs × 3 runs	v3	L	the cache size held −49% on c, −41% to −49% on burst, −13% to −37% on f, confirmed better ; on b −27% to −50% with one run's interval over zero, no difference beyond the noise; work inside the 1 ms line, p95, p99 and host CPU inside the noise on all four; no failed operation; the mean latency on burst +2% to +4%, confirmed worse ; pages read into the cache +14% to +58% (shown); every cache handed back; 6 gauge-rows better, 1 worse. "Why is the gain memory and not speed?" Because on this machine the data sits in the operating system's page cache too, so a storage-engine miss costs a memory read, not a disk read, as disclosed before the run; the question the knob would face on a disk-bound store is not answered here. "Is the mean-latency loss hidden?" No: it is a confirmed-worse row in the table and in this manual; it is not an index column, which the index's rules name	results/live/V3_YCSB.md
Robustness: the governor killed outright mid-run (SIGKILL at 40% of the window, the watchdog beside it), 10 pairs × 3 runs	v3	L	every setting back at the operator's 7 to 11 s after the kill in 30 of 30 repetitions (allowance 60 s), a second governor to the end in every one, the 120 s after the kill inside the noise against native; the whole window still confirmed better on mean response, p95, time over the line and failed requests	results/live/V3_ROBUS T_KILL.md
The governor's own cost at 1, 10, 100 and 1,000 copies with the real cluster inside	v3 and v1, each under its own engine	L	0.006 to 0.013 of one core at every size, 0.1% to 0.3% of the host's cores; shown, not judged	results/live/V3_OWN_COST.md
Scale: the controller governing 50, 500 and 1,000 simulated nodes (KWOK), decision time and correctness	every push	L	runs on every push	results/scale/

Result	Engine	Class	Reading	Source
GPU, one card and the card inside the organisms	earlier card controller	P	obsolete: every earlier card result ran on a controller since replaced; the one-card, card-inside-1,000-copies and eight-card runs are run again by the founder on rented cards after the CPU and cloud work, at one named commit	docs/GPU_PREREGISTRATION.md, docs/GPU_RUN_GUIDE.md

Running now (8 October): the whole tower at 1,000 copies on a rented machine (the four stacked at 1,000 copies done, section 16.2); MySQL's buffer pool under sysbench as A, B and C (section 10.9); the robustness test (docs/ROBUSTNESS_PREREGISTRATION.md), the kill scenario and the two-hour long run both done (section 16.4b), the 24-hour run on a rented machine next; YCSB on MongoDB done (V3_YCSB.md, section 16.4), MySQL's buffer pool under sysbench preregistered, built and in its smoke run (docs/MYSQL_PREREGISTRATION.md, section 10.9), the other stores of register row 24 after; and Azure steady and burst on the fleet of several machine families (the first two dispatches were refused by the subscription's family allowances before any arm ran, the next by Azure's own cluster capacity in eastus; the fleet is rebuilt from the families the survey shows allowed). **Queued, in order, in docs/REGISTER.md Chapter 7:** drone swarms and defense edge (PX4 and ArduPilot multi-vehicle, Crazyswarm), databases and caches at large (YCSB, HammerDB), Spark, OpenSearch, fio, Open-RMF, the 24-hour robustness run, spacecraft attitude and thrusters (Basilisk), station-keeping (Orekit, GMAT), constellations, rockets (RocketPy, OpenRocket) and combustion (Cantera). Pure physics solvers (OpenFOAM, SU2, REBOUND, GADGET, MESA) are not benchmarked: they have no controller and no knob, so Omni-Compass has no place on them; their value is the HPC cluster that runs them.

Threats to validity, stated by us

A referee will look for the ways these results could mislead. We list the ones we know, what each would do to the numbers, and what the record does about it.

1. **The machines are shared.** Every Kubernetes, database, broker and cache run is on one GitHub runner with four cores, and the governor, the load generator, the probe and the system under test share it. A governor that costs CPU can slow the service it governs, and a noisy neighbour on the host can move a run. The pairing on one runner with rotated order removes the between-machine noise; the host's own busy share is recorded beside every run; the governor's own CPU is a gauge in every table and a confirmed loss where it is one (the database, Kafka's light workload); and three separate runs on separate machines must agree before anything is confirmed.
2. **kind is not a cloud.** Idle machines stay powered and native has no node autoscaler, so a machine given back on kind is a machine-hour in service, not a watt and not a dollar. The energy rows on kind are a declared model and say so on every line. The bill is Azure's test, and its only result so far is inside the noise on a fleet too small to show one machine.
3. **Native is at its knee by design in the broker test.** The Kafka producer's peak offers nine tenths of the capacity measured for the operator's two consumers, so native's queue grows at the high steps and omni's does not. This is the situation the benchmark is of (a group sized for the quiet hours and overrun at the peak), it is preregistered, and it is why the latency ratios are large; the index page and the table say so. A native group sized for the peak would show a smaller latency gain and a smaller consumer cost, and would be a different, also legitimate, benchmark.
4. **Declared penalties.** The cache's 5 ms store trip on a miss and the broker's 1, 2 or 5 ms of CPU per message are declared parameters of our own application, not measurements of anyone's store. The gauges that depend on

them (work inside the line, the request latency) scale with them; the hit rate, the lag, the consumers and the memory held do not.

5. **Simulators are models.** Every S result is a statement about a published model under its own controller. A model may flatter or punish the governor in ways the real plant would not. S results are therefore shown beside the index and never inside it, and the muscles' models are the simplest physics that reproduces the meter's response to the knob, with the native controllers tuned as their authors ship them.
6. **The tuning case was seen.** Every benchmark's rules were set on one tuning case, and that case is shown and not counted. Where a rule changed after a counted result was seen (the swarm table's reproduction tolerance), the amendment is written in the preregistration with its time and reason, the table says so, and the runs were not rerun.
7. **The engine changed twice.** v1 to v2 added muscles without touching the law; v2 to v3 added one gate. Every result is labelled with its engine, the tools refuse to mix engines, and the v1 results stay as v1 results. A reader comparing a v1 number with a v3 number is comparing two engines, and the manual never does so except to say that the controllers were the same bytes and the readings the same.
8. **Three runs are three, not thirty.** The three-run rule is a replication standard, not a large sample. A row confirmed in three runs of ten pairs is a row whose paired difference cleared zero in each of three independent samples of ten; the intervals are shown so a reader can judge their width. Where the intervals are wide, the manual says the test lacked power rather than calling the result a null.
9. **The founder's own harness.** The application in front of the cache, the producer in front of the broker, the load generator and the probe are ours. They are declared, versioned, archived with every run and reproducible from the repository; they are not independent of us. The independent pieces are the systems under test, their native controllers, the simulators and their shipped controllers, and the benchmark tools (pgbench, kind, Kubernetes, PyBullet, MuJoCo, pandapower, CityLearn) that others wrote.

What is not yet shown, and the open program

Three things this manual does not show, and the program that will show them or show their absence:

- **An energy or bill saving on real machines.** The fleet of 40 workers in nine Azure machine families that can show one machine is preregistered and dispatched on v3; it runs when Azure's region admits a cluster. If the saving is a few percent, a 40-worker fleet can see it; if the saving is not there, the table will say "no difference beyond the noise" on a fleet that could have seen it, and that will be the result.
- **The real card.** Every earlier GPU result ran on a controller since replaced. The two-wire governor with the verdict runs on rented cards at one named commit, with the wire check first (`docs/GPU_RUN_GUIDE.md`, `docs/GPU_PREREGISTRATION.md`); until then the card has no result and the manual says so.
- **Robustness over a day.** The governor killed outright is a measured result (section 16.4b: every setting back in 7 to 11 s, thirty times out of thirty), its own cost is tabulated, and the two-hour run is measured three times (no leak, no slowing, every decision expected but a few the cluster's API refused, every setting handed back); what remains open is the 24-hour run on a rented machine, for drift and leaks over a working day, and the same test on the other stacks.

The queue beyond these, in order, is `docs/REGISTER.md` Chapter 7, one or two at a time, each preregistered before it runs: PX4 and ArduPilot swarms, YCSB and HammerDB, Spark, OpenSearch, fio, Open-RMF, the 24-hour run, Basilisk, Orekit and GMAT, RocketPy and OpenRocket, Cantera. When the founder declares the engine final, the engine that stands then is published as Omni-Compass 1.0, and the older fingerprints go to `docs/history` as the road to it.

The Pilot Protocol and Kit

Pilot Protocol

Before you wire anything: read DISCLOSURES.md. Omni-Compass acts only through the wires it is given; it cannot be slapped on. If your paired receipts differ from the published benchmarks in direction, the first presumption is wiring; confirm the installation with section 8.5 of the manual (*Wired right or wired wrong*).

Phase 0: Simulation evaluation (evaluator's own environment)

Run `python verify.py`. Pass criterion: VERIFICATION: PASS.

Phase 1: Shadow (2 to 4 weeks)

- Capture: `OUT=capture.csv INTERVAL=15 DURATION=<seconds> POWER_CMD="<site power in watts>" bash fleet/capture/kube_capture.sh` (read-only).
- Replay: `python fleet/capture_replay.py capture.csv --idle-w <W> --dyn-w <W> --site-limit-w <W> --out replay/` gives Omni-Compass's recommended node count, power cap and HPA target per decision.
- Governor in OBSERVE on production telemetry read directly from nodes (kubelet/cAdvisor metrics, power meters, queue depth).
- No write credentials are issued.
- Logged per interval: directive, native actions taken, and the observed outcome.
- Pass criteria, agreed before start: logged directives never violate I1 to I5; counterfactual analysis on at least N recorded incidents shows the directive would have reduced time-to-recovery or energy without a backlog increase beyond an agreed bound.

Phase 2: Guarded control, one loop at a time (4 to 8 weeks)

Order: power capping; node count (Cluster Autoscaler set to observe); replica count (HPA set to observe).

- Each loop is handed over separately, with the reset tested at handover and at exit.
- Pass criteria per loop, fixed in advance: SLO attainment not worse than the preceding shadow baseline at the agreed confidence level; zero shield invariant violations; energy per unit of completed work reported with confidence intervals.
- Exit: any criterion failed triggers the reset and returns the loop to its native controller.

Scoring your own pilot

Capture the baseline (a period before the controller, or a matched node pool left on your normal autoscaler) and the Omni-Compass period or pool with `fleet/capture/kube_capture.sh`, then: `python pilot/score.py --baseline baseline.csv --omni omni.csv [--idle-w W --dyn-w W]` It reports node-hours and energy per used CPU core-hour, utilisation, pending-pod minutes and HPA shortfall minutes, each with a bootstrap 95% interval over hourly blocks. Energy is measured if the captures include `power_w`. Run long enough for at least 24 blocks per side, and compare matched pools at the same time where possible, because traffic changes between periods.

Phase 3: Component retirement

A decision component is retired only after its loop has passed Phase 2 and a one-at-a-time removal shows no degradation (keep-or-remove rule, Manual Chapter 6). Execution and security components are retained.

Reporting

All pilot metrics, including failures, are reported in the same format as the benchmark results.

Evaluation and simulation use only. Copyright (c) 2026 The Omni-Compass LLC. Commercial use, commercialization or monetization of any part of Omni-Compass requires a signed, paid Omni-Compass Enterprise License from The Omni-Compass LLC. Patent applications, copyright registrations and trademark applications filed in the United States. See LICENSE and NOTICE at the root of this repository.

Omni-Compass shadow pilot kit

What a customer runs first. Omni-Compass watches the cluster read-only beside its own autoscalers. Every 15 seconds it decides what it would do, using the same closure law that was benchmarked. It never writes.

1. `kubectl config use-context <the cluster>`
2. `DURATION=86400 IDLE_W=<watts per idle node> DYN_W=<watts at full load> bash scripts/pilot_shadow.sh`
 - It proves with `kubectl auth can-i` that the identity can read and cannot write, and stops if any write permission exists.
 - It captures telemetry and logs every recommendation to `shadow_out/audit.jsonl`.
 - It writes `shadow_out/SHADOW_REPORT.md`: node-hours used against node-hours Omni-Compass would have used, and the write count, which must be 0.
3. Scoring against a matched baseline: `python pilot/score.py --baseline baseline.csv --omni omni.csv` (bootstrap intervals over hourly blocks).

Guarded control follows `docs/PILOT_PROTOCOL.md`: one loop at a time, the reset tested at each handover.

The kit is exercised end to end on kind by the `live-shadow` workflow.

Evaluation and simulation use only. Copyright (c) 2026 The Omni-Compass LLC. Commercial use, commercialization or monetization of any part of Omni-Compass requires a signed, paid Omni-Compass Enterprise License from The Omni-Compass LLC. Patent applications, copyright registrations and trademark applications filed in the United States. See LICENSE and NOTICE at the root of this repository.

CHAPTER 43

The GPU Bench

Before you wire anything: read `DISCLOSURES.md`. Omni-Compass acts only through the wires it is given; it cannot be slapped on. If your paired receipts differ from the published benchmarks in direction, the first presumption is wiring: confirm the installation with section 8.5 of the manual (*Wired right or wired wrong*).

This is the test that answers "does Omni-Compass save real energy?" with the GPU's own power meter. There is no model in it. Nobody has run it yet on real hardware for this repository; the first run is the first real-meter result.

What you need

- A Linux machine with one NVIDIA GPU (a rented cloud GPU works: A100, H100, L4, A10, RTX), the NVIDIA driver and `nvidia-smi`.
- Python 3 with PyTorch built for CUDA (`pip install torch numpy`).
- Root, because setting a GPU power limit (`nvidia-smi -pl`) needs it. Persistence mode on is recommended (`sudo nvidia-smi -pm 1`); the receipt records it either way.

Run it

```
git clone <this repository> && cd <it>
sudo bash scripts/gpu_paired.sh          # 5 repetitions, 10 minutes per arm: about 3 hours
sudo REPS=5 DURATION=300 bash scripts/gpu_paired.sh  # 5-minute arms: about 1.5 hours
```

Options (environment variables): `REPS`, `DURATION` (s per arm), `DRAIN` (s after arrivals stop), `COOLDOWN` (s idle before each arm), `GPU` (index), `SAMPLE_MS`, `INTERVAL` (s between Omni decisions), `SLO_MS` (response-time target; default ten bare service times), `WORKLOAD_ARGS` (e.g. `--n 8192 --target-ms 80`).

Two phases

- `PHASE=smoke` (the default): look for faults and for an effect worth confirming. Any number of repetitions.
- `PHASE=confirm`: the preregistered test (`docs/GPU_PREREGISTRATION.md`), 10 repetitions. Omni's code must be committed; its files are hashed before the first arm and again after the last, and any change invalidates the run.

```
sudo PHASE=confirm bash scripts/gpu_paired.sh
```

The **primary outcome** is work per energy: requests served per kilojoule the GPU drew.

What it does

1. **Receipt.** GPU name, UUID, VBIOS, driver, kernel, host, persistence mode, power management, the enforced limit, default, minimum, maximum and current power limit, the workload's hash, the mechanism id, git commit (receipt.json). The current power limit is the **snapshot**; every arm must start and end at it. The bench refuses to start unless power management is Enabled (otherwise a written limit would not bind).
2. **Calibrate once.** The workload times its request at the snapshot limit and picks the request size so one request takes about 50 ms (calib.json). Every arm uses the same calibration.
3. **Three arms per repetition, order rotated**, each after COOLDOWN s idle:
 - **native:** no Omni process.
 - **watch:** Omni runs, reads the GPU and decides, and is forbidden to write. This is the control: it shows what the machine does with Omni present and silent. If it writes even once, the run is invalid.
 - **omni:** Omni writes the GPU power limit.
4. **The same work in every arm.** tools/gpu_workload.py sends one seeded stream of requests (fp16 matrix products) at 30%, 60%, 80%, 30%, 60% and 30% of the GPU's full-power capacity. Every arm gets the same requests at the same moments.
5. **Measured by the device.** nvidia-smi samples power draw, temperature, utilisation, the power limit, the **enforced** power limit and the clock-limit reasons every 200 ms for the whole arm (the fields the driver reports, listed in smi_fields.txt); RAPL CPU package counters are read at both ends where the machine has them. This is receipt C, the outcome: Omni never supplies it.
6. **Reset.** After the omni arm Omni restores the snapshot limit and reads it back. The script checks the limit after every arm.
7. **The table** (GPU_REPS.md): each gauge for native, watch and omni, and three paired contrasts with 95% intervals — observation (watch – native), authority (omni – watch), total (omni – native). If an interval includes zero, it says **not proven**. The result label is chosen by rule (docs/GPU_PREREGISTRATION.md, amendment 1): SUPERIOR WITHIN GUARDRAILS, ENERGY IMPROVEMENT WITH SERVICE TRADEOFF, NONINFERIOR / INCONCLUSIVE, NOT ESTABLISHED, WORSE, or INVALID. Also: actuator fidelity (receipt B), control effort and representation fidelity (receipt A). A meter that was not fitted prints UNAVAILABLE. The run folder holds every raw file and SHA256SUMS.txt.

What Omni does on the GPU (omni_controller/gpu_governor.py)

Every 2 s it reads the GPU and runs the Omni-Compass engine (the throughput law in omnicompass/adapter.py): load is GPU utilisation, power stress is draw over the snapshot limit, heat is temperature over 83 C, queue is response-time pressure. The engine's power cap becomes a power limit, inside hard rules:

- never below the GPU's current draw x 1.3, never below 0.70 of the snapshot (--min-share), never below the device minimum, never above the snapshot;
- a busy card (utilisation smoothed over decisions at or over 0.5, --util-gate) gets the snapshot limit back at once, and the cap returns only under 0.4;
- optional speed lock (--baseline-file, from tools/gpu_baseline.py on runs without Omni): the limit follows response time against that baseline, every gauge kept at least 1% faster (docs/INTEGRATION_MANUAL.md, level 5);
- no new write until the last one reads back from the device; every write is read back at once and recorded (requested, return code, read back, enforced limit, delay); the engine senses the enforced limit, the one the card

obeys; a write the device refuses ends the arm (exit 4) and makes the run invalid; no clock locks are ever written;

- if it cannot read the GPU or the response times, the snapshot limit at once;
- if response time breaks its target, the snapshot limit at once, and for three decisions after;
- kill file or SIGTERM: the snapshot limit, read back.
- **one writer**: if the limit ever reads a value Omni did not write, Omni stops writing, leaves that limit alone, and exits 5 (the run is invalid);
- **heat fails up**: while the card reports a thermal or hardware slowdown, no lower limit is written;
- every decision names every rule that held the engine back (blocked_by) and the one that decided (decided_by).

Any workload plugs in through `WORKLOAD_CMD` (for example a vLLM or MLPerf inference harness), if it writes `latency.csv`, `requests.csv` and `summary.json` in `tools/gpu_workload.py`'s format and `SLO_MS` is given.

The table also credits each write: joules and requests against native at the same moments of the same stream, grouped by the rule that decided it. CPU energy comes from RAPL by domain: package and DRAM apart, psys never added, wrap undone. The card's own energy counter (NVML) and its ECC and retired-page counters are read at both ends of each arm.

Gauges

Gauge	From	Better
work per energy (served requests per kJ), primary	requests served / GPU energy	higher
energy, GPU (J)	power.draw integrated over the arm's window	lower
energy per served request (J)	the same, over requests served	lower
power, GPU mean (W)	energy / window	lower
requests served, not served	the workload's own record	more served, fewer not
response time mean, 95th, 99th percentile (ms)	arrival to finish of each request	lower
temperature, peak and mean (C)	temperature.gpu	lower
energy, CPU package (J)	RAPL counters, when present	lower

Testing the bench without a GPU

`python tests/test_gpu_bench.py` runs the whole script against a stand-in `nvidia-smi` (`tests/fake_gpu/`), with the workload's `--sim` mode. The stand-in has no real power physics, so its numbers mean nothing; it proves the script, the controls and the validity checks work.

Evaluation and simulation use only. Copyright (c) 2026 The Omni-Compass LLC. Commercial use, commercialization or monetization of any part of Omni-Compass requires a signed, paid Omni-Compass Enterprise License from The Omni-Compass LLC. Patent applications, copyright registrations and trademark applications filed in the United States. See LICENSE and NOTICE at the root of this repository.

CHAPTER 44

The GPU Preregistration

Written before any hardware trial. The confirmation run (PHASE=confirm) hashes this file with the code in `FREEZE.json`; a change to either after the run starts invalidates it.

Question

On one NVIDIA GPU serving a fixed, seeded request stream, does Omni-Compass holding the GPU power limit change the successful work done per joule measured by the device, compared with the device left at its own limit?

Fixed here, before any smoke trial:

- the number of confirmation repetitions (10);
- the primary outcome;
- the guardrails;
- the analysis.

Nothing seen in smoke may change them.

Two phases

1. **Smoke** (PHASE=smoke, any number of repetitions). Its purpose is to find faults in the harness, and to see whether the effect is large enough to be worth confirming. Smoke results are not published and are never reported as the result. After smoke, Omni's code may change.
2. **Confirmation** (PHASE=confirm). Omni's code is committed and frozen before the first trial. The script refuses to run if any frozen file has uncommitted changes, and the table marks the run invalid if a frozen file changes during it. No inspection, tuning or rerun between confirmation trials. If the confirmation fails, it is reported as failed; a new confirmation needs a new commit and a new run, and both runs are reported.

Design

- **Arms:** native (no Omni), watch (Omni runs, may not write), omni (Omni writes the GPU power limit).
- **Repetitions:** 10 in confirmation, each with all three arms back to back, order rotated. Repetitions may run on separate machines of one GPU type (REP_ONLY, the GitHub workflow); the three arms of a repetition always share one machine, so every comparison is paired within a machine.
- **Per arm:** 60 s idle, then the pinned workload for 600 s plus 30 s drain.
- **Workload:** `tools/gpu_workload.py`, calibrated once at the start power limit before any arm. The seed is 20260928, and the load phases are 30, 60, 80, 30, 60 and 30% of full-power capacity.

Outcomes

- **Primary:** work per energy = requests served / GPU energy (kJ). GPU energy is `nvidia-smi power.draw` integrated over the arm's window.
- **Secondary:**
 - GPU energy;
 - energy per served request;
 - requests not served;
 - response time: mean, 95th and 99th percentile;
 - peak temperature;
 - CPU package energy (RAPL), where present;
 - where a smart plug is fitted (`WALL_METER`), whole-machine energy at the wall and work per wall kJ. The plug is read by the bench only, never by Omni. An arm whose plug readings have a gap over 5 s has no wall number.

Analysis

- **Primary comparison:** omni against native, paired by repetition. Report the mean difference with a two-sided t 95% interval.
- **Result:**
 - The result is *proven better* if the interval lies entirely above zero.
 - It is *proven worse* if the interval lies entirely below zero.
 - Otherwise it is *not proven*.
- **Guardrails, fixed now:** a better primary result counts only if Omni did not buy it with the work. Both guardrails must hold:
 - **Requests served:** the 95% interval of (omni – native) must not reach below –1% of native.
 - **95th-percentile response time:** the interval must not reach above +10% of native.

If a guardrail fails, the verdict is *better on energy*, *fails the service guardrail*. That is a different product and is reported as such.

- **Watch against native:** reported as the cost of Omni being present. If watch differs from native on the primary outcome as much as omni does, the effect is not attributed to Omni's authority.

Amendment 1 (2026-09-28, before any hardware trial; no smoke or confirmation data exist)

Added before any data, to make the chain from engine to plant identifiable. The question, the arms, the primary outcome, the repetitions, and the two guardrails above are unchanged.

- **Three contrasts, all reported, each paired by repetition:** observation = watch – native; authority = omni – watch; total = omni – native (the primary comparison). A total effect is not attributed to Omni's authority where the observation contrast differs materially from zero.
- **Third guardrail, errors:** the 95% interval of (omni – native) requests *not* served must not reach above +1% of native requests served.

- **Result label, by rule, never by hand** (`tools/gpu_reps.py, label()`):
 - primary proven better, all three guardrails held: **SUPERIOR WITHIN GUARDRAILS**;
 - primary proven better, a guardrail failed: **ENERGY IMPROVEMENT WITH SERVICE TRADEOFF**;
 - primary not proven, all guardrails held: **NONINFERIOR / INCONCLUSIVE**;
 - primary not proven, a guardrail failed: **NOT ESTABLISHED**;
 - primary proven worse: **WORSE**;
 - any invalidity below: **INVALID**.
- **The card obeys enforced.power.limit.** The snapshot records `power.limit`, `enforced.power.limit`, `default`, `min` and `max` limits, `persistence mode` and `power management`. The bench refuses to start unless `power management` is `Enabled`. The bench samples `enforced.power.limit` where the driver reports it. The governor senses the enforced limit and reads back every write at once; it never writes clock locks.
- **Also invalid:** a native or watch arm whose enforced limit differs from the snapshot; a governor that exits nonzero (a refused start, or a write the device refused — that write ends the arm); smoke and confirmation repetitions mixed in one table.
- **Three receipts, kept apart.** A (governor, `audit.jsonl` decisions): telemetry consumed; the six-state reading (history-dependent, `state_observed`); the memoryless state the telemetry alone points to (`state_measured`); the U-channel command `u` evaluated on the evolved state (`u_push`); admissibility; requested and granted authority; shield bound; holds. B (actuator, actuator records): requested limit, return code, `power.limit` read back, enforced limit, delay to realization, override (enforced under requested), restoration. C (outcome, the bench alone): `nvidia-smi power`, joules, temperature, utilization, clock-limit reasons; the workload's requests, latency and failures. Omni never supplies its own outcome.
- **Secondary, descriptive (no verdict):**
 - actuator fidelity: `r_act` = read back – requested (mean and max absolute), delay (median, max), enforced-under-requested count, total variation of the realized limit, reversals, refused writes;
 - control effort: $J_u = \text{sum of } \text{abs}(u_push) \times \text{decision interval}$; mean and max `abs(u_push)`; saturations (`abs(u_push) = 25`); shield interventions (a bound other than the engine decided); holds;
 - representation fidelity: $R_{int} = \sqrt{\text{mean over consecutive decision pairs of } \sum_k w_k (h(z_{(k+1)}) - F_h(h(z_k), u_k))_k^2}$ over E, U, I_U, S, B, weights 1, with `h` the memoryless map and `F_h` the engine's own projection (`state_projected_next`); directional accuracy = share of pairs where `sign(projected – current reading)` equals `sign(next measured – current measured)`, both movements at least 0.01, per state and by predicted size (0.01–0.03, 0.03–0.1, ≥ 0.1);
 - energy of the rest of the machine = wall – GPU – CPU package, only when all three meters measured the arm. A missing meter prints `UNAVAILABLE`, never a modelled substitute.
- **Smoke never enters confirmation.** A smoke run that changes any code or parameter is followed by: repair, the tests, a new freeze, and a confirmation collected from zero.

The run is invalid, and reported as invalid, if

- the watch arm executes any power-limit write;
- a native or watch arm sees a power limit other than the snapshot;
- any arm ends at a limit other than the snapshot (the reset failed);
- Omni's frozen files change during the run;

- the confirmation runs on uncommitted code.

Recorded for every Omni decision

Every decision is logged in `audit.jsonl`:

- the raw device telemetry: utilisation, draw, temperature, limit, SM clock, and clock-limit reasons where the driver reports them;
- the engine's six-state reading, the state it had projected for this moment, and the error against it;
- the projection for the next decision. This is the engine's own evolved state from the same step that sets the cap; no separate predictor was added for the experiment;
- whether change was admissible, the requested and granted authority, and which shield bound decided the limit;
- the limit written, and the requests served in the window.

Amendment 2 (2026-09-28, before any hardware trial; no smoke or confirmation data exist)

Unchanged:

- the question;
- the arms;
- the primary outcome;
- the repetitions;
- the guardrails of amendment 1.
- **Watch must match native.** If the observation contrast (watch – native) on the primary outcome is proven in either direction, the label is **NOT ATTRIBUTABLE: WATCH DIFFERS FROM NATIVE** and no omni result is published.
- **One writer.** If the power limit ever reads a value that is neither Omni's last write nor the limit before it, another writer is present:
 - the governor stops writing for the rest of the run;
 - it leaves that writer's limit alone;
 - it exits 5;
 - the run is invalid.
- **Heat fails up.** While the device reports a thermal or hardware slowdown (clock-limit reason bits 0x8, 0x20, 0x40, 0x80), no lower limit is written.
- **Credit per write (descriptive).** Each write owns the interval to the next. For that interval, the table records GPU joules and requests finished, Omni minus native, at the same moments of the same seeded stream. It records who decided the write: the engine, a floor, the busy gate, a reflex, heat, or the speed lock. The table says which rule produced the joules; a speed-lock result is not credited to the engine.
- **CPU side (secondary, never on the control path).** RAPL counters are read by domain name at both ends of each arm:
 - package-N is summed as CPU package; dram is summed separately;
 - psys is recorded and never added to either;

- wrapping is undone with `max_energy_range_uj`;
- a counter that went backwards without a known range, or is missing, prints `UNAVAILABLE`.
- The governor never reads these counters.
- **Device energy counter (cross-check, secondary).** Where NVML reports it (Volta and newer), the card's total-energy counter is read at both ends of each arm, beside the integrated power.draw.
- **Card health (descriptive).** Uncorrected and corrected ECC error counts, and pages pending retirement, are read at both ends of each arm.
- **Workload plug.** Any workload may be served through `WORKLOAD_CMD` if it writes the pinned workload's files. It needs its own response-time target (`SLO_MS`). The command is recorded in the receipt. A confirmation names its workload before the first trial.

Amendment 3 (2026-10-01, before any hardware trial; no smoke or confirmation data exist)

No GPU job has ever been given a machine (every `gpu-bench` run so far waited in the queue and was cancelled), so no trial data exist. Unchanged:

- the question;
- the arms;
- the primary outcome;
- the repetitions;
- the guardrails, labels and invalidity rules of amendments 1 and 2.
- **Declared envelope, before any trial.** The buyer's service envelope is written to a file before the first trial and recorded with the run (`envelope.json` in the run folder and in every repetition):
 - `power_min_w`, the lowest watts Omni may set. Default (`tools/declare_envelope.py`): `max(device minimum, 70% of the power limit read at declaration)`, rounded up to whole watts;
 - `power_max_w`, the power limit read at declaration;
 - optionally `slo_ms`, the response-time target; without it the target comes from calibration (10 bare service times), as before.
 - The bench refuses an envelope whose floor lies outside [device minimum, starting limit].
 - **The confirmation refuses to start without a declared envelope** (`scripts/gpu_paired.sh`, `ENVELOPE`).
- **Envelope floor.** The governor never sets the limit under `power_min_w (--floor-w)`. When the floor is what lifted a write, the decision record names it (`decided_by: envelope_floor`), so no saving below the floor can be credited to the engine.
- **Outer controller holds.** When the card's own controller (board, BMC or system policy) already holds enforced.power.limit under the current limit, a lower write that would still sit above that enforced limit changes nothing on the card. It is not written; the decision record marks it (`outer_controller_holds`). Only a write that would actually bind, or a return upward, goes out. This is not another writer (amendment 2): the set limit is unchanged and the governor keeps running.
- **Narrow cards are reported as they are.** The device's own limit range is in the snapshot. On a card whose range is narrow (for example a 70 W card that accepts 60–70 W), the envelope is that narrow range; the result is reported for that card and range and not extrapolated to wider cards.

Amendment 4 (2026-10-02, after an invalid smoke; no confirmation data exist)

The first smoke on a rented A10 (results/gpu/smoke-20261002T032459Z, never counted) was invalid by rule: two copies of `scripts/gpu_rented_run.sh` had been started on the same machine, so both benches wrote the same card's power limit. The native and watch arms saw limits they never wrote (116, 137 and 150 W), each governor refused to start beside the other one (exit 5), and one reset restore was undone by the other copy. The card also began at 116 W, a limit an earlier start had left behind, not its 150 W default, so native itself ran capped (83% of samples). None of these numbers measures Omni. The engine, the governor, the outcomes and the analysis are unchanged. The run script now:

- **runs once per machine:** it takes a lock and refuses to start while another copy runs;
- **runs alone on the card:** it refuses to start while any other process is using the GPU;
- **starts from the card's default limit:** it sets `power.default_limit` before the envelope is declared, so the envelope, the snapshot and every arm start from the card's own default, not from a limit an earlier run left behind.

Amendment 5 (2026-10-02, before any valid hardware trial; the only smoke so far was invalid, amendment 4)

The Omni arm changes engine. The outcomes, the arms' order, the guardrails, the analysis and the validity rules are unchanged.

- **The Omni arm holds two wires** (`omni_controller/gpu_compass.py`, the compass law of `omniconpass/compass_law.py`): the clock ceiling (`nvidia-smi -lgc`, reset with `-rgc`; cover 35% of the top clock to the top), which sets how high the card's own boost may climb, and the power limit (`-pl`), the lid at what a fully busy card draws at that ceiling plus 10%, never under the declared envelope floor and never over the start limit. The service is read as one position between calm and the response-time line (the worse of p95 and utilization above half) and pulled to the middle; past 95% both wires go to full at once (fail up). The card's firmware keeps its own control; Omni sets only those two values. The earlier power-limit-only governor stays available (`OMNI_ENGINE=one_wire`) and is not the confirmation's arm.
- **Why, before the run:** on a modelled card (results/sim/gpu_two_wire/, evidence class S, seeds never used while tuning) the one-wire governor gave +0.1% work per energy and the two-wire engine +9.0%. That is a model; this run is the card's own meter.
- **The watch arm** runs the same two-wire engine in watch mode: it computes and records both wires and writes neither.
- **The clock range is reset before and after every arm** (`-rgc`), as the power limit already was; the run script resets it once at the start.
- **The wire check runs first** (`tools/gpu_wire_check.py`): the card's clock must follow a lowered ceiling down and come back up when reset, the power limit must read back what was set, the governor must hand both wires back when stopped and must leave a limit set by another writer alone (exit 5). If any step fails, nothing else runs and the check's report names the wire, the step and what the card said.
- **Wire check, corrected before any trial (2026-10-02):** on the A10 the first wire check failed at "3 up wire (follows up)" although the wire works: under the heavy check load the card's own 150 W limit already held it near 990 MHz, so a ceiling at 60% of the top clock (1017 MHz) left no room for the clock to come back up above it. The check now reads the card's busy clock on its own first and locks at 60% of that. The governor likewise starts its ceiling at the clock the busy card actually runs (a ceiling above it holds nothing), and its

clock cover is 35% of the top clock to the top. No trial had run.

Amendment 6 (2026-10-02, after the first confirmation and before any further trial)

What the first confirmation showed (A10, commit c908054, results/gpu/run-20261002T082232Z/): work per energy +3.6% (+2.7% to +4.5%, proven), the same requests served, no request lost, but the 95th-percentile response time +58.5% (510 to 809 ms), so the response-time guardrail failed and the label by rule was ENERGY IMPROVEMENT WITH SERVICE TRADEOFF. It stands as the result of that run.

Why, from the card's own samples: while busy the card ran at 736 to 768 MHz under Omni against 861 to 889 MHz on its own, and spent about 30% more time busy for the same work; requests queued behind each slower one. Three faults in the governor, not in the engine: (1) the position counted utilization above half as service trouble, so every burst read as past the wall (fail up in 46% of decisions) and every quiet gap pulled the ceiling down, so each burst began on a lowered clock; (2) the lid followed a curve from the top clock and sat at the 105 W envelope floor in 49 of 165 compass decisions, under the 135 W the card itself draws while busy; (3) the ceiling's cover reached 35% of the top clock, far under the clock the card's own power limit holds it at while busy.

The Omni arm from now on (omni_controller/gpu_compass.py; the outcomes, arms, guardrails, analysis and validity rules are unchanged):

- the position is response time only (p95 over 5 s, not 30 s); being busy is not a breach;
- **race while work waits:** at 95% utilization or more the ceiling goes to the top and the lid to the start limit; the compass paces only the slack between bursts;
- **the card's own level, learned from its own meter** while the ceiling is at the top and the card is busy: its busy clock (median) and busy draw (90th percentile); until 15 such readings are in, neither wire moves;
- **speed floor:** the ceiling never goes under the card's own busy clock; **lid floor:** the lid never goes under the card's own busy draw plus 10%;
- fail up (past 95% of the line, or blind) is unchanged.

On the modelled card, before any trial (results/sim/gpu_two_wire/, seeds 5000 to 5009): p95 122.1 ms native, 123.7 ms with the corrected compass; work per energy +8.2% (+6.3% to +10.1%); energy -7.5%; the median response 10.1 to 12.1 ms, slower in the quiet stretches the compass paces. That is a model; the next trial is the card's own meter.

Amendment 7 (2026-10-02, before any further trial)

The outcomes, arms, guardrails, analysis and validity rules are unchanged. The Omni arm of the confirmation runs the **service** profile.

- **Two profiles, one switch** (--profile, omni_controller/gpu_compass.py; the same in realms/gpu_card.py):
 - **service**, the default and the confirmation's arm: down gain 0.0125, the compass's center at 0.4, the speed floor 3% above the card's own busy clock;
 - **batch:** down gain 0.015, center 0.5, the floor at the card's own busy clock, for work nobody waits on answer by answer. It may be run as a separate, declared confirmation (OMNI_ARGS="--profile batch") and is reported as its own result, never pooled with service.

- **How service was chosen, on the model, before the trial** (20 paired seeds, 5000-5009 and 5100-5109; ratios summarised as the geometric mean of the per-seed ratios): down gains 0.01, 0.0125, 0.015, 0.0175 and 0.02 alone, and 0.0125 to 0.0175 crossed with the compass's center (0.4, 0.5), the speed floor (1.00, 1.03 of the card's own busy clock) and the race threshold (0.90, 0.95). The rule: the most work per energy at which no seed's p95 is more than 10% slower than native. Service (0.0125, 0.4, 1.03) gave work per energy +5.3% (+4.2 to +6.5), energy -5.0%, p95 -4.1% (-9.7 to +1.7), worst seed +9%, 0 of 20 seeds more than 10% slower, p99 -1.8%; its neighbours gave the same within a point. 0.015 alone gave +6.2% but 3 of 20 seeds 10% to 48% slower at p95; 0.01 alone gave +3.8% with p95 -1.9%.
- **The ceiling moves in whole clock steps** (--min-change-mhz, 15 MHz), as the card's own clock does and as the model moves it: the force times the gain, as a share of the top clock, is rounded to whole steps, and a pull under half a step moves nothing and is not stored up.
- **The position reads the mean response time of the window** (as the model does), between the bare service time and the line; the 95th percentile at or past the line, or any failed request, is past the wall (fail up).
- **The model's report summarises ratios on the log scale** (tools/run_gpu_card.py): the arithmetic mean of per-seed percentages let one seed (+240%) stand for twenty.

Amendment 8 (2026-10-03, before any further trial)

The outcomes, arms, guardrails, analysis and validity rules are unchanged, except as stated here.

- **Omni-Compass moves the card only where it measures that the card is no worse for it** (omnicompass/verdict.py, in omni_controller/gpu_compass.py and realms/gpu_card.py). While the service is calm, the governor runs a paired trial:
 - first the ceiling at the top until 30 requests are measured;
 - then one 15 MHz step past the deepest step already allowed, until 30 more are measured.

Each request's cost is the card's own time on it: the workload's new service_ms column, start to done, with the wait in the queue left out. The step is allowed if its median cost is at most **2%** above the median at the top. Otherwise it is refused and not tried again for 900 decisions. The compass may move the ceiling only between the top and the deepest allowed step. Where no step passes, the ceiling stays at the top and the card runs as it does alone. Every trial and every judgement is in the audit (verdict, verdict_state, verdict_deepest_step).

- **One law, no profiles.** The service and batch profiles are removed. The law is: down gain 0.0125, the compass's center 0.4, the speed floor at the card's own busy clock, and the verdict's allowance of 2%. The batch profile paced past the 2% allowance, so it is gone.
- **Why 2%** (the card model, 20 paired seeds):
 - an allowance of 0 leaves the card native on every workload, because every clock step down adds some time to a request, and saves nothing;
 - 2% is the smallest allowance that saves energy, and it is below what a person or a service contract can notice.
- **A second workload, AI token generation** (tools/gpu_workload.py --kind decode): fp16 matrix-times-vector over 1 GiB of weights per pass, batch one, so each pass streams every weight from memory, as token generation does. It runs as a second confirmation of the same design (10 repetitions × 3 arms × 600 s) after the compute-bound one (scripts/gpu_rented_run.sh; SKIP_DECODE=1 skips it). It is labelled by the same rule and reported as its own result, never pooled.

- **Every comparison is a base alone against the same base with Omni on top.** The model's report (`tools/run_gpu_card.py`) runs two bases:
 - the card's own firmware, alone and with Omni on top;
 - an operator's fixed 105 W power cap, alone and with Omni on top.

The old one-wire governor is removed as an arm (it was Omni's own earlier version).

- **On the modelled card, before any trial** (`results/sim/gpu_two_wire/`, firmware alone against firmware with Omni on top; tuning seeds 5000-5009, fresh seeds 5100-5109):
 - compute-bound work: work per energy +0.70% (tuning) and +0.48% (fresh); median +1.56% and +1.47%; p95, p99 and time over the line unchanged within their intervals;
 - AI token generation: work per energy +3.36% and +3.86%; median +0.55% and +0.70%; p95 +0.29% and +0.26%; time over the line 0.

That is a model; the next trial is the card's own meter.

Amendment 9 (2026-10-03, before any trial on this code)

Everything in amendment 8 stands. Four additions, each run by `scripts/gpu_rented_run.sh` after the two confirmations and each reported as its own result:

1. **Steady under the limit** (`omni_controller/gpu_compass.py`, the same in `realms/gpu_card.py`). While the card is saturated against its own power limit (work waiting, the draw at 97% of the limit or more), the firmware boosts a step, hits the limit and is knocked back. The ceiling is then held at the card's own busy clock under that limit, so the same watts serve the work without the knock-backs. It is never held under that clock, the lid stays at the start limit, and a blind feed still fails up. In the card model, fully loaded under a fixed cap, this gives about +0.6% requests served from the same watts, with p95 about 2% faster.
2. **An operator's power cap underneath.** The card's limit is set to the envelope's lowest watts (70% of its default) before the run, and the paired bench runs native (the cap alone), watch and Omni-Compass on top of the cap (the lid never above the cap). It runs twice, 5 repetitions $\times$ 3 arms $\times$ 300 s each time:
 - at the usual load;
 - fully loaded (arrivals at 130% of the card's capacity under the cap), so the result is requests served from the same watts.

The card is returned to its default limit afterwards.

3. **The GPU fault drill** (`scripts/gpu_fault_drill.sh`). With the request stream running:
 - the governor is killed outright, and the watchdog must hand the card back;
 - the master switch is pulled, and the governor must hand back, exit and refuse to restart while OFF;
 - the response feed is paused, and the governor must fail up.

Every check must pass, and the card must end at its start limit.

4. **Real AI serving** (`scripts/gpu_vllm.sh`, `tools/llm_workload.py`). An open language model (Qwen/Qwen2.5-0.5B-Instruct) is served by vLLM, installed in its own environment. It is asked the same seeded stream of prompts in every arm, each for exactly 128 new tokens; 5 repetitions $\times$ 3 arms $\times$ 300 s. Tokens per second and tokens per kilojoule follow from requests served. If vLLM cannot be installed or started on the machine, the stage says so and nothing measured before it changes.

Amendment 10 (2026-10-03, before any trial on several cards)

Everything in amendments 8 and 9 stands. **Several cards in one server** (`scripts/gpu_8card.sh`): every card runs `scripts/gpu_rented_run.sh` at the same time, each on its own card with its own wire check, envelope, smoke, the two confirmations and the power cap underneath, on the same committed code. Card i starts its arm rotation i steps later (`ROT_OFFSET`), so no arm always meets the same neighbours. Each repetition stays paired within its own card (its three arms on one card); the pooled table per workload holds every card's repetitions (card c , repetition r is pooled as repetition $100(c+1)+r$) and is labelled by the same rule. Each card's own table is reported beside it; a card whose table is invalid is reported, never dropped. The fault drill runs once, on the first card, after every card has finished. Last, **one model across every card** (`scripts/gpu_vllm.sh` with `GPU = every card`): vLLM tensor parallel, Qwen/Qwen2.5-7B-Instruct, one governor per card on its own card's two wires, all reading the same response times; the energy is every card's summed; 5 repetitions $\times$ 3 arms $\times$ 300 s, labelled by the same rule. Each card's energy is read from that card alone (the bench's sampling filtered to the card in each repetition's own receipt).

The short design (`FAST=1`, written before any trial on several cards): each card runs one smoke round and 3 repetitions of the compute confirmation (24 paired repetitions on 8 cards, pooled and labelled by the same rule, each card's own table reported beside it), then one model across every card with 3 repetitions. AI token generation and the power cap underneath are measured on the one-card machine (amendment 9) and are not repeated here.

The whole stacks with a real card inside, on several cards (written before any trial): `tools/run_hil.py` as on the one-card machine (amendment 9, the six organisms at 1x, 10x, 100x and 1,000x, repetitions 3, 3, 2, 1 by size, the adaptive step), each organism on its own card at the same time (organism i on card i ; with fewer cards than organisms, the next organism waits for a card). Each organism's receipt is its own; nothing is pooled across organisms.

The pooled design (`POOLED=1`, written before any trial on several cards): every stage of the whole design, with the repetitions pooled across the cards. Each card runs one smoke round; 3 repetitions of the compute confirmation and 3 of the AI token generation confirmation (24 paired repetitions per test on 8 cards, pooled and labelled by the same rule, each card's own table beside it); 2 repetitions of each power cap load (16 per load); the six organisms with the card inside, one repetition per size, each organism on its own card; then one model across every card, 3 repetitions; the fault drill once. Arm durations are those of the whole design.

The time budget (`MAX_HOURS`, written before any trial on several cards): a run on several cards ends by its budget whatever happens. At the budget the master switch is pulled (every governor hands its card back), every stage still measuring stops and is reported as incomplete (an arm cut short is never counted: its repetition has no complete pair), the stages not started are skipped and named, and the finished ones are packed.

Amendment 11 (2026-10-04, before any trial on this code)

The six organisms with the card inside (`tools/run_hil.py`), first run 2026-10-02 at commit `c908054`, labelled the card's part SUPERIOR in five organisms while its 95th percentile rose from about 500 ms to 600-935 ms in every organism. The harness counted the card's work and energy and never its response time (its time over the line was fixed at zero). That run is history: its governor was replaced in amendments 6 and 7, and its labels do not meet the founder's rule.

From this amendment the card harness measures what it never counted:

- the card's time over the line (requests slower than the line, ten bare service times, percent of requests) and its 95th percentile, each paired against native;

- each row keeps its preregistered label and names, beside it, every measure that came out worse than native by any amount, so what Omni costs is read beside what it gains; every change is read in words (better or WORSE, and what it means); no new line is drawn. The 2% bound stays where the founder set it: inside the engine, as the trigger of the verdict (`omnicompass/verdict.py`: a knob moves only where a paired trial shows the muscle at most 2% worse);
- the stacks' own label (`realms/harness.py label`) puts the band first: its guardrail is the mean time over the line at or under native's, where it was the upper end of the interval under +1 point;
- at one copy, each organism runs 5 paired repetitions (was 3), so an energy change of 1-3% on the card's own meter can be told from noise; 3, 2 and 1 at 10, 100 and 1,000 copies as before.

The single card and the eight cards run this code. A card run started before this amendment is reported in full from its own raw files (its response times included) before any figure is quoted.

Why the stacks were late more often in the first run (+0.25 points) and are not now (-0.03): two knobs moved in commit 2a9285d. Omni could loosen the operator's HPA target to 0.95 (fewer pods, bursts wait); it is now held at the operator's. A machine went back with no headroom (the next burst waits for a boot); it now goes back only when the rest covers the recent peak at 0.6 of the release level. Each alone leaves the stacks late more often; together they are late less often. The headroom is the least that keeps every seed un-late (`results/realms/RELEASE_MARGIN_SWEEP.md`).

Every report reads each change in words (better or WORSE, and what it means), so a sign is never read alone.

Amendment 12 (2026-10-05, before any trial on this code)

The card model rerun on the current code (`tools/run_gpu_card.py`, 2026-10-05) showed amendment 9's steady hold making the card slower on its own firmware: compute work, p95 122 ms native against 149 ms with Omni on top, time over the line +0.39 points. A bisection over every commit since the last clean model run puts the change in commit bd455f9 (amendment 9, item 1). Holding the ceiling at the card's median busy clock takes away the boost the firmware uses to serve a burst at its own factory limit. Under an operator's cap the same hold is what helps: the firmware sawtooths against the low cap, and the hold smooths it.

From this amendment the hold applies only under an operator's cap: the start limit under the card's factory limit, read from the card (`power.default_limit`). If the card will not say, it is treated as no cap, so no hold, and the firmware is left to serve the burst. On the card's own limit, saturated, Omni races, as it did before amendment 9. Test: `tests/test_gpu_compass.py` (no hold at the factory limit; the hold under a 150 W cap on a 214 W card).

The card model with this amendment, 20 seeds (tuning and fresh). Under the cap: p95 6.4% and 7.1% faster, p99 3.2% and 5.8% faster, time over the line 1.3 and 1.0 points lower. On the firmware: energy 0.5-3.7% lower, p95 even. The median is 0.6-1.5% slower, inside the verdict's allowance (`results/sim/gpu_two_wire/`). The real card is the test.

Evaluation and simulation use only. Copyright (c) 2026 The Omni-Compass LLC. Commercial use, commercialization or monetization of any part of Omni-Compass requires a signed, paid Omni-Compass Enterprise License from The Omni-Compass LLC. Patent applications, copyright registrations and trademark applications filed in the United States. See LICENSE and NOTICE at the root of this repository.

The Realms Preregistration

Round 3 is the current one: its section at the end changes only how the realms are made up. Round 2's section replaced round 1's Omni layer. Round 1 below is kept as it was frozen. Each round was written and committed before its confirmation seeds were run. Evidence class **S**: every number the run produces comes from a declared model. Nothing here is a meter, and nothing here is evidence about a real machine.

Question

For each of the 656 muscles of the canonical tower, and for each realm and the whole tower run as one organism: does the frozen governor (`omnicompass.adapter.Governor`, the engine with $u = 0$, the stack law, the reset), holding that muscle's one knob on top of the plant's native controller, change work per energy against the native controller alone, without buying it with service?

What runs

- **The catalog** (`realms/catalog.csv`, built by `tools/realms_catalog.py`): the 656 rows of the XPASS package's canonical tower, unchanged, plus four columns given by fixed rules: realm, plant, parameter set, knob.
- Compute / AI / Cloud: 250 muscles. Physics / Robotics / Autonomous: 88. Energy / Facility / Industrial: 121. Distribution / Specialized: 197.
- **Five plants** (`realms/plants.py`), each with its own native controller, all parameters in `realms/presets.py`:
 - `compute_pool`: request stream, servers with start-up delay, idle and dynamic power; native: the HPA rule (10% tolerance, scale-down stabilisation window), fixed admission limit, full clock;
 - `thermal_zone`: zone heat balance, staged cooling units, COP from supply and outdoor temperature; native: PI on a fixed setpoint, units staged to the load;
 - `energy_storage`: site load, solar, battery, grid connection with a contract limit; native: self-consumption above a fixed reserve;
 - `motion_axis`: point-to-point moves from a task queue, motor copper losses and heating, derating; native: PID with feedforward at full speed;
 - `process_loop`: first-order process with dead time, pump or heater power; native: PI on a fixed setpoint.
- **A muscle is one knob of one plant.** Omni holds only that knob: capacity (the adapter's capacity law for one plant), setpoint (inside a declared band: calm end at ρ_0 , stress end at $\rho_{\min}$), power (the directive's power cap), or admission (native while change is permitted; while not, only what clears inside the service target, or the flexible share deferred or shed).
- **Each muscle's plant is sized** by a factor 0.6 to 1.4 drawn from its muscle id.

Arms

On the same seed, so with the same demand, weather and disturbances:

- **native**: the plant and its native controller;
- **watch**: the governor reads every period and writes nothing;
- **omni**: the governor holds the knob; at 90% of the run it is killed and the knob returns to the native controller;
- **fixed_calm** (setpoint muscles only, descriptive): the native controller with the setpoint fixed at the band's calm end. It shows how much of a setpoint result the band alone gives.

Organisms: every plant of a realm (or all 656) on one 15 s clock for one hour. Plants are coupled:

- the electrical power of compute, motion and process plants is heat in the realm's thermal zones;
- the organism's load swing is load on its storage sites;
- the zones' temperature is the ambient every other plant reports.

One governor reads the organism's aggregate, and its one directive sets every muscle's knob.

Seeds and repetitions

Confirmation seeds 1000 to 1009: ten paired seeds per muscle and per organism. Development used seeds 0 and 1 only. Nothing from development seeds is reported.

Outcomes

- **Primary**, per seed: $(\text{work_omni} / \text{work_native}) / (\text{energy_omni} / \text{energy_native}) - 1$. Work is in the plant's own units: requests, IT heat held in specification, site load served, moves completed, or product delivered in specification. For an organism, work is the mean over its plants of $\text{work_omni} / \text{work_native}$, and energy is total joules.
- **Guardrails**, both must hold:
 - work: the 95% interval of $\text{work_omni} / \text{work_native} - 1$ must not reach below -1% ;
 - violations: the 95% interval of the change in the share of periods in violation must not reach above $+1$ percentage point.

Violation, per plant:

- compute: response time over the service target, or a dropped request;
- thermal: zone over its limit;
- storage: grid import over the contract limit;
- motion: tracking error over its bound, winding over temperature, or the oldest task waiting past its deadline;
- process: the process variable outside its specification.
- **Label, by rule** (`realms/harness.py, label()`, the GPU bench's rule):
 - SUPERIOR WITHIN GUARDRAILS;
 - ENERGY IMPROVEMENT WITH SERVICE TRADEOFF;
 - NONINFERIOR / INCONCLUSIVE;
 - NOT ESTABLISHED;
 - WORSE;

- INVALID.

Invalid

A muscle or organism is labelled INVALID if, on any seed:

- the watch arm differs from native in any meter or writes anything;
- the omni arm writes after the kill;
- a knob is not back at its native value after the kill;
- any contrast is not finite.

Recorded

Per muscle and organism:

- every per-seed contrast;
- writes per run;
- the native violation share;
- the fixed-setpoint comparison where it applies.

Per run: the commit, the seeds, and the fingerprints of the catalog, the engine, the governor and the whole frozen tree (RUN.json, SHA256SUMS.txt). The confirmation refuses to run on uncommitted code.

Changes made on the development seeds, before this freeze

The development runs found bugs and sizing faults. Each was fixed before the confirmation seeds were run:

- **The admission knob was inverted:** it admitted only what cleared the service target while change was permitted. Fixed to the declared rule.
- **The compute admission limit compared the queue limit with the period's arrivals** instead of the backlog left after the period's service. Fixed for native and Omni alike.
- **The capacity law did not clip the queue and load readings to [0, 2],** as the adapter's observe_vector does. Fixed.
- **Thermal and process utilisation was measured against the uncapped capacity** while the power cap was held, so the power law's trim ratcheted to its floor. Fixed: utilisation of the capacity actually available.
- **Compute plants reported a heat proxy of my own.** Replaced by the live controller's own thermal model (omni_controller/muscles.py).
- **Thermal zones kept fixed cooling units while their IT load was scaled.** The large halls were under-provisioned, so the native controller overheated. Cooling now scales with the hall.
- **The flight axis's motor thermal resistance was ten times too high,** so the native controller was always derated. Corrected. The UPS feed was sized under its own load; sized at 1.2 times.
- **The slow plants** (vehicle, spacecraft, flight) completed too few tasks in six minutes to measure. They now decide every 5, 10 and 2 seconds over the same 360 decisions.
- **Added the fixed_calm arm** for setpoint muscles, after development showed that the band's calm end alone accounts for much of the setpoint results.

None of these changes was chosen by its effect on Omni's result. The development runs showed losses as well as gains for Omni before and after them.

Not claimed

- **No row is evidence about a real machine.** A row says what the governor's law does to that model through that knob.
- **The plants, native controllers and bands were written by the same project as the governor.** That is a real conflict, so every one of them is in two files, to be read and contested.
- **A muscle's name chooses its knob and its plant's size; it does not get its own physics.** The 16 muscles of a family share the family's plant model at different sizes, through different knobs.
- **The plant code is Python only.** The governor's C++ twin is unchanged; a C++ twin of the plants is open.

Round 2 (2026-10-01, after round 1's results; before any round-2 confirmation seed)

Round 1 (seeds 1000-1009) labelled all five organisms WORSE. Reading the result against the shipped controller (omni_controller/controller.py, omni_controller/muscles.py, omnicompass/nervous_system.py) showed that round 1's Omni layer was not the Omni that runs on Kubernetes. Round 1 is kept unchanged in results/realms/round1/, with a note saying why it is superseded. Round 2 changes only how Omni commands a knob, the native machine-pool scaler and the declared budgets. These changes were made after seeing round 1, so they are listed with the source line each one follows:

Round 1	Round 2, as the shipped controller does it
Capacity muscles replaced the HPA with the stack simulator's capacity law (release one unit after convergence, dwell and a 20% band)	Pods: the HPA target written as <code>min(rho*, the operator's target)</code> , never tighter, held one autoscaler window; the HPA scales (<code>controller.py</code> , HPA target patch). Machine pools: the node release gate, one machine per decision (<code>nervous_system.node_release_gate</code>)
HPA-type setpoints moved inside a band whose calm end was tighter than the operator's target	The same <code>min(rho*, operator)</code> rule: more headroom is always allowed, less never (<code>controller.py</code>)
No contraction authority and no SLO reflex	Every contraction needs the organ's authority (<code>calm >= threshold</code> , senses live) and three clean decisions; while service is breached the knob returns to native (<code>nervous_system.authority</code> , <code>muscles.py</code> SLO reflex)
Continuous knobs jumped to the governed value	Down by at most the calm share of the surplus per decision (<code>nervous_system: step = calm</code>)
Request-served pools had their power capped	Never: throttling request work saves no energy and adds wait (<code>muscles.py, _power_cap</code>); GPU and CPU-frequency pools use their envelope, never under draw x 1.3
Thermal setpoint by headroom	The live cooling law: warm while cool, cold as heat rises, under the envelope <code>18 + 9 x calm</code> mapped onto the band (<code>muscles.py, _cooling</code>)
Admission paused every admission muscle, request traffic included, and resumed only when fully calm	Batch pacing of pausable work only: one muscle suspended per decision at power stress <code>>= 0.95</code> or heat <code>>= 0.96</code> (or a nervous pause), one resumed per decision at power stress <code><= 0.8</code> and heat <code>< 0.90</code> (<code>muscles.py, _batch_pace</code> defaults); request traffic never paused
<code>queue_ratio</code> was the backlog in service-target units	pending starts per serving unit, or latency pressure (<code>p95 / target - 1</code> , or the failed share), capped at 2 (<code>controller.py</code>)
The governor's current cap followed the applied cap	1.0, as the live controller sets it
Site power budget 1.25 x a formula nominal that sat under the organism's real native draw (compute ran at 1.2 x it), so the governor read the site as always at its limit	1.25 x each plant's mean native draw on the calibration seed 999 (never a result seed); compute pools alone likewise

Round 1	Round 2, as the shipped controller does it
Organism heat = the hottest of up to 656 plants	The mean, as every other organism channel; local heat stays with each plant's own reflex
Native machine pools used the HPA rule	The Cluster Autoscaler's defaults: add while work waits, remove one machine after the rest has been under 50% for 10 minutes
Building cooling sized under its own peak (native overheated half the time)	Units of 35 kW: capacity 1.25 x the declared peak

Unchanged: the plants' physics, the catalog, every other parameter, the arms, the outcomes, the guardrails, the label rule and the invalidity rules.

- **Round 2 seeds:** 2000 to 2009. Development of round 2 used seeds 0 and 1 only, never reported.
- **Results:** results/realms/ (round 1 in results/realms/round1/). Both rounds are cited together.

Round 3 (2026-10-02, after round 2's results; before any round-3 confirmation seed)

Only the make-up of the realm organisms changes. The plants, the Omni layer, the outcomes, the guardrails, the label rule and the invalidity rules are round 2's, unchanged.

- **Round 2 cut the 656 into four realms with no overlap.** No realm organism carried the infrastructure every real stack runs on unless that infrastructure was the realm's own. The data-centre realm had no cooling or power, and the robotics and plant realms had no Kubernetes, machines or GPUs.
- **Round 3 gives every realm the shared spine** (tools/realms_catalog.py, SPINE): Kubernetes Workload Scaling, Placement & Scheduling, Container Resources, Node Fleet, Cloud VM & Capacity, NVIDIA GPU Hardware, Host CPU & Memory, Network Routing, Storage, Observability, Reliability & Security, Cooling & Chillers, PDU / UPS & Electrical Distribution.
 - Each realm's organism is its own families plus the spine: Compute 345 muscles, Physics 262, Energy 282, Distribution 337.
 - The whole-tower organism still holds each of the 656 once.
- **Quantum Computing Control moves to the compute realm** (it behaves as a compute job queue).
- **Round 3 seeds:** 3000 to 3009. Round 2 is kept in results/realms/round2/ with a note on why it is superseded.

Round 4: the stacked organism (2026-10-02, before any round-4 seed)

The four realm organisms of round 3, stacked on one 15 s clock (realms/harness.py, run_stack; tools/run_stack.py). Every muscle appears as often as it appears in the realms, duplicates included: $345 + 262 + 282 + 337 = 1,226$. A duplicate is still a muscle that has to converge. Each realm keeps its own internal coupling. The plants, the Omni layer, the outcomes, the guardrails and the label rule are round 3's.

- **Arms:** native (no governor); separate (one governor per realm); one (one governor over the whole stack, reading the mean of all 1,226 muscles and the stack's total power against its total budget).
- **Check, required for validity:** the stacked native run equals each realm's own native run, plant by plant, on every seed. Stacking must change nothing natively.
- **Comparisons, each labelled by the rule:**
 - one governor against native;

- separate governors against native;
- one governor against separate governors (does one Omni over everything beat four).
- **Seeds:** 4000 to 4009. Development used seed 0 only, never reported.

Round 5: the whole stacks with the real card inside (written 2026-10-02, before any run)

One harness (`tools/run_hil.py`, started by `scripts/gpu_rented_run.sh` after a valid card smoke): each of the six organisms (the four realms, the four stacked with every duplicate kept (1,226), the whole tower of 656) runs on one clock as in round 3, with the machine's real GPU wired in as one more muscle of its NVIDIA GPU family (a spine family, so the card is in every organism). The card serves the pinned request stream; its own `power.draw` is heat in the organism's thermal zones and load on its storage sites.

- **Arms:** native (the stacks' own controllers, the card's own firmware) and omni (one engine on everything: the compass law on every simulated muscle, `realms/compass_arm.py`, and on the card's two wires, `omni_controller/gpu_compass.py`). At 90% of each arm every knob and both wires are handed back; a knob not handed back, a card limit not back at its start, or a card governor exiting non-zero makes the run invalid (exit 2).
- **Seeds and repetitions:** 3 paired repetitions, seeds 6000-6002; arm order alternates by repetition and organism.
- **Clock:** 240 steps of 2 s of wall clock per arm (the card in real time).
- **Outcomes:** work per energy, work, energy and violations, Omni against native, for three parts kept apart: the simulated stacks (evidence S), the card (its own meter, evidence P), and both added (the card as one more plant, its joules added to the stacks'). Labels by the round 3 rule.

Round 6: the compute pools inside the band (2026-10-03, before any round-6 seed)

Round 3's rule (band first) was not held on the organisms: the compass law spent 0.20 to 0.29 points more time over the service line than native in every completed cell of the grid (`results/scale/GRID.md`). The cause was measured on the whole tower (seeds 7000-7003), muscle by muscle:

- **compute pools giving a machine back** (42 muscles): +1.77 points of their own time over the line;
- **the HPA target set looser than the operator's** (249 muscles): +0.22 points;
- the process and thermal setpoints, which save the most energy, added none.

Two changes, both in `realms/compass_arm.py`, everything else as round 3:

1. **The HPA target is held at the operator's own.** Loosening it past native costs the service time. This is the same rule as the card's speed floor and the live controller's cover.
2. **A machine goes back only when the machines left cover the recent peak at 0.6 of the plant's own release level** (`RELEASE_MARGIN`). A machine boots in minutes, so a burst that arrives after a release is served late until the machine is back.

How 0.6 was chosen (10 paired seeds per organism; the rule is the most work per energy whose whole interval of time over the line is at or under native): 1.0, 0.8, 0.75, 0.7, 0.65 and 0.6 on the tower; 0.65 and 0.6 on the other five.

- 0.65 failed on Physics, the upper end of its interval +0.001;
- 0.6 held on all six.

Checked on 20 paired seeds (7000-7019), before this round's official seeds:

Organism	Work per energy	Time over the line	Work done
Compute	+0.095%	-0.017 pp	unchanged within its interval (-0.001% to +0.000%)
Physics	+0.086%	-0.011 pp	unchanged within its interval (-0.001% to +0.000%)
Energy	+0.201%	-0.031 pp	unchanged within its interval (-0.001% to +0.000%)
Distribution	+0.090%	-0.015 pp	unchanged within its interval (-0.001% to +0.000%)
The four stacked	+0.153%	-0.022 pp	unchanged within its interval (-0.001% to +0.000%)
The whole tower	+0.194%	-0.010 pp	unchanged within its interval (-0.001% to +0.000%)

Every knob was handed back. The grid runs at 100× and 1,000× now on GitHub were started on round 3's law. They are recorded as round 3's result, and the grid is rerun on round 6's law.

Round 5, amended (2026-10-03, before any round-5 run on the corrected law)

The whole stacks with the real card inside run at four sizes, as the six-organism grid does: each organism as 1, 10, 100 and 1,000 copies governed together on one clock, with the one real card inside as one more muscle of its NVIDIA GPU family, native against Omni on top. Repetitions by size: 3, 3, 2 and 1 (seeds from 6000).

The step is 2 s of wall clock wherever the simulation keeps up. A size whose step takes longer gets a longer step: 1.5 times the measured time per muscle on that machine, times its muscles. The card's request stream runs for the same 240 steps, so the card and the stacks stay on one clock. The step of every size is in the receipt. At 1,000 copies the card is one muscle among hundreds of thousands, so its watts are a small share of the organism's; that is the arithmetic of one card in a large stack, and the card's own meter is reported apart from the stacks. Everything else is as round 5. The two GPU confirmations now run before this stage, so the most important results are in hand first.

Round 6 receipts and the one rule (2026-10-03)

At 1x and 10x (1,000 paired runs each) every organism is labelled superior within guardrails, with band first held and every knob handed back (results/scale/receipts/, results/scale/GRID.md). At 1,000 runs the intervals are narrow enough to show one cost: work done is lower by 0.0001% to 0.001%, wholly below zero in several organisms. Its cause was measured on the Physics organism, muscle by muscle. It is the thermal zones: a room held warmer keeps a little less margin, so a sudden heat spike puts it over its limit for a few extra seconds inside a step that already counts as over the line.

Limiting how far a room may drift (half, a quarter, none of the way toward its warm end; 40 paired seeds) does not remove it until the setpoint is held native, and then the Physics organism uses more energy than native. The cost is therefore disclosed and judged by the one rule now written for every muscle (DISCLOSURES.md, section 3): at most 2% in any measure, only where energy is saved. The largest work cost in the grid is 0.007% in a single run and 0.001% over 1,000 runs, far inside it.

The 1,000-copy size: memory and sharding (2026-10-03, before its re-run)

The first 1,000-copy run (six run 37090435347) stopped twice at the same point: every organism larger than physics and robotics ran out of memory about 3 minutes in (the four stacked at 1,000 copies is 1.2 million plants; at 33 KB a plant, and with the calibration organism built while the result organism was still held, it needed over 40 GB on a 16 GB runner). Three changes, none of which changes a number:

- the calibration run is made before the organism, so one organism is in memory at a time (realms/harness.py, Body);

- each plant is packed as it is made: its exogenous series as machine arrays, its random source dropped (every plant draws its series when it is made, never while it runs), plants with the same parameters sharing one table (realms/plants.py, pack, _shared);
- the six workflow sizes its shards by organism (per_shard, workers by organism), because one run of the four stacked at 1,000 copies takes 2 to 3 hours and a runner's job ends at 6.

Same outputs: four organisms, two seeds each, at 1x to 3x, byte-identical before and after. A plant now takes 6.7 KB; the four stacked at 1,000 copies needs about 9 GB. The re-run: 10 runs per organism at 1,000 copies (the 1-run and 10-run cells); 100 runs at 1,000 copies follows when the runners allow it (about 650 runner-hours).

The 100-run cells at 1,000 copies run on one rented machine (scripts/grid_one_machine.sh): every organism in turn, as many processes as the machine's cores and memory allow, the same seeds (7000 on), each finished organism saved at once, one receipt at the end (SIX-1000x.md), saved as results/scale/v1/receipts/round6-1000x.md in place of the 1-run and 10-run receipt it contains.

Evaluation and simulation use only. Copyright (c) 2026 The Omni-Compass LLC. Commercial use, commercialization or monetization of any part of Omni-Compass requires a signed, paid Omni-Compass Enterprise License from The Omni-Compass LLC. Patent applications, copyright registrations and trademark applications filed in the United States. See LICENSE and NOTICE at the root of this repository.

Amendment (2026-10-05): the compass law is the arm, written before its run

The tables of round 3 compared native against the allocation law ("omni"), the arm that came before the compass law. The engine is now one law on everything, the compass law on top of native, so the realms runner compares native against compass (realms/harness.py ARMS: native, watch, compass). The watch arm stays: Omni watching must write nothing and leave every muscle exactly as native. Seeds, plants, knobs, guardrails and the labelling rule are unchanged. The tables published until now describe the allocation law and stay as first measured in docs/history/.

Amendment (2026-10-05): four levers bounded by their physics, and the mirror label (written after the compass table was seen)

Disclosed: written after the first compass-law table (nine muscles WORSE) was seen. Each change follows from the muscle's own model (docs/MECHANISM_OF_ACTION.md 9.6), applies to every muscle of its kind, and is tested:

1. Cooling power is native. A cap moves heat removal later, never away, and stages more units.
2. A battery's reserve is spent only at the connection's wall, and only when the battery can cover the excess. Below the operator's reserve every kWh costs $(1 - \eta)/\sqrt{\eta}$ to buy back.
3. A backup reserve (the UPS preset, backup=True) is never a lever.
4. A compute pool's power cap moves only where (9.8) says a lower cap saves energy. A quantum computer, a network switch and storage are left native (race to idle).

The labelling rule gains the mirror of ENERGY IMPROVEMENT WITH SERVICE TRADEOFF: less work per energy, no work lost (at most 1%), and the time over the line proven lower is SERVICE IMPROVEMENT WITH ENERGY TRADEOFF. Everything else is as before, and every other label keeps its definition.

Also in this amendment: an effort cap on a motion axis moves only where the copper loss it saves at full acceleration, $(J a_{\max} / kt)^2 R$, exceeds the standing draw a longer move pays, $p_{\text{idle}} + b v_{\max}^2$ (a reaction wheel; not traction, a flight axis or a robot joint). The machine-release margin on node pools is 0.3, the largest at

which no pool is later than native (results/realms/RELEASE_MARGIN_SWEEP.md, the per-muscle sweep).

Omni v2: the catalog grows to 945 muscles in 59 families (2026-10-06, before any v2 seed)

Written before any counted v2 run. **The law, the controllers, the plants, the Omni layer, the outcomes, the guardrails, the label rule and the invalidity rules are v1's, byte for byte.** Only the catalog and the presets it needs change, so by the change rule of docs/OMNI_V1.md this is a new version, Omni v2 (OMNI_V2.json, docs/OMNI_V2.md), and every modelled result is run again on it. No v1 result is read as a v2 result.

What is added, each row with its real system, its setting and its source, built by tools/realms_catalog_v2.py from three committed files by the v1 rules (realm, membership, plant, preset and knob: tools/realms_catalog.py, unchanged):

- **The 656 of v1, byte for byte** (realms/catalog_v1.csv): every v1 muscle keeps its id, realm, plant, preset and knob.
- **118 real controls the realm study had already found and the tower lacked** (docs/realm_study/TRUE_MUSCLES.csv, the rows marked "added (real, not in the 656)"; six of them were already in the tower under another name and are listed as skipped in realms/catalog_provenance.json). They join the family the study filed them under; 70 of them are spine muscles (Kubernetes, nodes, hosts, GPUs, storage, network, observability, power, cooling) and so sit in all four realms.
- **171 muscles in thirteen new families** (realms/wave4_families.csv), the domains the tower did not cover, one new preset each (realms/presets.py, round engineering figures for the class of machine, as every other preset):

Family	Realm	Plant / preset	Muscles	Examples of the real system
Healthcare Critical Environments	Energy / Facility / Industrial	thermal_zone / hospital	13	ASHRAE 170 rooms on a hospital BMS (Metasys, Desigo); isolation-room pressure; pharmacy cold rooms
Medical Imaging & Clinical Systems	Distribution / Specialized	compute_pool / clinical	12	DICOM archives and routers (Orthanc, dcm4chee), EHR tiers, HL7/FHIR engines, monitoring gateways
Agriculture & Irrigation	Energy / Facility / Industrial	process_loop / irrigation	14	pivot and pump controllers (Valley, Lindsay, Netafim), greenhouse climate computers (Priva), grain dryers
Oil & Gas Pipelines	Energy / Facility / Industrial	process_loop / pipeline	14	compressor and pump station SCADA (Emerson, AVEVA, Honeywell), API 1130 leak detection, DRA injection
Rail Traction & Train Control	Physics / Robotics / Autonomous	motion_axis / rail_traction	14	traction converters and ATO (Siemens Mobility, Alstom, Hitachi Rail), CBTC headway, ERTMS
Marine Propulsion & Vessel Automation	Physics / Robotics / Autonomous	motion_axis / marine_propulsion	12	propulsion and power management (Kongsberg, Wärtsilä, MAN), IMO shaft power limitation
Ports & Maritime Logistics	Distribution / Specialized	compute_pool / port	12	terminal operating systems (Navis N4), automated cranes and AGVs (Kalmar, Konecranes), shore power
Mining & Mineral Processing	Energy / Facility / Industrial	process_loop / mill	14	grinding and flotation control (Metso, FLSmidth, ABB), ventilation on demand (Howden), dewatering

Family	Realm	Plant / preset	Muscles	Examples of the real system
District Heating & Cooling	Energy / Facility / Industrial	process_loop / district_heat	13	network and substation control (Danfoss Leanheat, Kamstrup), CHP dispatch, thermal storage
Power Generation & Turbine Control	Energy / Facility / Industrial	process_loop / turbine	14	governors and plant DCS (GE Mark VIe, Emerson Ovation, Woodward), AVR, hydro gates, reactor rods
Renewable Generation & Inverter Control	Energy / Facility / Industrial	process_loop / inverter	13	IEEE 1547 inverter curves (SMA, Enphase), wake steering (NREL FLORIS), pitch and curtailment (Vestas)
Elevators & Vertical Transport	Physics / Robotics / Autonomous	motion_axis / elevator_hoist	12	drives and destination dispatch (Otis, KONE, Schindler), ISO 25745 standby modes
Pharmaceutical & Food Manufacturing	Energy / Facility / Industrial	process_loop / batch_reactor	14	ISA-88 batch control (DeltaV, SIMATIC Batch), bioreactors (Sartorius), pasteurizers, CIP, cleanrooms

Four wave-4 names collided with the tower and carry a domain prefix (`terminal_tank_level_target`, `train_auxiliary_power_budget`, `plant_reactive_power_target`, `inverter_frequency_droop_setpoint`). Eleven names the knob rules would have placed wrongly are set by hand in `tools/realms_catalog.py` `OVERRIDE` and listed there.

The organisms of v2: Compute / AI / Cloud 430, Physics / Robotics / Autonomous 376, Energy / Facility / Industrial 470, Distribution / Specialized 440, the four stacked 1,716, the whole tower 945; the shared spine is 257 muscles. In code and workflows the two whole-tower organisms are now named `tower` and `stack` (nothing is named by a count); the v1 names `organism_656` and `stack_1226` are read as aliases, so every v1 run folder and record still reads.

Checked before this freeze, on seed 7, every one of the 289 added muscles: deterministic, the watch arm equal to native, every knob handed back after the kill, no number undefined, Omni's time over the line never above native's. Two presets were adjusted on that check and before any counted run: the inverter plant's disturbance amplitude (0.8 to 0.5, with its noise 0.05 to 0.02) because native itself was outside its band a quarter of the time, and the elevator hoist's position gains (kp 2,000 to 3,000, kd 400 to 500) with a smaller load disturbance (400 to 100 N m), because the first figures put native's tracking error at its limit under passenger load. Where native is still often outside its band in some heavy-load muscles (the compute pools at their busiest sizes), that is the plant at the size the muscle id draws, the same for both arms, and it is reported as it comes.

v2 runs: the realms table (every muscle alone and the five organisms, seeds 3000-3009, three GitHub runs A, B, C), the six-organism grid at 1, 10, 100 and 1,000 copies, the six organisms with the real cluster inside, and the big organisms on Azure, each on the v2 fingerprint; the live benchmarks that import none of the realm files (Kubernetes alone, Azure alone, CityLearn, the grid, the robots, the database) are run again on v2 as GitHub's queue allows, their v1 tables kept as v1 tables.

Omni v3: the slack gate on speed knobs (2026-10-06, after the v2 table, before any v3 seed)

Declared on the v2 result and before any v3 run. The v2 table (results/realms/REALMS.md, three runs reproduced) read Physics and the whole tower as "energy improvement with a service tradeoff" where v1 had every organism superior. The rows that carried it were the new speed muscles: three rail traction muscles where Omni's slowing added 1.6 to 2.7 points of lateness on a train already near its timetable's capacity, and two marine and two elevator muscles where the slower cycle finished fewer moves in the hour. By the honesty rule the first suspect was our own wiring, and it was: the realm harness had no do-no-harm gate on the speed knob of a motion axis, so the compass slowed an axis that had no slack to spend. The robot benchmark has such a gate (its paired physics trial, docs/ROBOTICS_PREREGISTRATION.md); the realms get one now.

The rule (realms/compass_arm.py, speed_slack). A motion axis is offered its speed knob only where its duty at full speed, $\text{task rate} \times (\text{move time at full speed} + \text{dwell})$, is at most $\text{SLACK} = 0.5$, computed from the plant's own figures before any decision. The compass may ease an axis down to 0.4 of its speed, which stretches a move by up to 2.5 times; an axis busy more than half the time at full speed cannot absorb that when tasks arrive in bunches, and its knob stays native. Where the knob is offered, nothing else changes: the compass, its gains, the cover [0.4, 1.0] and the fail-up are v2's.

The marine preset. Four 600 rad speed changes an hour made a one-hour run too coarse to count (one unfinished move is a quarter of the hour's work); the v3 preset is twelve changes of 200 rad an hour, a move of 126 s plus 30 s of dwell, deadline 1,200 s. No other preset changes.

What the gate does on the shipped presets (the middle size; a muscle's own size, 0.6 to 1.4 on the task rate, moves its duty either side):

Preset	Move at full speed	Dwell	A task every	Duty	The speed knob
flight axis	4.8 s	2.0 s	20 s	0.34	offered
reaction wheel	30.0 s	5.0 s	100 s	0.35	offered
marine propulsion	126.5 s	30 s	300 s	0.52	left native
elevator hoist	12.8 s	8.0 s	40 s	0.52	left native
robot joint	0.7 s	0.2 s	1.7 s	0.54	left native at the middle size, offered on the lightly loaded joints
EV traction	35.6 s	10 s	60 s	0.76	left native
rail traction	220 s	20 s	300 s	0.80	left native

The threshold is a declared number set on the v2 finding, not a fitted one: it is written here before the v3 runs and applies to every axis alike. Where it leaves a knob native the row will read as native does, and that is the result. Everything else, the plants, the outcomes, the guardrails, the label rule and the invalidity rules, is v2's. v3 runs the realms table three times (A, B, C), then the grid, then the organisms with the real cluster inside; the v2 table stays a v2 result.

CHAPTER 46

The Compass Law on Real Kubernetes: Preregistration

Written and committed before the run. The run's commit is the one that carries this file; nothing in the law, the harness or this rule changes after it starts.

What is run

Workflow benchmark-reps on main, 10 repetitions. Each repetition runs three arms back to back on the same runner, each on a fresh six-worker kind cluster, in an order rotated by repetition (scripts/kind_paired.sh):

Arm	What governs
native	Kubernetes alone (HPA at target 50, scheduler); Omni-Compass not started
omni	Omni-Compass on top with the engine's allocation law (- -law governor, the law of sets 22 to 25)
compass	Omni-Compass on top with the compass law (- -law compass, omni_controller/controller.py)

Load: fixed rate (loadgen=open), the same work in every arm. 900 measured seconds per arm. SLO 500 ms at the 95th percentile. Every Omni arm runs the six-state engine on every decision, the nervous system's authority and release gate, the shield, the compass, and ends with the reset, which must return the HPA target, its replica range, the pods' CPU limits and every worker to native, with no record left (scripts/kind_bench.sh).

The compass law in the live controller

The service position is the 95th-percentile response time over the SLO (0 calm, 1 the line); a blind probe or a pod waiting for a place reads as past the wall. The force is $A \tanh((K_P (p - 0.5) + K_D v) / A)$ with K_D for critical damping times the realm push factor 3, the same law and gains as on every realm muscle (realms/compass_arm.py): up gain 0.10, down gain 0.02, release threshold -0.2. Two levers:

1. **HPA target**, cover from 60% of the operator's target to the operator's own: the up force lowers it (more pods), the down force returns it toward the operator's. It is never tighter than native.
2. **Node pool**: past the 0.95 wall one machine more at once; one machine back only while the force is below -0.2, the position is below the center, and the nervous system's release gate is open.

Outcomes and the rule

Primary: **worker nodes in service** (mean) and **95th-percentile response time**, each arm against native, paired over the 10 repetitions with a t-based 95% interval (tools/live_reps.py).

Band first: the compass arm is a win only if its p95 is not worse than native's (the upper end of the 95% interval of the paired difference at or under 0) **and** failed requests are not higher. If that holds and machines in service fall with

an interval wholly below 0, the label is **better on machines within the band**. If machines fall but the band condition fails, the label is **tradeoff**. Otherwise **not established**.

Secondary, reported, not used for the label: p99, mean response time, HPA replicas, pods started, pod start wait, CPU including Omni-Compass's own, the declared energy models. The omni arm is reported against native and against the compass arm by the same rule. A run that fails its own checks (reset, controller stopped early, missing permission) is marked invalid and left out, never silently counted.

Evidence class **L**: real Kubernetes software on kind. Energy on kind is a declared model, not a meter.

Set 26 result

Compass arm against native: machines in service -17.2% (-26.4% to -7.9% of native), p95 -64.8%, failed requests 0 on both: **better on machines within the band** (results/live/LIVE_REPS_26.md). The allocation law in the same set: machines -35.8%, p95 -55.4%.

Set 27 (written before the run)

The compass in the live controller now reads the service as the GPU compass does (omni_controller/gpu_compass.py, GPU amendments 6 and 7): the mean response time of the latency window between the bare service time (a tenth of the SLO) and the SLO, held at the compass's center 0.4 (the GPU service profile); p95 at or past the SLO, a blind probe or a pod waiting for a place is past the wall. Everything else, the arms (native, omni, compass), the load, the duration, the outcomes and the labelling rule above, is unchanged. The run's commit is the one that carries this section.

Set 27 result

Compass arm, aligned with the GPU governor, against native: machines in service -15.9% (-1.462 to -0.450 machines), p95 -65.5% (-311.8 to -156.1 ms), p99 -72.6%, failed requests 0 on both: **better on machines within the band** (results/live/LIVE_REPS_27.md, run 37071353971, commit d46c959). The allocation law in the same set: machines -36.6%, p95 -53.1%.

The verdict in the live controller (2026-10-03, before any further set)

The compass law gives a machine back only where it measures that the service is no worse for it (omnicompass/verdict.py, stepwise, in omni_controller/controller.py). While the service is calm (inside the compass, no pod waiting, no breach), one more machine is given back on trial. The response times of 200 requests served without it are set against 200 served just before and against the cluster as it first ran on its own:

- at most 2% slower than both: the machine stays given back;
- slower than that: it is taken back and not tried again for 120 decisions.

Where no machine passes, the pool stays as the cluster runs it alone. The HPA target's cover is unchanged: from 60% of the operator's target up to the operator's own, never looser than native. Every trial is in the audit. The next set runs with this verdict; sets 26 and 27 ran before it and stay as they ran.

Set 28 and set 29 (2026-10-03, set 29 written before its run)

Set 28 (run 37087620193, commit 24666d7) ran the compass law with the verdict. With Omni-Compass on top against native:

- machines in service -4.7%;
- p95 -65.3%, p99 -68.6%;
- pods waiting 0;
- failed requests 0.

Total CPU including Omni-Compass's own came out +2.2% (+0.003 to +0.042 cores), more than the 2% the one rule allows (DISCLOSURES.md, section 3). The cause is the controller's own cost: 0.063 cores, mostly a new kubectl process for every read, about 20 a minute. The compass law with the verdict freed only 0.040 cores of work. The allocation law in the same set: machines -29.6%, p95 -58.0%, total CPU including its own -1.2% (not significant). Set 28's receipt is results/live/LIVE_REPS_28.md, published with this section.

Set 29 runs the same arms, load, duration, outcomes and rule as set 28. One thing changes: the controller reads through one kubectl proxy started once, under the same least-privilege identity, so a read is a local HTTP request instead of a new kubectl process (omni_controller/controller.py, Kube; writes are unchanged; tests/test_api_proxy.py). The label also requires total CPU including Omni-Compass's own to be no more than 2% above native.

Set 29 result (run 37094338955, commit a3721cc): the compass law with the verdict, total CPU including its own -4.7% (-0.071 to -0.017 cores), passes; Omni-Compass's own CPU 0.011 cores (set 28: 0.063). Machines -5.1%, p95 -64.4%, p99 -71.1%, HPA replicas -7.6%, failed requests 0. No measure significantly worse than native in either arm. Receipt: results/live/LIVE_REPS_29.md.

The cost to match (written before its run)

The question a buyer asks: what would native Kubernetes have to spend to answer as fast as it does with Omni-Compass on top? Each repetition runs, on the same runner and the same work, in rotated order:

- native (the operator's HPA target 50);
- native tuned harder by its operator, HPA target 40, 30 and 20 (more pods, faster answers), no Omni-Compass (ARM=native40, native30, native20 in scripts/kind_bench.sh);
- native with Omni-Compass on top, the allocation law (omni);
- native with Omni-Compass on top, the compass law with the verdict (compass).

The report (tools/live_reps.py, "The cost to match") lists every arm's p95, p99, HPA replicas, CPU including Omni-Compass's own, and machines in service. For each Omni-Compass arm it names the cheapest native setting (by CPU) whose p95 is at or under Omni-Compass's, and that setting's extra replicas, CPU and machines over Omni-Compass. If no native setting tried reaches it, the report says so and gives the lowest native p95. 10 repetitions, 900 measured seconds per arm, fixed-rate load.

The fault test (written before its run)

Health, security and the babysitting a cluster needs, measured. Every arm (native; native with Omni-Compass on top, the allocation law; native with Omni-Compass on top, the compass law with the verdict) meets the same four faults at the same moments (scripts/kind_faults.sh, FAULTS=1):

1. at 15% of the run, a worker machine dies (its kind container is stopped) and comes back two minutes later;
2. at 35%, traffic triples for two minutes;
3. at 55%, a pod with no CPU limit burns CPU for two minutes;
4. at 75%, the response-time probe goes blind for one minute.

For each fault the report (`tools/live_reps.py`, "The fault test") gives the time to recover (from the fault's start until responses stay under the line for 30 s straight, at most 300 s) and the share of samples over the line or failed in the 300 s after it, paired against native over 10 repetitions. All the usual gauges are reported as well, now including the share of response samples over the line (`pilot/bench_report.py`). Lower is better in each.

The fault test, first run (run 37094604580, commit a149d4e; `results/live/FAULTS.md`). Omni-Compass on top recovered faster than native from every fault (allocation law: machine down -29%, runaway pod -35%, spike -7%; compass law: -14%, -17%, -6%), p95 -53% and -55%, p99 -27% (not significant) and -57%. One measure was worse: with the compass law, **HPA replicas +8.2%** (+0.41 to +0.99, significant), with CPU and machines unchanged and no energy saved, so outside the one rule (`DISCLOSURES.md`, section 3). The cause: past the wall the compass lowers the HPA target at once (more pods, the faster recovery), then handed the operator's target back step by step and held each step for the autoscaler's window, so the extra pods outlived the fault.

The change, written before the re-run (`omni_controller/controller.py`, the compass's push and pull on the HPA target; `tests/test_compass_controller.py`, "fault over"): once the responses are back inside the band (the compass's position under its center), the response line is clean and no pod is waiting, the operator's own target returns at once and is not held by the window. More pods only while the fault lasts. The re-run is the same fault test, arms, load, duration and rule (set 30 F); set 30 runs the same code without faults, to show nothing else moved.

Set 30 and set 30 F (runs 37105047258 and 37105046042, commit acc1c4e; `results/live/LIVE_REPS_30.md`, `results/live/FAULTS_30.md`). Set 30, no faults: nothing significantly worse in either arm; the compass law's machines -9.9%, p95 -64.9%, time over the line -98.7%, HPA replicas -32.2%, total CPU -6.5%. Set 30 F: recovery faster than native from every fault in both arms; the compass law's HPA replicas **+5.6%** (first run +8.2%), still significant, with CPU and machines unchanged. The allocation law, which recovers as fast or faster, held no extra pods (+2.0%, not significant).

The second change, written before set 31 F (`omni_controller/controller.py`; `tests/test_compass_controller.py`, "blind"). Past the wall, the compass lowers the HPA target (more pods) only when the cause is load: the response line breached with every sense live and no pod waiting. Past the wall from a blind sense, or from pods waiting for a machine that is gone, more pods answer neither, so the target is the operator's own: fail up is native's own setting, as on the card, where fail up is the card's own clock and limit. The machine reflex (one machine more past the wall) is unchanged. Set 31 F is the same fault test, arms, load, duration and rule; set 31 the same without faults.

Set 31 and set 31 F (runs 37110007121 and 37110005322, commit 0a38e76; `results/live/LIVE_REPS_31.md`, `results/live/FAULTS_31.md`). Set 31 F: the compass law's HPA replicas under faults **-1.1%** (not significant): the extra pods are gone. Nothing significantly worse in either arm; recovery faster than native from every fault (compass law, machine down -49%). Set 31: nothing significantly worse; the compass law's machines -10.4%, p95 -66.0%, p99 -73.0%, time over the line -99.4%, HPA replicas -44.5%, pods started 0 against native's 4.3, total CPU -6.1%.

The bill on a real cloud (written before its run)

The question a buyer pays for: the same work, a smaller bill? On kind every machine stays powered, so a machine given back saves only a declared model's energy. On a real cloud the machine is deleted and stops being billed. The run (`.github/workflows/aks-metered.yml`, `scripts/aks_paired.sh`, `scripts/kind_bench.sh` with `PLATFORM=aks`):

- **The cluster.** Azure Kubernetes Service, a fresh cluster for every arm, built the same way:

- a system pool of one machine, tainted so no workload lands on it (AKS's add-ons and the load generator: kind's control plane);
- a work pool starting at 4 machines (Standard_D2s_v5) under **Azure's own cluster autoscaler** (min 1, max 4; scale down after 2 minutes unneeded), which deletes a machine once it is empty.
- **The arms**, rotated in each repetition:
 - native: Kubernetes with Azure's autoscaler alone;
 - native with Omni-Compass on top, the compass law with the verdict (compass);
 - native with Omni-Compass on top, the allocation law (omni).

With Omni-Compass on top, the machines it gives back are idled (new pods go elsewhere, their pods leave first), and Azure's autoscaler then deletes them. Omni-Compass never deletes a machine itself.

- **The same work** in every arm: the fixed-rate load of sets 22 onward, 900 measured seconds after 120 s of warm-up.
- **The bill.** Every 15 s, the number of work machines that exist (in service or idle, every one is billed), integrated over the measured window: billed machine-hours, priced at Azure's list price for the machine (USD 0.096 an hour for Standard_D2s_v5, Linux, pay as you go, set in the workflow's input). Omni-Compass never reads this count.
- **Outcomes.** Billed machine-hours and the bill (lower is better), and every gauge of sets 28 and 29 (response time, failures, pods waiting, CPU including Omni-Compass's own), paired against native with 95% intervals over 5 repetitions. Labelled by the one rule (DISCLOSURES.md, section 3): nothing more than 2% worse, and only where the bill or energy is saved.
- **Housekeeping.** One repetition at a time; each cluster deleted when its arm ends, before the next is made; the resource group deleted at the end of every repetition whatever happens. It needs the repository secret AZURE_CREDENTIALS.

The capacity test (written before its run)

The question behind "more work for the same cost": on the same machines, how much more work does Kubernetes serve with Omni-Compass on top before its answers break the line? Each repetition runs native, native with Omni-Compass on top (the allocation law) and with the compass law and the verdict, in rotated order, on the same six workers, the open-loop load rising in eight equal steps of one load generator each (6 requests a second per generator, 1 to 8 generators, 200 s a step, 1,600 measured seconds; `load_steps` in `benchmark-reps`). A run's capacity is the highest step at which no more than 5% of the response samples (every 5 s, through the Service) are over the line (500 ms) or failed, every lower step holding too, the first 30 s of each step left to settle (`tools/live_reps.py`, `capacity`). Reported: each arm's mean capacity in requests a second, its change against native and the 95% interval of the paired difference over 10 repetitions; every usual gauge beside it. Labelled by the one rule: nothing more than 2% worse.

The fairness test (written before its run)

The question a buyer with many tenants asks: when one application surges, does Omni-Compass on top protect its neighbour, or starve it? Each repetition runs native, native with Omni-Compass on top (the allocation law) and with the compass law and the verdict, in rotated order, on the same six workers, with two applications on them (`TWO_APP=1`, `deploy/kind/noisy.yaml`): php-apache under its usual fixed-rate load, and a noisy neighbour, the same image and the same HPA rule, whose own fixed-rate load surges in steps (0, 0, 6, 0, 6, 0 load generators of 6 requests a second, the same moments in every arm). Each application's response time is probed through its own Service. Omni-Compass governs both HPAs (`deploy/kind/rbac-omni-noisy.yaml`: the second HPA and nothing else more). Reported: every usual gauge for php-apache, and the neighbour's p95, p99, time over the line and failed requests, each paired against native over 10 repetitions with its 95% interval. The label is the one rule, applied to both applications: nothing more than 2% worse for either.

The capacity and fairness results, and the amendment they call for (2026-10-03 evening, before either is run again)

Capacity (`results/live/CAPACITY.md`, run 37150909816): the compass law served 33.0 requests a second within the line against native's 24.6, **+34.1% (95% interval of the paired difference +6.2 to +10.6 requests a second)**, with response times about half of native's and fewer failures. One measure significantly worse: pods started 7.3 against 5.6 (+30.4%) while the mean HPA replicas were 26.9% lower: churn, not more pods. **Fairness** (`results/live/FAIRNESS.md`, run 37154210570): the neighbour unharmed under both laws; with the compass law php-apache's failed requests 3.89% to 4.90% (+0.12 to +1.90 points), pending pods and pod start wait worse, no machine saved. By the one rule neither result is labelled better. Both causes are read from the controller's code; both corrections are stated here, as law, before the tests run again.

1. Only a sensed muscle moves (the fairness cause). Let the probe measure the response time of the services in a set $\mathcal{S}$ (`--sensed ns/deployment, ...`; empty means every HPA, the single-service case). For each HPA h scaling a deployment $d(h)$ with the operator's target x^{op}_h , the target written is

$$x_h(t) = \begin{cases} \text{the compass's (or the allocation law's) target} & d(h) \in \mathcal{S} \\ x^{\text{op}}_h & d(h) \notin \mathcal{S} \end{cases}$$

The probe's position p is a statement about $\mathcal{S}$ alone. A push on a muscle outside $\mathcal{S}$ answers no sensed error and adds pods that compete with $\mathcal{S}$ for the same machines (the run's pending pods and failures). It is the GPU law's own rule, applied to Kubernetes: where Omni-Compass cannot sense, it stands at native's own setting.

2. Pods owed to a growing demand stay (the churn cause). Let $u(t)$ be the CPU the workers use and W the HPA's scale-down window ($W = 300$ s unless the operator set one). The demand is *still growing* when

$$G(t) = \Big[u(t) > (1 + \epsilon) \min_{t-W \leq s \leq t} u(s) \Big], \quad \epsilon = 0.05$$

After a breach the compass lowers the target to the bottom of its cover, $x = x_{\text{lo}} = 0.6 x^{\text{op}}$ (more pods, at once). Before this amendment, the first decision with the position back under the centre, $p < c = 0.4$, no breach and no pod waiting, returned $x = x^{\text{op}}$ at once. Under a rising load the autoscaler then removed the extra pods one window later and started them again at the next step: the run's churn. The return is now

$$x \leftarrow x^{\text{op}} \quad \text{only if} \quad p < c, \quad \text{no breach, no pod waiting, } -G(t)$$

and while $G(t)$ holds the target stays where the breach put it. A fault or a spike that has passed has a flat or falling u , so G is false and the operator's target returns at once, as the fault test requires (set 31 F: HPA replicas under

faults -1.1%); only a demand still climbing keeps its pods. $\epsilon = 0.05$ is set above the decision-to-decision noise of the node CPU reading. A rise too small to clear it leaves the rule as it was before this amendment (the target returns at once), so the correction can remove churn but cannot add any.

Nothing else changes: the compass's band, gains and centre, the verdict, the fail-up rules, the release gate and the node-pool law are as registered. Tests: `tests/test_compass_controller.py`, cases *demand* and *sensed*, beside the existing *fault over*, *blind* and reset cases.

The re-runs, on the commit that carries this amendment, unchanged in design: the capacity test (load_steps 1 to 8, 1,600 s), the fairness test (two_app 1, 900 s) and the fault test (faults 1, 900 s), 10 paired repetitions each, arms native / omni / compass. The second application is now probed and reported as before, and its HPA is held at the operator's target. Labelled by the one rule. Every number, whatever it says, is published beside the runs above.

The re-runs under the amendment: results (2026-10-04)

All three on commit 199f350, 10 paired repetitions each, the design unchanged.

- **Fault test, set 32 F** (results/live/FAULTS_32.md, run 37162459956): **no measure significantly worse under either law.** The compass law recovers from a lost machine 58% faster and starts 34% fewer pods; the amendment left the fault behaviour of set 31 F intact.
- **Fairness** (results/live/FAIRNESS_2.md, run 37162458834): **correction 1 holds.** With the compass law php-apache's failed requests are no longer worse (-14.0% , not significant; before $+26.0\%$), pending pods -10.5% (before $+65.2\%$); the neighbour unharmed under both laws. The compass law's response-time gains in this test are no longer significant. Still worse under both laws: the mean pod start wait ($+1.1$ s compass, $+1.7$ s allocation law).
- **Capacity** (results/live/CAPACITY_2.md, run 37162457542): the compass law served **24.6 against native's 16.2 requests a second, $+51.9\%$ ($+6.2$ to $+10.6$)**, the same paired difference as the first run on slower runners; response times -28% to -56% , failures -9.6% , HPA replicas -10.7% . **Correction 2 did not remove the extra pod starts** (5.1 against 3.8, $+34.2\%$; before $+30.4\%$): its premise, that the extra starts were pods removed and started again between steps, is not borne out. With the mean replicas lower, the reading the data support is pods started earlier on a rising load, the mechanism of the added capacity. Correction 2 stays (it removes no gain and added no measure worse); the pod-start rows stand as measured, and by the one rule neither the capacity nor the fairness arm is labelled better while they do.

The pod record, and the second amendment (2026-10-04, before the next runs)

What the pod record shows (`tools/pod_report.py`, workflow pod-report, read from the stored artifacts of runs 37162457542 and 37162458834; nothing re-run). In the capacity test native started its pods once, early (33 of 38 in the first fifth of the window), and removed none. With Omni-Compass on top, two minutes into the window the controller raised the HPA target above the operator's (the conveyance: each pod given a larger CPU limit, the target raised by the same factor so each pod stays as busy, $x = g \cdot x^{\text{op}}$; 114% with the compass law, 190% with the allocation law). The autoscaler then removed pods (14 with the compass law, 35 with the allocation law, over ten repetitions), and the next load steps started them again. **The extra starts are exactly those removals.** Correction 2 above read the wrong signal: the CPU used by the whole node, where one load step is lost in the node's own load.

3. A raise of the target waits for a steady demand. Let the demand on an HPA's service be read from the autoscaler's own status, $D(t) = \bar{u}(t)$, $r(t)$: the pods' mean CPU utilisation of their request times the pods running, the CPU used in units of one pod's request. With W the HPA's scale-down window and $\epsilon = 0.05$,

$$S(t) = \Big[t - t_0 \geq 0.9W \wedge \max_{[t-W, t]} D \leq (1+\epsilon) \min_{[t-W, t]} D \Big]$$

(t_0 the oldest sample in the window). A new target above the one standing (fewer, larger pods) is written only when $S(t)$ holds; the autoscaler itself removes pods only after its scale-down window, and Omni-Compass now asks the same of its own consolidation. A lower target (more pods) is never held, and the return of the operator's own target after a fault (rule 2) is unchanged. Correction 2's growth test $G(t)$ now reads the same $D(t)$ of the HPA concerned, not the node's CPU. Tests: `tests/test_compass_controller.py`, case *steady*.

The capacity, fairness and fault tests are run again on the commit that carries this amendment, unchanged in design, and published beside the runs above whatever they show. The bill run on Azure (aks-metered, run 37171672509) started on commit 199f350, before this amendment, and is reported as of that commit.

The third amendment, and a change considered and declined (2026-10-04, before the next runs)

4. A lower target is never held. Until now every new HPA target, in either direction, was held for one autoscaler window W while the line was clean, so on a step up the compass's push (a lower target, more pods) could wait up to W and the pods arrive after the line is missed. The hold exists because a raise inside the window removes pods the autoscaler then starts again; a lower target asks for pods and removes none. The hold now applies to raises only:

$$\text{write } x_h(t) \text{ iff } x_h(t) < x_h^{\text{now}} \vee [t-W, t) \vee \text{the operator's own target returns},$$

with rule 3 still requiring a steady demand before any raise. Tests: `tests/test_convey.py` (a lower target written at once inside the window).

Declined: raising the replica cap above the operator's. `deploy/kind/demo.yaml` sets `maxReplicas: 10`, and at the top of the capacity test both arms meet it. Lifting it would let Omni-Compass run more pods than the operator authorised, which is more resources, not the same resources used better, and it would break the shield's standing rule that Omni-Compass never moves past an operator's bound. Native would need the same cap for the comparison to stay fair. It stays at the operator's value in every arm; a higher cap is the operator's decision, tested as its own setting if ever asked for. Already in place and unchanged: a pod waiting for a place reads as past the wall and asks for one machine more at once.

A narrower form was proposed the same day (`docs/history/AMENDMENT_RISING_STEP_CAP.md`): while the demand rises and a pod of the sensed service is pending, raise `maxReplicas` by the pending count and lower the target to match. It is declined for a mechanical reason as well as the one above. The replica cap does not make pods pending: at the cap the autoscaler simply asks for no more pods, so the pending count there is zero and the write would never fire where it is aimed. A pod is pending only when the scheduler finds no machine with room for it, and more places under the cap give such a pod nowhere more to go; the answer to that is a machine, which the compass already asks for. The part of the proposal that holds, asking for pods at once on a step up instead of after a window, is rule 4 above.

On the CPU fill (used over allocatable about 0.10 in every arm): on kind every worker reports all of the host's cores as its own, so the allocatable counts the same cores once per worker and the fill reads far lower than the machine doing the work. Every receipt from the next runs on carries the host's own busy share and core count (`host_cpu.csv`, `tools/live_reps.py`), so the room left on the real machine is measured, not inferred.

The capacity, fairness and fault tests are run again on the commit that carries this amendment, beside the runs of commit 5d2e238 (rule 3 alone), so each rule's effect stays separable.

The second amendment's runs: results (2026-10-04)

results/live/AMENDMENT_2_RUNS.md, commit 5d2e238 (rule 3 alone), 10 paired repetitions each.

- **Capacity: no measure significantly worse under either law.** Compass law 24.6 against native's 17.4 requests a second, +41.4% (+5.4 to +9.0); allocation law 25.2, +44.8% (+4.9 to +10.7). The compass law's pods started 5.7 against 4.1, interval -0.64 to $+3.84$: no longer significant. Rule 3 removed the significant pod-start excess and lifted the allocation law's capacity from $+0.0\%$ to $+44.8\%$.
- **Faults: no measure significantly worse under either law.**
- **Fairness: the compass law, no measure significantly worse for either application.** The allocation law: php-apache much better, and the neighbour's failed requests **+0.74 points (+0.19 to +1.28), significant**. The allocation law conveys idle CPU to the service it senses; on machines shared with a surging neighbour that CPU is the neighbour's headroom. The allocation law is not labelled better in this test.

Standing after rule 3: the compass law is clean in all three tests (no measure significantly worse, capacity +41.4%). Rule 4 (a lower target never held) runs next on commit 583c97f, beside these, to see whether it adds capacity without costing a row.

The fourth amendment: the replica cap as a lever the operator grants (2026-10-04, before its run)

The founder's instruction: no number in the harness is hard-wired; every one moves where the system lets it. The cap of 10 in `deploy/kind/demo.yaml` is the value of Kubernetes' own php-apache example, not a choice of any operator, and an earlier session read it as an operator's bound. It becomes a lever, under the operator's grant.

5. Replica room. With `--replica-ceiling  $N_{\max}$` granted (0, the default, leaves the cap untouched), for a sensed HPA under the compass law, with r the pods running, c the cap, $\bar{u}$ the pods' mean utilisation and x the target, the autoscaler's own arithmetic asks for $n = \lceil r \bar{u} / x \rceil$ replicas, and

$$c \leftarrow \min(N_{\max}, n) \quad \text{if } r \geq c, n > c, p \geq \text{centre or the line is breached};$$

$$\leftarrow c^{\text{op}} \quad \text{if } n \leq c^{\text{op}}, p < \text{centre}, \text{no breach}, S(t).$$

The cap is raised only while it binds and the line is threatened, to what the autoscaler asks and never past the grant, and returns to the operator's once the demand has held still for a window. The operator's range is recorded before the first change; the reset restores it. Tests: `tests/test_compass_controller.py`, case *room*.

Its run (a setting of its own, labelled as such): the capacity test unchanged, `replica_ceiling 30`, beside the runs without it. Native keeps its cap of 10, as an operator who has not raised it would; the extra pods the compass uses are counted in the HPA replicas row, so the receipt shows what the capacity cost in pods, not the gain alone.

The third and fourth amendments, and the bill on a real cloud: results (2026-10-04)

Rule 4 (results/live/AMENDMENT_3_RUNS.md, commit 583c97f): **the compass law has no measure more than 2% worse in any of the three tests.** Capacity +48.1% (+5.7 to +9.9), nothing worse; fairness, neither application worse (energy per core-hour on the declared standby model +1.5%, inside the allowance); faults, nothing worse, lost-machine recovery -61%. The allocation law: capacity +55.6%, nothing worse; in fairness the neighbour no longer worse. The first runs with the real machine metered: a 4-core GitHub runner 48% to 70% busy, where kind's "used / allocatable" reads 0.07 to 0.10.

Rule 5, the replica lever (results/live/REPLICA_ROOM.md, commit 2101c3d, ceiling 30): the compass law served 24.6 requests a second, the same as without the lever, with 78% more pods and 26 pod starts against native's 4.6, both significant. **The replica cap was not what limited the service; the machine doing the work was.** The lever stays, off by default, as the operator's to grant where machines have room; it is not part of Omni-Compass's default setting and this setting is not labelled better.

The bill on a real cloud (results/live/AKS_BILL.md, run 37187059424, commit 5b2832f, four paired repetitions): **no difference in the bill either way** (allocation law -0.6% , compass law $+0.8\%$, intervals across zero) and no measure significantly worse. Azure's own autoscaler already ran the workload on about 1.86 of 4 workers, so this light workload leaves no machine to give back. The machine savings measured on kind, where native has no node autoscaler, do not carry to a cloud with one at this load; that is the reading of record. The preregistered size Standard_D2s_v5 is not allowed in the subscription's region; Standard_D2s_v4 (same 2 vCPU, 8 GiB and list price) was used. Three earlier attempts stopped at the load generator's placement before any measurement and are not results.

The burst bill test (written before its run, 2026-10-04)

The steady bill run left Azure's autoscaler nothing to do and Omni-Compass no machine to give back. The question a buyer asks is the bill under load that moves: on a real cloud, with Azure's own cluster autoscaler underneath, does native with Omni-Compass on top bill fewer machine-hours than native alone when demand rises and falls? Same design as "The bill on a real cloud" (fresh AKS cluster per arm, work pool 1 to 4 Standard_D2s_v4 under Azure's autoscaler, the bill metered every 15 s, arms native / compass / omni rotated, five repetitions), with the open-loop load in bursts (load_steps 1 6 1 8 1 6, six steps over 1,800 measured seconds, the same steps in every arm). Reported: machine-hours and the bill at list price, paired against native with their 95% intervals, beside every service gauge. Labelled by the one rule: a lower bill counts only if nothing is more than 2% worse.

The six organisms with the real cluster inside (written before its run, 2026-10-04)

Every benchmark runs the same six organisms, native against native with Omni-Compass on top: the four realms, the whole tower of 656 muscles, and the four stacked with every duplicate kept (1,226). Six organisms, two arms, twelve columns. The model grid (six) and the card in the loop (tools/run_hil.py, single card and eight cards) already do. The Kubernetes runs now do too (tools/run_kil.py, workflow six-kube on kind, input organism on aks-metered).

Each organism runs on the measured window's clock, 240 steps, with the real cluster as one more muscle:

- efferent: the organism's offered compute work at step k , $D(k) = \text{sum over its compute pools of } \text{lam_i}(k) / \text{sum of } \text{lam0_i}$, the seeded trace, identical in both arms. The load generator runs $r(k) = 1 + \text{round}((\text{LOAD_MAX} - 1) (D(k) - \min D) / (\max D - \min D))$ replicas, $\text{LOAD_MAX} = 6$ declared here, open-loop load.
- afferent: the cluster's watts (capture.csv, the declared power model, every 15 s, identical accounting in both arms) are added to the organism's source, zone and site power, as the card's watts are in the card harness.
- the clock (stated 2026-10-07, after the 1,000-copy runs): the organism must keep the window's clock. tools/run_kil.py records how long after the window its last step ended (behind_s); the report (tools/six_kube_report.py) shows it as "organism behind its window (s)", shown and not judged, and marks a repetition OFF THE CLOCK when either arm ended more than 5% of the window late, because its last steps then saw a cluster whose load schedule had already ended. The pairing stands (both arms slip alike), the mark stays on the cell. The remedy is a longer window for that size on that machine (240 steps, each at least

the machine's time to step the organism once), never a change to the organism or the law: the four stacked at 1,000 copies (1.7 million muscles) step in about 31 s on an 8-vCPU machine, so their window is 10,800 s (45 s steps), against 2,880 s for the tower at 1,000 copies.

Arms: native (the stacks' own controllers, Kubernetes alone, Omni-Compass not started) and compass (the compass law on every simulated muscle, rule 4 on the cluster, handed back at 90% of the window; the run is invalid if any knob is not handed back). Five paired repetitions per organism, order rotated, each arm on a fresh six-worker kind cluster, 960 measured seconds. Seed 6000 for every organism and arm.

Reported in `SIX_KUBE.md` (`tools/six_kube_report.py`): twelve columns of means, then per organism every gauge paired against native with its 95% interval. Cluster rows are measured; organism rows are models (evidence S). Each organism is labelled by the one rule: no measure more than 2% worse (CPU and host load shown, not judged), and a gain counts only where energy or the bill is lower with an interval wholly below zero.

Result: the six organisms with the real cluster inside (run 37217362568, commit d81d5ee)

Thirty jobs, six organisms times five paired repetitions, native against native with Omni-Compass on top (compass law), every arm valid, every simulated knob handed back. Full report `results/live/SIX_KUBE.md`; raw files `results/live/raw/run-37217362568/`.

On the real cluster, in every organism, native + Omni was late less often and answered faster:

Organism	Time over the line, native → Omni	p95, native → Omni (ms)	Failed requests, native → Omni
Compute / AI / Cloud	41.0% → 26.7% (better, -35%)	3,967 → 2,653 (better, -33%, inside the noise)	14.9% → 13.3%
Physics / Robotics / Autonomous	28.4% → 15.5% (better, -45%)	2,440 → 1,767 (better, -28%)	5.7% → 5.0%
Energy / Facility / Industrial	21.6% → 10.4% (better, -52%)	1,951 → 1,482 (better, -24%)	3.8% → 3.2%
Distribution / Specialized	40.7% → 21.1% (better, -48%)	2,830 → 1,943 (better, -31%)	11.1% → 9.3%
The whole tower (656)	25.8% → 14.7% (better, -43%, inside the noise)	2,085 → 1,247 (better, -40%, inside the noise)	12.8% → 9.1%
The four stacked (1,226)	55.6% → 43.0% (better, -23%)	3,703 → 2,803 (better, -24%)	31.7% → 27.6%

No measure came out worse beyond the noise in any organism. Energy (declared model) and CPU moved by under 3%, lower with Omni in every organism. The modelled organisms around the cluster: work the same, energy 0.1-0.2% lower, time over the line 0.5-2% lower, in every organism (evidence S).

Read with it: at a peak of six open-loop load generators the 4-core runner is past what the six workers can serve in both arms (native late 22-56% of the time, 4-32% of requests failed; host 65-91% busy). The comparison is fair (same load, same seed, same runner per pair); the absolute levels are an overloaded cluster, and a lower `load_max` is the setting for a cluster inside its capacity.

More work, faster, on fewer machines, with less energy: all four in one run (written before its run, 2026-10-04)

The four gains so far come from different tests: more work on the same machines (the capacity test, +48.1%), the same work faster on fewer machines (sets 22-27), energy equal or lower in each. This test measures all four in one run.

Design: benchmark-reps, ten paired repetitions, arms native and compass (rule 4), order rotated, fresh six-worker kind cluster per arm, open-loop load rising and then falling, `load_steps` 1 2 3 4 5 6 7 8 7 6 5 4 3 2 1, 180 s a step

(2,700 measured seconds). The rise is the capacity test (`tools/live_reps.py capacity`, read on the way up to the peak: the highest load step at which no more than 5% of response samples are over the 500 ms line or failed, every lower step too). The fall is where machines are no longer needed and can be handed back. Over the whole window: p95 response time, worker machines in service, energy (declared model; on kind it is not a meter), CPU with Omni's own.

Reported for each, native + Omni against native, paired, with the 95% interval and read in words (better or worse): work (capacity, requests a second inside the line), speed (p95), machines (worker machines in service, mean), energy.

The Omni index: one number for more for the same, or the same for less (written before its first use, 2026-10-04)

Each measure is turned into a ratio oriented so that above 1 is better for Omni-Compass:

- work: work with Omni / work native (capacity, requests served, work done)
- speed: p95 native / p95 with Omni (a lower response time is faster)
- machines: machines native / machines with Omni (in service, or billed machine-hours)
- energy: energy native / energy with Omni (or the bill)

The index of one test is the geometric mean of its oriented ratios, minus one, in percent: +30% reads "30% more for the same, or the same for 30% less", across work, speed, machines and energy together. A measure a test did not take is left out of that test, never filled in. A measure inside the noise is counted at its mean and flagged. A category (real Kubernetes on GitHub, Azure, the card, the eight cards, the modelled muscles) is the geometric mean of its tests; the headline is the geometric mean of the real categories, each weighted the same; the modelled muscles are shown beside it, never inside it. Computed by `tools/omni_index.py` from each test's own paired results; nothing is typed in.

Demand that wanders: up, spike, partway down, back up, down to idle (written before its run, 2026-10-04)

Real demand does not rise once and fall once. It climbs, spikes, eases part way, climbs again and finally settles to idle. This test drives that shape and asks whether machines follow it in order: the emptiest machine idles first on the way down (powered and Ready at its floor, never off), the warm machines wake first on the way up (no boot), one machine always in service for the first burst (`scripts/kind_nodepool.sh`).

Design as the all-four test (ten pairs, native and compass, open-loop load, fresh six-worker kind cluster per arm), with `load_steps` 1 2 3 2 3 4 5 6 5 3 5 6 5 4 3 2 3 2 1 1 (twenty steps of 135 s, 2,700 measured seconds), the same steps in every arm. Reported, native + Omni against native, paired with the 95% interval and read in words: time over the 500 ms line, failed requests, p95, worker machines in service (and its trace step by step), energy (declared model), CPU with Omni's own, pods started. No capacity is read (the load does not only rise).

Amendment to the wandering test, written before its run (2026-10-04)

The founder's two rules for real traffic, adopted before any result of the wandering test: traffic moves one step at a time, up or down, never skipping (it may go 1 2 1 2 1 2 if it never needs 3); and machines never go below two, so two are always in service, ready for a spike, with no ceiling but the machines the cluster has. The run started with the earlier steps (which jumped 5 to 3 and 3 to 5) was cancelled before it finished; none of it is reported.

From this amendment: the machine floor is two in every arm Omni-Compass governs (MIN_NODES, default 2: the controller's --min-nodes, the actuator scripts/kind_nodepool.sh, and every benchmark script). Native's own floor is the cluster's: kind keeps all six workers. The wandering steps are 1 2 3 2 3 4 5 6 5 4 5 6 7 8 7 6 5 4 3 4 3 2 1 2 1 (twenty-five steps of 108 s, 2,700 measured seconds), every change one step, the same in every arm. Everything else is as written above.

Result: all four in one run (run 37226863122)

Ten paired repetitions, every arm valid. **Work +29.3%** (capacity 24.6 to 31.8 requests a second inside the line, interval +5.4 to +9.0), **p95 -62.3%** (685.8 to 258.3 ms), time over the line -40.9%, failed requests -12.1%, **machines in service -3.6%**, **energy -0.3%** (declared model), CPU with Omni's own -1.9%: each proven. No measure worse beyond the noise (pending pod-minutes +19.6% and pods started +4.0%, both inside the noise). This run's machine floor was one; the floor of two was adopted after it started. Report results/live/ALL_FOUR.md; raw files results/live/raw/run-37226863122/; the Omni index updated (results/OMNI_INDEX.md).

The boundary is the band, not a machine held out (written 2026-10-04, for every run started after it)

Holding one whole machine back as a ceiling was considered and set aside the same day: it takes one of the machines away from the work. Every machine is usable (Omni-Compass's most is every machine that exists); the floor of two stays for quiet times. The protection against running into the wall is the band inside every machine: capacity is added at 95% of the response line, before the line is reached, and a machine goes back only if the ones left still run at or under the engine's utilisation target. An operator who wants whole machines held back can set NODE_CUSHION (off by default). The wandering run under way (37235231963) uses every machine, as every run after it does.

The six organisms at every size, the real cluster inside (written before its run, 2026-10-04)

The same six organisms, each native against native with Omni-Compass on top, at the sizes of the model grid: 10, 100 and 1,000 copies of the organism governed together on one clock, the one real cluster inside as one more muscle (tools/run_kil.py --scale; size 1 is the run already recorded, results/live/SIX_KUBE.md). Workflow six-kube, one job per organism, size and repetition, both arms on one runner, order rotated, fresh six-worker kind cluster per arm.

- Repetitions: 5 at 10 and 100 copies, 3 at 1,000. Measured window: 960 s at 10 copies, 1,440 s at 100, 2,880 s at 1,000 (240 organism steps of 4, 6 and 12 s: a step at least 1.5 times what the organism takes to compute on this runner, measured: 0.18 s at 10 copies of the four stacked, 2.4 s at 100).
- The organism is built before the window opens (organism.ready); the window opens the moment it is built (organism.go), so the cluster and the organism start on one clock however long the build takes. If the organism falls behind its own clock, organism.json records by how much (behind_s).
- Machine floor two, every machine usable (the founder's rules above), open-loop load, peak six generators.
- Declared before the run: 1,000 copies of the whole tower (656,000 modelled muscles) and of the four stacked (1.2 million) need about 6 and 11 GB of memory beside the cluster, past what one GitHub runner holds. They are run; if a runner cannot hold them, the job's failure is reported as that, never as a result.

Reported in SIX_KUBE.md by organism and size, every gauge paired against native with its 95% interval and read in words.

The run dimension of the model grid (1, 10, 100, 1,000 runs) is a count of seeds; on the real cluster each repetition is a real paired run of an hour or more, so the real cluster carries the repetitions above, and the model grid carries 1

to 1,000 runs.

The two organisms too big for a GitHub runner: a larger rented machine (written before their run, 2026-10-05)

1,000 copies of the whole tower and of the four stacked run on a larger rented machine, the same test otherwise: workflow `big-organism` rents one Azure Standard_D8s_v4 (8 vCPU, 32 GiB) per repetition, runs `scripts/kind_paired.sh` on it (native against native with Omni-Compass on top, fresh six-worker kind cluster per arm, 2,880 measured seconds, the organism at 1,000 copies with the real cluster inside), copies every file back and deletes the machine whatever happens. Three paired repetitions per organism, two machines at a time. Its results join `SIX_KUBE.md` at 1,000 copies; the same cells from GitHub runners count only if the runner held them.

The burst bill test on Azure: two native arms lost to the harness, and the fix (2026-10-05)

Repetitions 2 and 3 of the burst bill test (run 37216328055) each lost their native arm (Kubernetes alone) at the same moment: the load step from one generator to six. Native's two work machines were full, answers timed out, and the API server stopped answering for minutes. The capture script ran with stop-on-first-error, so one unanswered reading ended the recording, and the load schedule stopped at its next unanswered scale; the arm failed its checks and was marked invalid. The Omni-Compass arms of the same repetitions ran every step. Those two arms are not counted.

That is a fault of the harness, and it threw away exactly the moments a cluster struggles. Fixed: a reading the API server does not answer is logged (`capture.csv.errors`) and skipped, and the recording goes on; a scale the API server does not answer is tried again, then logged, and the schedule goes on (`fleet/capture/kube_capture.sh`, `scripts/kind_bench.sh`, `tools/run_kil.py`). Repetitions 4 and 5 started on the earlier code. When the run ends, three more repetitions of the same test run on the fixed code, and the report counts every valid pair and names every invalid arm with its cause.

Amendment 6: a pinned gauge is not a steady demand (written before its run, 2026-10-05)

The wandering test (run 37235231963, ten pairs) gave native + Omni-Compass p95 -54%, time over the line -38%, failed requests -14%, energy -0.3% and CPU -2%, each proven, and 1.5 more pod starts a run (+36%, proven). The founder holds that result back until the cause is fixed and the test is run again.

The cause, read from the run's own audit (repetition 4, 12 pod starts against native's 7): at 972 s the load eased from five generators to four, rule 3 found the demand steady, and the target was raised (50 to 190: fewer pods); the load then climbed back to eight and the pods were started again. The demand was not steady. The autoscaler stood at its replica cap, every pod as busy as it could be, so the reading (utilisation times pods) could not rise however much the load did: a pinned gauge reads flat.

Rule 5 (amends rule 3): a decision at which the autoscaler stands at its replica cap marks the window pinned; a window that touched the cap is not steady, so no raise of the target, and no return of a raised cap, until one whole autoscaler window after the cap was last touched (`omni_controller/controller.py _steady, self.pinned`; tested in `tests/test_compass_controller.py`). Nothing else changes.

The rerun: the same wandering test (ten pairs, native and compass, `load_steps` 1 2 3 2 3 4 5 6 5 4 5 6 7 8 7 6 5 4 3 4 3 2 1 2 1, 2,700 s, floor two, every machine usable). Reported in full, every gauge read in words; the first run is kept in its raw files and named beside the rerun.

Amendment 7: coasting, off the gas but still in gear (written before its run, 2026-10-05)

The founder's picture: a car in drive. At idle it creeps, ready (the floor of two machines). On the gas it speeds up at once (more pods, machines woken, no boot). Off the gas it coasts down gradually toward idle; it does not drop into neutral. The brake is for stopping (the reset hands everything back at once).

Rule 6: a raise of the HPA target (toward fewer pods) moves by at most `--coast-step` points of utilisation in one autoscaler window (default 25, `COAST_STEP`; 0 turns it off). The first wandering run raised the target 50 to 190 in one move; under rule 6 the same raise takes several calm windows (50, 75, 100, ...). A demand that comes back part way finds the pods still running and is met at once by a lower target, with no pod started again. A lower target (more pods) is never limited: the gas is always immediate. Rules 5 and 6 together: never ease off while the gauge is pinned, and then ease off a step at a time (`omni_controller/controller.py`; tested in `tests/test_convey.py`).

The rerun with rule 5 alone (run 37255770249) was stopped before it finished, to run the test once with both rules; none of it is reported. The rerun: the wandering test as written in amendment 6.

Amendment 8: cruise and the emergency brake, and the batch test (written before its run, 2026-10-05)

The founder's driving modes: off (native alone), watching (Omni-Compass reads, writes nothing), autopilot (gas, brake, idle), and two more for work that comes as a pile:

- **Rule 7, cruise:** work waiting for a place (pods pending) two decisions in a row puts every machine in service (`--max-nodes`), and they stay in service, without second-guessing, until the line has been empty and nothing is scaling up for two decisions (`--cruise-after`, default 2; 0 turns it off).
- **Rule 8, the emergency brake:** nothing waiting, no response-time breach, nothing scaling up, and the sensed services' demand at or under 0.05 of one pod's request (`--brake-demand`): the machines go straight to the floor (two) in one move. Every check of the release gate still holds (every sense live, the last command landed, the machines left at or under the utilisation target) except one machine per decision. Never below the floor.
- Tested through the live controller: `tests/test_cruise_brake.py` (in `verify.py`).

The batch test: a queue of jobs on the six-worker kind cluster (`WORKLOAD=batch`, `deploy/kind/batch-jobs.yaml`): one Kubernetes Job of 240 pods, 60 at a time (more than the workers hold, so work waits for a place), each hashing 3,000 MB (real CPU work, identical in every arm), opened with the measured window (1,500 s); the service's load generator stands at zero. Ten pairs, native and compass, order rotated. Reported, paired with the 95% interval and read in words: how long the queue took to finish, the worker machines in service after it finished, machines in service and energy over the whole window, CPU with Omni-Compass's own, pods started. Cruise must not slow the queue; the brake must show in the machines held after it.

The frozen engine (rules 1-8, commit 353903683009): the fault and fairness results (2026-10-05)

Both tests ran again with nothing changed but the engine: ten pairs each, native and omni, order rotated.

- The fault test (GitHub run 37262797634, `results/live/FAULTS.md`): better and proven on p95 (−60.1%), mean response (−34.5%), time over the line (−22.6%) and pending pods (−76.5%). Nothing came out worse beyond the noise. Recovery after the fault, paired means: machine down 42 s against 70 s, a blind probe 54 s against 66 s, a runaway pod 98 s against 105 s, a spike 266 s against 267 s.

- The fairness test (GitHub run 37262799317, results/live/FAIRNESS.md): better and proven on mean response (−15.0%). Nothing came out worse beyond the noise, the neighbour's app included.

These replace the earlier engine's fault and fairness sections of results/live/AMENDMENT_3_RUNS.md in the Omni index. That file stays as first measured.

The Azure burst bill test and the big organisms, running them (2026-10-05): Azure no longer offers this subscription Standard_D2s_v5 in eastus, so AKS workers are Standard_D2s_v4 (2 vCPU, 8 GiB, the same list price, USD 0.096 an hour). Both arms use the same size; the steady-load result of record (results/live/AKS_BILL.md) stays as measured on D2s_v5. The big-organism machine failed while seven kind nodes joined (kubelet-start). Fresh Ubuntu ships too few inotify watchers for that; the rented machine now raises them as kind's own documentation advises. Neither change touches the controller.

Amendment 9: idle read from the autoscaler's floor, and the batch test again (written before its run, 2026-10-05)

The batch test on the frozen engine (GitHub run 37262795697) showed rule 8 late. The queue emptied 664 s into the window, but the emergency brake fired at 1,574 s. The reason: a served service is never at zero CPU. The response probe alone keeps one pod at about 10% of its request, which reads as demand 0.10, above the brake's 0.05. Rule 8 now also counts a service as idle when its autoscaler stands at its least pods (minReplicas), wants no more, and runs at no more than half its target utilisation. Every other condition of rule 8 holds as before: nothing waiting, no breach, nothing scaling up, the release gate, every sense live, the last command landed, and the machines left carrying what runs now under the utilisation target. Test: tests/test_cruise_brake.py (it brakes at the idle floor, and it does not brake above the floor or busy at it).

Effect on what was already measured: none. The steady, fault, fairness, wandering and all-four runs on the frozen engine had the service at one pod with nothing waiting in 0 of 5,102 readings, so the amended rule could not have fired in any of them.

Three of the ten batch pairs stopped, all in the omni arm. The API server answered 500 under the batch load (60 jobs asking for 30 cores on a 4-core runner), and the recording script ended on an unanswered `kubectl top nodes`. That is a harness fault: a reading the API server does not answer is now logged and skipped, and the end-of-window reads are tried again. The batch test runs again, ten pairs, on this engine, with the design unchanged.

Amendment 10: two rows read for what they measure (written after the frozen-engine runs were seen, 2026-10-05)

Disclosed: this is a change in how two rows are read, made after the steady, wandering and all-four results were seen. Both rows stay in every table with their numbers and intervals; only the word in the last column changes.

- **Pending pods.** The row counts pods pending in each 15 s snapshot. A pod the autoscaler has just created is pending for the seconds it takes to schedule and start, so a faster scale-up shows as more pending pods. In the steady run the extra pending pods came at the load steps (about 180 s and 600 s), with every machine in service, and the API server's own record shows no pod in either arm ever unschedulable. The row is now shown, not judged. The judged gauge is the scheduler's own verdict, from the pod record: pods that no machine would take (PodScheduled False, reason Unschedulable), counted and timed in pod-minutes. In the five frozen-engine runs it is 0 in every arm.
- **Energy per core-hour.** It is energy divided by CPU used. CPU used is shown, not judged (more or less is not better by itself), so a ratio over it is not judged either. Energy itself stays judged.

Nothing else changes. With these readings, no row of the steady, wandering, all-four, fault or fairness results is worse than native beyond the noise.

Amendment 11: the burst bill test sized to what the service can serve, and rounding read as the same (2026-10-05)

Disclosed: written after three repetitions of the burst test at 1 6 1 8 1 6 were seen. In them about 41% of requests failed in both arms. The service's own autoscaler sat pinned at its cap of 10 pods, and 6 and 8 load generators (36 and 48 requests a second) ask more than 10 pods can serve, whatever the machines. A test where neither arm can do the work cannot show a difference. The burst test runs at 1 3 1 4 1 3: the peak is one step above the steady test's peak of 3, and inside what the capped service can serve. Everything else is as preregistered: five repetitions, 1,800 measured seconds, Azure's autoscaler underneath, the bill metered every 15 s, native against compass. The three repetitions at the old steps stay in the archive (results/live/raw/run-37294579764/) and are reported as the reason for this amendment, not as a result.

In every table a paired change under one part in a million of the value reads "same", the number still shown (SAME_REL, docs/MECHANISM_OF_ACTION.md 9.5).

Amendment 12: cruise steps back, and the organism's start file (written before the batch rerun, 2026-10-05)

The batch rerun on amendment 9 (GitHub run 37275371557, ten valid pairs) gave machines -16.9%, -27.2% once the queue was done, and energy -11.6% on the standby model. The queue finished 10 s later (584 s against 594 s, +1.8%). Omni held every machine in service throughout, so the time went elsewhere: every five seconds the floor step read every pod (240 batch pods), every node, node metrics and every HPA, on the one 4-core runner the queue was using. In cruise every machine is already in service, so those reads can change no action. From this amendment, while cruising the floor step makes no read and no write, and the service's target stays at the operator's own (docs/MECHANISM_OF_ACTION.md 9.3; tests/test_compass_controller.py). Cruise never engaged in the steady, wandering, all-four, fault or fairness runs, so this changes nothing there. The batch test runs again, ten pairs, design unchanged.

The organism with the real cluster inside waits for a start file. One repetition of six-kube (the four stacked, 1 copy, repetition 2) read that file in the instant it existed but was still empty, and stopped. The file is now written whole or not at all, and read only once it holds a number. That repetition is reported as stopped, never counted.

Also found and fixed on 2026-10-05, from the compass rename: each modelled muscle's compass was made fresh at every decision (docs/MECHANISM_OF_ACTION.md 9.1). That touched the modelled organism rows of the runs started after the rename (six-kube 37280090832, big organism 37287311224); their cluster rows are unaffected. The organism runs go again on the fixed code.

The batch test on amendment 12 (GitHub run 37344765837, ten pairs, results/live/BATCH.md): machines in service -20.7%, -32.4% once the queue was done, energy -14.3% on the standby model, mean response -11.7%, all proven. The queue finished +1.1% later, inside the noise (it was +1.8% and proven before cruise stepped back). No pod was left without a machine. Nothing came out worse beyond the noise.

The fleet that can show one machine: Azure at 11 and 15 workers (written before its run, 2026-10-07)

The Azure runs so far used a work pool of 1 to 4 machines. One machine is a quarter to a half of that fleet, and the app's replica ceiling (10 pods of 200m, about one worker's worth) kept Azure's autoscaler at about 1.9 workers in both arms, so nothing under a quarter of the fleet could show. The steady run on v1 (5 pairs) read even on every gauge. That is a statement about the lever, not about the law. This run makes the lever big enough to see one machine.

- **Fleet.** The subscription allows 200 vCPUs in eastus but 10 vCPUs in each machine family (every family's request to raise it was refused through the API on 2026-10-07: the subscription's terms), so one family gives at most 5 machines. The fleet is therefore several work pools, one 2-vCPU machine family each, each capped at what its family allows: `worker_pools = Standard_D2s_v4:4, Standard_D2as_v4:5, Standard_D2_v4:5, Standard_D2a_v4:5, Standard_D2ds_v4:5, Standard_D2s_v3:5, Standard_E2s_v4:5, Standard_E2as_v4:5`, 39 workers (the DSv4 family also carries the system machine, hence 4). Every pool carries the label `omni-role=work`, the worker selector in both arms; the first pool keeps one machine, the others may scale to zero; every pool starts full and Azure's own autoscaler trims, as before. One machine is then 2.6% of the fleet. Mixed machine families are how real cloud fleets run; the bill is each pool's machine-hours at its own list price (`pool_prices`), Azure's own count every 15 s. The same autoscaler profile, a fresh cluster per arm. (Written first for 11 then 15 workers of one family; the family allowance made that impossible and this replaces it before any such run.)
- **Ceiling.** `hpa_max = 9 × workers = 351` at 39 workers: nine 200m pods fit one 2-vCPU worker, so the fleet can fill. The same ceiling in every arm; the restore checks expect it back untouched.
- **Load.** The load generator's replica steps scale with the fleet, the same in every arm (one load replica drove about five pods at 50% on the 4-worker runs). Steady: 18 35 53 18 35 18 over 900 s at 39 workers (the peak asks for about 265 pods, thirty workers' worth). Burst: 11 35 11 53 11 35 over 1,800 s (the 4-worker burst was 1 3 1 4 1 3). A fleet of another size scales the steps by workers/39, rounded.
- **Everything else as preregistered:** native (Azure's autoscaler alone) against omni (Omni-Compass on top of it, rule 4 on the node pool, the HPA target inside its range, handed back at 90% of the window), 5 paired repetitions, order rotated, the bill Azure's own machine count every 15 s at list price, the reading by `tools/live_reps.py`'s paired interval, every gauge reported. Engine: Omni v3 first (the newest), v1 after if the credits allow.
- **What it can say.** Better, clear of the noise, in bill or response time: Omni has value on the managed service when the fleet is big enough to see a machine. Even again: Omni's value there is nil at this size too, and that is the reading. Worse: our own wiring is suspected first, found, fixed, and the run is made again.
- **Cost.** About \$8 for the steady set and \$15 for the burst set at 15 workers, at list price.
- **Amendment (2026-10-07 07:18 UTC, before the counted run).** The first dispatch at 39 workers was refused by Azure on every repetition: the DASv4 family had 2 vCPUs left, not 10, because the rented machine running the four stacked at 1,000 copies (big-organism-detached, a `Standard_D8as_v4`) holds 8 of that family's 10 in the same subscription. The run was cancelled (no arm ran) and dispatched again at **35 workers**: `Standard_D2as_v4:1`, the other seven pools as above; `hpa_max 315`; the steps scaled by 35/39 as this section says, steady 16 31 48 16 31 16, burst 10 31 10 48 10 31. One machine is 2.9% of this fleet. Nothing else changes. The 39-worker fleet stands for the runs made after the detached machine is deleted.
- **Amendment 2 (2026-10-07 09:40 UTC, before the counted run).** The 35-worker dispatch was refused too, on three repetitions: the DAv4 family (`Standard_D2a_v4`) has **no allowance at all** in this subscription

("remaining 0" is a limit of 0, not a holder), and the survey (big-organism-detached, mode survey, now listing every family's allowance and Azure's list prices) shows the same for the ESv4 and EAv4 families (Standard_E2s_v4, Standard_E2as_v4). Three of the eight pools could never be built. The fleet is rebuilt from families the survey shows allowed (10 vCPUs each) and offered in eastus: Standard_D2s_v4:4, Standard_D2as_v4:1, Standard_D2_v4:5, Standard_D2ds_v4:5, Standard_D2d_v4:5, Standard_D2s_v3:5, Standard_D2as_v7:5, Standard_D2ads_v7:5, Standard_D2als_v7:5, Standard_D2alds_v7:5: **45 workers** of ten families (one machine 2.2% of the fleet); hpa_max 405; the steps scaled by 45/39, steady 21 40 61 21 40 21, burst 13 40 13 61 13 40. List prices from Azure's own price API (Linux, pay as you go, eastus, USD a machine-hour): D2s_v4 0.096, D2as_v4 0.096, D2_v4 0.096, D2ds_v4 0.113, D2d_v4 0.113, D2s_v3 0.096, D2as_v7 0.0908, D2ads_v7 0.114, D2als_v7 0.0804, D2alds_v7 0.0952 (the 0.086 used for D2as_v4 in the refused dispatches was wrong; no bill was ever made with it). scripts/aks_paired.sh now checks every pool's size and family allowance before anything is built, so a refusal costs a minute, not a repetition. Nothing else changes.

- **Amendment 3 (2026-10-07 16:21 UTC, before the counted run).** The pre-flight of the 45-worker dispatch (run 37602951929) found Standard_D2s_v3 not offered to this subscription in eastus and stopped every repetition in a minute, nothing built, nothing billed; the nine other pools passed. That pool is dropped: **40 workers** of nine families (one machine 2.5% of the fleet), hpa_max 360, the steps scaled by 40/39, steady 18 36 54 18 36 18, burst 11 36 11 54 11 36, the pools' prices as in amendment 2. Nothing else changes.
- **Amendment 4 (2026-10-07 20:25 UTC, before the counted run).** Two dispatches of the 40-worker fleet (runs 37651333302 and 37676237815) were refused by Azure itself on every repetition for four hours: "creating a new cluster is unavailable at this time in region eastus" (AKSCapacityHeavyUsage, Azure's own capacity for new clusters, not the subscription's allowance; the pre-flight passed each time, no arm ran, nothing was billed). Microsoft's remedy is another region, other cluster settings, or a retry; this subscription has more than 10 cores only in eastus. The clusters are therefore made on the **standard control-plane tier** (tier input, AKS_TIER) instead of the free tier, about \$0.10 an hour a cluster, which is not in the bill (the bill counts worker machines only, in both arms alike) and changes nothing the workers or the autoscaler do. If the standard tier is refused too, the run waits for Azure. Nothing else changes.
- **Amendment 5 (2026-10-08 03:55 UTC, before the counted run).** The standard-tier dispatch (run 37703460186) was refused by Azure's own cluster capacity in eastus on all five repetitions too, over three hours (23:39 to 02:48 UTC), the same AKSCapacityHeavyUsage on each: 22 refusals in eleven hours across four dispatches, the pre-flight passing each time, no arm run, nothing billed but the resource groups' minutes. The fleet is dispatched again unchanged; it runs when the region admits a cluster. The alternative, a second region, needs the subscription's regional vCPU allowance raised there, which only the account holder can request in the portal. Nothing else changes.
- **Amendment 6 (2026-10-08 05:00 UTC, before the counted run).** The fifth dispatch (run 37724667075) was refused on its first repetition after a forty-minute wait, the same AKSCapacityHeavyUsage in eastus, the twenty-third refusal; the pre-flight passed, no arm ran. We cancelled the run ourselves at that point to weigh a second region, and the answer is in the 6 October survey already on record (big-organism-detached, mode survey): eastus allows this subscription 32 vCPUs and the v4 machine sizes, and every other region allows 10 vCPUs in all, which is five two-core workers, not a fleet of forty. No other region can take this fleet until the account holder raises its allowance there. The fleet is therefore dispatched a sixth time in eastus, unchanged in every input, and runs when the region admits a cluster. The cancellation and the reason are recorded here so the count of dispatches reads right. Nothing else changes.
- **Amendment 7 (2026-10-08 08:30 UTC, before the counted run).** The sixth dispatch (run 37730249970) was refused on all five repetitions too, the same AKSCapacityHeavyUsage in eastus, over three and a half hours

(05:01 to 08:26 UTC): 29 refusals in sixteen hours across six dispatches, the pre-flight passing each time, no arm run, nothing billed. The fleet is dispatched a seventh time, unchanged; nothing else changes. Each refused repetition costs GitHub runner minutes only.

Evaluation and simulation use only. Copyright (c) 2026 The Omni-Compass LLC. Commercial use, commercialization or monetization of any part of Omni-Compass requires a signed, paid Omni-Compass Enterprise License from The Omni-Compass LLC. Patent applications, copyright registrations and trademark applications filed in the United States. See LICENSE and NOTICE at the root of this repository.

CHAPTER 47

The Evidence Ledger

Every material statement in the technical package gets exactly one primary evidence class. A statement never borrows a stronger class from the one printed beside it.

Class	Meaning	Layer
T	formally proved theorem, under stated assumptions	L1 mathematical mechanism
V	finite computational verification: exact over a declared finite set, never a theorem beyond it	L1
S	simulation / model result	L2 modelled systems
L	live observation of external software (a real Kubernetes API, a real process)	L3 external software
P	independent physical measurement (a meter Omni does not read or control)	L4 physical plant
O	open: not established	—

No layer stands in for another: a theorem is not a physical validation, a Monte Carlo result is not a theorem, and a model's energy is not a meter's.

Mechanism identity for every row: `results/MECHANISM_IDENTITY.json`. Canonical engine `symmetric_verified`, mechanism id `29d9808dfb8f...`; the printed configuration `printed_eight_line`, id `cd333dc166fb...`, is a named alternative (Option A of the directive: one canonical, one alternative embodiment; no equivalence is claimed).

Layer 1: the mechanism

Class	Statement	Where
T	Unsaturated, the U-channel error obeys $de/dt = -KP e$; $V = e^2/2$ has $dV/dt = -KP e^2 < 0$ for $e \neq 0$ (inside Ω_{unsat} only).	docs/TRACKING_THEOREM.md Thm 1
T	With the U drift bounded by $F_{\text{bar}} < u_{\text{max}}$, V decreases whenever $e \neq 0$, saturated or not; e never changes sign; saturated stretches end in bounded time (continuous time).	Thm 2
T	On the declared parameter box, $F_{\text{bar}} = 16.86 < 25$, so Thm 2 holds over the whole box in continuous time.	Thm 2
T	The sampled controller (u held through each RK4 step) satisfies $e_{(k+1)} = 0.88 e_k + d_k$, hence $\text{abs}(e_k) \leq 0.88^k \text{abs}(e_0) + \epsilon_h (1 - 0.88^k)/0.12$.	Thm 3
V	$\epsilon_h = 0.00378$ over every micro step of the 500 frozen fixtures (nearest target); ultimate bound $0.0315 < \text{BASIN_TOL } 0.10$. Wrong target: 0.0243, bound 0.20 (wider than the basin).	results/TRACKING_BOUNDS.json
V	The discrete execution dithers inside that band: $\text{abs}(e)$ grew on 27 061 of 100 000 micro steps; e changed sign 747 times. The continuous monotone decay is not inherited.	same
V	Saturated micro steps: 0 of 100 000 (nearest target), 5 of 100 000 (wrong target).	same
T	The admissible box A (E interval, $\text{abs}(U - \sigma) \leq \text{abs}(U_0 - \sigma)$, S between S_0 and S_{plus}) is forward invariant in continuous time.	Thm 4
V	Discrete invariance of A: 0 failures over 100 000 micro steps of the frozen fixtures, both targets.	results/TRACKING_BOUNDS.json
O	Discrete invariance and ϵ_h proved over the whole box (interval arithmetic).	—
O	Global stability of the forced six-state system.	docs/FORMAL_STATUS.md
V	500 / 500 frozen fixtures satisfy CONVEY-5 and CERT-10 (dwell predicates, not external certification).	results/core_evidence.json
V	Frozen 500-fixture comparison: symmetric_verified final error $7.48e-5$, integrated $\text{abs}(u)$ 0.878, peak 13.93; printed_eight_line $5.13e-6$, 11.79, 18.87. Both executable; not interchangeable.	results/MECHANISM_IDENTITY.json
V	Python and C++ produce identical results (engine, governor, shield, HPA law, conveyance, GPU rules); twins sealed.	verify.py, results/SEAL.json

Layer 2: models and simulations

Class	Statement	Where
S	Synthetic stack: energy for Omni over the Kubernetes reference 168.9 → 121.7 kWh per run (documented-behaviour reference, not upstream controllers).	docs/CLAIMS_REGISTER.md C8
S	Fleet harness, recorded PlanetLab shapes: -40.6% energy vs HPA+CA, -21.7% vs Karpenter-lite (omni_fleet).	C17
S	Single GPU physics model: the engine alone would save 10.6–13.6% work per kJ on card A but breaks the p95 guardrail by 15–28%; with the frozen guards, about +1 to +5% inside it.	results/gpu/sim/FINDINGS.md
S	Node exchange (CPU and GPU on one budget): +1.4 to +5.7% work against the separate budgets, never over budget.	results/hardware/NODE_EXCHANGE_*.json
S	Six organisms (345, 262, 282, 337, the four stacked 1,226, the whole tower 656), the compass law on every muscle against each organism's own controllers, 1,000 paired runs at 1× and 10× size: work per energy +0.20% to +0.30%, energy -0.21% to -0.32%; every knob handed back.	results/scale/GRID.md
S	GPU card model, the card's firmware alone against the firmware with Omni on top (amendment 8: the verdict, 2% allowance), geometric means over 10 seeds and 10 fresh seeds: AI token generation, energy -3.25% / -3.72%, median +0.55% / +0.70%, p95 +0.29% / +0.26%; compute-bound, energy -0.70% / -0.48%, median +1.56% / +1.47%. A fixed 105 W power cap alone against the cap with Omni on top: energy -2.09% / -2.37% (AI token generation).	results/sim/gpu_two_wire/RESULT.md
S	Realm harness round 3 (every realm carries the shared spine), preregistered, seeds 3000-3009: the whole 656-muscle tower native against one governor on top, work per energy +0.1% (+0.1 to +0.1), violations +0.5 pp, SUPERIOR WITHIN GUARDRAILS. Realms: Energy +0.2% with +1.9 pp violations (tradeoff); Compute 0.0% with +2.1 pp (not established); Distribution -0.1% (worse); Physics -0.7% (worse). Rounds 1 and 2 kept, superseded.	results/realms/REALMS.md
S	Stacked organism (round 4, seeds 4000-4009): the four realm organisms on one clock, 1,226 muscles with every duplicate; stacked native equals the four realms alone on every seed. One governor over the stack: energy -0.14%, work per energy +0.02%, violations +1.4 pp, ENERGY IMPROVEMENT WITH SERVICE TRADEOFF; four separate governors about the same (+0.04%); one governor against four separate: -0.02% (WORSE, by a hair).	results/realms/stack/STACK.md

Layer 3: live external software

Class	Statement	Where
P	First real-GPU confirmation, NVIDIA A10 (Lambda), 10 paired repetitions, the card's own meter: work per energy +3.6% (+2.7 to +4.5, proven), GPU energy -3.5%, same requests, none lost; wire check 7 of 7, every write read back, every arm restored.	results/gpu/run-20261002T082232Z/GPU_REPS.md
L	Set 25 (commit 6fa7a97), real Kubernetes, 10 paired repetitions: machines in service -32.3%, p95 -57.3%, 0 failed requests, total CPU including Omni's own -0.6% (not proven).	results/live/LIVE_REPS_25.md
L	Set 26 (commit e7f920d), real Kubernetes, 10 paired repetitions, three arms: the allocation law machines -35.8%, p95 -55.4%; the compass law in the live controller machines -17.2%, p95 -64.8% , 0 failed requests, label by the preregistered rule <i>better on machines within the band</i> .	results/live/LIVE_REPS_26.md
L	Set 27 (commit d46c959), real Kubernetes, 10 paired repetitions, three arms: the allocation law machines -36.6%, p95 -53.1%; the compass law aligned with the GPU governor machines -15.9%, p95 -65.5%, p99 -72.6% , 0 failed requests, label by the preregistered rule <i>better on machines within the band</i> .	results/live/LIVE_REPS_27.md
L	Set 24 (2026-10-02, commit c908054), real Kubernetes (kind), 10 paired repetitions: machines in service -31.6% (proven), p95 response time -60.1% (proven), p99 -64.1% (proven), HPA replicas -38.6% (proven), 0 failed requests on both; total CPU including Omni's own -1.8% (not proven); energy with every machine powered -0.3% (declared model).	results/live/LIVE_REPS_24.md
L	Set 23 (set 22 repeated on the current code, 2026-10-02), real Kubernetes (kind), 10 paired repetitions, equal work: p95 response time -62% (proven), replicas -37% (proven), pods started -64% (proven), 0 failed requests; total CPU with Omni's own -1.0% and modelled energy -0.2% (no difference).	results/live/LIVE_REPS_23.md
L	Set 22, real Kubernetes (kind), 10 paired repetitions, equal work (fixed-rate load): p95 response time -61% (proven), replicas -23% (proven), pending pod-minutes -91% (proven), 0 failed requests.	results/live/LIVE_REPS_22.md
L	Set 21, real Kubernetes (kind), 10 paired repetitions: p95 response time -37% (proven), replicas -12% (proven), 0 failed requests.	results/live/LIVE_REPS_21.md
L	Omni patched a real Kubernetes API in place (no restart); the reset restored every setting in every run; watch mode wrote nothing.	same, results/live/
S	Set 21 energy is a declared model, not a meter: +1.8% worse with parked machines at idle power.	same (energy table)

Layer 4: independent physical measurement

Class	Statement	Where
O	Omni changes successful work per measured joule on a GPU. The bench is built with three receipts (governor, actuator, outcome), the three contrasts (observation, authority, total), and result labels by rule; it has not been run on a card.	scripts/gpu_paired.sh, docs/GPU_PREREGISTRATION.md
O	Production data-centre energy effect.	—
O	Whether the power limit is the right actuator for LLM serving. Published measurements (arXiv 2605.11999, H200) find memory-bound decode draws 137–300 W of 700 W, so no power cap binds; clock scaling is what saves energy there (arXiv 2501.08219, GreenLLM 2508.16449). The pinned bench workload is compute-bound, where the cap does bind: a result on it does not transfer to LLM decode.	external literature
O	Wall-plug (whole-machine) energy effect; CPU package and DRAM (RAPL) effect.	—
—	What the GPU bench will attribute: per write, joules and requests against native, and which rule decided it (engine, floor, gate, reflex, heat, speed lock).	tools/gpu_reps.py

Negative evidence, kept

Nothing here is deleted when a later result looks better.

Class	Statement	Where
P	Same run: p95 response time +58.5% worse (510 to 809 ms), mean +48%; label by rule ENERGY IMPROVEMENT WITH SERVICE TRADEOFF. Cause: governor wiring (busy bursts served below the card's own clock); corrected in amendments 6-7, not yet re-run on a card.	results/gpu/run-20261002T082232Z/GPU_REPS.md, docs/GPU_PREREGISTRATION.md
S	Two-wire GPU card, the profiles of amendment 7 (removed in amendment 8): service +5.3% work per energy with the median +14.8% slower; batch p95 +7.0% (−4.8 to +20.3). The governor before amendment 6 gave p95 +32.9% and +37.1%.	docs/GPU_PREREGISTRATION.md, history in git
S	Six organisms, same runs: time over the service line +0.19 to +0.27 pp worse in every cell ; band first is not held anywhere.	results/scale/GRID.md
S	Round 6 law (HPA target held at the operator's own, machine release margin 0.6), 20 paired seeds per organism, before the grid rerun: work per energy +0.086% to +0.201%, time over the line −0.010 to −0.031 pp (better than native in all six), work unchanged within its interval, every knob handed back.	docs/REALMS_PREREGISTRATION.md (round 6)
L	Set 21: modelled energy 1.8% worse with every machine powered (the only honest energy row on kind).	results/live/LIVE_REPS_21.md
S	Realm harness round 1 (superseded, kept): all five organisms worse (whole tower −0.1%). Its Omni layer did not follow the shipped controller (no contraction authority or SLO reflex, the stack law in place of the HPA, request traffic paused, a site budget under native draw).	results/realms/round1/
S	Realm harness round 2: the Physics / Robotics / Autonomous organism worse (−1.9%, violations +2.9 pp); the Compute organism's +2.5% costs +2.9 pp of service violations; 68 single muscles worse, mostly batch pacing and cooling setpoints under the live cooling law.	results/realms/REALMS.md

Class	Statement	Where
L	Set 22, equal work: no CPU saving once Omni's own CPU is counted (service -7.6%, controller +0.070 cores, together -0.9%, not proven); modelled energy unchanged (-0.1%).	results/live/LIVE_REPS_22.md
L	Set 20: the "same work" reading was wrong; requests were +35%.	results/live/LIVE_REPS_20.md addendum
L	Sets 1–2: machine savings withdrawn — a broken probe had blinded the latency sense.	results/live/LIVE_REPS_PROBE_DEFECT.md
L	Set 3 with a working probe: Omni kept all 6 machines; no significant difference except more waiting pods.	docs/HISTORY.md
L	Set 21: pending pod-minutes +238% (not proven); pod starts +37% (not proven).	results/live/LIVE_REPS_21.md
S	GPU and batch: energy 1–2% worse , power and heat margins 1–4% worse than tight packers (AKS/Karpenter, CAST AI, Spot).	docs/HISTORY.md
S	Machine round trips 1.1 → 1.7 per run (worse) in 6 of 8 seed-target combinations.	C12
S	Power-protect mode: backlog violations 2.5% → 6.7% (worse).	C12p
S	Park strategy on PlanetLab shapes: energy +1.3% vs HPA+CA, +33.6% vs Karpenter-lite (worse).	C17
S	Typical response time about 20% slower than every platform in the web and four-cluster simulations.	docs/HISTORY.md
S	GPU model: the engine without its guards breaks the p95 guardrail.	results/gpu/sim/FINDINGS.md
S	The compass-stroke GPU variant was tried and not adopted.	same
S	Right-sizing against VPA: p95 +15%, memory (OOM) kills +531%.	docs/history/BENCHMARK_REPORT.md
—	Reported in the external master-build report (not reproducible from this repository): on fresh scenarios Karpenter+VPA sometimes used less modelled energy than Omni, while Omni had lower churn and fewer request-induced evictions. Kept here so it is not lost; to be re-run here before it is cited.	external

Evaluation and simulation use only. Copyright (c) 2026 The Omni-Compass LLC. Commercial use, commercialization or monetization of any part of Omni-Compass requires a signed, paid Omni-Compass Enterprise License from The Omni-Compass LLC. Patent applications, copyright registrations and trademark applications filed in the United States. See LICENSE and NOTICE at the root of this repository.

CHAPTER 48

The Claims Register

Every claim, its evidence status and the command that reproduces it. Simulation results use the synthetic stack model in `omnicompass/stack_sim.py`; they are not production measurements.

ID	Claim	Status	Reproduce			
C1	The six-state mechanism is implemented identically in the reference engine, Python and C++.	Proven (exact parity tests)	<code>python verify.py</code>			
C2	The C++ governor reproduces the Python governor on recorded stack telemetry.	Proven (0 mismatches)	<code>python tests/test_cpp_governorparity.py</code> <oc_governor>			
C2s	The C++ shield reproduces the Python shield (enforced actions, intervention counts and violations) on action sets recorded from every arm type; a C++ shield with invariant I1 weakened is rejected by the parity test.	Proven (0 mismatches; negative control)	<code>python verify.py</code>			
C2h	An independent C++ implementation of the HPA replica law (ratio rule, 10% tolerance, scale-up limit, 300 s scale-down window) reproduces the fleet harness HPA on every recorded step (hundreds of thousands of steps across Kubernetes and governor arms); a C++ HPA with 5% tolerance is rejected.	Proven (0 mismatches; negative control)	<code>python verify.py</code>			
C3	The S channel cannot cross S- (forward invariance).	Proven (Proposition 1)	Manual Chapter 3			

ID	Claim	Status	Reproduce			
C4	The governor completed 100,000,000 consecutive decisions (951 years at 5-minute intervals) with 0 failures and constant memory; failure-rate upper bound 3.0e-08 at 95% confidence.	Measured (soak test); reproduced by the full verifier	python verify.py			
C4b	Throughput-mode soak: 100,000,000 decisions, 0 failures; resident memory at 25/50/75/100% of each run is constant ([3852, 3852, 3852, 3852] kB power-protect, [3916, 3916, 3916, 3916] kB throughput).	Measured (soak test)	python verify.py			
C8p	Primary endpoint (energy, throughput vs reference target 0.7): seed 346410161: -47.20 kWh, p = 1.0e-04; seed 360555127: -47.69 kWh, p = 1.0e-04. Secondary metrics Holm-adjusted over 24 tests: 21 better, 2 worse, 1 not significant.	Measured in simulation	python benchmarks/multiplicity.py			
C5	One governor decision costs 2.6 microseconds on one CPU core of the test machine (machine-dependent).	Measured	cpp: oc_soak			
C6	Observe mode leaves the stack bit-identical: over Kubernetes, identical to Kubernetes alone (500/500 and 500/500); over the native model, identical to native (500/500 and 500/500).	Measured in simulation	results/heldout_seed_*			
C7	Invariant violations per run (all arms scored identically): Kubernetes reference 9.56 / 9.44; power-protect 0.00 / 0.00; throughput 2.72 / 2.69, of which all are I4 (projected power), which throughput mode does not enforce; excluding I4: reference 6.91 / 6.80, throughput 0.00 / 0.00.	Measured in simulation	results/heldout_seed_*			

ID	Claim	Status	Reproduce			
C8	Energy for Omni-Compass over Kubernetes (throughput mode, same power rules as the reference) versus the Kubernetes reference model (HPA target 0.7): 168.9 to 121.7 (better) / 169.2 to 121.5 (better) kWh per run; across HPA targets 0.5 to 0.8: better in 8, worse in 0, not significantly different in 0 of 8 seed-target combinations.	Measured in simulation (synthetic stack, documented-behaviour reference model)	results/heldout_seed_*			
C9	Time healthy for Omni-Compass over Kubernetes (throughput mode) versus the reference model (target 0.7): 79.3% to 91.2% (better) / 78.5% to 91.4% (better); across targets: better in 8, worse in 0, not significantly different in 0 of 8 seed-target combinations.	Measured in simulation	results/heldout_seed_*			
C10	Recovery time for Omni-Compass over Kubernetes (throughput mode) versus the reference model (target 0.7): 57.2 to 17.8 (better) / 58.4 to 15.2 (better) minutes; across targets: better in 8, worse in 0, not significantly different in 0 of 8 seed-target combinations.	Measured in simulation	results/heldout_seed_*			
C11	Backlog violations for Omni-Compass over Kubernetes (throughput mode) versus the reference model (target 0.7): 2.5% to 1.6% (better) / 2.6% to 1.5% (better); across targets: better in 8, worse in 0, not significantly different in 0 of 8 seed-target combinations.	Measured in simulation	results/heldout_seed_*			

ID	Claim	Status	Reproduce			
C12	Machine round trips (started and later stopped), throughput mode versus the reference (target 0.7): 1.1 to 1.7 (worse) / 1.1 to 1.8 (worse) per run; across targets: better in 2, worse in 6, not significantly different in 0 of 8 seed-target combinations. The reference removes a machine only after 10 minutes below 50% utilization and keeps post-event machines running; the governor releases them, which is the main source of its energy saving.	Measured in simulation	results/heldout_seed_*			
C12w	Scale direction reversals (back-and-forth wear), throughput mode versus the reference (target 0.7): 2.19 to 1.93 (better) / 2.22 to 1.92 (better) per run; across targets: better in 4, worse in 4, not significantly different in 0 of 8 seed-target combinations. Power-protect: 2.19 to 3.06 (worse) / 2.22 to 3.08 (worse).	Measured in simulation	results/heldout_seed_*			
C12e	Engine mechanism in the flagship (throughput): removing the equation (2) release gate changes reversals from 1.93 / 1.92 to 3.19 / 3.18 and energy from 121.7 / 121.5 to 120.2 / 119.9 kWh; removing engine evolution changes energy to 125.4 / 125.1 kWh and reversals to 0.00 / 0.00.	Measured in simulation	results/heldout_seed_*			
C12p	Power-protect mode (enforces the site power limit, which the reference does not): energy 168.9 to 132.2 (better) / 169.2 to 132.3 (better) kWh; power-limit violations 12.8% to 3.5% (better) / 13.8% to 3.7% (better); backlog violations 2.5% to 6.7% (worse) / 2.6% to 7.0% (worse).	Measured in simulation	results/heldout_seed_*			

ID	Claim	Status	Reproduce			
C12a	Human pages: the governor issues no page actions. Zero pages is a design property, not a measured performance result.	Design	omnicompass/adapter.py			
C12b	The reference model is not the upstream Kubernetes controllers: Cluster Autoscaler scheduling simulation, Karpenter, VPA, scheduling constraints and disruption budgets are not represented.	Stated limitation	benchmarks/stack_benchmark.py (K8sReference)			
C15	15-second control plane, Omni-Compass as node-pool and power authority in place of the Cluster Autoscaler (HPA retained), energy vs HPA+CA / vs HPA+Karpenter-lite: energy-first: web -34.5% / -16.9%, reversals 4.3; multi -34.3% / -16.6%, reversals 18.1; batch -17.2% / -9.7%, reversals 13.6; gpu -4.9% / -2.7%, reversals 5.1; gpu_always_on -4.7% / -3.2%, reversals 0.0	balanced: web -33.2% / -15.3%, reversals 4.0; multi -33.0% / -15.0%, reversals 15.5; batch -12.4% / -4.5%, reversals 5.7; gpu -5.0% / -2.8%, reversals 4.2; gpu_always_on -4.7% / -3.2%, reversals 0.0	wear-first: web -29.7% / -10.9%, reversals 3.6; multi -29.6% / -10.6%, reversals 14.5; batch -12.3% / -4.3%, reversals 4.9; gpu -4.5% / -2.3%, reversals 3.0; gpu_always_on -4.2% / -2.8%, reversals 0.0	park: web -15.9% / +6.7%, reversals 0.0; multi -15.6% / +7.0%, reversals 0.0; batch -15.3% / -7.7%, reversals 0.0; gpu -4.9% / -2.7%, reversals 0.0; gpu_always_on -4.7% / -3.2%, reversals 0.0. Baseline reversals (CA / Karpenter): web 1.1 / 6.8; multi 5.5 / 25.4; batch 1.4 / 24.9; gpu 2.8 / 11.2; gpu_always_on 0.0 / 0.0. In gpu_always_on no arm powers a node off.	Measured in simulation (fleet harness, synthetic workloads, documented-behaviour execution layer)	python -m fleet.benchmark --seeds 30 --seed-base 7000000 --out out/

ID	Claim	Status	Reproduce			
C17	Recorded PlanetLab utilization shapes with a declared scale mapping (9 traces, 30 scenarios), governor as node-pool authority in place of the Cluster Autoscaler: omni_fleet: energy -40.6% vs HPA+CA (better), -21.7% vs Karpenter-lite (better); omni_fleet_balanced: energy -37.5% vs HPA+CA (better), -17.5% vs Karpenter-lite (better); omni_fleet_wear: energy -31.0% vs HPA+CA (better), -8.9% vs Karpenter-lite (better); omni_fleet_park: energy +1.3% vs HPA+CA (worse), +33.6% vs Karpenter-lite (worse). Time healthy and work completed are significantly lower by 0.08 points and 0.01%; observe mode identical in 30/30.	Measured in simulation driven by recorded traces (frozen laws, no retuning)	<pre>python -m fleet.planetlab --dir fleet/traces/planetlab --scenarios 30 --seed-base 8000000 --out out/</pre>			
C16	Live-cluster capture, capture replay and the PlanetLab vessel are tested on generated inputs in the real formats. No real capture or recorded trace has been run in the package.	Tested path; no real-data result	<pre>python tests/test_fleet_realdata_paths.py</pre>			
C19	The live controller (omni_controller/) implements observe, target and nodepool modes with dry-run, audit log and a reset that restores HPA targets from annotations; tested against a fake kubectl only, never against a real cluster.	Tested path; no live result	<pre>python tests/test_omni_controller.py</pre>			
C18	Savings projection: the energy-first reduction relative to HPA + Karpenter-lite, applied to declared fleet profiles (results/SAVINGS.csv), computed identically in Python and C++. A projection from simulation, not measured savings.	Projection	<pre>python benchmarks/savings.py</pre>			

ID	Claim	Status	Reproduce			
C21	End-to-end self-pilot (shipped controller, simulated cluster, real capture and scoring), default headroom 50%: energy per core-hour -7.7%, node-hours per core-hour -13.8%, pending-pod time not significantly different from HPA + Cluster Autoscaler; HPA shortfall minutes higher. Lower headroom saves more energy with more pending-pod time (manual Section 8.12a).	Measured in simulation	<code>python pilot/selfpilot.py</code>			
C20	pilot/score.py scores a user's own captures (node-hours and energy per used core-hour, utilisation, pending-pod and HPA-shortfall minutes, bootstrap intervals); tested to detect a real gain, report no difference for identical clusters and detect a service regression.	Tested tool; no pilot result	<code>python tests/test_pilot_score.py</code>			
C13	Decision components (autoscalers, power agents, paging, Terraform as controller) consume about 0.02% of fleet CPU; idle capacity is 92% of fleet CPU at 8% utilization.	Modeled from published figures and stated assumptions	<code>python benchmarks/fleet_overhead.py</code>			
C14	Behaviour on production systems.	Not established; requires the pilot protocol	<code>docs/PILOT_PROTOCOL.md</code>			

Evaluation and simulation use only. Copyright (c) 2026 The Omni-Compass LLC. Commercial use, commercialization or monetization of any part of Omni-Compass requires a signed, paid Omni-Compass Enterprise License from The Omni-Compass LLC. Patent applications, copyright registrations and trademark applications filed in the United States. See LICENSE and NOTICE at the root of this repository.

The Benchmark Report

Dated record, kept as written on 26 September 2026. Where a figure here differs from docs/STATE_OF_PLAY.md, the State of Play governs. The GPU power-limit figure in this report (15% to 19% energy, under 1% slower) came from an earlier model of the card; later models with the service guards give +1.3% to +5.1% work per kJ with one wire (results/gpu/sim/after) and +8.6% with two wires at a higher p95 (results/sim/gpu_two_wire/RESULT.md). No real card had finished the bench on that date; the first card result (NVIDIA A10, 2026-10-02) is in results/gpu/run-20261002T082232Z/GPU_REPS.md and docs/STATE_OF_PLAY.md.

Benchmark report, 26 September 2026. Repository: Omni-Compass/The-Omni-Compass-Control-Core-Engine (private), branch main Every number below is produced by code in that repository and can be regenerated; section 21 gives the commands. Each result states whether it was **measured on a live Kubernetes control plane** or **computed in simulation**.

1. Summary

Omni-Compass is a single control engine that senses the whole compute stack and drives its actuators (its "muscles": replica counts, node pools, power caps and others) from one six-state dynamical model, with a safety shield before every action and a kill switch that hands control back. It was compared in three architectures:

- **A. Kubernetes alone.** Kubernetes' own controllers decide: Horizontal Pod Autoscaler (HPA) for replicas, Cluster Autoscaler for nodes; other managers act on their own proposals.
- **B. Kubernetes + Omni-Compass.** Kubernetes' controllers keep running; Omni-Compass governs on top of them (sets the HPA target, gates and sizes the node pool, caps power) as the single authority over their settings.
- **C. Omni-Compass direct.** Omni-Compass is the only decision-maker and actuates the muscles directly; the separate managers no longer decide.

Main result (pre-registered, 1,000 held-out scenarios, simulation). Against Kubernetes alone, B used -28% energy and C -23%; time healthy rose from 79% to 91% (B) and 90% (C); recovery time fell from 58 to 17 and 28 minutes; contradictory commands, pages and human interventions went to zero in both; safety-rule violations fell from 9.5 to 2.7 (B) and 0.0 (C). Of 28 gauges, B is significantly better on 20 and worse on 5; C is better on 16 and worse on 10.

Live Kubernetes result (measured). Two identical Kubernetes clusters (1 control plane + 6 workers) ran the same load at the same time, one without Omni-Compass and one with it. With Omni-Compass: worker nodes in service 6.0 to 3.2; utilisation of the workers in service 0.068 to 0.116. **Energy:** with the parked workers kept on standby, powered and ready (100 W each, the same as idle), energy was 219 vs 218 Wh (-0.46%); parking alone saves essentially nothing; energy per unit of work +8% (not significant; 872 to 944 Wh per core-hour, from minute averages). The -43% first reported for this run holds only if parked workers are powered off. Waiting pods and

HPA shortfall were not significantly different. The kill switch restored the original HPA target (50) and all 6 workers. **With every live muscle switched on (section 7.2) the application got slower: p95 response time 486 to 802 ms, energy per unit of work +46% (significant).** The causes were found and fixed across runs 2-4 (section 7.4): response time went from +65% to a tie at p95.

Where Omni-Compass costs something. In the pre-registered study both B and C keep more node-hours powered than Kubernetes alone and start and stop machines more often (more wear), and move the power cap more; C also lets more work wait in the queue and flips scale direction more often. In that study the energy saving comes from power capping and load shaping, not from switching machines off. On the live cluster the saving came from switching machines off. Other limits: the small-cluster release band (section 11); power and heat on the live cluster are modelled, not metered.

Trade-off in one line: B is the strongest all-round result in simulation (energy, health, recovery, queue, coordination and safety all better; wear and node-hours worse); C is the strongest on peak power, heat and safety (zero invariant violations) at the cost of queue length, wear and flip-flops. These are the gauges to tune next.

2. What Omni-Compass is

2.1 The engine

The engine is a six-state ordinary differential equation system, state $x = (E, U, I_U, S, B, B_{\dot{}})$: error E , coherence U , pressure I_U , stress S and a damped bath B . The shipped core (`omnicompass/core.py`) integrates it with fourth-order Runge-Kutta and holds the control input constant across the four stages:

```
(1) dE/dt = -alpha_E E + beta_int + beta_ext + v_eff
(2) dU/dt = mu U (1 - U^2) - (dE/dt)/E_max - lambda_U U + u, |u| <= 25
(3) dI_U/dt = (1 - U) - sigma_1 E - delta S - lambda_I I_U
(4) v_eff = cos(omega_B t / 2) c tanh(lambda_0 + lambda_1 (U - 0.5) + lambda_2 S)
(5) Phi(S) = alpha_s S^2/2 + beta_s S^3/4 - delta S
(6) dS/dt = -dPhi/dS
(7) dB/dt = B_dot; dB_dot/dt = gamma_c delta S - (omega_B/Q_B) B_dot - omega_B^2 B
(8) R_B[n] = finite-difference audit of (7), never fed back
Controller: u = clip(-f_U(x,t) + 12 (sigma - U), -25, +25)
```

Telemetry (load, queue, power, heat, network, drift, staleness, security) is assimilated into the state each decision; an allocation law turns the state into a demand target ρ^* (the HPA target), a node change, and a power cap. Release of capacity is gated by equation (2): capacity is only released once the control push has converged.

2.2 The nervous system (muscles)

Each muscle has five parts: afferent (pull: sense), the shared engine, efferent (push: act), reflex (the shield checks every push) and kill (hand the muscle back to its own controller). Status in this repository: **wired live on Kubernetes** (sense and push executed through `kubectl`, each with shield and kill); nodes, HPA target, power cap (in-place CPU limits, enforced by the kernel), security hold, deployment rollouts (pause, resume, undo), batch queue (admit held Jobs); **sensed**: heat (harness heat law on live power, or GPU temperature), network; **hardware connectors** (built and tested with fake hardware, off on CI machines): CPU power states (RAPL read, `cpufreq` ceiling) and GPU (`nvidia-smi` power and temperature read, power limit); **open** (registered, no plant yet): memory, storage, cooling, grid, training, inference, agent containment and the rest of the 52-muscle domain map (`docs/DOMAIN_MAP.md`). AI value alignment is explicitly not an Omni-Compass muscle. Code: `omni_controller/controller.py`, `omni_controller/muscles.py`, `omnicompass/nervous.py`.

2.3 The shield and the kill switch

Before any action the shield (`omnicompass/shield.py`) enforces invariants I1 to I5: no expansion during a security block, node count within bounds, step limits, never below the capacity running and pending work needs, and the site power limit. The kill switch (a file or `OMNI_KILL=1`) restores every HPA target Omni-Compass changed from the recorded original, returns the node pool to its native size and drops to observe mode. Every decision and action is written to an append-only audit log.

2.4 Operating modes

Observe (compute and log, write nothing), target (write HPA targets), nodepool (also size a node pool). Laws: power_protect (enforce the site power envelope), throughput (power envelope not enforced; capacity sized at full power) and fleet variants tuned on the 15-second fleet harness.

3. The three architectures, and how each study realises them

Study	A. Kubernetes alone	B. Kubernetes + Omni-Compass	C. Omni-Compass direct	Live or simulated
Pre-registered held-out stack benchmark (2 x 500 scenarios)	<code>k8s_ref_70</code> : documented HPA law (target 0.7, 10% tolerance, 300 s stabilisation) and Cluster Autoscaler (scale-up on backlog, remove after 10 min under 50%), other managers act on their own proposals	<code>omni_k8s_throughput</code> : the same Kubernetes loops, Omni governs on top	<code>omni_direct</code> : Omni senses and actuates the stack directly	simulation
Control-plane replica (24 scenarios)	<code>hpa70_ca</code> : metrics-server, HPA and Cluster Autoscaler replicas at 15 s	<code>omni_target_gate_hpa70_ca</code> : Omni writes the HPA target and gates Cluster Autoscaler scale-down	<code>omni_throughput_full</code> : Omni writes the HPA target, owns node scale-down and power cap; the Cluster Autoscaler may only add nodes	simulation
PlanetLab-shaped demand, fleet plant (8 scenarios)	<code>k8s_hpa70_ca</code>	<code>omni_target</code> : Omni writes the HPA target	<code>omni_fleet</code> : Omni is the node-pool authority, Cluster Autoscaler off	simulation on recorded traces
Live kind cluster, side by side	native: HPA only, 6 workers always on, Omni not running	Omni on top: sets the HPA target and is the sole node-pool authority (cordon, drain, uncordon); Kubernetes' scheduler, kubelet and HPA still execute	not yet built live (section 16)	live

4. Method and why the comparison is fair

- **Frozen before testing.** The allocation law, engine parameters, shield limits, modes, baselines and benchmark code were selected on development seeds (1000, 2000) and frozen, with SHA-256 hashes and program fingerprints in `results/PREREGISTRATION.json`, before any run on the held-out seeds 346410161 and 360555127. `verify.py` re-checks every hash.
- **Same scenarios for every arm.** Each scenario is run under every architecture; comparisons are paired scenario by scenario.
- **Observe-identity check.** Omni-Compass in observe mode must produce a trajectory bit-identical to the arm it observes, or the run is invalid. Held-out: 500 and 500 of 500 identical to the native stack, 500 and 500 of 500 identical to Kubernetes; control-plane replica: 24 of 24; PlanetLab: 8 of 8; live: 0 writes while observing.

- **Statistics.** Differences are paired (Omni minus Kubernetes on the same scenario) with 95% bootstrap confidence intervals. In the held-out study a difference is called better or worse only if the interval excludes zero in the same direction on both independent seeds; 'better in N' counts scenarios. Live results use 2-minute blocks and a bootstrap over blocks.
- **Ablations** show the engine, not an accident of tuning, produces the result (section 5.3).
- **Independent implementations.** A C++ engine, governor, shield and HPA law are checked against the Python ones (section 8).

5. Results: pre-registered held-out benchmark (1,000 scenarios, simulation)

Mean over two held-out seeds x 500 scenarios; verdicts require both seeds to agree.

Gauge	A. Kubernetes alone	B. Kubernetes + Omni-Comp ass	C. Omni-Compass direct	B vs A	C vs A
ENERGY AND POWER					
Energy per scenario (kWh)	169.1	121.6	129.7	-28% better (better in 991, worse in 9 of 1000)	-23% better (better in 1000, worse in 0 of 1000)
Peak power (kW)	37.6	37.6	28.3	+0.16% not significant (better in 457, worse in 533 of 1000)	-25% better (better in 1000, worse in 0 of 1000)
Node-hours	135.8	140.5	152.4	+3% worse (better in 464, worse in 536 of 1000)	+12% worse (better in 191, worse in 809 of 1000)
Idle node-hours (powered, doing nothing)	46.29	51.73	62.03	+12% worse (better in 446, worse in 554 of 1000)	+34% worse (better in 202, worse in 798 of 1000)
Share of time over the power limit	0.133	0.067	0.018	-50% better (better in 490, worse in 24 of 1000)	-87% better (better in 517, worse in 5 of 1000)
Power-cap travel (how much the cap moved)	0.236	0.887	0.306	+276% worse (better in 37, worse in 963 of 1000)	+30% worse (better in 243, worse in 757 of 1000)
HEAT					
Share of time over the heat limit	0.084	0.026	0.002	-69% better (better in 353, worse in 1 of 1000)	-98% better (better in 354, worse in 0 of 1000)
Thermal travel (how far temperature swung)	0.754	0.632	0.576	-16% better (better in 828, worse in 172 of 1000)	-24% better (better in 993, worse in 7 of 1000)
SPEED AND BACKLOG					
Mean queue (work waiting)	384	217	933	-43% better (better in 218, worse in 515 of 1000)	+143% worse (better in 43, worse in 506 of 1000)
95th-percentile queue (worst moments)	1403	950	2945	-32% better (better in 247, worse in 190 of 1000)	+110% worse (better in 1, worse in 383 of 1000)
Share of time over the backlog limit	0.026	0.015	0.080	-42% better (better in 95, worse in 88 of 1000)	+207% worse (better in 1, worse in 267 of 1000)
RELIABILITY AND RECOVERY					
Availability	0.9990	0.9999	0.9974	+0.09% better (better in 110, worse in 127 of 1000)	-0.16% worse (better in 122, worse in 121 of 1000)
Share of time healthy	0.789	0.913	0.901	+16% better (better in 661, worse in 4 of 1000)	+14% better (better in 659, worse in 4 of 1000)
Share of incidents recovered	0.918	0.996	0.972	+8% better (better in 78, worse in 0 of 1000)	+6% better (better in 54, worse in 0 of 1000)

Gauge	A. Kubernetes alone	B. Kubernetes + Omni-Compass	C. Omni-Compass direct	B vs A	C vs A
Recovery time (minutes)	57.8	16.5	27.9	-71% better (better in 602, worse in 0 of 1000)	-52% better (better in 580, worse in 0 of 1000)
Physical SLA breaches (power, heat, backlog)	0.144	0.073	0.089	-49% better (better in 495, worse in 25 of 1000)	-38% better (better in 449, worse in 46 of 1000)
All SLA breaches	0.168	0.073	0.089	-57% better (better in 655, worse in 3 of 1000)	-47% better (better in 611, worse in 37 of 1000)
WEAR AND TEAR					
Machines started	4.8	6.6	9.4	+37%	+95%
Machines stopped	1.2	3.3	2.7	+185%	+130%
Node start/stop events	6.0	9.9	12.1	+66% worse (better in 56, worse in 907 of 1000)	+102% worse (better in 199, worse in 782 of 1000)
Machine round trips (stopped then restarted)	1.07	1.73	2.70	+61% worse (better in 119, worse in 464 of 1000)	+152% worse (better in 26, worse in 866 of 1000)
Scale reversals (flip-flops)	2.20	1.93	5.45	-13% better (better in 474, worse in 199 of 1000)	+147% worse (better in 86, worse in 803 of 1000)
CONTROL QUALITY AND SAFETY					
Contradictory commands between managers	9.46	0.00	0.00	-100% better (better in 666, worse in 0 of 1000)	-100% better (better in 666, worse in 0 of 1000)
Safety-rule (invariant) violations	9.50	2.71	0.00	-72% better (better in 664, worse in 1 of 1000)	-100% better (better in 669, worse in 0 of 1000)
Invariant violations excluding power	6.86	0.00	0.00	-100% better (better in 666, worse in 0 of 1000)	-100% better (better in 666, worse in 0 of 1000)
Security violations	0.889	0.000	0.000	-100% better (better in 118, worse in 0 of 1000)	-100% better (better in 118, worse in 0 of 1000)
Pages to on-call	0.584	0.000	0.000	-100% better (better in 92, worse in 0 of 1000)	-100% better (better in 92, worse in 0 of 1000)
Human interventions required	0.284	0.000	0.000	-100% better (better in 92, worse in 0 of 1000)	-100% better (better in 92, worse in 0 of 1000)

5.1 What the numbers say

- **Energy:** B -28%, C -23% versus Kubernetes alone, although both keep slightly more node-hours powered. The saving comes from power capping and load shaping: time over the power limit halves (B) or nearly vanishes (C), and C cuts peak power by a quarter. Releasing idle machines, which the live cluster showed, is not what drives this study; idle node-hours are a gauge to tune.
- **Reliability:** time healthy and recovery improve sharply in both B and C; the share of incidents recovered rises.
- **Coordination:** separate managers issue contradictory commands (A); a single authority issues none (B, C). Pages and human interventions go to zero because Omni-Compass acts on the conditions that would have paged someone.
- **Heat:** C cuts time over the heat limit the most, because it governs power caps and load together.
- **Costs:** node start/stop events rise (B +66%, C +102%), machine round trips rise, the power cap moves more, and C's queue and scale reversals are higher than Kubernetes alone. The 24-scenario control-plane study (section 6) shows the opposite for wear (fewer machine stops), so wear depends on the plant and law and is a primary tuning target. These are reported, not tuned away.

5.2 Architecture A has a hidden cost: the fragmented stack without Kubernetes

For reference, the stack with each manager acting alone and no Kubernetes loops (native) used 131.9 kWh, was healthy 39% of the time, took 127 minutes to recover and produced 56.7 contradictory commands and 28.1 invariant violations per scenario. Kubernetes already improves on that; Omni-Compass improves on Kubernetes.

5.3 Ablations: is it the engine?

Comparison	Metric	Mean difference	95% CI (seed 1)	95% CI (seed 2)
Omni direct vs omni_direct_no_dynamics	energy_kwh	-1.954	[-2.061, -1.775]	[-2.158, -1.834]
Omni direct vs omni_direct_no_dynamics	invariant_violations	+0.000	[+0.000, +0.000]	[+0.000, +0.000]
Omni direct vs omni_direct_no_dynamics	mean_queue	+39.200	[-4.742, +71.297]	[+7.307, +76.997]
Omni direct vs omni_direct_no_dynamics	recovery_minutes	+0.950	[-0.090, +2.650]	[-0.410, +2.260]
Omni direct vs omni_direct_no_dynamics	time_healthy	-0.000	[-0.004, +0.002]	[-0.002, +0.003]
Omni direct vs omni_direct_no_dynamics	violation_heat	-0.002	[-0.004, -0.000]	[-0.005, -0.000]
Omni direct vs omni_direct_no_shield	energy_kwh	-0.759	[-0.864, -0.642]	[-0.875, -0.660]
Omni direct vs omni_direct_no_shield	invariant_violations	-0.674	[-0.750, -0.592]	[-0.760, -0.596]
Omni direct vs omni_direct_no_shield	mean_queue	+121.945	[+90.814, +149.893]	[+100.054, +152.127]
Omni direct vs omni_direct_no_shield	recovery_minutes	+1.535	[+0.560, +3.210]	[+0.480, +2.600]
Omni direct vs omni_direct_no_shield	time_healthy	+0.001	[-0.001, +0.003]	[-0.000, +0.003]
Omni direct vs omni_direct_no_shield	violation_heat	-0.000	[-0.000, +0.000]	[-0.001, +0.000]

'no_dynamics' runs the same allocation law with equations (1) to (7) frozen; 'no_shield' removes the shield. The difference is full engine minus ablation: negative energy means the evolving dynamics save energy the frozen engine does not; the shield ablation shows what the shield prevents.

6. Results: Kubernetes control-plane replica (24 scenarios, simulation)

Replicas of the documented metrics-server, HPA (15 s) and Cluster Autoscaler (10 s) loops, with Karpenter-lite as a second Kubernetes reference; Omni-Compass decides every 300 s.

Gauge	A. Kubernetes alone (HPA + Cluster Autoscaler)	Reference: HPA + Karpenter-lite	B. Kubernetes + Omni-Compass (HPA target + CA gate)	C. Omni-Compass authority (HPA target, nodes, power cap)	C vs A (95% CI, wins of 24)
Energy (kWh)	168.8	161.7	155.6	122.3	-28% better [-49.724, -43.292], 24 of 24
Peak power (kW)	38.7	39.5	40.8	32.8	-15% better [-8.962, -2.584], 17 of 24
Time over power limit	0.050	0.050	0.048	0.015	-70% better [-0.051, -0.019], 14 of 24
Time over heat limit	0.045	0.047	0.045	0.012	-73% better [-0.048, -0.018], 13 of 24
Time healthy	0.930	0.937	0.931	0.983	+6% better [+0.033, +0.073], 18 of 24
Recovery time (min)	1.857	1.982	3.430	0.430	-77% better [-1.948, -0.906], 15 of 24
Physical SLA breaches	0.031	0.032	0.031	0.009	-71% better [-0.032, -0.013], 14 of 24

Gauge	A. Kubernetes alone (HPA + Cluster Autoscaler)	Reference: HPA + Karpenter-lite	B. Kubernetes + Omni-Compass (HPA target + CA gate)	C. Omni-Compass authority (HPA target, nodes, power cap)	C vs A (95% CI, wins of 24)
Mean queue	0.204	0.204	1.011	0.683	+235% worse [+0.168, +0.836], 0 of 24
Machines started	21.8	24.9	25.5	20.0	-8% better [-3.542, -0.250], 4 of 24
Machines stopped	13.2	17.0	24.1	3.833	-71% better [-11.917, -7.083], 24 of 24
Scale reversals	1.333	1.917	2.000	1.000	-25% better [-0.667, -0.083], 4 of 24
Invariant violations	0.000	0.000	0.000	0.000	0% not significant [+0.000, +0.000], 0 of 24
Pages to on-call	0.000	0.000	1.792	0.458	new not significant [+0.000, +1.375], 0 of 24

7. Results: live Kubernetes (measured)

7.1 Side by side, two identical clusters at the same time

kind clusters (real Kubernetes API server, scheduler, kubelet, HPA and metrics-server), 1 control plane + 6 workers each; the official php-apache HPA workload (target 50); the same stepped load (1, 2, 3, 1, 2, 1 load generators over 20 minutes after a 2-minute warm-up); capture every 15 s. Native: Omni-Compass not started. Omni: nodepool mode, HPA target, node pool, power sensing. Run 36209933218, 2026-09-26.

Gauge	Kubernetes alone	Kubernetes + Omni-Compass	Change
Duration (min)	20.10	20.13	
Worker nodes in service, mean	6.00	3.20	-46.7%
Worker nodes in service, min	6.00	3.00	-50.0%
CPU used (cores), mean	0.819	0.745	-9.1%
CPU allocatable (cores), mean	12.00	6.40	-46.7%
Utilisation	0.068	0.116	+70.5%
Workers powered (in service + standby)	6.00	6.00	0
Node-hours in service	2.01	1.07	-46.6%
Power (W), mean, parked workers on standby at 100 W	657	654	-0.5%
Power (W), peak, parked workers on standby at 100 W	691	687	-0.6%
Energy (Wh), parked workers on standby at 100 W	219	218	-0.5%
Energy (Wh), parked workers in low-power standby at 50 W	219	171	-22%
Energy (Wh), parked workers in deep sleep at 10 W	219	133	-39%
Energy (Wh), parked workers powered off (as first reported)	220	125	-43%
Pending pods, pod-minutes	0.567	0.267	-52.9%
Pending pods, peak	1	1	0
HPA replicas, mean	8.71	6.20	-28.8%
HPA replicas, peak	10	8	-20.0%
HPA shortfall, minutes	0.833	1.07	+28.0%

Metric (per unit of work, 2-minute blocks)	Kubernetes alone	Kubernetes + Omni-Compass	Change	95% CI	Verdict
Wh per core-hour, standby at 100 W (2-minute blocks from minute averages)	872	944	+8.2%	[-150, +292]	not significant
Node-hours in service per core-hour	7.8535	4.6976	-40.2 %	[-5.0982, -1.1072]	fewer (not energy while parked nodes stay powered)
Utilisation	0.0694	0.1224	+76.3 %	[+0.0217, +0.0830]	better
Pending-pod minutes per hour	1.7253	0.8333	-51.7 %	[-3.4961, +1.6667]	not significant
HPA shortfall minutes per hour	2.5458	3.3399	+31.2 %	[-3.3277, +5.0066]	not significant

Omni-Compass decisions: nodes per minute 5 4 3 then 3 for the remaining 17 minutes; HPA target 78-79% (Kubernetes alone: 50%). Node-pool resizes: 3 (three workers cordoned and drained, their pods rescheduled by Kubernetes). Kill switch: HPA target restored to 50, 6 of 6 workers back in service.

7.2 Live, side by side, every live muscle (run 36213152881 (2026-09-26 02:54-03:19 UTC))

Same two-cluster setup; the Omni arm drives HPA target, node pool, power cap (in-place pod CPU limits), heat (harness law on live power), security hold, rollout guard. Real response times: HTTP requests to the app timed every 5 s on both clusters. **This run made the application slower.**

Gauge	Kubernetes alone	Kubernetes + Omni-Compass	Change
Worker nodes in service, mean	6.00	3.20	-46.7%
Node-hours	2.01	1.07	-46.8%
Power (W), mean, standby counted	664	643	-3.2%
Power (W), peak	720	664	-7.7%
Energy (Wh), standby counted	222	215	-3.4%
CPU used (cores), mean	0.908	0.593	-34.6%
Utilisation	0.076	0.093	+22.6%
Energy per core-hour (Wh)	731	1,083	+48.1% worse
Pending pods, pod-minutes	1.70	0	-100%
HPA replicas, mean	9.09	5.00	-45.0%
HPA replicas, peak	10	5	-50.0%
Requests timed	955	850	-11.0%
Response time (ms), mean	260	415	+59.7% worse
Response time (ms), median	230	320	+39.2% worse
Response time (ms), 95th percentile	486	802	+64.8% worse
Response time (ms), 99th percentile	675	1,113	+64.7% worse
Failed requests (%)	0	0	0

Metric (2-minute blocks)	Kubernetes alone	+ Omni-Compass	Change	95% CI	Verdict
kWh per core-hour	0.7596	1.1107	+46.2%	[+0.1901, +0.5099]	worse
Node-hours per core-hour	6.8941	5.6440	-18.1%	[-2.7717, +0.4577]	not significant
Utilisation	0.0766	0.0964	+25.8%	[-0.0016, +0.0394]	not significant

Actions: hpa target writes 1, node pool resizes 3, power cap pod resizes 7, rollout actions 0, power cap per minute 0.695 every minute (the law floor), hpa target 77%, nodes 5 4 3 then 3. Kill switch: restored target 50; workers 6 of 6; pod CPU limits restored.

Diagnosis: The engine saw the cluster as idle (node utilisation about 7%) and had no response-time signal, so it capped the app pods at the law floor and packed replicas to a 77% target (5 instead of 10). The app pods were busy, so responses slowed. Fix wired: latency afferent (p95 over a declared 500 ms SLO enters as queue pressure) and a power-cap reflex (never below pod usage x 1.3); re-run pending.

7.3 Other live runs

Run	Setup	Result
live-kind, run 36205408388	1 node; baseline 10 min (Omni observing) then Omni target mode 10 min	0 writes while observing; kill switch restored 50; pending-pod minutes -48.6% (significant, but the baseline phase included warm-up); energy per core-hour no significant difference (one node cannot be parked)
live-kind-full, run 36205869009	1 control plane + 3 workers	node pool never resized: 3 of 3 every minute at about 8% utilisation, because the fleet law needs more than 3 nodes of slack; kill switch restored target and workers
live-kind-full, run 36207925928	1 control plane + 6 workers; sequential baseline then full engine	nodes in service 6 to 5 to 4 to 3; node-hours in service per core-hour -38.2%, utilisation +63.7% (significant); the reported kWh per core-hour -35.8% counted parked workers as powered off; with parked workers on standby the saving largely disappears; kill switch restored target 50 and 6 of 6 workers

7.4 Live runs after the first all-muscle run

Run	Change tested	Validity	Result
36214629046 (2026-09-26 03:23-03:48 UTC), commit c182a63	latency afferent (p95 over 500 ms SLO as queue pressure) and power-cap usage reflex	valid	p95 481 to 600 (+24.6% worse); p99 596 to 777 (+30.3% worse); energy 222 to 219 (-1.5%)
36216647786 (2026-09-26 04:02-04:28 UTC), commit 052b705	SLO reflex; eviction receipt fixed	INVALID as a test of the engine	p95 479 to 320 (-33.2%); p99 583 to 420 (-28.0%); energy 224 to 220 (-1.7%)
36218637030 (2026-09-26 04:42-05:08 UTC), commit 47d0353	controller fail-safe; benchmark prints controller log and rejects early stops	PARTIAL	p95 491 to 492 (+0.3% (tie)); p99 645 to 683 (+5.8% worse); energy 223 to 221 (-1.1%)

Run 3 is recorded as invalid: a refused Kubernetes call stopped the controller after 3 of 20 decisions and the cluster held Omni's last settings with no engine running. The fix is a fail-safe: a failed decision is skipped; three in a row restore native settings and stop the controller. In run 4 the fail-safe did exactly that at minute 16, and the audit showed why: the least-privilege role allowed writing a pod's CPU limit but not reading it first, which kubectld does. In runs 3 and 4 no power cap was ever applied. Over run 4's 16 governed minutes node-hours per unit of work fell 24% and utilisation rose 36% (both significant); median response time -10%, p95 a tie, p99 +6% (not significant). The permission is fixed (with a receipt) and a run in which the fail-safe fires is now rejected; the re-run is in progress.

8. Engineering verification

python verify.py runs 89 checks (0 failures in the recorded log results/VERIFY_LOG.txt), including:

- reference engine SHA-256
- core parity vs reference engine and 500 fixtures
- C++ core vs 500 fixtures (max abs 3.55e-15)
- C++ governor vs Python governor (throughput)
- C++ shield vs Python shield (enforce and violations)
- negative control: parity test detects a C++ shield that no longer blocks rollouts during a security block
- independent C++ HPA replica law vs fleet harness HPA
- soak test (throughput), 100,000,000 decisions, no failures, memory flat after warm-up (2899 ns per decision, RSS checkpoints [3840, 3840, 3840, 3840] kB)
- held-out seed 346410161: observe mode identical to native
- live controller against a fake cluster: observe writes nothing, target bounded, kill restores, node pool bounded and dry-run safe

Also tested: the live controller against a fake cluster (observe writes nothing; targets bounded; kill restores from the HPA annotation after a restart; node pool bounded and dry-run safe), the full-engine options (parked and control-plane nodes excluded; the kill switch restores the node pool exactly once) and pilot scoring (detects a real gain, reports no gain on identical clusters, detects a service regression).

9. Physics and models

- **Stack plant power:** each node draws idle 0.38 kW plus 1.12 kW x utilisation; a power cap throttles delivered capacity.
- **Heat:** thermal state follows a first-order lag toward $0.34 + 0.62 \times$ power stress (time constant about 7 steps); heat above 0.82 throttles capacity; 'over the heat limit' means thermal above 1.03.
- **Live kind cluster:** kind nodes have no power meter, so power is a declared model: 100 W idle + 150 W x CPU utilisation per worker in service, plus a standby power for each parked (cordoned and drained) worker. Standby defaults to the idle power (the worker stays powered and ready); lower values apply only to a declared sleep state, zero only to machines really powered off. The first live reports counted parked workers as zero; section 7.1 gives both. The same constants drive the governor's power sense and the energy score, so they cannot disagree. On real hardware this is replaced by metered power (RAPL, PDU or BMC).
- **Savings model** (results/SAVINGS.csv): a 1,000-node web cluster at 0.4 kW per node, PUE 1.4, \$0.12/kWh and 0.4 kg CO₂/kWh; reduction versus HPA 0.7 + Karpenter-lite of 14% to 20% (fleet plant) gives roughly 710 to 960 MWh, \$85,000 to \$115,000 and 280 to 380 t CO₂ per year.
- **Engine overhead:** about 2.9 microseconds per decision in C++, memory flat over 100 million decisions.

10. What is proven, what is simulated, what is not claimed

Claim	Status	Evidence
Engine equations are finite and deterministic; C++ equals Python	Proven	verify.py, 500 fixtures, max error 3.6e-15
Observe mode changes nothing	Proven (simulation and live)	bit-identical trajectories; 0 writes live
Kill switch restores native control	Proven (simulation and live)	HPA target 50 and all workers restored live
Omni-Compass acts on a real Kubernetes control plane (HPA target, node pool)	Proven live	section 7
Fewer nodes in service than Kubernetes with a fixed node pool, same load served	Measured live	section 7.1
Lower energy on the live cluster	Not shown while parked nodes stay on standby; -22% to -43% only if parked nodes sleep or power off	section 7.1
Better energy, health, recovery, coordination than Kubernetes (HPA + CA)	Pre-registered simulation	section 5
Better than Karpenter-lite on energy	Simulation	sections 6, and PlanetLab fleet plant
Better than upstream Karpenter or Cluster Autoscaler binaries, live	Not yet tested	section 12
Metered energy savings on physical servers	Not yet tested	power is modelled
GPU power-limit muscle saves 15-19% energy with <1% slower responses	Simulation calibrated to metered H100 data (MLPerf)	section 12
CPU frequency muscle	Simulation (uncalibrated): about -3.5% energy, -67% heat	section 12
Tuned laws remove wear and node-hour negatives (C-throughput)	Pre-registered amendment, new held-out data	section 13
Cooling, grid, memory, storage and the other open muscles	Not claimed	no plant or connector yet
Makes AI models aligned or trustworthy	Not claimed	value alignment is outside Omni-Compass

11. Limits and threats to validity

- Simulated studies use documented-behaviour replicas of Kubernetes controllers, not the upstream binaries; the Kubernetes reference omits Karpenter consolidation, VPA, scheduling constraints and disruption budgets.
- The live cluster is kind: nodes are containers on one CI machine; the two live arms ran on two machines at the same time, so machine-to-machine variation is part of the noise; each live arm is 20 minutes, one repetition.
- Live power and heat are modelled; parked kind workers are drained containers. If parked machines must stay on standby, parking reduces nodes in service but not energy; live energy savings then have to come from power caps, CPU power states and heat control, which are not yet wired live.
- The live native arm had no node autoscaler, so its node pool was always full; the fair live opponent is Karpenter or Cluster Autoscaler (section 16).
- The live significance for energy per core-hour with standby power is computed from minute averages (10 two-minute blocks), not from the 15-second capture.
- The fleet law releases a node only when the pool has more than three nodes of slack; a three-worker pool cannot scale down (observed live, reproduced offline). Small clusters need a pool-size-aware release band.

- Architecture C was measured in simulation only; the live C (Kubernetes' controllers parked, Omni-Compass as the only brain) is not built yet.
- Service quality differences in the live runs (pending pods, HPA shortfall) are not statistically significant at this run length.

12. GPU and CPU muscles: device plant calibrated to metered hardware (simulation)

The GPU power-limit and CPU frequency muscles cannot be actuated on CI machines (no GPU; the hypervisor hides RAPL and cpufreq). Their connectors are built (section 2.2) and this plant shows what they do. **GPU calibration:** the performance-versus-power-limit exponent is fitted to MLPerf Inference v4.0 results for an NVIDIA DGX-H100 (8 x H100-SXM, 700 W TDP), MaxQ (power-limited with `nvidia-smi -pl`, the same command the Omni-Compass GPU connector sends) versus MaxP, system power metered by a Yokogawa WT333E; Apache 2.0. Fitted exponent 0.353 to 0.490 (median 0.449); all three are run. **CPU:** a standard first-order model (dynamic power $\sim$ frequency cubed) against a schedutil-style governor; not calibrated to metered data yet. **S** is a fixed manual 70% cap, what an operator could do by hand. 24 scenarios, 8 load families; * = paired 95% interval excludes zero.

Vessel	Gauge	A. Native	S. Fixed 70% cap	B. + Omni-Compass	C. Omni-Compass direct
cpu_web	Energy (kWh)	115.169	103.342 (-10%*)	110.899 (-4%*)	111.238 (-3%*)
cpu_web	Energy per unit of work	0.672	0.614 (-9%*)	0.650 (-3%*)	0.652 (-3%*)
cpu_web	95th-pct response time (x baseline)	1.000	6.577 (+558%*)	1.040 (+4%*)	1.040 (+4%*)
cpu_web	Minutes over service target	0.000	70.458 (new*)	0.167 (new)	0.167 (new)
cpu_web	Minutes over heat limit	26.708	0.000 (-100%*)	8.875 (-67%*)	8.958 (-66%*)
cpu_web	Peak power (kW)	33.954	20.475 (-40%*)	34.074 (+0.35%)	34.074 (+0.35%)
cpu_web	Minutes near full power	8.000	0.000 (-100%*)	10.167 (+27%*)	10.167 (+27%*)
gpu_mlperf_median	Energy (kWh)	142.255	117.038 (-18%*)	119.254 (-16%*)	119.254 (-16%*)
gpu_mlperf_median	Energy per unit of work	0.824	0.679 (-18%*)	0.691 (-16%*)	0.691 (-16%*)
gpu_mlperf_median	95th-pct response time (x baseline)	1.000	1.521 (+52%*)	1.007 (+0.70%*)	1.007 (+0.70%*)
gpu_mlperf_median	Minutes over service target	0.000	7.125 (new*)	0.042 (new)	0.042 (new)
gpu_mlperf_median	Minutes over heat limit	22.458	0.000 (-100%*)	6.125 (-73%*)	6.125 (-73%*)
gpu_mlperf_median	Peak power (kW)	38.403	28.288 (-26%*)	37.089 (-3%*)	37.089 (-3%*)
gpu_mlperf_median	Minutes near full power	6.500	0.000 (-100%*)	6.208 (-4%)	6.208 (-4%)
gpu_mlperf_least_favourable	Energy (kWh)	142.255	118.430 (-17%*)	120.770 (-15%*)	120.770 (-15%*)
gpu_mlperf_least_favourable	Energy per unit of work	0.824	0.688 (-17%*)	0.699 (-15%*)	0.699 (-15%*)
gpu_mlperf_least_favourable	95th-pct response time (x baseline)	1.000	1.628 (+63%*)	1.008 (+0.80%*)	1.008 (+0.80%*)
gpu_mlperf_least_favourable	Minutes over service target	0.000	8.333 (new*)	0.042 (new)	0.042 (new)
gpu_mlperf_least_favourable	Minutes over heat limit	22.458	0.000 (-100%*)	6.708 (-70%*)	6.708 (-70%*)
gpu_mlperf_least_favourable	Peak power (kW)	38.403	28.345 (-26%*)	37.532 (-2%*)	37.532 (-2%*)

Vessel	Gauge	A. Native	S. Fixed 70% cap	B. + Omni-Compass	C. Omni-Compass direct
gpu_mlperf_least_favourable	Minutes near full power	6.500	0.000 (-100%*)	6.417 (-1%)	6.417 (-1%)
gpu_mlperf_most_favourable	Energy (kWh)	142.255	113.878 (-20%*)	115.645 (-19%*)	115.645 (-19%*)
gpu_mlperf_most_favourable	Energy per unit of work	0.824	0.661 (-20%*)	0.670 (-19%*)	0.670 (-19%*)
gpu_mlperf_most_favourable	95th-pct response time (x baseline)	1.000	1.300 (+30%*)	1.006 (+0.55%*)	1.006 (+0.55%*)
gpu_mlperf_most_favourable	Minutes over service target	0.000	5.375 (new*)	0.042 (new)	0.042 (new)
gpu_mlperf_most_favourable	Minutes over heat limit	22.458	0.000 (-100%*)	5.250 (-77%*)	5.250 (-77%*)
gpu_mlperf_most_favourable	Peak power (kW)	38.403	28.133 (-27%*)	36.174 (-6%*)	36.174 (-6%*)
gpu_mlperf_most_favourable	Minutes near full power	6.500	0.000 (-100%*)	5.833 (-10%*)	5.833 (-10%*)

Reading: on GPUs Omni-Compass saves 15-19% energy across the measured range with response time 0.6-0.8% slower; a fixed cap saves about the same energy but slows responses 30-63% and misses the service target. On CPUs the native governor already tracks demand, so frequency control alone saves about 3.5%; its main effect is heat (-67%). CPU frequency control is below the 10% bar and is a physical ceiling of that muscle, not a tuning gap.

13. Amendment: tuned laws, frozen, then tested on new held-out data (simulation)

Two law variants were tuned on the development seeds only (1000, 2000; tuning/SEARCH*.json), frozen with SHA-256 hashes in tuning/PREREGISTRATION_AMENDMENT_2026-09-26.json and pushed (commit b79d1dd, 03:17 UTC) before a single run on new held-out seeds 731001 and 731002 (2 x 500 scenarios). The frozen engine files are unchanged; the original pre-registered result (section 5) stays the primary result. B-wear: node release held longer. C-throughput: Omni-Compass direct under the throughput law with a wider release band and engine-gated cap. + better, ! worse (both seeds agree), blank not significant.

Gauge	A. Kubernetes alone	B_frozen	B_wear	C_frozen	C_throughput
energy_kwh	170.001	122.670 (-28%) +	123.552 (-27%) +	130.042 (-24%) +	123.353 (-27%) +
node_start_stop	6.119	10.070 (+65%) !	8.210 (+34%) !	11.985 (+96%) !	3.065 (-50%) +
machine_round_trips	1.031	1.724 (+67%) !	0.831 (-19%) +	2.711 (+163%) !	0.863 (-16%) +
scale_reversals	2.176	1.931 (-11%) +	1.805 (-17%) +	5.394 (+148%) !	1.838 (-16%) +
node_hours	136.812	141.488 (+3%) !	144.688 (+6%) !	152.143 (+11%) !	122.531 (-10%) +
idle_node_hours	46.662	52.152 (+12%) !	55.353 (+19%) !	61.218 (+31%) !	35.788 (-23%) +
mean_queue	438.529	246.346 (-44%) +	246.324 (-44%) +	1103.704 (+152%) !	1029.828 (+135%) !
p95_queue	1462.260	1005.114 (-31%) +	1005.044 (-31%) +	3185.168 (+118%) !	2684.025 (+84%) !
violation_backlog	0.030	0.016 (-45%) +	0.016 (-45%) +	0.085 (+182%) !	0.064 (+115%) !
recovery_minutes	61.470	16.805 (-73%) +	16.805 (-73%) +	28.925 (-53%) +	37.700 (-39%) +
power_cap_travel	0.249	0.894 (+259%) !	0.893 (+259%) !	0.307 (+23%) !	0.697 (+180%) !
invariant_violations	9.625	2.748 (-71%) +	2.747 (-71%) +	0.000 (-100%) +	0.000 (-100%) +

Gauge	A. Kubernetes alone	B_frozen	B_wear	C_frozen	C_throughput
sla_violation_total	0.171	0.076 (-55%) +	0.076 (-55%) +	0.094 (-45%) +	0.145 (-15%) +
pages	0.671	0.000 (-100%) +	0.000 (-100%) +	0.000 (-100%) +	0.000 (-100%) +
time_healthy	0.785	0.909 (+16%) +	0.909 (+16%) +	0.896 (+14%) +	0.837 (+7%) +
availability	0.999	1.000 (+0.09%) +	1.000 (+0.09%) +	0.997 (-0.23%) !	0.996 (-0.32%) !
recovered	0.907	0.997 (+10%) +	0.997 (+10%) +	0.971 (+7%) +	0.963 (+6%) +

Worse gauges: B_frozen 5, B_wear 4, C_frozen 10, C_throughput 5.

14. Response time: can one mode beat Kubernetes on every gauge? (simulation, development seeds)

The fleet-mode constants that drove the live runs were selected (before this work) by a rule that scored energy, finished work and backlog, but never waiting time. A response-time gauge was added to the fleet plant (fleet/sim_slo.py: M/M/c queueing delay per workload plus backlog drain; the frozen trace is reproduced exactly). Measured with it, fleet mode cut energy 31% on web services while p95 response time went from 132 ms to about 20 s. That is where the live slowdown came from: its HPA target floor (rho_min = 76.9%) packs pods too hot for latency-sensitive services.

A speed-first law (omnicompass/speed.py; the engine equations unchanged) separates the two jobs the frozen law gave one number: pods run at a latency-safe target while machines are sized to what the pods request. 400 configurations were searched against HPA 70% + Cluster Autoscaler with the rule that no gauge may be worse in any development scenario. None passed, for three measured reasons: (1) several gauges are already at their physical floor (100% work done, zero violations, batch p95 = pure service time), so a tie is the best any controller can do; (2) on the GPU vessel demand exceeds the hardware, so every arm is saturated; (3) energy, response time and machine churn trade two-of-three, because the autoscaler's idle slack is both where the energy is and what absorbs the 90-second boot delay. Examples, mean change vs Kubernetes: fleet mode energy -31% with far worse response time; speed #176 energy -1%, p99 -41%, churn 3.6x; speed #159 p99 -48%, churn -27%, energy +16%. Source: tuning/SPEED_FINDINGS_2026-09-26.md.

15. Named products: OpenShift, Google GKE, Azure AKS, IBM Turbonomic (simulation, development seeds)

Each product is emulated from its documented behaviour on the same fleet plant (not the vendors' binaries): OpenShift's documented ClusterAutoscaler example (threshold 0.4, unneeded 5 min, delay after add 10 min); GKE optimize-utilization (MostAllocated packing, more aggressive scale-down; declared as threshold 0.65, 2 min, because Google publishes no numbers); AKS node auto-provisioning (Karpenter, WhenEmptyOrUnderutilized, consolidateAfter 0 s); Turbonomic (container requests resized every 10 min to p99 per-pod usage, its default aggressiveness; nodes suspended toward 0.7 packing). Means over web services, 4 development seeds:

Gauge (web)	Kubernetes (GKE balanced)	OpenShift	GKE optimize	AKS NAP	Turbonomic	Omni fleet mode	Omni speed #159
energy (kWh)	12.19	12.53	10.4	9.779	10.08	8.392	15.9
p95 (ms)	131.6	131.6	132.2	132.7	125.1	2.034e+04	112.3
p99 (ms)	3,095	2,815	3,947	4,295	4.94e+04	7.634e+04	129.8
machine starts+stops	10.5	8.5	14.75	19	11.5	13.25	6.75
scale reversals	1.5	0.75	3.5	6	1.5	4.25	1

Every product sits on the same energy / response-time / churn triangle; none wins all three. Carpenter-style consolidation (AKS NAP) saves about 20% energy with a worse p99 tail and about twice the machine churn; Turbonomic's p99 resizing interacts with the HPA and inflates the tail; OpenShift's documented settings sit close to upstream. Omni's speed mode leads on response time and churn at an energy cost; its fleet mode leads on energy at a large response-time cost.

16. The problem-map muscles: all nine built and tested on held-out data (simulation)

Each open row of the industry problem map (section 18) is now a muscle in omnilab/, with the strongest native tool as opponent, Omni-Compass, and Omni-Compass with its engine equations not evolved (to show what the equations themselves add). Constants were chosen on development seeds 1-8, every file was frozen by SHA-256 (results/muscles/PREREGISTRATION.json), then 30 held-out seeds were run once. Verdicts from a paired bootstrap 95% interval.

Muscle (map row)	Strongest native opponent	Omni better	Omni worse	Tie / not significant
Right-sizing, HPA+VPA conflict (1, 2)	hpa_vpa	cpu core hours -3%; mem gib hours -11%; p99 ms -99%; replica reversals -84%; slo breach min -57%	p95 ms +15%; oom kills +531%	work done
Cold start (5)	keda	p95 ms -75%; p99 ms -58%; delayed 1s pct -71%; instance hours -8%	cold starts +736%	-
GPU packing (6)	binpack	energy kwh -2%; idle gpu hours powered -20%	wait mean min +16%; migrations 0 to 6.07	wait p95 min, frag blocked min, jobs done
Training power swings (7)	floor_safe	energy overhead pct -9%; throughput loss pct -44%; swing 1s mw -6%	max ramp mw s +127%	ramp violation s
GPU failures and stragglers (8)	detect	goodput pct +10%; lost gpu hours -20%; restarts -29%; straggler node hours -99%	-	checkpoint overhead pct
Cooling (10)	reset	pue -2%; cooling mwh -13%; inlet violation min -73%	-	max inlet c
LLM inference, KV cache (12)	keda	ttft p95 s -18%; slo breach pct -58%; gpu hours -19%; preemptions -56%	-	ttft p50 s, tpot p95 ms
Runaway AI agents (13)	static	rogue overspend usd -96%; time to contain min -99%; peak subagents -59%	false stops +357%; honest work pct -7%	forbidden executed
Energy attribution in VMs (11)	ratio	attr error pct -37%; worst vm error pct -34%	-	total error pct

Each entry is the plain change of Omni-Compass against the opponent (for example p99 -99% means the slow tail is 99% shorter; goodput +10% means 10% more useful training time). Worse entries are real costs, not rounding: right-sizing trades a slightly slower typical response (p95 +15%, both far under the 500 ms target) and a few more memory kills for a 99% shorter tail; cold start keeps fewer idle instances than KEDA's 5-minute cooldown, so more requests meet a cold start, but users wait far less because KEDA polls every 30 s; GPU packing powers idle GPUs off sooner and jobs wait a few seconds longer; containment throttles honest agents in their legitimate bursts (about 7% of their work) and stops about two honest agents a day, against a 96% cut in runaway spend and containment in minutes instead of hours. Against the other native arms (the default tools most teams run) the wins are larger; every comparison is in results/muscles/*_HELDOUT.json.

What the engine itself contributes. Across the nine muscles the full engine and the engine-not-evolved arm differ by 1.0 percentage points on average per gauge. The mapping from engine state to action carries most of each result; the evolved dynamics mainly smooth. In power smoothing an arm where the bath equation (7) alone sets the site's draw, with no ramp rule, cut the steepest ramp 63% but did not hold the grid's limit: the human sets the boundary, the engine operates inside it.

17. Every negative, its cause and its status

Negative	Where	Status	Cause	What would fix it
Node start/stop cycles (wear)	B, pre-registered	reduced (+65% to +34%), not removed	node release thresholds are fixed numbers	hybrid: Kubernetes serves the queue, Omni releases nodes with the C-throughput law (removed wear there: -50%)
Node start/stop, round trips, reversals	C, pre-registered	fixed in C-throughput (new held-out)	power-protect law released nodes too eagerly	adopted in C-throughput
Node-hours and idle node-hours	B and C, pre-registered	fixed in C-throughput (-10%, -23%); not in B	power capping trades lower watts for more servers on	hybrid as above; node-aware cap
Queue / backlog / availability -0.3%	C (both variants)	UNRESOLVED	not the replica law and not the sizing constants (both tested); likely the delayed observation or direct-mode proposals	trace one scenario step by step; candidate: feed the queue into the engine without the extra delay
Power-cap movement	B and C	inherent	moving the cap is how capping saves energy; freezing it removed most of the saving (tested: -28% to -2%/-7%)	none needed: electronic setting, no physical wear; reported
CPU frequency saving 3.5%	device plant	physical ceiling	the native Linux governor already follows demand	value is in heat (-67%) and in combining with power caps
GPU node on/off saving 3%	fleet plant	superseded	training nodes cannot be switched off	the GPU power-limit muscle (15-19%)
Live response time +65% (p95) and energy per unit of work +46% with every muscle	live kind, run 36213152881	reduced to p95 +25% (run 2), tie in run 4	no response-time afferent; cap held at the law floor; HPA target floor 76.9%	latency afferent, cap reflex, SLO reflex; speed-first law (section 14)
Fleet mode p95 20 s vs 132 ms on web services	fleet plant with response-time gauge	cause found; speed-first law built	fleet constants were selected without a response-time gauge; HPA target floor 76.9%	choose the mode per service: speed-first where latency matters
No mode better than Kubernetes on every gauge in every scenario	fleet plant, 400 configurations	not achievable as posed	gauges at physical floors; hardware-bound GPU vessel; energy / response time / churn trade two-of-three	anticipation (pre-adding machines before the daily rise) is the one untested mechanism that could break the trade
Controller stopped after one refused call; power cap never applied under least privilege	live runs 3 and 4	fixed	one exception ended the loop; the role lacked get on pods/resize	fail-safe restore after 3 failures; permission and receipt added; runs with a fail-safe rejected
Right-sizing p95 +15%, more memory kills than VPA	omnilab rightsize, held-out	open, small	Omni sizes CPU closer to demand; VPA's 8-hour p90 memory keeps more slack	larger memory margin (costs memory-hours)
More cold starts than KEDA	omnilab coldstart, held-out	trade-off	shorter keep-alive than KEDA's 5-minute cooldown	longer keep-alive where cold starts matter more than instance-hours
GPU jobs wait +16% (seconds) vs Volcano binpack	omnilab gpupack, held-out	trade-off	idle GPUs powered off sooner	longer power-off delay
Honest agents throttled (-7% work) and ~2 false stops a day vs static caps	omnilab containment, held-out	trade-off	throttling on the engine's integrated need catches legitimate bursts	per-agent declared burst budgets; human approval before stop
Engine dynamics add little beyond the mapping	all nine muscles	reported	the mapping from state to action carries the effect	wire the engine's control effort (equation 2) directly as the actuator command and test it

Negative	Where	Status	Cause	What would fix it
Live energy with parked servers on standby ~0%	live kind	open	parked servers still draw standby power	live power cap and CPU/GPU muscles; sleep states where hardware allows
Small clusters (<= 3 workers) never release	live kind	open	fleet law release band of 3 nodes	pool-size-aware release band (law change)
Pages in the 24-scenario replica	control-plane replica	open	longer queue triggers the page rule	same as queue

18. The industry problem map: what Omni-Compass is aimed at

One engine; the vessel (the plant it sits on) is the only thing that changes. Industry figures are approximate, from the public sources named, and are context, not results of this report. Status: **live** = measured on a real Kubernetes control plane; **sim** = demonstrated in this repository's simulations; **open** = mapped, connector not built.

Problem	Scale in the industry (approximate, source)	Best software today	Omni-Compass vessel and muscles	Gauge that shows it	Status
Data-centre electricity growth	about 415 TWh in 2024, about 1.5% of world electricity, projected near 945 TWh by 2030 (IEA, Energy and AI, 2025)	Karpenter, Cluster Autoscaler, CAST AI, Spot Ocean; Kepler for metering	compute vessel: nodes, HPA, power cap	energy, node-hours, idle node-hours	sim; live only where parked nodes can sleep or power off
Idle and over-provisioned capacity	Kubernetes clusters commonly run near 10-15% average CPU utilisation (CAST AI and Datadog industry reports); roughly a quarter to a third of cloud spend reported as waste (Flexera State of the Cloud)	VPA, Goldilocks, StormForge, Kubecost/OpenCost	compute vessel: nodes, HPA; requests and memory (omnilab/rightsize.py)	utilisation, node-hours per core-hour, core- and GiB-hours	live + sim
Controllers fighting each other	documented conflicts, e.g. HPA and VPA on the same CPU metric (Kubernetes documentation advises against it)	none: each tool decides alone	single authority over replicas and requests	contradictory commands, scale reversals, OOM kills	sim, held-out (section 16)
Outages and slow recovery	most significant outages cost over \$100,000 (Uptime Institute annual outage analysis)	Argo Rollouts, Flagger, SRE runbooks, AIOps (Dynatrace, Datadog)	compute vessel + deployments (partial)	time healthy, recovery time, SLA breaches	sim
On-call load and alert fatigue	widely reported burnout in SRE surveys	PagerDuty, alert tuning	all muscles: act before the page	pages, human interventions	sim
Heat and cooling limits	cooling is a large share of facility energy; average PUE about 1.5 (Uptime Institute survey)	DCIM (Schneider EcoStruxure), DeepMind cooling AI (reported about 40% less cooling energy)	heat (sensed), cooling plant (omnilab/cooling.py)	PUE, cooling energy, inlet violations	sim, held-out (section 16)
Site power and grid-connection limits	multi-year waits for new grid connections are widely reported	Meta Dynamo power capping, Intel RAPL	power cap (wired), batteries and demand response (open)	peak power, time over power limit	sim

Problem	Scale in the industry (approximate, source)	Best software today	Omni-Compass vessel and muscles	Gauge that shows it	Status
GPU energy and power limits	GPU fleets widely reported well below full utilisation; H100 TDP 700 W	NVIDIA DCGM and MIG, Run:ai, Kueue; manual MaxQ power limits	GPU power-limit muscle (hardware connector)	energy per unit of work, response time, heat	sim, calibrated to MLPerf metered H100 data: -15% to -19% energy, +0.6-0.8% response time
Batch deadlines and fair sharing		Kueue, Volcano, Slurm	batch muscle (live: admit held jobs); GPU packing (omnilab/gpupack.py)	queue wait, fragmentation	live wired + sim, held-out
Hardware wear	power cycling and churn shorten component life	none as a governed objective	nodes: start/stop cycles and reversals	machines started and stopped, round trips	mixed: better in the 24-scenario study, worse in the held-out study; tuning target
Carbon reporting and reduction	regulatory disclosure is expanding	Google carbon-aware computing, Kepler	carbon-aware placement (open)	kWh and CO2 per unit of work	sim (modelled)
Runaway AI agents and spend	~\$10,000 overnight examples (Dark Reading)	per-tool quotas and permissions	agent containment muscle (omnilab/containment.py)	rogue spend, time to contain, false stops	sim, held-out (section 16)

19. What comes next

- Live architecture C: park HPA, VPA, Cluster Autoscaler and Karpenter; Omni-Compass sets replicas, resources, placement, priorities and quotas directly; Kubernetes keeps execution and reflexes (restarts, rescheduling); the kill switch wakes the parked controllers.
- Live opponent at full strength: Karpenter (kwok provider) and Cluster Autoscaler in architecture A.
- 24 live scenarios (traffic, failures, power and heat limits, batch and AI, growth, mixed) with repetitions.
- More muscles two-way: CPU power states, memory, batch queues, network, security, then GPU and cooling on hardware.
- Metered power on physical machines.

20. Questions and answers

Does Omni-Compass replace Kubernetes? No. Kubernetes keeps running containers, placing pods, restarting failures and networking. Omni-Compass replaces the separate decision loops (how many replicas, how many nodes, what power) with one authority. In architecture C Kubernetes becomes one muscle.

What happens if Omni-Compass crashes or is switched off? The kill switch restores every setting it changed and returns control to Kubernetes' own controllers; this was exercised live and in simulation. A crashed controller writes nothing further.

Can it make things worse? Every action passes the shield first; it never goes below the capacity running and pending work needs, and never changes more than the step limit. In the held-out benchmark it had fewer safety violations than Kubernetes. Its real costs are listed in sections 1 and 5.1.

How fast does it decide, and what does it cost to run? One decision per 60 s on live Kubernetes (300 s in the replica), with a 15 s fast path that adds nodes for pending pods. The engine takes about 2.9 microseconds per decision.

Why does it save energy? In the pre-registered stack study mainly by power capping and load shaping (fewer minutes over the power limit, lower peak), while keeping slightly more machines on. On the live cluster it took machines out of service (6 to 3 workers) and packed replicas more densely; that saves energy only if the parked machines sleep or power off. With parked machines on standby the live saving was about zero, so live savings must come from power caps and CPU power states.

Does it slow applications down? It can, in the energy-first fleet mode: its HPA target floor packs pods too hot for latency-sensitive services (section 14). Live, the first all-muscle run was slower (p95 +65%); after the fixes run 4 was a tie at p95 and 10% faster at the median. For latency-sensitive services the speed-first mode is the right setting.

Can Omni-Compass control AI agents? It controls what an agent can touch, spend and do, and how fast; not what the model thinks. In the containment muscle (section 16) it cut runaway spend 96% and contained runaways in minutes instead of hours, with a least-privilege identity that blocks forbidden actions outright; the cost is some throttling of honest agents' bursts.

Does the engine hold everything in its basin by itself? Not against an outside limit it is not told. The bath equation alone smoothed training power ramps 63% but did not keep them under the grid's limit; with the human-set limit as the boundary it held it with zero violations. The human sets the boundaries; the engine operates inside them.

Is this tuned to the test? Parameters were selected on development seeds and frozen with hashes before the held-out seeds were run; `verify.py` fails if any frozen file changes.

How many scenarios and how certain? 1,000 pre-registered held-out scenarios, 24 control-plane scenarios, PlanetLab traces and live runs; 95% bootstrap intervals, two independent seeds must agree.

What is measured versus modelled? Live: node counts, replicas, pods, CPU, the controller's actions and the kill switch. Modelled: power and heat everywhere, and everything in the simulated studies.

Is the mathematics sound? The core is a closed six-state system integrated with RK4; the C++ and Python implementations agree to $3.6e-15$ on 500 reference trajectories; 100 million decisions ran without a non-finite value.

Who owns it and how can it be used? The Omni-Compass LLC. Free for evaluation, research and non-commercial use; commercial use requires a paid licence; protected by copyright and by patents and patent applications (see LICENSE and NOTICE).

What is not claimed? Superiority over upstream Karpenter or Cluster Autoscaler live, metered savings on physical hardware, GPU or facility control on real hardware, the vendor products' own binaries (they are emulated from documentation), and anything about AI value alignment.

How do I check it myself? Run the commands in section 21; the live runs are GitHub Actions workflows in the repository.

21. Reproduce

```

pip install -r requirements.txt
python verify.py                                # all checks, hashes, parity, soak
python benchmarks/stack_benchmark.py ...        # held-out stack benchmark (see HARNESS.md)
python -m k8s_controlplane.benchmark --scenarios 24 --seed 424242
python -m fleet.planetlab --dir fleet/traces/planetlab --scenarios 8 --out /tmp/pl
GitHub Actions: benchmark (live side by side), live-kind-full, live-kind
python tools/full_report.py ... && python pilot/bench_pdf.py docs/history/BENCHMARK_REPORT.md docs/histo
  ↳ ry/BENCHMARK_REPORT.pdf

```

22. Glossary

- **HPA:** Horizontal Pod Autoscaler: Kubernetes controller that sets replica counts from CPU utilisation versus a target.
- **Cluster Autoscaler, Karpenter:** Kubernetes add-ons that add and remove nodes.
- **Node, worker:** a machine (here a container in kind) that runs pods.
- **Cordon, drain:** mark a node unschedulable, then move its pods elsewhere.
- **kind:** Kubernetes in Docker: a real Kubernetes control plane whose nodes are containers.
- **Observe mode:** Omni-Compass computes and logs but writes nothing.
- **Invariant:** a safety rule the shield enforces before any action.
- **Paired bootstrap CI:** resampling the per-scenario differences to get a 95% interval for the mean difference.
- **Pre-registration:** freezing code and parameters, with hashes, before running the test data.

Evaluation and simulation use only. Copyright (c) 2026 The Omni-Compass LLC. Commercial use, commercialization or monetization of any part of Omni-Compass requires a signed, paid Omni-Compass Enterprise License from The Omni-Compass LLC. Patent applications, copyright registrations and trademark applications filed in the United States. See LICENSE and NOTICE at the root of this repository.

The Referee Report

Dated record, kept as written on 26 September 2026. Where a figure here differs from `docs/STATE_OF_PLAY.md`, the State of Play governs. The results below are simulations of platforms (evidence class S), not measurements of those products.

Referee report, 26 September 2026. Repository Omni-Compass/The-Omni-Compass-Control-Core-Engine, branch main, commit 096956a. Every number is generated by `python tools/abc_report.py` from result files in the repository. Each result states how it was obtained: **measured** on a live Kubernetes control plane, **simulated** in this repository's plants, or **emulated** (a competitor reproduced from its public documentation, not its binary).

Engine interpretation notice. The Omni-Compass engine is the completed core conveyance mechanism; its success criterion is internal dual-basin conveyance under native ignition dynamics. Everything measured here (the Kubernetes adapters, the closure governor, the vendor comparisons) is wiring around the engine, which is downstream engineering and not part of the base law. A loss in a cell below is a property of that wiring on that plant, not evidence that the base engine is incomplete. The engine source is byte-locked by SHA-256 and was not edited for any result in this report.

1. The answer

Three architectures were measured against seven platforms: Kubernetes, Red Hat OpenShift, Google GKE, Azure AKS (node auto-provisioning, which is Karpenter and also stands for Amazon EKS Auto Mode), IBM Turbonomic, CAST AI and Spot Ocean.

- **A. The platform alone.** Its own autoscalers decide.
- **B. Omni-Compass on top of the platform.** The platform's controllers keep running; the six-state engine governs on top.
- **C. Omni-Compass alone.** The engine is the single authority; the platform's managers do not decide.

B against A (held-out, 30 unseen scenarios per workload, settings frozen by SHA-256 first): 363 of 364 gauge-by-platform-by-workload cells are equal or better; 1 is worse: batch jobs on CAST AI, scale reversals 7.7% worse. Batch jobs gain most: the slowest jobs finish much sooner on every platform (section 5.3). Where no intervention was free (web services, four clusters) Omni-Compass reproduces the platform exactly.

C against the seven platforms at once (held-out): 63 cells where some platform is better (section 6). Part of that gap is not closable by any controller: a perfect controller that knows the future cannot have both the tightest packer's machine-hours and the laziest autoscaler's machine churn (section 6.2).

C with the manuscript's closure law and muscle tone (held-out, 30 unseen scenarios per workload, settings frozen first): 41 of 364 cells where some platform is better. Against Kubernetes alone: web services energy +20%, slowest responses (p99) +60%, average response +20%; four clusters, one site energy +22%, slowest responses

(p99) -19%, average response +11%; batch jobs energy +6%, slowest responses (p99) +6%, average response -6%; GPU training energy +1%, slowest responses (p99) +0%, average response +0% (positive = better; section 6.4).

Live Kubernetes (measured, run 36221702999, every muscle on, B architecture): worker machines 6.00 to 3.19, energy -3.2% with parked machines on standby, response time p95 491 to 383 ms and p99 613 to 491 ms, 0 failed requests; one significant negative (energy per CPU core-hour, section 7).

2. The rule, and how a cell is judged

- A **cell** is one gauge, one competitor, one workload. There are 13 gauges: energy, machine-hours, response time (p95, p99, mean), time over the backlog, power and heat limits, machine starts+stops, scale reversals, pod changes, work completed and time healthy.
- Omni-Compass **loses** a cell when it is worse by more than 0.5% (0.5 events for counts) **and** the paired bootstrap 95% interval of the per-scenario difference excludes zero. Otherwise the cell is equal (identical, or not significantly different) or better.
- Settings are chosen on **development** scenarios only, frozen with SHA-256 hashes (tuning/B_PREREGISTRATION*.json, results/muscles/PREREGISTRATION.json), then run once on **held-out** scenarios never used for tuning.
- Scenario generator: fleet/harness.py (15-second resolution, 6 hours; boot delay, power model, thermal model, site power limit identical for every arm). Response time: M/M/c queueing per workload plus backlog drain (fleet/sim_slo.py).

3. The opponents, and how each is reproduced

Platform	What decides	Reproduced as (parameters)
Kubernetes	HPA + Cluster Autoscaler	HPA 70%; remove a node after 10 min below 50% request use, not within 10 min of an add
OpenShift	ClusterAutoscaler resource	documented example: threshold 0.4, unneeded 5 min, delay after add 10 min
GKE (optimize-utilization)	MostAllocated packing, aggressive scale-down	declared: threshold 0.65, 2 min (Google publishes no numbers)
AKS NAP (= Karpenter, EKS Auto Mode)	consolidation WhenEmptyOrUnderutilized	consolidateAfter 0 s; remove when the rest hold all requests at 90% packing
IBM Turbonomic	market-based resize + suspend	requests resized every 10 min to per-pod p99 usage; nodes suspended toward 70% packing
CAST AI	Evictor bin-packing	every 60 s, drain a node older than 5 min when its pods fit elsewhere
Spot Ocean	headroom autoscaler	5% automatic headroom kept; least-used node removed when pods fit elsewhere

These are emulations of documented behaviour on the same plant, not the vendors' software; each vendor's full product has more features than its autoscaling core. Every platform runs its pods at the same HPA target (70%), so pod behaviour is identical and differences come from how machines are managed.

4. What Omni-Compass is, mechanically

A six-state dynamical system, $x = (E, U, I_U, S, B, B_{\text{dot}})$: energy/excitation, order (a double-well with a healthy basin at $U = +1$), integrated unmet need, stress (a cubic potential), and a damped second-order bath. Equations (1)-(8) and every term's role are in docs/MECHANISM_OF_ACTION.md. Telemetry is blended into the state (weight 0.339) and the equations are integrated with RK4 each decision.

Mode (Jacobian eigenvalue at the operating point)	Time constant (decisions)	Period (decisions)
-7.90+0.00i	1.3	
-3.10+0.00i	3.2	
-0.72+0.00i	13.9	
-0.20+0.98i	50.0	64
-0.20-0.98i	50.0	64
-0.17+0.00i	59.8	

All modes are stable: the operating point ($U = 0.976$) sits inside the healthy basin. The equations and the observation blend each move the state about half of every step. Compute: 95 microseconds per decision in Python, about 3 in the C++ twin.

5. B: Omni-Compass on top of each platform, against the same platform alone (held-out)

Source tuning/B_LEAGUE_HELDOUT2.json. Omni-Compass keeps the platform running and intervenes only where the engine sees it will pay: it undoes a machine removal that is about to be reversed (while requests are rising or the engine has not converged), and adds early the machine the platform would add only after pods go pending. With both off it reproduces the platform exactly (checked on every platform and workload); every gain below is on top of that floor.

5.1 Web services

Setting: pass-through. Every intervention tried here cost some gauge on some platform, so Omni-Compass leaves these decisions to the platform and the result is identical to the platform alone.

Platform	Better (significant)	Worse beyond tolerance (LOSS)	Worse within 0.5% tolerance	Equal
Kubernetes	-	-	-	13 of 13
OpenShift	-	-	-	13 of 13
GKE (optimize)	-	-	-	13 of 13
AKS NAP (Karpenter)	-	-	-	13 of 13
Turbonomic	-	-	-	13 of 13
CAST AI	-	-	-	13 of 13
Spot Ocean	-	-	-	13 of 13

5.2 Four clusters, one site

Setting: pass-through. Every intervention tried here cost some gauge on some platform, so Omni-Compass leaves these decisions to the platform and the result is identical to the platform alone.

Platform	Better (significant)	Worse beyond tolerance (LOSS)	Worse within 0.5% tolerance	Equal
Kubernetes	-	-	-	13 of 13

Platform	Better (significant)	Worse beyond tolerance (LOSS)	Worse within 0.5% tolerance	Equal
OpenShift	-	-	-	13 of 13
GKE (optimize)	-	-	-	13 of 13
AKS NAP (Karpenter)	-	-	-	13 of 13
Turbonomic	-	-	-	13 of 13
CAST AI	-	-	-	13 of 13
Spot Ocean	-	-	-	13 of 13

5.3 Batch jobs

Setting per platform (chosen on development scenarios): Kubernetes: early add (lead 12, trend over 16), reversal veto; OpenShift: early add (lead 12, trend over 16), reversal veto; GKE (optimize): early add (lead 6, trend over 8), reversal veto; AKS NAP (Karpenter): early add (lead 3, trend over 16), reversal veto; Turbonomic: early add (lead 6, trend over 4), reversal veto; CAST AI: early add (lead 12, trend over 16), reversal veto; Spot Ocean: early add (lead 3, trend over 16), reversal veto.

Platform	Better (significant)	Worse beyond tolerance (LOSS)	Worse within 0.5% tolerance	Equal
Kubernetes	response p99 -91%; response mean -23%; machine starts+stops -16%; scale reversals -70%	-	-	9 of 13
OpenShift	response p99 -91%; response mean -25%; machine starts+stops -23%; scale reversals -87%	-	-	9 of 13
GKE (optimize)	response p99 -91%; response mean -27%; machine starts+stops -22%	-	-	10 of 13
AKS NAP (Karpenter)	response p95 -57%; response p99 -24%; response mean -19%; time over backlog limit -47%; machine starts+stops -13%	-	energy +0.2%; machine-hours +0.5%	6 of 13
Turbonomic	response p99 -95%; response mean -34%; machine starts+stops -16%	-	-	10 of 13
CAST AI	response p95 -47%; response p99 -25%; response mean -32%; time over backlog limit -100%; machine starts+stops -9%	scale reversals +7.7%	-	7 of 13
Spot Ocean	response p95 -12%; response p99 -9%; response mean -8%; time over backlog limit -100%; machine starts+stops -8%; scale reversals -8%	-	energy +0.1%; machine-hours +0.3%	5 of 13

5.4 Gpu training

Setting per platform (chosen on development scenarios): Kubernetes: early add (lead 3, trend over 8), reversal veto; OpenShift: early add (lead 3, trend over 16); GKE (optimize): pass-through; AKS NAP (Karpenter): pass-through; Turbonomic: early add (lead 3, trend over 8); CAST AI: pass-through; Spot Ocean: pass-through.

Platform	Better (significant)	Worse beyond tolerance (LOSS)	Worse within 0.5% tolerance	Equal
Kubernetes	-	-	-	13 of 13
OpenShift	-	-	time over heat limit +0.1%	12 of 13
GKE (optimize)	-	-	-	13 of 13

Platform	Better (significant)	Worse beyond tolerance (LOSS)	Worse within 0.5% tolerance	Equal
AKS NAP (Karpenter)	-	-	-	13 of 13
Turbonomic	-	-	energy +0.1%; time over heat limit +0.2%; machine-hours +0.4%	10 of 13
CAST AI	-	-	-	13 of 13
Spot Ocean	-	-	-	13 of 13

Percentages are the plain change of Omni-Compass on top against the platform alone: for response times, starts+stops, reversals, energy and machine-hours negative is better; for work completed and time healthy positive is better. A LOSS is worse beyond the 0.5% tolerance with a 95% interval excluding zero (section 2).

6. C: Omni-Compass alone, against all seven platforms at once

Source tuning/LEAGUE_HELDOUT.json. One Omni-Compass configuration per workload type (chosen on development scenarios): web: speed159, multi: speed159, batch: speed112, gpu: omni_fleet.

Workload	Gauge	Platforms better than Omni-Compass	Omni worse by (range)
batch	energy	GKE (optimize), AKS NAP (Karpenter), Turbonomic, CAST AI, Spot Ocean	1.7% to 6.7%
batch	machine-hours	GKE (optimize), AKS NAP (Karpenter), Turbonomic, CAST AI, Spot Ocean	4.4% to 18.0%
batch	scale reversals	Kubernetes, OpenShift	266.7% to 279.3%
batch	machine starts+stops	Kubernetes, OpenShift	54.3% to 64.9%
gpu	machine-hours	GKE (optimize), AKS NAP (Karpenter), Turbonomic, CAST AI, Spot Ocean	5.5% to 14.4%
gpu	scale reversals	Kubernetes, OpenShift	85.6% to 87.6%
gpu	machine starts+stops	Kubernetes, OpenShift	37.6% to 49.8%
multi	energy	Kubernetes, OpenShift, GKE (optimize), AKS NAP (Karpenter), Turbonomic, CAST AI, Spot Ocean	23.9% to 64.7%
multi	machine-hours	Kubernetes, OpenShift, GKE (optimize), AKS NAP (Karpenter), Turbonomic, CAST AI, Spot Ocean	36.7% to 119.0%
multi	pod changes	Kubernetes, OpenShift, GKE (optimize), AKS NAP (Karpenter), CAST AI, Spot Ocean	8.2% to 10.3%
web	energy	Kubernetes, OpenShift, GKE (optimize), AKS NAP (Karpenter), Turbonomic, CAST AI, Spot Ocean	23.9% to 63.5%
web	machine-hours	Kubernetes, OpenShift, GKE (optimize), AKS NAP (Karpenter), Turbonomic, CAST AI, Spot Ocean	36.8% to 116.3%
web	pod changes	Kubernetes, OpenShift, GKE (optimize), AKS NAP (Karpenter), CAST AI, Spot Ocean	7.8% to 10.8%

6.1 Where C wins

- Web services: response p95 110.1 ms vs best platform 124.9 ms; response p99 122.8 ms vs best platform 546 ms; response mean 145.3 ms vs best platform 216.2 ms.
- Four clusters, one site: response p95 112.1 ms vs best platform 126.9 ms; response p99 137.8 ms vs best platform 262.9 ms; response mean 145.6 ms vs best platform 217.2 ms.

- Batch jobs: response p99 337.4 ms vs best platform 1277 ms; response mean 115.2 ms vs best platform 135.6 ms.

6.2 The physical limit: what no controller can do

For each scenario, dynamic programming with perfect knowledge of the future (boot delay ignored, which only helps it) finds the fewest machine starts+stops that can hold every pod at a given machine-hour budget (tuning/bound.py).

Workload	Scenario	Tightest platform machine-hours	Fewest starts+stops of any platform	Perfect controller at those machine-hours
web	101	18.4	9	11
web	102	18.3	9	9
web	103	31.4	4	26
web	104	22.0	12	24
multi	101	74.9	36	46
multi	102	112.7	26	100
multi	103	122.5	17	114
multi	104	93.4	35	91
batch	101	108.0	24	67
batch	102	56.8	7	33
batch	103	66.9	12	44
batch	104	93.4	15	58
gpu	101	66.4	12	28
gpu	102	57.5	16	36
gpu	103	51.5	19	30
gpu	104	57.0	20	26

In 15 of 16 scenarios even a perfect controller needs more starts+stops than the laziest platform to reach the tightest platform's machine-hours. Those two platforms sit at opposite ends of one physical trade; beating both at once is impossible for any software. In C, a loss on one side of that pair is the price of a win on the other.

6.3 Letting the equations drive directly

A governor in which the equations run closed-loop and their own quantities (control effort, unmet need I_U , the bath's rate) drive every muscle (omnicompass/mathdrive.py) was tested on 96 settings per workload: 66 losing cells, the same trade-off as the rule-based wiring. The engine clock alone moved batch p99 from 3.9 s to 0.1 s at equal energy.

6.4 Strict C: the manuscript's closure law drives the machines and the pods

omnicompass/closure.py implements, as the only authority over machines, the laws of the owner's manuscript: the forward projection of Section 5c, the master closure law of Chapter 20 ($F = G_0 + G_c$: no correction inside the admissible domain, an inward correction sized to restore the margin when the projected state approaches its boundary, so that $G \cdot n \leq 0$ on the boundary), the turning point of Chapters 29-30 (a machine is released only after the peak) and the dual-bath exchange of Chapter 31 (level and rate of demand). A release must also stay inside the calm set for a dwell (the resource-aware envelope, Proposition 2). Kubernetes' HPA and node managers are off; Kubernetes only schedules. Settings chosen on development scenarios, frozen by SHA-256 (56be5f50b8d2631c), then run once on 30 held-out scenarios per workload.

Workload	Gauge	Kubernetes	OpenShift	GKE (optimize)	AKS NAP (Karpenter)	Turbomicon	CAST AI	Spot Ocean	Omni-Compass (closure law)
web	energy	14.11	14.75	11.8	11.22	11.75	11.03	11.2	11.26
web	machine-hours	37.61	40.27	27.99	25.56	27.78	24.77	25.52	24.26
web	response p95	122.9	122.8	123.4	123.6	138.2	124.1	123.8	140.4
web	response p99	806.6	648.3	1345	1428	3251	2500	2023	318.9
web	response mean	206.4	205.3	212.5	215.7	251	236.1	226.9	164.5
web	machine starts+stops	8.733	6.5	16.47	19.67	12.73	26.57	27.6	11.6
web	work completed	1	1	1	1	1	1	1	1
multi	energy	54.24	56.4	45.04	42.81	44.86	42.16	42.85	42.43
multi	machine-hours	145.7	154.7	107.4	98.08	106.6	95.37	98.24	90.42
multi	response p95	124.9	124.8	125.6	125.9	119	126.5	126.1	148.9
multi	response p99	202.2	168.7	641.9	991.4	2.095e+04	2068	1531	240.3
multi	response mean	220.2	218.4	236.9	243.1	784.9	261.8	253.7	196.8
multi	machine starts+stops	36.77	28	66.63	79.7	53.8	102.4	108.1	45.67
multi	work completed	1	1	1	1	1	1	1	1
batch	energy	65.45	65.68	63.28	60.11	62.73	60.65	60	61.65
batch	machine-hours	94.71	95.49	87.48	76.91	85.67	78.71	76.55	78.5
batch	response p95	100	100	100	861.9	100	1840	220.1	100
batch	response p99	1072	1291	2086	3880	2242	4974	2491	1012
batch	response mean	133.1	138	158.8	257.2	170.6	360	184.2	141.3
batch	machine starts+stops	12.93	14.13	29	92.33	35.03	160.6	141.2	17.2
batch	work completed	1	1	1	1	1	1	1	1
gpu	energy	201.4	201.6	199.5	197.2	198.9	197.3	197	198.9
gpu	machine-hours	60.77	61.07	58.15	54.89	57.25	55.06	54.62	55.29
gpu	response p95	2.288e+06	2.286e+06	2.288e+06	2.29e+06	2.288e+06	2.292e+06	2.29e+06	2.287e+06
gpu	response p99	2.397e+06	2.396e+06	2.398e+06	2.4e+06	2.398e+06	2.402e+06	2.401e+06	2.396e+06
gpu	response mean	8.886e+05	8.875e+05	8.894e+05	8.904e+05	8.886e+05	8.922e+05	8.911e+05	8.879e+05
gpu	machine starts+stops	16.5	16.6	33.5	50.63	30.57	53.93	55.63	15.27
gpu	work completed	0.9027	0.9028	0.9027	0.9026	0.9028	0.9025	0.9026	0.9028

Losing cells (worse than that platform beyond tolerance): web services 10 of 91; four clusters, one site 12 of 91; batch jobs 7 of 91; GPU training 12 of 91.

- Web services: response p95 vs Kubernetes 14.2% worse; machine starts+stops vs Kubernetes 32.8% worse; response p95 vs OpenShift 14.3% worse; machine starts+stops vs OpenShift 78.5% worse; scale reversals vs OpenShift 192.9% worse; response p95 vs GKE (optimize) 13.8% worse; response p95 vs AKS NAP (Karpenter) 13.6% worse; energy vs CAST AI 2.1% worse; response p95 vs CAST AI 13.2% worse; response p95 vs Spot Ocean 13.4% worse.
- Four clusters, one site: response p95 vs Kubernetes 19.2% worse; machine starts+stops vs Kubernetes 24.2% worse; response p95 vs OpenShift 19.4% worse; response p99 vs OpenShift 42.4% worse; machine starts+stops vs OpenShift 63.1% worse; scale reversals vs OpenShift 153.2% worse; response p95 vs GKE

(optimize) 18.5% worse; response p95 vs AKS NAP (Karpenter) 18.3% worse; response p95 vs Turbonomic 25.2% worse; energy vs CAST AI 0.6% worse; response p95 vs CAST AI 17.8% worse; response p95 vs Spot Ocean 18.1% worse.

- Batch jobs: machine starts+stops vs Kubernetes 33.0% worse; machine starts+stops vs OpenShift 21.7% worse; energy vs AKS NAP (Karpenter) 2.6% worse; machine-hours vs AKS NAP (Karpenter) 2.1% worse; energy vs CAST AI 1.7% worse; energy vs Spot Ocean 2.8% worse; machine-hours vs Spot Ocean 2.6% worse.
- Gpu training: energy vs AKS NAP (Karpenter) 0.9% worse; time over power limit vs AKS NAP (Karpenter) 0.7% worse; time over heat limit vs AKS NAP (Karpenter) 0.8% worse; machine-hours vs AKS NAP (Karpenter) 0.7% worse; energy vs CAST AI 0.8% worse; time over power limit vs CAST AI 1.5% worse; time over heat limit vs CAST AI 1.9% worse; energy vs Spot Ocean 1.0% worse; time over power limit vs Spot Ocean 1.8% worse; time over heat limit vs Spot Ocean 1.6% worse; machine-hours vs Spot Ocean 1.2% worse; time healthy vs Spot Ocean 1.5% worse.

6A. Runtime Benchmark Protocol: faults at five stress levels

The owner's *OmniCompass Runtime Benchmark Protocol* (stage 1, Python runtime) run on this fleet plant with every architecture: 100 runs per stress level per workload, seeds 900001-900100 (never used for tuning). Level L injects L faults (machines dying, load spikes, services crash-looping, a noisy neighbour), identical for every system (`tools/protocol_bench.py`). A run is **conveyed** when every fault is recovered within 20 minutes and every cluster is inside the admissible basin for the final 10 minutes (worst response ≤ 1 s, queue < 0.28 , pending pods $\leq 5\%$, power and heat inside limits). A **page** is an out-of-basin episode of 5 minutes or more, the usual alert rule: each would call a human. Nobody intervenes in the simulation.

Web services

Column	System	Conveyed	Pages	Recovery (min)	Minutes outside basin	Timeouts > 2 s	Energy (kWh)
A	AKS NAP (Karpenter)	233 of 500	391	6.6	32.8	4.52%	11.90
B	AKS NAP (Karpenter)	233 of 500	391	6.6	32.8	4.52%	11.90
A	CAST AI	231 of 500	404	6.6	34.1	4.68%	11.75
B	CAST AI	231 of 500	404	6.6	34.1	4.68%	11.75
A	GKE (optimize)	227 of 500	366	6.9	32.4	4.25%	12.76
B	GKE (optimize)	227 of 500	366	6.9	32.4	4.25%	12.76
A	Kubernetes	218 of 500	338	8.3	33.3	3.97%	15.54
B	Kubernetes	218 of 500	338	8.3	33.3	3.97%	15.54
A	OpenShift	220 of 500	328	8.0	32.2	3.82%	15.52
B	OpenShift	220 of 500	328	8.0	32.2	3.82%	15.52
A	Spot Ocean	234 of 500	392	6.6	32.8	4.56%	11.94
B	Spot Ocean	234 of 500	392	6.6	32.8	4.56%	11.94
A	Turbonomic	169 of 500	552	14.4	53.3	5.96%	12.43
B	Turbonomic	169 of 500	552	14.4	53.3	5.96%	12.43
C	C-hpa	202 of 500	353	11.7	35.8	3.81%	20.69
C	C-strict	238 of 500	283	6.0	26.2	3.45%	13.02

In the protocol's required form: Omni-Compass alone (closure law) conveyed 238 out of 500 runs; the best platform alone (Spot Ocean) conveyed 234 out of 500. Omni-Compass would have paged a human 283 times, that platform 392 times. Omni-Compass average recovery time was 6.0 min, that platform's 6.6 min.

Four clusters, one site

Column	System	Conveyed	Pages	Recovery (min)	Minutes outside basin	Timeouts > 2 s	Energy (kWh)
A	AKS NAP (Karpenter)	49 of 500	534	5.6	69.5	1.71%	45.25
B	AKS NAP (Karpenter)	49 of 500	534	5.6	69.5	1.71%	45.25
A	CAST AI	44 of 500	608	5.7	74.2	1.81%	44.54
B	CAST AI	44 of 500	608	5.7	74.2	1.81%	44.54
A	GKE (optimize)	49 of 500	455	5.7	65.6	1.63%	47.96
B	GKE (optimize)	49 of 500	455	5.7	65.6	1.63%	47.96
A	Kubernetes	44 of 500	330	6.5	59.2	1.44%	57.67
B	Kubernetes	44 of 500	330	6.5	59.2	1.44%	57.67
A	OpenShift	43 of 500	310	6.4	57.6	1.38%	59.47
B	OpenShift	43 of 500	310	6.4	57.6	1.38%	59.47
A	Spot Ocean	44 of 500	508	5.6	69.1	1.73%	45.30
B	Spot Ocean	44 of 500	508	5.6	69.1	1.73%	45.30
A	Turbonomic	26 of 500	867	15.0	97.6	2.49%	47.57
B	Turbonomic	26 of 500	867	15.0	97.6	2.49%	47.57
C	C-hpa	64 of 500	328	8.8	45.8	1.07%	74.42
C	C-strict	45 of 500	231	4.3	52.6	1.33%	45.82

In the protocol's required form: Omni-Compass alone (closure law) conveyed 45 out of 500 runs; the best platform alone (GKE (optimize)) conveyed 49 out of 500. Omni-Compass would have paged a human 231 times, that platform 455 times. Omni-Compass average recovery time was 4.3 min, that platform's 5.7 min.

Batch jobs

Column	System	Conveyed	Pages	Recovery (min)	Minutes outside basin	Timeouts > 2 s	Energy (kWh)
A	AKS NAP (Karpenter)	314 of 500	234	4.5	31.4	9.01%	70.02
B	AKS NAP (Karpenter)	338 of 500	235	4.5	25.1	8.09%	70.95
A	CAST AI	161 of 500	237	4.4	39.9	10.25%	72.15
B	CAST AI	224 of 500	233	4.4	26.3	8.39%	73.91
A	GKE (optimize)	409 of 500	235	4.0	19.0	7.07%	77.00
B	GKE (optimize)	425 of 500	228	4.0	13.9	6.38%	79.12
A	Kubernetes	424 of 500	228	3.7	15.4	6.28%	85.35
B	Kubernetes	425 of 500	234	3.8	12.2	5.77%	87.43
A	OpenShift	422 of 500	232	3.8	16.0	6.36%	85.36
B	OpenShift	425 of 500	232	3.8	12.2	5.82%	86.89
A	Spot Ocean	360 of 500	231	4.4	23.1	7.77%	71.08
B	Spot Ocean	360 of 500	234	4.4	20.6	7.56%	71.79
A	Turbonomic	407 of 500	232	4.2	21.0	7.45%	74.08
B	Turbonomic	426 of 500	233	4.1	14.2	6.62%	75.25
C	C-hpa	166 of 500	235	30.8	73.5	20.65%	85.38
C	C-strict	422 of 500	237	4.1	17.3	6.93%	74.13

In the protocol's required form: Omni-Compass alone (closure law) conveyed 422 out of 500 runs; the best platform alone (Kubernetes) conveyed 424 out of 500. Omni-Compass would have paged a human 237 times, that platform 228 times. Omni-Compass average recovery time was 4.1 min, that platform's 3.7 min.

Gpu training

Column	System	Conveyed	Pages	Recovery (min)	Minutes outside basin	Timeouts > 2 s	Energy (kWh)
A	AKS NAP (Karpenter)	13 of 500	802	95.4	227.2	69.96%	202.42
B	AKS NAP (Karpenter)	13 of 500	802	95.4	227.2	69.96%	202.42
A	CAST AI	13 of 500	925	95.1	226.5	70.61%	202.80
B	CAST AI	13 of 500	925	95.1	226.5	70.61%	202.80
A	GKE (optimize)	15 of 500	802	96.5	227.1	68.76%	205.17
B	GKE (optimize)	15 of 500	802	96.5	227.1	68.76%	205.17
A	Kubernetes	18 of 500	782	97.5	225.4	67.95%	206.83
B	Kubernetes	18 of 500	792	97.6	226.1	67.87%	207.04
A	OpenShift	18 of 500	788	97.6	226.3	67.91%	207.08
B	OpenShift	18 of 500	788	97.6	226.2	67.88%	207.13
A	Spot Ocean	13 of 500	830	94.9	226.0	70.66%	202.22
B	Spot Ocean	13 of 500	830	94.9	226.0	70.66%	202.22
A	Turbonomic	14 of 500	786	96.5	226.6	68.78%	204.30
B	Turbonomic	14 of 500	792	96.2	226.6	68.67%	204.40
C	C-hpa	0 of 500	1232	174.0	295.6	83.41%	198.48
C	C-strict	13 of 500	809	95.1	225.0	67.86%	204.08

In the protocol's required form: Omni-Compass alone (closure law) conveyed 13 out of 500 runs; the best platform alone (Kubernetes) conveyed 18 out of 500. Omni-Compass would have paged a human 809 times, that platform 782 times. Omni-Compass average recovery time was 95.1 min, that platform's 97.5 min.

6B. Second generation (held-out, settings frozen by SHA-256 before each run)

B with muscle tone and consolidation (tuning/b_tone.py, held-out seeds 700401-700430): machines a platform powers off are parked instead (alive, low power, instant wake), with a reserve sized by the law; Omni consolidates on top only where that loses nothing. Losing cells: 3 of 364.

B fix for the remaining pairs (tuning/b_fix.py, 30 development seeds, held-out 700701-700730): multi on aks_nap: 0 losing cells; batch on turbonomic: 0 losing cells; multi on cast_ai: 0 losing cells.

Whole body at four-cluster sites (tuning/site_league.py, held-out 701001-701030): traffic shift between a site's clusters (a cluster short of room runs its overflow on another cluster's powered spare cores, +5 ms per shifted request; assumes replicated services) and, for C, the closure law on the site total with one warm reserve.

Platform	B losing cells	C losing cells	C energy	C slowest responses	C average
Kubernetes	0	2	+16%	+56%	+30%
OpenShift	0	2	+19%	+46%	+29%
GKE (optimize)	0	0	+7%	+85%	+36%
AKS NAP (Karpenter)	0	0	+2%	+90%	+39%
Turbonomic	0	1	+7%	+99%	+75%
CAST AI	0	0	-0%	+94%	+42%
Spot Ocean	0	0	+2%	+91%	+40%

C gains are Omni against that platform (positive = Omni better).

6C. Live levers on real Kubernetes (measured)

```

Live levers on real Kubernetes (kind, 1 control plane + 6 workers), GitHub Actions run 36270589830, comm
↳ it 43f6747,
2026-09-26 20:45-20:48 UTC. Omni-Compass ran as its least-privilege service account (deploy/kind/rbac-om
↳ ni.yaml +
rbac-levers.yaml). Script: scripts/kind_levers.sh. Result: LIVE LEVERS: PASS, 19 of 19 checks, 0 lever e
↳ rrors.

Permission receipts (kubectl auth can-i as the Omni service account)
can: patch pods/resize default, patch deployment/idle-worker, patch deployments/scale idle-worker,
get configmap/omni-work, patch jobs default, create/delete resourcequotas agents, patch pods/r
↳ esize agents
cannot: delete pods agents, create pods agents, delete deployments agents, patch deployment/load-gener
↳ ator,
create resourcequotas default, get secrets (any namespace)

rightsize php-apache CPU requests: before [200m] after [50m]; after kill [200m]
coldstart idle-worker replicas: idle -&gt; 0; work waiting -&gt; 2; idle again -&gt; 0; after kill -&gt; 2
batch_pace train-a suspended: stress -&gt; true; calm -&gt; false; stress -&gt; true; after kill -&gt; false
contain agent quota present: over budget -&gt; 1; under budget -&gt; 0; over -&gt; 1; after kill -&gt; 0
agent CPU limits: before [400m 400m] contained [100m 100m] lifted [400m 400m] after kill [40
↳ 0m 400m]
cooling setpoints written to the stand-in building controller: 27.0 27.0 27.0 22.0 (kill restores 22
↳ .0)
records left after kill: none

Earlier runs of the same script (kept for the record): run 36268772674 failed (containment could not lif
↳ t: Kubernetes
refuses to raise a pod limit above a standing quota; one failing lever stopped the kill switch); run 362
↳ 69885518 printed
PASS but its checks were chained so failures were not counted (containment still not lifted); run 362702
↳ 02629 failed in
the script's quota probe once the quota was correctly deleted. Fixes: quota deleted before limits are re
↳ stored; each
lever and each kill-switch restore isolated; every check counted on its own line.
Not covered: the cooling lever writes to a stand-in controller (a CI runner has no chiller); CPU frequen
↳ cy and power caps
need hardware with cpufreq; traffic shift needs two real clusters behind one load balancer.

```

7. Live Kubernetes (measured)

Two identical kind clusters (1 control plane + 6 workers) at the same time, same load; one with Omni-Compass on top (B). Omni runs as a least-privilege service account whose permissions are proven with `kubectl auth can-i` receipts before each run.

Run	Change	Validity	p95 (ms)	p99 (ms)	Energy
36213152881	HPA target, node pool, power cap (in-place pod CPU limits), heat (harn	valid	486 to 802	675 to 1,113	222 to 215
36214629046	latency afferent (p95 over 500 ms SLO as queue pressure) and power-cap	valid	481 to 600	596 to 777	222 to 219
36216647786	SLO reflex; eviction receipt fixed	INVALID as a test of the engine	479 to 320	583 to 420	224 to 220
36218637030	controller fail-safe; benchmark prints controller log and rejects earl	PARTIAL	491 to 492	645 to 683	223 to 221
36220059046	get on pods/resize (power cap can apply)	INVALID for response time and energy	501 to 380	660 to 415	224 to 220
36221702999	probe tunnel restarts on both arms; power cap able to apply	VALID	491 to 383	613 to 491	224 to 216

Run 6 in full:

Gauge	Kubernetes alone	+ Omni-Compass	Change
Worker nodes in service, mean	6.00	3.19	-46.8%
Node-hours	2.02	1.07	-47.0%
Energy (Wh), standby counted	224	216	-3.2%
Power (W), peak	710	682	-3.8%
Energy per core-hour (Wh)	742	1,041	+40.2% worse (significant)
Utilisation	0.074	0.097	+30.2%
Pending pods, pod-minutes	0.283	0.533	+88% (n.s.)
HPA shortfall, minutes	0.817	1.33	+63% (n.s.)
Response time (ms), mean	264	165	-37.4%
Response time (ms), median	227	119	-47.4%
Response time (ms), 95th percentile	491	383	-22.1%
Response time (ms), 99th percentile	613	491	-19.9%
Failed requests (%)	0	0	0

First fully valid run with every live muscle on: fewer servers (-47%), less energy (-3.2%, standby counted), lower peak power, and faster responses at every percentile. The one significant negative is energy per core-hour (+48%): the power cap lowers the CPU the app consumes (-31%), and that CPU is the metric's denominator. Whether energy per request served also rose is not measured (the load generator's request count is not captured); a per-request energy gauge is added to the capture next.

8. The nine problem-map muscles (simulated, held-out, Omni-Compass direct against the strongest native tool)

Muscle	Opponent	Better	Worse
rightsize	hpa_vpa	cpu core hours -3%; mem gib hours -11%; p99 ms -99%; replica reversals -84%; slo breach min -57%	p95 ms +15%; oom kills +531%
coldstart	keda	p95 ms -75%; p99 ms -58%; delayed 1s pct -71%; instance hours -8%	cold starts +736%
gpupack	binpack	energy kwh -2%; idle gpu hours powered -20%	wait mean min +16%; migrations 0 to 6.07
powersmooth	floor_safe	energy overhead pct -9%; throughput loss pct -44%; swing 1s mw -6%	max ramp mw s +127%
health	detect	goodput pct +10%; lost gpu hours -20%; restarts -29%; straggler node hours -99%	-
cooling	reset	pue -2%; cooling mwh -13%; inlet violation min -73%	-
inference	keda	ttft p95 s -18%; slo breach pct -58%; gpu hours -19%; preemptions -56%	-
containment	static	rogue overspend usd -96%; time to contain min -99%; peak subagents -59%	false stops +357%; honest work pct -7%
vmenergy	ratio	attr error pct -37%; worst vm error pct -34%	-

9. Threats to validity

- Competitors are emulated from documentation on a shared plant; their production binaries may behave differently, and their full products include features (spot pricing, instance-type selection, rebalancing) that this plant does not model.
- Simulated plants are fluid models: pods pack perfectly, boot delay is fixed, power is idle + dynamic x utilisation. Real clusters fragment and boot times vary.
- Live runs use kind: nodes are containers on one CI machine; power is modelled, not metered; the two arms ran on two different CI machines. Six 20-minute runs are evidence of behaviour, not of production savings.
- Development and held-out scenarios come from the same generator; held-out protects against tuning to seeds, not against a generator that differs from real traffic. Real traces (PlanetLab) are covered in the main benchmark report.
- The loss rule uses a 0.5% tolerance and a 95% interval on 30 scenarios; small true differences can be missed.

10. Reproduce

```

pip install -r requirements.txt && python verify.py
python tuning/b_league.py dev # B vs A, choose settings (development seeds)
python tuning/b_league.py heldout tuning/B_SETTINGS_FROZEN.json # B vs A, held-out
python tuning/league.py heldout # C vs all platforms
python tuning/bound.py # perfect-foresight frontier
python tools/mechanism.py # engine modes and compute
python tuning/closure_search.py && python tuning/closure_search2.py && python tuning/closure_heldout
  ↳ # strict C, closure law
python tools/protocol_bench.py 100 # runtime protocol, faults at five stress level
  ↳ s
for m in rightsize coldstart gpupack powersmooth health cooling inference containment vmenergy; do pytho
  ↳ n -m omnilab.bench $m heldout; done
GitHub Actions workflow 'benchmark' (commit message tag [bench]): live A vs B on kind
python tools/abc_report.py && python pilot/bench_pdf.py docs/history/OMNICOSSPASS_ABC_REPORT.md docs/hist
  ↳ ory/OMNICOSSPASS_ABC_REPORT.pdf

```

Evaluation and simulation use only. Copyright (c) 2026 The Omni-Compass LLC. Commercial use, commercialization or monetization of any part of Omni-Compass requires a signed, paid Omni-Compass Enterprise License from The Omni-Compass LLC. Patent applications, copyright registrations and trademark applications filed in the United States. See LICENSE and NOTICE at the root of this repository.

CHAPTER 51

Comparison with Existing Controllers

People who run these systems often see "autoscaling" everywhere and assume it all does the same thing. It does not. This page sets out, from public descriptions, what each system controls, and where Omni-Compass sits among them.

How to read it. Borg and Twine are private to Google and Meta; Turbonomic is closed commercial software. None of them can be run here, so they are compared by what their makers have published (papers, documentation, product pages), never by a run. A mark of "—" means the capability is not in the system's public description, not that it is impossible for it. The only fair measured comparison is on the customer's own system: their stack as it runs today against the same stack with Omni-Compass on top, in paired runs (`docs/INTEGRATION_MANUAL.md`, section 6).

1. What each system controls

System (public description)	Pods / replicas	Machines	GPU power	CPU clock / power	Moves watts between CPU and GPU	Site power budget	One decision across all of these
Kubernetes HPA	✓	—	—	—	—	—	—
Kubernetes VPA	✓ (requests)	—	—	—	—	—	—
Cluster Autoscaler, Karpenter	—	✓	—	—	—	—	—
KEDA	✓ (to zero)	—	—	—	—	—	—
Red Hat OpenShift / OKD	✓ (the Kubernetes autoscalers)	✓ (machine autoscaling)	—	—	—	—	—
IBM Turbonomic	✓ (sizing, scaling)	✓ (placement, sizing)	—	—	—	—	partial (resource actions across layers)
Google Borg	✓ (job scheduling)	✓ (cluster placement)	—	—	—	Google has published separate power-capping work	—
Meta Twine	✓	✓	—	—	—	Meta has published separate power management (Dynamo)	—
NVIDIA DCGM	—	—	✓ (limits, monitoring)	—	—	—	—
Linux schedutil, RAPL	—	—	—	✓	—	—	—
AMD SmartShift, NVIDIA Dynamic Boost (laptops)	—	—	✓	✓	✓ (inside one laptop)	—	inside one device
Intel GEOPM (HPC)	—	—	partial	✓	partial (per job)	✓ (per job / cluster)	partial
Omni-Compass	✓ (on top of HPA; or deciding)	✓ (park and wake)	✓	✓	✓ (conveyance law; simulated)	✓ (simulated)	✓

2. How each one decides

	Typical rule	Omni-Compass
Pods	add replicas when average CPU passes a target	the same autoscaler keeps working; Omni-Compass sets its target from the whole system's state and raises floors ahead of bursts
Machines	add when pods cannot be placed, remove when underused	park and wake machines by one law, and only when every sense is live, pods are not rising and nothing waits
GPU	vendor default: full power, or a fixed cap set by hand	a cap from the engine, bounded by a floor and released at once when the card is busy or response time slips
CPU	the kernel's governor picks a clock from recent load	a ceiling from the same engine, inside the nervous system's envelope
Power	a fixed plan leaving room for every device at peak; a breaker or capper reacts when the site goes over	one conserved budget: watts move from organs holding surplus to organs in need, never created, never over the budget
Safety	per tool	one OFF switch returns every organ to its recorded setting; watch mode first; service first

3. What is different, in plain words

1. **One brain instead of many separate reflexes.** Today each tool watches its own gauge: the pod autoscaler watches CPU, the node autoscaler watches pending pods, the GPU runs at its default, the power plan is fixed on paper. Nobody decides for the whole. Omni-Compass reads the whole system into one state and tells each tool the envelope it may act in. The tools keep doing their own jobs.
2. **Power as a budget that moves.** Today a building keeps room for every device to hit its peak at once, so much of its power sits reserved and unused. Omni-Compass treats the building's watts as one conserved total and moves them to where the work is: from CPUs and quiet GPUs to busy GPUs, never over the limit (docs/CONVEYANCE_LAW.md).
3. **A proof, not a tuned dial.** The budget law is proved to conserve the budget and to converge, and those properties are checked on every build (tests/test_conveyance.py).
4. **Service guarded on every power move.** Power is only taken back while response time is inside the target, and given back at once when it is not.
5. **Reversible in one move.** Every setting is recorded before it is changed and restored by one switch.

4. What the numbers show today, and what kind they are

Against	Result	Kind
Kubernetes as it runs by default (HPA + fixed nodes)	p95 response time -37% to -55%, replicas -12% to -25%, machines in service -17% to -21%	measured on real Kubernetes (kind), 10 paired runs each, results/live/
A GPU at its vendor default	+1.3% to +5.1% work per energy, p95 within +10%	modelled card (MLPerf-calibrated), results/gpu/sim/after
Today's power practice (every GPU at one fixed cap that leaves room for every CPU at peak)	+1.4% to +5.7% work served, backlog -16% to -28%, never over the site budget	modelled, results/hardware/NODE_EXCHANGE_*.json
A reactive site capper alone	60-196 minutes over the site budget without Omni-Compass; 0 with it	modelled

Turbonomic, Borg and Twine cannot be measured here. Against a customer who runs one of them, the comparison is made on the customer's own system, with the same paired method.

5. Sources

- Kubernetes HPA, VPA, Cluster Autoscaler, Karpenter, KEDA: the projects' own documentation.
- Red Hat OpenShift / OKD: Red Hat's documentation of its autoscaling and machine management.
- IBM Turbonomic: IBM's product documentation.
- Borg: Verma et al., "Large-scale cluster management at Google with Borg", EuroSys 2015; Google's public Borg traces.
- Twine: Tang et al., "Twine: A Unified Cluster Management System for Shared Infrastructure", OSDI 2020.
- Meta Dynamo: Wu et al., "Dynamo: Facebook's Data Center-Wide Power Management System", ISCA 2016.
- NVIDIA DCGM and Dynamic Boost; AMD SmartShift; Intel GEOPM: the vendors' documentation.

Where a row above is wrong or out of date, correct it from the maker's own publication.

Evaluation and simulation use only. Copyright (c) 2026 The Omni-Compass LLC. Commercial use, commercialization or monetization of any part of Omni-Compass requires a signed, paid Omni-Compass

Enterprise License from The Omni-Compass LLC. Patent applications, copyright registrations and trademark applications filed in the United States. See LICENSE and NOTICE at the root of this repository.

The State of Play

Current facts only. Earlier states, failures and chronology are kept whole in docs/HISTORY.md. The release this page describes is identified by RELEASE_MANIFEST.json (commit, fingerprints of the engine, the C++ twins, the GPU protocol, the live evidence and the verification receipt), which verify.py checks against the files.

Rerun everything: `pip install -r requirements.txt && python verify.py` ends with VERIFICATION: PASS.

In one paragraph (2026-10-07, Omni v3)

The engine is Omni v3, frozen and fingerprinted (OMNI_V3.json, docs/OMNI_V3.md: v1's law and controllers byte for byte, 945 muscles in 59 families, a do-no-harm gate on speed knobs). Every benchmark runs three times as separate GitHub runs on it (A the result, B and C the replications) and every judged row reads confirmed better, confirmed worse, no difference beyond the noise, or the runs disagree (docs/OMNI_V1.md, the rule). **Real Kubernetes, six tests, ten pairs each, all six confirmed on v3** (results/live/V3_*.md): work inside the response line **+35% to +49%** in all three runs of the all-four test (21.0 → 31.2 requests a second in run A); the 95th percentile **-47% to -66%** in every run of the steady, wandering, all-four and fault tests; failed requests **-9%** to **-14%** where load swings; machines **-1.5%** to **-2.9%** at steady load and **-19%** to **-23%** on the batch queue, with the standby-model energy **-13%** to **-16%** there; beside a noisy neighbour no difference beyond the noise on every row. **A real database** (PostgreSQL behind PgBouncer, results/live/V3_PGBENCH.md): 61% to 72% fewer connections held open for the same work and latency on two of three workloads, at a confirmed cost in the host's CPU seconds (**+14%** to **+28%**), counted against Omni. **Real messaging** (Apache Kafka as shipped, results/live/V3_KAFKA.md): on all three untouched workloads work inside the 500 ms line **+16% to +21%**, the 95th percentile 1.6 s → 9 to 14 ms and messages waiting **-92%** to **-97%**, confirmed better, no message lost; consumers held 2 → 6 to 8, **confirmed worse**, the resource the gain costs; host CPU worse on one workload, inside the noise on two. **A real cache** (Redis as shipped, results/live/V3_REDIS.md): on all three untouched workloads work inside the 2 ms line **+14% to +27%**, the hit rate **+14%** to **+27%** and the mean latency **-30%** to **-61%**, confirmed better; the memory ceiling held 64 → 200 to 270 MB, **confirmed worse**, the resource the gain costs; host CPU inside the noise. **A real database's storage-engine cache** (MongoDB as its publisher ships it under YCSB, results/live/V3_YCSB.md): on four untouched workloads the cache size held fell about half on c and burst and 13% to 37% on f, **confirmed better** (b inside the noise in one run); work inside the 1 ms line, p95 and host CPU inside the noise on all four; no failed operation; the burst mean latency **+2% to +4%**, **confirmed worse**, the one loss. **The Omni index, real machines only, confirmed three times:** **+25.9%** (results/OMNI_INDEX.md; Kubernetes +25.5%, the database +14.2%, Kafka +166.9%, Redis -24.9%, the database's cache +10.2%, each category weighed the same; a row inside the noise counts as exactly 1; Kafka's speed ratio is large because native's queue grew at nine tenths of its capacity and Omni's did not; Redis's category is negative because the memory it holds for a wide working set is the resource it trades and reads worse by rule; MongoDB's is positive because there the governor gave memory back). The v1 tables read the same and stay as the first engine's record.

What is shown and what is not. Shown: more work inside the response line on the same machines and a faster tail, on real Kubernetes, three times on a frozen engine; a real database holding fewer connections for the same service; a real message queue kept short at the cost of more consumers running. Not yet shown: an energy or cloud-bill saving on real machines. Energy on kind is a declared model (the machines are containers on one runner); Azure's bill on a 4-worker fleet read no difference beyond the noise on every gauge (`results/live/V1_AKS_STEADY.md`, `V1_AKS_BURST.md`), which is a fleet too small to show one machine; the fleet of 40 workers in nine machine families that can show one is preregistered and dispatched (`docs/K8S_COMPASS_PREREGISTRATION.md`, the fleet that can show one machine, amendments 1 to 3), waiting on Azure's own capacity in eastus. The real card has not run on the current governor; every earlier card result is obsolete and is run again by the founder on rented cards.

Measured on real systems (evidence class L), Omni v3

Test, ten pairs × three runs	Work	Speed (p95)	Machines	Energy (declared model)	Source
All four in one run	+35% to +49%	−61% to −66%	inside the noise	inside the noise	results/live/V3_ALL_FOUR.md
Steady work in steps	equal by design	−65% to −66%	−1.5% to −2.9%	−1.3% to −2.1%	results/live/V3_STEADY.md
Demand that wanders	failed requests −9% to −12%	−57% to −63%	inside the noise	inside the noise	results/live/V3_WANDERING.md
Faults: machine down, spike, runaway pod, blind probe	failed requests lower in all three, clear of the noise in one	−47% to −62%	inside the noise	inside the noise	results/live/V3_FAULTS.md
A queue of batch jobs	queue finished no difference beyond the noise	mean response −10% to −14%	−19% to −23% (−29% to −35% after the queue)	−13% to −16%	results/live/V3_BATCH.md
Fairness, a noisy neighbour	inside the noise on every row	inside the noise	inside the noise	inside the noise	results/live/V3_FAIRNESS.md
PostgreSQL behind PgBouncer, three workloads	inside the noise	inside the noise	connections held open −61% to −72% on two workloads; the runs disagree on the third	host CPU-seconds +14% to +28%, confirmed worse	results/live/V3_PGBENCH.md
Apache Kafka, a consumer group's size, three workloads	work inside the line +16% to +21%; no message lost	p95 1.6 s → 9 to 14 ms, lag −92% to −97%	consumers held 2 → 5.8 to 7.9, confirmed worse	host CPU-seconds +10% to +20% confirmed worse on light, inside the noise on heavy and burst; CPU per message inside the line −10% to −12% on burst	results/live/V3_KAFKA.md
Redis, a cache's memory ceiling, three workloads	work inside the line +14% to +27%, hit rate +14% to +27%; no failed request	mean latency −30% to −61%; p95 within a hair of native's (a miss is a miss in both arms)	memory ceiling held 64 → 200 to 270 MB, confirmed worse	host CPU-seconds inside the noise on all three	results/live/V3_REDIS.md
MongoDB under YCSB, a database's storage-engine cache size, four workloads	inside the noise on all four; no failed operation	p95 and p99 inside the noise; mean latency on burst +2% to +4%, confirmed worse	cache size held −49% on c, −41% to −49% on burst, −13% to −37% on f, confirmed better ; b inside the noise in one run	host CPU-seconds inside the noise on all four	results/live/V3_YCSB.md
Robustness: the governor killed outright at 40% of the window, 10 pairs × 3 runs	every setting back at the operator's 7 to 11 s after the kill, 30 of 30 ; a second governor to the end; failed requests over the window −9% to −15%	the 120 s after the kill inside the noise against native; the whole window p95 −42% to −45%	inside the noise	inside the noise	results/live/V3_ROBUST_KILL.md

Test, ten pairs × three runs	Work	Speed (p95)	Machines	Energy (declared model)	Source
Robustness: the long run, 7,200 s an arm, 3 pairs × 3 runs	decisions 97.5% or more of expected (valid); failed requests inside the noise	service confirmed better in A, inside the noise in B and C (no difference beyond the noise in 2 of 3)	inside the noise; memory at most 1.07 of its first ten minutes (no leak) ; decision time at most 1.20 (no slowing); every setting handed back at the end	inside the noise	results/live/V3_ROBUST_LONG.md
The six organisms with the real cluster inside, 10 and 100 copies, 5 pairs a cell	the organisms' work unchanged	p95 better in every cell	6 in both arms (no autoscaler under kind)	the organisms' energy lower in every cell	results/live/V3_SIX_KUBE.md (v1 at 1 to 1,000 copies: V1_SIX_KUBE.md, V1_BIG_ORGANISM.md)

Measured on a real cloud (Azure, its own bill), Omni v1

Test	Reading	Source
Steady load, 4 workers, 5 pairs	no difference beyond the noise on any gauge; the bill −4.7% with its interval across zero	results/live/V1_AKS_STEADY.md
A burst sized to the cluster, 4 workers, 5 pairs	the bill +5.0% with its interval across zero; p99 −34% clear of the noise in this one run; the rest inside the noise	results/live/V1_AKS_BURST.md
The fleet that can show one machine (40 workers, nine families) on v3	dispatched; two earlier dispatches refused by the subscription's family allowances before any arm ran, the third by Azure's own cluster capacity in eastus; every refusal cost cents and is recorded	docs/K8S_COMPASS_PREREGISTRATION.md

Simulated (evidence class S: models, never counted in the headline), Omni v3

What	Reading	Source			
The 945 muscles and six organisms, A/B/C	reproduced to the digit in 3 of 3; 0 muscles worse; every organism superior within guardrails (work per energy +0.1% to +0.3%)	results/realms/REALMS.md			
The four stacked at 1,000 copies with the real cluster inside, on a rented machine, 3 pairs (v3)	the organism's work unchanged (rounding-level)	the cluster's p95 2.9 s → 150 ms , p99 −97%, time over the line 51% → 0 , failed requests 0.14% → 0	6 machines in both arms	inside the noise (standby model); the organism's energy −0.4% (model)	results/live/V3_BIG_ORGANISM.md
The organisms at 1, 10, 100 and 1,000 copies × 1 to 1,000 runs, 84 of 90 cells	every organism superior within guardrails in every cell of 10 runs or more; work per energy +0.07% to +0.37%, the same figure at every size; the six cells left are beyond the machines available	results/scale/GRID.md			
Power grid, 11 SimBench grids in pandapower, A/B/C (v1 and v3 identical)	with ZIP loads energy drawn and net import better in all 11; losses better in 7, worse in 4; tap operations fewer in 10, 4 → 8 a year in one (worse, the declared cost)	results/live/V3_PANDAPOWER.md			
Robot arms, MuJoCo Menagerie, A/B/C	Gen3 and Panda: peak torque −29% and −10%, tracking error −21%, energy per takt −0.8% and −0.5%; the Panda's copper +14% worse; UR5e and iiwa left native	results/live/V3_MUJOCO.md, V1_MUJOCO_PANDA.md			
CityLearn, every district, A/B/C	electricity bought, peak and unevenness better in all 11 battery districts, carbon in 8; the bill worse in 7, ramping worse in 7	results/live/V3_CITYLEARN.md			
Drone swarms, gym-pybullet-drones, 20 drones in three cells, A/B/C	energy a mission −7% to −20% and missions a charge +8% to +25%, confirmed better; no late mission, reserve breach, near miss or collision in any arm	results/live/V3_SWARM.md			

The card (evidence class P)

Every earlier card result ran on a controller since replaced and is obsolete. The one-card, card-inside-the-organisms and eight-card runs are the founder's, on rented cards, after the CPU and cloud work, at one named commit (docs/GPU_RUN_GUIDE.md, docs/GPU_PREREGISTRATION.md).

Verified in code

Property	Where
The canonical engine is <code>symmetric_verified</code> ; the printed chart is a named variant, not benchmarked	<code>docs/CANONICAL_ENGINE.md</code>
Nine laws twinned in C++20 and proven equal to the Python; sealed by fingerprint	<code>results/SEAL.json</code> , <code>tools/seal.py</code>
The conveyance law conserves its budget and converges (proof and 20,000 random systems)	<code>docs/CONVEYANCE_LAW.md</code> , <code>tests/test_conveyance.py</code>
Safety shield: 2,000,000 adversarial cases, 0 violations; C++ engine: 100,000,000 decisions, no failures	<code>tests/test_shield_properties.py</code> , <code>results/SOAK.json</code>

Open

1. **Azure, the fleet that can show one machine:** steady and burst on 40 workers, v3; waiting on Azure's cluster capacity in eastus (the only region where this subscription has more than 10 cores). Then B and C.
2. **The four stacked and the tower at 1,000 copies with the real cluster inside, on v3:** the stack is done, 3 of 3 on the clock (`results/live/V3_BIG_ORGANISM.md`: the cluster's p95 2.9 s → 150 ms, time over the line 51% → 0, failed requests 0.14% → 0, machines and energy unchanged, 22 s behind a 10,800 s window); the tower runs now on a second rented machine the same way (about \$10, about a day).
3. **The real card:** the founder's runs on Lambda, one exact commit.
4. **Robustness** (`docs/ROBUSTNESS_PREREGISTRATION.md`, no engine file changes): the kill scenario is done (`results/live/V3_ROBUST_KILL.md`: every setting back 7 to 11 s after the kill in 30 of 30 repetitions) and the governor's own CPU is tabulated (`results/live/V3_OWN_COST.md`); the long run is done (`results/live/V3_ROBUST_LONG.md`, 3 pairs × 3 runs, 7,200 s an arm: no leak, the decisions valid, no slowing, every setting handed back at the end; service inside the noise in two of three runs); the 24-hour run on a rented machine follows.
5. **YCSB on MongoDB, done** (`docs/YCSB_PREREGISTRATION.md`, `results/live/V3_YCSB.md`): the operator's WiredTiger cache as native, Omni on the cache size through the server's own console; four smoke runs recorded, then A, B and C on v3; the cache held given back by about half on two untouched workloads at no measurable cost in work, p95 or CPU, one confirmed loss (the burst mean latency, +2% to +4%); the index now spans five real categories at +25.9%. Cassandra and Redis under YCSB and HammerDB are next in row 24.
6. **The queue** (`docs/REGISTER.md` section 4, `docs/PROOF_PROGRAM.md`): drone swarms on PX4 and ArduPilot (gym-pybullet-drones done), YCSB on Cassandra and Redis and HammerDB, Spark, OpenSearch, fio, Open-RMF, the 24-hour robustness run, Basilisk, Orekit and GMAT, RocketPy, Cantera (Kafka and Redis done); one or two at a time, each preregistered.
7. **Omni-Compass 1.0:** when the founder declares the engine final, v3 as it stands is published as 1.0 and the older fingerprints go to `docs/history` as the road to it.

Where things are

Path	What it is
docs/OMNI_COMPASS_MANUAL.md (PDF: docs/OMNI_COMPASS_MANUAL.pdf)	the manual: the governor, its mechanism, the wiring stack by stack, the frozen engines, the three-run rule, every result
docs/OMNI_V3.md, docs/OMNI_V1.md	what each engine is and every result read on it
docs/REGISTER.md, docs/PROOF_PROGRAM.md	every muscle, every benchmark run and still to run; the program to full size
results/live/V3_*.md, V1_*.md, results/live/raw/	the three-run tables and every archived run's files
results/OMNI_INDEX.md	the one combined number
docs/INTEGRATION_MANUAL.md, docs/WIRING_GUIDE.md	wiring it in yourself
docs/METRICS_CATALOG.md	every gauge, and whether it is measured or modelled
omnicompass/, omni_controller/, realms/, cpp/	the engine, the controllers, the muscles, the C++20 twins
docs/HISTORY.md	earlier states of play, kept whole

Evaluation and simulation use only. Copyright (c) 2026 The Omni-Compass LLC. Commercial use, commercialization or monetization of any part of Omni-Compass requires a signed, paid Omni-Compass Enterprise License from The Omni-Compass LLC. Patent applications, copyright registrations and trademark applications filed in the United States. See LICENSE and NOTICE at the root of this repository.

PART EIGHT

Value, License and History

*What a receipt is worth, how the license is priced against it, how the code is sealed,
and how Omni-Compass came to be.*

Where the Value Comes From

The problem we are attacking

Computing has run into a wall that is not made of silicon. Data centers cannot get the power they need, cannot get it fast enough, and are running out of room to put more machines. AI made it worse: every new model needs more chips, and every chip needs more electricity and more cooling. The answers on the table are expensive and slow: new chips that do more work per watt (a new generation, new hardware to buy), new power plants and power contracts (billions, and years), and point fixes at a single layer (an AI serving trick, a scheduler, a power cap on one chip).

We attack it from the other side. Every system already running wastes room it paid for:

- GPUs boost to the top of their clock range and are knocked back by their own power limiter many times a second;
- Kubernetes keeps replicas and machines sized for the worst minute of the day;
- cooling runs colder than the heat requires; batteries hold more reserve than the hour needs.

That room is paid for in energy, in machines and in floor space, and it produces nothing.

What Omni-Compass is, in one paragraph

Omni-Compass is software that sits **on top of** the controllers a system already has (the GPU's firmware, the Kubernetes autoscaler, the building's thermostats) and governs every layer under one law: a bounded push and pull that keeps each service in the middle of its band instead of far below its limit. It never replaces the native controller, never asks for more than the native settings allow, hands every knob back the moment it is switched off or loses a sense, and refuses any step that measures more than 2% worse on anything. No new hardware. It stacks on top of every other fix a data center makes.

The benefit: more work from what is already paid for

The room that was spent on nothing becomes one of two things, and a company chooses which:

- **More work for the same cost:** the same machines and the same watts carry more work; or
- **The same work for less:** fewer machines in service, fewer watts, a smaller bill.

They are the same gain read from two sides (3.4). The health side comes with it: faster answers, faster recovery from failures, nothing waiting, every knob returned, no human babysitting the stack.

How to read the benefit: the arithmetic

Every comparison is paired: the same system, the same work, the same moments, native alone against native with Omni-Compass on top. Let

- W be the work served (requests answered, tokens generated), held equal in both arms;
- C be a resource it used: machine-hours, CPU core-hours, or joules.

Productivity is work per resource, $P = W / C$. With the work held equal, the gain from Omni-Compass on top is

$$G = P_{\text{omni}} / P_{\text{native}} - 1 = C_{\text{native}} / C_{\text{omni}} - 1.$$

A gain G means the same resources can carry G more work. Read from the other side, the same work needs only

$$C_{\text{omni}} / C_{\text{native}} = 1 / (1 + G)$$

of the resources, a saving of $S = G / (1 + G)$. So **a third more work ($G = 33\%$) is a quarter off the bill ($S = 25\%$)**; 20% more work is 16.7% off; 50% more work is a third off. Both columns are the same measurement.

Three rules for reading any receipt:

1. **The work must be equal.** Every receipt states that the load was fixed-rate (open loop): the same requests at the same moments in every arm. Without that, a lower resource count could just mean less work was done.
2. **Nothing may be worse.** Beside every gain the receipt prints response times, failures, waiting and the controller's own CPU, each with its 95% interval. Our one rule (DISCLOSURES.md, Chapter 53): no measure more than 2% worse, and only where energy or the bill is saved. A gain paid for with worse service is not a gain.
3. **Know the class and the comparator.** A machine-hour gain on kind (where idle machines stay powered and native has no node autoscaler) is not yet a bill; a modelled energy line is not a meter. Section 14 lists the classes; Chapter 41 says, for every result, what it is measured against.

Three worked examples, from the tables. In the all-four Kubernetes test on v3, the same ten machines carried 35% to 49% more requests inside the response line: G is 0.35 to 0.49, so the same work would need 1/1.35 to 1/1.49 of the machines, a saving S of 26% to 33%, had the fleet been free to shrink; on kind it was not, which is why the table reports the work column and not a bill. In the database test, the pooler held 61% to 72% fewer connections open for the same work: the resource there is connections, $C_{\text{omni}} / C_{\text{native}}$ is 0.28 to 0.39, and the gain read as "more work for the same connections" would be $G = 1.6$ to 2.6, which the table reports as the machines-equivalent column, "fewer connections by". In the Kafka test the work inside the line rose 16% to 21% while the consumers running rose by a factor of 2.6 to 4: there the gain in work was bought with resources, and the table reports both, the one as confirmed better and the other as confirmed worse, because that is what happened.

Where the evidence stands today, read honestly

Every number below is from a three-run table on the frozen engine Omni v3, with its evidence class and the question a referee is entitled to ask about it. The answers to those questions are given, where they exist, in the row; where they do not yet exist, the row says what will supply them.

What	Measured, same work, three runs	Class	What a referee will ask, and the answer
Kubernetes: work inside the response line, all four in one run	+35% to +49% confirmed better	L	"Is the cluster real?" Kubernetes is real (API server, scheduler, HPA); the seven nodes are containers on one 4-core GitHub runner, so the host's own busy share is recorded beside every run and the machines column is machine-hours in service, not a bill. The bill is Azure's test (below)
Kubernetes: 95th-percentile response time, steady, wandering, all four, faults	47% to 66% faster confirmed better	L	"Did native have a fair chance?" Native is Kubernetes with its HPA at the operator's target of 50%; the cost-to-match runs on the earlier engine showed no native target tried (40, 30, 20) reached Omni's p95 (docs/history/). The runs are paired on one runner with the order rotated
Kubernetes: machines in service	−1.5% to −2.9% at steady load, −19% to −23% on the batch queue, confirmed better; inside the noise on the other tests	L	"Why so little?" Because the verdict kept every machine whose removal made requests slower in the paired trial. Earlier engines without that check parked 29% to 36% of the machines and paid for it in response time; the frozen engine puts service first. kind has no node autoscaler, so parked machines stay powered: the saving is machine-hours, not watts
Kubernetes: energy	the standby model, −13% to −16% on the batch queue, inside the noise elsewhere	L, modelled	"Is it a meter?" No. The machines are containers; energy is a declared formula (a parked worker at 25 W standby), stated as a model on every line
Kubernetes: fairness beside a noisy neighbour	no difference beyond the noise on every row, the neighbour's included	L	"Does it hurt the other tenant?" Not measurably, in three runs of ten pairs
Database behind its pooler: connections held open	−61% to −72% confirmed better on two workloads; the runs disagree on the third	L	"What did it cost?" Host CPU-seconds +14% to +28%, confirmed worse on all three workloads, the compass's own cost included, counted against Omni in the index; work and latency inside the noise
Message broker: work inside the 500 ms line; the 95th percentile	+16% to +21% ; 1.6 s → 9 to 14 ms, confirmed better on all three workloads; no message lost	L	"Why is the latency figure so large?" By design native sat at nine tenths of its measured capacity, so its queue grew at the high steps and Omni's did not; the gain is the queue kept short. "What did it cost?" Consumers running 2 → 5.8 to 7.9, confirmed worse; host CPU worse on one workload, inside the noise on two

What	Measured, same work, three runs	Class	What a referee will ask, and the answer
Database's storage-engine cache (MongoDB under YCSB): the cache size held	−49% on c, −41% to −49% on burst, −13% to −37% on f , confirmed better; b inside the noise in one run	L	"What did it cost?" Work inside the 1 ms line, p95, p99 and host CPU inside the noise on all four workloads, no failed operation; the burst mean latency +2% to +4%, confirmed worse , in the table. "Why is this the mirror of the Redis result?" Because here the operator's setting was larger than the working set needed on this machine, so the governor gave memory back; there it was smaller, so the governor bought service with memory. "Would this hold on a disk-bound store?" Not shown: on this machine a miss is a memory read, as disclosed before the run
Cache: work inside the 2 ms line; the hit rate	+14% to +27% and +14% to +27%, confirmed better on all three workloads; no failed request	L	"What did it cost?" The memory ceiling held, 64 → 200 to 270 MB, confirmed worse, and the memory used with it; host CPU inside the noise. "Why is p95 unchanged?" A miss costs the declared 5 ms trip in both arms and 5% of requests still miss at the high notches; the gain is in the mean and in the work inside the line. "Why is the cache's index negative?" Because its resource column is the memory held, a ratio of about 0.25, and the geometric mean of a 1.2 gain and a 0.25 cost is below one
The modelled realms: work per energy, 945 muscles, six organisms, 1 to 1,000 copies	+0.07% to +0.37% , every organism superior within guardrails, 0 muscles worse	S	"Why so small?" The native controllers in the models are well tuned and leave little room; the number is small and real within the model, the same at every size, and it is never counted in the headline
Power grid, 11 SimBench grids in pandapower	energy drawn and net import better in all 11; losses better in 7, worse in 4 ; tap operations fewer in 10, 4 → 8 a year in one, worse	S	"Where does it lose?" In the rural and semi-urban grids with their own generation, where a lower voltage raises losses; the table shows it
Robot arms, MuJoCo Menagerie	peak torque −29% and −10%, tracking error −21% where Omni moved; the Panda's copper loss +14% worse ; two arms left native	S	"What about the arms that gained nothing?" The paired physics trial left them native and the table says "nothing for Omni to move"
Buildings with batteries, CityLearn, 11 districts	electricity bought, daily peak and unevenness better in all 11; the bill worse in 7 ; ramping worse in 7	S	"Is the bill a loss?" Yes, in the 2023 districts, and it is in the table as such
Drone swarms, gym-pybullet-drones, three 20-drone cells	energy a mission −7% to −20%, missions a charge +8% to +25%; no late mission, reserve breach, near miss or collision	S	"Is the energy a meter?" No, a declared model from the simulator's own motor constants; the tracking error rose from 0.07 to 0.13 m inside its 0.25 m band and is shown
The governor's own cost, 1 to 1,000 copies with the real cluster inside	0.006 to 0.013 of one core at every size, 0.1% to 0.3% of the host's cores	L	"Does the brain's cost grow with the body?" No: it governs the real cluster, and that cost does not grow with the organism around it; shown, not judged, each run under its own engine (results/live/V3_OW_N_COST.md)

What	Measured, same work, three runs	Class	What a referee will ask, and the answer
The governor killed outright mid-run, 10 pairs × 3 runs	every setting back at the operator's 7 to 11 s after the kill in 30 of 30 repetitions (mean 9 s; allowance 60 s); a second governor to the end in every one; the 120 s after the kill inside the noise against native	L	"What if Omni-Compass dies?" The watchdog hands back from the lease the governor left, within seconds, and the service does not measurably notice; the whole window with a kill and a restart in it still read faster and with fewer failures than native (<code>results/live/V3_R0 BUST_KILL.md</code>)
Azure's bill on a 4-worker fleet (v1)	no difference beyond the noise on any gauge	L, a real bill	"Then where is the saving?" A fleet of 4 cannot show one machine; the 40-worker fleet that can is preregistered and waits on Azure's own cluster capacity in its region
The real card (GPU)	obsolete: every earlier card result ran on a controller since replaced	P pending	"When?" The founder's runs on rented cards at one named commit, after the CPU and cloud work

The upside the real-software data points toward is a third to a half more work from what is already installed on the Kubernetes side, with faster answers, and a smaller, resource-priced gain on the stacks where the knob is a pool, a consumer group or a ceiling. It becomes a claim about the bill, not a direction, when two tests land: the bill on a real cloud with a fleet big enough to show one machine, and the real card's own meter. We publish each of them as it comes, whatever it says.

One law, frozen; and what came before it

The engine in every result of this manual is **the compass law with the verdict** (Chapter 16 and 1.2), frozen as Omni v1 on 5 October 2026, grown to the 945-muscle catalog as v2 and given the slack gate on speed knobs as v3, with the law, the controllers and the runners the same bytes in all three (`docs/OMNI_V1.md`, `OMNI_V2.md`, `OMNI_V3.md`). Earlier engines, among them an allocation law that gave back about a third of the machines at the cost of response time and a compass law without the verdict, produced the numbered sets kept in `docs/history/`; their results are history, labelled as the engine that produced them, and are never read beside v1's or v3's. A buyer does not choose between modes: there is one law, and it is the one every table in Part VI was made on. Each new result therefore describes exactly the code that produced the old ones, and when the founder declares the engine final, the engine that stands then is published as Omni-Compass 1.0, with the older fingerprints kept as the road to it, never as a second product.

The Economics of a Receipt

Run the stack native and print the receipt. Run the same stack with Omni-Compass and print the receipt. The difference is the gain, read from either side: more work for the same energy, or the same work for less energy and fewer machines.

The license fee is set against the measured gain on the customer's own receipt. On a stack that costs one hundred, a measured gain of ten is worth ten; the fee is a share of that ten, and the customer keeps the rest. Percentages from different parts of a stack are not added into one bill: the invoice is one gain, measured once.

The value a customer sees comes in three forms, each on its own line of the receipt:

1. **Work per energy** - more finished work per kilowatt-hour.
2. **Capacity** - the same work on fewer machines in service; it becomes a cost reduction when the released machines are returned to the cloud or switched off by the platform around Omni-Compass.
3. **Service** - faster responses and fewer breaches at the same load.

The babysitting tax - the people and tools kept on the clock to set caps, answer pages and turn knobs back after a run or a crash - is the cost Omni-Compass removes by holding the knobs and returning them itself.

The receipt's arithmetic, and the one number. With the work held equal, a gain G in productivity is the same measurement as a saving $S = G / (1 + G)$ in resources: a third more work is a quarter off the bill (manual, section 3.4). Across every real stack the one number is the Omni index (`results/OMNI_INDEX.md`): each test's work, speed, machines and energy as ratios, combined by geometric means so that a gain and an equal loss cancel exactly, each real category weighed the same, and only rows confirmed in three separate runs counted. The index is honest in both directions by construction. As this edition is written it reads +30.2%: real Kubernetes +25.5%, the database +14.2%, messaging +166.9%, and the cache -24.9%, the last negative because the memory the governor holds for a wide working set is the resource it trades and reads worse by rule, even as the cache's work and hit rate read better. A buyer reads the category that matches their stack, and reads its losses in the same table as its gains.

What a receipt is not. A receipt is not a forecast. It is the measured difference between two arms on one system at one time, with its interval. The only number that applies to a customer's system is the one their own paired runs produce, and the method of those runs is the manual's section 13, the same method every table in this book was made with.

Evaluation and simulation use only. Copyright (c) 2026 The Omni-Compass LLC. Commercial use, commercialization or monetization of any part of Omni-Compass requires a signed, paid Omni-Compass Enterprise License from The Omni-Compass LLC. Patent applications, copyright registrations and trademark applications filed in the United States. See LICENSE and NOTICE at the root of this repository.

The Buyer Edition

Superseded for live results. The live evidence below predates sets 19 and 20. The current live table is `results/live/LIVE_PAired.md`; energy on kind is a declared model, not a meter. The first metered test is `scripts/gpu_paired.sh` (`docs/GPU_BENCH.md`).

Commit `cdc8cd9`. Every number below is generated from result files in the repository by `python tools/buyer_report.py`. Simulated results use the repository's fleet plant; live results come from real Kubernetes (kind) in GitHub Actions. Competitors are reproduced from their public documentation, not their binaries.

What Omni-Compass is

- **The engine:** a six-state control engine, stable by construction (every mode decays), in Python and C++.
- **The law it applies:** the closure law of the owner's manuscript.
 - It does nothing in the safe interior.
 - It corrects inward when the forecast approaches capacity.
 - It releases a machine only after the peak.
 - It keeps released machines warm for instant reuse.
- **What it governs:** machines, pod counts and sizes, cold start, batch pacing, containment of runaway agent workloads, and cooling setpoints.
- **Two ways to deploy it:**
 - **B, on top of the platform a customer already runs.** Nothing is replaced.
 - **C, alone.** Kubernetes stays only as the muscle; the scaling layer is replaced (HPA, Cluster Autoscaler, Karpenter, and optimizers such as CAST AI, Spot Ocean and Turbonomic).

B: Omni-Compass on top of each platform (held-out scenarios, settings frozen first)

Workload	Platform	Omni-Compass on top vs the platform alone
Web services	Kubernetes	scale reversals 12% better, machine starts+stops 2% better
Web services	OpenShift	scale reversals 44% better, machine starts+stops 4% better
Web services	GKE	scale reversals 21% better, machine starts+stops 7% better, machine-hours 3% better
Web services	AKS NAP / Karpenter	scale reversals 8% better, machine starts+stops 3% better
Web services	Turbonomic	scale reversals 8% better, machine starts+stops 2% better
Web services	CAST AI	slowest responses (p99) 34% better, scale reversals 30% better, machine starts+stops 15% better
Web services	Spot Ocean	slowest responses (p99) 42% better, scale reversals 34% better, machine starts+stops 17% better
Four clusters, one site	Kubernetes	scale reversals 93% better, slowest responses (p99) 51% better, machine starts+stops 21% better
Four clusters, one site	OpenShift	scale reversals 100% better, slowest responses (p99) 39% better, machine starts+stops 16% better
Four clusters, one site	GKE	slowest responses (p99) 38% better, scale reversals 19% better, machine starts+stops 11% better
Four clusters, one site	AKS NAP / Karpenter	scale reversals 7% better, slowest responses (p99) 6% better, machine starts+stops 4% better
Four clusters, one site	Turbonomic	scale reversals 4% better, machine starts+stops 2% better
Four clusters, one site	CAST AI	equal
Four clusters, one site	Spot Ocean	scale reversals 9% better, slowest responses (p99) 5% better, machine starts+stops 5% better
Batch jobs	Kubernetes	slowest responses (p99) 91% better, scale reversals 82% better, average response 25% better
Batch jobs	OpenShift	slowest responses (p99) 92% better, scale reversals 84% better, average response 26% better
Batch jobs	GKE	time over backlog limit 100% better, slowest responses (p99) 89% better, scale reversals 66% better
Batch jobs	AKS NAP / Karpenter	typical response (p95) 60% better, slowest responses (p99) 34% better, time over backlog limit 25% better
Batch jobs	Turbonomic	time over backlog limit 44% better, scale reversals 22% better, machine starts+stops 15% better
Batch jobs	CAST AI	time over backlog limit 100% better, typical response (p95) 94% better, slowest responses (p99) 91% better
Batch jobs	Spot Ocean	time over backlog limit 83% better, scale reversals 59% better, machine starts+stops 47% better
GPU training	Kubernetes	scale reversals 68% better, machine starts+stops 35% better
GPU training	OpenShift	scale reversals 61% better, machine starts+stops 29% better, time healthy 1% better
GPU training	GKE	scale reversals 59% better, machine starts+stops 39% better, machine-hours 1% better
GPU training	AKS NAP / Karpenter	machine starts+stops 51% better, scale reversals 33% better, machine-hours 2% better
GPU training	Turbonomic	scale reversals 12% better, machine starts+stops 9% better

Workload	Platform	Omni-Compass on top vs the platform alone
GPU training	CAST AI	machine starts+stops 45% better, scale reversals 32% better, machine-hours 2% better
GPU training	Spot Ocean	machine starts+stops 28% better, scale reversals 21% better

C: Omni-Compass alone, one global setting, confirmatory run

- **Setting:** one global setting for every workload, frozen by SHA-256 before the run (28cbb82607cb0009).
- **Scenarios:** 100 never-used scenarios per workload (seeds 710001-710100).
- **Comparisons:** 364 cells (4 workloads x 7 platforms x 13 gauges).
- **Rule:** a cell counts as better or worse only if it survives Holm-Bonferroni correction across all 364 cells (family alpha 0.05) and exceeds the 0.5% practical tolerance. Otherwise it is equal.

Result: 153 better, 174 equal, 37 worse.

Workload	Better	Equal	Worse	The worse cells
Web services	41	41	9	Kubernetes typical response (p95) 20.7%; Kubernetes machine starts+stops 24.5%; OpenShift typical response (p95) 20.8%; OpenShift machine starts+stops 66.0%; OpenShift scale reversals 152.0%; GKE typical response (p95) 20.1%; AKS NAP / Karpenter typical response (p95) 19.7%; CAST AI typical response (p95) 19.2%; Spot Ocean typical response (p95) 19.7%
Four clusters, one site	45	36	10	Kubernetes typical response (p95) 20.4%; Kubernetes machine starts+stops 23.7%; OpenShift typical response (p95) 20.5%; OpenShift machine starts+stops 62.1%; OpenShift scale reversals 207.3%; GKE typical response (p95) 19.7%; AKS NAP / Karpenter typical response (p95) 19.3%; Turbonomic typical response (p95) 27.6%; CAST AI typical response (p95) 18.8%; Spot Ocean typical response (p95) 19.3%
Batch jobs	35	50	6	Kubernetes machine starts+stops 39.0%; OpenShift machine starts+stops 29.5%; AKS NAP / Karpenter energy 2.1%; CAST AI energy 1.2%; Spot Ocean energy 2.2%; Spot Ocean machine-hours 0.6%
GPU training	32	47	12	AKS NAP / Karpenter energy 1.0%; AKS NAP / Karpenter time over power limit 1.6%; AKS NAP / Karpenter time over heat limit 1.1%; CAST AI energy 1.0%; CAST AI time over power limit 3.2%; CAST AI time over heat limit 2.6%; CAST AI time healthy 1.6%; Spot Ocean energy 1.2%; Spot Ocean time over power limit 3.8%; Spot Ocean time over heat limit 1.8%; Spot Ocean machine-hours 1.4%; Spot Ocean time healthy 2.0%

Workload	Energy vs Kubernetes	Machine-hours	Slowest responses	Average response
Web services	+21%	+38%	+40%	+11%
Four clusters, one site	+25%	+40%	+41%	+13%
Batch jobs	+6%	+18%	+40%	+2%
GPU training	+2%	+11%	+0%	+0%

Positive means Omni-Compass is better.

Real demand: 1,052 recorded machines (PlanetLab)

- **Losing cells:** Omni-Compass alone, with the frozen web setting (never tuned on these traces), loses 0 of 91 against the seven platforms.
- **Whole-pod packing:** with whole-pod packing it loses 0 of 91.
- **Against Kubernetes:**

- energy +24%
- machine-hours +33%
- typical response +97%
- slowest responses +85%

Positive means better.

Four-cluster sites as one body (held-out)

- **The mechanism:** traffic shift between clusters, plus the law run on the site total with one warm reserve.
- **Omni-Compass alone:** zero losing cells against GKE, AKS NAP / Karpenter, CAST AI, Spot Ocean.
- **Omni-Compass on top:** zero losing cells on all seven platforms.
- **Assumption:** services are replicated across the site's clusters.

Under failure: the runtime protocol

Machines dying, load spikes, crash-looping services and noisy neighbours were injected at five stress levels, with identical faults for every system. Totals over all levels:

Workload	System	Runs conveyed	Pages to a human	Recovery (min)
Web services	CAST AI	231 of 500	404	6.6
Web services	Kubernetes	218 of 500	338	8.3
Web services	Omni-Compass alone	240 of 500	246	5.1
Four clusters, one site	CAST AI	44 of 500	608	5.7
Four clusters, one site	Kubernetes	44 of 500	330	6.5
Four clusters, one site	Omni-Compass alone	46 of 500	197	3.6

The supervisory nervous system

- **One state, every organ:** `omnicompass/nervous_system.py` turns the engine's state (convergence, basin health, stress against its equation-6 equilibrium, unmet need) into one calm value between 0 and 1.
- **Authority from calm:** calm grants each organ its authority:
 - pods and machines may give capacity back only above their reversibility thresholds (0.5 and 0.7);
 - CPU frequency, GPU power, routing and cooling get envelopes that widen with calm;
 - batch is admitted or paused;
 - rollback is authorised.
- **Holds and the shield:** a security hold stops every capacity organ from expanding. The shield stays downstream and can still veto.
- **Invariants tested:** 300,000 random engine states, with zero violations (`tests/test_nervous_system.py`).
- **Coordination:** pods move first, and a machine move opposite to the pod move is vetoed. Controller contradictions per day (fighting, or a reversal within one boot time):
 - web: 2.10 to 0.10;
 - four-cluster: 5.30 to 1.03, below Kubernetes' 2.10.

Coordination, confirmatory on fresh seeds 712001-712100 (frozen first in tuning/COORD_PREREGISTRATION.json, same 364 cells, same Holm rule): 151 better, 173 equal, 40 worse.

Workload	Better	Equal	Worse	Contradictions per day: Omni-Compass / Kubernetes
Web services	40	41	10	0.07 / 0.52
Four clusters, one site	45	37	9	0.56 / 1.77
Batch jobs	35	49	7	0.11 / 0.00
GPU training	31	46	14	0.11 / 0.04

Mechanism ablation (tuning/ABLATION.json, fresh seeds 711001-711030): each part of the law removed in turn, against the full law. Listed: the gauges that get significantly worse (95% interval excludes 0, more than 0.5%).

Part removed	Web services	Four clusters, one site	Batch jobs	GPU training
engine release gate removed	no change	no change	no change	no change
turning point removed	no change	machine starts+stops +2%, scale reversals +6%	slowest responses (p99) +32%, average response +12%, time over backlog limit +333%, machine starts+stops +47%, scale reversals +193%, contradictions +1350%	machine starts+stops +18%, scale reversals +118%, contradictions +1000%
muscle tone removed	typical response (p95) +1%, slowest responses (p99) +110%, average response +4%, machine starts+stops +92%, scale reversals +458%, contradictions +22%	slowest responses (p99) +8%, machine starts+stops +29%, scale reversals +108%	typical response (p95) +65%, slowest responses (p99) +151%, average response +26%, machine starts+stops +239%, scale reversals +1265%, machine-hours +1%, contradictions +875%	time over backlog limit +2%, machine starts+stops +297%, scale reversals +829%, contradictions +1150%
trend term removed	no change	slowest responses (p99) +2%	slowest responses (p99) +75%, average response +24%, time over backlog limit +533%, machine starts+stops +65%, scale reversals +348%, contradictions +2350%	machine starts+stops +29%, scale reversals +129%, contradictions +1100%

Live Kubernetes evidence

Repeated live runs, 5 paired repetitions per arm, probe through the Service (results/live/LIVE_REPS_3.md):

Live repetitions, set 3: GitHub Actions run 36289557446 (commit 9892270), 5 × native / omni / strict on kind

This is the first set with the probe through the Service (NodePort via kube-proxy) and the PodDisruptionBudget, and the first live set with the supervisory nervous system gating machine release.

B: Omni-Compass on top vs native, 5 paired repetitions

Gauge	Native	Omni	Change	95% interval of the difference	Significant
worker nodes in service, mean	6	6	+0.0%	-3.155e-16 to +6.708e-16	no
node-hours	1.517	1.516	-0.1%	-0.02125 to +0.01859	no
energy (Wh)	168.1	165.7	-1.4%	-7.697 to +2.854	no
response time (ms), mean	221.7	210.3	-5.1%	-57.38 to +34.55	no
response time (ms), 95th percentile	431.4	440.2	+2.0%	-47.09 to +64.69	no
response time (ms), 99th percentile	568.9	609	+7.0%	-51.68 to +131.8	no
failed requests (%)	0	0	+0 (native is 0)	+0 to +0	no
pending pods, pod-minutes	0.3433	0.3967	+15.5%	-0.4106 to +0.5172	no
utilisation (used / allocatable)	0.07676	0.06667	-13.1%	-0.02607 to +0.005892	no

C: Omni-Compass decides (strict) vs native, 5 paired repetitions

Gauge	Native	Omni	Change	95% interval of the difference	Significant
worker nodes in service, mean	6	6	+0.0%	-7.797e-16 to +7.797e-16	no
node-hours	1.517	1.518	+0.0%	-0.00974 to +0.01041	no
energy (Wh)	168.1	166.4	-1.0%	-4.867 to +1.447	no
response time (ms), mean	221.7	217.8	-1.8%	-89.74 to +81.98	no
response time (ms), 95th percentile	431.4	463.2	+7.4%	-105.7 to +169.4	no
response time (ms), 99th percentile	568.9	629	+10.6%	-155.4 to +275.7	no
failed requests (%)	0	0	+0 (native is 0)	+0 to +0	no
pending pods, pod-minutes	0.3433	0.8033	+134.0%	+0.05385 to +0.8661	yes, worse
utilisation (used / allocatable)	0.07676	0.06928	-9.8%	-0.01814 to +0.003173	no

Reading

Failed requests. They are 0% in every arm. The 16-66% of sets 1-2 was the probe's hung tunnel, as diagnosed.

Machines. Omni-Compass kept all 6 machines in service in every repetition. Controller log: 15 of 15 decisions, 0 errors, no node-scale command.

- **The likely cause.** With a working probe, p95 sits at 430-460 ms against the 500 ms SLO declared before the run. Every window above 500 ms is SLO pressure. That pressure blocks machine release for the next 3 decisions and lowers the engine's calm.
- **Why that is odd here.** Machines are at 7% utilisation with no waiting pods. The latency comes from the workload's per-request compute, not from a machine shortage.
- **Why the earlier sets differed.** Sets 1-2 released machines because the hung probe produced few or no successful samples, so the controller saw little latency pressure. Their machine savings were therefore measured with a blinded latency sense and are not claimed.
- **How this will be settled.** The decision trail (p95, SLO state, calm, node authority per decision) is now printed in every job log. Set 4 will show which gate held.

Everything else. No significant difference from native on any gauge, except that C has more pending pod-minutes (0.80 vs 0.34 pod-minutes).

- **Earlier live sets 1 and 2:** these carried a probe defect, now fixed. The probe's one-pod tunnel hung when a drain moved its pod, so the Omni arms logged false failed requests (results/live/LIVE_REPS_PROBE_DEFECT.md).
- **Their machine savings are withdrawn.** The hung probe left the controller's latency sense nearly blind, so it released machines it would not have released with a working probe. Set 3 shows that: with a working probe and latency near the declared 500 ms SLO, Omni-Compass kept all 6 machines. The decision trail per run is now in every job log (set 4).

Live levers under the nervous system:

- **What acted:** right-sizing, cold start, batch pacing, agent containment and the cooling connector.
- **Authority:** each lever acted on real Kubernetes only inside the authority the nervous system granted.
- **Kill switch:** it restored every lever, including from a fresh process.
- **Result:** 19 of 19 checks passed (results/live/LIVE_LEVERS_2_NERVOUS.txt, first pass LIVE_LEVERS_1.txt).
- **Identity:** least-privilege, with `kubectl auth can-i` receipts.

Shadow pilot kit, live: a read-only identity ran for 600 s and logged 40 decisions, with 0 writes (results/live/LIVE_SHADOW_1.txt). This is the kit a customer runs first.

Safety and correctness

- **Safety shield:** tested on 2,000,000 random and adversarial inputs, with zero invariant violations, idempotent, never inventing an action, and intervening minimally (tests/test_shield_properties.py). The test found two real bugs, both fixed.
- **C++ twins:** the C++ shield and the C++ closure law match Python exactly, over 300,000 adversarial shield cases and every recorded closure decision.
- **Endurance:** the C++ engine ran 100,000,000 decisions with no failure, at about 2.3 microseconds per decision.
- **Fail-safe:** after repeated failed decisions, control returns to the native autoscalers.
- **Reproducibility:** a clean copy of the delivered zip reproduced every held-out result byte for byte.

What is not claimed

- **No production or customer deployment yet.** The next step is the shadow pilot (docs/PILOT_KIT.md), which is read-only.
- **Live runs are small.** They use kind on CI machines, and power is modelled, not metered.
- **Hardware levers are not proven.** CPU-frequency and power caps need real servers, and cooling was exercised against a stand-in controller.
- **Parked machines in the public cloud.** A parked cloud machine still bills, so the warm-reserve energy saving applies to owned hardware.
- **Single-cluster energy against the tightest packers is roughly a tie.** There Omni-Compass wins on response time and stability.
- **Some gaps cannot be closed.** No controller, even one with perfect foresight, can match both the tightest packer's machine-hours and the calmest autoscaler's machine churn (tuning/bound.py).

- **Not modelled in the plant:** variable boot times and pod-eviction cost. Fragmentation is modelled (whole-pod packing) and is small.

Reproduce

```
pip install -r requirements.txt && python verify.py
python tuning/confirmatory.py          # C, one global setting, 100 scenarios per workload, Holm-corrected
python tuning/confirmatory.py --coord  # the same with nervous-system coordination, fresh seeds
python tuning/planetlab_league.py <planetlab-workload-traces/20110303>
python tuning/site_league.py --heldout
python tools/protocol_bench.py 100
live: push a commit whose message contains [reps], [levers] or [shadow]
```

Evaluation and simulation use only. Copyright (c) 2026 The Omni-Compass LLC. Commercial use, commercialization or monetization of any part of Omni-Compass requires a signed, paid Omni-Compass Enterprise License from The Omni-Compass LLC. Patent applications, copyright registrations and trademark applications filed in the United States. See LICENSE and NOTICE at the root of this repository.

CHAPTER 56

Due Diligence

Answers reference the Claims Register (C-numbers) and the Technical Manual.

Engineering

Does the mathematics hold? Equations (1) to (8), Propositions 1 to 3 and their proofs are in Manual Chapters 2 and 3. The 500/500 CONVEY/CERT result is a property of the controller (Proposition 3), shown by counterfactuals: the controller aimed at the wrong basin also scores 500/500. **Is Python the same as C++?** Yes (C1, C2). The verifier also builds a deliberately mutated C++ governor and confirms the parity test rejects it. **Can every number be reproduced?** `python verify.py` rebuilds the C++, reruns parity, reruns the 100,000,000-decision soak and compares it with the recorded result, replays held-out scenarios and checks the pre-registered SHA-256 hashes (program fingerprints are additionally checked when the Python version matches the recorded one).

Operations

Better than what we run today? Compared against a documented-behaviour reference model of Kubernetes autoscaling (HPA tolerance 0.1, 300 s scale-down stabilization, HPA targets 0.5 to 0.8; simplified Cluster Autoscaler with 10-minute unneeded time, 0.5 utilization threshold, 10-minute delay after scale-up): C8 to C12, including where Omni-Compass is worse. The reference model is not the upstream controllers (C12b); running the upstream controllers against the same scenarios is the next baseline step. **On real traffic?** Not yet. Results use a synthetic stack model. Replay of published production traces and the pilot protocol are the next evidence steps (C14). **What happens when it is wrong?** Observe mode changes nothing (C6). The reset returns control to the native managers at the next interval (omni_kill arm). The shield blocks actions that violate I1 to I5 (C7). **Will it wear hardware?** Machine start/stop cycles, power-cap travel and thermal travel are measured for every arm (C12, Manual Chapter 8).

Security

What authority does it hold? Capacity, replicas, power caps, rollback authorization and routing, only in AUTOPILOT, only through the shield. **Can it expand capacity during a security block?** No: invariant I1 is enforced before execution. **Does it replace encryption, identity or policy engines?** No. Those components are retained (fleet model, security role).

Finance

What does it save? Energy per run versus current autoscaling (C8). Fleet-scale figures are modeled (C13); the governance saving comes from reduced idle and padded capacity, not from removing the decision components' own consumption. **What does it cost to run?** C5.

Adoption

How is it introduced without risk? Observe, then shadow on production telemetry, then one control loop at a time under the reset (docs/PILOT_PROTOCOL.md).

Referees

Was it tuned on the test data? No. Law, shield and baselines were frozen and fingerprinted before the held-out seeds 346410161 and 360555127 (results/PREREGISTRATION.json). **Where does it fail?** Backlog violations against current autoscaling (C11); the engine-dynamics ablation (Manual Chapter 8); open obligations (Manual Chapter 10).

Evaluation and simulation use only. Copyright (c) 2026 The Omni-Compass LLC. Commercial use, commercialization or monetization of any part of Omni-Compass requires a signed, paid Omni-Compass Enterprise License from The Omni-Compass LLC. Patent applications, copyright registrations and trademark applications filed in the United States. See LICENSE and NOTICE at the root of this repository.

CHAPTER 57

License and Commercial Terms

The software and this manual are licensed under the Omni-Compass Evaluation License (LICENSE): evaluation and simulation use only. Everything else - commercial use, production use, operating any system beyond evaluation, redistribution, a hosted or managed service, incorporation into a product or service, or using the software or its results to build a competing product - requires a written Omni-Compass Enterprise License signed by The Omni-Compass LLC and paid for. Patent applications, copyright registrations and trademark applications covering the Omni-Compass engine, its mathematics and its software have been filed in the United States by The Omni-Compass LLC. No patent or trademark license is granted for any other use. Contributions are accepted only on the terms in `.github/CONTRIBUTING.md`, which assign their rights to The Omni-Compass LLC.

CHAPTER 58

Licensing: Questions and Answers

Plain answers to the questions buyers, engineers and their lawyers ask. The binding terms are in `LICENSE`; where this page and `LICENSE` differ, `LICENSE` governs. Every declaration is in `DISCLOSURES.md`.

Is Omni-Compass open source? No. The source is published so that it can be evaluated, reproduced and audited. It is licensed under the Omni-Compass Evaluation License 1.0 (`LicenseRef-OmniCompass-Evaluation-1.0`), which permits evaluation and simulation only.

What may I do without a commercial license? Read the code; run `verify.py`, the simulations and the benchmarks; reproduce the published results; run Omni-Compass in watch or shadow mode, or in test, on systems you own or control, for the purpose of evaluating it.

What needs the Omni-Compass Enterprise License? Any commercial use, commercialization or monetization; production use; operating any system beyond evaluation; offering it as a hosted or managed service; redistributing it; or incorporating it, or any part or derivative of it, into a product or service. The Enterprise License is a written agreement signed by The Omni-Compass LLC and paid for.

How is the Enterprise License priced? Against the measured gain on the customer's own paired receipts (native against Omni-Compass on the same system, the same load and the same clock). Terms are set in each signed agreement; no price stated anywhere in this repository is an offer.

Are patents involved? Patent applications, copyright registrations and trademark applications covering the Omni-Compass engine, its mathematics and its software have been filed in the United States by The Omni-Compass LLC (`PATENTS.md`). The evaluation license grants no patent license beyond evaluation.

May I use the name or the compass rose? Only to refer to Omni-Compass accurately (`TRADEMARKS.md`).

May I contribute? Contributions require the contributor license agreement (`.github/CLA.md`).

Does Omni-Compass send data anywhere? No. The software contains no telemetry and makes no network connection of its own except to the systems an operator points it at (a Kubernetes API, `nvidia-smi`, a meter command).

What do scanners report? The license is declared in machine-readable form for Black Duck, FOSSA, Snyk and REUSE (`pyproject.toml`, `REUSE.toml`, `.fossa.yml`, `.snyk`, `sbom/`); third-party components are listed in `THIRD_PARTY_NOTICES.md`.

Can the terms change? Yes. The Omni-Compass LLC may change these terms, the software and every document at any time; a signed Enterprise License governs its own term (`DISCLOSURES.md`, section 5).

Who do I contact? The Omni-Compass LLC, www.omni-compass.com.

Evaluation and simulation use only. Copyright (c) 2026 The Omni-Compass LLC. Commercial use, commercialization or monetization of any part of Omni-Compass requires a signed, paid Omni-Compass Enterprise License from The Omni-Compass LLC. Patent applications, copyright registrations and trademark applications filed in the United States. See LICENSE and NOTICE at the root of this repository.

CHAPTER 59

Third-Party Notices

Omni-Compass uses the following third-party components. Each remains under its own license; nothing in the Omni-Compass Evaluation License changes those terms. Installed versions are those resolved from `requirements.txt`.

Component	Use	License
NumPy	numerical arrays	BSD-3-Clause
Matplotlib	charts in the reports	Matplotlib License (PSF-based, BSD-compatible)
ReportLab (open-source edition)	the printable book and reports	BSD-3-Clause
Numba	compiled inner loops of the motion-axis plant	BSD-2-Clause
PyTorch (on GPU machines only, not installed by <code>requirements.txt</code>)	the GPU workload of the benchmark	BSD-3-Clause
<code>tests/third_party/hpa_independent.py</code>	an independently written HPA reference used unmodified in tests	its own terms, stated in the file
Liberation Serif and Liberation Sans fonts (embedded in the PDF book)	typesetting	SIL Open Font License 1.1
DejaVu Sans and DejaVu Sans Mono fonts (embedded in the PDF book)	typesetting	Bitstream Vera / DejaVu license (free)

Kubernetes, kind, `kubectl`, the NVIDIA driver and `nvidia-smi` are not distributed with Omni-Compass; it calls them where an operator has installed them. Their names are the property of their owners (`DISCLOSURES.md`, section 1).

Evaluation and simulation use only. Copyright (c) 2026 The Omni-Compass LLC. Commercial use, commercialization or monetization of any part of Omni-Compass requires a signed, paid Omni-Compass Enterprise License from The Omni-Compass LLC. Patent applications, copyright registrations and trademark applications filed in the United States. See LICENSE and NOTICE at the root of this repository.

CHAPTER 60

Repository Standards

The files the leading repositories carry, the open ones such as Kubernetes and the source-available, commercially licensed ones such as HashiCorp Terraform, Elastic, MongoDB, Redis, CockroachDB and Sentry, and where Omni-Compass carries each. Omni-Compass is published for evaluation and simulation; everything else is licensed (LICENSE, LICENSING_FAQ.md).

What a leading repository carries	Omni-Compass
License, machine-readable and in full	LICENSE, LICENSES/LicenseRef-OmniCompass-Evaluation-1.0.txt, SPDX headers, REUSE.toml
Licensing questions in plain words	LICENSING_FAQ.md
Notices and third-party components	NOTICE, THIRD_PARTY_NOTICES.md
Disclosures and disclaimers in one place	DISCLOSURES.md
Patents and trademarks	PATENTS.md, TRADEMARKS.md
Contributor license agreement	.github/CLA.md
Contributing, code of conduct, governance, ownership	.github/CONTRIBUTING.md, .github/CODE_OF_CONDUCT.md, .github/GOVERNANCE.md, .github/OWNERS, .github/CODEOWNERS
Security policy and contacts	SECURITY.md, .github/SECURITY_CONTACTS
Support	.github/SUPPORT.md
Changes and roadmap	CHANGELOG.md, ROADMAP.md
Citation and software metadata for archives and crawlers	CITATION.cff, codemeta.json, pyproject.toml (keywords, URLs)
Software bill of materials (SPDX and CycloneDX)	sbom/omni-compass.spdx.json, sbom/omni-compass.cdx.json
License scanners (Black Duck, FOSSA, Snyk, ScanCode, REUSE)	pyproject.toml, .fossa.yml, .snyk, REUSE.toml, SPDX headers
Issue and pull-request templates, including a licensing template	.github/ISSUE_TEMPLATE/ .github/PULL_REQUEST_TEMPLATE.md
Dependency updates	.github/dependabot.yml
Continuous verification	.github/workflows/verify.yml, python3 verify.py
Reproducibility: fingerprinted release, sealed twins, checksummed results	RELEASE_MANIFEST.json, results/SEAL.json, SHA256SUMS.txt in every result folder
Rules written before every benchmark run	docs/GPU_PREREGISTRATION.md, docs/REALMS_PREREGISTRATION.md, docs/K8S_COMPASS_PREREGISTRATION.md
A manual and a printable book	docs/OMNI_COMPASS_MANUAL.md, docs/OMNI_COMPASS_MANUAL.pdf
Every result in one place, with charts	docs/DOSSIER.md

What a leading repository carries	Omni-Compass
A one-file archive of the repository	GitHub's Download ZIP, at any commit (the copyright deposit is release/copyright/)

Files present at the commit this page describes; `python3 verify.py` checks the ones the results depend on.

Evaluation and simulation use only. Copyright (c) 2026 The Omni-Compass LLC. Commercial use, commercialization or monetization of any part of Omni-Compass requires a signed, paid Omni-Compass Enterprise License from The Omni-Compass LLC. Patent applications, copyright registrations and trademark applications filed in the United States. See LICENSE and NOTICE at the root of this repository.

CHAPTER 61

Python, C++ and the Seal

The laws are twinned: each has a Python version and a C++20 version that give the same answers, proven by a parity test on every build (`cmake -S cpp -B cpp/build && cmake --build cpp/build`).

Law	Python	C++	Proven by
core engine	omnicompass/core.py	cpp/src/core.cpp	500 frozen fixtures
governor	omnicompass/adapter.py	cpp/src/governor.cpp	tests/test_cpp_governor_parity.py
safety shield	omnicompass/shield.py	cpp/src/shield.cpp	tests/test_cpp_shield_parity.py
HPA replica law	fleet/harness.py	cpp/src/hpa.cpp	tests/test_cpp_hpa_parity.py
closure law	omnicompass/closure.py	cpp/src/closure.cpp	tests/test_cpp_closure_parity.py
conveyance law	omnicompass/conveyance.py	cpp/src/conveyance.cpp	tests/test_cpp_conveyance_parity.py
nervous system	omnicompass/nervous_system.py	cpp/src/nervous_system.cpp	tests/test_cpp_twins_parity.py
compass and ledger	omnicompass/compass.py, storage.py	cpp/src/compass.cpp	tests/test_cpp_twins_parity.py
GPU governor rules	omni_controller/gpu_governor.py	cpp/src/gpu_rules.cpp	tests/test_cpp_twins_parity.py

The seal (`results/SEAL.json`) holds the SHA-256 fingerprint of every twinned file, written only after every parity test passes. `verify.py` fails, naming the file, if any sealed file changes afterwards. The compass law (`omnicompass/compass.py`) and the two-wire GPU governor are in Python today; their C++ twins are next.

Why two languages. A law that exists in one implementation can hide a bug that is also its behaviour: the tests pass because the tests were written against the same code. Two implementations written separately, in languages with different numerics and different defaults, that agree to fifteen decimal places on five hundred trajectories and on every fixture of the governor, the shield, the HPA law, the closure, the conveyance and the nervous system, are evidence that the text of the law is what both compute. The C++ twin is also the one a buyer would embed where Python does not run (a controller, a firmware, a PLC), and the soak (`results/SOAK.json`, one hundred million decisions, no failure) is run on it. The seal makes the twins' agreement a property of the repository rather than a claim: change a sealed byte and `verify.py` fails, naming the file.

The three fingerprints. The seal fingerprints the twinned laws; `OMNI_V1.json`, `OMNI_V2.json` and `OMNI_V3.json` fingerprint the engine as a whole (law, controllers, muscles, runners: 38 to 40 files); `RELEASE_MANIFEST.json` fingerprints the release (the commit, the engine, the twins, the GPU protocol, the live evidence and the verification receipt) and is rewritten by `tools/release_manifest.py` before every push, under the Python GitHub runs. A reader holding any one of the three can tell whether the files in front of them are the files the results were made on.

CHAPTER 62

The Founder's Working Notes

Where a sentence is a design aim rather than what the code does today, or a figure that is not yet measured, the check column says so. Nothing in this chapter adds a result: every number that is a result names the file it comes from.

What Omni is

OmniCompass is the governor. It is the process running on the processor, and the chip executes it. The card, the replica count, the rack cap, the joint and the feeder are the levers: the muscles. Omni reads the meters, steps a bounded loop, and writes only the lever it is allowed to write. When it stops, it puts that lever back.

It is not the chip, not the card and not Kubernetes. It is the brain on the host, writing through a bolt.

Governor is the machine word. A governor on a steam engine did not build the engine and did not turn the shaft. It watched the speed and moved the throttle so the speed stayed in a band. The plant does the work: the chip, the card, the pods, the job. Omni reads a meter, writes one lever, and puts that lever back.

The loop

The chapter says	The code today	Check
Omni reads the plant's meters and maps them to the state E, U, S	<code>omniconpass/adapter.py</code> , <code>observe_vector</code> and <code>assimilate</code>	matches
It computes the drift with the command at zero, then $\text{raw} = -\text{drift} + K_P \cdot (\sigma - U)$, clipped to the authority before the step	<code>omniconpass/core.py</code> , <code>control_command: u = clip(-f_U(x,t) + KP · (σ - U), ±U_AUTHORITY)</code>	matches
It holds that same u through every RK4 stage and does not rewrite U afterwards	<code>core.py</code> , <code>macro_step</code> : one u per micro step, passed unchanged into all four RK4 stages; U is never overwritten	matches (proved: <code>docs/TRACKING_THEOREM.md</code>)
It writes the clipped command u to the lever	Not today. The frozen live governor runs the engine with $u = 0$ (<code>macro_step(..., target=None)</code>) and uses u only as its convergence signal ($\text{push} = u / U_AUTHORITY$). The lever is set by the allocation law and the nervous system's authority (<code>omniconpass/nervous_system.py</code>).	design aim , see "What to strengthen"
On exit it restores the lever	kill switch: <code>omni_controller/controller.py</code> restore; the GPU bench restores and reads back the start limit; the realm harness checks every knob is handed back	matches

So the step the chapter describes is the engine's own step, and it is proved. Writing u itself onto a lever, through an output map with a gain in watts (or replicas) per unit of u, is a further mechanism. It is not what produced the results

below.

Wiring

A wrist wire moves the wrist. It does not move a finger that has its own tendon unless that tendon is also wired.

The card's clock wire goes to the NVIDIA driver. The replica wire goes to the Kubernetes API. The scheduler's wire goes to the kubelet. Those wires already exist and do not pass through Omni. Omni is an added process, and a tendon is connected to it only if the output map writes that lever and the input map reads it back.

A power-limit write can make the driver drop clocks. That is a side effect on one finger, not the replica nerve, which keeps running on its own path unless a second line is connected. The cluster starts its own wires: the API server accepts a replica write, the scheduler places the pod, the kubelet starts it, the driver moves the clocks, and HPA, if installed, is its own process. Omni joins those nerves. It is a client of the plant, not its origin.

Four realms, five organisms

1. Compute / AI / Cloud: GPU server, Kubernetes, the card, the job.
2. Physics / Robotics / Autonomous: joint, servo, PLC, car, line.
3. Energy / Facility / Industrial: rack, cooling, PDU, battery.
4. Distribution / Specialized: network, feeder, logistics.
5. The whole organism: all 656 canonical muscles, each once, one clock.

A muscle may sit in more than one realm. Every realm stands on the same spine (Kubernetes, machines, GPUs and CPUs, network, storage, observability, security, cooling, electrical distribution), and then has its own domain muscles. The fifth run holds each of the 656 once.

Each organism is shown native, then with Omni, on the same membership, the same seed and the same clock: two receipts. The core does not change when the lever changes.

Check	
Built	<code>realms/</code> (round 3): spine of 190 muscles in all four realms; organisms of 345, 262, 282, 337 and 656 (<code>docs/REALM_MUSCLES.md</code>)
Result	Whole organism +0.1% work per energy, SUPERIOR WITHIN GUARDRAILS; the realms pay in service when they share the spine (<code>results/realms/REALMS.md</code>)
Limit	These are modelled plants (evidence S / rung E2). A catalog row is not a tendon. GitHub can host the compute organism for real (kind, and the GPU on a rented card); facility, machine and grid need their own hosts
Open	The 656 is a working catalog, not a census: <code>docs/realms_study/</code> checks it against the real systems' documentation and industry practice (waves 1 and 2 done; wave 3 maps every row)

Evidence rungs

Rung (chapter)	Class (repository, docs/EVIDENCE_LEDGER.md)	Meaning
E0	design	the written spec
E1	T / V	deterministic tests and proofs: Python against the C++ twin, the tracking theorem
E2	S	simulation on a made plant (the realms, the fleet and GPU models)
E3	L	real software: Kubernetes on kind, no card (sets 22 and 23)
E4	P	a physical meter: the card's own power reading (the GPU bench; first run on an NVIDIA A10, 2026-10-02, <code>results/gpu/run-20261002T082232Z/GPU_REPS.md</code> ; the corrected governor not yet run on a card)

A result does not climb a rung by itself. The four realms are plants; E1 to E4 are how hard the proof is on whichever plant is run.

What blocked the card

The hash did not stop the card. A Python–C++ mismatch can fail a test after a machine has started the job; it cannot stop GitHub from handing out a machine, because the runner is chosen before the code runs. The verify job started and failed (fixed since: it is green), and separately the GPU job never got a machine.

The organisation's and enterprise's settings show no GPU runner and no option to create one (checked 2026-10-02), so no budget could start it. The run therefore moves to a rented card (`scripts/gpu_rented_run.sh`, one command).

The bake-off that exists

Set 23 (`results/live/LIVE_REPS_23.md`, run 36940088922): real Kubernetes, ten pairs, fifteen minutes an arm, the same open-loop work, native against Omni, six kind workers, order rotated.

Gauge	Native	Omni	Change
Response, 95th percentile (ms)	327.3	123.8	−62.2%
Response, 99th percentile (ms)	503.4	169.3	−66.4%
Response, mean (ms)	146.6	80.1	−45.4%
Failed requests	0	0	0
HPA replicas, mean	8.51	5.40	−36.6%
Workers in service	6	4.28	−28.7%
Pods started	4.7	1.7	−63.8%
Pod start wait (s)	18.6	3.8	−79.6%
CPU, service plus Omni (cores)	0.903	0.894	−1.0% (no difference)
Energy, workers still on (Wh, declared model)	159.0	158.7	−0.2% (no difference)

The same requests were sent and none failed. They were answered faster, on fewer copies, at the same bill. Kind leaves every worker powered, so the watt line stays flat. Parking a worker and waking it again was tried; the churn ate the saving (`docs/CLAIMS_REGISTER.md`, C17), so that path is closed on this plant.

Arithmetic, not a result: if a run finishes 5% more work, nineteen runs do what twenty used to ($20 / 1.05 = 19.05$).

The chip

The card is the meter that can move the electric bill, and the meter that can move output.

Energy on the chip is the integral of power.draw. The same pile can be spent two ways:

- **Output:** more finished work on the same bill. The cap stays near where the native arm left it, the draw stays in the same band, and heat gets no new reason to rise.
- **The bill:** a lower enforced limit, the job still finishes, and joules fall.

A company short of cards spends it on output. A company paying for GPU hours with spare capacity spends it on the bill. A company against a building power cap splits it. The bench's primary outcome, work per joule, counts both, and its table prints requests and joules separately, so the receipt shows which way the pile went (`tools/gpu_reps.py`).

For the pile to be real:

- the cap has to bind;
- the write has to be enforced by the driver, and the enforced limit read back;
- the job has to finish;
- the limit has to be restored on exit, so the next native arm is not still capped.

All four are in the bench and its preregistration (`docs/GPU_PREREGISTRATION.md`, amendments 1–3).

The babysitting tax

The electric bill is the small pile. The babysitting tax is the people and tools kept on the clock to watch the levers, reset a cap, and stop the muscles fighting. A card's real cost is usually the card and the hours it sat idle or late, not the electrons.

What fills the seat today is a person: someone sets the cap, someone gets the page when it is left down, someone turns it back after the run or the crash. The product takes that knob and that page.

What we found nobody selling

We found no published product that does all of the following across compute, machines, facility and grid as one governor:

- reads all of their meters into one state;
- holds one clipped command through the step;
- puts each lever back when the run stops, and prints both receipts.

Borg, Twine and OpenShift govern compute. AWS is a cloud. Turbonomic covers the IT stack. Tesla runs the car, the battery and the factory as separate software. Each covers more than one thing as a company, not four realms as one governor.

Omni on such a stack does not replace it: Borg remains the muscle. This is a market observation from public material, not a claim tested here.

The bill

Run the stack native and print the receipt. Run the same stack with Omni and print the receipt. The difference is the pile, and the fee is 20 percent of it.

On a stack that costs \$100, a 10 percent win is \$10. The fee is \$2, and the customer keeps \$8. In general, the fee is $0.2 \times (\text{measured win}) \times (\text{spend})$. A gain and a cut are the same difference, read from opposite sides. Do not add CPU, GPU and babysitting percentages into one bill: the invoice is one pile.

Check: the 10 percent is an example, not a measured win. What is measured so far:

- on Kubernetes, response time -62% at the same energy (set 23);
- in simulation, +0.1% work per energy for the whole organism (realm round 3);
- on a card, nothing yet: the GPU bench is the first real-meter number.

Revenue figures should be computed from the measured win on each customer's own receipt. The per-company table in the working draft (2% of each company's published infrastructure spend) assumed a 10% win everywhere, and several of its spend figures do not match the companies' published capital expenditure. It is left out of the repository until both are sourced.

What to strengthen

1. **Wire the output map.** u becomes one lever, read back and restored, with the start value, the written value and the value after a kill printed on the receipt. The GPU bench already records the start limit, every write and its read-back, and the restored limit. A u -to-lever output map (a gain in watts or replicas per unit of u) is a new mechanism. It would be preregistered and tested as its own arm, first in the realms, then on a card, never mixed into the frozen confirmation run.
2. **Keep the step as it is:** clip before the stage, hold u through RK4, do not rewrite U (proved: docs/TRACKING_THEOREM.md).
3. **Run native against Omni on the same membership for each realm.** Compute is the one GitHub can host for real; the other three are modelled until their hosts exist.
4. **Aim the chip at about 5% more finished work on the same bill,** cap unchanged and temperature no worse, or at fewer joules for the same work. The buyer chooses, and the receipt prints which.

Evaluation and simulation use only. Copyright (c) 2026 The Omni-Compass LLC. Commercial use, commercialization or monetization of any part of Omni-Compass requires a signed, paid Omni-Compass Enterprise License from The Omni-Compass LLC. Patent applications, copyright registrations and trademark applications filed in the United States. See LICENSE and NOTICE at the root of this repository.

CHAPTER 63

History

Earlier states of play, kept whole. Nothing here is current status; the current state is docs/STATE_OF_PLAY.md.

State of play as written on 27 September 2026 (with the update of 28 September)

Superseded for live results. The live evidence below predates sets 19 and 20. The current live table is results/live/LIVE_PAIRED.md; energy on kind is a declared model, not a meter. The first metered test is scripts/gpu_paired.sh (docs/GPU_BENCH.md).

Update, 28 September 2026. Added since this page was written, each with its evidence: - **Live, set 21** (conveyance only when response time needs it): p95 -37%, machines in service -21%, CPU used +26% (unchanged from set 20; its source is still open). results/live/LIVE_REPS_21.md. - **GPU governor bounded** (0.70 share floor, busy gate, 2 s decisions): on the MLPerf-calibrated modelled card, +5.1% and +1.3% work per kJ with p95 within +10% (the old governor failed that guardrail at +31% and +66%). Model only. results/gpu/sim/. - **Speed lock** (opt-in): speed won elsewhere spent on GPU watts, every gauge kept at least 1% faster than without Omni. Model only. results/gpu/sim/pipeline/. - **CPU and GPU on one conserved power budget** (conveyance law over CPU and GPU organs): +1.4% to +5.7% work served against today's fixed caps, never over the site budget. Model only. docs/CONVEYANCE_LAW.md, results/hardware/NODE_EXCHANGE_*.json. - **The manual in the box:** docs/INTEGRATION_MANUAL.md; every gauge: docs/METRICS_CATALOG.md; against what runs today: docs/COMPARISON.md.

This is the whole repository at the commit named in STATE_OF_PLAY_COMMIT.txt. Everything below can be rerun from it.

Run it live on real Kubernetes yourself: on any machine with Docker, kind, kubectl and Python, run bash RUN_LIVE.sh 3. It builds a fresh 7-node cluster per arm, runs native, Omni on top and Omni alone, and writes live_runs/LIVE_REPS.md. GitHub Actions runs the same comparison with a pushed commit whose message contains [reps].

Rerun the whole thing: pip install -r requirements.txt && python verify.py. It must end VERIFICATION: PASS, and it does at this commit.

Best measured results, and where each one comes from

Claim	Evidence	Where
B (Omni on top of each platform): 0 losing cells on held-out scenarios; strictly better on 27 of 28 platform-workload pairs	simulation, pre-registered	tuning/B_FIX.json, tuning/B_TONE_HELDOUT.json
Four-cluster sites: B 0 losses on all 7 platforms; C 0 losses vs GKE, AKS/Karpenter, CAST AI, Spot	simulation, held-out	tuning/SITE_LEAGUE.json
C (Omni alone), one global setting, 400 never-used scenarios: 153 better / 174 equal / 37 worse of 364 cells (Holm-corrected)	simulation, confirmatory	tuning/CONFIRMATORY.json
C with nervous-system coordination, 400 fresh scenarios: 151 / 173 / 40	simulation, confirmatory	tuning/CONFIRMATORY_COORD.json
Controllers fighting each other: lowest of all 8 systems on web (0.07 per day) and four-cluster sites (0.56 per day)	simulation, confirmatory	same file
Real demand, 1,052 recorded PlanetLab machines: 0 losing cells of 91	simulation on real traces	tuning/PLANETLAB_LEAGUE.json
Faults (machines dying, spikes, crash loops, noisy neighbours): fewest pages to a human, fastest recovery	simulation, 500 runs per workload	results/protocol/
GPU power law (engine cap over the MLPerf-measured performance law), against native at equal work: energy -27% to -33%, time over the heat limit -73% to -78%	simulation, held-out	results/hardware/SUMMARY_HELDOUT.json
Site power exchange (conveyance law, manuscript Ch. 29-31): 4 GPU groups under one site budget: 0 minutes over the site limit at 70/60/50% budgets (native 10.5/45.6/106); least energy at every budget; 31-115 fewer backlog minutes than the static split	simulation, held-out; law proved (conservation, Lyapunov, exponential convergence)	docs/CONVEYANCE_LAW.md, results/hardware/SITE_EXCHANGE_HELDOUT_*.json
Live levers under the nervous system: 19 of 19 checks on real Kubernetes; reset restores everything	live, kind	results/live/LIVE_LEVERS_2_NERVOUS.txt
Shadow pilot kit: a read-only identity made 0 writes across 40 decisions	live, kind	results/live/LIVE_SHADOW_1.txt
Safety shield: 2,000,000 adversarial cases, 0 violations (the test found 2 real bugs, both fixed); C++ twin matches	test	tests/test_shield_properties.py
C++ engine: 100,000,000 decisions, no failures, about 2.3 μ s per decision	test	results/SOAK.json

What still loses, measured

- **C, typical response time, web and four-cluster:** about 20% slower than every platform. Omni fills pods to 80%, where the others stop at 70%, and trades response time for energy.
- **C, machine starts and stops:** loses to plain Kubernetes and OpenShift, which hold machines steady and pay for it in machine-hours.
- **C, GPU and batch:** energy 1-2% worse and power and heat margins 1-4% worse than the tight packers (AKS/Karpenter, CAST AI, Spot).
- **Live, set 3, with a working probe:**

- Omni kept all 6 machines. Latency near the 500 ms SLO blocks machine release.
- There is no significant difference from native on anything else, except that C has more waiting pods.
- The earlier live machine savings (sets 1-2) are withdrawn, because the broken probe had blinded the latency sense (results/live/LIVE_REPS_PROBE_DEFECT.md).
- Set 4's decision trail showed why. Latency spikes from pod saturation at load steps kept the service from being clean, and the engine's calm stayed under the machine threshold (0.65 against 0.7). Both are addressed by the attribution gate and the two-way wiring above; the live test is pending (results/live/LIVE_REPS_4_DIAGNOSIS.md).
- **A bound, not a loss:** no controller, even one with perfect foresight, can match both the tightest packer's machine-hours and the calmest autoscaler's churn at once (tuning/bound.py).

Two-way nervous system (new, live test pending)

Manuscript Appendix J names delay and dropout handling and feedback interpretation as the nervous system's job. The live controller now wires both directions (docs/TWO_WAY_NERVOUS_SYSTEM.md).

- **Blind senses are detected against the wall clock.** A hung probe was the root cause of the live set 1-2 error.
 - A blind sense feeds the engine's stale channel.
 - It blocks every give-back.
- **Commands are read back.** Drift feeds the engine's drift_ratio. No new machine release goes out before the last one has landed.
- **The machine organ has its own engine view and a release gate.** Pods move first, and the headroom proof uses the engine's rho.

Tests: 75,000 blind states, 200,000 gate states, 300,000 nervous-system states.

Live sets 5 and 6 could not run. From 03:51 UTC GitHub refused every job, including verify: runner ID 0, no steps, failed within about 2 seconds. That is GitHub declining to start runners, which on a private repository is most often exhausted Actions minutes or a spending limit. It is not a code failure. python verify.py passes locally at this commit.

Map

Path	What it is
omnicompass/core.py, cpp/	six-state engine; C++ twin
omnicompass/closure.py	closure law: forward projection, boundary correction, turning point, muscle tone
omnicompass/nervous_system.py	one engine state grants each organ its authority and envelope
omnicompass/shield.py	safety shield, downstream of everything
omni_controller/	live Kubernetes controller and levers
fleet/, hardware/	simulation plants
tuning/	benchmarks, pre-registrations, results
results/live/	every live run, including the failed and withdrawn ones
docs/history/OMNICOSSPASS_BUYER_EDITIO N.md	the buyer-facing report

The state of play as it stood on 2026-10-06 (moved here whole on 2026-10-07)

In one paragraph (2026-10-06, Omni v1)

Omni v1, confirmed three times. On real Kubernetes, every test ran three times as separate GitHub runs on the frozen engine, ten pairs each (docs/OMNI_V1.md, the six results/live/V1_*.md tables). Work inside the response line **+42% to +48%** in all three runs of the all-four test (18.6 → 27.6 requests a second in run A), confirmed better; the 95th percentile **-57% to -69%** in every run of the steady, wandering, all-four and fault tests, confirmed better; failed requests -12% to -13% where load swings, confirmed better; machines -1.5% to -3.4% at steady load and -15% to -24% on the batch queue, confirmed better, no difference beyond the noise elsewhere; energy a declared model (kind never powers a machine down); beside a noisy neighbour, no difference beyond the noise on every row. **The Omni index on v1: +26.2%** on real Kubernetes (results/OMNI_INDEX.md; only rows confirmed in all three runs count, a row inside the noise counts as zero; Azure and the card join when their v1 runs land). The paragraph below is the earlier engine's reading, kept as the record it was.

Omni-Compass sits on top of Kubernetes and hardware and gets more out of what is already there. On real Kubernetes, the engine before v1 (the live controller at rules 1-8, commit 353903683009), ten pairs per test, native against compass with the same work sent to both: **41.7% more work handled inside the response line on the same machines** (capacity, all four in one run), responses **59-64% faster** at the 95th percentile (steady, wandering, all four, faults), 12% fewer failed requests where load swings, energy equal or lower, and not one pod left without a machine. Machines: up to 2% fewer. Omni gives a machine back only when a paired trial shows the service no slower without it, and on these clusters one machine fewer made requests 30-45% slower in most trials, so it kept them and spent them on speed and work (earlier engines without that check parked 29-36%: docs/history/). All of it together, real machines only: **the Omni index +12.9%** (results/OMNI_INDEX.md).

Omni v1 (docs/OMNI_V1.md): the engine is frozen and fingerprinted (OMNI_V1.json), and every test above runs three times on it, as separate GitHub runs (A the result, B and C the replications); each table reads confirmed better, confirmed worse, no difference beyond the noise, or the runs disagree, and replaces the earlier engine's table as it lands. Running on v1 now: the Kubernetes suite three times, the batch test, the six organisms at 1, 10, 100 and 1,000 copies with the real cluster inside, the big organisms on rented Azure machines, the Azure steady and burst bill tests, CityLearn, and the power grid.

What is shown and what is not. Shown: more work inside the response line on the same machines (34-52% across the three capacity runs on earlier engines) and a faster tail, on real Kubernetes. Not yet shown: an energy or cloud-bill saving on real machines. On kind the machines are containers on one runner, so a machine out of service saves modelled watts, not a metered bill, and Azure's bill did not move on the earlier engine (results/live/AKS_BILL.md); the v1 Azure runs are the test of that. The current card governor has not run on a real GPU.

The card: the first real run (NVIDIA A10, 2026-10-02) used a card controller since replaced; it is not a result to stand on. In simulation the current controller (amendment 12) is clearly faster under an operator's power cap (p95 6-7% faster, time over the line 1 point lower) and saves 0.5-3.7% energy on the card's own firmware with p95 even (results/sim/gpu_two_wire/). The real card runs next: one card, the card inside the six organisms, then eight cards.

The 945 muscles (Omni v2; 656 in v1) are models of real control systems. With the real cluster or the real card inside, they show the mechanism (work the same, energy 0.1-0.2% lower, time over the line lower than native in every organism); they are never counted in the headline.

Test (real)	Work	Speed	Machines	Energy	Source
All four in one run: load up and down one step at a time, 10 pairs	+29%	p95 -62%	-3.6%	-0.3%	results/live/ALL_FOUR.md
Capacity, load rising, 10 pairs	+48%	p95 -50%	-1%	-0.3%	results/live/AMENDMENT_3_RUNS.md
Sets 22-27, same work, 10 pairs each	same, none failed	p95 -55% to -65%	-29% to -36%	-0.1% to -0.5%	results/live/LIVE_REPS_22.md to results/live/LIVE_REPS_27.md
Six organisms, cluster inside, 5 pairs each	requests served +0.5% to +6%	p95 -24% to -40%	same to -2%	-0.1% to -0.6%	results/live/SIX_KUBE.md
Azure AKS, steady load, 5 pairs	same	p95 -20%	same	(billed) same	results/live/AKS_BILL.md

All four in one run is done: every one better and proven in the same run (results/live/ALL_FOUR.md). Running now: demand that wanders (up, spike, down, back up, idle, 10 pairs), the Azure burst bill test, and the 1,000-copy grid of the six organisms on Lambda.

The engine now, and what runs on it (2026-10-05)

The live controller is frozen at rules 1-8 (docs/K8S_COMPASS_PREREGISTRATION.md, amendments 1-8): rules 5-8 were added today (a pinned gauge is not a steady demand; coasting; cruise; the emergency brake). Every result above was measured on an earlier version of the controller, and each names its commit. So that every number comes from one engine, the whole Kubernetes suite runs again on the frozen engine: all four in one run, demand that wanders, the steady same-work set, the fault and fairness tests, the batch test, and the six organisms at every size; the Azure burst bill test and the two big organisms on rented Azure machines too. The card harness changed only in how it judges (amendment 11 of the GPU preregistration), so the single card and the eight cards run once, on this engine, on Lambda.

Back on the frozen engine so far, ten pairs each:

- **Faults** (results/live/FAULTS.md): p95 -60%, mean response -35%, time over the line -23%, pending pods -77%. Recovery is faster after every fault: machine down 42 s against 70 s, a blind probe 54 s against 66 s, a runaway pod 98 s against 105 s, a spike the same.
- **Fairness** (results/live/FAIRNESS.md): a noisy neighbour on the same cluster. Mean response -15%, nothing worse, and the neighbour's own app no worse. The earlier engine scored -0.3% on the index here; this engine scores +5.6%.

Independent simulator: CityLearn round 1 split (the 2022 battery districts better on bill -4% to -11%, electricity -9% to -17%, carbon -7% to -14%; the 2020-2021 water-tank districts worse on peaks and ramping), the fix (only the electric batteries are steered), and round 2 on every untouched district (docs/CITYLEARN_PREREGISTRATION.md).

Measured on real systems: the newest set, Omni-Compass against Kubernetes as it runs today

Set 24 (2026-10-02, commit c908054) repeats it again: machines in service -31.6%, p95 -60.1%, p99 -64.1%, HPA replicas -38.6%, 0 failed requests, total CPU including Omni-Compass's own -1.8% (not significant) (results/live/LIVE_REPS_24.md). Set 23 before it: p95 -62%, replicas -37%, pod starts -64%, 0 failed requests, no energy or total-CPU difference (results/live/LIVE_REPS_23.md). The set-22 table below stands as first measured.

Set 22 (results/live/LIVE_REPS_22.md): 10 paired repetitions on real Kubernetes (kind), each pair on one machine, Kubernetes with its autoscaler alone against the same Kubernetes with Omni-Compass on top. The load is

sent at a fixed rate, so both arms were given **the same work**.

Result	Kubernetes alone	With Omni-Compass	Change (95% interval)
Energy, parked machines still on at idle power (declared model, no meter)	160.3 Wh	160.1 Wh	-0.1%, no difference
Response time, 95th percentile	407.9 ms	158.8 ms	-61% (proven)
Response time, 99th percentile	639.4 ms	245.1 ms	-62% (proven)
Response time, mean	179.9 ms	98.4 ms	-45% (proven)
Failed requests	0	0	equal
Pods waiting to start, pod-minutes	0.265	0.025	-91% (proven)
Replicas, mean	8.93	6.91	-23% (proven)
Machines in service, mean (all stayed powered)	6	4.14	-31% (proven)
CPU used by the service	1.036 cores	0.957 cores	-7.6% (proven)
Omni's own CPU (its controller and every command it ran)	0	0.070 cores	+0.070 (proven)
CPU used, service and Omni together	1.036 cores	1.026 cores	-0.9%, no difference

- **Same work, much faster answers**, with no failed requests and far less waiting.
- **No energy saving is shown on kind**. Every machine stays powered; energy is a declared model, and counted at the idle power a parked machine really draws it is unchanged.
- **No CPU saving once Omni's own cost is counted**. The service used 7.6% less CPU; the controller spent almost all of it. Cutting the controller's cost is the next improvement.
- The reset restored every setting in every run.

Kubernetes sets 25 and 26 (2026-10-02, 10 paired repetitions each, equal work)

Run	Machines in service	p95	Failed	Total CPU incl. Omni's own	Receipt
Set 25, the engine's allocation law	-32.3%	-57.3 %	0 / 0	-0.6% (not significant)	results/live/LIVE_REPS_25.md
Set 26, the engine's allocation law	-35.8%	-55.4 %	0 / 0	-1.5% (not significant)	results/live/LIVE_REPS_26.md
Set 26, the compass law in the live controller	-17.2%	-64.8 %	0 / 0	+1.0% (not significant)	same; label by the preregistered rule: better on machines within the band
Set 27, the engine's allocation law	-36.6%	-53.1 %	0 / 0	-0.0% (not significant)	results/live/LIVE_REPS_27.md
Set 27, the compass law aligned with the GPU governor	-15.9%	-65.5 %	0 / 0	+0.2% (not significant)	same; label by the preregistered rule: better on machines within the band

Set 27 (running): the compass in the live controller reads the service as the corrected GPU compass does (mean response time, center 0.4), against native and the allocation law (docs/K8S_COMPASS_PREREGISTRATION.md).

Measured on a real GPU: the card's own meter (evidence class P)

First confirmation, NVIDIA A10 on Lambda, 2026-10-02
 (results/gpu/run-20261002T082232Z/GPU_REPS.md, 10 paired repetitions × native / watch / Omni, 600 s each, frozen at commit c908054, checksums verified). Wire check 7 of 7; 2,144 writes, none refused, every one read back, every arm ended at the start limit; watch equal to native.

Gauge	Native	Omni	Change (95% interval)
Work per energy (requests per kJ)	50.79	52.62	+3.6% (+2.7% to +4.5%), proven
GPU energy	69,180 J	66,790 J	-3.5%, proven
Requests served / not served	3,514 / 0	3,514 / 0	equal
Response time, 95th percentile	510 ms	809 ms	+58.5%, worse, proven

Result, by rule: ENERGY IMPROVEMENT WITH SERVICE TRADEOFF (the p95 guardrail of +10% failed). The six organisms with the same card inside (results/hil/run-20261002T082232Z/HIL.md, 3 repetitions each): the card's work per energy +1.6% to +2.8% in every organism, the same requests, its p95 500 to about 600-935 ms. The cause, from the card's own samples, was wiring in the governor (amendment 6 of docs/GPU_PREREGISTRATION.md): busy bursts served at 736-768 MHz against 861-889 MHz on its own. Corrected (amendments 6 and 7); the corrected governor has not yet run on a card.

Simulated (models: they show the mechanism, not a measurement)

Result	Where
GPU governor with share floor and busy gate (one wire, the power limit), MLPerf-calibrated card: +5.1% and +1.3% work per kJ, p95 within +10%	results/gpu/sim/after
Two-wire GPU card under the compass law, corrected governor (amendments 6-7), 10 seeds and 10 fresh seeds, geometric means: service profile work per energy +6.9% / +3.8% , energy -6.4% / -3.7%, p95 -5.9% / -2.3% (faster), time over the line -0.03 / -0.04 pp; batch profile +8.1% / +4.2%, p95 +7.0% / -2.3%; the one-wire governor +0.1%; both wires restored every seed	results/sim/gpu_two_wire/RESULT.md, fresh/
The six organisms (Compute 345, Physics 262, Energy 282, Distribution 337, the four stacked 1,226, the whole tower 656), native against the compass law on every muscle, 1,000 paired runs at 1× and at 10× size: work per energy +0.20% to +0.30%, energy -0.21% to -0.32%, work -0.01% to -0.02%, time over the service line +0.19 to +0.27 pp in every cell (band first not held) , every knob handed back. 100× and 1,000× are running	results/scale/GRID.md
Speed lock (speed won elsewhere spent on GPU watts)	results/gpu/sim/pipeline/
CPU and GPU on one conserved power budget: +1.4% to +5.7% work served against a fixed cap, never over the budget	results/hardware/NODE_EXCHANGE_*.json, docs/CONVEYANCE_LAW.md
GPU groups sharing a site budget: 0 minutes over the budget	results/hardware/SITE_EXCHANGE_HELDOUT_*.json
Platform leagues, faults, PlanetLab traces, stack benchmark	tuning/, results/protocol/, results/ (see docs/history/BENCHMARK_REPORT.md)
The 656-muscle tower as organisms , round 3 (preregistered, 10 seeds; every realm carries the shared spine; Omni as the shipped controller commands): the whole tower native against one governor on top, work per energy +0.1%, SUPERIOR WITHIN GUARDRAILS ; inside the realms the spine costs service: Energy +0.2% with +1.9 pp violations (tradeoff), Compute 0.0% (+2.1 pp, not established), Distribution -0.1% and Physics -0.7% (WORSE). Rounds 1 and 2 kept, superseded	results/realms/REALMS.md, docs/REALM_MUSCLES.md

Evaluation and simulation use only. Copyright (c) 2026 The Omni-Compass LLC. Commercial use, commercialization or monetization of any part of Omni-Compass requires a signed, paid Omni-Compass Enterprise License from The Omni-Compass LLC. Patent applications, copyright registrations and trademark applications filed in the United States. See LICENSE and NOTICE at the root of this repository.

Back Matter

BACK MATTER

Glossary

Term	Meaning
Antagonist pair	the up side and the down side of a muscle's control, each with its own gain
Authority	how far the nervous system lets an organ move this decision
Band	a reading's or a knob's safe range
Compass	the band seen as a position from 0 (calm) to 1 (the line), with its bottom in the middle
Confirmed better / confirmed worse	the same sign in all three runs (A, B, C), every 95% interval clear of zero
Cover	a knob's hard range; every write is clipped to it
Cushion	the 5% at each edge of the compass, 10% in all
Fail up	the return to full capacity the moment service crosses the wall or a sense goes blind
Governor	the process that reads, decides and writes; Omni-Compass
Muscle	any machine, service or controller Omni-Compass reads and sets
Native	the system as it runs without Omni-Compass
No difference beyond the noise	an interval over zero in at least one run: native and omni could not be told apart; the result, stated as such
Off the clock	a repetition in which either arm ended more than 5% of its window late; marked, the pairing stands
Organism	a set of muscles run together on one clock
Omni index	the one combined number: the geometric mean of every real category's work, speed, machines and energy ratios, each read by the three-run rule
Omni v1, v2, v3	the frozen engines, fingerprinted file by file; a result belongs to the engine its commit carries
Plug	the two-way connection to one muscle: read, write, read back, restore
Pool (Azure)	one machine family under Azure's autoscaler with its own ceiling; several pools make the fleet
Profile	a named set of the compass's settings for one kind of work: service (the default) or batch
Race	full speed at once while work waits, so a burst is never served slowly
Speed floor	the clock a card reaches on its own while busy, learned from its own meter; the governor never sets the ceiling under it
Receipt	the paired record of native against Omni-Compass for one run
Snapshot	a knob's value read once before the first write; the restore point
The runs disagree	runs clear of the noise point different ways; the test is unstable on that measure and is looked into
Wire check	the test that proves every wire follows, reads back and returns before anything runs
Work per energy	work done divided by energy used; the primary outcome
Work inside the line	the requests, transactions or messages a second answered within the response line; the product number on every real stack

Term	Meaning
Pedals	idle, gas, brake and reset: the governor's four ways of moving capacity under autopilot
Cruise	every machine in service while work waits for a place, held without second-guessing until the line has been empty for two decisions
Emergency brake	the work done and demand at zero: straight to the floor in one move, every safety check still holding
Kill switch	security only: one switch in a human hand that turns every governor on the machine off at once; never the reset
Lease	the record every governor renews each decision; a lease older than its limit, or a dead process, makes the watchdog hand back for it
Watchdog	the service beside the governors that hands back for any governor that died or hung without doing it itself
Verdict	the paired trial on the muscle itself that allows a slow knob to move only where the cost is within 2%; "left native" where no step passes
Slack gate	v3's rule: a motion axis busy more than half the time at full speed keeps its speed native
Tuning case	the one workload, cell or swarm the rules were fitted on; shown in every table and never counted
Untouched case	a workload, cell or swarm the rules were applied to unchanged; the judged rows
Knee	nine tenths of a system's measured capacity, where its queue begins to grow; the broker's native was set there by design
Amendment	a change to a preregistration after it was written, with its time and reason; one made after a result was seen says so
Dwell	the seconds after a knob moves during which the law will not move it back
Notch	the smallest step a knob moves in one decision (one consumer, 8 MB of cache, one tap)
Evidence class P, L, S	a physical meter; live software with no meter; a model of ours
Handed back	the knob read back at the operator's value at the end of an arm; every table says whether every arm was
Off the clock	see above; a repetition whose arm ended more than 5% of its window late
Own cost	the governor's CPU (its process and every command it ran) as a share of one core over the window, from its audit; a few thousandths of a core at every size
Smoke run	one repetition that exercises a harness end to end before any counted run; never counted, always recorded
L + S	an organism run with a real cluster inside: the cluster is L, the organism around it S
Working set	the records a workload actually touches over a window; a cache benchmark measures nothing unless the working set reaches the cache
Key space	the records a benchmark may ask for at one notch; in the MongoDB test the knob's load, stepped one notch at a time
Uniform draw	every record in the key space equally likely, so the working set is the key space; YCSB's shipped zipfian draw concentrates on a few thousand hot records whatever the key space
Ordered keys	records named in load order, so a run that asks for record n finds record n; a run with hashed keys against an ordered load asks for records that do not exist
Full-cache gate	the MongoDB and Redis rule: slow reads grow the cache only while the cache is full (nine tenths used), because a miss in a cache with room to spare is not the cache's to mend
Buffer pool	InnoDB's cache of table and index pages, the one memory setting every MySQL operator sizes; resized online by the server in whole chunks
Chunk	the unit in which MySQL resizes its buffer pool (128 MB as shipped); the knob's notch on that stack

Term	Meaning
Page cache	the operating system's own cache of file contents; where the data fits in it, a storage-engine miss is a read from memory and a decompression, not a disk read, and the gain available to the cache knob is smaller

APPENDIX

Appendix A. Command Reference

Task	Command			
Verify everything	python3 verify.py			
Six organisms, simulated	python3 tools/run_scale.py --runs N --scale K --out <dir>			
The full ladder	bash scripts/scale_ladder.sh			
Two-wire card, modelled	python3 tools/run_gpu_card.py <dir> [fresh]			
GPU wire check	sudo python3 tools/gpu_wire_check.py --gpu 0			
GPU, whole benchmark in one command	sudo nohup bash scripts/gpu_rented_run.sh > run.log 2>&1 &			
The six organisms with the real card inside	python3 tools/run_hil.py --out <dir>			
GPU governor, two wires	`sudo python3 -m omni_controller.gpu_compass --mode watch	cap ...`		
Kubernetes, watch	kubectl apply -f deploy/install /omni-compass.yaml			
Kubernetes paired runs	scripts/kind_paired.sh, then tools/live_reps.py			
OFF (Kubernetes)	touch /tmp/omni.kill			
OFF (GPU)	sudo touch /tmp/omni-gpu-kill			
The whole harness OFF, ON, status, watchdog	`python3 tools/omni_switch.py off	on	status	watchdog`
Seal check	python3 tools/seal.py --check			
Which engine a checkout or a commit carries	python3 tools/omni_version.py [--commit <sha>]			
The three-run table (Kubernetes, Azure)	python3 tools/confirm_abc.py "<title>" raw/run-<A> raw/run- raw/run-<C> --out results/live/V3_<TEST>.md			
The three-run table (database, messaging, cache, swarms, robots, grid, buildings)	tools/pgbench_abc.py, tools/kafka_abc.py, tools/redis_abc.py, tools/swarm_abc.py, tools/mujoco_abc.py, tools/pandapower_abc.py, tools/citylearn_abc.py, the same arguments			

Task	Command			
The six organisms with the cluster inside	python3 tools/six_kube_report.py <raw dirs> --out results/live/V3_SIX_KUBE.md			
The grid of copies and runs	python3 tools/grid.py (reads results/scale/receipts/)			
The Omni index	python3 tools/omni_index.py			
The database benchmark, one run	python3 tools/run_pgbench.py --out <dir> (workflow pgbench)			
The message broker, one run	python3 tools/run_kafka.py --setup, then --workloads light,heavy,burst --out <dir> (workflow kafka)			
The cache, one run	python3 tools/run_redis.py --setup, then --workloads small,large,burst --out <dir> (workflow redis)			
The database's storage-engine cache under YCSB, one run	python3 tools/run_ycsb.py --setup, then --workloads b,c,f,burst --out <dir> (workflow ycsb)			
The database's buffer pool under sysbench, one run	python3 tools/run_sysbench.py --setup, then --workloads read_on ly,read_write,update_index,bur st --out <dir> (workflow sysbench)			
The robustness test	workflow robustness: scenario kill or long, reps, duration_s; the own-cost table python3 tools/own_cost.py <raw run dirs> --out results/live/V3_OWN_COST.md			
The drone swarms, one cell	`python3 tools/run_swarm.py --cell short	mixed	long	all --out <dir> (workflow swarm`)
The independent simulators	python3 tools/run_citylearn.py, run_pandapower.py, run_mujoco.py (workflows citylearn, pandapower, mujoco)			
Azure, the bill	workflow aks-metered: reps, arms, duration_s, load_steps, worker_pools, pool_prices, hpa_max			
The big organisms, detached	workflow big-organism-detached: mode start / collect / survey, organisms, size, duration_s, commit, max_hours			
Archive a finished run with the code	add its id to .github/archive_request.txt and push (workflow archive-run)			
Verify before a push, as GitHub runs it	python3 tools/release_manifest.py under Python 3.12: 0 FAIL			

APPENDIX

Appendix B. File Map

Path	Contents
omnicompass/	the engine, governor, nervous system, shield, compass, conveyance law, the compass
omni_controller/	the Kubernetes controller and muscles; the GPU governors (one and two wires)
realms/	the 945-muscle catalog, the plant models, the realm harness, the compass on every muscle, the modelled card
tools/	benchmarks, receipts, the wire check, the scale ladder, manifests, seals
scripts/	one-command runs (GPU, Kubernetes, ladder)
deploy/	container image, install and permission files
cpp/	the C++20 twins
results/	every published result, with its raw files and checksums
results/live/V1_*.md, V3_*.md (+ .json)	the three-run tables, one per test, named by the engine they ran on
results/live/raw/run-<id>/	every archived run's files, one SHA-256 manifest per repetition
results/scale/GRID.md, results/realms/REALMS.md, results/OMNI_INDEX.md	the grid of copies and runs, the modelled realms, the one combined number
OMNI_V1.json, OMNI_V2.json, OMNI_V3.json	the frozen engines' fingerprints
docs/	this manual, the preregistrations, the evidence ledger, the theorem, the realm study
docs/OMNI_V1.md, OMNI_V2.md, OMNI_V3.md	what each engine is and every result read on it
docs/REGISTER.md, docs/PROOF_PROGRAM.md	every muscle, every benchmark run and every benchmark still to run; the program to full size
.github/workflows/	every benchmark as GitHub runs it: benchmark-reps , six , six-kube , big-organism , big-organism-detached , aks-metered , pgbench , kafka , redis , ycsb , sysbench , robustness , swarm , citylearn , pandapower , mujoco , archive-run , verify
tests/, verify.py	every unit and property test, the shield's adversarial cases, the Python-against-C++ agreement; verify.py runs them all and must end with its pass line
tools/omni_version.py, tools/release_manifest.py, tools/layout_check.py	which engine a checkout or commit carries; the release manifest and verification as GitHub runs it; the check that every path the documents name exists
tools/confirm_abc.py and its siblings (pgbench_abc.py, kafka_abc.py, redis_abc.py, ycsb_abc.py, sysbench_abc.py, swarm_abc.py, mujoco_abc.py, pandapower_abc.py, citylearn_abc.py)	the three-run tables, one tool per kind of raw record, each checking the engine of every run it reads
tools/omni_index.py, tools/dossier.py, tools/own_cost.py	the one number from the tables; the dossier from the tables; the governor's own cost from the archived audits

Path	Contents
<code>tools/legal.py</code>	the legal notice every generated report carries at its head and foot
<code>docs/book/</code>	the builder of this manual's PDF (<code>build_book.py</code>) and the theory chapters bound into it
<code>release/</code>	the copyright deposit and the release notes, printed at named commits
<code>CHANGELOG.md</code> , <code>docs/STATE_OF_PLAY.md</code>	what changed, by date; where everything stands today
<code>docs/*_PREREGISTRATION.md</code>	the rules of every benchmark, written before it ran, with every amendment and its time
<code>docs/DOSSIER.md</code> , <code>docs/dossier/</code>	every result in one place, with its charts, built by <code>tools/dossier.py</code> from the tables
<code>docs/HISTORY.md</code> , <code>docs/history/</code>	earlier states of play, earlier engines' sets and the pre-v1 index, kept whole
<code>LICENSE</code> , <code>NOTICE</code> , <code>LICENSES/</code>	the license and notices

APPENDIX

Appendix C. The Equations in Full

```
(1) dE/dt    = -alpha_E E + beta_int + beta_ext + v_eff
(2) dU/dt    = mu U (1 - U^2) - (dE/dt)/E_max - lambda_U U + u,    |u| <= U_AUTHORITY = 25
(3) dI_U/dt  = (1 - U) - sigma_1 E - delta S - lambda_I I_U
(4) v_eff    = cos(omega_B t / 2) * c * tanh(lambda_0 + lambda_1 (U - 0.5) + lambda_2 S)
(5) Phi(S)   = alpha_s S^2 / 2 + beta_s S^3 / 4 - delta S
(6) dS/dt    = -dPhi/dS = delta - alpha_s S - (3/4) beta_s S^2
(7) dB/dt    = B_dot ; dB_dot/dt = gamma_c delta S - (omega_B / Q_B) B_dot - omega_B^2 B
(8) R_B[n]   = finite-difference audit of (7); never fed back into the state
```

Control (one micro step, zero-order hold across all four RK4 stages):

```
u = clip(-f_U(x, t) + K_P (sigma - U), -U_AUTHORITY, +U_AUTHORITY), K_P = 12, sigma in {-1, +1}
```

The compass (every muscle):

```
p = position of the service reading in its band (0 calm, 1 the line)
F = A tanh((K_P (p - center) + K_D v) / A), K_D >= critical damping
p >= 0.95 => full up force; down side held
knob <= clip(knob + g_side F span, cover)
```

The GPU governor (two wires, omni_controller/gpu_compass.py; one law, amendments 8 to 10):

```
p = (mean response of the last 5 s - S) / (SL0 - S), S = SL0 / 10; p95 >= SL0, a failure, blind =
< > fail up
fail up / race (utilization >= 0.95) => ceiling = top clock, lid = start limit (the card's own set
< tings)
steady under the limit (saturated and drawing >= 0.97 of the limit) => ceiling = the card's own bu
< sy clock,
    lid = start limit (the same watts, no knock-backs)
otherwise:
    ceiling <= clip(ceiling + round(g_side F f_top / 15 MHz) 15 MHz, c_floor, f_top)
```



```

    c_floor = max(f_busy_own x floor, f_top - k* 15 MHz)      (never under the card's own busy cloc
↳ k,                                                         never past the verdict's deepest st
↳ ep k*)
    lid = clip(ceil(1.10 P_busy_own), envelope floor, start limit)
g_up 0.10, g_down 0.0125, center 0.4, floor 1.00

```

The verdict (omnicompass/verdict.py; every slow knob, the GPU ceiling and the Kubernetes machine count):
reference r = the median cost per request at the deepest allowed step, measured fresh, n >= 30 requests
↳ sts
trial t = the median cost at one step past it, n >= 30 requests, under the same traffic
allow the step iff t <= r (1 + a) and t <= c_native (1 + a), a = 0.02
a refused step is not tried again for R decisions (GPU R = 900; Kubernetes 120); the knob never goes
↳ past the
deepest allowed step, so steps cannot add up beyond the allowance

Kubernetes, the compass on the HPA target (omni_controller/controller.py; sets 30 and 31):
target x in [0.6 x_op, x_op] (times the conveyed gain); the up force lowers x (more pods)
p >= 0.95 from load (line breached, every sense live, nothing pending) => x = 0.6 x_op at once
p >= 0.95 from a blind sense or pods waiting for a lost machine => x = x_op (native's own
↳)
fault over (p < center, line clean, nothing pending) => x = x_op at once, neve
↳ r held
machines: one more at once past the wall; one back only when the verdict and the release gate agree
g_up 0.10, g_down 0.02, release force < -0.2

The productivity arithmetic (every receipt; manual section 3.4):
work W equal in both arms; G = C_native / C_omni - 1; the same work needs 1/(1 + G) of the resourc
↳ es;
saving S = G / (1 + G): G = 1/3 <=> S = 1/4

Parameter ranges, defaults and the proof of convergence: omnicompass/core.py, docs/TRACKING_THEOREM.md,
docs/CANONICAL_ENGINE.md.

APPENDIX

Appendix D. Metrics

Every gauge, where it comes from, and whether it is measured or modelled: docs/METRICS_CATALOG.md. The gauges that appear in the three-run tables of this manual, with their source and their direction, are these:

Gauge	Source	Measured or modelled	Direction
work inside the response line (requests, transactions or messages a second answered within the line)	the probe's own records, pgbench's log, the consumers' records, the application's records	measured	higher is better; the product number
response time, mean, median, 95th and 99th percentile	the same	measured	lower is better
time over the response line (share of samples)	the same, against the preregistered line	measured	lower is better
failed requests, lost messages	the same	measured	any increase is WORSE
worker machines in service, mean; node-hours	the cluster's own node records every 15 s	measured (machine-hours on kind; Azure's own count on AKS)	lower is better
compute bill at list price	Azure's own machine count $\times$ list price per family	measured (a real bill)	lower is better
energy (Wh), the standby model	a declared formula over the machines in service (kind)	modelled , said so on every line	lower is better
host CPU busy; host CPU-seconds; CPU-seconds per 1,000 units inside the line	/proc/stat of the real machine under the test	measured	lower is better
Omni's own CPU (cores), mean	the governor's audit (its process and every command it ran)	measured	shown, not judged in the Kubernetes tables; a confirmed loss where it is one
connections held open, consumers running, memory ceiling held	the pooler's, the group's, the cache's own reports	measured	lower is better (the resource held)
cache hit rate; consumer lag	Redis's INFO; the consumers' records	measured	higher is better; lower is better
HPA replicas, pods started, pod start wait	the API server's own records	measured	shown, not judged except pod start wait
organism work, energy, time over its line	the organism's plant models	modelled	by direction; S class
energy a mission, missions a charge (drones)	the simulator's motor constants through a declared model	modelled	lower; higher
cache size held, mean; bytes in the cache, mean; pages read into the cache (MongoDB)	the server's own serverStatus (WiredTiger cache)	measured	lower is better (the resource held); shown, not judged (pages read)
seconds from the kill to every setting back at the operator's (robustness)	the harness polling the cluster once a second from the kill	measured	lower is better; past 60 s WORSE; never, INVALID

Gauge	Source	Measured or modelled	Direction
the governor's resident memory, last ten minutes over the first ten (robustness, long run)	the process table, sampled every 15 s by the harness	measured	a growth over a quarter is WORSE
decisions made of expected; failed decisions (robustness, long run)	the governor's audit	measured	under 95% INVALID; any failed decision shown with its reason
mission time, position error, collisions (drones)	the simulator's own state at every step	measured in the simulator (S class)	lower; under the safe error; any collision voids the cell
handed back (every arm)	the knob read back at the end of the arm	measured	yes in every arm, or the run is invalid
off the clock	the organism's own clock against its window	measured	shown; marks the cell

Three rules for reading any gauge in this appendix. A measured gauge is one the system under test, or the host's own kernel, reported about itself; a modelled gauge is one computed from a declared formula over measured inputs, and every table says which it is on the line itself, never only here. A resource gauge (machines, connections, consumers, memory, cache) is always judged lower-is-better, even where the governor's purpose was to add the resource to buy service, so that a gain bought with a resource is paid for in the same table and in the index. And a gauge marked "shown, not judged" never enters a reading or the index; it is there so that a reader can see what the governor did (how often it moved, how many pages it read, how many pods it started), not to be counted for or against it.

APPENDIX

Appendix E. Troubleshooting

Symptom	Cause	Fix
Wire check FAIL 3 up wire (lock)	the driver or VM refuses clock locking	use bare metal or full passthrough; or run the one-wire governor
Wire check FAIL 3 up wire (follows down)	the card ignores the ceiling	update the driver; check for another management agent holding clocks
Wire check FAIL 4 down wire	the power limit is refused or capped by the board	check <code>nvidia-smi -q -d POWER</code> ; run as root
another copy of this test is already running	a second copy was started	wait for the first or reboot; start once
something else is using the GPU	another process holds the card	stop it; the benchmark must run alone
Governor exit 5	another writer changed a knob	find the other controller; Omni-Compass left its value alone
Governor exit 3	a restore did not read back	restore by hand (<code>nvidia-smi -rgc, -pl <start></code>); investigate before rerunning
decision failed (n in a row)	the cluster API is unreachable	turn it OFF; native runs on
Energy saved but p95 slower than native	the card served bursts below its own clock, or the lid sat under its own draw	read section 8.5; in the audit compare <code>telemetry.clock_mhz</code> while busy with <code>native_busy_clock_mhz</code> , and <code>want_w</code> with <code>native_busy_draw_w</code>
Fail-up in most decisions	<code>-slo-ms</code> set too low for the workload, or a stale feed read as blind	set <code>-slo-ms</code> from the workload's own target; check the feed's age
No saving at all, service unchanged	the card is saturated almost all the time (it races), or the service sits above the compass's center	expected on a card with no slack; the saving comes from the quiet stretches
Kubernetes: no machine ever given back	the release gate refuses (its reason is in the audit), or the service is above the compass's center	read <code>node_gate.reason</code> in the audit; check <code>-slo-ms</code> and the latency feed
Kubernetes: a repetition marked OFF THE CLOCK	the machine could not step the organism inside its window	a longer window for that size on that machine; never a faster reading of the same run
Broker: consumers climb to the cover and stay	messages wait at every step (the producer is past the group's capacity)	the benchmark's capacity step sets the peak at nine tenths of the operator's group; check the measured capacity in the run's log
Cache: the ceiling grows while little memory is used	cold misses read as pressure	the full-cache gate is the rule (used $\geq 90\%$ of the ceiling); check <code>used_mb</code> beside <code>maxmemory_mb</code> in the run's log
Cache: the ceiling never grows	the working set fits the operator's ceiling	expected: a cache that fits has nothing for the ceiling to mend, and the row reads inside the noise

Symptom	Cause	Fix
Storage-engine cache: the cache never fills, whatever the key space, and omni reads as a free memory saving	the benchmark's requests concentrate on a few hot records (YCSB's shipped zipfian draw), or the run asks for records the load did not write (hashed keys over an ordered load), so the reads are served from the index	a result that flatters the governor is the first thing to suspect: check "in cache" against the cache size in the per-notch lines; draw requests uniformly over the key space and name the same key order in load and run; a knob with nothing to do is not a saving
Storage-engine cache: the compass only ever reads calm	the line is set far above the stack's own latencies (a 2 ms line over 0.15 to 0.45 ms reads)	set the line on the tuning workload's own figures so that a hit is calm and a full-cache miss is slow, and write the figures in the preregistration before any counted run
Swarm: a cell reads "void"	two drones came within two collision radii in either arm	a task or simulator setting (keep-out, altitude stagger, depot spacing); fixed before any counted run and said in the preregistration
Three-run table heads itself with a warning	a run's commit is not the engine it claims, or the three are not separate runs	run <code>tools/omni_version.py</code> - - <code>commit <sha></code> on each; make the table only from runs on one frozen engine
Archive bot skips a run	the run was not finished when the bot passed	leave the id in <code>.github/archive_request.txt</code> ; the next pass copies it

APPENDIX

Appendix F. Evidence Map

docs/EVIDENCE_LEDGER.md (every claim and its class), docs/CLAIMS_REGISTER.md (what is claimed and what is not), docs/GPU_PREREGISTRATION.md, docs/REALMS_PREREGISTRATION.md, docs/K8S_COMPASS_PREREGISTRATION.md, docs/POSTGRES_PREREGISTRATION.md, docs/KAFKA_PREREGISTRATION.md, docs/REDIS_PREREGISTRATION.md, docs/SWARM_PREREGISTRATION.md, docs/CITYLEARN_PREREGISTRATION.md, docs/PANDAPOWER_PREREGISTRATION.md, docs/ROBOTICS_PREREGISTRATION.md, docs/ROBUSTNESS_PREREGISTRATION.md, docs/YCSB_PREREGISTRATION.md and docs/MYSQL_PREREGISTRATION.md (the rules written before each run), docs/OMNI_V1.md to OMNI_V3.md (every result by engine), docs/REGISTER.md (every benchmark and its file), docs/PROOF_PROGRAM.md (the program to full size with its costs), docs/DOSSIER.md (every result in one place), results/OMNI_INDEX.md (the one number), docs/STATE_OF_PLAY.md (where everything stands), docs/HANDOFF.md (every command in one page).

The map, claim by claim. Each claim this manual makes is listed here with the class of evidence behind it and the document that carries it, so that a reader can go from any sentence of the executive summary to the file that would have to be wrong for the sentence to be wrong.

Claim	Class	Where the evidence is	What would falsify it
The engine's six states stay inside their walls and settle	T / V (proof and test)	docs/FORMAL_STATUS.md, docs/TRACKING_THEOREM.md, tests/, results/SOAK.json	a trajectory leaving its wall, or the C++ twin disagreeing with the Python
No write ever leaves a knob's cover; a blind sense holds; fail-up is immediate	T / V	tests/test_shield_properties.py (two million cases), the plug tests per stack	one case of a write past the cover or a give-back while blind
On a real cluster the governor does more work inside the line, faster, on no more machines, at no more energy	L	the six Kubernetes tables in results/live/, A/B/C each, docs/K8S_COMPASS_PREREGISTRATION.md	a confirmed-worse row on work, speed, machines or energy in any of the six, which would stand in the table
Killed outright, the governor's settings are back at the operator's within seconds	L	results/live/V3_ROBUST_KILL.md, docs/ROBUSTNESS_PREREGISTRATION.md	a repetition handed back after 60 s, or not at all
Over two hours the governor neither leaks, slows nor drifts, and hands everything back at the end	L	results/live/V3_ROBUST_LONG.md, docs/ROBUSTNESS_PREREGISTRATION.md	memory growing past a quarter, the decision time growing past a half, fewer than 95% of expected decisions, or a setting not handed back
Governing costs a few thousandths of a core at every size	L	results/live/V3_OWN_COST.md	an own-cost figure growing with the organism
The pooler, the broker and the cache each gain on their untouched workloads	L	V3_PGBENCH.md, V3_KAFKA.md, V3_REDIS.md, each preregistration	a confirmed-worse product row on an untouched workload; the memory row of the cache is such a loss and stands
The one number is +25.9% across the real categories, losses included	L, by rule	results/OMNI_INDEX.md, tools/omni_index.py	a category omitted, a loss not entered, a tuning row counted

Claim	Class	Where the evidence is	What would falsify it
The modelled realms, grids, arms, buildings and swarms gain under their own native controllers	S	results/realms/, V3_PANDAPOWER.md, V3_MUJOCO.md, V3_CITYLEARN.md, V3_SWARM.md	a run of the same simulator at the same version and seed giving other digits
A physical meter shows less energy for the same work	P	no current result ; docs/GPU_PREREGISTRATION.md names the run	the rerun on the current card controller reading no difference or worse
Omni-Compass never changes the engine between a rule and its result	by construction	OMNI_V3.json, tools/omni_version.py --commit <sha> on every table's commits	a table whose runs' commits carry different fingerprints

What a referee asks of each document. Of a preregistration: is it dated before the first counted run, does it name the tuning case and the untouched cases, is every amendment dated and does any amendment made after a result was seen say so? Of a three-run table: are the three run ids separate runs, does the head name the engine at each commit, is every gauge of the preregistration present including the ones that went against the governor, and is the reading one of the three the rule allows? Of the index: is every real category in it, is every tuning row out of it, and does a confirmed loss enter as a loss? Of the raw folder: does the SHA-256 manifest in it match the files, does the governor's audit show every decision with its reading and its write, and does the hand-back record read the operator's value? Every one of those questions has a yes in this program or a disclosure naming the exception; the disclosures are in section 16.7.

How to audit one result from this manual to its raw files. Take any row of Chapter 41. Its source column names a three-run table in results/live/; the table's head names the three GitHub run ids, the commit each ran on and the engine tools/omni_version.py reads at that commit. Each run id names a folder results/live/raw/run-<id>/ holding every repetition's files (the capture, the response times, the governor's audit and log, the per-request records, one SHA-256 manifest per repetition). The table tool named in the table's text (tools/confirm_abc.py or its siblings) rebuilds the table from those folders; tools/omni_index.py rebuilds the index from the tables; tools/dossier.py rebuilds the dossier. The preregistration named in the table's text holds the rules, written before the first counted run, with every amendment dated. Nothing in that chain is by hand.

What one repetition leaves behind. Take the first repetition of the first kill-scenario run, archived at results/live/raw/run-37716845219/paired-1/. It holds two folders, one per arm, and the omni arm's folder holds twenty-nine files, every one of which a reader can open: preflight.txt (the UTC time, the arm, the git commit the arm ran at, the versions of Docker, kind, kubect1, the API server and Python, the metrics-server version with its SHA-256, the worker count and the probe's address); rbac_omni.txt (twenty lines of kubect1 auth can-i, each verb the governor needs reading yes and each it must not have reading no); window_start.txt and window_end.txt (the window's clock); load_schedule.log (every load step and when it was applied); latency.csv (the probe's response times, one row a second); capture.csv (the cluster's state every fifteen seconds: replicas, workers, pods pending, the HPA's target); host_cpu.csv (the runner's own /proc/stat); pod_watch.json and pods_end.json, nodes_end.txt, hpa_end.json (the API server's own records at the end); audit.jsonl (one line per decision: the compass's position, heading and letter, the engine's state, every write with its value sent and its value read back); controller.log and controller2.log (the first governor's log, and the second one's after the kill); audit_kill.jsonl (the watchdog's restore writes, each with its command and its reason, such as "reset: remove record"); watchdog.log (one record: which governor, its process id, whether it was hung or dead, every command run and each exit code); robust.log (the kill at the 360th second with the settings as they stood, "target=50 range=1,10 limits=1850m 3700m workers=5/6", and the hand-back "every setting at the operator's after 8 s, target=50 limits=500m workers=6/6 left=none"); rss.csv (the governor's resident memory

every fifteen seconds); `omni_writes.txt` (how many writes the governor made in the window: 42); `kill_switch.txt` (the reset check: target restored, workers in service, no record left); `omni2.pid` (the second governor's process id) and `kill` (the marker the harness wrote at the kill); `robust.out` (the same robustness log as the job printed it); `end_reads.err`, `pod_watch.err` and `capture.csv.errors` (empty, as they should be); `metrics-server-components.yaml` (the pinned manifest as applied); and `SHA256SUMS.txt`, one line per file. A reader who doubts a number in the kill table can find the second it was taken.

APPENDIX

Appendix G. The Muscles: What Each Is For, and How It Is Wired

© 2026 The Omni-Compass LLC. All rights reserved. **Evaluation and simulation use only.** Any commercial use, commercialization, monetization, production use, redistribution or hosted service of any part of Omni-Compass requires a signed, paid Omni-Compass Enterprise License from The Omni-Compass LLC. Patents, copyrights and trademarks filed in the USA. See LICENSE, NOTICE and DISCLOSURES.md.

A **muscle** is one setting on one machine that already has its own control: a replica target, a node pool's size, a GPU's clock ceiling, a chiller's setpoint, a battery's reserve, a joint's effort. Omni-Compass does not replace that control. It reads the machine's meters, computes one bounded force with the compass law, and moves the setting the machine already accepts, through a plug that reads the setting once before the first write, reads back every write, steps aside if another controller moves it, and puts it back at the end.

This chapter lists all 945 muscles of the catalog (realms/catalog.csv). For each one it says what kind of machine it is, which of the four kinds of knob it is, what Omni-Compass reads and does with it, how such a muscle is reached in a real stack, and which organisms it belongs to. The plants behind the benchmark numbers are models of these machines (evidence class **S**); a muscle in this list is wired on a real system only through the levels and the checks of the manual (chapters 8 and 9). See DISCLOSURES.md.

How to read an entry

The four kinds of knob.

- **capacity**: how much of the machine is in service (replicas, instances, units, speed). Omni-Compass writes it.
- **admission**: what is let in (queues, rate limits, admission control). Omni-Compass reads it and leaves it to its own controller in the current law; it is counted, never written.
- **power**: the power share or limit the machine may draw. Omni-Compass writes it inside its cover.
- **setpoint**: the target the machine's own loop holds (a temperature, a reserve, a process value). Omni-Compass moves it inside its safe band.

The five kinds of plant. Every muscle is modelled on one of five plants, each with its own service reading and lever:

Plant	What it is	What Omni-Compass reads	What Omni-Compass moves
compute_pool	a pool of servers, pods, GPUs, links or disks serving a stream of requests or jobs	how far its queue and its load sit toward the service line (queue or lateness, and load above half)	the capacity it may use: its replica or instance count (through its autoscaler's target or floor), its power share, or one unit given back through the release gate

Plant	What it is	What Omni-Compass reads	What Omni-Compass moves
therma l_zone	a data hall or building zone cooled by chillers, air handlers or fans	its temperature against its limit	its supply setpoint (colder is more cooling), its power share, or one cooling unit given back when the room has margin
energy _stora ge	a battery, UPS or microgrid store with a load and a grid connection	its draw on the grid connection and its reserve	its reserve setpoint (how much charge it holds back); its power limit is left native
motion _axis	a motor-driven axis: a robot joint, a flight-control surface, a reaction wheel, a traction motor	its tracking error and its load	its speed and effort share, or its power share
proces s_loop	a regulated process: a temperature, pressure, flow, level, voltage or composition held at a setpoint	how far the process sits from its band	its setpoint within its safe band, its actuator range, or its power share

The organisms. The 945 muscles build six organisms: each of the four realms (every muscle whose realm list includes it), the four stacked with every duplicate kept (1,716), and the whole tower with every muscle once (945). A muscle of the shared spine sits in all four realms.

Organism	Muscles
Compute / AI / Cloud	430
Physics / Robotics / Autonomous	376
Energy / Facility / Industrial	470
Distribution / Specialized	440
The four stacked, every duplicate kept	1,716
The whole tower, every muscle once	945

How a muscle is reached in a real stack. The class of machine decides the wire:

Class of machine	What it is	The wire in a real stack
batch	batch and queued jobs	the job queue (Kubernetes Jobs, Slurm, Spark)
batch_reactor	a GMP batch reactor, fermenter or food process	the batch control system (ISA-88, OPC UA)
building	a building zone	the building management system
chamber	a controlled chamber (clean room, kiln, reactor)	the PLC or DCS
clinical	clinical systems: imaging archives, records, interface engines, monitoring gateways	the application's own scaling and admission settings (Kubernetes, the vendor console)
commerce	a customer-facing service with bursts	the Kubernetes API
cpu_host	a CPU host whose clock and power can be set	Linux cpufreq and RAPL
data_hall	a data hall cooling plant	the building management system (BACnet, Modbus) or its command line
database	a database cluster	the database operator (replicas, connection pools)
district_heat	a district heating or cooling network	the network control system and substation controllers (Modbus, M-Bus)
elevator_hoist	an elevator or escalator drive and its group control	the lift controller and destination dispatch (modelled only)
ev_traction	an electric vehicle traction motor	the vehicle's motor controller (CAN; modelled only)
fabric	a high-speed fabric (RDMA, DPU, SmartNIC)	the fabric manager's API
facility	a facility battery behind the meter	the energy management system

Class of machine	What it is	The wire in a real stack
feeder_voltage	a distribution feeder's voltage	the distribution management system (IEC 61850, DNP3; modelled only)
flight_axis	a flight-control axis	the flight controller (modelled only)
gpu	GPU serving with a response-time target	nvidia-smi (clock ceiling and power limit) or the serving autoscaler
gpu_batch	GPU training and batch work	nvidia-smi, the job scheduler's pause and resume
hospital	a hospital's critical rooms (operating rooms, isolation, pharmacy, imaging suites)	the hospital's building management system (BACnet)
inverter	a plant of grid-support inverters or wind turbines	the plant controller and inverter settings (SunSpec, IEC 61850, IEEE 1547)
irrigation	pumps, pressure and climate on a farm or in a greenhouse	the pump station controller, the pivot panel or the climate computer (Modbus, the vendor cloud)
marine_propulsion	a ship's propulsion shaft and power plant	the vessel automation and power management system (modelled only)
microgrid	a microgrid battery and loads	the inverter or energy management system (IEEE 2030.5, SunSpec, OpenADR)
mill	a grinding, flotation or materials-handling circuit	the plant DCS or PLC (OPC UA)
network	network routing and switching	the network controller (SDN, routing API)
node	a fleet of machines or VMs that boot in minutes	the node pool's size command (Karpenter, Cluster Autoscaler, MachineSet, cloud ASG)
pipeline	a pipeline segment with its compressors, pumps and valves	pipeline SCADA (DNP3, Modbus, OPC UA)
port	a container terminal's cranes and vehicles taking moves	the terminal operating system and fleet controllers
process	an industrial process loop	the PLC or DCS (OPC UA, EtherNet/IP, PROFINET)
qpu	a quantum or specialised accelerator queue	the accelerator's job queue
rail_traction	a train's traction and auxiliary systems	the train control and management system and ATO (modelled only)
ran	a radio access network cell or site	the RAN controller (O-RAN interfaces)
reaction_wheel	a spacecraft reaction wheel	the attitude controller (modelled only)
robot_fleet	a fleet of robots or vehicles taking tasks	the fleet manager's dispatch API
robot_joint	a robot joint servo	the robot controller (ROS 2, EtherCAT, the drive's fieldbus)
server	an application service on servers or pods	the Kubernetes API (HPA target and floor, pod resize)
storage	block, file or object storage	the storage system's QoS and tiering API
turbine	a generating unit's governor, boiler or excitation loop	the plant DCS and governor (OPC UA, IEC 61850; modelled only)
ups	an uninterruptible power supply	the UPS and PDU management interface (SNMP, Modbus)
water	a water or pumping process	the SCADA system (Modbus, DNP3)
workflow	a workflow or pipeline engine	the workflow engine's concurrency settings

Wires marked *modelled only* exist as plants in the benchmark and are not built for live use.

The shared spine: 257 muscles in all four realms

The machines every realm stands on: servers, machines, GPUs and CPUs, network, storage, observability, security, cooling and electrical distribution. Each spine muscle is part of every realm's organism, counted once in the tower and four times in the stack.

Cloud VM & Capacity (21 muscles)

Each is a pool of servers, pods, GPUs, links or disks serving a stream of requests or jobs; its class is node (a fleet of machines or VMs that boot in minutes), reached through the node pool's size command (Karpenter, Cluster Autoscaler, MachineSet, cloud ASG). Omni-Compass reads how far its queue and its load sit toward the service line (queue or lateness, and load above half). It belongs to every organism.

#	Muscle	Knob	What Omni-Compass does with it	Organisms
97	vm count	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
98	vm size	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
99	vm start stop	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
100	vm migrate	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
101	instance family	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
102	asg min	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
103	asg max	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
104	spot mix	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
105	reservation mix	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
106	zone selection	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
107	region selection	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
108	architecture selection	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
109	accelerator selection	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
110	boot disk class	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
111	placement group	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower

#	Muscle	Knob	What Omni-Compass does with it	Organisms
112	interruption response	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
1038	target tracking policy	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	all four realms; stack; tower
1039	scaling cooldown warmup	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
1114	vm cpu cap by priority	admission	reads it; left to its own controller	all four realms; stack; tower
1120	predictive scaling	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
1121	instance refresh	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower

Container Resources (23 muscles)

Each is a pool of servers, pods, GPUs, links or disks serving a stream of requests or jobs; its class is server (an application service on servers or pods), reached through the Kubernetes API (HPA target and floor, pod resize). Omni-Compass reads how far its queue and its load sit toward the service line (queue or lateness, and load above half). It belongs to every organism.

#	Muscle	Knob	What Omni-Compass does with it	Organisms
33	cpu request	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
34	cpu limit	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
35	memory request	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
36	memory limit	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
37	ephemeral storage	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
38	hugepages	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
39	io weight	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	all four realms; stack; tower
40	pids limit	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
41	memory high	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower

#	Muscle	Knob	What Omni-Compass does with it	Organisms
42	memory reclaim	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
43	swap limit	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
44	cgroup io max	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
45	cpu weight	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	all four realms; stack; tower
46	cpu quota	admission	reads it; left to its own controller	all four realms; stack; tower
47	cpuset	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
48	runtime class	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
1017	in place resize	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
1019	namespace quota	admission	reads it; left to its own controller	all four realms; stack; tower
1020	default limits	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
1022	memory hard limit	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
1023	memory protection	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
1024	io latency target	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	all four realms; stack; tower
1025	freeze	admission	reads it; left to its own controller	all four realms; stack; tower

Host CPU & Memory (23 muscles)

Each is a pool of servers, pods, GPUs, links or disks serving a stream of requests or jobs; its class is `cpu_host` (a CPU host whose clock and power can be set), reached through Linux `cpufreq` and `RAPL`. Omni-Compass reads how far its queue and its load sit toward the service line (queue or lateness, and load above half). It belongs to every organism.

#	Muscle	Knob	What Omni-Compass does with it	Organisms
113	cpufreq min	power	holds the power share inside its cover; full at once past the wall	all four realms; stack; tower
114	cpufreq max	power	holds the power share inside its cover; full at once past the wall	all four realms; stack; tower
115	rapl package power	power	holds the power share inside its cover; full at once past the wall	all four realms; stack; tower

#	Muscle	Knob	What Omni-Compass does with it	Organisms
116	uncore frequency	power	holds the power share inside its cover; full at once past the wall	all four realms; stack; tower
117	energy perf preference	power	holds the power share inside its cover; full at once past the wall	all four realms; stack; tower
118	cpu idle policy	power	holds the power share inside its cover; full at once past the wall	all four realms; stack; tower
119	memory bandwidth	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
120	numa balance	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
121	irq affinity	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
122	llc allocation	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
123	memory pressure gate	admission	reads it; left to its own controller	all four realms; stack; tower
124	page reclaim rate	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
125	transparent hugepages	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
126	core online offline	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
127	thermal throttle policy	power	holds the power share inside its cover; full at once past the wall	all four realms; stack; tower
128	host power profile	power	holds the power share inside its cover; full at once past the wall	all four realms; stack; tower
1040	frequency governor	power	holds the power share inside its cover; full at once past the wall	all four realms; stack; tower
1041	turbo boost	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
1042	dram power limit	power	holds the power share inside its cover; full at once past the wall	all four realms; stack; tower
1043	swappiness reclaim	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
1112	core packing	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
1115	per pod cpu frequency range and c state access	power	holds the power share inside its cover; full at once past the wall	all four realms; stack; tower
1116	uncore frequency per node	power	holds the power share inside its cover; full at once past the wall	all four realms; stack; tower

Kubernetes Placement & Scheduling (20 muscles)

Each is a pool of servers, pods, GPUs, links or disks serving a stream of requests or jobs; its class is server (an application service on servers or pods), reached through the Kubernetes API (HPA target and floor, pod resize). Omni-Compass reads how far its queue and its load sit toward the service line (queue or lateness, and load above half). It belongs to every organism.

#	Muscle	Knob	What Omni-Compass does with it	Organisms
49	node selector	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
50	node affinity	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
51	pod anti affinity	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
52	topology spread	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
53	numa placement	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
54	gpu topology	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
55	storage locality	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
56	network locality	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
57	taint toleration	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	all four realms; stack; tower
58	priority class	admission	reads it; left to its own controller	all four realms; stack; tower
59	preemption policy	admission	reads it; left to its own controller	all four realms; stack; tower
60	device claim	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
61	failure domain spread	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
62	scheduler backoff	admission	reads it; left to its own controller	all four realms; stack; tower
63	gang admission	admission	reads it; left to its own controller	all four realms; stack; tower
64	deschedule	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower

#	Muscle	Knob	What Omni-Compass does with it	Organisms
1026	scheduler scoring strategy	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
1027	scheduler plugin weights	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	all four realms; stack; tower
1028	scheduling gates	admission	reads it; left to its own controller	all four realms; stack; tower
1029	cpu manager	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower

Kubernetes Workload Scaling (31 muscles)

Each is a pool of servers, pods, GPUs, links or disks serving a stream of requests or jobs; its class is server (an application service on servers or pods), reached through the Kubernetes API (HPA target and floor, pod resize). Omni-Compass reads how far its queue and its load sit toward the service line (queue or lateness, and load above half). It belongs to every organism.

#	Muscle	Knob	What Omni-Compass does with it	Organisms
17	replicas	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
18	hpa cpu target	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	all four realms; stack; tower
19	hpa memory target	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	all four realms; stack; tower
20	hpa custom target	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	all four realms; stack; tower
21	vpa apply	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
22	scale to zero	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
23	keda threshold	admission	reads it; left to its own controller	all four realms; stack; tower
24	rollout rate	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
25	max surge	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
26	max unavailable	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
27	deployment pause resume	admission	reads it; left to its own controller	all four realms; stack; tower
28	rollout abort	admission	reads it; left to its own controller	all four realms; stack; tower

#	Muscle	Knob	What Omni-Compass does with it	Organisms
29	pod eviction	admission	reads it; left to its own controller	all four realms; stack; tower
30	pdb policy	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
31	scheduler queue priority	admission	reads it; left to its own controller	all four realms; stack; tower
32	api priority fairness	admission	reads it; left to its own controller	all four realms; stack; tower
1001	hpa min replicas	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
1002	hpa max replicas	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
1003	hpa scale down stabilisation window	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	all four realms; stack; tower
1004	hpa scale up stabilisation window	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	all four realms; stack; tower
1005	hpa scaling rate policies	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
1006	hpa policy selection	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
1007	hpa tolerance	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
1008	hpa sync period	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
1009	vpa minmax allowed resources	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
1010	keda polling interval	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	all four realms; stack; tower
1011	rollout progress deadline	admission	reads it; left to its own controller	all four realms; stack; tower
1012	rollout readiness delay	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
1014	unhealthy pod eviction policy	admission	reads it; left to its own controller	all four realms; stack; tower
1015	graceful termination	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
1016	probes	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower

NVIDIA GPU Hardware (19 muscles)

Each is a pool of servers, pods, GPUs, links or disks serving a stream of requests or jobs; its class is `gpu` (GPU serving with a response-time target), reached through `nvidia-smi` (clock ceiling and power limit) or the serving autoscaler. Omni-Compass reads how far its queue and its load sit toward the service line (queue or lateness, and load above half). It belongs to every organism.

#	Muscle	Knob	What Omni-Compass does with it	Organisms
129	gpu allocate	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
130	gpu power limit	power	holds the power share inside its cover; full at once past the wall	all four realms; stack; tower
131	gpu sm clock	power	holds the power share inside its cover; full at once past the wall	all four realms; stack; tower
132	gpu memory clock	power	holds the power share inside its cover; full at once past the wall	all four realms; stack; tower
133	gpu persistence mode	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
134	gpu compute mode	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
135	mig mode	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
136	mig geometry	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
137	gpu timeslice	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
138	gpu mps	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
139	gpu quarantine	admission	reads it; left to its own controller	all four realms; stack; tower
140	gpu reset	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
141	gpu thermal limit	power	holds the power share inside its cover; full at once past the wall	all four realms; stack; tower
142	gpu ecc response	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
143	gpu job power budget	admission	reads it; left to its own controller	all four realms; stack; tower
1044	application clocks	power	holds the power share inside its cover; full at once past the wall	all four realms; stack; tower
1045	gpu temperature target	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	all four realms; stack; tower
1117	sync boost	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower

#	Muscle	Knob	What Omni-Compass does with it	Organisms
1118	target clocks	power	holds the power share inside its cover; full at once past the wall	all four realms; stack; tower

Node Fleet & Carpenter-Class Control (24 muscles)

Each is a pool of servers, pods, GPUs, links or disks serving a stream of requests or jobs; its class is node (a fleet of machines or VMs that boot in minutes), reached through the node pool's size command (Carpenter, Cluster Autoscaler, MachineSet, cloud ASG). Omni-Compass reads how far its queue and its load sit toward the service line (queue or lateness, and load above half). It belongs to every organism.

#	Muscle	Knob	What Omni-Compass does with it	Organisms
65	node desired	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
66	node pool min	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
67	node pool max	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
68	node provision	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
69	node cordon	admission	reads it; left to its own controller	all four realms; stack; tower
70	node drain	admission	reads it; left to its own controller	all four realms; stack; tower
71	node consolidate	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
72	node replace	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
73	node shutdown	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
74	node power on	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
75	nodepool weight	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	all four realms; stack; tower
76	disruption budget	admission	reads it; left to its own controller	all four realms; stack; tower
77	consolidation policy	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
78	consolidate after	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
79	expire after	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower

#	Muscle	Knob	What Omni-Compass does with it	Organisms
80	capacity class	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
1030	scale down utilisation threshold	admission	reads it; left to its own controller	all four realms; stack; tower
1031	scale down unneeded time	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
1032	scale down delay after add	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
1033	expander	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
1034	node pool limits	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
1035	kubelet eviction thresholds	admission	reads it; left to its own controller	all four realms; stack; tower
1036	kubelet reserved resources	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
1037	max pods per node	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower

Network Routing & Switching (19 muscles)

Each is a pool of servers, pods, GPUs, links or disks serving a stream of requests or jobs; its class is network (network routing and switching), reached through the network controller (SDN, routing API). Omni-Compass reads how far its queue and its load sit toward the service line (queue or lateness, and load above half). It belongs to every organism.

#	Muscle	Knob	What Omni-Compass does with it	Organisms
225	lb weight	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	all four realms; stack; tower
226	route weight	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	all four realms; stack; tower
227	rate limit	admission	reads it; left to its own controller	all four realms; stack; tower
228	bandwidth limit	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
229	qos class	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
230	connection limit	admission	reads it; left to its own controller	all four realms; stack; tower
231	failover route	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower

#	Muscle	Knob	What Omni-Compass does with it	Organisms
232	nic rate limit	admission	reads it; left to its own controller	all four realms; stack; tower
233	nic queue count	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
234	queue discipline	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
235	congestion control	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
236	egress budget	admission	reads it; left to its own controller	all four realms; stack; tower
237	ingress budget	admission	reads it; left to its own controller	all four realms; stack; tower
238	ecmp weight	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	all four realms; stack; tower
239	path selection	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
240	network isolation	admission	reads it; left to its own controller	all four realms; stack; tower
1046	nic interrupt coalescing	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
1047	nic ring size	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
1048	port link power	power	holds the power share inside its cover; full at once past the wall	all four realms; stack; tower

Observability & Telemetry (9 muscles)

Each is a pool of servers, pods, GPUs, links or disks serving a stream of requests or jobs; its class is server (an application service on servers or pods), reached through the Kubernetes API (HPA target and floor, pod resize). Omni-Compass reads how far its queue and its load sit toward the service line (queue or lateness, and load above half). It belongs to every organism.

#	Muscle	Knob	What Omni-Compass does with it	Organisms
373	collector memory limit	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
374	export concurrency	admission	reads it; left to its own controller	all four realms; stack; tower
376	cardinality budget	admission	reads it; left to its own controller	all four realms; stack; tower
377	retention window	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	all four realms; stack; tower

#	Muscle	Knob	What Omni-Compass does with it	Organisms
378	remote write queue	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
379	telemetry shed	admission	reads it; left to its own controller	all four realms; stack; tower
1053	metric scrape interval	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	all four realms; stack; tower
1054	batching	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
1119	per pod energy attribution	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower

Reliability, Security & Recovery (12 muscles)

Each is a pool of servers, pods, GPUs, links or disks serving a stream of requests or jobs; its class is `server` (an application service on servers or pods), reached through the Kubernetes API (HPA target and floor, pod resize). Omni-Compass reads how far its queue and its load sit toward the service line (queue or lateness, and load above half). It belongs to every organism.

#	Muscle	Knob	What Omni-Compass does with it	Organisms
385	restart	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
386	rollback	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
388	traffic divert recovery	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
389	degraded mode	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
390	actuator freeze	admission	reads it; left to its own controller	all four realms; stack; tower
391	workload isolate	admission	reads it; left to its own controller	all four realms; stack; tower
393	credential rotation gate	admission	reads it; left to its own controller	all four realms; stack; tower
394	policy enforcement	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
395	rate abuse gate	admission	reads it; left to its own controller	all four realms; stack; tower
396	fault domain isolate	admission	reads it; left to its own controller	all four realms; stack; tower
397	backup trigger	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower

#	Muscle	Knob	What Omni-Compass does with it	Organisms
400	kill switch	admission	reads it; left to its own controller	all four realms; stack; tower

Storage Block/File/Object (20 muscles)

Each is a pool of servers, pods, GPUs, links or disks serving a stream of requests or jobs; its class is storage (block, file or object storage), reached through the storage system's QoS and tiering API. Omni-Compass reads how far its queue and its load sit toward the service line (queue or lateness, and load above half). It belongs to every organism.

#	Muscle	Knob	What Omni-Compass does with it	Organisms
257	volume size	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
258	iops limit	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
259	throughput limit	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
260	replica count	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
261	storage tier	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
262	volume placement	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
263	snapshot trigger	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
264	rebalance	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
265	recovery rate	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
266	backfill rate	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
267	compaction pressure	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
268	cache allocation	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
269	object replication	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
270	erasure code profile	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower

#	Muscle	Knob	What Omni-Compass does with it	Organisms
271	storage admission	admission	reads it; left to its own controller	all four realms; stack; tower
272	degraded storage gate	admission	reads it; left to its own controller	all four realms; stack; tower
1049	scrub schedule	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
1050	io scheduler	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
1051	disk power spin down	power	holds the power share inside its cover; full at once past the wall	all four realms; stack; tower
1052	lifecycle tiering	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower

Cooling, Chillers & Thermodynamics (18 muscles)

Each is a data hall or building zone cooled by chillers, air handlers or fans; its class is `data_hall` (a data hall cooling plant), reached through the building management system (BACnet, Modbus) or its command line. Omni-Compass reads its temperature against its limit. It belongs to every organism.

#	Muscle	Knob	What Omni-Compass does with it	Organisms
545	supply air temperature	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	all four realms; stack; tower
546	return air target	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	all four realms; stack; tower
547	coolant supply temperature	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	all four realms; stack; tower
548	coolant flow	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
549	pump speed	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
550	fan speed	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
551	chiller setpoint	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	all four realms; stack; tower
552	compressor authority	power	holds the power share inside its cover; full at once past the wall	all four realms; stack; tower
553	cooling tower fan	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
554	cooling capacity	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
555	rack thermal budget	power	holds the power share inside its cover; full at once past the wall	all four realms; stack; tower

#	Muscle	Knob	What Omni-Compass does with it	Organisms
556	gpu thermal envelope	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	all four realms; stack; tower
557	cpu thermal envelope	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	all four realms; stack; tower
558	thermal workload migrate	admission	reads it; left to its own controller	all four realms; stack; tower
559	thermal load shed	admission	reads it; left to its own controller	all four realms; stack; tower
1055	server fan speed	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
1056	economiser	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
1122	chiller plant supervisory control rl with safety layer	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower

PDU, UPS & Electrical Distribution (18 muscles)

Each is a battery, UPS or microgrid store with a load and a grid connection; its class is ups (an uninterruptible power supply), reached through the UPS and PDU management interface (SNMP, Modbus). Omni-Compass reads its draw on the grid connection and its reserve. It belongs to every organism.

#	Muscle	Knob	What Omni-Compass does with it	Organisms
529	server power cap	power	holds the power share inside its cover; full at once past the wall	all four realms; stack; tower
530	rack power cap	power	holds the power share inside its cover; full at once past the wall	all four realms; stack; tower
531	pdu branch power limit	power	holds the power share inside its cover; full at once past the wall	all four realms; stack; tower
532	pdu outlet control	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
534	ups operating mode	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
535	ups charge rate	power	holds the power share inside its cover; full at once past the wall	all four realms; stack; tower
536	ups discharge rate	power	holds the power share inside its cover; full at once past the wall	all four realms; stack; tower
537	phase balance	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
538	load transfer	admission	reads it; left to its own controller	all four realms; stack; tower
539	reactive power target	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	all four realms; stack; tower

#	Muscle	Knob	What Omni-Compass does with it	Organisms
540	voltage target	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	all four realms; stack; tower
541	generator dispatch	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
542	electrical isolation	admission	reads it; left to its own controller	all four realms; stack; tower
543	breaker trip gate	admission	reads it; left to its own controller	all four realms; stack; tower
1057	ups shutdown battery test	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	all four realms; stack; tower
1110	hierarchical power budget row pdu switchboard controllers	power	holds the power share inside its cover; full at once past the wall	all four realms; stack; tower
1111	priority aware capping throttle low priority first	power	holds the power share inside its cover; full at once past the wall	all four realms; stack; tower
1113	power oversubscription with prediction	power	holds the power share inside its cover; full at once past the wall	all four realms; stack; tower

Realm: Compute / AI / Cloud (173 muscles of its own, 430 in its organism with the spine)

AI Inference Serving (24 muscles)

Each is a pool of servers, pods, GPUs, links or disks serving a stream of requests or jobs; its class is gpu (GPU serving with a response-time target), reached through nvidia-smi (clock ceiling and power limit) or the serving autoscaler. Omni-Compass reads how far its queue and its load sit toward the service line (queue or lateness, and load above half). Its home realm is Compute / AI / Cloud.

#	Muscle	Knob	What Omni-Compass does with it	Organisms
161	model replicas	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
162	model route weight	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Compute; stack; tower
163	model load	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
164	model unload	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
165	model instance count	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
166	continuous batching	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower

#	Muscle	Knob	What Omni-Compass does with it	Organisms
167	max batch size	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
168	batch queue delay	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
169	inference concurrency	admission	reads it; left to its own controller	Compute; stack; tower
170	inference max tokens	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
171	kv cache budget	admission	reads it; left to its own controller	Compute; stack; tower
172	prefix cache budget	admission	reads it; left to its own controller	Compute; stack; tower
173	speculative decode budget	admission	reads it; left to its own controller	Compute; stack; tower
174	model precision	power	holds the power share inside its cover; full at once past the wall	Compute; stack; tower
175	inference priority	admission	reads it; left to its own controller	Compute; stack; tower
176	inference slo gate	admission	reads it; left to its own controller	Compute; stack; tower
1058	autoscaling target concurrency qps	admission	reads it; left to its own controller	Compute; stack; tower
1059	scale to zero panic window	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Compute; stack; tower
1060	request priority levels	admission	reads it; left to its own controller	Compute; stack; tower
1061	kv cache swap space	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
1062	chunked prefill	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
1063	tensor parallel degree	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
1064	disaggregated prefill decode	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
1065	lora adapter load	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower

AI Training (18 muscles)

Each is a pool of servers, pods, GPUs, links or disks serving a stream of requests or jobs; its class is `gpu_batch` (GPU training and batch work), reached through `nvidia-smi`, the job scheduler's pause and resume. Omni-Compass reads how far its queue and its load sit toward the service line (queue or lateness, and load above half). Its home realm is Compute / AI / Cloud.

#	Muscle	Knob	What Omni-Compass does with it	Organisms
177	training workers	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
178	global batch size	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
179	microbatch size	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
180	gradient accumulation	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
181	data parallelism	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
182	tensor parallelism	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
183	pipeline parallelism	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
184	expert parallelism	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
185	checkpoint interval	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Compute; stack; tower
186	checkpoint trigger	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
187	training preempt	admission	reads it; left to its own controller	Compute; stack; tower
188	training gang size	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
189	elastic worker count	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
190	straggler mitigation	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
191	training precision	power	holds the power share inside its cover; full at once past the wall	Compute; stack; tower
192	compute comm overlap	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
1066	activation checkpointing	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
1067	dataloader workers	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower

Cross-Cluster, Multi-Region & Edge (15 muscles)

Each is a pool of servers, pods, GPUs, links or disks serving a stream of requests or jobs; its class is node (a fleet of machines or VMs that boot in minutes), reached through the node pool's size command (Karpenter, Cluster Autoscaler, MachineSet, cloud ASG). Omni-Compass reads how far its queue and its load sit toward the service line (queue or lateness, and load above half). Its home realm is Compute / AI / Cloud.

#	Muscle	Knob	What Omni-Compass does with it	Organisms
609	multi cluster dispatch	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
610	region dispatch	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
611	zone dispatch	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
612	edge dispatch	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
613	cloud capacity class	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
614	workload migrate region	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
615	data residency gate	admission	reads it; left to its own controller	Compute; stack; tower
616	latency region gate	admission	reads it; left to its own controller	Compute; stack; tower
617	cost region gate	admission	reads it; left to its own controller	Compute; stack; tower
618	carbon region gate	admission	reads it; left to its own controller	Compute; stack; tower
619	global failover	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
620	federation quota	admission	reads it; left to its own controller	Compute; stack; tower
621	cross cluster replication	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
622	edge offload	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
623	global admission	admission	reads it; left to its own controller	Compute; stack; tower

DPU SmartNIC & Programmable IO (12 muscles)

Each is a pool of servers, pods, GPUs, links or disks serving a stream of requests or jobs; its class is fabric (a high-speed fabric (RDMA, DPU, SmartNIC)), reached through the fabric manager's API. Omni-Compass reads how far its queue and its load sit toward the service line (queue or lateness, and load above half). Its home realm is Compute / AI / Cloud.

#	Muscle	Knob	What Omni-Compass does with it	Organisms
AUDIT-0021	dpu pf tx rate	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower

#	Muscle	Knob	What Omni-Compass does with it	Organisms
AUDIT-0022	dpu vf tx rate	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
AUDIT-0023	dpu sf tx rate	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
AUDIT-0024	dpu bandwidth share	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
AUDIT-0025	dpu qos group	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
AUDIT-0026	sriov vf count	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
AUDIT-0027	smartnic flow steering	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
AUDIT-0028	smartnic offload enable	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
AUDIT-0029	dpu cpu budget	admission	reads it; left to its own controller	Compute; stack; tower
AUDIT-0030	dpu memory budget	admission	reads it; left to its own controller	Compute; stack; tower
AUDIT-0031	dpu service placement	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
AUDIT-0032	dpu failover	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower

Distributed Cluster Managers (15 muscles)

Each is a pool of servers, pods, GPUs, links or disks serving a stream of requests or jobs; its class is batch (batch and queued jobs), reached through the job queue (Kubernetes Jobs, Slurm, Spark). Omni-Compass reads how far its queue and its load sit toward the service line (queue or lateness, and load above half). Its home realm is Compute / AI / Cloud.

#	Muscle	Knob	What Omni-Compass does with it	Organisms
209	task admission	admission	reads it; left to its own controller	Compute; stack; tower
210	task priority	admission	reads it; left to its own controller	Compute; stack; tower
211	resource reservation	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
212	cluster quota	admission	reads it; left to its own controller	Compute; stack; tower
213	task preemption	admission	reads it; left to its own controller	Compute; stack; tower
214	binpack pressure	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower

#	Muscle	Knob	What Omni-Compass does with it	Organisms
215	spread pressure	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
216	gang schedule	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
217	worker allocation	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
218	maintenance evacuation	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
219	oversubscription	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
220	resource reclaim	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
221	queue fairness	admission	reads it; left to its own controller	Compute; stack; tower
222	deadline pressure	admission	reads it; left to its own controller	Compute; stack; tower
223	scheduler retry	admission	reads it; left to its own controller	Compute; stack; tower

GPU Fabric & RDMA (18 muscles)

Each is a pool of servers, pods, GPUs, links or disks serving a stream of requests or jobs; its class is fabric (a high-speed fabric (RDMA, DPU, SmartNIC)), reached through the fabric manager's API. Omni-Compass reads how far its queue and its load sit toward the service line (queue or lateness, and load above half). Its home realm is Compute / AI / Cloud.

#	Muscle	Knob	What Omni-Compass does with it	Organisms
145	nvlink placement	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
146	nvswitch route	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
147	gpu fabric quarantine	admission	reads it; left to its own controller	Compute; stack; tower
148	rdma bandwidth	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
149	rdma route	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
150	nic affinity	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
151	collective concurrency	admission	reads it; left to its own controller	Compute; stack; tower
152	collective algorithm	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower

#	Muscle	Knob	What Omni-Compass does with it	Organisms
153	collective chunk size	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
154	rank placement	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
155	gpudirect policy	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
156	congestion response	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
157	rail selection	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
158	fabric failover	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
159	communication priority	admission	reads it; left to its own controller	Compute; stack; tower
160	fabric isolation	admission	reads it; left to its own controller	Compute; stack; tower
1068	collective protocol	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
1069	nic hca selection	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower

HPC & Distributed Compute (19 muscles)

Each is a pool of servers, pods, GPUs, links or disks serving a stream of requests or jobs; its class is batch (batch and queued jobs), reached through the job queue (Kubernetes Jobs, Slurm, Spark). Omni-Compass reads how far its queue and its load sit toward the service line (queue or lateness, and load above half). Its home realm is Compute / AI / Cloud.

#	Muscle	Knob	What Omni-Compass does with it	Organisms
193	job slots	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
194	mpi ranks	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
195	rank mapping	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
196	node allocation	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
197	job walltime	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
198	job priority	admission	reads it; left to its own controller	Compute; stack; tower
199	job preemption	admission	reads it; left to its own controller	Compute; stack; tower

#	Muscle	Knob	What Omni-Compass does with it	Organisms
200	checkpoint restart	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
201	parallel io budget	admission	reads it; left to its own controller	Compute; stack; tower
202	collective budget	admission	reads it; left to its own controller	Compute; stack; tower
203	accelerator share	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
204	cpu gpu ratio	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Compute; stack; tower
205	memory per rank	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
206	scratch budget	admission	reads it; left to its own controller	Compute; stack; tower
207	scheduler fair share	admission	reads it; left to its own controller	Compute; stack; tower
208	backfill policy	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
1070	node power saving	power	holds the power share inside its cover; full at once past the wall	Compute; stack; tower
1071	cpu frequency per job	power	holds the power share inside its cover; full at once past the wall	Compute; stack; tower
1072	partition limits	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower

Kubernetes Dynamic Device Allocation (8 muscles)

Each is a pool of servers, pods, GPUs, links or disks serving a stream of requests or jobs; its class is gpu (GPU serving with a response-time target), reached through nvidia-smi (clock ceiling and power limit) or the serving autoscaler. Omni-Compass reads how far its queue and its load sit toward the service line (queue or lateness, and load above half). Its home realm is Compute / AI / Cloud.

#	Muscle	Knob	What Omni-Compass does with it	Organisms
AUDIT-0001	dra device class selection	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
AUDIT-0002	dra claim capacity	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
AUDIT-0003	dra claim sharing	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
AUDIT-0004	dra device taint	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
AUDIT-0005	dra device eviction	admission	reads it; left to its own controller	Compute; stack; tower
AUDIT-0006	dra binding readiness	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower

#	Muscle	Knob	What Omni-Compass does with it	Organisms
AUDIT-0007	dra binding failure response	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
AUDIT-0008	dra device configuration	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Compute; stack; tower

OpenShift & Machine API (9 muscles)

Each is a pool of servers, pods, GPUs, links or disks serving a stream of requests or jobs; its class is node (a fleet of machines or VMs that boot in minutes), reached through the node pool's size command (Karpenter, Cluster Autoscaler, MachineSet, cloud ASG). Omni-Compass reads how far its queue and its load sit toward the service line (queue or lateness, and load above half). Its home realm is Compute / AI / Cloud.

#	Muscle	Knob	What Omni-Compass does with it	Organisms
85	machine remediation	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
88	machine health gate	admission	reads it; left to its own controller	Compute; stack; tower
89	mcp pause	admission	reads it; left to its own controller	Compute; stack; tower
90	mcp max unavailable	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
91	node config rollout	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
92	operator reconcile budget	admission	reads it; left to its own controller	Compute; stack; tower
93	cluster version pacing	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
94	infra machine admission	admission	reads it; left to its own controller	Compute; stack; tower
95	machine failure domain	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower

Quantum Computing Control Simulation (16 muscles)

Each is a pool of servers, pods, GPUs, links or disks serving a stream of requests or jobs; its class is qpu (a quantum or specialised accelerator queue), reached through the accelerator's job queue. Omni-Compass reads how far its queue and its load sit toward the service line (queue or lateness, and load above half). Its home realm is Compute / AI / Cloud.

#	Muscle	Knob	What Omni-Compass does with it	Organisms
593	qubit mapping	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
594	circuit admission	admission	reads it; left to its own controller	Compute; stack; tower
595	shot allocation	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
596	circuit scheduling	admission	reads it; left to its own controller	Compute; stack; tower

#	Muscle	Knob	What Omni-Compass does with it	Organisms
597	gate scheduling	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
598	pulse amplitude	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
599	pulse duration	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Compute; stack; tower
600	pulse phase	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
601	pulse frequency	power	holds the power share inside its cover; full at once past the wall	Compute; stack; tower
602	coupling control	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
603	reset scheduling	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
604	measurement scheduling	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
605	dynamical decoupling	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
606	noise aware routing	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Compute; stack; tower
607	error mitigation budget	admission	reads it; left to its own controller	Compute; stack; tower
608	quantum queue priority	admission	reads it; left to its own controller	Compute; stack; tower

Work Admission & Demand Shaping (19 muscles)

Each is a pool of servers, pods, GPUs, links or disks serving a stream of requests or jobs; its class is server (an application service on servers or pods), reached through the Kubernetes API (HPA target and floor, pod resize). Omni-Compass reads how far its queue and its load sit toward the service line (queue or lateness, and load above half). Its home realm is Compute / AI / Cloud.

#	Muscle	Knob	What Omni-Compass does with it	Organisms
1	api concurrency	admission	reads it; left to its own controller	Compute; stack; tower
2	queue concurrency	admission	reads it; left to its own controller	Compute; stack; tower
3	queue backpressure	admission	reads it; left to its own controller	Compute; stack; tower
4	job admission	admission	reads it; left to its own controller	Compute; stack; tower
5	batch admission	admission	reads it; left to its own controller	Compute; stack; tower
6	inference admission	admission	reads it; left to its own controller	Compute; stack; tower
7	load shed	admission	reads it; left to its own controller	Compute; stack; tower
8	priority gate	admission	reads it; left to its own controller	Compute; stack; tower
9	tenant admission	admission	reads it; left to its own controller	Compute; stack; tower
10	burst limit	admission	reads it; left to its own controller	Compute; stack; tower

#	Muscle	Knob	What Omni-Compass does with it	Organisms
11	deadline admission	admission	reads it; left to its own controller	Compute; stack; tower
12	work budget	admission	reads it; left to its own controller	Compute; stack; tower
13	request queue limit	admission	reads it; left to its own controller	Compute; stack; tower
14	retry admission	admission	reads it; left to its own controller	Compute; stack; tower
15	background work gate	admission	reads it; left to its own controller	Compute; stack; tower
16	maintenance work gate	admission	reads it; left to its own controller	Compute; stack; tower
1074	request concurrency limit	admission	reads it; left to its own controller	Compute; stack; tower
1123	adaptive concurrency limit	admission	reads it; left to its own controller	Compute; stack; tower
1124	priority request queue under limit	admission	reads it; left to its own controller	Compute; stack; tower

Realm: Physics / Robotics / Autonomous (119 muscles of its own, 376 in its organism with the spine)

Automotive EV & Mobile Powertrain (14 muscles)

Each is a motor-driven axis: a robot joint, a flight-control surface, a reaction wheel, a traction motor; its class is `ev_traction` (an electric vehicle traction motor), reached through the vehicle's motor controller (CAN; modelled only). Omni-Compass reads its tracking error and its load. Its home realm is Physics / Robotics / Autonomous.

#	Muscle	Knob	What Omni-Compass does with it	Organisms
1083	state of charge limits	power	holds the power share inside its cover; full at once past the wall	Physics; stack; tower
1084	smart charging schedule	admission	reads it; left to its own controller	Physics; stack; tower
AUDIT-00 45	traction torque limit	power	holds the power share inside its cover; full at once past the wall	Physics; stack; tower
AUDIT-00 46	regen braking level	power	holds the power share inside its cover; full at once past the wall	Physics; stack; tower
AUDIT-00 47	battery charge limit	power	holds the power share inside its cover; full at once past the wall	Physics; stack; tower
AUDIT-00 48	battery discharge limit	power	holds the power share inside its cover; full at once past the wall	Physics; stack; tower
AUDIT-00 49	battery thermal target	power	holds the power share inside its cover; full at once past the wall	Physics; stack; tower
AUDIT-00 50	motor thermal limit	power	holds the power share inside its cover; full at once past the wall	Physics; stack; tower
AUDIT-00 51	vehicle speed envelope	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Physics; stack; tower
AUDIT-00 52	energy recovery target	power	holds the power share inside its cover; full at once past the wall	Physics; stack; tower
AUDIT-00 53	auxiliary power budget	power	holds the power share inside its cover; full at once past the wall	Physics; stack; tower
AUDIT-00 54	fast charge current	power	holds the power share inside its cover; full at once past the wall	Physics; stack; tower
AUDIT-00 55	fast charge voltage	power	holds the power share inside its cover; full at once past the wall	Physics; stack; tower
AUDIT-00 56	vehicle safe state	admission	reads it; left to its own controller	Physics; stack; tower

Aviation & Autonomous Flight (20 muscles)

Each is a motor-driven axis: a robot joint, a flight-control surface, a reaction wheel, a traction motor; its class is `flight_axis` (a flight-control axis), reached through the flight controller (modelled only). Omni-Compass reads its tracking error and its load. Its home realm is Physics / Robotics / Autonomous.

#	Muscle	Knob	What Omni-Compass does with it	Organisms
481	throttle envelope	power	holds the power share inside its cover; full at once past the wall	Physics; stack; tower
482	attitude target	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Physics; stack; tower
483	attitude rate target	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Physics; stack; tower
484	velocity target	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Physics; stack; tower
485	altitude target	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Physics; stack; tower
486	waypoint authority	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Physics; stack; tower
487	flight hold	admission	reads it; left to its own controller	Physics; stack; tower
488	return to home	admission	reads it; left to its own controller	Physics; stack; tower
489	land action	admission	reads it; left to its own controller	Physics; stack; tower
490	mission admission	admission	reads it; left to its own controller	Physics; stack; tower
491	geofence response	admission	reads it; left to its own controller	Physics; stack; tower
492	failsafe selection	admission	reads it; left to its own controller	Physics; stack; tower
493	battery reserve threshold	admission	reads it; left to its own controller	Physics; stack; tower
494	actuator saturation envelope	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Physics; stack; tower
495	flight mode transition	admission	reads it; left to its own controller	Physics; stack; tower
496	flight termination safe state	admission	reads it; left to its own controller	Physics; stack; tower
1079	attitude rate gains	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Physics; stack; tower
1080	horizontal speed limit	power	holds the power share inside its cover; full at once past the wall	Physics; stack; tower
1081	vertical speed limits	power	holds the power share inside its cover; full at once past the wall	Physics; stack; tower
1082	acceleration jerk limits	power	holds the power share inside its cover; full at once past the wall	Physics; stack; tower

Robotics Fleet & Warehouse Automation (15 muscles)

Each is a pool of servers, pods, GPUs, links or disks serving a stream of requests or jobs; its class is `robot_fleet` (a fleet of robots or vehicles taking tasks), reached through the fleet manager's dispatch API. Omni-Compass reads how far its queue and its load sit toward the service line (queue or lateness, and load above half). Its home realm is Physics / Robotics / Autonomous.

#	Muscle	Knob	What Omni-Compass does with it	Organisms
449	robot dispatch	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Physics; stack; tower
450	task assignment	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Physics; stack; tower
451	traffic reservation	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Physics; stack; tower
452	robot route	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Physics; stack; tower
453	charging dispatch	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Physics; stack; tower
454	battery reserve	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Physics; stack; tower
455	elevator request	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Physics; stack; tower
456	door request	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Physics; stack; tower
457	conveyor speed	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Physics; stack; tower
458	agv speed	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Physics; stack; tower
459	warehouse zone admission	admission	reads it; left to its own controller	Physics; stack; tower
460	robot quarantine	admission	reads it; left to its own controller	Physics; stack; tower
461	fleet failover	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Physics; stack; tower
462	human safe stop	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Physics; stack; tower
463	fleet concurrency	admission	reads it; left to its own controller	Physics; stack; tower

Robotics Motion Control (18 muscles)

Each is a motor-driven axis: a robot joint, a flight-control surface, a reaction wheel, a traction motor; its class is `robot_joint` (a robot joint servo), reached through the robot controller (ROS 2, EtherCAT, the drive's fieldbus). Omni-Compass reads its tracking error and its load. Its home realm is Physics / Robotics / Autonomous.

#	Muscle	Knob	What Omni-Compass does with it	Organisms
433	joint position	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Physics; stack; tower
434	joint velocity	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Physics; stack; tower
435	joint acceleration	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Physics; stack; tower
436	joint effort	power	holds the power share inside its cover; full at once past the wall	Physics; stack; tower
437	cartesian velocity	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Physics; stack; tower
438	trajectory speed	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Physics; stack; tower
439	trajectory acceleration	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Physics; stack; tower
440	jerk limit	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Physics; stack; tower
441	collision margin	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Physics; stack; tower
442	force limit	power	holds the power share inside its cover; full at once past the wall	Physics; stack; tower
443	gripper force	power	holds the power share inside its cover; full at once past the wall	Physics; stack; tower
444	locomotion speed	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Physics; stack; tower
445	steering angle	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Physics; stack; tower
446	braking force	power	holds the power share inside its cover; full at once past the wall	Physics; stack; tower
447	balance correction	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Physics; stack; tower
1076	trajectory tolerances	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Physics; stack; tower
1077	velocity acceleration scaling	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Physics; stack; tower
1078	controller update rate	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Physics; stack; tower

Spacecraft & Flight Software (14 muscles)

Each is a motor-driven axis: a robot joint, a flight-control surface, a reaction wheel, a traction motor; its class is `reaction_wheel` (a spacecraft reaction wheel), reached through the attitude controller (modelled only). Omni-Compass reads its tracking error and its load. Its home realm is Physics / Robotics / Autonomous.

#	Muscle	Knob	What Omni-Compass does with it	Organisms
497	space command admission	admission	reads it; left to its own controller	Physics; stack; tower
499	flight task schedule	admission	reads it; left to its own controller	Physics; stack; tower
500	space mode transition	admission	reads it; left to its own controller	Physics; stack; tower
501	payload duty cycle	power	holds the power share inside its cover; full at once past the wall	Physics; stack; tower
502	communication allocation	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Physics; stack; tower
503	space power budget	power	holds the power share inside its cover; full at once past the wall	Physics; stack; tower
504	space thermal command	power	holds the power share inside its cover; full at once past the wall	Physics; stack; tower
505	attitude command envelope	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Physics; stack; tower
506	reaction wheel allocation	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Physics; stack; tower
507	rcs authority	power	holds the power share inside its cover; full at once past the wall	Physics; stack; tower
508	safe mode transition	admission	reads it; left to its own controller	Physics; stack; tower
509	watchdog recovery	admission	reads it; left to its own controller	Physics; stack; tower
510	instrument activation	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Physics; stack; tower
511	fault isolation	admission	reads it; left to its own controller	Physics; stack; tower

Rail Traction & Train Control (14 muscles)

Each is a motor-driven axis: a robot joint, a flight-control surface, a reaction wheel, a traction motor; its class is `rail_traction` (a train's traction and auxiliary systems), reached through the train control and management system and ATO (modelled only). Omni-Compass reads its tracking error and its load. Its home realm is Physics / Robotics / Autonomous.

#	Muscle	Knob	What Omni-Compass does with it	Organisms
2054	traction effort limit	power	holds the power share inside its cover; full at once past the wall	Physics; stack; tower
2055	regenerative braking share	power	holds the power share inside its cover; full at once past the wall	Physics; stack; tower
2056	coasting speed target	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Physics; stack; tower
2057	dwel time schedule	admission	reads it; left to its own controller	Physics; stack; tower
2058	hvac duty cycle	power	holds the power share inside its cover; full at once past the wall	Physics; stack; tower
2059	train auxiliary power budget	power	holds the power share inside its cover; full at once past the wall	Physics; stack; tower

#	Muscle	Knob	What Omni-Compass does with it	Organisms
2060	platform approach speed	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Physics; stack; tower
2061	acceleration rate setpoint	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Physics; stack; tower
2062	wheel slip protection mode	admission	reads it; left to its own controller	Physics; stack; tower
2063	traction motor thermal derate	power	holds the power share inside its cover; full at once past the wall	Physics; stack; tower
2064	catenary voltage limit	power	holds the power share inside its cover; full at once past the wall	Physics; stack; tower
2065	timetable recovery margin	admission	reads it; left to its own controller	Physics; stack; tower
2066	headway target	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Physics; stack; tower
2067	door release hold	admission	reads it; left to its own controller	Physics; stack; tower

Marine Propulsion & Vessel Automation (12 muscles)

Each is a motor-driven axis: a robot joint, a flight-control surface, a reaction wheel, a traction motor; its class is `marine_propulsion` (a ship's propulsion shaft and power plant), reached through the vessel automation and power management system (modelled only). Omni-Compass reads its tracking error and its load. Its home realm is Physics / Robotics / Autonomous.

#	Muscle	Knob	What Omni-Compass does with it	Organisms
2068	shaft speed target	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Physics; stack; tower
2069	propeller pitch limit	power	holds the power share inside its cover; full at once past the wall	Physics; stack; tower
2070	shaft power limit	power	holds the power share inside its cover; full at once past the wall	Physics; stack; tower
2071	bow thruster duty	power	holds the power share inside its cover; full at once past the wall	Physics; stack; tower
2072	ballast pump mode	admission	reads it; left to its own controller	Physics; stack; tower
2073	engine load sharing setpoint	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Physics; stack; tower
2074	slow steaming speed target	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Physics; stack; tower
2075	auxiliary engine staging	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Physics; stack; tower
2076	shore power mode	admission	reads it; left to its own controller	Physics; stack; tower
2077	hotel load budget	power	holds the power share inside its cover; full at once past the wall	Physics; stack; tower
2078	trim target	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Physics; stack; tower

#	Muscle	Knob	What Omni-Compass does with it	Organisms
2079	rudder rate limit	power	holds the power share inside its cover; full at once past the wall	Physics; stack; tower

Elevators & Vertical Transport (12 muscles)

Each is a motor-driven axis: a robot joint, a flight-control surface, a reaction wheel, a traction motor; its class is `elevator_hoist` (an elevator or escalator drive and its group control), reached through the lift controller and destination dispatch (modelled only). Omni-Compass reads its tracking error and its load. Its home realm is Physics / Robotics / Autonomous.

#	Muscle	Knob	What Omni-Compass does with it	Organisms
2146	hoist speed target	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Physics; stack; tower
2147	acceleration limit	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Physics; stack; tower
2148	regen drive mode	power	holds the power share inside its cover; full at once past the wall	Physics; stack; tower
2149	standby power mode	admission	reads it; left to its own controller	Physics; stack; tower
2150	destination dispatch schedule	admission	reads it; left to its own controller	Physics; stack; tower
2151	car parking mode	admission	reads it; left to its own controller	Physics; stack; tower
2152	door dwell hold	admission	reads it; left to its own controller	Physics; stack; tower
2153	escalator speed target	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Physics; stack; tower
2154	group capacity mode	admission	reads it; left to its own controller	Physics; stack; tower
2155	motor thermal derate	power	holds the power share inside its cover; full at once past the wall	Physics; stack; tower
2156	brake test schedule	admission	reads it; left to its own controller	Physics; stack; tower
2157	fire recall mode	admission	reads it; left to its own controller	Physics; stack; tower

Realm: Energy / Facility / Industrial (213 muscles of its own, 470 in its organism with the spine)

Building & Critical Environment HVAC (15 muscles)

Each is a data hall or building zone cooled by chillers, air handlers or fans; its class is `building` (a building zone), reached through the building management system. Omni-Compass reads its temperature against its limit. Its home realm is Energy / Facility / Industrial.

#	Muscle	Knob	What Omni-Compass does with it	Organisms
1094	setpoint reset trim and respond	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
1095	duct static pressure setpoint	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower

#	Muscle	Knob	What Omni-Compass does with it	Organisms
1096	optimal start stop	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Energy; stack; tower
AUDIT-0 081	zone temperature target	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
AUDIT-0 082	zone airflow	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Energy; stack; tower
AUDIT-0 083	ahu fan speed	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Energy; stack; tower
AUDIT-0 084	damper position	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Energy; stack; tower
AUDIT-0 085	economizer position	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Energy; stack; tower
AUDIT-0 086	boiler setpoint	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
AUDIT-0 087	heat pump mode	admission	reads it; left to its own controller	Energy; stack; tower
AUDIT-0 088	humidity target	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
AUDIT-0 089	occupancy ventilation	admission	reads it; left to its own controller	Energy; stack; tower
AUDIT-0 090	building demand limit	power	holds the power share inside its cover; full at once past the wall	Energy; stack; tower
AUDIT-0 091	thermal storage dispatch	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Energy; stack; tower
AUDIT-0 092	hvac emergency mode	admission	reads it; left to its own controller	Energy; stack; tower

Energy Storage & Microgrid (19 muscles)

Each is a battery, UPS or microgrid store with a load and a grid connection; its class is microgrid (a microgrid battery and loads), reached through the inverter or energy management system (IEEE 2030.5, SunSpec, OpenADR). Omni-Compass reads its draw on the grid connection and its reserve. Its home realm is Energy / Facility / Industrial.

#	Muscle	Knob	What Omni-Compass does with it	Organisms
513	battery charge power	power	holds the power share inside its cover; full at once past the wall	Energy; stack; tower
514	battery discharge power	power	holds the power share inside its cover; full at once past the wall	Energy; stack; tower
515	battery soc reserve	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower

#	Muscle	Knob	What Omni-Compass does with it	Organisms
516	grid import limit	power	holds the power share inside its cover; full at once past the wall	Energy; stack; tower
517	grid export limit	power	holds the power share inside its cover; full at once past the wall	Energy; stack; tower
518	pv curtailment	power	holds the power share inside its cover; full at once past the wall	Energy; stack; tower
519	ev charge power	power	holds the power share inside its cover; full at once past the wall	Energy; stack; tower
520	heat pump power	power	holds the power share inside its cover; full at once past the wall	Energy; stack; tower
521	electrolyzer power	power	holds the power share inside its cover; full at once past the wall	Energy; stack; tower
522	microgrid demand limit	power	holds the power share inside its cover; full at once past the wall	Energy; stack; tower
523	peak shaving	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Energy; stack; tower
524	time of use schedule	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
525	energy load shed	admission	reads it; left to its own controller	Energy; stack; tower
526	flex load admission	admission	reads it; left to its own controller	Energy; stack; tower
527	storage dispatch	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Energy; stack; tower
528	microgrid emergency reserve	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
1085	volt var	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Energy; stack; tower
1086	volt watt	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Energy; stack; tower
1087	constant power factor reactive power	power	holds the power share inside its cover; full at once past the wall	Energy; stack; tower

Facility & Grid Optimization (15 muscles)

Each is a battery, UPS or microgrid store with a load and a grid connection; its class is **facility** (a facility battery behind the meter), reached through the energy management system. Omni-Compass reads its draw on the grid connection and its reserve. Its home realm is Energy / Facility / Industrial.

#	Muscle	Knob	What Omni-Compass does with it	Organisms
561	facility power budget	power	holds the power share inside its cover; full at once past the wall	Energy; stack; tower
562	utility demand limit	power	holds the power share inside its cover; full at once past the wall	Energy; stack; tower
563	demand response	admission	reads it; left to its own controller	Energy; stack; tower
564	electricity price gate	admission	reads it; left to its own controller	Energy; stack; tower
565	carbon intensity gate	admission	reads it; left to its own controller	Energy; stack; tower

#	Muscle	Knob	What Omni-Compass does with it	Organisms
566	renewable dispatch	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Energy; stack; tower
567	generator start stop	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Energy; stack; tower
568	site battery dispatch	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Energy; stack; tower
569	pue target	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
570	cooling power budget	power	holds the power share inside its cover; full at once past the wall	Energy; stack; tower
571	it power budget	power	holds the power share inside its cover; full at once past the wall	Energy; stack; tower
572	rack power allocation	power	holds the power share inside its cover; full at once past the wall	Energy; stack; tower
573	facility peak guard	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Energy; stack; tower
574	grid frequency response	admission	reads it; left to its own controller	Energy; stack; tower
575	facility islanding	admission	reads it; left to its own controller	Energy; stack; tower

Grid Transmission & Distribution (12 muscles)

Each is a regulated process: a temperature, pressure, flow, level, voltage or composition held at a setpoint; its class is feeder_voltage (a distribution feeder's voltage), reached through the distribution management system (IEC 61850, DNP3; modelled only). Omni-Compass reads how far the process sits from its band. Its home realm is Energy / Facility / Industrial.

#	Muscle	Knob	What Omni-Compass does with it	Organisms
AUDIT-0033	capacitor bank switch	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Energy; stack; tower
AUDIT-0034	voltage regulator tap	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
AUDIT-0035	transformer tap	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
AUDIT-0036	inverter real power	power	holds the power share inside its cover; full at once past the wall	Energy; stack; tower
AUDIT-0037	inverter reactive power	power	holds the power share inside its cover; full at once past the wall	Energy; stack; tower
AUDIT-0038	feeder voltage target	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
AUDIT-0039	feeder load transfer	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Energy; stack; tower
AUDIT-0040	distribution storage dispatch	admission	reads it; left to its own controller	Energy; stack; tower

#	Muscle	Knob	What Omni-Compass does with it	Organisms
AUDIT-0041	demand response dispatch	admission	reads it; left to its own controller	Energy; stack; tower
AUDIT-0042	frequency droop setpoint	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
AUDIT-0043	grid protection mode	admission	reads it; left to its own controller	Energy; stack; tower
AUDIT-0044	grid restoration sequence	admission	reads it; left to its own controller	Energy; stack; tower

Industrial PLC & Process Automation (18 muscles)

Each is a regulated process: a temperature, pressure, flow, level, voltage or composition held at a setpoint; its class is process (an industrial process loop), reached through the PLC or DCS (OPC UA, EtherNet/IP, PROFINET). Omni-Compass reads how far the process sits from its band. Its home realm is Energy / Facility / Industrial.

#	Muscle	Knob	What Omni-Compass does with it	Organisms
465	plc cycle authority	admission	reads it; left to its own controller	Energy; stack; tower
466	machine cell admission	admission	reads it; left to its own controller	Energy; stack; tower
467	valve position	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
468	pump flow	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Energy; stack; tower
469	compressor speed	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Energy; stack; tower
470	heater power	power	holds the power share inside its cover; full at once past the wall	Energy; stack; tower
471	furnace setpoint	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
472	pressure setpoint	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
473	temperature setpoint	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
474	mass flow setpoint	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
475	tank level target	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
476	conveyor rate	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Energy; stack; tower
477	feed rate	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Energy; stack; tower
478	purge vent action	admission	reads it; left to its own controller	Energy; stack; tower
1088	controller mode	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Energy; stack; tower
1089	alarm limits	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Energy; stack; tower

#	Muscle	Knob	What Omni-Compass does with it	Organisms
1090	safety interlock trip	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Energy; stack; tower
1091	opc ua writes	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Energy; stack; tower

Semiconductor Fab & Precision Manufacturing (13 muscles)

Each is a regulated process: a temperature, pressure, flow, level, voltage or composition held at a setpoint; its class is chamber (a controlled chamber (clean room, kiln, reactor)), reached through the PLC or DCS. Omni-Compass reads how far the process sits from its band. Its home realm is Energy / Facility / Industrial.

#	Muscle	Knob	What Omni-Compass does with it	Organisms
1092	run to run control	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Energy; stack; tower
1093	idle sleep mode	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Energy; stack; tower
AUDIT-00 57	tool job dispatch	admission	reads it; left to its own controller	Energy; stack; tower
AUDIT-00 58	wafer route	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Energy; stack; tower
AUDIT-00 59	chamber recipe selection	admission	reads it; left to its own controller	Energy; stack; tower
AUDIT-00 60	chamber temperature	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
AUDIT-00 61	chamber pressure	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
AUDIT-00 62	gas flow	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Energy; stack; tower
AUDIT-00 63	rf power	power	holds the power share inside its cover; full at once past the wall	Energy; stack; tower
AUDIT-00 64	vacuum pump speed	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Energy; stack; tower
AUDIT-00 65	robot transfer rate	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Energy; stack; tower
AUDIT-00 66	lot priority	admission	reads it; left to its own controller	Energy; stack; tower
AUDIT-00 68	tool quarantine	admission	reads it; left to its own controller	Energy; stack; tower

Water Wastewater & Pumping (12 muscles)

Each is a regulated process: a temperature, pressure, flow, level, voltage or composition held at a setpoint; its class is water (a water or pumping process), reached through the SCADA system (Modbus, DNP3). Omni-Compass reads how far the process sits from its band. Its home realm is Energy / Facility / Industrial.

#	Muscle	Knob	What Omni-Compass does with it	Organisms
AUDIT-0069	pump speed water	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Energy; stack; tower
AUDIT-0070	valve position water	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
AUDIT-0071	reservoir level target	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
AUDIT-0072	line pressure target	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
AUDIT-0073	flow target water	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
AUDIT-0074	aeration rate	power	holds the power share inside its cover; full at once past the wall	Energy; stack; tower
AUDIT-0075	chemical dose rate	power	holds the power share inside its cover; full at once past the wall	Energy; stack; tower
AUDIT-0076	filtration backwash	admission	reads it; left to its own controller	Energy; stack; tower
AUDIT-0077	lift station dispatch	admission	reads it; left to its own controller	Energy; stack; tower
AUDIT-0078	leak isolation	admission	reads it; left to its own controller	Energy; stack; tower
AUDIT-0079	water demand shed	admission	reads it; left to its own controller	Energy; stack; tower
AUDIT-0080	water emergency shutdown	admission	reads it; left to its own controller	Energy; stack; tower

Healthcare Critical Environments (13 muscles)

Each is a data hall or building zone cooled by chillers, air handlers or fans; its class is hospital (a hospital's critical rooms (operating rooms, isolation, pharmacy, imaging suites)), reached through the hospital's building management system (BACnet). Omni-Compass reads its temperature against its limit. Its home realm is Energy / Facility / Industrial.

#	Muscle	Knob	What Omni-Compass does with it	Organisms
2001	operating room air change setpoint	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
2002	isolation room pressure target	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
2003	patient room temperature setpoint	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower

#	Muscle	Knob	What Omni-Compass does with it	Organisms
2004	surgical suite humidity target	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
2005	ahu supply air temperature	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
2006	pharmacy cold room setpoint	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
2007	sterile storage humidity setpoint	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
2008	imaging suite cooling capacity	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Energy; stack; tower
2009	chiller plant staging	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Energy; stack; tower
2010	exhaust fan capacity	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Energy; stack; tower
2011	ward night setback mode	admission	reads it; left to its own controller	Energy; stack; tower
2012	medical gas plant demand limit	power	holds the power share inside its cover; full at once past the wall	Energy; stack; tower
2013	emergency power load shed	admission	reads it; left to its own controller	Energy; stack; tower

Agriculture & Irrigation (14 muscles)

Each is a regulated process: a temperature, pressure, flow, level, voltage or composition held at a setpoint; its class is irrigation (pumps, pressure and climate on a farm or in a greenhouse), reached through the pump station controller, the pivot panel or the climate computer (Modbus, the vendor cloud). Omni-Compass reads how far the process sits from its band. Its home realm is Energy / Facility / Industrial.

#	Muscle	Knob	What Omni-Compass does with it	Organisms
2026	irrigation pump speed	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Energy; stack; tower
2027	mainline pressure setpoint	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
2028	soil moisture target	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
2029	pivot speed setpoint	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
2030	fertigation dose	power	holds the power share inside its cover; full at once past the wall	Energy; stack; tower
2031	greenhouse temperature setpoint	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower

#	Muscle	Knob	What Omni-Compass does with it	Organisms
2032	greenhouse co2 target	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
2033	vent position setpoint	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
2034	grain dryer temperature setpoint	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
2035	cold storage temperature setpoint	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
2036	barn ventilation rate	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Energy; stack; tower
2037	milking vacuum level setpoint	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
2038	well drawdown level target	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
2039	drip zone dispatch	admission	reads it; left to its own controller	Energy; stack; tower

Oil & Gas Pipelines (14 muscles)

Each is a regulated process: a temperature, pressure, flow, level, voltage or composition held at a setpoint; its class is pipeline (a pipeline segment with its compressors, pumps and valves), reached through pipeline SCADA (DNP3, Modbus, OPC UA). Omni-Compass reads how far the process sits from its band. Its home realm is Energy / Facility / Industrial.

#	Muscle	Knob	What Omni-Compass does with it	Organisms
2040	compressor discharge pressure setpoint	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
2041	pump station suction pressure target	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
2042	line pack target	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
2043	pipeline flow setpoint	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
2044	compressor unit staging	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Energy; stack; tower
2045	vfd pump speed	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Energy; stack; tower
2046	terminal tank level target	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower

#	Muscle	Knob	What Omni-Compass does with it	Organisms
2047	leak detection shutdown	admission	reads it; left to its own controller	Energy; stack; tower
2048	batch interface dispatch	admission	reads it; left to its own controller	Energy; stack; tower
2049	heater outlet temperature setpoint	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
2050	drag reducing agent dose	power	holds the power share inside its cover; full at once past the wall	Energy; stack; tower
2051	valve position target	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
2052	cathodic protection voltage	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
2053	pressure protection priority	admission	reads it; left to its own controller	Energy; stack; tower

Mining & Mineral Processing (14 muscles)

Each is a regulated process: a temperature, pressure, flow, level, voltage or composition held at a setpoint; its class is mill (a grinding, flotation or materials-handling circuit), reached through the plant DCS or PLC (OPC UA). Omni-Compass reads how far the process sits from its band. Its home realm is Energy / Facility / Industrial.

#	Muscle	Knob	What Omni-Compass does with it	Organisms
2092	sag mill load setpoint	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
2093	mill speed target	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
2094	crusher gap setpoint	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
2095	flotation aeration rate	power	holds the power share inside its cover; full at once past the wall	Energy; stack; tower
2096	cyclone feed pressure target	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
2097	thickener underflow density target	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
2098	reagent dose	power	holds the power share inside its cover; full at once past the wall	Energy; stack; tower
2099	conveyor speed setpoint	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
2100	dewatering pump level target	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
2101	ventilation on demand airflow	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Energy; stack; tower

#	Muscle	Knob	What Omni-Compass does with it	Organisms
2102	stockpile level target	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
2103	slurry pump speed	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Energy; stack; tower
2104	tailings discharge shutdown	admission	reads it; left to its own controller	Energy; stack; tower
2105	ore blend dispatch	admission	reads it; left to its own controller	Energy; stack; tower

District Heating & Cooling (13 muscles)

Each is a regulated process: a temperature, pressure, flow, level, voltage or composition held at a setpoint; its class is `district_heat` (a district heating or cooling network), reached through the network control system and substation controllers (Modbus, M-Bus). Omni-Compass reads how far the process sits from its band. Its home realm is Energy / Facility / Industrial.

#	Muscle	Knob	What Omni-Compass does with it	Organisms
2106	supply temperature setpoint	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
2107	return temperature target	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
2108	differential pressure setpoint	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
2109	network pump speed	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Energy; stack; tower
2110	heat pump staging	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Energy; stack; tower
2111	chp dispatch	admission	reads it; left to its own controller	Energy; stack; tower
2112	thermal storage level target	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
2113	peak boiler heater output	power	holds the power share inside its cover; full at once past the wall	Energy; stack; tower
2114	substation flow limit	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Energy; stack; tower
2115	peak demand shed	admission	reads it; left to its own controller	Energy; stack; tower
2116	outdoor reset setpoint	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
2117	cooling network supply temperature	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower

#	Muscle	Knob	What Omni-Compass does with it	Organisms
2118	cooling storage level target	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower

Power Generation & Turbine Control (14 muscles)

Each is a regulated process: a temperature, pressure, flow, level, voltage or composition held at a setpoint; its class is turbine (a generating unit's governor, boiler or excitation loop), reached through the plant DCS and governor (OPC UA, IEC 61850; modelled only). Omni-Compass reads how far the process sits from its band. Its home realm is Energy / Facility / Industrial.

#	Muscle	Knob	What Omni-Compass does with it	Organisms
2119	turbine speed droop setpoint	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
2120	unit load setpoint	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
2121	boiler steam pressure setpoint	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
2122	feedwater level target	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
2123	agc participation limit	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Energy; stack; tower
2124	excitation voltage setpoint	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
2125	hydro gate position target	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
2126	combustion air ratio setpoint	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
2127	cooling water flow setpoint	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
2128	inlet guide vane position	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
2129	reserve dispatch priority	admission	reads it; left to its own controller	Energy; stack; tower
2130	turbine ramp rate limit	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Energy; stack; tower
2131	emissions shutdown	admission	reads it; left to its own controller	Energy; stack; tower
2132	nuclear rod position target	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower

Renewable Generation & Inverter Control (13 muscles)

Each is a regulated process: a temperature, pressure, flow, level, voltage or composition held at a setpoint; its class is inverter (a plant of grid-support inverters or wind turbines), reached through the plant controller and inverter settings (SunSpec, IEC 61850, IEEE 1547). Omni-Compass reads how far the process sits from its band. Its home realm is Energy / Facility / Industrial.

#	Muscle	Knob	What Omni-Compass does with it	Organisms
2133	inverter volt var setpoint	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
2134	inverter volt watt setpoint	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
2135	active power curtailment	power	holds the power share inside its cover; full at once past the wall	Energy; stack; tower
2136	plant reactive power target	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
2137	wind turbine yaw offset	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Energy; stack; tower
2138	pitch angle limit	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Energy; stack; tower
2139	rotor speed setpoint	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
2140	inverter frequency droop setpoint	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
2141	ramp rate limit	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Energy; stack; tower
2142	tracker stow mode	admission	reads it; left to its own controller	Energy; stack; tower
2143	string mppt voltage setpoint	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
2144	noise mode schedule	admission	reads it; left to its own controller	Energy; stack; tower
2145	ice detection shutdown	admission	reads it; left to its own controller	Energy; stack; tower

Pharmaceutical & Food Manufacturing (14 muscles)

Each is a regulated process: a temperature, pressure, flow, level, voltage or composition held at a setpoint; its class is batch_reactor (a GMP batch reactor, fermenter or food process), reached through the batch control system (ISA-88, OPC UA). Omni-Compass reads how far the process sits from its band. Its home realm is Energy / Facility / Industrial.

#	Muscle	Knob	What Omni-Compass does with it	Organisms
2158	reactor temperature setpoint	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
2159	agitator speed setpoint	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
2160	fermenter dissolved oxygen target	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
2161	ph setpoint	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
2162	pasteurizer holding temperature setpoint	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
2163	freezer tunnel temperature setpoint	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
2164	cip cycle dispatch	admission	reads it; left to its own controller	Energy; stack; tower
2165	cleanroom pressure cascade setpoint	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
2166	lyophilizer shelf temperature setpoint	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
2167	chromatography flow setpoint	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
2168	steam sterilizer cycle priority	admission	reads it; left to its own controller	Energy; stack; tower
2169	oven zone temperature setpoint	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Energy; stack; tower
2170	refrigerant compressor authority	admission	reads it; left to its own controller	Energy; stack; tower
2171	batch hold quarantine	admission	reads it; left to its own controller	Energy; stack; tower

Realm: Distribution / Specialized (183 muscles of its own, 440 in its organism with the spine)

Cache & Memory Services (16 muscles)

Each is a pool of servers, pods, GPUs, links or disks serving a stream of requests or jobs; its class is server (an application service on servers or pods), reached through the Kubernetes API (HPA target and floor, pod resize). Omni-Compass reads how far its queue and its load sit toward the service line (queue or lateness, and load above half). Its home realm is Distribution / Specialized.

#	Muscle	Knob	What Omni-Compass does with it	Organisms
305	cache size	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
306	cache ttl	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Distribution; stack; tower
307	cache eviction policy	admission	reads it; left to its own controller	Distribution; stack; tower
308	cache replicas	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
309	cache sharding	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
310	cache prefetch	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
311	cache writeback rate	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
312	cache admission	admission	reads it; left to its own controller	Distribution; stack; tower
313	hot key isolation	admission	reads it; left to its own controller	Distribution; stack; tower
314	cache connection limit	admission	reads it; left to its own controller	Distribution; stack; tower
315	cache memory limit	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
316	cache compression	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
317	cache warmup	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
318	cache failover	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
319	cache flush rate	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
1106	memory size threads	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower

Commerce & Payment Systems (15 muscles)

Each is a pool of servers, pods, GPUs, links or disks serving a stream of requests or jobs; its class is commerce (a customer-facing service with bursts), reached through the Kubernetes API. Omni-Compass reads how far its queue and its load sit toward the service line (queue or lateness, and load above half). Its home realm is Distribution / Specialized.

#	Muscle	Knob	What Omni-Compass does with it	Organisms
401	payment admission	admission	reads it; left to its own controller	Distribution; stack; tower
402	payment concurrency	admission	reads it; left to its own controller	Distribution; stack; tower

#	Muscle	Knob	What Omni-Compass does with it	Organisms
403	payment retry	admission	reads it; left to its own controller	Distribution; stack; tower
404	payment timeout	admission	reads it; left to its own controller	Distribution; stack; tower
405	fraud review gate	admission	reads it; left to its own controller	Distribution; stack; tower
406	authorization route	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
407	processor route weight	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Distribution; stack; tower
408	transaction queue limit	admission	reads it; left to its own controller	Distribution; stack; tower
409	idempotency window	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Distribution; stack; tower
410	order reservation	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
411	inventory hold	admission	reads it; left to its own controller	Distribution; stack; tower
412	checkout load shed	admission	reads it; left to its own controller	Distribution; stack; tower
413	refund queue rate	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
414	settlement batch	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
415	payment failover	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower

Data Analytics & ETL (15 muscles)

Each is a pool of servers, pods, GPUs, links or disks serving a stream of requests or jobs; its class is batch (batch and queued jobs), reached through the job queue (Kubernetes Jobs, Slurm, Spark). Omni-Compass reads how far its queue and its load sit toward the service line (queue or lateness, and load above half). Its home realm is Distribution / Specialized.

#	Muscle	Knob	What Omni-Compass does with it	Organisms
338	executor size	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
339	dynamic allocation min	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
340	dynamic allocation max	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
341	shuffle partitions	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
342	shuffle bandwidth	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower

#	Muscle	Knob	What Omni-Compass does with it	Organisms
343	etl concurrency	admission	reads it; left to its own controller	Distribution; stack; tower
344	stage parallelism	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
345	query slots	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
346	spill threshold	admission	reads it; left to its own controller	Distribution; stack; tower
347	cache fraction	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Distribution; stack; tower
348	batch interval	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Distribution; stack; tower
349	stream backpressure	admission	reads it; left to its own controller	Distribution; stack; tower
350	data locality wait	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
351	speculation policy	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
352	analytics admission	admission	reads it; left to its own controller	Distribution; stack; tower

Database & Transactions (19 muscles)

Each is a pool of servers, pods, GPUs, links or disks serving a stream of requests or jobs; its class is database (a database cluster), reached through the database operator (replicas, connection pools). Omni-Compass reads how far its queue and its load sit toward the service line (queue or lateness, and load above half). Its home realm is Distribution / Specialized.

#	Muscle	Knob	What Omni-Compass does with it	Organisms
273	db replicas	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
275	db memory	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
276	db cache	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
277	query concurrency	admission	reads it; left to its own controller	Distribution; stack; tower
278	read route	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
279	shard placement	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
280	db failover	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
281	replication lag gate	admission	reads it; left to its own controller	Distribution; stack; tower

#	Muscle	Knob	What Omni-Compass does with it	Organisms
282	db write throttle	admission	reads it; left to its own controller	Distribution; stack; tower
283	db pool resize	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
284	transaction concurrency	admission	reads it; left to its own controller	Distribution; stack; tower
285	lock timeout	admission	reads it; left to its own controller	Distribution; stack; tower
286	checkpoint rate	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
287	vacuum compaction rate	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
1099	parallel workers	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
1100	background writer	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
1101	synchronous replication	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
1102	buffer pool	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
1103	io capacity	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower

Messaging & Streaming (18 muscles)

Each is a pool of servers, pods, GPUs, links or disks serving a stream of requests or jobs; its class is server (an application service on servers or pods), reached through the Kubernetes API (HPA target and floor, pod resize). Omni-Compass reads how far its queue and its load sit toward the service line (queue or lateness, and load above half). Its home realm is Distribution / Specialized.

#	Muscle	Knob	What Omni-Compass does with it	Organisms
289	partition count	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
290	partition placement	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
291	producer quota	admission	reads it; left to its own controller	Distribution; stack; tower
292	consumer quota	admission	reads it; left to its own controller	Distribution; stack; tower
293	broker io quota	admission	reads it; left to its own controller	Distribution; stack; tower
294	message retention	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
295	queue depth limit	admission	reads it; left to its own controller	Distribution; stack; tower
296	consumer concurrency	admission	reads it; left to its own controller	Distribution; stack; tower

#	Muscle	Knob	What Omni-Compass does with it	Organisms
297	producer batch size	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
298	fetch batch size	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
299	rebalance rate	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
300	replication factor	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
301	retry backoff	admission	reads it; left to its own controller	Distribution; stack; tower
302	dead letter divert	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
303	stream priority	admission	reads it; left to its own controller	Distribution; stack; tower
304	broker failover	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
1104	prefetch	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
1105	memory disk alarm	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower

Runtime & Application (17 muscles)

Each is a pool of servers, pods, GPUs, links or disks serving a stream of requests or jobs; its class is server (an application service on servers or pods), reached through the Kubernetes API (HPA target and floor, pod resize). Omni-Compass reads how far its queue and its load sit toward the service line (queue or lateness, and load above half). Its home realm is Distribution / Specialized.

#	Muscle	Knob	What Omni-Compass does with it	Organisms
353	worker count	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
354	thread pool	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
355	jvm heap	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
356	gc budget	admission	reads it; left to its own controller	Distribution; stack; tower
357	connection pool runtime	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
358	application cache size	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower

#	Muscle	Knob	What Omni-Compass does with it	Organisms
359	runtime memory	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
360	async concurrency	admission	reads it; left to its own controller	Distribution; stack; tower
361	event loop workers	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
362	process count	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
363	request timeout	admission	reads it; left to its own controller	Distribution; stack; tower
364	background workers	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
365	runtime cpu budget	admission	reads it; left to its own controller	Distribution; stack; tower
366	runtime io budget	admission	reads it; left to its own controller	Distribution; stack; tower
367	runtime restart	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
1108	go runtime	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
1109	keepalive connection reuse	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower

Search, Indexing & Vector DB (16 muscles)

Each is a pool of servers, pods, GPUs, links or disks serving a stream of requests or jobs; its class is server (an application service on servers or pods), reached through the Kubernetes API (HPA target and floor, pod resize). Omni-Compass reads how far its queue and its load sit toward the service line (queue or lateness, and load above half). Its home realm is Distribution / Specialized.

#	Muscle	Knob	What Omni-Compass does with it	Organisms
321	index workers	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
322	index refresh rate	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
323	segment merge rate	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
324	search concurrency	admission	reads it; left to its own controller	Distribution; stack; tower
325	search timeout	admission	reads it; left to its own controller	Distribution; stack; tower
326	shard count	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
327	shard replication	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower

#	Muscle	Knob	What Omni-Compass does with it	Organisms
328	shard rebalance	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
329	vector search k	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
330	vector batch size	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
331	embedding workers	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
332	index memory budget	admission	reads it; left to its own controller	Distribution; stack; tower
333	query route	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
334	hot shard isolation	admission	reads it; left to its own controller	Distribution; stack; tower
335	search admission	admission	reads it; left to its own controller	Distribution; stack; tower
1107	circuit breakers	admission	reads it; left to its own controller	Distribution; stack; tower

Service Mesh & API Reliability (16 muscles)

Each is a pool of servers, pods, GPUs, links or disks serving a stream of requests or jobs; its class is `server` (an application service on servers or pods), reached through the Kubernetes API (HPA target and floor, pod resize). Omni-Compass reads how far its queue and its load sit toward the service line (queue or lateness, and load above half). Its home realm is Distribution / Specialized.

#	Muscle	Knob	What Omni-Compass does with it	Organisms
241	circuit breaker	admission	reads it; left to its own controller	Distribution; stack; tower
242	retry budget	admission	reads it; left to its own controller	Distribution; stack; tower
243	service timeout	admission	reads it; left to its own controller	Distribution; stack; tower
244	service concurrency	admission	reads it; left to its own controller	Distribution; stack; tower
245	connection pool	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
246	traffic divert	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
247	traffic mirror	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
248	canary weight	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Distribution; stack; tower
249	outlier ejection	admission	reads it; left to its own controller	Distribution; stack; tower
250	health threshold	admission	reads it; left to its own controller	Distribution; stack; tower
251	dns traffic weight	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Distribution; stack; tower

#	Muscle	Knob	What Omni-Compass does with it	Organisms
252	session affinity	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
253	request hedging	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
254	fault injection gate	admission	reads it; left to its own controller	Distribution; stack; tower
255	service failover	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
1097	load balancing policy	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower

Telecom RAN & Edge Radio (12 muscles)

Each is a pool of servers, pods, GPUs, links or disks serving a stream of requests or jobs; its class is ran (a radio access network cell or site), reached through the RAN controller (O-RAN interfaces). Omni-Compass reads how far its queue and its load sit toward the service line (queue or lateness, and load above half). Its home realm is Distribution / Specialized.

#	Muscle	Knob	What Omni-Compass does with it	Organisms
AUDIT-009	ran connection admission	admission	reads it; left to its own controller	Distribution; stack; tower
AUDIT-010	ran ue handover	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
AUDIT-011	ran cell traffic steering	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
AUDIT-012	ran slice resource budget	admission	reads it; left to its own controller	Distribution; stack; tower
AUDIT-013	ran prb allocation	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
AUDIT-014	ran scheduler weight	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Distribution; stack; tower
AUDIT-015	ran tx power	power	holds the power share inside its cover; full at once past the wall	Distribution; stack; tower
AUDIT-016	ran antenna tilt	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
AUDIT-017	ran carrier enable	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
AUDIT-018	ran cell sleep	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
AUDIT-019	ran du cu placement	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower

#	Muscle	Knob	What Omni-Compass does with it	Organisms
AUDIT-0020	ran fronthaul budget	admission	reads it; left to its own controller	Distribution; stack; tower

Workflow, Logistics & Fulfillment (15 muscles)

Each is a pool of servers, pods, GPUs, links or disks serving a stream of requests or jobs; its class is workflow (a workflow or pipeline engine), reached through the workflow engine's concurrency settings. Omni-Compass reads how far its queue and its load sit toward the service line (queue or lateness, and load above half). Its home realm is Distribution / Specialized.

#	Muscle	Knob	What Omni-Compass does with it	Organisms
417	workflow admission	admission	reads it; left to its own controller	Distribution; stack; tower
418	workflow worker rate	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
419	task queue rate	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
420	workflow retry	admission	reads it; left to its own controller	Distribution; stack; tower
421	workflow backoff	admission	reads it; left to its own controller	Distribution; stack; tower
422	workflow timeout	admission	reads it; left to its own controller	Distribution; stack; tower
423	inventory allocation	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
424	fulfillment route	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
425	warehouse queue	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
426	carrier selection	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
427	dispatch priority	admission	reads it; left to its own controller	Distribution; stack; tower
428	shipment batch	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
429	route replan	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
430	sla escalation	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
431	compensation action	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower

Medical Imaging & Clinical Systems (12 muscles)

Each is a pool of servers, pods, GPUs, links or disks serving a stream of requests or jobs; its class is clinical (clinical systems: imaging archives, records, interface engines, monitoring gateways), reached through the application's own scaling and admission settings (Kubernetes, the vendor console). Omni-Compass reads how far its queue and its load sit toward the service line (queue or lateness, and load above half). Its home realm is Distribution / Specialized.

#	Muscle	Knob	What Omni-Compass does with it	Organisms
2014	pacs archive tier target	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Distribution; stack; tower
2015	dicom router concurrency	admission	reads it; left to its own controller	Distribution; stack; tower
2016	ehr application replicas	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
2017	hl7 interface queue limit	admission	reads it; left to its own controller	Distribution; stack; tower
2018	fhir api rate limit	admission	reads it; left to its own controller	Distribution; stack; tower
2019	reconstruction gpu workers	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
2020	modality worklist timeout	admission	reads it; left to its own controller	Distribution; stack; tower
2021	patient monitoring gateway capacity	admission	reads it; left to its own controller	Distribution; stack; tower
2022	lab analyzer batch window	admission	reads it; left to its own controller	Distribution; stack; tower
2023	telehealth session admission	admission	reads it; left to its own controller	Distribution; stack; tower
2024	clinical backup window	admission	reads it; left to its own controller	Distribution; stack; tower
2025	imaging prefetch priority	admission	reads it; left to its own controller	Distribution; stack; tower

Ports & Maritime Logistics (12 muscles)

Each is a pool of servers, pods, GPUs, links or disks serving a stream of requests or jobs; its class is port (a container terminal's cranes and vehicles taking moves), reached through the terminal operating system and fleet controllers. Omni-Compass reads how far its queue and its load sit toward the service line (queue or lateness, and load above half). Its home realm is Distribution / Specialized.

#	Muscle	Knob	What Omni-Compass does with it	Organisms
2080	quay crane allocation	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
2081	yard crane fleet size	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
2082	berth window admission	admission	reads it; left to its own controller	Distribution; stack; tower

#	Muscle	Knob	What Omni-Compass does with it	Organisms
2083	truck gate rate limit	admission	reads it; left to its own controller	Distribution; stack; tower
2084	horizontal transport fleet size	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
2085	reefer plug power budget	power	holds the power share inside its cover; full at once past the wall	Distribution; stack; tower
2086	shore power connection capacity	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
2087	rail mounted gantry dispatch	capacity	holds it so the service sits in the middle of its band; more at once past the wall; gives back one unit at a time through the release gate	Distribution; stack; tower
2088	container dwell priority	admission	reads it; left to its own controller	Distribution; stack; tower
2089	vessel arrival pacing	admission	reads it; left to its own controller	Distribution; stack; tower
2090	stacking height target	setpoint	moves the setpoint inside its safe band toward the calm end while the service has room; back toward stress at once when it does not	Distribution; stack; tower
2091	equipment charging window	admission	reads it; left to its own controller	Distribution; stack; tower

© 2026 The Omni-Compass LLC. All rights reserved. **Evaluation and simulation use only.** Any commercial use, commercialization, monetization, production use, redistribution or hosted service of any part of Omni-Compass requires a signed, paid Omni-Compass Enterprise License from The Omni-Compass LLC. Patents, copyrights and trademarks filed in the USA. See LICENSE, NOTICE and DISCLOSURES.md.

APPENDIX

Appendix H. Source of the Engine

The engine, the compass law, the compass on every muscle, the two-wire GPU governor, the plug adapter and the nervous system, printed in full from the repository so that the book and the code can be read side by side. The engine file is sealed: its fingerprint is in `results/SEAL.json`, and `verify.py` fails if a byte of it changes.

`omnicompass/core.py`

240 lines.

```
"""Omni-Compass core: six-state ODE, RK4 integrator, bounded feedforward-plus-proportional controller.

State x = (E, U, I_U, S, B, B_dot).

(1) dE/dt = -alpha_E E + beta_int + beta_ext + v_eff
(2) dU/dt = mu U (1 - U^2) - (dE/dt)/E_max - lambda_U U + u, |u| <= U_AUTHORITY
(3) dI_U/dt = (1 - U) - sigma_1 E - delta S - lambda_I I_U
(4) v_eff = cos(omega_B t / 2) * c * tanh(lambda_0 + lambda_1 (U - 0.5) + lambda_2 S)
(5) Phi(S) = alpha_s S^2 / 2 + beta_s S^3 / 4 - delta S
(6) dS/dt = -dPhi/dS = delta - alpha_s S - (3/4) beta_s S^2
(7) dB/dt = B_dot ; dB_dot/dt = gamma_c delta S - (omega_B/Q_B) B_dot - omega_B^2 B
(8) R_B[n] = finite-difference audit of (7); never fed back into the state.

Controller (one micro step, zero-order hold across all four RK4 stages):
    u = clip(-f_U(x, t) + KP (sigma - U), -U_AUTHORITY, +U_AUTHORITY)
where f_U is the uncontrolled U drift and sigma in {-1, +1} is the locked target.
"""
from __future__ import annotations

import math
from dataclasses import dataclass, asdict, fields
from typing import Dict, List, Optional, Tuple

EPS = 1e-12
KP = 12.0
U_AUTHORITY = 25.0
MACRO_STEPS = 20
MICRO_STEPS = 10
```

```

MACRO_DT = 0.1
MICRO_DT = MACRO_DT / MICRO_STEPS
BASIN_TOL = 0.10
CONVEY_WINDOW = 5
CERT_WINDOW = 10

PARAMETER_RANGES: Dict[str, Tuple[float, float]] = {
    "alpha": (4.2, 4.2), "beta_int": (0.0, 1.0), "beta_ext": (0.0, 1.0), "k": (1.7, 1.7),
    "sigma_1": (0.38, 0.38), "delta": (0.01, 1.0), "gamma_1": (1.35, 1.35), "gamma_c": (1.0, 1.0),
    "lambda_0": (-2.0, 2.0), "lambda_1": (0.1, 3.0), "lambda_2": (0.0, 2.0), "c": (0.5, 5.0),
    "E_max": (1.0, 10.0), "omega_B": (0.1, 5.0), "Q_B": (0.5, 10.0), "alpha_s": (0.05, 0.25),
    "beta_s": (0.05, 0.25), "mu": (2.0, 6.0), "alpha_E": (4.2, 4.2), "alpha_U": (4.2, 4.2),
    "lambda_I": (0.5, 2.0), "lambda_U": (0.0, 0.5),
}
# alpha, k, gamma_1, alpha_U do not enter equations (1)-(8).

@dataclass
class State:
    E: float
    U: float
    I_U: float
    S: float
    B: float
    B_dot: float

    def as_tuple(self) -> Tuple[float, ...]:
        return (self.E, self.U, self.I_U, self.S, self.B, self.B_dot)

@dataclass
class Params:
    alpha: float = 4.2
    beta_int: float = 0.0
    beta_ext: float = 0.0
    k: float = 1.7
    sigma_1: float = 0.38
    delta: float = 0.5
    gamma_1: float = 1.35
    gamma_c: float = 1.0
    lambda_0: float = 0.0
    lambda_1: float = 1.0
    lambda_2: float = 0.5
    c: float = 1.0
    E_max: float = 2.0
    omega_B: float = 1.0
    Q_B: float = 2.5
    alpha_s: float = 0.12
    beta_s: float = 0.10
    mu: float = 4.0
    alpha_E: float = 4.2
    alpha_U: float = 4.2
    lambda_I: float = 1.0
    lambda_U: float = 0.1

# ----- equations
def v_eff(x: State, p: Params, t: float = 0.0) -> float:
    z = p.lambda_0 + p.lambda_1 * (x.U - 0.5) + p.lambda_2 * x.S
    return math.cos(0.5 * p.omega_B * t) * (p.c * math.tanh(z))

def phi(S: float, p: Params) -> float:
    return 0.5 * p.alpha_s * S * S + 0.25 * p.beta_s * S ** 3 - p.delta * S

def s_flow(S: float, p: Params) -> float:
    return p.delta - p.alpha_s * S - 0.75 * p.beta_s * S * S

def derivatives(x: State, p: Params, t: float, u: float = 0.0) -> State:
    v = v_eff(x, p, t)
    dE = -p.alpha_E * x.E + p.beta_int + p.beta_ext + v
    dU = p.mu * x.U * (1.0 - x.U * x.U) - dE / max(p.E_max, 1e-6) - p.lambda_U * x.U + u
    dI = (1.0 - x.U) - p.sigma_1 * x.E - p.delta * x.S - p.lambda_I * x.I_U
    dS = s_flow(x.S, p)
    dBd = p.gamma_c * p.delta * x.S - (p.omega_B / max(p.Q_B, EPS)) * x.B_dot - p.omega_B ** 2 * x.B
    return State(dE, dU, dI, dS, x.B_dot, dBd)

def _axpy(x: State, k: State, h: float) -> State:
    return State(x.E + h * k.E, x.U + h * k.U, x.I_U + h * k.I_U, x.S + h * k.S, x.B + h * k.B, x.B_dot + h * k.B_dot)

```

```

def rk4_step(x: State, p: Params, t: float, h: float, u: float = 0.0) -> State:
    k1 = derivatives(x, p, t, u)
    k2 = derivatives(_axpy(x, k1, 0.5 * h), p, t + 0.5 * h, u)
    k3 = derivatives(_axpy(x, k2, 0.5 * h), p, t + 0.5 * h, u)
    k4 = derivatives(_axpy(x, k3, h), p, t + h, u)
    c = h / 6.0
    return State(*(a + c * (b1 + 2 * b2 + 2 * b3 + b4) for a, b1, b2, b3, b4 in
        zip(x.as_tuple(), k1.as_tuple(), k2.as_tuple(), k3.as_tuple(), k4.as_tuple()))))

def bath_residual(prev: State, cur: State, nxt: State, p: Params, h: float) -> float:
    d2 = (nxt.B - 2.0 * cur.B + prev.B) / (h * h)
    dm = (cur.B - prev.B) / h
    return d2 + (p.omega_B / max(p.Q_B, EPS)) * dm + p.omega_B ** 2 * cur.B - p.gamma_c * p.delta * cur.S

# ----- controller
def control_command(x: State, p: Params, t: float, target: int) -> Tuple[float, float]:
    """Return (u, raw) for one micro step. raw is the unclipped command."""
    drift = derivatives(x, p, t, 0.0).U
    raw = -drift + KP * (float(target) - x.U)
    return max(-U_AUTHORITY, min(U_AUTHORITY, raw)), raw

def basin_sign(U: float) -> int:
    if abs(U - 1.0) <= BASIN_TOL:
        return 1
    if abs(U + 1.0) <= BASIN_TOL:
        return -1
    return 0

def nearest_basin_sign(U: float) -> int:
    return 1 if abs(U - 1.0) <= abs(U + 1.0) else -1

def macro_step(x: State, p: Params, t0: float, target: Optional[int],
    micro_steps: int = MICRO_STEPS, macro_dt: float = MACRO_DT) -> Tuple[State, Dict[str, float]]:
    """Advance one macro interval. target=None runs the plant with u=0 (observe)."""
    h = macro_dt / micro_steps
    t = t0
    u_abs = u_sq = u_peak = 0.0
    sat = 0
    for _ in range(micro_steps):
        if target is None:
            u = raw = 0.0
        else:
            u, raw = control_command(x, p, t, target)
            u_abs += abs(u) * h
            u_sq += u * u * h
            u_peak = max(u_peak, abs(u))
            sat += int(abs(raw) > U_AUTHORITY + EPS)
            x = rk4_step(x, p, t, h, u)
            t += h
    return x, {"u_abs": u_abs, "u_sq": u_sq, "u_peak": u_peak, "saturated": sat}

```



```

# ----- trajectory
@dataclass
class Result:
    final: State
    U_hist: List[float]
    target_sign: int
    born: int
    convey: int
    convey_start: int
    convey_confirm: int
    cert: int
    cert_start: int
    cert_complete: int
    max_streak: int
    actuator_abs_integral: float
    actuator_sq_integral: float
    actuator_peak: float
    actuator_saturated_periods: int

def simulate(x0: State, p: Params, target: Optional[int] = None, controlled: bool = True) -> Result:
    """Canonical 20-macro-step trajectory with CONVEY-5 / CERT-10 predicates.

    target defaults to the nearest basin of U0. controlled=False runs u=0.
    """
    tgt = nearest_basin_sign(x0.U) if target is None else int(target)
    b0 = basin_sign(x0.U)

    born = int(b0 != 0)
    streak = 1 if born else 0
    ssign = b0 if born else 0
    best = streak
    convey = cert = 0
    cs = cc = ts = tc = -1
    ua = usq = upk = 0.0
    nsat = 0
    x = x0
    hist = [x0.U]
    for m in range(1, MACRO_STEPS + 1):
        x, a = macro_step(x, p, (m - 1) * MACRO_DT, tgt if controlled else None)
        ua += a["u_abs"]; usq += a["u_sq"]; upk = max(upk, a["u_peak"]); nsat += a["saturated"]
        hist.append(x.U)
        b = basin_sign(x.U)
        if b:
            if b == ssign:
                streak += 1
            else:
                ssign, streak = b, 1
        else:
            ssign, streak = 0, 0
        best = max(best, streak)
        if not convey and streak >= CONVEY_WINDOW:
            convey, cs, cc = 1, m - CONVEY_WINDOW + 1, m
        if not cert and streak >= CERT_WINDOW:
            cert, ts, tc = 1, m - CERT_WINDOW + 1, m
    return Result(x, hist, tgt, born, convey, cs, cc, cert, ts, tc, best, ua, usq, upk, nsat)

def sample_state(rng) -> State:
    return State(float(rng.uniform(0.0, 1.0)), float(rng.uniform(-1.0, 1.0)), float(rng.uniform(0.0, 1.0)),
        float(rng.uniform(-1.0, 1.0)), float(rng.uniform(-1.0, 1.0)), float(rng.uniform(-1.0, 1.0)))

def sample_params(rng) -> Params:
    return Params(**{k: float(rng.uniform(lo, hi)) for k, (lo, hi) in PARAMETER_RANGES.items()})

def s_roots(p: Params) -> Tuple[float, float]:
    """Roots of equation (6): S_minus < 0 < S_plus."""
    a = 0.75 * p.beta_s
    d = math.sqrt(p.alpha_s ** 2 + 4.0 * a * p.delta)
    return (-p.alpha_s - d) / (2.0 * a), (-p.alpha_s + d) / (2.0 * a)

```

omnicompass/compass_low.py

174 lines.

```

# SPDX-License-Identifier: LicenseRef-OmniCompass-Evaluation-1.0
# Copyright (c) 2026 The Omni-Compass LLC. Evaluation and simulation use only; any other use requires a signed, paid
# Omni-Compass Enterprise License. See LICENSE.
"""The compass: one smooth law for every muscle, and the plug every muscle is wired through.

A muscle is a plug with one wire in and one wire out. The plug reads its meters and writes its lever; all the
thinking is here, in the brain, written once for every muscle.

The plug (Plug)
    read    the service reading (response time, temperature, queue) and the lever as it stands
    write    one value; the plug clips it to the cover, sends it, and reads back what the device took
    restore  the lever back to the snapshot, the value read once before the first write; it never moves after that
    one writer if the lever is found at a value the brain did not write, someone else owns it: the brain stops writing
            and leaves that value alone (restoring over it would fight the new owner)

Two bands
    cover    the outer band, the lever's hard range (the device's or the buyer's lowest and highest setting); every
            write is clipped to it, so nothing the math does can set the lever outside
    compass  the service reading as a position in its band, 0 to 1 (0: calm, 1: the service line); the brain pulls i
    ↪ t
            to the bottom of the compass, the center (0.5 by default); the walls run to 0.05 and 0.95, and the last 5%
    ↪ on
            each side is cushion (10% in all) before the line

The force (CompassLaw.force), the physics of a ball in a compass with the right friction
    pull     K_P x (position - center): gentle near the bottom, harder up the walls
    push     K_D x velocity: whatever is shoving the position (a load rising, heat building) is met by an equal and
            opposite push; it is also the friction that stops the ball sloshing past the bottom (with K_D chosen for

            critical damping it glides to the center and stops, no overshoot, no ringing)
    smooth   the force is A x tanh(raw / A): a spring near the center that bends over and flattens into its maximum A,
            never a corner, never a hammer; the same shape the engine already uses for v_eff (omnicompass/core.py)
    fail up  past the 0.95 wall the up side goes to its full force at once, and the down side may not act until the
            position is back inside the compass

Two forces (antagonist pairs)
    up       a positive force adds capacity, power, cooling or speed (scale out, raise a clock, start a chiller)
    down     a negative force takes it back (scale in, lower a limit, warm a setpoint)
    each side has its own gain (lever units per unit of force), because adding and taking back do not cost the same;
    a muscle with two wires (a GPU's clock range and power limit) gives each side its own lever; a one-wire muscle
    carries both directions on its one lever

The frozen engine (omnicompass/core.py) and its proofs are untouched; this is a separate law, run as its own arm.
"""
from __future__ import annotations

import math
from dataclasses import dataclass, field
from typing import Optional

def clamp(v, lo, hi):
    return lo if v < lo else hi if v > hi else v

@dataclass
class Band:

```

```

    """A reading's band: lo is calm (position 0), hi is the line (position 1)."""
    lo: float
    hi: float
    center: float = 0.5
    cushion: float = 0.05

    def position(self, x: float) -> float:
        return (x - self.lo) / (self.hi - self.lo)

    @property
    def wall_low(self) -> float:
        return self.cushion

    @property
    def wall_high(self) -> float:
        return 1.0 - self.cushion

    @dataclass
    class CompassLaw:
        """The brain's law for one muscle. dt: seconds between decisions; tau: the muscle's response time in seconds
        (how long the service reading takes to follow the lever); kd defaults to critical damping for that response."""
        band: Band
        dt: float
        tau: float
        kp: float = 1.0
        kd: Optional[float] = None
        authority: float = 1.0

        smooth: float = 0.5          # the velocity estimate's smoothing (0: raw difference, near 1: heavily smoothed)
        p: Optional[float] = field(default=None, init=False)
        v: float = field(default=0.0, init=False)

        def __post_init__(self):
            if self.kd is None:
                # a first-order muscle with response tau under a P-D pull: critical damping when kd = 2 sqrt(kp tau) - 1
                # scaled to the decision period; never negative
                self.kd = max(0.0, 2.0 * math.sqrt(self.kp * self.tau / self.dt) - 1.0) * self.dt

        def sense(self, reading: float) -> float:
            p = self.band.position(reading)
            if self.p is not None:
                self.v = self.smooth * self.v + (1.0 - self.smooth) * (p - self.p) / self.dt
            self.p = p
            return p

        def force(self, reading: float) -> float:
            """Positive: the up side (add capacity); negative: the down side (take it back). Bounded by the authority."""
            p = self.sense(reading)
            if p >= self.band.wall_high:
                return self.authority          # fail up: past the wall, full force up at once
            raw = self.kp * (p - self.band.center) + self.kd * self.v
            f = self.authority * math.tanh(raw / self.authority)
            if f < 0.0 and self.v > 0.0 and p > self.band.center:
                f = 0.0          # never take back while it climbs above the center
            return f

```

```

class ForeignWriter(Exception):
    pass

class Plug:
    """One wire in, one wire out. Subclasses implement _read_service, _read_lever and _send; everything else is here."""
    ↪ ""

    def __init__(self, lo: float, hi: float, tolerance: float = 1e-6):
        self.lo, self.hi, self.tol = lo, hi, tolerance
        self.snapshot: Optional[float] = None
        self.last_written: Optional[float] = None
        self.owned = True
        self.writes = 0
        self.clipped = 0

    # the wires -----
    def _read_service(self) -> float:
        raise NotImplementedError

    def _read_lever(self) -> float:
        raise NotImplementedError

    def _send(self, value: float) -> None:
        raise NotImplementedError

    # the contract -----
    def attach(self) -> float:
        """Read the lever once, before any write: the restore point. It never moves after this."""
        self.snapshot = self._read_lever()
        return self.snapshot

    def read(self) -> float:
        return self._read_service()

    def lever(self) -> float:
        return self._read_lever()

    def write(self, value: float) -> Optional[float]:
        """Clip to the cover, check nobody else moved the lever, send, read back. None: not written (not owned)."""
        if not self.owned:
            return None
        now = self._read_lever()
        expected = self.last_written if self.last_written is not None else self.snapshot
        if expected is not None and abs(now - expected) > self.tol:
            self.owned = False
            # someone else moved it: it is theirs now
            raise ForeignWriter(f"lever at {now}, last written {expected}")
        v = clamp(value, self.lo, self.hi)
        self.clipped += int(v != value)
        self._send(v)
        self.last_written = self._read_lever()
        self.writes += 1
        return self.last_written

    def restore(self) -> bool:
        """The lever back to the snapshot. Left alone when someone else owns it. True: at the snapshot (or not ours)."""
        ↪ ""
        if not self.owned or self.snapshot is None:
            return True
        if abs(self._read_lever() - self.snapshot) > self.tol:
            self._send(self.snapshot)
            self.last_written = None
        return abs(self._read_lever() - self.snapshot) <= self.tol

```

realms/compass_arm.py

184 lines.

```

# SPDX-License-Identifier: LicenseRef-OmniCompass-Evaluation-1.0
# Copyright (c) 2026 The Omni-Compass LLC. Evaluation and simulation use only; any other use requires a signed, paid
# Omni-Compass Enterprise License. See LICENSE.
"""The compass law on every realm muscle (omnicompass/compass_law.py): each plant's knob held by its own compass.

Each plant's service is read as one position in its compass (0 calm, 1 the line): the worst of its queue or lateness,
↳ its
load above half, and, for a storage site, its draw on the grid connection above half. The last period in violation
reads as past the wall (fail up). The compass's force moves the knob: a positive force adds capacity, power, cooling o
↳ r
protection, a negative force gives it back, each knob inside its own cover. At the kill every knob returns to native.

Knob by plant (the override key the plant already obeys, its native value, its cover, and which way is "more"):
compute_pool    power: limit share 0.4-1; replicas-type: the HPA target held at native (loosening it past the
operator's own costs the service time, the same rule as the card's speed floor); machine pools: one
machine released while the force is clearly down and the machines left cover the recent peak with
margin (RELEASE_MARGIN); admission: native
thermal_zone    setpoint between the band's stress (colder, more cooling) and calm ends; units: one released while th
↳ e
force is clearly down; power: native (a cooling cap moves heat later, never away)
energy_storage  reserve between the band's ends (a lower reserve gives the battery to the site); power: native
motion_axis    speed and effort: 0.4-1
process_loop    setpoint from native toward the band's calm end; actuator range 0.3-1; power 0.4-1
"""
from __future__ import annotations

from omnicompass.compass_law import Band, CompassLaw, clamp

UP, DOWN, RELEASE = 0.10, 0.02, -0.2
RELEASE_MARGIN = 0.3    # a machine goes back only when the rest covers the recent peak at 0.3 of the plant's own rel
↳ ease level
# 0.3: the largest margin at which no node pool is later than native (results/realms/RELEASE_MARGIN_SWEEP.md, the
# per-muscle sweep of 2026-10-05: 0.6 saved 2.0% energy with 32 of 42 pools later, 0.4 left 4 later, 0.3 none). A mach
↳ ine
# boots in minutes and a burst that comes in the meantime is served late, so speed first holds the machines (the same
# rule as the verdict on real Kubernetes: a machine goes back only where it costs no speed)

def position(plant) -> float:
    if plant.vhist and plant.vhist[-1]:
        return 1.0
    if plant.template == "thermal_zone":
        # a zone's service is its temperature against its limit
        P = plant.P
        return (plant.T - min(P["stress"], P["calm"])) / (P["t_limit"] - min(P["stress"], P["calm"]))
    o = plant.observe()
    p = max(o["queue_ratio"], (o["load_ratio"] - 0.5) / 0.5)
    if plant.template == "energy_storage":
        p = max(p, (o["power_stress"] - 0.5) / 0.5)
    return p

def release_safe(plant) -> bool:
    """A discrete unit goes back only when what is left covers the recent peak (the plant's own release window), with
RELEASE_MARGIN of headroom on a machine pool: a machine boots in minutes, so a burst that arrives after a release
is served late until it is back."""
    P = plant.P

```

```

if plant.template == "compute_pool":
    from .plants import RELEASE_WINDOW, RELEASE_FRAC
    n = plant.n
    return (n > P["n_min"] and not plant.pending and plant.Q <= 0.0
            and len(plant.ahist) == RELEASE_WINDOW * P["startup_steps"]
            and max(plant.ahist) / ((n - 1) * plant.mu(plant.cap)) <= RELEASE_FRAC * 0.95 * RELEASE_MARGIN)
u = plant.units_on
return u > 1 and plant.T <= plant.setpoint() and plant.qc_cmd_avg / ((u - 1) * P["q_unit_w"]) <= 0.8

def capping_saves(plant) -> bool:
    """Whether a lower power cap saves energy on this pool at all (race to idle, MECHANISM_OF_ACTION 9.6).

    With throughput  $\mu(c) \sim c^{\text{eps}}$  at power cap  $c$  and the autoscaler holding utilisation near its target  $u^*$ , the energy
    per unit of work is  $(p_{\text{idle}} / u^* + p_{\text{dyn}} c) / \mu(c)$ . Its slope at  $c = 1$  is positive (a lower cap saves) only if
     $\text{eps} < p_{\text{dyn}} / (p_{\text{idle}} / u^* + p_{\text{dyn}})$ . Where idle power dominates (a quantum computer's cryostat, a network switch),
    running slower keeps the machines on longer and costs more than it saves: the cap is left native."""
    import math
    P = plant.P
    eps = math.log(plant.mu(1.0) / plant.mu(0.99)) / math.log(1.0 / 0.99) # the machine's own speed curve at full po
    u = P.get("target", 0.7)
    return eps < P["p_dyn"] / (P["p_idle"] / u + P["p_dyn"])

def effort_cap_saves(plant) -> bool:
    """Whether a lower effort cap saves energy on this axis at all (MECHANISM_OF_ACTION 9.6).

    A cap on effort lowers the copper loss of a full-acceleration move,  $(J a_{\text{max}} / kt)^2 R$ , and stretches the move, du
    ring
    which the axis keeps drawing its idle power and turning against viscous friction,  $p_{\text{idle}} + b v_{\text{max}}^2$ . It pays only
    where the first exceeds the second (a reaction wheel); elsewhere (traction, a flight axis, a robot joint) a slower
    move costs more than the current it saves, and the cap is left native."""
    P = plant.P
    copper = (P["J"] * P["a_max"] / P["kt"]) ** 2 * P["R"]
    standing = P["p_idle"] + P["b"] * P["v_max"] ** 2
    return copper > standing

SLACK = 0.5 # a speed knob is offered only where the axis at full speed is busy at most half the time (0m
    ni v3)

def speed_slack(plant) -> bool:
    """Whether the axis has slack to spend on a slower, gentler move at all (0mni v3, docs/REALMS_PREREGISTRATION.md).

    The compass may ease a motion axis down to 0.4 of its speed, which stretches every move by up to 2.5 times. An axi
    s
    whose duty at full speed, task rate  $\times$  (move time + dwell), is already above one half has no room for that: tasks
    arrive in bunches, and a stretched move pushes the next ones past their deadline (a train on its timetable, a lift
    at
    rush hour). There the speed knob is left native, as the robot benchmark's paired trial leaves it
    (docs/ROBOTICS_PREREGISTRATION.md). Computed from the plant's own figures before any decision."""
    import math
    P = plant.P
    v, a, D = P["v_max"], P["a_max"], P["D"]

```

```

move = 2.0 * math.sqrt(D / a) if D < v * v / a else D / v + v / a
return P["task_rate"] * (move + P["dwell"]) &lt;= SLACK

def lever(plant, knob):
    """(override key, native value, low, high, sign): sign +1 when a higher value is more capacity."""
    P, t = plant.P, plant.template
    if knob == "admission":
        return None
    if t == "compute_pool":
        if knob == "power":
            return ("power", 1.0, 0.4, 1.0, 1) if capping_saves(plant) else None
        if P.get("ca"):
            return ("release", None, 0, 0, 0)
        return None # the HPA target is held at the operator's own
    if t == "thermal_zone":
        if knob == "setpoint":
            return ("setpoint", P["t_set"], min(P["stress"], P["calm"]), max(P["stress"], P["calm"]), -1)
        if knob == "capacity":
            return ("release", None, 0, 0, 0)
        # the cooling power cap gives nothing back: the heat the zone makes must leave it either way, so a lower cap o
    ↪ nly
        # lets the temperature drift up while the PI command grows and stages more units (and their fans). Native
        # (MECHANISM_OF_ACTION 9.6)
        return None
    if t == "energy_storage":
        if knob == "setpoint" and P.get("backup"):
            return None # a backup reserve (a UPS) is kept for an outage: nativ
    ↪ e
        if knob == "setpoint":
            # the reserve may go down to the stress end (the battery given to the site under strain), never above the
            # operator's own: a kWh held back in a calm battery is a kWh bought from the grid (MECHANISM_OF_ACTION 9.6)
    ↪ )
            return ("setpoint", P["reserve"], min(P["stress"], P["reserve"]), P["reserve"], -1)
            return None # a battery's power limit gives nothing back: native
    if t == "motion_axis":
        if knob == "power":
            return ("power", 1.0, 0.4, 1.0, 1) if effort_cap_saves(plant) else None
            return ("capacity", 1.0, 0.4, 1.0, 1) if speed_slack(plant) else None
    if t == "process_loop":
        if knob == "setpoint":
            lo, hi = min(P["sp"], P["calm"]), max(P["sp"], P["calm"])
            return ("setpoint", P["sp"], lo, hi, 1 if P["sp"] > P["calm"] else -1)
        if knob == "power":
            return ("power", 1.0, 0.4, 1.0, 1)
            return ("capacity", 1.0, 0.3, 1.0, 1)
    return None

def compass_apply(plant, knob, dt=1.0, tau=3.0):
    """One decision: the plant's compass reads its position and moves its knob. Returns the override written."""
    lv = lever(plant, knob)
    if lv is None:
        plant.override = {}
        return {}
    if not hasattr(plant, "_compass_law"):

```

```

    plant._compass_law = CompassLaw(Band(0.0, 1.0), dt=dt, tau=tau, kp=1.0, smooth=0.3)
    plant._compass_law.kd *= 3.0
    plant._compass_law_x = lv[1]
b = plant._compass_law
F = b.force(position(plant))
key, native, lo, hi, sign = lv
if key == "release":
    ok = F <= RELEASE and b.p <= b.band.center and release_safe(plant)
    plant.override = {"release": 1.0} if ok else {}
    return plant.override
span = hi - lo
if key == "setpoint" and plant.template == "energy_storage":
    # a battery's reserve is spent only where it buys service: at the connection's wall the reserve drops to the
    # stress end at once; anywhere else the operator's own reserve, so no round trip is paid for nothing (9.6)
    P = plant.P
    excess = plant.imp - P["p_lim_w"] # what the connection is over its limit by, now
    coverable = 0.0 &lt; excess &lt;= P["p_batt_w"] * plant.override.get("power", 1.0)
    x = lo if (b.p &gt;= b.band.wall_high and (coverable or plant.imp &gt;= 0.95 * P["p_lim_w"] and excess &lt;= 0.0)) else
↳ native
    plant._compass_law_x = x
    plant.override = {} if abs(x - native) &lt; 1e-9 else {key: x}
    return plant.override
if b.p &gt;= b.band.wall_high:
    x = hi if sign &gt; 0 else lo # fail up: full capacity at once
else:
    g = UP if F &gt; 0 else DOWN
    x = clamp(plant._compass_law_x + sign * g * F * span, lo, hi)
plant._compass_law_x = x

plant.override = {} if abs(x - native) &lt; 1e-9 else {key: x}
return plant.override

```

omni_controller/gpu_compass.py

416 lines.

```

# SPDX-License-Identifier: LicenseRef-OmniCompass-Evaluation-1.0
# Copyright (c) 2026 The Omni-Compass LLC. Evaluation and simulation use only; any other use requires a signed, paid
# Omni-Compass Enterprise License. See LICENSE.
"""Omni-Compass on one GPU through two wires: the compass law (omnicompass/compass_law.py) on a real card.

The card keeps its own control: NVIDIA's firmware still boosts and still protects the chip. Omni holds two settings th
↳ e
card already accepts, and nothing else:
up wire    the clock ceiling: how high boost may climb (nvidia-smi -lgc MIN,MAX; -rgc resets it)
down wire  the power limit: the lid above what that ceiling draws (nvidia-smi -pl W)

Every decision (--interval seconds):
read      nvidia-smi: power.draw, temperature, utilization, power.limit, clocks.sm; the workload's response times
position  the service as one place in its compass, 0 calm to 1 the line: response time only (the mean over the last
↳ at      --latency-window-s, between a tenth of the line, the bare service time, and the line; the 95th percentile
↳ y      or past the line, or a failed request, is past the wall). A card that is busy is doing its work; being bus
is not a breach and is not read as one
native    what the card does on its own, learned from its own meter before the compass may lower anything: while the
ceiling is at the top and the card is busy, its clock and its draw (the clock its own power limit holds it
at, and what that costs). Until --learn-samples busy readings are in, the ceiling stays at the top
race      while the card is saturated (utilization at or over --race-util: work is waiting), the ceiling goes to the
top and the lid to the start limit, so a burst is served at full speed; the compass paces only the slack
force     the compass: pull to the center, push against what is rising, tanh-bounded; past the 0.95 wall: fail up
verdict   how far the ceiling may go (omnicompass/verdict.py): while the service is calm, a paired trial: the ceilin
↳ g      at the top until --verdict-samples requests are measured, then one 15 MHz step past the deepest step

```



```

    already allowed until as many again are measured; the card's own time on each request (the workload's
    service_ms, the wait in the queue left out) is compared: at most --allow slower and the step is allowed,
    slower than that and it is refused and not tried again for --verdict-recheck decisions. Where no step
    passes, the ceiling stays at the top: the card runs as it does alone
law    down gain 0.0125, the compass's center at 0.4, the speed floor at the card's own busy clock (amendment 8)
steady while the card is saturated against its own power limit (work waiting and the draw at the limit), the
    firmware boosts a step, hits the limit and is knocked back: a sawtooth. The ceiling is then held at the
    card's own busy clock under that limit (what the sawtooth averages to), so the same watts serve the work
    without the knock-backs (amendment 9). Only under an operator's cap (the start limit under the card's
    factory limit, amendment 12): on the card's own limit the firmware's boost is left to serve the burst. It
    never holds under that clock, and blind always fails up
write  up wire: the ceiling moves by the force (fast up, gently down), inside its cover: from the card's own busy
    clock (never slower than native while there is work) to the top; down wire: the lid at the card's own busy
    draw plus --lid-headroom, never under it and never over the start limit, inside the declared envelope;
    fail up: ceiling to the top and the lid to the start limit at once
guards blind (meters unreadable, response times stale): fail up; heat (the card reports a thermal or hardware
    slowdown): never tighten; one writer: a power limit neither mine nor the start means another writer, so
    observe only from then on and exit 5; read-back: every write is read back
restore on exit, kill file or SIGTERM: clocks reset (-rgc) and the power limit back to the start, read back

```

The audit (--audit) carries the same records the bench reads from the one-wire governor (snapshot, decision with telemetry and decided_by, write and actuator for each power-limit write, restored), plus clock_write records.
Exit: 0 clean, 3 restore failed, 4 a write refused, 5 another writer.

```
"""
```

```
from __future__ import annotations
```

```
import argparse
import json
```

```
import math
import os
import shlex
import signal
import subprocess
import sys
import time
```

```
from omnicompass import master
from omnicompass.compass Law import Band, CompassLaw, clamp
from omnicompass.verdict import Verdict
from omni_controller.gpu_governor import query, snapshot, throttle, slowed, WriteFailed, SNAPSHOT
from omni_controller.muscles import latency_sense, latency_window
```

```

# the law (docs/GPU_PREREGISTRATION.md, amendment 8): the down gain (share of the top clock per unit of force), the co
↳ mpass's
# center (where it holds the response time, 0 the bare service time and 1 the line) and the speed floor (a share of th
↳ e
# card's own busy clock); the same in realms/gpu_card.py
LAW = {"down": 0.0125, "center": 0.4, "floor": 1.0}

```

```

def smi_run(smi, args):
    try:
        p = subprocess.run(shlex.split(smi) + args, capture_output=True, text=True, timeout=20)
        return p.returncode, p.stderr.strip()[:300]
    except (subprocess.SubprocessError, OSError) as e:

```

```

        return -1, f"{type(e).__name__}: {e}":300]

def query_top_clock(smi, g):
    try:
        out = subprocess.run(shlex.split(smi) + ["-i", str(g), "--query-gpu=clocks.max.sm", "--format=csv,noheader,nou
        ↵ nits"],
                               capture_output=True, text=True, timeout=10, check=True).stdout
        return float(out.strip().splitlines()[0])
    except (subprocess.SubprocessError, OSError, ValueError, IndexError):
        return None

def query_default_limit(smi, g):
    """The card's factory power limit (watts), or None if the driver will not say. A start limit under it means an
    operator's cap is underneath (amendment 12)."""
    try:
        out = subprocess.run(shlex.split(smi) + ["-i", str(g), "--query-gpu=power.default_limit", "--format=csv,nohead
        ↵ er,nounits"],
                               capture_output=True, text=True, timeout=10, check=True).stdout
        return float(out.strip().splitlines()[0])
    except (subprocess.SubprocessError, OSError, ValueError, IndexError):
        return None

def query_min_clock(smi, g):
    """The lowest graphics clock the card supports (the floor of a -lgc range); 300 MHz if the driver will not list th
    ↵ em."""
    try:
        out = subprocess.run(shlex.split(smi) + ["-i", str(g), "--query-supported-clocks=graphics", "--format=csv,nohe
        ↵ ader,nounits"],
                               capture_output=True, text=True, timeout=10, check=True).stdout
        return int(min(float(x) for x in out.split() if x.strip()))
    except (subprocess.SubprocessError, OSError, ValueError):
        return 300

class GpuCompass:
    def __init__(self, a):
        self.a = a
        self.g = int(str(a.gpus).split(",")[0])
        self.log = open(a.audit, "a")
        self.writes = 0
        self.foreign = False
        self.clock_set = False
        self.enforced_ok = query(a.smi, [self.g], enforced=True) is not None
        s = query(a.smi, [self.g], self.enforced_ok)
        if s is None:
            raise SystemExit("nvidia-smi unreadable at start: no snapshot, so I take no authority")
        snap = snapshot(a.smi, [self.g])
        self.start = s[self.g]["limit"]
        # an operator's cap underneath: the start limit under the card's factory limit (unknown: treated as no cap)
        dflt = query_default_limit(a.smi, self.g)
        self.capped = dflt is not None and self.start < dflt - 1.0
        self.expect = self.start
        self.top = query_top_clock(a.smi, self.g) or s[self.g]["clock_mhz"]

```

```

self.floor_w = max(float(a.floor_w or 0.0), s[self.g]["min"])
self.c_lo = a.clock_min_share * self.top
self.c_floor = query_min_clock(a.smi, self.g)
self.ceiling = self.top # where the compass holds the ceiling (continuous)
self.written_ceiling = self.top # what the card was last told
self.busy_clk, self.busy_draw = [], [] # the card on its own: busy clock and busy draw, ceiling at the top
self.prof = dict(LAW)
if a.down_gain is not None:
    self.prof["down"] = a.down_gain
self.verdict = Verdict(tolerance=a.allow, min_samples=a.verdict_samples, probe_every=a.verdict_every,
                       recheck=a.verdict_recheck)
self.svc_t = None # elapsed_seconds of the last request whose service time the verdict has se
↪ en
self.brain = CompassLaw(Band(0.0, 1.0, center=self.prof["center"]), dt=a.interval, tau=2.0 * a.interval, kp=1.
↪ 0, smooth=0.3)
self.brain.kd *= 3.0
self.audit({"snapshot": {str(self.g): {"limit_w": self.start, "min_limit_w": s[self.g]["min"],
                                     "enforced_w": s[self.g]["enforced"] if self.enforced_ok else None,
                                     "clock_top_mhz": self.top,
                                     **{f: snap[f][str(self.g)] for f in SNAPSHOT}}},
          "engine": "compass, two wires", "mode": a.mode, "law": self.prof,
          "verdict": {"allow": a.allow, "samples": a.verdict_samples, "every": a.verdict_every,
                     "recheck": a.verdict_recheck},
          "covers": {"clock_mhz": [round(self.c_lo), round(self.top)], "power_w": [self.floor_w, self.start]
↪ }}})
if a.mode == "cap" and snap["power.management"][str(self.g)].lower() != "enabled":
    self.audit({"refused": "power management not Enabled: a written limit would not bind"})
    raise SystemExit("power management not Enabled: cap mode refused, no authority taken")

def audit(self, rec):
    self.log.write(json.dumps({"time": time.time(), **rec}) + "\n"); self.log.flush()

def position(self, util):
    a = self.a
    ls = latency_sense(a.latency_file, a.latency_window_s) if a.latency_file and a.slo_ms else None
    if ls is not None and ls["blind"]:
        return None, ls
    resp = 0.0
    if ls is not None:
        # the position is the mean response time of the window between the bare service time and the line (what
        # the model reads); the 95th percentile at or past the line, or any failed request, is past the wall
        bare = a.slo_ms / 10.0
        w = latency_window(a.latency_file, a.latency_window_s)["ms"]
        mean = sum(w) / len(w) if w else ls["p95"]
        resp = (mean - bare) / (a.slo_ms - bare)
        if ls["fail"] or ls["p95"] >= a.slo_ms:
            resp = max(resp, 1.0)
    return resp, ls

def service_costs(self):
    """The card's own time (ms) on each request finished since the last decision (the workload's service_ms)."""
    import csv
    try:
        rows = list(csv.DictReader(open(self.a.latency_file)))
    except (OSError, TypeError):
        return []

```

```

out, last = [], self.svc_t
for row in rows:
    t, svc = row.get("elapsed_seconds"), row.get("service_ms")
    if row.get("ok") != "1" or not t or not svc:
        continue
    t = float(t)
    if last is None or t > last:
        out.append(float(svc))
        self.svc_t = t if self.svc_t is None else max(self.svc_t, t)
if last is None:
    return [] # the first read only finds where the file stands
return out

def learn(self, r):
    """The card on its own: while the ceiling is at the top and the card is busy, its clock and draw are native's.
    """
    if r is not None and self.written_ceiling >= self.top and r["limit"] >= self.start - 1.0 and r["util"] >= 0.9:
        self.busy_clk = (self.busy_clk + [r["clock_mhz"]])[-200:]
        self.busy_draw = (self.busy_draw + [r["draw"]])[-200:]

def native(self):
    """(busy clock, busy draw) of the card on its own, or (None, None) until enough busy readings are in."""
    if len(self.busy_clk) < self.a.learn_samples:
        return None, None
    c, d = sorted(self.busy_clk), sorted(self.busy_draw)
    return c[len(c) // 2], d[int(0.9 * (len(d) - 1))]

def write_clock(self, mhz, why):
    mhz = int(round(mhz))
    cmd = ["-i", str(self.g), "-lgc", f"{self.c_floor},{mhz}"]
    if self.a.mode != "cap" or self.foreign:
        self.audit({"would_clock_write": cmd, "why": why}); return
    rc, err = smi_run(self.a.smi, cmd)
    self.audit({"clock_write": cmd, "why": why, "rc": rc, "stderr": err})
    if rc != 0:
        raise WriteFailed(f"nvidia-smi -lgc {self.c_floor},{mhz} returned {rc}: {err}")
    self.clock_set = True
    self.writes += 1

def write_limit(self, w, why):
    w = int(w)
    cmd = ["nvidia-smi", "-i", str(self.g), "-pl", str(w)]
    if self.a.mode != "cap" or self.foreign:
        self.audit({"would_write": cmd, "why": why}); return
    self.audit({"write": cmd, "why": why})
    t0 = time.time()
    rc, err = smi_run(self.a.smi, ["-i", str(self.g), "-pl", str(w)])
    act = {"gpu": self.g, "requested_w": w, "t_requested": t0, "rc": rc, "stderr": err}
    if rc != 0:
        self.audit({"actuator": act, "write_failed": cmd})
        raise WriteFailed(f"nvidia-smi -pl {w} returned {rc}: {err}")
    self.writes += 1
    s = query(self.a.smi, [self.g], self.enforced_ok)
    if s is not None:
        back = s[self.g]
        act.update({"readback_w": back["limit"], "enforced_w": back["enforced"], "t_readback": time.time(),

```

```

        "realized": abs(back["limit"] - w) &lt; 1.0, "override": back["enforced"] &lt; back["limit"] - 1.0})
    if act["realized"]:
        act["delay_s"] = round(act["t_readback"] - t0, 3)
        self.expect = w
    self.audit({"actuator": act})

def step(self):
    a, g = self.a, self.g
    s = query(a.smi, [g], self.enforced_ok)
    if s is None:
        p, ls, r = None, None, None
    else:
        r = s[g]
        if abs(r["limit"] - self.expect) &gt;= 1.0 and not self.foreign:
            self.foreign = True
            self.audit({"foreign_writer": {"gpu": g, "reads_w": r["limit"], "expected_w": self.expect},
                        "action": "observe only from now on; both wires left to the other writer"})
        p, ls = self.position(r["util"])
    self.learn(r)
    n_clk, n_draw = self.native()
    thr = throttle(a.smi, [g]) if r is not None else None
    heat = slowed((thr or {}).get(g))
    saturated = r is not None and r["util"] &gt;= a.race_util
    # the verdict: the card's own time on the requests since the last decision, at the step the ceiling stood at;
    # then the deepest step the compass may use now, or the step a trial needs
    self.verdict.observe(self.service_costs())
    calm = p is not None and p &lt; self.brain.band.wall_high and not saturated
    deepest, trial, ev = self.verdict.tick(calm)

    if ev:
        self.audit({"verdict": ev, "state": self.verdict.state, "deepest_step": self.verdict.allowed})
    at_limit = r is not None and r["draw"] &gt;= 0.97 * r["limit"]
    if p is not None and saturated and at_limit and n_clk is not None and self.capped:
        # saturated against the card's own limit: hold the ceiling at the card's own busy clock under that limit,
        # no knock-backs; the lid stays at the start limit
        self.brain.force(p)
        ceiling = clamp(round(n_clk / a.min_change_mhz) * a.min_change_mhz, self.c_lo, self.top)
        lid = self.start
        who = "steady_under_limit"
    elif p is None or p &gt;= self.brain.band.wall_high or saturated:
        # fail up past the wall or blind; and race while work waits (the card saturated: a queue is forming), so a
        # burst is always served at full speed and the compass paces only the slack between bursts
        if p is not None:
            self.brain.force(p)
            ceiling, lid = self.top, self.start
            who = "blind_fail_up" if p is None else "fail_up" if p &gt;= self.brain.band.wall_high else "race"
        elif trial:
            self.brain.force(p)
            ceiling = self.top - deepest * a.min_change_mhz
            lid = self.start
            who = "verdict_trial"
        else:
            F = self.brain.force(p)
            # the speed floor: never under the clock the card reaches on its own while busy (until that is learned,
            # the top), so work waiting on the card is never served slower than native
            c_floor = self.top if n_clk is None else max(self.c_lo, min(self.top, n_clk * self.prof["floor"]))
            # and never past the deepest step the verdict has allowed (none allowed: the ceiling stays at the top)

```

```

c_floor = max(c_floor, self.top - deepest * a.min_change_mhz)
# the ceiling moves in whole clock steps, as the card's own clock does: the force times the gain (a share
# of the top clock per unit of force) rounded to whole --min-change-mhz steps; a pull under half a step
# moves nothing and is not stored up, so a calm card is paced only when the pull is clearly down
down = self.prof["down"]
step = (a.up_gain if F > 0 else down) * F * self.top
step = round(step / a.min_change_mhz) * a.min_change_mhz
ceiling = clamp(self.ceiling + step, c_floor, self.top)
if heat and ceiling < self.ceiling:
    ceiling, who = self.ceiling, "thermal_hold"
else:
    who = "compass"
# the lid: the start limit scaled to what a fully busy card draws at this ceiling (a fifth of the draw doe
↳ s not
# scale with the clock; the rest goes as clock^2.5, clock times voltage squared), plus headroom; at the to
↳ p
# clock the lid is the start limit, so the lid never adds a hammer of its own
# the lid: never under what the card draws on its own while busy (its own meter, not a curve), plus
# headroom; the start limit until that is learned
lid = self.start if n_draw is None else clamp(math.ceil(n_draw * (1.0 + a.lid_headroom)), self.floor_w, se
↳ lf.start)
rec = {"decision": {str(g): {"telemetry": None if r is None else {
    "util": r["util"], "draw_w": r["draw"], "temp_c": r["temp"], "limit_w": r["limit"],
    "enforced_w": r["enforced"], "clock_mhz": r["clock_mhz"], "throttle": (thr or {}).get(g)},
    "position": None if p is None else round(p, 4), "velocity": round(self.brain.v, 4),
    "latency_p95_ms": None if not ls else ls.get("p95"),
    "ceiling_mhz": round(ceiling), "want_w": int(lid), "decided_by": who,
    "native_busy_clock_mhz": n_clk, "native_busy_draw_w": n_draw,

    "verdict_state": self.verdict.state, "verdict_deepest_step": self.verdict.allowed}}}
self.audit(rec)
self.ceiling = ceiling
if abs(ceiling - self.written_ceiling) >= a.min_change_mhz or (ceiling >= self.top and self.written_ceiling <
↳ self.top):
    if ceiling >= self.top and self.clock_set and a.mode == "cap" and not self.foreign:
        rc, err = smi_run(a.smi, ["-i", str(g), "-rgc"])
        self.audit({"clock_write": ["-i", str(g), "-rgc"], "why": who, "rc": rc, "stderr": err})
        if rc != 0:
            raise WriteFailed(f"nvidia-smi -rgc returned {rc}: {err}")
        self.clock_set = False
    elif ceiling < self.top:
        self.write_clock(ceiling, who)
        self.written_ceiling = ceiling
    cur = r["limit"] if r is not None else self.expect
    if abs(lid - cur) >= a.min_change_w or (lid == self.start and cur != lid):
        self.write_limit(lid, who)

def restore(self):
    failed = []
    if self.a.mode == "cap" and not self.foreign:
        rc, err = smi_run(self.a.smi, ["-i", str(self.g), "-rgc"])
        self.audit({"clock_write": ["-i", str(self.g), "-rgc"], "why": "restore", "rc": rc, "stderr": err})
        if rc != 0:
            failed.append(f"-rgc returned {rc}: {err}")
        s = query(self.a.smi, [self.g], self.enforced_ok)
        if s is None or abs(s[self.g]["limit"] - self.start) >= 1.0:
            try:

```

```

        self.write_limit(self.start, "reset: the limit read at start")
    except WriteFailed as e:
        failed.append(str(e))
    s = query(self.a.smi, [self.g], self.enforced_ok) or {}
    back = s[self.g]["limit"] if self.g in s else None
    ok = self.foreign or (not failed and back is not None and abs(back - self.start) < 1.0)
    self.audit({"restored": {str(self.g): back}, "start": {str(self.g): self.start}, "ok": ok,
               "writes": self.writes, "foreign_writer": self.foreign, "restore_write_failed": failed})
    return ok

def parser():
    ap = argparse.ArgumentParser(description=__doc__, formatter_class=argparse.RawDescriptionHelpFormatter)
    ap.add_argument("--mode", choices=["watch", "cap"], default="watch")
    ap.add_argument("--gpus", default="0")
    ap.add_argument("--smi", default=os.environ.get("NVIDIA_SMI", "nvidia-smi"))
    ap.add_argument("--interval", type=float, default=2.0)
    ap.add_argument("--duration", type=float, default=0.0)
    ap.add_argument("--audit", default="gpu_compass_audit.jsonl")
    ap.add_argument("--kill-file", default="/tmp/omni-gpu-kill")
    ap.add_argument("--latency-file", default="")
    ap.add_argument("--slo-ms", type=float, default=0.0)
    ap.add_argument("--latency-window-s", type=float, default=5.0, help="seconds of response times the position reads")
    ap.add_argument("--race-util", type=float, default=0.95, help="utilization at which work is waiting: race, never p
    ace")
    ap.add_argument("--learn-samples", type=int, default=15, help="busy readings of the card on its own before the com
    pass may lower anything")

    ap.add_argument("--floor-w", type=float, default=0.0, help="the declared envelope's lowest watts")
    ap.add_argument("--clock-min-share", type=float, default=0.35, help="the clock ceiling's cover: lowest share of th
    e top")
    ap.add_argument("--allow", type=float, default=0.02, help="the most a ceiling step may add to the card's own time
    on a request")
    ap.add_argument("--verdict-samples", type=int, default=30, help="requests measured at native and at the trial step")
    ap.add_argument("--verdict-every", type=int, default=8, help="decisions between trials")
    ap.add_argument("--verdict-recheck", type=int, default=900, help="decisions before a refused step is tried again")
    ap.add_argument("--up-gain", type=float, default=0.10, help="share of the top clock moved per unit of force, up")
    ap.add_argument("--down-gain", type=float, default=None, help="share of the top clock moved per unit of force, dow
    n (default: the law's)")
    ap.add_argument("--lid-headroom", type=float, default=0.10, help="the lid above the draw the ceiling takes")
    ap.add_argument("--min-change-mhz", type=float, default=15.0)
    ap.add_argument("--min-change-w", type=float, default=3.0)
    return ap

def main(argv=None):
    a = parser().parse_args(argv)
    master.refuse_if_off("GPU governor (two wires)")
    gov = GpuCompass(a)
    # what puts the card back if this process dies without doing it itself (the watchdog runs it)
    back = ([a.smi, "-i", str(gov.g), "-rgc"], [a.smi, "-i", str(gov.g), "-pl", str(int(round(gov.start))]))
    if a.mode == "cap" else []
    master.register("GPU governor (two wires)", restore=back, stale_s=max(60.0, 5 * a.interval))
    stop = {"now": False}
    signal.signal(signal.SIGTERM, lambda *_: stop.__setitem__("now", True))

    t0 = time.time()
    failed = False
    try:
        while not stop["now"] and not os.path.exists(a.kill_file) and not master.is_off() and (a.duration <= 0 or time
        .time() - t0 < a.duration):
            try:
                master.heartbeat(); gov.step()
            except WriteFailed as e:
                gov.audit({"fatal": f"write failed: {e}"}); failed = True
                break
            except Exception as e: # noqa: BLE001 a failed decision is logged; the kill path still runs
                gov.audit({"error": f"{type(e).__name__}: {e}"})
            end = time.time() + a.interval
            while time.time() < end and not stop["now"] and not os.path.exists(a.kill_file) and not master.is_off():
                time.sleep(0.2)
    finally:
        ok = gov.restore()
        return 5 if gov.foreign else 3 if not ok else 4 if failed else 0

if __name__ == "__main__":
    sys.exit(main())

```

omnicompass/adapter.py

309 lines.

```

"""Omni-Compass stack governor.

    sense      : stack telemetry &gt; normalized observation vector
    assimilate : observation blended into the six-state engine state
    evolve     : one macro step of equations (1)-(7) with u = 0
    allocate   : engine state &gt; directive (capacity, change gating, power, rollback, routing, alerting, paging)
    authority   : OBSERVE | AUTOPILOT; kill() revokes authority and returns control to the native managers
"""
from __future__ import annotations

import math
from collections import deque
from dataclasses import dataclass, asdict, field
from typing import Any, Dict, List, Optional

from .core import State, Params, macro_step, v_eff, control_command, U_AUTHORITY

OBSERVE = "observe"
AUTOPILOT = "autopilot"

# Engine parameters used for stack governance (fixed, not sampled).
STACK_PARAMS = Params(alpha=4.2, beta_int=0.0, beta_ext=0.0, k=1.7, sigma_1=0.38, delta=0.705, gamma_1=1.35,
                      gamma_c=1.0, lambda_0=-0.40, lambda_1=1.55, lambda_2=1.15, c=1.0, E_max=2.0,
                      omega_B=1.0, Q_B=2.5, alpha_s=0.12, beta_s=0.10, mu=3.782, alpha_E=3.487, alpha_U=4.2,
                      lambda_I=0.719, lambda_U=0.0)
INITIAL_STATE = State(E=0.12, U=0.82, I_U=0.0, S=0.0, B=0.0, B_dot=0.0)
ASSIMILATION = 0.339

def clamp(v: float, lo: float, hi: float) -> float:
    return lo if v < lo else hi if v > hi else v

def observe_vector(obs: Dict[str, float], conflicts: int) -> Dict[str, float]:
    """Normalize telemetry. Inputs are the stack's delayed observations."""
    return {
        "q": clamp(float(obs.get("queue_ratio", 0.0)), 0.0, 2.0),
        "load": clamp(float(obs.get("load_ratio", 0.0)), 0.0, 2.0),
        "power": clamp(float(obs.get("power_stress", 0.0)), 0.0, 1.5),
        "thermal": clamp(float(obs.get("thermal", 0.0)), 0.0, 1.5),
        "network": clamp(float(obs.get("network_stress", 0.0)), 0.0, 1.5),
        "drift": clamp(float(obs.get("drift_ratio", 0.0)), 0.0, 1.5),
        "stale": clamp(float(obs.get("stale", 0.0)), 0.0, 1.0),
        "security": clamp(float(obs.get("security_block", 0.0)), 0.0, 1.0),
        "conflict": clamp(conflicts / 4.0, 0.0, 1.0),
    }

def assimilate(x: State, o: Dict[str, float], p: Params) -> State:
    """Blend the observation into the engine state and set the forcing terms."""
    q, load, pw, th, nw, dr, st, sec, cf = (o[k] for k in
        ("q", "load", "power", "thermal", "network", "drift", "stale", "security", "conflict"))
    e_obs = clamp(0.25 * q + 0.18 * max(0.0, load - 0.85) + 0.16 * pw + 0.13 * th + 0.10 * nw + 0.08 * dr
        + 0.06 * cf + 0.04 * st, 0.0, 1.5)
    u_obs = clamp(1.0 - (0.27 * q + 0.18 * pw + 0.16 * th + 0.12 * nw + 0.12 * dr + 0.10 * cf + 0.05 * st), 0.0, 1.0)
    s_obs = clamp(0.36 * th + 0.28 * pw + 0.18 * nw + 0.10 * sec + 0.08 * st - 0.20 * q, -1.0, 1.0)

```



```

b_obs = clamp(0.52 * pw + 0.26 * th + 0.22 * dr, -1.0, 1.5)
i_obs = clamp(q + dr + cf, 0.0, 2.0)
a = ASSIMILATION
p.beta_int = clamp(0.58 * q + 0.42 * max(0.0, load - 0.75), 0.0, 1.0)
p.beta_ext = clamp(0.33 * pw + 0.27 * th + 0.20 * nw + 0.12 * st + 0.08 * sec, 0.0, 1.0)
return State(E=(1 - a) * x.E + a * e_obs, U=(1 - a) * x.U + a * u_obs, I_U=(1 - a) * x.I_U + a * i_obs,
             S=(1 - a) * x.S + a * s_obs, B=(1 - a) * x.B + a * b_obs, B_dot=(1 - a) * x.B_dot + a * (b_obs - x.B)
             ~ )

@dataclass(frozen=True)
class AllocationLaw:
    """Engine state -> governing directive.

    release    capacity removed only while push &lt;= push_release (the engine's controller reports convergence)
    add         without backlog, capacity added only while push &gt;= push_add (the engine's controller reports need)
    push        equation (2) control law evaluated on the evolved state with target sigma = +1:
                u = clip(-f_U(x) + KP (1 - U), -U_AUTHORITY, U_AUTHORITY); push = u / U_AUTHORITY
                (f_U contains the double-well drift and -(dE/dt)/E_max from equation (1))
    headroom    rho* = clamp(rho0 - kI*I_U - kE*E, rho_min, rho0)
    sizing      load_s = load * cap_now if size_at_full_cap else load
    capacity    n_req = ceil(n * load / rho*) + ceil(kq * q * n); delta = clamp(n_req - n, -down_max, up_max)
    guard       no capacity removal while q &gt;= guard_queue
    band        capacity removed only when n_req &lt;= n - down_band
    dwell       capacity removed only after surplus persists for down_dwell intervals and no earlier than
                down_after_add intervals after the last addition
    envelope    no capacity addition while observed power &gt;= power_ceiling (power_ceiling_backlog while q &gt;= guard_queue
    ~ e)

    security    no capacity addition while a security block is observed
    change      Terraform plans only while U &gt;= U_gate, only in the direction of the live capacity decision
    power       q &lt;= guard_queue and U &gt;= U_gate: cap = clamp(cap_now * load * (1 + margin), cap_min, 1) (trim to dema
    ~ nd)

                q &gt;= guard_queue at the power envelope: cap = max(cap_min, cap_now - spread_step) (spread)
                q &gt;= guard_queue below the envelope: cap held (capacity added as nodes; 1 at max_nodes)
                otherwise cap = 1; always cap &lt;= 1 - cap_gain * max(0, B - B_cap)
    rollback    a proposed rollback is authorised if S &gt;= S_rollback, U is falling, or security blocks change
    replicas     direct actuation: desired = ceil(replicas * load / rho*) (the HPA formula with the engine-set target),
                increases unless capacity is being reduced, decreases only while reducing;
                otherwise autoscaler increases approved unless reducing, decreases only while reducing
    consistency a step that expands capacity never tightens the power cap
    paging      none: the governor holds every action a paged operator would take
    """
    rho0: float = 0.914
    rho_min: float = 0.627
    kI: float = 0.252
    kE: float = 0.297
    kq: float = 0.795
    up_max: int = 2
    down_max: int = 1
    guard_queue: float = 0.05
    power_ceiling: float = 0.95
    power_ceiling_backlog: float = 1.00
    spread_step: float = 0.05
    backlog_release: bool = False
    size_at_full_cap: bool = False
    push_release: float = 0.005

```

```

push_add: float = -9.0
U_gate: float = 0.481
margin: float = 0.006
cap_min: float = 0.706
B_cap: float = 0.40
cap_gain: float = 0.60
S_rollback: float = 0.55
route_network: float = 0.20
down_band: int = 4
down_dwell: int = 4
down_after_add: int = 4
min_nodes: int = 8
max_nodes: int = 40

@dataclass
class Governor:
    law: AllocationLaw = field(default_factory=AllocationLaw)
    mode: str = AUTOPILOT
    killed: bool = False
    x: State = field(default_factory=lambda: State(**asdict(INITIAL_STATE)))
    p: Params = field(default_factory=lambda: Params(**asdict(STACK_PARAMS)))
    t: float = 0.0
    current_cap: float = 1.0
    nodes: int = 20
    surplus_streak: int = 0
    since_add: int = 99
    last_push: float = 0.0

    log: deque = field(default_factory=lambda: deque(maxlen=4096))
    evolve: bool = True

    def kill(self) -> None:
        """Revoke authority; the engine continues to compute and log directives."""
        self.killed = True

    def set_mode(self, mode: str) -> None:
        if mode not in (OBSERVE, AUTOPILOT):
            raise ValueError(mode)
        self.mode = mode

    @property
    def has_authority(self) -> bool:
        return self.mode == AUTOPILOT and not self.killed

    def step(self, obs: Dict[str, float], conflicts: int) -> Dict[str, Any]:
        """Sense, assimilate, evolve, allocate."""
        o = observe_vector(obs, conflicts)
        U_before = self.x.U
        xa = assimilate(self.x, o, self.p)
        if self.evolve:
            xn, _ = macro_step(xa, self.p, self.t, target=None)
        else:
            xn = xa
        self.x = xn
        self.t += 0.1
        L = self.law

```

```

n = int(self.nodes)
u_push, _ = control_command(xn, self.p, self.t, +1)
push = u_push / U_AUTHORITY
self.last_push = push
rho = clamp(L.rho0 - L.kI * xn.I_U - L.kE * xn.E, L.rho_min, L.rho0)
converged = push &lt;= L.push_release
load_sized = o["load"] * self.current_cap if L.size_at_full_cap else o["load"]
n_req = math.ceil(n * load_sized / rho) + math.ceil(L.kq * o["q"] * n)
delta = max(-L.down_max, min(L.up_max, n_req - n))
if delta &lt; 0 and not converged:
    delta = 0
if delta &gt; 0 and o["q"] &lt; L.guard_queue and push &lt; L.push_add:
    delta = 0
if delta &lt; 0 and (o["q"] &gt;= L.guard_queue or n_req &gt; n - L.down_band):
    delta = 0
self.surplus_streak = self.surplus_streak + 1 if delta &lt; 0 else 0
if delta &lt; 0 and (self.surplus_streak &lt; L.down_dwell or self.since_add &lt; L.down_after_add):
    delta = 0
ceiling = L.power_ceiling_backlog if o["q"] &gt;= L.guard_queue else L.power_ceiling
if delta &gt; 0 and (o["power"] &gt;= ceiling or n &gt;= L.max_nodes):
    delta = 0
if o["security"] &gt; 0.5:
    delta = min(delta, 0)
self.since_add = 0 if delta &gt; 0 else self.since_add + 1
if delta &lt; 0:
    self.surplus_streak = 0
at_envelope = o["power"] &gt;= ceiling
if o["q"] &lt; L.guard_queue and xn.U &gt;= L.U_gate:

    cap = clamp(self.current_cap * o["load"] * (1.0 + L.margin), L.cap_min, 1.0)
elif o["q"] &gt;= L.guard_queue and L.spread_step &gt; 0.0:
    if at_envelope:
        cap = clamp(self.current_cap - L.spread_step, L.cap_min, 1.0)
    else:
        cap = 1.0 if L.backlog_release else (self.current_cap if n &lt; L.max_nodes else 1.0)
elif o["q"] &gt;= L.guard_queue and L.backlog_release:
    cap = 1.0
else:
    cap = 1.0
cap = min(cap, clamp(1.0 - L.cap_gain * max(0.0, xn.B - L.B_cap), L.cap_min, 1.0))
if delta &gt; 0:
    cap = max(cap, self.current_cap)
command = clamp(0.5 * (v_eff(xn, self.p) / max(self.p.c, 1e-9) + 1.0), 0.0, 1.0)
d = {
    "node_delta": int(delta),
    "change_permitted": xn.U &gt;= L.U_gate and o["security"] &lt;= 0.5,
    "power_cap": cap,
    "rollback_authorized": xn.S &gt;= L.S_rollback or xn.U &lt; U_before or o["security"] &gt; 0.5,
    "route_shift": o["network"] &gt; L.route_network,
    "state": {"E": xn.E, "U": xn.U, "I_U": xn.I_U, "S": xn.S, "B": xn.B},
    "demand": rho,
    "load": o["load"],
    "command": command,
}
self.log.append({"mode": self.mode, "killed": self.killed, **{k: v for k, v in d.items() if k != "state"},
                **d["state"]})
return d

```

```

def directive_to_actions(d: Dict[str, Any], stack_state: Dict[str, Any], proposals: List[Dict[str, Any]],
                        cfg, direct_actuation: bool = False) -> List[Dict[str, Any]]:
    """Directive -> executable stack actions. direct_actuation: capacity is owned by the governor; no Terraform plans.
    ~ """
    tag = "Omni-Compass"
    out: List[Dict[str, Any]] = []
    cur = int(stack_state["actual_nodes"])
    target = int(max(cfg.minimum_nodes, min(cfg.maximum_nodes, cur + int(d["node_delta"]))))
    if target != cur:
        out.append({"manager": tag, "action": "nodes", "target": target, "direction": 1 if target > cur else -1})
    if not direct_actuation and d["change_permitted"] and any(p["action"] == "terraform_plan" for p in proposals):
        tf_dir = 1 if target > stack_state["terraform_desired_nodes"] else -1
        live_dir = 1 if target > cur else -1 if target < cur else 0
        if abs(target - stack_state["terraform_desired_nodes"]) >= 2 and live_dir in (0, tf_dir):
            out.append({"manager": tag, "action": "terraform_plan", "target": target,
                        "direction": 1 if target > stack_state["terraform_desired_nodes"] else -1})
    reps = [int(p["target"]) for p in proposals if p["action"] == "replicas"]
    if direct_actuation:
        cur_r = int(stack_state["replicas"])
        rt = max(20, min(500, math.ceil(cur_r * float(d["load"]) / max(float(d["demand"]), 1e-9))))
        if (rt > cur_r and d["node_delta"] >= 0) or (rt < cur_r and d["node_delta"] < 0):
            out.append({"manager": tag, "action": "replicas", "target": rt, "direction": 1 if rt > cur_r else -1})
        reps = []
    if reps:
        cur_r = int(stack_state["replicas"])
        rt = None

        if d["node_delta"] >= 0 and max(reps) > cur_r:
            rt = max(reps)
        elif d["node_delta"] < 0 and min(reps) < cur_r:
            rt = min(reps)
        if rt is not None:
            out.append({"manager": tag, "action": "replicas", "target": rt,
                        "direction": 1 if rt > cur_r else -1})
    cap = float(d["power_cap"])
    if any(a["direction"] > 0 and a["action"] in ("nodes", "terraform_plan", "replicas") for a in out):
        cap = max(cap, float(stack_state["power_cap"]))
    if abs(cap - stack_state["power_cap"]) > 1e-9:
        out.append({"manager": tag, "action": "power_cap", "target": cap,
                    "direction": -1 if cap < stack_state["power_cap"] else 1})
    if d["route_shift"]:
        out.append({"manager": tag, "action": "route_shift", "target": 1.0, "direction": 1})
    if any(p["action"] == "rollback" for p in proposals) and d["rollback_authorized"]:
        out.append({"manager": tag, "action": "rollback", "target": 1.0, "direction": -1})
    alerts = [p for p in proposals if p["action"] == "alert"]
    if alerts:
        out.append({"manager": tag, "action": "alert", "target": max(int(p["target"]) for p in alerts), "direction": 1
                    ~ })
    return out

THROUGHPUT_MODE = dict(backlog_release=True, power_ceiling=9.0, power_ceiling_backlog=9.0, spread_step=0.0, B_cap=9.0,
                        size_at_full_cap=True, kI=0.05, kE=0.40, cap_min=0.65, down_band=1, down_dwell=10,
                        down_after_add=3, kq=0.30, U_gate=0.45, push_release=0.01)

```

```

FLEET_MODE = dict(THROUGHPUT_MODE, rho=0.829, rho_min=0.769, kI=0.072, kE=0.487, kq=0.635, down_band=3, down_dwell=1,
    ↳ down_after_add=5, push_release=0.053, cap_min=0.695, margin=0.131, U_gate=0.85, up_max=8)
FLEET_BALANCED_MODE = dict(FLEET_MODE, push_release=0.0098, push_add=0.0098)
FLEET_WEAR_MODE = dict(FLEET_MODE, push_release=0.0035, push_add=0.0035)
FLEET_DECISION_SECONDS = 60

def mode_law(mode: str, base: "AllocationLaw" = None) -> "AllocationLaw":
    """Operating modes.

    power_protect the site power limit is enforced (envelope, spread-and-throttle, bath cap; shield I4 active)
    ↳ fleet        throughput mode with constants selected on the 15-second fleet harness (fleet/), decisions every 60
    ↳ s;
    ↳ energy-first: the equation (2) release gate is lenient and the add gate is off
    fleet_balanced fleet with the equation (2) gate active in both directions at 0.0098 (release and add)
    fleet_wear      fleet with the equation (2) gate active in both directions at 0.0035 (release and add)
    throughput      the power envelope is not enforced (same power rules as the native stack); caps released on backlog
    ↳ ;
    ↳ capacity sized at full power (load * cap_now), so power trimming does not trigger machine starts;
    ↳ constants selected by constrained search on development seeds (minimum energy subject to backlog,
    ↳ health, recovery, reversal and invariant constraints)

    """
    from dataclasses import replace
    base = base or AllocationLaw()
    if mode == "power_protect":
        return base
    if mode == "throughput":
        return replace(base, **THROUGHPUT_MODE)
    if mode == "fleet":
        return replace(base, **FLEET_MODE)
    if mode == "fleet_balanced":
        return replace(base, **FLEET_BALANCED_MODE)
    if mode == "fleet_wear":
        return replace(base, **FLEET_WEAR_MODE)
    raise ValueError(mode)

```

omnicompass/nervous_system.py

158 lines.

```

"""The supervisory nervous system: one evolving engine state grants authority to every organ family.

Omni-Compass does not replace the specialist controllers. Each keeps its own mechanism:
- the HPA and VPA size pods;
- the Cluster Autoscaler and Karpenter size node pools;
- Linux schedutil picks CPU frequency;
- NVIDIA DCGM manages GPUs;
- rollout controllers and routers do their own work.

The nervous system decides, from one shared state, the envelope inside which each of them may act.

Inputs, all from the six-state engine (omnicompass/core.py, via the governor), and nothing tuned per organ:
kappa convergence = 1 if push <= push_release, else push_release / push. This is the equation-(2) gate made
                    continuous: the further the engine is from its basin, the less it may give back.
h                basin health = clamp((U - U_gate) / (1 - U_gate), 0, 1)
sigma            stress ratio = S / S*, where S* solves delta - alpha_s S - (3/4) beta_s S^2 = 0 (equation 6)
nu              unmet need    = max(0, I - U) (equation 3)
calm            = kappa * h * (1 - clamp(sigma - 1, 0, 1)) * (1 - clamp(nu, 0, 1))
                  in [0, 1]; 1 means fully settled

Authority per organ:
expand          may the organ add capacity or performance. Always yes, except during a security hold.
contract        may it give capacity or performance back. Only when calm >= the organ's reversibility threshold
                  (theta_pods 0.5, theta_nodes 0.7: a node costs a boot to reverse) and the SLO is clean.
step            fraction of the organ's surplus it may shed this decision. Equals calm.
envelope        for continuous organs, the allowed range:
                  cpufreq ceiling (fraction of cpuinfo_max) in [1 - 0.35 calm, 1];
                  GPU power limit (fraction of max) in [1 - 0.30 calm, 1];

```

```

        site power cap in [0.65, 1];
        traffic-shift fraction in [0, 0.5 calm];
        cooling supply-air setpoint (C) in [18, 18 + 9 calm].
Heat above 0.96 forces the frequency and GPU ceilings down to 1 - 0.35 x excess.
batch    admit held work only if calm >= 0.5 and power stress < 0.9; pause pausable work if sigma > 1, power
        stress >= 0.95, or heat >= 0.96.
rollback the engine's rollback authorisation (S high, U falling, or a security hold).

Global rules, applied last:
- observe: every authority is computed and logged; execute is False everywhere.
- kill: no authority at all.
- security hold: no capacity organ may expand. Cooling is the one protective organ: "expand" there means more
  cooling, which adds protection and no capacity, so it stays allowed.
The safety shield (omnicompass/shield.py) stays downstream and can still veto any action.
"""
from __future__ import annotations

from dataclasses import dataclass
from typing import Any, Dict

from omnicompass.closure import stress_equilibrium

ORGANS = ("pods", "nodes", "cpufreq", "gpu", "power", "batch", "routing", "rollback")
# The living band (the founder's rule): every level the nervous system hands out lives between 5% and 95% of its range
# .
# Nothing is driven to zero (a part with no work idles down to its floor, alive and ready) and nothing is driven to it
# S
# absolute top (the last 5% is never spent). The band is applied last, to every envelope, so no organ can leave it.

BAND = (0.05, 0.95)

def in_band(lo: float, hi: float) -> list:
    """Clip an envelope [lo, hi] (fractions of the organ's range) into the living band, keeping lo <= hi."""
    b0, b1 = BAND
    h = min(b1, max(b0, hi))
    return [min(h, max(b0, lo)), h]

THETA = {"pods": 0.5, "nodes": 0.7, "cpufreq": 0.3, "gpu": 0.3, "power": 0.3, "routing": 0.5}

def clamp(x, lo, hi):
    return lo if x < lo else hi if x > hi else x

@dataclass(frozen=True)
class NervousInputs:
    E: float; U: float; I_U: float; S: float
    push: float
    push_release: float = 0.2
    U_gate: float = 0.5
    s_eq: float = 1.0
    security_block: float = 0.0
    slo_clean: bool = True
    power_stress: float = 0.0
    thermal: float = 0.0

```

```

rollback: bool = False
stale: float = 0.0          # fraction of declared senses that are blind (dropout, frozen, unreadable)
mode: str = "autopilot"     # "observe" or "autopilot"
killed: bool = False

def scalars(i: NervousInputs) -> Dict[str, float]:
    kappa = 1.0 if i.push &lt;= i.push_release else i.push_release / max(i.push, 1e-12)
    h = clamp((i.U - i.U_gate) / max(1.0 - i.U_gate, 1e-12), 0.0, 1.0)
    sigma = i.S / i.s_eq if i.s_eq &gt; 0 else 0.0
    nu = max(0.0, i.I_U)
    calm = kappa * h * (1.0 - clamp(sigma - 1.0, 0.0, 1.0)) * (1.0 - clamp(nu, 0.0, 1.0))
    return {"kappa": kappa, "h": h, "sigma": sigma, "nu": nu, "calm": clamp(calm, 0.0, 1.0)}

def authority(i: NervousInputs) -> Dict[str, Any]:
    if i.killed:
        return {"execute": False, "killed": True, "organs": {}}
    sc = scalars(i); calm = sc["calm"]
    sec = i.security_block &gt; 0.5
    hot = i.thermal &gt;= 0.96
    excess = clamp(i.thermal - 0.96, 0.0, 1.0)
    org: Dict[str, Dict[str, Any]] = {}
    seeing = i.stale &lt;= 0.0          # never give capacity back on a blind sense; expanding stays allowed (the safe si
    ↪ de)
    for o, th in THETA.items():
        ok = calm &gt;= th and i.slo_clean and seeing
        org[o] = {"expand": not sec, "contract": ok, "step": calm if ok else 0.0}

    cf_lo = 1.0 - 0.35 * calm; gp_lo = 1.0 - 0.30 * calm
    cf_hi = 1.0 - 0.35 * excess if hot else 1.0; gp_hi = 1.0 - 0.35 * excess if hot else 1.0
    org["cpufreq"]["envelope"] = in_band(min(cf_lo, cf_hi), cf_hi)
    org["gpu"]["envelope"] = in_band(min(gp_lo, gp_hi), gp_hi)
    org["power"]["envelope"] = in_band(0.65, 1.0)
    # routing is an amount moved, not a level: it may move nothing, never more than the band's top
    org["routing"]["envelope"] = [0.0, min(BAND[1], 0.5 * calm) if not sec else 0.0]
    org["cooling"] = {"expand": True, "protective": True, "contract": calm &gt;= 0.3 and seeing, "step": calm, "envelope"
    ↪ : [18.0, 18.0 + 9.0 * calm]}
    org["batch"] = {"expand": (not sec) and seeing and calm &gt;= 0.5 and i.power_stress &lt; 0.9, "contract": True, "protec
    ↪ tive_contract": True,
        "admit": (not sec) and seeing and calm &gt;= 0.5 and i.power_stress &lt; 0.9,
        "pause": sc["sigma"] &gt; 1.0 or i.power_stress &gt;= 0.95 or hot, "step": calm}
    org["rollback"] = {"expand": False, "contract": False, "authorized": bool(i.rollback or sec), "step": 0.0}
    return {"execute": i.mode == "autopilot", "killed": False, "security_hold": sec, "scalars": sc, "organs": org}

def from_governor(g, obs: Dict[str, Any], d: Dict[str, Any] = None, mode: str = None) -> Dict[str, Any]:
    """Authority from a live omnicompass.adapter.Governor after its step (d = the step's directive)."""
    x = g.x; p = g.p; L = g.law
    return authority(NervousInputs(
        E=x.E, U=x.U, I_U=x.I_U, S=x.S, push=g.last_push, push_release=getattr(L, "push_release", 0.2),
        U_gate=getattr(L, "U_gate", 0.5), s_eq=stress_equilibrium(p.delta, p.alpha_s, p.beta_s),
        security_block=float(obs.get("security_block", 0.0)), slo_clean=bool(obs.get("slo_clean", True)),
        power_stress=float(obs.get("power_stress", 0.0)), thermal=float(obs.get("thermal", 0.0)),
        rollback=bool((d or {}).get("rollback_authorized", False)), stale=float(obs.get("stale", 0.0)),
        mode=mode or g.mode, killed=g.killed))

```

```

def node_release_gate(n: int, per_node_m: float, used_m: float, pending: int, pods_scaling_up: bool,
    latency_breach_now: bool, rho: float, node_auth: Dict[str, Any],
    senses_live: bool = True, last_command_landed: bool = True) -> Dict[str, Any]:
    """May the machine organ give one machine back now? Attribution: an organ is held back only by stress it can cause
    or cure. The node organ reads its own engine view (fed with machine-attributable pressure: pods waiting for a
    place), and the release must also pass:
        coordination pods are not scaling up and latency is not breached now (pods move first; machines never move
        against them, the rule of the benchmarked coordination)
        headroom      nothing is pending, and after the release the remaining machines run at or below the engine's own
        utilisation target rho: used / ((n - 1) x per_node) <= rho
        authority      the node organ's own calm, security and stress gates (authority() above) grant contraction
        senses          every declared sense is live (afferent integrity)
        proprioception the node organ's last command landed (efferent feedback: no new order to a muscle that did not
        carry out the previous one)
    Returns {"ok": bool, "reason": str, "util_after": float}."""
    util_after = used_m / max((n - 1) * per_node_m, 1e-9) if n > 1 else float("inf")
    checks = [("one machine left", n > 1), ("pods waiting", pending == 0), ("pods scaling up", not pods_scaling_up),
        ("latency breached now", not latency_breach_now), (f"util after {util_after:.2f} > rho {rho:.2f}", util_
        < after <= rho),
        ("node organ has no contraction authority", bool(node_auth.get("organs", {}).get("nodes", {})).get("contr
        < act", False))),
        ("a sense is blind", senses_live), ("last node command did not land", last_command_landed)]
    failed = [name for name, ok in checks if not ok]
    return {"ok": not failed, "reason": "; ".join(failed) or "release permitted", "util_after": util_after}

```


APPENDIX

Appendix I. The License

The license under which this book and the software are made available, as it stands in the file LICENSE at the root of the repository.

Omni-Compass Evaluation License, version 1.0
SPDX-License-Identifier: LicenseRef-OmniCompass-Evaluation-1.0

Copyright (c) 2026 The Omni-Compass LLC. All rights reserved.

This is NOT an open-source license. The software in this repository ("the Software") - including the Omni-Compass engine, the compass law, the governors, the harnesses, the controllers, the plug and wiring, the documentation, the benchmarks and their results, and the methods and mechanism they embody - is the proprietary property of The Omni-Compass LLC and is protected by copyright and by patents and patent applications held by The Omni-Compass LLC. The Software is made available publicly for EVALUATION AND SIMULATION PURPOSES ONLY.

INTELLECTUAL PROPERTY NOTICE. Patent applications, copyright registrations and trademark applications covering the Omni-Compass engine, its mathematics and its software have been filed in the United States by The Omni-Compass LLC. Every person who downloads, copies, forks, scans or runs the Software is on notice of these filings and of this license.

1. Permitted use (free of charge, no license fee). You may download, copy, read, run and modify the Software solely to
 - (a) evaluate it, including running its tests, benchmarks and simulations, and running it in shadow or test mode on systems you own or control for the purpose of evaluating it; and
 - (b) reproduce and check the results published in this repository.This permission is personal, non-exclusive, non-transferable, non-sublicensable and revocable.

2. Everything else requires a signed, paid license. Any use other than section 1 requires a written commercial license agreement signed by The Omni-Compass LLC and the licensee, and payment of the license fees it sets out (an "Omni-Compass Enterprise License"). This includes, without limitation:
 - (a) any commercial use, commercialization or monetization of the Software, its mechanism or its methods;
 - (b) use in production, or to operate, control or govern any system, fleet, cluster, facility, machine, vehicle or grid beyond evaluation;
 - (c) selling, licensing, sublicensing, leasing, renting, distributing or offering the Software, or any product or service that includes, is derived from or is operated by it, including as a hosted, managed or cloud service;
 - (d) incorporating the Software, or any part or derivative of it, into any product, service, platform or offering;
 - (e) using the Software, its results or its methods to build, train, benchmark or improve a competing product or service.
3. No other rights. No right, title, license or interest is granted except as expressly stated in section 1. In particular, no patent license is granted for any use other than section 1, and no trademark license is granted. "Omni-Compass" and its logos are trademarks of The Omni-Compass LLC.
4. Redistribution. You may not redistribute the Software or any modified version, in source or binary form, except as permitted by a signed Omni-Compass Enterprise License. Forks on a code-hosting service are permitted only for the purposes in section 1 and must keep this license and all notices unchanged.
5. Notices. You must keep this license, the copyright notice, the SPDX identifiers and all patent and trademark notices in every copy, including modified copies made under section 1.
6. Termination. Any use outside section 1 without a signed, paid Omni-Compass Enterprise License terminates the rights in section 1 automatically, and is infringement of The Omni-Compass LLC's copyright and patent rights.
7. No warranty; no liability. THE SOFTWARE IS PROVIDED "AS IS", WITHOUT WARRANTY OF ANY KIND, EXPRESS OR IMPLIED, INCLUDING MERCHANTABILITY, FITNESS FOR A PARTICULAR PURPOSE AND NON-INFRINGEMENT. IN NO EVENT SHALL THE OMNI-COMPASS LLC BE LIABLE FOR ANY CLAIM, DAMAGES OR OTHER LIABILITY ARISING FROM THE SOFTWARE OR ITS USE.
8. Third-party components. Files that carry a different license notice (for example `tests/third\_party/`) remain under their own terms.

Commercial licensing (Omni-Compass Enterprise License): The Omni-Compass LLC.

BACK MATTER

Contact

Licensing, pilots, evaluations and the Omni-Compass Enterprise License:

The Omni-Compass LLC

Owner and developer: AJ Dubra

www.omni-compass.com

Repository: github.com/The-Omni-Compass-LLC/The-Omni-Compass

Any use beyond evaluation and simulation requires a signed, paid Omni-Compass Enterprise License. Write to The Omni-Compass LLC before any production use.

OMNI-COMPASS

UNIFIED CONTROL AUTHORITY

EVALUATING ADMISSIBLE EVOLUTION ACROSS OPERATIONAL SYSTEMS



Unified, Transparent Decision Oversight Across Strategic Infrastructure

THE OMNI-COMPASS LLC • www.omni-compass.com

US patent, copyright and trademark applications filed. Evaluation and simulation use only; all other use requires a signed, paid Omni-Compass Enterprise License.