Omni-Compass benchmark: Kubernetes native vs Kubernetes + Omni-Compass

Two identical kind clusters (1 control plane + 6 workers), same add-ons, workload (php-apache + HPA, target 50), load schedule and capture. **Native**: Omni-Compass not running. **Omni**: Omni-Compass drives the HPA target, the node pool (cordon/drain/uncordon), the power cap (in-place CPU limits on the app pods, enforced by the kernel), heat (harness law on live power), security (hold signal) and the rollout guard. Response times are real HTTP requests timed every 5 s. Power is a declared model (idle 100 W + 150 W x utilisation per worker in service; parked workers on standby at idle power), not a meter.

Gauges

Gauge	Native	Omni-Compass	Change
duration (min)	20.08	20.23	
worker nodes in service, mean	6.00	6.00	+0.0%
worker nodes in service, min	6.00	6.00	+0.0%
worker nodes in service, max	6.00	6.00	+0.0%
node-hours	2.01	2.02	+0.7%
power (W), mean	627	642	+2.4%
power (W), peak	651	686	+5.3%
energy (Wh)	210	216	+3.1%
energy, parked workers still on at idle power (Wh)	210	216	+3.1%
CPU used (cores), mean	0.833	1.23	+47.9%
CPU allocatable (cores), mean	24.00	24.00	+0.0%
utilisation (used / allocatable)	0.035	0.051	+47.9%
energy per core-hour (Wh)	753	521	-30.8%
node-hours per core-hour	7.20	4.87	-32.4%
pending pods, pod-minutes	0.267	0.267	+0.0%
pending pods, peak	1.00	1.00	+0.0%
HPA replicas, mean	8.95	9.33	+4.2%
HPA replicas, peak	10.00	10.00	+0.0%
HPA shortfall (desired > current), minutes	0.533	1.05	+96.9%
requests timed	1,080	1,100	+1.9%
time over the response line (% of samples)	0.093	0.000	-100.0%
response time (ms), mean	115	94.93	-17.1%
response time (ms), median	103	82.00	-20.2%
response time (ms), 95th percentile	218	152	-30.1%
response time (ms), 99th percentile	299	192	-35.7%
failed requests (%)	0.000	0.000	0

Significance (2-minute blocks: native 10, omni 10; bootstrap 95% interval)

Metric	Native	Omni-Compass	Change	95% CI of difference	Verdict
node_hours_per_core_hour	7.8700	5.5202	-29.9%	[-4.0384, -0.6568]	GAIN
kwh_per_core_hour	0.8187	0.5852	-28.5%	[-0.3930, -0.0678]	GAIN
utilisation	0.0342	0.0506	+47.7%	[+0.0043, +0.0296]	GAIN
pending_pod_minutes_per_hour	0.7680	0.7619	-0.8%	[-2.3040, +2.2857]	EVEN, no significant difference
hpa_shortfall_minutes_per_hour	1.5360	3.0416	+98.0%	[-1.5603, +5.3212]	EVEN, no significant difference

What Omni-Compass did

- decisions: 20 (one per minute)
- HPA target writes: 0
- node-pool resizes: 3
- scheduling-floor adds (pending pods): 0
- power-cap pod resizes (in place): 0
- rollout actions: 0
- power cap decided per minute: 0.695
- heat (thermal) per minute: 0.358 0.393 0.421 0.446 0.469 0.489 0.505 0.52 0.535 0.547 0.557 0.563 0.568 0.572 0.577 0.582 0.585 0.588 0.589 0.591
- nodes decided per minute: 6 5 4 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3
- HPA target decided per minute (%): 77 77 77 77 77 77 77 77 77 77 77 77 77 77 77 77 77 77 77 77

Reset

```
restored target: 50
workers in service: 6 of 6
Omni records left after kill: none
```

Minute by minute

min	nodes native	nodes omni	power W native	power W omni	replicas native	replicas omni	pending native	pending omni	CPU native	CPU omni
0	6.00	6.00	615	616	5.00	5.00	0.000	0.250	0.507	0.570

min	nodes native	nodes omni	power W native	power W omni	replicas native	replicas omni	pending native	pending omni	CPU native	CPU omni
1	6.00	6.00	624	630	5.75	7.00	0.250	0.000	0.727	0.956
2	6.00	6.00	615	623	6.00	7.25	0.000	0.000	0.521	0.694
3	6.00	6.00	621	633	6.00	8.00	0.000	0.000	0.675	1.03
4	6.00	6.00	632	651	9.00	9.75	0.000	0.000	0.976	1.47
5	6.00	6.00	632	645	9.00	10.00	0.000	0.000	0.939	1.34
6	6.00	6.00	628	647	9.00	10.00	0.000	0.000	0.903	1.40
7	6.00	6.00	643	672	9.75	10.00	0.000	0.000	1.27	2.03
8	6.00	6.00	645	673	10.00	10.00	0.000	0.000	1.32	2.08
9	6.00	6.00	644	675	10.00	10.00	0.000	0.000	1.28	2.12
10	6.00	6.00	635	676	10.00	10.00	0.000	0.000	1.07	2.12
11	6.00	6.00	616	626	10.00	10.00	0.000	0.000	0.537	0.809
12	6.00	6.00	616	622	10.00	10.00	0.000	0.000	0.547	0.726
13	6.00	6.00	619	636	10.00	10.00	0.000	0.000	0.620	1.06
14	6.00	6.00	632	645	10.00	10.00	0.000	0.000	0.939	1.31
15	6.00	6.00	631	648	10.00	10.00	0.000	0.000	0.980	1.40
16	6.00	6.00	633	643	10.00	10.00	0.000	0.000	0.994	1.26
17	6.00	6.00	623	634	10.00	10.00	0.000	0.000	0.734	1.05
18	6.00	6.00	614	624	10.00	10.00	0.000	0.000	0.471	0.761
19	6.00	6.00	616	622	10.00	10.00	0.000	0.000	0.560	0.737

Limits

- kind nodes are containers sharing one CI machine; timings and CPU are noisier than real servers.
- The two arms ran on two different CI machines at the same time; machine-to-machine variation is part of the noise.
- Energy is modelled from utilisation and nodes in service with the declared constants above.